

HARVARD
Kenneth C. Griffin



**GRADUATE SCHOOL
OF ARTS AND SCIENCES**

THESIS ACCEPTANCE CERTIFICATE

The undersigned, appointed by the
Department of Statistics
have examined a dissertation entitled
The Statistics of Discrimination
presented by Johann D. Gaebler
candidate for the degree of Doctor of Philosophy and hereby
certify that it is worthy of acceptance.

Signed by:
Signature Sharad Goel
EF59EA38FAAE43F...

Typed name: Prof. Sharad Goel

Signed by:
Signature Natesh Pillai
BD3188B9E2294D3...

Typed name: Prof. Natesh Pillai

DocuSigned by:
Signature Xiao-Li Meng
F95F8E22E5AD4C7...

Typed name: Prof. Xiao-Li Meng

Date: April 8, 2026

The Statistics of Discrimination

A dissertation presented by

Johann D. Gaebler

to

The Department of Statistics

in partial fulfillment of the requirements
for the degree of Doctor of Philosophy
in the subject of Statistics

Harvard University
Cambridge, Massachusetts

April 8th, 2026

© 2026 Johann D. Gaebler
All rights reserved.

The Statistics of Discrimination

ABSTRACT

Policymakers, regulators, courts, and the public increasingly rely on statistical evidence to understand and address human and machine bias in domains like lending, policing, hiring, and medicine. Empirical studies measuring discrimination, however, face both definitional and data challenges. First, mathematical definitions of bias often imperfectly capture legal and intuitive notions of discrimination, and seemingly natural measures can have counterintuitive implications that would harm the groups they are intended to protect. Second, discrimination is often most concerning in settings where data are limited: Key information about decisions, such as the factors on which they were based, is often unavailable or difficult to measure. This dissertation develops theory, methods, and applications across a variety of domains aimed at addressing these twin challenges in the measurement of discrimination.

Part I studies algorithmic discrimination. Machine decisions are simpler to analyze than human decisions in important respects—decision inputs are structured and usually observable, and the same inputs reliably yield the same decisions—but the meaning of algorithmic “discrimination” is often unclear. Chapter 2 critiques popular mathematical definitions of algorithmic fairness that seek to equalize model performance metrics like false positive rates across race, gender, or other characteristics. Because these measures suffer from inframarginality, a subtle statistical limitation, equalizing them typically requires making suboptimal decisions, including for protected groups. Chapters 3, 4, and 5 develop alternative approaches to measuring discrimination in algorithmic decisions. In contrast to the approaches critiqued in Chapter 2, Chapter 3 advocates for a consequentialist approach to fairness, proposing an algorithm that efficiently reduces inequality while balancing competing policy goals in lending. Chapters 4 and 5 draw on legal frameworks governing discrimination. Chapter 4 develops a risk-based measure of disparate impact, along with a regression-based framework for estimating it in observational settings. Motivated by emerging AI regulations, Chapter 5 adapts correspondence experiments, a behavioral science method for detecting discrimination, to audit large language models used in hiring.

Part II extends the methods of Part I to human decision-making. Unlike algorithms, human decisions are often poorly documented, difficult to reconstruct, and costly to manipulate, hindering the measurement of discrimination. We focus on applications to criminal justice, where these problems are particularly acute. Chapter 6 develops the robust outcome test, a method for measuring disparate impact in settings where little is known about the factors driving decisions, rendering the regression-based approach of Chapter 4 infeasible. By combining “benchmark” and “outcome” tests, standard but individually flawed methods of detecting discrimination, the robust outcome test yields strong statistical guarantees under a nonparametric assumption that empirical evidence suggests often holds in practice. Paralleling the correspondence experiments of Chapter 5, Chapter 7 develops a causal framework for estimating discrimination from observational data in multi-stage decision processes like prosecutorial charging decisions occurring downstream of potentially discriminatory earlier decisions like racially biased arrests. Finally, an alternative response to data limitations is leveraging novel data sources. Chapter 8 uses large-scale telematics data, second-by-second records of driving behavior for hundreds of thousands of individuals, to benchmark racial disparities in traffic enforcement against actual differences in speeding, disentangling discriminatory enforcement from differences in underlying behavior.

Contents

Title Page	i
Copyright	ii
Abstract	iii
Table of Contents	v
Acknowledgments	x
Dedication	xiii
Citations to Previous Work	xiv
1 Introduction	1
1.1 Background	3
1.1.1 Case Study: NYPD Stop-and-Frisk	4
1.1.2 Legal and Economic Theories of Discrimination	9
1.1.3 Empirical Approaches to Measuring Discrimination	15
1.2 Overview and Summary of Contributions	22
1.2.1 Part I: Algorithmic Discrimination	22
1.2.2 Part II: Human Discrimination	25
I Algorithmic Discrimination	28
2 The Measure and Mismeasure of Fairness	29
2.1 Introduction	29
2.2 Mathematical Definitions of Fairness	32
2.2.1 Formal Setting	33
2.2.2 Limiting the Effect of Decisions on Disparities	35
2.2.3 Limiting the Effect of Attributes on Decisions	37
2.3 Equitable Decisions in the Absence of Externalities	42
2.3.1 Utility, Risk, and Threshold Rules	42
2.3.2 The Problem of Inframarginality	44
2.3.3 The Problem with Fairness through Unawareness	48
2.4 Equitable Decisions in the Presence of Externalities	54
2.4.1 The Geometry of Fair Decision Making	55

2.4.2	A Formal Theory of Fairness in the Presence of Externalities	57
2.5	A Path Forward	64
2.5.1	Competing Notions of Ethical Decision Making: Process vs. Outcomes	65
2.5.2	Designing Equitable Algorithms	68
2.5.3	Case Study: An Algorithm to Allocate Limited Medical Resources	75
2.6	Conclusion	80
3	Fair Allocation through Selective Information Acquisition	83
3.1	Introduction	83
3.2	Related Work	85
3.3	A Model of Screening and Allocation	86
3.3.1	A stylized example	89
3.4	Finding Optimal Policies	90
3.5	Experiments	92
3.5.1	Synthetic data	93
3.5.2	Empirical credit data	96
3.6	Discussion	97
4	Measuring Disparate Impact in Human and Machine Decisions	98
4.1	Statistically Assessing Disparate Impact	102
4.2	An Application to Policing	104
4.2.1	Estimating risk-adjusted disparities	105
4.2.2	Sensitivity analysis	107
4.3	Discussion	109
4.4	Materials and Methods	112
5	Auditing Large Language Models for Race and Gender Disparities	113
5.1	Background	114
5.2	Empirical Analysis and Results	117
5.2.1	LLMs for Candidate Evaluation	117
5.2.2	Assessing Adverse Impact Ratios	119
5.2.3	Assessing Correspondence Experiments	120
5.2.4	Sensitivity to Prompt Variation and Context	122
5.3	Discussion	122
II	Human Discrimination	126
6	A Simple, Statistically Robust Test of Discrimination	127
6.1	Statistical Guarantees	130
6.1.1	Formal Setup	130
6.2	Assessing Monotonicity	133
6.2.1	Empirical evaluation	133
6.2.2	Simulation study	134
6.3	An Application to Police Stops	136
6.4	Discussion	139
7	A Causal Framework for Observational Studies of Discrimination	142
7.1	A Motivating Example	143

7.2	A Measure of Discrimination	146
7.2.1	A formal model of discrimination	147
7.2.2	Estimating the CDE	149
7.2.3	An alternative measure of discrimination	153
7.3	Assessing Second-Stage Discrimination in a Stylized Scenario	157
7.3.1	The data-generating process	158
7.3.2	Estimating the CDE	163
7.3.3	Simulation results	165
7.4	An Empirical Analysis of Prosecutorial Charging Decisions	168
7.5	Discussion	173
8	Measuring Racial and Ethnic Disparities in Traffic Enforcement with Large-Scale Telematics Data	175
8.1	Data and Results	177
8.2	Discussion	184
8.3	Materials and Methods	186
A	Appendix to Chapter 2	190
A.1	Path-specific Counterfactuals	190
A.2	Constructing Causally Fair Policies	190
A.3	A Stylized Example of College Admissions	195
A.4	Proof of Theorem 2.10	195
A.5	Proof of Proposition 2.15	197
A.6	Prevalence and the Proof of Theorem 2.17	204
A.6.1	Shyness and Prevalence	204
A.6.2	Outline	207
A.6.3	Convexity, Complete Metrizability, and Universal Measurability	209
A.6.4	Shy Sets and Probes	220
A.6.5	Proof of Theorem 2.17	232
A.6.6	Proof of Theorem 2.9	245
A.6.7	Proof of Corollary 2.18	247
A.6.8	General Measures on $\mathcal{K}$	249
A.7	Theorem 2.11 and Related Results	253
A.7.1	Extension to Continuous Covariates	253
A.7.2	A Markov Chain Perspective	258
A.8	Proofs of Propositions 2.19 and 2.20	261
A.8.1	Beta Distributions and Stochastic Dominance	261
A.8.2	Proof of Proposition 2.19	264
A.8.3	Proof of Proposition 2.20	268
B	Appendix to Chapter 3	271
B.1	Proofs	271
B.2	Optimized Procedure	275
C	Appendix to Chapter 4	278
C.1	Theories of Discrimination	278
C.1.1	Disparate treatment and disparate impact	278

C.1.2	A formal illustration of included-variable bias	280
C.1.3	Proofs	283
C.2	Sensitivity analysis	288
C.2.1	Introduction	288
C.2.2	Theory	296
C.3	Estimation of the Non-Parametric Estimand by Linear Regression	337
C.4	Accounting for factors beyond risk	340
C.5	Additional figures	342
D	Appendix to Chapter 5	350
D.1	A Note on Terminology	350
D.2	Data Redaction and Processing	350
D.3	Synthetic Application Generation	351
D.3.1	Prompting	353
D.4	Additional Results and Notes	353
D.4.1	Manipulation Check	353
D.4.2	Prompt Variations	355
D.4.3	Impact of Blinding	358
D.4.4	Prediction of Race and Gender from Redacted Materials	358
D.4.5	Additional Adverse Impact Ratio Analyses	359
D.4.6	Alternative Analysis of Name Effects	360
D.4.7	Selection Ratios	361
D.5	Models	361
E	Appendix to Chapter 6	374
E.1	Mathematical Appendix	374
E.1.1	Proof of Theorem 6.1	374
E.1.2	General Utilities and Quasi-Rational Decision Makers	377
E.1.3	Risk-Decision Curves	381
E.1.4	Stochastic Orderings	383
E.1.5	Proof of correctness of the robust outcome test	390
E.1.6	A Note on Sampling Error	394
E.2	Data and Risk Estimation	395
E.2.1	Police Stops	396
E.2.2	Lending	396
E.2.3	Recidivism	397
E.2.4	Contraband	398
E.2.5	Admissions	399
E.3	Simulation Study	399
E.3.1	Violations of the MLRP	399
E.3.2	Differences in Base Rates	400
F	Appendix to Chapter 7	408
F.1	A Comparison to Alternative Ignorability Conditions	408
F.2	Analysis of a Restricted Family of Distributions	420
F.3	Extending Theorem 7.4 to Allow for Continuous Covariates	424
F.4	Summary Statistics of Prosecution Dataset	427

G Appendix to Chapter 8	430
References	451
Chapter 1: Introduction	451
Chapter 2: The Measure and Mismeasure of Fairness	457
Chapter 3: Fair Allocation through Selective Information Acquisition	467
Chapter 4: Measuring Disparate Impact in Human and Machine Decisions	470
Chapter 5: Auditing Large Language Models for Race & Gender Disparities: Implications for Artificial Intelligence–Based Hiring	477
Chapter 6: A Simple, Statistically Robust Test of Discrimination	482
Chapter 7: A Causal Framework for Observational Studies of Discrimination	488
Chapter 8: Measuring Racial and Ethnic Disparities in Traffic Enforcement with Large- Scale Telematics Data	492

Acknowledgments

To complete my Ph.D., I depended on many, many people for inspiration, guidance, friendship, support, and distraction. It would be impossible to adequately thank all the people who have helped me (and at times, carried me) along the way.

I would not have had the confidence to embark on a Ph.D. without the encouragement and mentorship of many teachers who fostered my curiosity and encouraged me to pursue research long before I wrote my first paper. I want to particularly thank William Wiswall, who was an early inspiration; Bruce Wade and Peter Hinow, who patiently introduced me to proof-based mathematics at UWM; Cesar Silva at Williams and Warren Goldfarb and W. Hugh Woodin at Harvard for introducing me to logic and mathematical research; and Joe Blitzstein for encouraging me to pursue (what felt like) a risky pivot to a new field.

I would never have had the ability to pursue *this* Ph.D. without the many wonderful researchers and engineers who patiently mentored me at the Stanford Computational Policy Lab. I would like to especially thank those of you I was lucky enough to work with directly: Phoebe Barghouty and Cheryl Philips for their tireless efforts collaborating on my very first research project; Joe Nudell, Dan Jenson, and Amy Shoemaker for many—hopefully fun but probably frustrating—hours spent helping me learn the ropes of coding and data science; and Marissa Gerchick and Keniel Yao for helping with some of my early (and often disorganized) projects.

I looked up to all of the students I was lucky to know at SCPL. Thank you to Will Cai, Alex Chohlas-Wood, Sam Corbett-Davies, Josh Grossman, Jongbin Jung, Allison Koenecke, Jerry Lin, Jan Overgoor, Hao Sheng, and Camelia Simoiu. I was fortunate to see many of you make the leap from staff to student (or, like Alex, straddle the divide!) before me. I am grateful to have had the opportunity to add footnotes to the many fascinating questions you asked and answered. And, of course, thank you to the many MS&E students who were subjected to our often over-lively conversations, especially Nikhil Garg, Hannah Li, and Lodewijk Gelauff.

I started my Ph.D. at Stanford, and wouldn't have made it far without my cohort at ICME: Thank you to Trevor Maxfield, Ayya Alieva, Xun Tang, and Leah Reeder for braving all six qualifying exams together with me. Thank you also to Stanford ICME and the Knight-Hennessy Scholars Program for supporting me during that period, and to the many other students and researchers I was fortunate to overlap with at Stanford.

The statisticians at Harvard are easily the smartest and the most fun I've encountered, and I've been incredibly fortunate to spend the second half of my Ph.D. surrounded by them. Thank you to my cohort here—Nowell Closser, Kyla Chasalow, Alan Chung, Nathan Cheng, Matthew Esmaili, Yuzhou Lin, Souhardya Sengupta, Philip O'Sullivan, Ben Schiffer, Xiaodong Yang—as well as to the many more senior students who welcomed us into the fold, especially James Bailie, Ritwik Bhaduri, Jonathan Che, David Ham, Biyanka Liang, Cory McCartan, and Adela Ye. I want to also thank everyone involved in the Harvard Statistics Consulting Service, especially Lucas Janson for letting me participate in and eventually run the consulting service, and the many statistical consultants

who tolerated me as their partner over the years.

At the Computational Policy Lab and the Shorenstein Center at Harvard, I've benefited enormously from working with many talented colleagues, particularly Charlotte Tuminelli and Evangelos Kassos. Thank you also to Jamie Hamilton, Jane Lee, Sarah Vicol, Emily Lemmerman, Mark Kogo, and Muskan Walia for putting up with—or, for the luckier among you, narrowly avoiding—my research mentorship. I learned a great deal from all of you, for which I am very grateful.

Thank you to the Stone Program for supporting my final year here at Harvard and for giving me the opportunity to meet and learn from the other scholars in my cohort: Gaby Aboulafia, Aditi Bhowmick, Aidan Connaughton, Brein Mosely, Clare Suter, Hannah Craig, Jessica Urzúa, Kiran Misra, Major Eason, and Zheng Ma. Thank you also to Danny Schneider, David Weil, Maya Sen, Jennifer Hochschild, and Deirdre Bloome for three fascinating and extremely valuable semesters in the Stone Inequality Proseminar.

I want to thank Jacob Jameson and Madison Coots for honing their skills with me and my wife Yunhan at the most important and prestigious academic pursuit: bar trivia. For the last decade, the Angry Face Association has been prepared at any time of day or night to field Q/C/Cs or explain gauge theory. Thank you Greg Parker, Makinde Ogunnaike, Kevin Yang, Jake McNamara, and Jenny Walsh for being some of my closest friends for so many years—I'm sorry to spread us even further across the country. And, of course, thank you to the many other people who have provided welcome reminders of the world outside research: Michael Savarese, Tim King, Kyle Laracey, Anand Rajesh, Sam King, Jonah Kallenbach, and many other friends near and far.

I would likely be out of a job next year if it weren't for the efforts of many people who helped me navigate the academic job market. Thank you to Nikhil Garg, Andrew Gelman, Swati Gupta, David Ham, Aziz Huq, Hannah Li, Susan Murphy, Julian Nyarko, Ravi Shroff, and Sonny Tambe for your advice and support.

The insights and patience of the brilliant researchers I've been fortunate to collaborate with have been one of the key reasons for whatever success I've had in my Ph.D. Thank you to my co-authors Chris Avery, Phoebe Barghouty, Guillaume Basse, Will Cai, Sam Corbett-Davies, Nikhil Garg, Jennifer Hill, Aziz Huq, Calvin Isley, Shanto Iyengar, Jongbin Jung, Justin Kaashoek, Tobias Konitzer, Samuel Madden, Marc Meredith, Hamed Nilforoshan, Cheryl Philips, Lisa Pinals, David Rothschild, Ravi Shroff, Sonny Tambe, Sabina Tomkins, Sarah Vicol, Sean J. Westwood, and Keniel Yao. Thank you also to my near-coauthors Michael Zanger-Tishler and Julian Nyarko for many conversations that have changed my thinking.

Thank you to my committee members Xiao-Li Meng and Natesh Pillai for so many interesting conversations about teaching and statistics, and for creating opportunities for me to grow as a scholar both inside and outside Harvard. And thank you especially to my advisor, Sharad Goel, to whom I owe so much intellectually, personally, and professionally. Despite my woeful lack of any sort of qualification and a pathetic job pitch, Sharad gave me the opportunity to join SCPL as a junior data scientist nearly a decade ago, introducing me to applied statistics and public policy. I can only hope that I embody a fraction of the curiosity, integrity, work ethic, and commitment to scientific truth that he modeled over our many years working together.

My family has been subjected to me pretending to be a Ph.D. ever since I was a kid. Thank you, Mom and Dad, for a warm home, an endless supply of hopelessly outdated reference books, an incredible education (encyclopedia notwithstanding), a hope of making the world a little better, and an unshakeable (and arguably contradictory) belief that solving math problems and writing papers for a handful of people is a worthwhile way to spend six years—or even a lifetime. And thank you to Carrie, Karl, Gideon, Lauren, Annie, and Lou for inspiring and (at times, begrudgingly) encouraging my academic aspirations. After thirty years, you would think you're entitled to some kind of break, but I'm afraid that actually having a Ph.D. is unlikely to help the situation. Thank you also to Ting

and Su for raising the most wonderful person I know and welcoming me into your family. You have always treated me like one of your own—a feeling I’ve experienced also with your brothers, sisters, nieces, nephews, and entire extended families. While the doctors in our family—Uncle David, Aunt Mary, 二舅, John, Katie, Ashley, Nathaniel, Gideon, Harry—blazed the trail I’m on, many more of you have encouraged and inspired me. Thank you in particular to Uncle Tim, whose love of science and learning I idolized from a young age. (And sorry to Carrie and the real doctors—you are very inspiring too.)

Most of all, thank you to my wife, Yunhan. None of this could have happened without you. You have supported my lifelong dream of becoming an absentminded, bespectacled professor, inspired me to keep going on many occasions when I would much rather have given up, provided invaluable input on all things personal and professional, cared for me when I was sick, moved across the country, and made many other sacrifices great and small. Perhaps if I had gotten a Ph.D. in English literature instead, I would be able to express my feelings without embarrassing myself, but given your likely reaction to the prospect of doing *another* Ph.D., let me instead just say: I love you.

To Yunhan.

Citations to Previous Work

Chapter 2 is based on Corbett-Davies, Gaebler, Nilforoshan, Shroff, and Goel (2023):

Sam Corbett-Davies, Johann D. Gaebler, Hamed Nilforoshan, Ravi Shroff, and Sharad Goel. The measure and mismeasure of fairness: A critical review of fair machine learning. *Journal of Machine Learning Research*, 24(312):1–117, 2023.

Replication code and data are available at <https://github.com/jgaeb/measure-mismeasure>.

Chapter 3 is based on Cai, Gaebler, Garg, and Goel (2020):

William Cai, Johann D. Gaebler, Nikhil Garg, and Sharad Goel. Fair allocation through selective information acquisition. In *Proceedings of the AAAI/ACM Conference on AI, Ethics, and Society*, pages 22–28, 2020.

Chapter 4 is based on Jung, Corbett-Davies, Gaebler, Shroff, and Goel (2025):

Jongbin Jung, Sam Corbett-Davies, Johann D. Gaebler, Ravi Shroff, and Sharad Goel. Measuring disparate impact in human and machine decisions. *Proceedings of the National Academy of Sciences*, 2025. Forthcoming.

Replication code and data are available at <https://github.com/jgaeb/rar-repro>. The `rar` package is available on CRAN.

Chapter 5 is based on Gaebler, Goel, Huq, and Tambe (2024):

Johann D. Gaebler, Sharad Goel, Aziz Huq, and Prasanna Tambe. Auditing large language models for race & gender disparities: Implications for artificial intelligence-based hiring. *Behavioral Science & Policy*, 10(2):46–55, 2024. doi: 10.1177/23794607251320229.

Replication code and data are available at <https://github.com/jgaeb/llm-audit>.

Chapter 6 is based on Gaebler and Goel (2025):

Johann D. Gaebler and Sharad Goel. A simple, statistically robust test of discrimination. *Proceedings of the National Academy of Sciences*, 122(10):e2416348122, 2025. doi: 10.1073/pnas.2416348122.

Replication code and data are available at <https://github.com/jgaeb/outcomepp>.

Chapter 7 is based on Gaebler, Cai, Basse, Shroff, Goel, and Hill (2022):

Johann D. Gaebler, William Cai, Guillaume Basse, Ravi Shroff, Sharad Goel, and Jennifer Hill. A causal framework for observational studies of discrimination. *Statistics and Public Policy*, 9(1):26–48, 2022. doi: 10.1080/2330443X.2021.2024778.

Replication code and data are available at <https://github.com/stanford-policylab/gcbsgh-rep>.

Chapter 8 is based on Cai, Gaebler, Kaashoek, Pinals, Madden, and Goel (2022):

William Cai, Johann D. Gaebler, Justin Kaashoek, Lisa Pinals, Samuel Madden, and Sharad Goel. Measuring racial and ethnic disparities in traffic enforcement with large-scale telematics data. *PNAS Nexus*, 1(4):pgac144, 2022. doi: 10.1093/pnasnexus/pgac144.

Replication code and data are available at <https://github.com/stanford-policylab/telematics>.

Chapter 1

Introduction

Few issues in America have been the subject of greater or more long-standing public debate than discrimination. In recent years, high-profile empirical research has sparked national controversy about racial discrimination in policing (e.g., Fryer Jr, 2019), swayed Supreme Court decisions (Arcidiacono, 2018; Card, 2018), and reshaped public perceptions of the fairness of the medical system (Obermeyer et al., 2019), the judicial system (Angwin et al., 2022), hiring (Kline et al., 2022; Bertrand and Mullainathan, 2004), and the algorithmic systems responsible for an ever-broadening array of high-stakes decisions (Buolamwini and Gebru, 2018). Policymakers, judges, journalists, and civic leaders increasingly rely on statistical studies of discrimination to evaluate potentially concerning biases and set policy. The growing influence of these studies in many areas of public life gives both researchers and the public a stake in ensuring that discrimination research is empirically sound and can speak to pressing social and policy issues.

In practice, two important statistical challenges complicate empirical studies of discrimination. The first challenge is conceptual: Defining appropriate statistical measures of discrimination is subtle for two reasons. First, common legal and intuitive conceptions of fairness are often both ambiguous—allowing for conflicting statistical interpretations—and depend on theoretical constructs that are difficult to capture statistically, such as a decision maker’s motivation for an action. For this reason, purely statistical evidence rarely suffices to establish discrimination in legal cases, as we discuss in detail below. Second, many natural-seeming statistical measures of discrimination have counterintuitive properties, harming the same groups decision makers aim to protect. The second

challenge is that key data are often unavailable or difficult to incorporate into statistical analyses. Applicants’ performance in job interviews, for instance, may go unrecorded, hampering studies of hiring discrimination. Moreover, only hired candidates typically generate data on job performance, often a difficult outcome even to quantify. And in many cases, researchers cannot directly measure key quantities, like a suspect’s probability of carrying a weapon, even in principle.

Compounding these challenges, humans are no longer the only decision makers capable of discrimination. Many of the leading legal and economic theories of discrimination were developed in an era when people made essentially all important decisions, from hiring and pretrial detention to medical diagnosis and college admissions. Today, algorithmic systems increasingly aid—or even replace—humans in many of these domains, raising concerns that algorithmic bias could replicate the harms human prejudices once caused. This shift has spawned a broad range of research, often under the umbrella of “algorithmic fairness” (Barocas et al., 2023), that aims to develop new methodological tools for detecting and correcting bias applicable to algorithms’ ever-expanding domain.

This thesis aims to address these challenges through theory, methods, and applications.

- *Theory:* We clarify the implications of different statistical definitions of discrimination and their correspondence to legal, economic, and intuitive notions of discrimination, evaluating their statistical properties and social consequences.
- *Methods:* We develop methods for measuring discrimination in challenging data settings, including settings where the most important inputs to decisions are not observed.
- *Applications:* We use these methods to study discrimination in substantively important applied contexts, especially policing and criminal justice.

In Part I, we study discrimination in algorithmic decisions. Algorithmic decisions feature definite, easily recorded inputs and outputs—meaning that researchers can typically reconstruct decisions exactly after the fact—but foreground theoretical issues surrounding mathematical definitions of bias. In Part II, we apply these insights to human discrimination, a more methodologically challenging setting where decision processes are less mechanically clear and the data documenting them often sparse. Despite their superficial dissimilarity, these complementary efforts help to reveal conceptual unities underlying the statistical study of human and algorithmic bias.

In the course of these investigations, we consider human and algorithmic discrimination in a range of high-stakes domains: medicine (Chapter 2), lending (Chapters 3 and 6), hiring (Chapter 5), pretrial detention (Chapter 6), admissions (Chapters 2 and 6), prosecutorial charging decisions (Chapter 7), and policing (Chapters 4, 6, 7, and 8). These analyses draw on a variety of synthetic, semi-synthetic, and real data, including novel lending, hiring, and policing datasets. While grounded in specific questions from these domains, the theories and methods we introduce apply broadly to many forms of discrimination. We hope that these insights prove useful to researchers studying discrimination and to the policymakers, judges, journalists, and civic leaders who depend on their work to promote equality.

The remainder of Chapter 1 is divided into two parts. In Section 1.1, we review common statistical methods for measuring discrimination and the legal and economic theories of discrimination unifying the approaches developed in later chapters. We ground this discussion in the New York City Police Department’s “Stop-and-Frisk” program, one of the most well-studied and practically consequential cases in the statistical literature on discrimination. With this background in place, in Section 1.2, we give an overview of the specific contributions of this work, summarizing the contents of Chapters 2 through 8.

1.1 Background

Since at least the 1950s, economists and legal scholars have developed theories of discrimination that bear a sometimes murky relationship to the statistical techniques commonly used to measure discrimination. Here we review these theories, highlighting the key statistical challenges that arise when trying to measure the notions of discrimination they lay out. We frame our discussion around the New York City Police Department’s “Stop-Question-Frisk” program, which we introduce in Section 1.1.1. In Section 1.1.2, we introduce the main theories of discrimination in law (disparate impact and disparate treatment) and economics (statistical discrimination and taste-based discrimination). Finally, in Section 1.1.3, we give an overview of the principal statistical approaches to measuring discrimination and discuss their connections with these legal and economic conceptions.

1.1.1 Case Study: NYPD Stop-and-Frisk

Discrimination in policing—and criminal justice more broadly—is one of the most extensively studied topics in the empirical discrimination literature (e.g., Eberhardt, 2019; Knox et al., 2020; O’Flaherty and Sethi, 2019; Goel et al., 2016; Knowles et al., 2001; Rehavi and Starr, 2014; Anwar and Fang, 2006; Gelman et al., 2007; Grogger and Ridgeway, 2006; Pierson et al., 2020). This interest reflects both the technical difficulty of studying policing, where officers often must act quickly and records of decision making are frequently incomplete, and policing’s central role in historical and contemporary American racial inequality. The New York City Police Department’s (NYPD) “Stop-Question-Frisk” (SQF) program, also known as “stop-and-frisk,” sparked a multiyear debate over police discrimination, with statistical analyses of officer bias playing a prominent role in the litigation and political pressure that ultimately curtailed it. (We analyze discrimination under SQF in detail in Chapter 4.) We begin by reviewing the program’s background and structure, before sketching two intuitive statistical strategies for investigating possible discrimination.

The Fourth Amendment of the U.S. Constitution proscribes “unreasonable searches and seizures,” requiring law enforcement officers to both limit the invasiveness of searches and conduct them only with sufficient justification. Under *Terry v. Ohio* (1968), police officers are permitted briefly to detain individuals who they have “reasonable suspicion” are engaged in criminal activity. If officers also have reasonable suspicion that the suspect is armed and dangerous, they are further permitted to pat down, or “frisk,” the suspect’s outer clothing for weapons to protect officer safety. More lenient than the “probable cause” standard applicable to searches generally, reasonable suspicion is met when officers point to “specific and articulable facts” supporting their actions (*Terry v. Ohio*, 1968).

As shown in the left panel of Figure 1.1, in the early 2000s, NYPD began regularly conducting widespread *Terry* stops under SQF. The scope of the program expanded rapidly, with officers conducting nearly 700,000 *Terry* stops and around 400,000 frisks in 2011, at SQF’s peak. After carrying out a *Terry* stop, officers were required to fill out a standardized UF-250 form recording the circumstances of the encounter and the basis for the encounter.¹ UF-250 forms recorded the time

¹NYPD substantially altered data reporting formats in 2017, when search rates had declined significantly from their 2011 peak. To maintain comparability across the most important period in the program’s evolution, we restrict our analysis to 2003–2016.




Figure 1.1: *The scope of the New York City Police Department’s “Stop-Question-Frisk” (SQF) program and disparities in police-civilian contact. The x-axis in both panels indicates year. Left: The number of Terry stops conducted under SQF between 2003 and 2016. Between 2012 and 2013, the year *Floyd v. City of New York* was decided, the number of stops fell by almost two thirds. Right: The breakdown of stops by suspect race. Across the 14-year period, Black pedestrians consistently made up around 50% of stops, followed by Hispanic pedestrians (approx. 30%), and White pedestrians (approx. 10%).*

and location of stops, their duration, demographic information about the suspect, circumstances justifying the stop (e.g., fitting a criminal description or engaging in furtive movements), and any associated outcomes, including frisks, searches, arrests, use of force, and the recovery of weapons or other contraband.

While the scale of NYPD’s SQF program was initially not well known publicly, the program became the subject of widespread outcry beginning in 2007 following the release of data documenting the program’s rapid expansion—a roughly five-fold increase in stops between 2002 and 2006. In addition to raising concerns that officers were stopping and frisking pedestrians without reasonable suspicion, the data documented substantial disparities in the races of the pedestrians officers were stopping. The population of New York City at the time was 25% Black, 27% Hispanic, and 35% White (U.S. Census Bureau, 2007),² while 52% of stops were of Black pedestrians, 30% of Hispanic pedestrians, and 11% of White pedestrians; see the right panel of Figure 1.1. These disparities fueled allegations that officers were conducting *Terry* stops in a discriminatory way, unfairly targeting non-White New Yorkers, and ultimately prompting litigation against NYPD. Relying in part on statistical analyses of SQF pedestrian stops, the U.S. District Court for the Southern District of New York ruled in *Floyd v. City of New York* (2013) that NYPD had violated suspects’ Fourth

²We refer specifically to non-Hispanic White individuals by “White” throughout.

and Fourteenth Amendment rights, “disproportionate[ly] and discriminator[ily] stopping . . . blacks and Hispanics in violation of the Equal Protection Clause.” Contemporaneously with the ruling in *Floyd*,³ NYPD reduced the scale of its *Terry* stop practices, cutting the number of stops by roughly 50%–75% for four consecutive years between 2012 and 2016.

1.1.1.1 Comparing Decision Rates

The simplest statistical strategy for detecting discrimination—and the principal strategy employed by the expert witnesses in *Floyd*—is to compare *decision rates*, i.e., how often members of each group receive a given decision. In the context of SQF, one relevant question is whether stop rates differ across groups. This approach is often referred to as a *benchmark test*, since it benchmarks decision rates for one group against another (Neil and Winship, 2019).

Benchmark tests are often complicated by what is commonly known as the *denominator problem* (Walker, 2001). In the context of NYPD’s SQF program, for instance, while UF-250 forms document the number of stops that took place, the number of cases where officers did *not* make a stop—and hence the relevant denominator when computing stop rates—is more difficult to estimate. One avenue, adopted by expert witnesses for both the plaintiffs and defendants in *Floyd*, is to normalize the number of stops that took place in a particular area using some estimate of the relevant population. However, even apart from estimation difficulties, isolating the appropriate population is itself fraught: In *Floyd*, experts for the defense normalized by race-specific crime counts in a given area, while experts for the plaintiffs normalized by both local population demographics and crime rates, leading to substantively different conclusions about the racial disproportionality of SQF stops. (We reconsider the denominator problem in Chapter 8, developing a novel technological approach for directly measuring the underlying violation rate.)

A common alternative strategy is to focus instead on decisions where both the numerator and denominator in the decision rate are well understood. Because officers can only frisk pedestrians after they have been stopped, analyses of discrimination by police sometimes study frisk (or search) decisions rather than stop decisions (e.g., Pierson et al., 2020). The left panel of Figure 1.2 shows

³Judge Shira A. Scheindlin, who presided over the district court proceedings, was removed from the case in October 2013 following concerns about the appearance of impartiality, and her ruling was initially stayed pending the Bloomberg administration’s appeal, which the de Blasio administration dropped in early 2014, leaving Scheindlin’s ruling in effect.




Figure 1.2: *Frisk and weapon recovery rates in the New York City Police Department’s “Stop-Question-Frisk” program. The x-axis in both panels indicates year, and color indicates race. Left: Between 2003 and 2016, frisk rates rose gradually for all groups by approximately 15 percentage points. Frisk rates are highest for Black suspects, followed by Hispanic suspects, and then White suspects. Right: The proportion of frisks resulting in the recovery of a weapon varies roughly inversely with the number of frisks performed, with weapon recovery rates reaching less than 3% for all groups in 2006–2007, before climbing to between 7% and 10% by 2016, when the total number of stops had fallen by nearly two orders of magnitude relative to the peak in 2011. Across the period, weapon recovery rates are highest for White suspects, followed by Hispanic, and then Black suspects.*

frisk rates by race in the SQF data. Across the 14 years of data represented, the share of White suspects who are frisked varies between roughly 25% and 50%, but is consistently around 15 percentage points lower than the comparable proportion for Black and Hispanic suspects. As in the comparison of stop rates, however, important difficulties attend the interpretation of this difference. Most importantly, the comparison of unadjusted frisk rates does not account for other legitimate explanations for the discrepancy. For instance, if police encounter Black and Hispanic suspects under circumstances where officers are actually more likely to have reasonable suspicion, such as in the vicinity of crime scenes, frisk rates could differ even absent discrimination. To pass legal muster, analyses of discrimination must generally account for alternative benign explanations for observed differences in decision rates, a problem we return to below.

1.1.1.2 Comparing Error Rates

An alternative to benchmark tests also considered in *Floyd* is to compare how often decisions are “correct” (or, equivalently, how often they are in error) across different groups. For instance, in the context of SQF, under *Terry*, a frisk of a pedestrian is successful if the officer recovers a weapon.

The right panel of Figure 1.2 illustrates the success rate of frisks across race groups. While frisks yielding weapons are rare across all race groups, frisks of White individuals succeed under this definition more often than frisks of racial minority suspects nearly every year across the 14 years of data, in many years by margins of 50%–100%. One interpretation of these results, which we explore in detail in Section 1.1.3.3 below, is that officers are more careful when frisking White suspects, whom officers frisk only when relatively confident they possess a weapon.

Comparing success rates in this manner is often referred to as an *outcome test* (Neil and Winship, 2019), since it compares the rate of successful outcomes across groups. Though motivated differently, comparisons of this and other error rates have become one of the most popular approaches to studying algorithmic discrimination in computer science. (We review this literature below in Section 1.1.2.3 and in greater detail in Chapter 2.) To make the connection clear, one can imagine a predictive algorithm that assesses the context of a *Terry* stop—such as the time, location, occurrence of crimes in the vicinity, and other factors similar to those recorded in UF-250 forms—and makes recommendations to an officer about whether a suspect is likely to be armed, similar to some existing recommendations for improving police search practices Goel et al. (2017). On this view, differences in error rates indicate that algorithms like the frisk aid described above function less well or otherwise harm members of certain groups. For instance, in a well-known investigation of the COMPAS recidivism prediction algorithm in Broward County, Florida, Angwin et al. (2022) found that false positive rates (i.e., the proportion of defendants who did not recidivate within the next two years labeled “high risk” by the COMPAS algorithm) were roughly twice as high for Black defendants as White defendants (45% versus 23%, respectively) while roughly the opposite pattern held for false negative rates. Similar investigations of facial recognition software (Raji and Buolamwini, 2019; Buolamwini and Gebru, 2018), mortgage lending algorithms (Martinez and Kirchner, 2021), speech recognition software (Koencke et al., 2020), medical diagnostic tools (Movva et al., 2023), and other algorithmic systems have raised concerns of widespread algorithmic bias.

Like comparisons of decision rates, comparisons of error and success rates raise subtle statistical challenges of their own. As a simple illustration, imagine implementing a universal frisk policy, which—while certain to violate the reasonable suspicion standard—would eliminate concerns that officers unfairly target certain groups when conducting frisks. (Bias in stop decisions would remain a concern, but does not bear directly on this particular analysis; see Chapter 7 for detailed discussion.)

Under this policy, the success rate—or equivalently, the error rate—of frisks would still almost certainly differ between groups; a simple estimation strategy based on the risk-adjusted regression framework developed in Chapter 4 suggests that frisks of White suspects would yield weapons 30%–50% more often than frisks of Black or Hispanic suspects.⁴ In other words, naïvely comparing the success rate of frisks across groups would suggest discrimination even under this tautologically nondiscriminatory frisk policy. We discuss the general statistical phenomenon responsible for this counterintuitive result in Section 1.1.3.3 below.

1.1.2 Legal and Economic Theories of Discrimination

Statistical analysis made it possible in *Floyd* to characterize the millions of *Terry* stops conducted by NYPD. In reaching its ruling, the Court nonetheless relied on a variety of other evidence, including testimony from NYPD officers and administrators, as well as plaintiffs who had experienced *Terry* stops, official statements of NYPD policy, and other forms of evidence, exposing both technical and conceptual gaps between the experts’ statistical analyses and the demands of antidiscrimination law. To understand these gaps, which principled statistical methods for measuring discrimination aim to close, we review the main legal and economic theories of discrimination. Lastly, in Section 1.1.2.3, we briefly compare them to other conceptions of bias from the computer science literature.

1.1.2.1 Discrimination in the Law

In the American legal tradition, there are two primary doctrines of discrimination: *disparate treatment* and *disparate impact*. Disparate treatment prohibits differential treatment on the basis of a protected attribute, like race, gender, or religion, whereas disparate impact prohibits even facially neutral policies that have unjustified adverse impacts on protected groups. The two doctrines differ not only in the kinds of conduct they prohibit, but also in the range of domains they regulate and

⁴Using the full set of covariates available in UF-250 forms, we estimate the probability that frisked individuals are in possession of a weapon. Using data from frisked individuals in each year, we fit gradient boosted tree models using **XGBoost** (Chen and Guestrin, 2016) trained and tuned on data from prior years, downweighting older observations. We restrict our analysis to 2008–2016, the range during which there is sufficient prior data to carry out our model fitting procedure. We then use Platt scaling to calibrate the predictions on the held-out years of data. After imputing risks for all stopped suspects, we estimate that frisks of White suspects yield weapons in 2.3% of cases, versus 1.6% of frisks of Black suspects and 1.8% of frisks of Hispanic suspects. See Chapter 4 for a more detailed treatment of this estimation problem.

in the legal tests used to establish whether violations have occurred.

Disparate treatment. Disparate treatment, the legal doctrine at issue in *Floyd*, is the most common and generally most familiar doctrine of discrimination law. While often grouped under the broader banner of “disparate treatment,” U.S. law actually contains distinct constitutional and statutory prohibitions on differential treatment on the basis of protected characteristics. The equal protection guarantees of the Fifth and Fourteenth Amendments of the U.S. Constitution prohibit the government from acting with a “discriminatory purpose,” i.e., an intention to harm or demean a protected group (*Washington v. Davis*, 1976). In *Personnel Administrator of Massachusetts v. Feeney* (1979), the Supreme Court held that

“Discriminatory purpose” ... implies more than intent as volition or intent as awareness of consequences. ... It implies that the decisionmaker ... selected or reaffirmed a particular course of action at least in part “because of,” not merely “in spite of,” its adverse effects upon an identifiable group.

As a consequence, Equal Protection jurisprudence has permitted the government to adopt policies based on protected characteristics like race or gender in limited circumstances when necessary to further certain government interests, such as male-only registration for the draft, or, until 2023, affirmative action in education (*SFFA v. Harvard*, 2023).⁵ In practice, because discriminatory intent is often difficult to establish directly, courts often consider a variety of evidence of intent, including the impact of the policy itself (*Vill. of Arlington Heights v. Metro. Hous. Dev. Corp.*, 1977). In *Floyd*, for instance, in addition to violations of Fourth Amendment protections against unreasonable search and seizure, the Court found that NYPD had violated equal protection based on both anecdotal and statistical evidence that NYPD had treated White and non-White individuals differently in its SQF program.

Beyond constitutional protections, federal statutory provisions, such as Title VII of the Civil Rights Act, the Equal Credit Opportunity Act, and the Fair Housing Act, also enjoin private actors like employers, lenders, and landlords from discriminating in hiring, lending, and housing, for

⁵Equal protection analysis applies different levels of scrutiny depending on the characteristic in question: so-called “strict scrutiny” in the case of race or religion, “intermediate scrutiny” in the case of gender, and “rational-basis review” when protected characteristics do not enter the analysis.

example, “on the basis of race, color, religion, national origin, sex or marital status, or age” (Equal Credit Opportunity Act, 1974).^{6,7} The usual legal standard applied in disparate treatment cases is that of “but for” causation: Policies treating two people of different races in otherwise identical circumstances differently constitute disparate treatment. In recent years, the Supreme Court has articulated this standard in explicitly counterfactual language:

If the employer intentionally relies in part on an individual employee’s sex when deciding to discharge the employee—put differently, if changing the employee’s sex would have yielded a different choice by the employer—a statutory violation has occurred (*Bostock v. Clayton County*, 2020).

Importantly, even if a protected attribute is not a “but-for” cause of the employment outcome in the sense described above, disparate treatment still results in some circumstances when the attribute is a “motivating factor” that influences the decision process (Civil Rights Act of 1991, 1991; *Price Waterhouse v. Hopkins*, 1989).

Disparate impact. Policies not intended to harm a particular group can still result in illegal discrimination under the disparate impact doctrine, which aims to remove “built-in headwinds” for minority groups (*Griggs v. Duke Power Co.*, 1971). Disparate impact occurs when a policy or practice results in unjustified adverse consequences for a protected group, even if it is formally blind to race, gender, or other protected characteristics. Unlike disparate treatment, which has constitutional underpinnings, disparate impact law derives force from statutes like Title VII of the Civil Rights Act at the federal level and, at the state level, specific civil-rights laws in states like California, Illinois, and New York. These standards regulate a narrower range of areas including lending, housing, and employment.

For instance, in *Griggs v. Duke Power Co.* (1971), the Supreme Court held that the Duke Power Company’s requirement that employees possess a high school diploma and pass a pair of employment

⁶In strict usage, “disparate treatment” refers to violations of these and other statutory provisions, as distinguished from “equal protection violations.”

⁷The range of characteristics protected—and the manner in which they are protected—differs between the constitutional provisions discussed above and federal statutory provisions like Title VII and between different statutory provisions. For instance, ECOA covers age and marital status, while Title VII covers neither.

tests to obtain certain high-paying positions had resulted in illegal disparate impact. Despite the policy’s formal neutrality toward employees’ races, White employees were almost ten times more likely to qualify for transfer and promotion to the high-paying positions. Drawing on evidence from the lower courts that the diploma and testing requirements bore no material relationship to job success, the justices reasoned that

The [Civil Rights] Act proscribes not only overt discrimination but also practices that are fair in form, but discriminatory in operation. The touchstone is business necessity. If an employment practice which operates to exclude Negroes cannot be shown to be related to job performance, the practice is prohibited.

Importantly, the mere existence of disparities does not constitute disparate impact. Although the Court had not found that Duke Power had a discriminatory purpose in adopting its promotion policy,⁸ Duke Power’s employment practices were still illegal because (1) they adversely impacted Black employees, and (2) the adverse impact did not arise from some business necessity. In the context of SQF, a similar analysis might consider the extent to which recovering weapons for officer safety—the sole purpose for frisks authorized by *Terry*—justifies differences in frisks across race groups. Since policing generally does not fall within the narrow range of activities covered by disparate impact law, evidence of unjustified disparities would not have sufficed in *Floyd*; nevertheless, the question of SQF’s disparate impact is both scientifically and substantively important, and we consider it in detail in Chapter 4.

1.1.2.2 Discrimination in Economics

Like American jurists, economists have focused on two related but distinct conceptions of discrimination: *statistical discrimination*, when decision makers treat groups differently in ways that track statistical differences between groups, and *taste-based discrimination*, when a decision maker forgoes profit (or some other form of utility) to avoid transacting with members of a particular group. Although disanalogous to disparate treatment and disparate impact in some regards, statistical discrimination in many contexts raises disparate treatment concerns, and taste-based discrimination,

⁸Prior to *Griggs*, the Duke Power Company had an overtly discriminatory policy requiring Black employees to work in the lowest-paid department. Duke Power adopted the policy at issue in *Griggs* the same day Title VII of the Civil Rights Act, under which the earlier policy was plainly illegal, came into effect.

despite the intent-related implications of its name, helps clarify the statistical logic of testing for disparate impact.

Statistical discrimination. Statistical discrimination (Arrow, 1973; Phelps, 1972) refers to when decision makers rely on protected characteristics like race or gender in settings with imperfect information. In the absence of foreknowledge of which policyholders will have automobile accidents or require expensive medical care, insurance companies, for instance, estimate the probability that these events will occur based on historical accident and claim rates for similar policyholders, pricing premiums accordingly. Insurers engage in statistical discrimination by charging higher health insurance premiums to, for instance, older policyholders and male drivers because, even conditional on the other covariates available for modeling, those populations have higher expected costs. Statistical discrimination is distinguished by the fact that this differential treatment has the purpose of maximizing profit (or some other measure of utility): An auto insurer, for instance, that does not statistically discriminate on the basis of gender necessarily sacrifices some profit, since it must either (1) undercharge male drivers relative to the amount it is likely to pay them in claims, or (2) charge female drivers higher premiums than their actual risk demands, potentially losing customers.

Despite its tight connection to advancing business interests, statistical discrimination has a complex legal status. For instance, statistical discrimination on the basis of race is illegal in hiring, where race can never be a “bona fide occupational qualification” (Title VII of the Civil Rights Act of 1964, 1964). Discrimination on the basis of race and gender is illegal in medical insurance pricing (Patient Protection and Affordable Care Act, 2010) and lending (Equal Credit Opportunity Act, 1974). At the same time, the efficiency gains from statistical discrimination are seen as justifying disparate treatment in other contexts. In medicine, for example, statistical discrimination is widely practiced, with a variety of common instruments for predicting risk of lung cancer, diabetes, kidney function, birth complications, and other conditions (Coots et al., 2025). Limited statistical discrimination is even legally permitted in law enforcement in some contexts. For instance, while it is impermissible for ancestry to be the sole factor motivating a traffic stop, border patrol agents can legally initiate traffic stops based in part on the apparent ancestry of a vehicle’s occupants (United

States v. Brignoni-Ponce, 1975).

Taste-based discrimination. In contrast to statistical discrimination, where decision makers account for race, gender, or other protected characteristics in the interest of maximizing profit, *taste-based discrimination* (Becker, 1957) refers to when decision makers sacrifice profits to avoid transacting with members of certain groups—or behave as if they are avoiding doing so. For instance, prior to the passage of the Fair Housing Act in 1968, mortgage lenders in the U.S. often explicitly denied loans to Black borrowers on the basis of race, even when such loans could be profitable (Rothstein, 2017). In Becker (1957)’s original formulation,

If an individual has a “taste for discrimination,” he must act *as if* he were willing to pay something . . . to be associated with some persons instead of others.

While taste-based discrimination like that exhibited by post-war American lenders was deliberate, decision makers can engage in taste-based discrimination out of ignorance rather than animus, and as a result of decisions that are formally blind to race, gender, or any other protected attribute.⁹ For example, one of the stop justifications provided on UF-250 forms, as noted above, is engagement in furtive movements. In *Floyd*, the Court found that furtive movements were rarely informative about whether suspects were engaged in criminal activity—arrests resulted more frequently when the “furtive movements” box was unchecked than when it was—but that officers frequently cited furtive movements as a basis for conducting a *Terry* stop. Moreover, the rates at which suspects were marked as having engaged in furtive movements differed substantially between race groups, with Black suspects indicated as having engaged in furtive movements 20% more often than White suspects. To the extent that officers frisked suspects engaging in furtive movements more often out of a sincere (and race-neutral) but erroneous belief that those individuals were more likely to possess weapons, non-White suspects were over-frisked relative to the actual risk they posed, resulting in taste-based discrimination. In other words, relative to a more accurately calibrated frisk policy, officers would have paid a “price,” in the form of lower overall weapon recovery rates, to frisk additional non-White suspects. As the parallels between this example and the example illustrating

⁹Contemporary economists have occasionally interpreted taste-based discrimination as requiring certain mental states (e.g., Bertrand et al., 2005). Because of its close connection to disparate impact law, here we adhere to Becker’s original definition.

disparate impact above highlight, taste-based discrimination and the disparate impact doctrine are closely connected. (The disparate impact doctrine is not generally understood to require rational or optimized behavior, an implication of Becker’s formulation of taste-based discrimination.) In Chapters 4 and 6, we develop these links in greater quantitative detail.

1.1.2.3 Algorithmic Fairness

While both intellectually and substantively important, the economic and legal theories of discrimination described above are not the only approaches to understanding discrimination that scholars have developed. In recent years, a large literature in computer science has conceptualized *bias*, particularly in algorithms, through the lens of unequal error rates (e.g., Zafar et al., 2017b; Chouldechova, 2017; Dwork et al., 2012; Hardt et al., 2016; Kleinberg et al., 2017; Kusner et al., 2017; Zafar et al., 2017); this perspective has been increasingly adopted in recent work in other areas, such as political science (e.g., Imai and Jiang, 2020), the law (e.g., Hellman, 2020), and economics (e.g., Arnold et al., 2022), with some authors linking error-rate disparities more directly to the legal conceptions of discrimination discussed above. However, the conception of differences in error rates as a form of bias has been challenged on both normative and pragmatic grounds. Normatively, concern about differences in error rates is difficult to justify in light of common accounts of why discrimination matters. In particular, concern with error rate disparities is difficult to square with the intuition that discrimination is bad because it harms the disadvantaged group: As we study in detail in Chapter 2, decision makers in fact must typically harm disadvantaged groups to achieve error-rate parity. More broadly, scholars have pointed out that error-rate parity typically requires decision makers to violate anticlassification norms (Corbett-Davies et al., 2017) and is incompatible with the burden-shifting frameworks used to assess disparate treatment and disparate impact in court (e.g., Grossman et al., 2024a).

1.1.3 Empirical Approaches to Measuring Discrimination

Building on the brief motivation of benchmark and outcome tests in Section 1.1.1, we now categorize and review the existing statistical approaches to measuring discrimination. We evaluate these approaches using the theories of discrimination outlined in Section 1.1.2, highlighting both technical

and substantive shortcomings of popular methods. Section 1.1.3.1 considers regression-based approaches, principally benchmark tests. Section 1.1.3.2 examines correspondence experiments, a design-based variant of the benchmark test. Finally, Section 1.1.3.3 reviews the outcome test literature.

1.1.3.1 Ordinary Regression

Most statistical analyses of discrimination build on the benchmark test approach discussed in Section 1.1.1.1, using ordinary linear or logistic regression (e.g., Rehavi and Starr, 2014) to measure the “influence” of protected characteristics like race or gender on decisions net of other potential explanatory factors.¹⁰ For instance, in the context of NYPD’s SQF program, one might regress whether officers frisked a suspect against the other information recorded in UF-250 forms, including suspects’ demographics. One would then typically interpret the fitted coefficient on suspect race as the extent to which officers were more likely to frisk Black suspects than White suspects net of the other factors considered in the analysis. From this perspective, benchmark tests of the type considered by the experts in *Floyd*, where analysts adjust for population demographics, can be understood as an attempt to render the (adjusted) populations sufficiently similar across groups to achieve balance on other relevant dimensions.

A regression analysis of this type identifying racial differences in frisk decisions would nevertheless not shed light on officers’ intentions. As noted above, in practice, courts typically regard the intentions behind racially disparate treatment with extreme suspicion, and even “unintentional” or “implicit” biases would likely constitute a disparate treatment violation (Huq, 2019). The more pressing statistical problem is whether a particular regression analysis captures a sufficient range of factors to render two comparable members of different race groups similarly situated in the appropriate sense. Disparate treatment analysis seeks to understand the extent to which officers made different frisk decisions “because of” race (see Section 1.1.2.1), and so the existence of a potential explanatory factor not accounted for in the regression analysis is a threat to its validity. To mitigate omitted variable bias, researchers generally adjust for a wide variety of potentially relevant factors, such as the location and time at which stops occurred, historical crime patterns in the

¹⁰For the difficulties associated with discussion of causal effects of protected attributes like race or gender, see Section 1.1.3.2 below.

vicinity of the stop, the beat or precinct in which a stop occurred to allow for precinct-specific policy differences, and other information of the kind found in UF-250 forms (Goel et al., 2016; Gelman et al., 2007). Chapter 7 considers the problem of assessing disparate treatment with regression methods in detail. In general, statistical adjustment can almost never fully address concerns that omitted variables have skewed results, though in some cases technological solutions can allow for the direct measurement of relevant factors; see, e.g., Chapter 8.

In contrast to typical statistical best practices, injudiciously adjusting for a broad range of covariates can skew estimates of discrimination. As an extreme example, in regression analyses of disparate treatment, researchers generally avoid adjusting for close proxies of race or gender, like hair or eye color, since differential treatment on the basis of eye color would likely be legally tantamount to differential treatment on the basis of race itself. In analyses of disparate impact, this concern is more acute to the extent that adjusting for factors unrelated to legitimate business or government purposes can similarly mask discrimination. In the context of *Griggs*, for instance, a regression analysis adjusting for high school diploma attainment would likely correctly identify that the policy had not resulted in disparate *treatment*, but would mask the larger pattern of disparate *impact*. We study this phenomenon, known as *included variable bias* (Ayres, 2010, 2005), in detail in Chapter 4, proposing a regression and sensitivity analysis framework addressing included and omitted variable bias.

1.1.3.2 Correspondence experiments

Regression approaches to disparate treatment like those explored in Chapter 7 can be thought of as observational methods for trying to understand what would have happened, for instance, during a *Terry* stop where the suspect was of a different race but circumstances were otherwise the same. While such a comparison cannot actually be carried out for any particular individual, *correspondence experiments*, also called “audit studies,” represent a design-based approach to investigating disparate treatment that comes closer to realizing this ideal. In a correspondence experiment, researchers experimentally manipulate cues signaling an individual’s race, gender, or some other protected characteristic to alter how decision makers like police officers perceive the individual. Through randomization, researchers can isolate the impact of the manipulation, achieving balance across

other potentially relevant factors.

Correspondence experiments have a long history both in academic literature and public policy. The earliest field experiments, conducted in the 1960s and 1970s, randomly assigned matched actors of different races to attempt to rent housing, apply for jobs, or obtain services. Both academic researchers (Daniel, 1968) and government agencies like the Department of Housing and Urban Development (Wienk, 1979) carried out large-scale experiments. However, isolating race or gender from other aspects of the manipulation of replacing one actor with another—such as behavior or mannerisms (Heckman, 1998)—is challenging. Researchers today generally focus on more limited manipulations in information-restricted settings following Bertrand and Mullainathan (2004), where the authors identified discrimination against Black job candidates in Chicago by replacing “White-sounding” names (e.g., “Emily” or “Greg”) with “Black-sounding” names (e.g., “Lakisha” or “Jamal”) on job applications. Correspondence experiments have grown increasingly relevant to understanding algorithmic discrimination, following the advent of machine learning techniques for leveraging high-dimensional, unstructured inputs like resumes directly in algorithmic decision systems. We extend correspondence experiments to this setting, studying algorithmic hiring discrimination, in Chapter 7.

While they effectively mitigate omitted variable bias in a formal sense, correspondence experiments have important limitations. Like regression analyses of discrimination, correspondence experiments cannot isolate a decision maker’s intent. Moreover, concrete manipulations tend to raise interpretive difficulties. For instance, some scholars have criticized the particular name manipulations used in Bertrand and Mullainathan (2004) for confounding the effects of race and class (Gaddis, 2017). Since it is rarely possible to isolate signals that solely communicate an individual’s race or gender—indeed, since such signals arguably do not exist even in principle (Sen and Wasow, 2016)—concerns about this kind of confounding are endemic to correspondence experiments. Correspondence experiments also raise practical difficulties. For instance, it may be difficult to alter the gender presentation of a resume listing a degree from an all-women’s college without either (1) producing a resume containing strongly conflicting gender signals, or (2) altering a decision maker’s perception of potentially relevant information, like the prestige of an applicant’s undergraduate institution. Related concerns have been raised about settings where decision makers make multiple decisions in distinct stages and separately manipulating their perception of race may be difficult (Knox et al., 2020). More

broadly, a substantial literature in statistics and related disciplines has documented the difficulties associated with causal inference involving manipulations of race or other social categories (Hu and Kohler-Hausmann, 2020; Sen and Wasow, 2016; Greiner and Rubin, 2011; Holland, 1986).

1.1.3.3 Outcome Tests

Rather than directly analyzing decisions, outcome tests infer discrimination from the outcomes of decisions. Becker (1993) originally proposed outcome tests as an alternative to benchmark test approaches then current in the analysis of lending discrimination:

Recent studies on whether banks discriminate in their mortgage lending against blacks and other minorities compare the likelihood of getting a loan for minority and white applicants who are similar in incomes, credit backgrounds, and other available characteristics. . . . Unfortunately, these studies do not use the correct procedure for assessing whether banks discriminate. . . . If banks discriminate against minority applicants, they should earn greater profits on the loans actually made to them than on those to whites.

The intuition behind this observation is simple: Profit-seeking lenders engaged in taste-based discrimination will lend to less creditworthy (and hence less profitable) White loan applicants in lieu of more creditworthy (and hence more profitable) Black loan applicants, behaving as if not lending to Black loan applicants compensates them for their lost profits. As a result, a discriminatory lender will realize more profits on loans to Black borrowers and less on loans to White borrowers than a purely profit-maximizing lender.

Because outcome tests sidestep omitted variable bias, unlike benchmark tests, scholars have widely applied them to study discrimination in lending (Berkovec et al., 1994, 1998), hiring (Pope and Sydnor, 2011), publication (Chilton et al., 2020; Smart and Waldfogel, 1996), elections (Anzia and Berry, 2011; Green et al., 2009), and criminal justice (Knowles et al., 2001).¹¹ As discussed in Section 1.1.1.2, in the context of NYPD’s SQF program, the outcome test suggests that officers discriminate against non-White pedestrians, since frisks of White pedestrians yield contraband more often, implying that White pedestrians are frisked only when officers are more confident that they

¹¹Some researchers, e.g., Berkovec et al. (1994), also conduct regression-based outcome tests that regress outcomes against demographics, adjusting for decision-relevant factors.

possess weapons.

Much like benchmark tests, however, outcome tests suffer from their own critical statistical shortcoming, *inframarginality*. In the utility-optimizing context of Section 1.1.2.2, decision makers act based on expected utility. For instance, a profit-maximizing lender will lend to the applicants promising the highest expected returns, and a police officer maximizing the number of weapons their frisks yield (given some constraint on the number of frisks they perform) will frisk those individuals with the highest *ex ante* probability of having a weapon, based on the available information observed while conducting the stop. Inframarginality refers to the fact that error-rate statistics, like weapon recovery rates, depend both on the standards that decision makers apply and on the distribution of risk (or, more generally, expected utility) within a given population.

Figure 1.3 depicts a simple mathematical example illustrating how inframarginality undermines outcome tests. Here the x -axis represents risk, i.e., the *ex ante* probability that an individual possesses a weapon, and the y -axis shows the density of example risk distributions for two groups. The solid vertical lines represent the thresholds where officers consider members of a given group sufficiently risky to warrant frisking, with the darkly and lightly shaded regions representing suspects who are or are not frisked, respectively. Different thresholds imply the existence of taste-based discrimination, since a policy maximizing weapon recovery across the groups would take the form of a shared threshold for both groups. Since frisking an individual with risk r yields a weapon with probability r , the weapon recovery rate is the conditional expectation of the risk distributions, shown by the dotted vertical lines. In the left panel of Figure 1.3, the weapon recovery rate for the blue group is 68%, versus 64% for the red group, indicating discrimination against the red group under the outcome test. This is correct, given the officers' policy of frisking all members of the red group with a risk of at least 35% (a 75% search rate) and all members of the blue group with a risk of at least 55% (a 70% search rate). In the scenario depicted in the right panel of Figure 1.3, the red and blue groups have precisely identical search rates and weapon recovery rates, meaning that the outcome test yields the same conclusion, i.e., discrimination against the red group. In fact, however, there is discrimination against the blue group, whose members are searched at a lower threshold (35%) than members of the red group (55%). In other words, despite precisely opposite taste-based discrimination in the two scenarios—favoring the blue group on the left and the red group on the right—the outcome test yields identical conclusions.




Figure 1.3: An illustration of the problem of inframarginality in outcome tests. The x-axis indicates risk (i.e., ex ante probability of possessing a weapon given an officer’s information prior to conducting a frisk), and the y-axis indicates the density of the distribution of risk in two groups. Dark vertical lines represent frisk thresholds (i.e., the level of risk above which officers frisk individuals from a given group), and dashed vertical lines represent the implied weapon recovery rate for the given frisk policies. In both panels, officers frisk 75% of members of the red group and 70% of members of the blue group, and recover weapons in 64% of frisks of members of the red group and 68% of frisks of members of the blue group. Despite being observably identical to a statistical analyst who knows only the frisk rates and weapon recovery rates, officers exhibit taste-based discrimination against members of the red group (threshold: 35%) relative to the blue group (threshold: 55%) in the left panel, but against members of the blue group (threshold: 35%) relative to the red group (threshold: 55%) in the right panel.

Since Becker originally proposed the outcome test, the problem of inframarginality has attracted considerable attention. Scholars have proposed rational-choice models where inframarginality does not arise in equilibrium (Knowles et al., 2001), leveraged differences in the groups that decision makers themselves belong to (Alesina and La Ferrara, 2014; Anwar and Fang, 2006), used differences in leniency along with random assignment of decision makers to identify the risks of marginal individuals (Arnold et al., 2018), or developed parametric models of risk distributions and decision maker behavior to circumvent inframarginality. In addition to strong theoretical assumptions and data requirements, existing methods also assume decision makers behave rationally, though, as we investigate in detail in Chapter 4, real-world decision makers often do not. In Chapter 6, we develop a simple alternative test of discrimination which addresses inframarginality and decision maker quasi-rationality in many cases.

1.2 Overview and Summary of Contributions

Measuring discrimination raises three kinds of difficulties: theoretical questions, especially about the meaning of discrimination and how to capture it in statistical estimands; methodological issues, both general (e.g., omitted variable bias) and discrimination-specific (e.g., included variable bias and inframarginality); and applied problems, such as gathering relevant data in benchmark tests and defining appropriate manipulations in correspondence experiments. Existing statistical methods for measuring discrimination provide only partial solutions to these challenges. Motivated by these limitations, in Chapters 2 through 8, we develop new theory and methods to address various aspects of these and other statistical issues arising in the measurement of discrimination, divided into two parts. Part I covers algorithmic discrimination, where theoretical issues are especially acute, and Part II human discrimination, where methodological challenges are more prominent. Below we briefly introduce the contents and contributions of each chapter.

1.2.1 Part I: Algorithmic Discrimination

Chapters 2 through 5 address a range of problems arising in the measurement of discrimination in algorithms. Chapter 2 first considers mathematical definitions of fairness developed in the algorithmic fairness literature, categorizing them into two broad conceptual families. We then demonstrate

a variety of theoretical results, showing under general conditions that in both constrained and unconstrained settings, implementing these fairness definitions as design principles actually harms the same groups they were intended to protect. Chapter 2 concludes with positive suggestions for improving equity in algorithmic decisions. Chapter 3 expands on this theme, developing a framework for improving algorithmic decisions through the acquisition of additional data for certain high-leverage individuals when model performance differs between groups, a classic problem in algorithmic fairness. Chapter 4 considers the problem of measuring disparate impact in algorithmic and human decisions, and introduces the risk-adjusted regression framework. We first develop a measure of disparate impact closely related to the theory of taste-based discrimination, before introducing a novel two-stage regression strategy to estimate it that addresses included variable bias. To mitigate potential omitted variable bias, we also introduce a novel sensitivity analysis tailored to our regression strategy. Finally, in Chapter 5, we study the adaptation of correspondence experiments to algorithmic settings, investigating whether realistic LLM-based hiring pipelines exhibit bias against minority candidates, an approach now widely used to audit language models for potential bias.

Chapter 2. The field of fair machine learning aims to ensure that decisions guided by algorithms are equitable. Over the last decade, several formal, mathematical definitions of fairness have gained prominence. In this chapter we first assemble and categorize these definitions into two broad families: (1) those that constrain the effects of decisions on disparities; and (2) those that constrain the effects of legally protected characteristics, like race and gender, on decisions. We then show, analytically and empirically, that both families of definitions typically result in strongly Pareto dominated decision policies. For example, in the case of college admissions, adhering to popular formal conceptions of fairness would simultaneously result in lower student-body diversity and a less academically prepared class, relative to what one could achieve by explicitly tailoring admissions policies to achieve desired outcomes. In this sense, requiring that these fairness definitions hold can, perversely, harm the very groups they were designed to protect. In contrast to axiomatic notions of fairness, we argue that the equitable design of algorithms requires grappling with their context-specific consequences, akin to the equitable design of policy. We conclude by listing several

open challenges in fair machine learning and offering strategies to ensure algorithms are better aligned with policy goals.

Chapter 3. Public and private institutions must often allocate scarce resources under uncertainty. Banks, for example, extend credit to loan applicants based in part on their estimated likelihood of repaying a loan. But when the quality of information differs across candidates (e.g., if some applicants lack traditional credit histories), common lending strategies can lead to disparities across groups. In this chapter, we consider a setting in which decision makers—before allocating resources—can choose to spend some of their limited budget further screening select individuals. We present a computationally efficient algorithm for deciding whom to screen that maximizes a standard measure of social welfare. Intuitively, decision makers should screen candidates on the margin, for whom the additional information could plausibly alter the allocation. We formalize this idea by showing the problem can be reduced to solving a series of linear programs. Both on synthetic and real-world datasets, this strategy improves utility, illustrating the value of targeted information acquisition in such decisions. Further, when there is social value for distributing resources to groups for whom we have a priori poor information—like those without credit scores—our approach can substantially improve the allocation of limited assets.

Chapter 4. Empirical analyses have grown increasingly important in discrimination litigation with the greater availability of detailed data on individuals and decisions. A popular analytic strategy is to estimate disparities after adjusting for observed covariates, typically with a regression model, in hopes of ferreting out discriminatory intent. This approach, however, is ill-suited to auditing algorithms that are now commonly used to aid decisions, which typically do not include race or other legally protected factors as inputs. Motivated by legal understandings of disparate impact, we introduce a new approach, which aims to measure “unjustified” disparities in both human and machine decisions. Our method, which we call *risk-adjusted regression*, proceeds in three steps. In the first step, we combine all available information in a machine learning model to estimate the value, or inversely, the risk, of taking a certain action, such as approving a loan application or hiring a job candidate. Second, we measure disparities in decisions after adjusting for

these risk estimates alone. Finally, in the third step, we assess the sensitivity of results to potential mismeasurement of risk. We demonstrate this approach on a detailed dataset of 2.2 million police stops of pedestrians in New York City, and show that traditional statistical tests of discrimination can substantially understate the magnitude of (risk-adjusted) racial disparities.

Chapter 5. Rapid advances in artificial intelligence (AI), including large language models (LLMs) with abilities that rival those of human experts on a wide array of tasks, are reshaping how people make important decisions. At the same time, critics worry that LLMs may inadvertently discriminate against some groups. To address these concerns, recent regulations call for auditing the LLMs used in important decisions such as hiring. But neither current regulations nor the scientific literature offers clear guidance on how to conduct these audits. In this chapter, we propose and investigate one approach for auditing algorithms: correspondence experiments, a widely applied tool for detecting bias in human judgments. We applied this method to a range of LLMs instructed to rate job candidates using a novel dataset of job applications for K–12 teaching positions in a large American public school district. By altering the application materials to imply that candidates are members of specific demographic groups, we measured the extent to which race and gender influenced the LLMs’ ratings of the candidates’ suitability. We found moderate race and gender disparities, with the models slightly favoring women and non-White candidates. This pattern persisted across several variations in our experiment. It is unclear what might be driving these disparities, but we hypothesize that they stem from post-training efforts, which are part of the LLM training process and intended to correct biases in these models. We conclude by discussing the limitations of correspondence experiments for auditing algorithms.

1.2.2 Part II: Human Discrimination

Chapters 6 through 8 build on the theory and techniques developed to study algorithmic discrimination in Part I, adapting them to the more challenging setting of human discrimination. A pervasive challenge in the study of human discrimination, particularly in domains like policing, is the limited availability of data documenting potential explanatory factors like the circumstances of a police stop or the details of a job interview or loan application. Building on the framework

developed in Chapter 4, in Chapter 6, we develop a simple statistical twist on Becker’s outcome test that addresses the problem of inframarginality in a wide variety of settings without imposing substantial additional data requirements. We show that despite their individual fallibility, carefully combining benchmark and outcome tests retains the simplicity of the original approaches while yielding surprisingly strong statistical guarantees under a common non-parametric assumption that likely holds approximately in real-world data. In Chapter 7, we consider the problem of measuring disparate treatment in multi-stage settings in the presence of discrimination at some earlier stage, such as measuring bias in prosecutorial charging decisions when arrest decisions may be biased in unknown ways. Despite concerns that post-treatment bias may skew estimates, we show that for appropriately defined estimands, second-stage disparate treatment can be measured with simple regression estimates under plausible real-life conditions. We then apply these observations to the study of disparate treatment in charging decisions in a large U.S. prosecutor’s office. Finally, in Chapter 8, we study discrimination in traffic stops, using telematics data to compare stop rates to the actual incidence of speeding in 10 U.S. cities. We find that while across the ten cities enforcement is only moderately more concentrated in minority neighborhoods overall relative to underlying violation rates, geographic concentration in individual cities has the potential to lead to disparate impacts.

Chapter 6. In observational studies of discrimination, the most common statistical approaches consider either the rate at which decisions are made (benchmark tests) or the success rate of those decisions (outcome tests). Both tests, however, have well-known statistical limitations, sometimes suggesting discrimination even when there is none. Despite the fallibility of the benchmark and outcome tests individually, here we prove a surprisingly strong statistical guarantee: under a common non-parametric assumption, at least one of the two tests must be correct; consequently, when both tests agree, they are guaranteed to yield correct conclusions. We present empirical evidence that the underlying assumption holds approximately in several important domains, including lending, education, and criminal justice—and that our hybrid test is robust to the moderate violations of the assumption that we observe in practice. Applying this approach to 2.8 million police stops across California, we find evidence of widespread racial discrimination.

Chapter 7. In studies of discrimination, researchers often seek to estimate a causal effect of race or gender on outcomes. For example, in the criminal justice context, one might ask whether arrested individuals would have been subsequently charged or convicted had they been a different race. It has long been known that such counterfactual questions face measurement challenges related to omitted-variable bias, and conceptual challenges related to the definition of causal estimands for largely immutable characteristics. Another concern, which has been the subject of recent debates, is post-treatment bias: many studies of discrimination condition on apparently intermediate outcomes, like being arrested, that themselves may be the product of discrimination, potentially corrupting statistical estimates. There is, however, reason to be optimistic. By carefully defining the estimand—and by considering the precise timing of events—we show that a primary causal quantity of interest in discrimination studies can be estimated under an ignorability condition that is plausible in many observational settings. We illustrate these ideas by analyzing both simulated data and the charging decisions of a prosecutor’s office in a large county in the United States.

Chapter 8. Past studies have found that racial and ethnic minorities are more likely than White drivers to be pulled over by the police for alleged traffic infractions, including a combination of speeding and equipment violations. It has been difficult, though, to measure the extent to which these disparities stem from discriminatory enforcement rather than from differences in offense rates. Here, in the context of speeding enforcement, we address this challenge by leveraging a novel source of telematics data, which include second-by-second driving speed for hundreds of thousands of individuals in 10 major cities across the United States. We find that time spent speeding is approximately uncorrelated with neighborhood demographics, yet, in several cities, officers focused speeding enforcement in small, demographically non-representative areas. In some cities, speeding enforcement was concentrated in predominantly non-White neighborhoods, while, in others, enforcement was concentrated in predominantly White neighborhoods. Averaging across the 10 cities we examined, and adjusting for observed speeding behavior, we find that speeding enforcement was moderately more concentrated in non-White neighborhoods. Our results show that current enforcement practices can lead to inequities across race and ethnicity.

Part I

Algorithmic Discrimination

Chapter 2

The Measure and Mismeasure of Fairness

2.1 Introduction

In banking, criminal justice, medicine, and beyond, consequential decisions are often informed by machine learning algorithms (Barocas and Selbst, 2016; Berk, 2012; Chouldechova et al., 2018; Shroff, 2017). As the influence and scope of algorithms increase, academics, policymakers, and journalists have raised concerns that these tools might inadvertently encode and entrench human biases. Such concerns have sparked tremendous interest in developing *fair* machine-learning algorithms, and, accordingly, a plethora of formal fairness criteria have been proposed in the computer science community (Carey and Wu, 2022; Cleary, 1968; Darlington, 1971; Galhotra et al., 2022; Imai et al., 2020; Mhasawade and Chunara, 2021; Woodworth et al., 2017; Wu et al., 2019; Zafar et al., 2017b; Zhang and Bareinboim, 2018; Zhang et al., 2017; Berk et al., 2021; Chiappa, 2019; Chouldechova, 2017; Chouldechova and Roth, 2020; Coston et al., 2020; Dwork et al., 2012; Hardt et al., 2016;

Acknowledgments. We thank Guillaume Basse, Sander Beckers, Hana Chockler, Alex Chohlas-Wood, Madison Coots, Avi Feller, Josh Grossman, Joe Halpern, Jennifer Hill, Aziz Huq, David Kent, Keren Ladin, Julian Nyarko, Emma Pierson, Ravi Sojitra, and Michael Zanger-Tishler for helpful conversations. This paper is based on work by Corbett-Davies et al. (2023) and Nilforoshan et al. (2022). H.N was supported by a Stanford Knight-Hennessy Scholarship and an NSF Graduate Research Fellowship under Grant No. DGE-1656518. J.G was supported by a Stanford Knight-Hennessy Scholarship. R.S. was supported by the NSF Program on Fairness in AI in Collaboration with Amazon under the award “FAI: End-to-End Fairness for Algorithm-in-the-Loop Decision Making in the Public Sector,” no. IIS-2040898. S.G. was supported by a grant from the Harvard Data Science Initiative. Any opinions, findings, conclusions, or recommendations expressed in this material are those of the authors and do not necessarily reflect the views of the NSF or Amazon. Reproduction materials are available at <https://github.com/jgaeb/measure-mismeasure>.

Imai and Jiang, 2020; Kilbertus et al., 2017; Kleinberg et al., 2017; Kusner et al., 2017; Loftus et al., 2018; Nabi and Shpitser, 2018; Wang et al., 2019; Zafar et al., 2017; Corbett-Davies et al., 2017). Here we synthesize and critically examine the statistical properties of popular formal fairness approaches as well as the consequences of enforcing them. Using both theory and empirical evidence, we argue that these approaches, when used as algorithmic design principles, can often cause more harm than good. In contrast to popular axiomatic approaches to algorithmic fairness, we advocate for a consequentialist perspective that directly grapples with the difficult policy trade-offs inherent to many algorithmically guided decisions.

We begin, in Section 2.2, by proposing a two-part taxonomy of formal fairness definitions. Our first category of definitions encompasses those that consider the effects of decisions on disparities. Imagine, for example, designing an algorithm to guide decisions for college admissions. Under the principle that fair algorithms should have comparable performance across demographic groups (Hardt et al., 2016), one might check that among applicants who were ultimately academically “successful” (e.g., who eventually earned a college degree, either at the institution in question or elsewhere), the algorithm would recommend admission for an equal proportion of candidates across race groups. Our second category of definitions encompasses those that seek to limit both the direct and indirect effects of one’s group membership on decisions. Following the principle that decisions should be agnostic to legally protected attributes like race and gender (cf. Dwork et al., 2012), one might mandate that these features not be provided to the algorithm. Further, because one’s race might impact earlier educational opportunities, and hence test scores, one might require that admissions decisions are robust to the effect of race along such causal paths.

These formalizations of fairness have considerable intuitive appeal. It can feel natural to exclude protected characteristics in a drive for equity; and one might understandably interpret disparities in error rates as indicating problems with an algorithm’s design or with the data on which it was trained. However, in Sections 2.3 and 2.4, we show that both classes of algorithmic fairness definitions suffer from deep statistical limitations. For example, for natural families of utility functions—like those that prefer both higher academic preparedness and more student-body diversity—we prove that common fairness criteria almost always, in a measure theoretic sense, lead to strongly Pareto

dominated decision policies.¹ In particular, in our running college admissions example, adhering to several of the popular conceptions of fairness we consider would simultaneously result in lower student-body diversity and a less academically prepared class, relative to what one could attain by explicitly tailoring admissions policies to achieve desired outcomes. In fact, under one prominent definition of fairness, we prove that the induced policies require simply admitting all applicants with equal probability, irrespective of one’s academic qualifications or group membership. These formal fairness criteria are thus often at odds with policy goals, and, perversely, can harm the very same groups one ostensibly sought to protect by developing and adopting axiomatic notions of fairness.

How, then, can we ensure algorithms are fair? There are no easy solutions, but we conclude in Section 2.5 by offering several observations and suggestions for designing more equitable algorithms. Most importantly, we believe it is critical to acknowledge and tackle head-on the substantive trade-offs at the heart of many decision problems. For example, when creating a college admissions policy, one must necessarily make difficult choices that balance competing priorities. Formal fairness axioms are poor tools for engaging with these challenging issues. Our overarching exhortation is thus to recognize algorithms as encoding policy choices, and to accordingly tailor their design.

Contributions. To summarize, we offer three main contributions. First, we survey the fairness literature, describing existing fairness definitions and organizing them into a two-part taxonomy. Our categorization of formal fairness definitions proposed in the computer science literature highlights their connections to influential legal and economic notions of discrimination. Second, we lay out a consequentialist framework for designing equitable algorithms. Our framework is motivated by viewing algorithmic fairness as a policy objective rather than as a technical problem. This approach exposes the statistical and normative limitations of many popular formal fairness definitions. Finally, we apply our consequentialist framework to develop a positive vision for addressing problems of fairness and equity in algorithm design.

Much of the content we present synthesizes and builds on research that we and our collaborators have conducted over the last several years (Cai et al., 2020; Chohlas-Wood et al., 2023a,b; Koenecke et al., 2023; Corbett-Davies et al., 2017). In particular, we draw heavily on two papers by Corbett-Davies et al. (2023) and Nilforoshan et al. (2022). In addition to synthesis, we broaden the formal

¹A policy is strongly Pareto dominated if there is an alternative feasible policy that is preferred under every utility function in the family (cf. Section 2.4.2).

theoretical results presented in this line of work and offer new, concrete illustrations of our theoretical arguments. Some of the results and arguments we present date back five years, and the field of algorithmic fairness has since moved forward in many ways. For example, in the intervening time, there has been increasing recognition of the shortcomings of popular formal fairness definitions (Barocas et al., 2019). Nevertheless, we believe our message is as relevant as ever. For instance, within the research community, new algorithmic fairness definitions are regularly introduced that, while different in some respects, frequently suffer from the same statistical and conceptual limitations as the notions we survey here. In the broader world, policymakers, algorithm designers, journalists, and advocates often still evaluate algorithms—and accordingly influence decisions—by turning to these formal fairness definitions without necessarily appreciating their shortcomings. For example, proposed legislation in Idaho sought to require that pretrial risk assessment algorithms have equal error rates across groups (Idaho H.B. 118, 2019). Although the proposed bill was never passed, it illustrates the ways in which these formal measures have garnered significant attention beyond the academic community.

The call to build equitable algorithms will only grow over time as automated decisions become even more widespread. As such, it is imperative to address limitations in past formulations of fairness, to identify best practices moving forward, and to outline important open research questions. By synthesizing and critically examining recent developments in fair machine learning, we hope to help both researchers and practitioners advance this increasingly influential field.

2.2 Mathematical Definitions of Fairness

We start by assembling and categorizing definitions of algorithmic fairness into a two-part taxonomy: those that seek to limit the effect of decisions on disparities, and those that seek to limit the effect of protected attributes like race or gender on the decisions themselves. We first introduce formal notation and concrete examples of decision problems in which one might seek to apply these fairness definitions, before reviewing prominent examples of both approaches in turn.

2.2.1 Formal Setting

Consider a population of individuals with observed covariates X , drawn i.i.d. from a set $\mathcal{X} \subseteq \mathbb{R}^n$ with distribution $\mathcal{D}_X$. Further suppose that $A \in \mathcal{A}$ describes one or more discrete protected attributes, such as race or gender, which can be derived from X (i.e., $A = \alpha(X)$ for some function α). Each individual is subject to a binary decision $D \in \{0, 1\}$, determined by a (randomized) rule $d(x) \in [0, 1]$, where $d(x) = \Pr(D = 1 \mid X = x)$ is the probability of receiving a positive decision, $D = 1$.^{2,3} Given a budget b with $0 < b \leq 1$, we require the decision rule to satisfy $\mathbb{E}[D] \leq b$. Finally, we suppose that each individual has some associated binary outcome Y . In some cases, we will be concerned with the causal effect of the decision D on Y , in which case we imagine that there exist two potential outcomes, $Y(0)$ and $Y(1)$, corresponding to what happens to the individual depending on whether they receive a negative or positive decision.⁴

To make our discussion concrete, we imagine two running examples corresponding to this formal setting: diabetes screening and college admissions. As we discuss in detail below, these two examples differ in the extent to which there is agreement about the ultimate value of different decision policies, which in turn impacts our mathematical analysis. Diabetes is a common and serious health condition that afflicts many American adults. If caught early, it is often possible to avoid some of the most significant consequences of the disease through, for example, changes to one’s diet and physical routine. A blood test can be used to determine whether an individual has diabetes, but as with many screening tools, there are risks and inconveniences associated with screening (e.g., a patient may need to take time off from work). In particular, if an individual were *certain* that they did not have diabetes, then they would prefer not to undergo screening. Our goal is to design an equitable screening policy $d(x)$ to determine which patients have ($Y = 1$) or do *not* have ($Y = 0$) diabetes, based on a set of covariates X . For example, following Aggarwal et al. (2022), the screening decision may be based on a patient’s age, body mass index (BMI) and race. (Those authors argue that consideration of race, while controversial, leads to more precise and equitable estimates of diabetes risk, a point we return to in Section 2.3.3.) We further imagine the budget b equals 1, corresponding

²That is, $D = \mathbb{1}_{U_D \leq d(X)}$, where U_D is an independent uniform random variable on $[0, 1]$.

³By “positive,” we simply mean the decision D is greater than zero, without ascribing any normative position to the decision. Individuals may or may not have a preferences for “positive” decisions in this sense.

⁴As is implicit in our notation, we assume that there are no spillover effects between units (Imbens and Rubin, 2015).

to the fact that everyone could be screened in principle.

Our second example concerns college admissions. Here, the population of interest is applicants to a particular college, and the decision D is the admissions committee’s binary admissions decision. To simplify our exposition, we assume all admitted students attend the school. In this setting, the covariates X may, for example, consist of an applicant’s test score and race $A \in \{a_0, a_1\}$, and Y is a binary variable that indicates college graduation (i.e., degree attainment). In contrast to our diabetes example, here we imagine that the decision itself may affect the outcomes. Specifically, $Y(1)$ and $Y(0)$ describe whether an applicant *would* attain a college degree if admitted to or if rejected from the school we consider, respectively. Note that $Y(0)$ is not necessarily zero, as a rejected applicant may attend—and graduate from—a different university. Further, in this case we set the budget b to be less than one to reflect the fact that the admissions committee has limited resources and is unable to admit every candidate.

As mentioned above, a key distinction between these two examples is the extent to which stakeholders may agree on the value of different potential decision policies. For example, in college admissions, there may be significant disagreement on how to balance competing priorities, such as academic preparedness and class diversity.⁵ Admissions committees may seek to increase both dimensions, but there is often an inherent trade-off, particularly since there is a limit on the number of students that can be admitted by the college (i.e., $b < 1$). Our diabetes example, in contrast, reflects a setting where there is ostensibly broader agreement on the value of different decision policies. Indeed, since there is effectively no limit on the number of diabetes tests that can be administered (i.e., $b = 1$), we can model the value of a decision policy as the sum of each individual’s value for being screened.⁶ In Sections 2.3 and 2.4, we in turn examine the structure of equitable decision making in the absence and presence of such trade-offs. First, though, we introduce several formal fairness criteria.

⁵In some jurisdictions, explicit considerations of racial diversity may be prohibited. For instance, a recent U.S. Supreme Court case bars colleges from explicitly considering race in admissions; however, colleges may consider “an applicant’s discussion of how race affected the applicant’s life” (SFFA v. Harvard, 2023). U.S. colleges may also consider other forms of diversity, such as economic or geographic diversity.

⁶In the case of infectious diseases—which involve greater externalities—there is again often disagreement about the value of different screening and vaccination policies. Paulus and Kent (2020) similarly draw a distinction between *polar* settings (in which parties have competing interests, like our admissions example) and *non-polar* settings (where there is broad alignment, as in our diabetes example).

2.2.2 Limiting the Effect of Decisions on Disparities

A popular class of fairness definitions requires that error rates (e.g., false positive and false negative rates) are equal across protected groups (Hardt et al., 2016).⁷ We refer to these definitions as examples of “classification parity,” meaning that some given measure of classification error is equal across groups defined by attributes such as race and gender. In particular, we include in this definition any measure that can be computed from the two-by-two confusion matrix tabulating the joint distribution of decisions D and outcomes Y for a group. Berk et al. (2021) enumerate seven such statistics, including false positive rate, false negative rate, precision, recall, and the proportion of decisions that are positive. The proportion of positive decisions is not, strictly speaking, a measure of “error”, but we nonetheless include it under classification parity since it can be computed from a confusion matrix. We also include the area under the ROC curve (AUC), a popular measure among practitioners examining the fairness of algorithms (Skeem and Lowenkamp, 2016).

Two of the above measures—the proportion of decisions that are positive, and the false positive rate—have received considerable attention in the machine learning community (Agarwal et al., 2018; Blum and Stangl, 2019; Calders and Verwer, 2010; Edwards and Storkey, 2016; Feldman et al., 2015; Jung et al., 2020a; Kamiran et al., 2013; Pedreshi et al., 2008; Zafar et al., 2017c; Zemel et al., 2013; Chouldechova, 2017; Hardt et al., 2016; Zafar et al., 2017).

Definition 2.1. We say that *demographic parity* holds when⁸

$$D \perp\!\!\!\perp A. \tag{2.1}$$

Definition 2.2. We say that *equalized false positive rates* holds when

$$D \perp\!\!\!\perp A \mid Y = 0. \tag{2.2}$$

In our running diabetes example, demographic parity means that the proportion of patients who are screened for the disease is equal across race groups. Similarly, in our college admissions example, demographic parity means an equal proportion of students is admitted across race groups. Equalized

⁷Some work relaxes strict equality of error rates or other metrics to requiring only that the difference be at most some fixed ϵ (e.g., Nabi and Shpitser, 2018). For ease of exposition, we consider strict equality throughout, though we emphasize that the spirit of the critique we develop applies also in cases where fairness constraints are approximate, rather than exact.

⁸We use the notation $X \perp\!\!\!\perp Y$ throughout to mean that the random variables X and Y are independent.

false positive rates, in our diabetes example, means that among individuals who in reality do not have diabetes—and thus for whom screening, *ex post*, would not have been beneficial—screening rates are equal across race groups.⁹

Causal analogues of these definitions have also recently been proposed (Imai et al., 2020; Mishler et al., 2021; Coston et al., 2020; Imai and Jiang, 2020), which require various conditional independence conditions to hold between the potential outcomes, protected attributes, and decisions.¹⁰ Below we list three representative examples of this class of fairness definitions: counterfactual predictive parity (Coston et al., 2020), counterfactual equalized odds (Mishler et al., 2021; Coston et al., 2020), and conditional principal fairness (Imai and Jiang, 2020).¹¹

Definition 2.3. We say that *counterfactual predictive parity* holds when

$$Y(1) \perp\!\!\!\perp A \mid D = 0. \quad (2.3)$$

In our college admissions example, counterfactual predictive parity means that among rejected applicants, the proportion who would have attained a college degree, had they been accepted, is equal across race groups. (For our diabetes example, because the screening decision does not affect whether a patient actually has diabetes, $Y(0) = Y(1) = Y$, and so counterfactual predictive parity, as well as the causal definitions below, reduce to their non-causal analogues).

Definition 2.4. We say that *counterfactual equalized odds* holds when

$$D \perp\!\!\!\perp A \mid Y(1). \quad (2.4)$$

In our running college admissions example, counterfactual equalized odds is satisfied when two conditions hold: (1) among applicants who would graduate if admitted (i.e., $Y(1) = 1$), students are

⁹In our college admissions example, the decision D impacts the outcome Y . One could, in theory, apply the definition of error rate parity above to that case by recognizing that $Y = Y(D)$. However, that interpretation does not seem aligned with the original intent of the definition. We instead discuss the admissions example in the context of the explicitly causal definitions of fairness below.

¹⁰In the literature on causal fairness, there is at times ambiguity between “predictions” $\hat{Y} \in \{0, 1\}$ of Y and “decisions” $D \in \{0, 1\}$. Following past work (e.g., Kusner et al., 2017; Wang et al., 2019; Corbett-Davies et al., 2017), here we focus exclusively on decisions, with predictions implicitly impacting decisions but not explicitly appearing in our definitions.

¹¹Our subsequent analytical results extend in a straightforward manner to structurally similar variants of these definitions (e.g., requiring $Y(0) \perp\!\!\!\perp A \mid D = 1$ or $D \perp\!\!\!\perp A \mid Y(0)$, variants of counterfactual predictive parity and counterfactual equalized odds, respectively).

admitted at the same rate across race groups; and (2) among applicants who would not graduate if admitted (i.e., $Y(1) = 0$), students are again admitted at the same rate across race groups.

Definition 2.5. We say that *conditional principal fairness* holds when

$$D \perp\!\!\!\perp A \mid Y(0), Y(1), W, \quad (2.5)$$

where, for some function ω on $\mathcal{X}$, $W = \omega(X)$ describes a reduced set of the covariates X . When W is constant (or, equivalently, when we do not condition on W), this condition is called *principal fairness*.

In the college admissions example, conditional principal fairness means that “similar” applicants—where similarity is defined by the potential outcomes and covariates W —are admitted at the same rate across race groups.

2.2.3 Limiting the Effect of Attributes on Decisions

An alternative framework for understanding fairness considers the effects of protected attributes on decisions. This approach can be understood as codifying the legal notion of disparate treatment (Zafar et al., 2017; Goel et al., 2017)—which we discuss further in Section 2.5.1. Perhaps the simplest way to limit the effects of protected attributes on decisions is to require that the decisions do not explicitly depend on them, what some call “fairness through unawareness” (cf. Dwork et al., 2012).

Definition 2.6. Suppose that the covariates can be partitioned into the protected attributes and all other covariates, i.e., that $\mathcal{X} = \mathcal{X}_u \times \mathcal{A}$, where $\mathcal{X}_u$ consists of “unprotected” attributes. Then, we say that *blinding* holds when, for all $a, a' \in \mathcal{A}$ and $x_u \in \mathcal{X}_u$,

$$d(x_u, a) = d(x_u, a'). \quad (2.6)$$

In our running diabetes example, blinding holds when the screening decision depends solely on factors like age and BMI, and, in particular, does not depend on the patient’s race. We similarly say college admissions decisions satisfy blinding when the decisions depend on factors like test scores and extracurricular activities, but not race.

Blinding is closely tied to the notion of *calibration*, the requirement that, conditional on the estimated probability of some outcome (such as graduation from college or having diabetes), the

outcome is independent of group membership. For example, among people with an estimated diabetes risk of 1%, calibration would require that the proportion of individuals who actually have diabetes be the same across groups. Many authors treat calibration as a kind of fairness constraint—in particular, to ensure that the meaning of estimated risks do not differ across groups—and it has received considerable attention in the fairness literature (e.g., Hébert-Johnson et al., 2018; Rothblum and Yona, 2022). We note, though, that miscalibration is equivalent to blindness in practice. In particular, when estimation error is small, risk estimates that are allowed to depend on group membership are calibrated; conversely, risk estimates that are blind to group membership typically are miscalibrated—an empirical phenomenon shown and discussed in Figure 2.3 below. Because of this close relationship, we do not treat calibration as a separate fairness constraint, but we do discuss calibration and its relationship to blinding in detail in Sections 2.3.3.1 and 2.5.2.

In contrast to blinding—in which race and other protected attributes are barred from being an explicit input to a decision rule—the causal versions of this idea consider both the direct and indirect effects of protected attributes on decisions (Mhasawade and Chunara, 2021; Wu et al., 2019; Zhang and Bareinboim, 2018; Zhang et al., 2017; Kilbertus et al., 2017; Kusner et al., 2017; Nabi and Shpitser, 2018; Wang et al., 2019). For example, even if decisions only directly depend on test scores, race may indirectly impact decisions through its effects on educational opportunities, which in turn influence test scores. In this vein, a decision rule is deemed fair if, at a high level, decisions for individuals are the same in “(a) the actual world and (b) a counterfactual world where the individual belonged to a different demographic group” (Kusner et al., 2017).¹² This idea can be formalized by requiring that decisions remain the same in expectation even if one’s protected characteristics are counterfactually altered, a condition known as counterfactual fairness (Kusner et al., 2017).

¹²Conceptualizing a general causal effect of an immutable characteristic such as race or gender is rife with challenges, the greatest of which is expressed by the mantra, “no causation without manipulation” (Holland, 1986). In particular, analyzing race as a causal treatment requires one to specify what exactly is meant by “changing an individual’s race” from, for example, White to Black (Gaebler et al., 2022; Hu and Kohler-Hausmann, 2020). Such difficulties can sometimes be addressed by considering a change in the *perception* of race by a decision maker (Greiner and Rubin, 2011)—for instance, by changing the name listed on an employment application (Bertrand and Mullainathan, 2004), or by masking an individual’s appearance (Chohlas-Wood et al., 2021; Goldin and Rouse, 2000; Grogger and Ridgeway, 2006; Pierson et al., 2020).




Figure 2.1: A causal DAG illustrating a hypothetical process for college admissions. Under path-specific fairness, one may require, for example, that race does not affect decisions along the path highlighted in red.

Definition 2.7. *Counterfactual fairness* holds when

$$\mathbb{E}[D(a') \mid X] = \mathbb{E}[D \mid X], \quad (2.7)$$

where $D(a')$ denotes the decision when one’s protected attributes are counterfactually altered to be any $a' \in \mathcal{A}$.

In our running college admissions example, this means that for each group of observationally identical applicants (i.e., those with the same values of X , meaning identical race and test score), the proportion of students who are actually admitted is the same as the proportion who would be admitted if their race were counterfactually altered.

Counterfactual fairness aims to limit all direct and indirect effects of protected traits on decisions. In a generalization of this criterion—termed path-specific fairness (Wu et al., 2019; Zhang et al., 2017; Chiappa, 2019; Nabi and Shpitser, 2018)—one allows protected traits to influence decisions along certain causal paths but not others. For example, one may wish to allow the direct consideration of race by an admissions committee to implement an affirmative action policy, while also guarding against any indirect influence of race on admissions decisions that may stem from cultural biases in standardized tests (Williams, 1983).

The formal definition of path-specific fairness requires specifying a causal DAG describing relationships between attributes (both observed covariates and latent variables), decisions, and outcomes. In our running example of college admissions, we imagine that each individual’s observed covariates are the result of the process illustrated by the causal DAG in Figure 2.1. In this graph,

an applicant’s race A influences the educational opportunities E available to them prior to college; and educational opportunities in turn influence an applicant’s level of college preparation, M , as well as their score on a standardized admissions test, T , such as the SAT. We assume the admissions committee only observes an applicant’s race and test score so that $X = (A, T)$, and makes their decision D based on these attributes. Finally, whether or not an admitted student subsequently graduates (from any college), Y , is a function of both their preparation and whether they were admitted.¹³

To formalize path-specific fairness, we start by defining, for the decision D , path-specific counterfactuals, a general concept in causal DAGs (cf. Pearl, 2001). Suppose $\mathcal{G} = (\mathcal{V}, \mathcal{U}, \mathcal{F})$ is a causal model with nodes $\mathcal{V}$, exogenous variables $\mathcal{U}$, and structural equations $\mathcal{F}$ that define the value at each node V_j as a function of its parents $\wp(V_j)$ and its associated exogenous variable U_j . (See, for example, Pearl (2009a) for further details on causal DAGs.) Let $V_1, \dots, V_m$ be a topological ordering of the nodes, meaning that $\wp(V_j) \subseteq \{V_1, \dots, V_{j-1}\}$ (i.e., the parents of each node appear in the ordering before the node itself). Let Π denote a collection of paths from node A to D . Now, for two possible values a and a' for the variable A , the path-specific counterfactuals $D_{\Pi, a, a'}$ for the decision D are generated by traversing the list of nodes in topological order, propagating counterfactual values obtained by setting $A = a'$ along paths in Π , and otherwise propagating values obtained by setting $A = a$. (In Algorithm A.1, we formally define path-specific counterfactuals for an arbitrary node—or collection of nodes—in the DAG.)

To see this idea in action, we work out an illustrative example, computing path-specific counterfactuals for the decision D along the single path $\Pi = \{A \rightarrow E \rightarrow T \rightarrow D\}$ linking race to the admissions committee’s decision through test score, highlighted in red in Figure 2.1. We describe the distribution of $D_{\Pi, a, a'}$ generatively, formally showing how to produce a draw from this distribution. To start, we draw values U_E^* , U_M^* , U_T^* , U_D^* of the exogenous variables. Now, the first column in Table 2.1 corresponds to draws V^* for each node V in the DAG, where we set A to a , and then propagate that value as usual. The second column corresponds to draws $\bar{V}^*$ of path-specific counterfactuals, where we set A to a' , and then propagate the counterfactuals only along the path

¹³In practice, the racial composition of an admitted class may itself influence degree attainment, if, for example, diversity provides a net benefit to students (Page, 2007). Here, for simplicity, we avoid consideration of such peer effects.

$A^* = a$	$\bar{A}^* = a'$
$E^* = f_E(A^*, U_E^*)$	$\bar{E}^* = f_E(\bar{A}^*, U_E^*),$
$M^* = f_M(E^*, U_M^*)$	$\bar{M}^* = f_M(E^*, U_M^*),$
$T^* = f_T(E^*, M^*, U_T^*)$	$\bar{T}^* = f_T(\bar{E}^*, M^*, U_T^*)$
$D^* = f_D(A^*, T^*, U_D^*)$	$\bar{D}^* = f_D(A^*, \bar{T}^*, U_D^*)$

Table 2.1: Computing path-specific counterfactuals for the DAG in Figure 2.1. The first column corresponds to draws V^* for each node V , where we set A to a , and then propagate that value as usual. The second column corresponds to draws $\bar{V}^*$ of path-specific counterfactuals, where we set A to a' , and then propagate the counterfactuals only along the path $A \rightarrow E \rightarrow T \rightarrow D$.

$A \rightarrow E \rightarrow T \rightarrow D$. In particular, the value for the test score $\bar{T}^*$ is computed using the value of $\bar{E}^*$ (since the edge $E \rightarrow T$ is on the specified path) and the value of M^* (since the edge $M \rightarrow T$ is not on the path). As a result of this process, we obtain a draw $\bar{D}^*$ from the distribution of $D_{\Pi, a, a'}$.

Path-specific fairness formalizes the intuition that the influence of a sensitive attribute on a downstream decision may, in some circumstances, be considered “legitimate” (i.e., it may be acceptable for the attribute to affect decisions along certain paths in the DAG). For instance, an admissions committee may believe that the effect of race A on admissions decisions D which passes through college preparation M is legitimate, whereas the effect of race along the path $A \rightarrow E \rightarrow T \rightarrow D$, which may reflect access to test prep or cultural biases of the tests, rather than actual academic preparedness, is illegitimate. In that case, the admissions committee may seek to ensure that the proportion of applicants they admit from a certain race group remains unchanged if one were to counterfactually alter the race of those individuals along the path $\Pi = \{A \rightarrow E \rightarrow T \rightarrow D\}$.

Definition 2.8. Let Π be a collection of paths, and, for some function ω on $\mathcal{X}$, let $W = \omega(X)$ describe a reduced set of the covariates X . *Path-specific fairness*, also called Π -*fairness*, holds when, for any $a' \in \mathcal{A}$,

$$\mathbb{E}[D_{\Pi, A, a'} \mid W] = \mathbb{E}[D \mid W]. \quad (2.8)$$

In the definition above, rather than a particular counterfactual level a , the baseline level of the path-specific effect is A , i.e., an individual’s actual (non-counterfactually altered) group membership (e.g., their actual race). We have implicitly assumed that the decision variable D is a descendant

of the covariates X . In particular, without loss of generality, we assume D is defined by the structural equation $f_D(x, u_D) = \mathbb{1}_{u_D \leq d(x)}$, where the exogenous variable $U_D \sim \text{UNIF}(0, 1)$, so that $\Pr(D = 1 \mid X = x) = d(x)$. If Π is the set of all paths from A to D , then $D_{\Pi, A, a'} = D(a')$, in which case, for $W = X$, path-specific fairness is the same as counterfactual fairness.

2.3 Equitable Decisions in the Absence of Externalities

In many decision-making settings, the decision maker is free to make the optimal decision for each individual, without consideration of spillover effects or other externalities. For instance, in our diabetes screening example, one could, in principle, screen all patients if that course of action were medically advisable.

To investigate notions of fairness in these settings, we first introduce a framework for utilitarian decision analysis. Specifically, we consider in this section situations in which there is broad agreement on the utility of different potential courses of action. (In the subsequent section, we consider cases where stakeholders disagree on the precise form of the utility.) In this setting, “threshold rules” maximize utility. We then describe the statistical phenomenon of inframarginality, a property that is endemic to fairness definitions that seek to enforce some form of classification parity. In particular, we discuss, both informally and mathematically, why inframarginality almost surely—in a measure theoretic sense—renders optimal decision making incompatible with classification parity. Finally, we discuss blinding. In parallel to our discussion of classification parity, we see that in many important settings, the information loss associated with, e.g., removing protected information from a predictive model, results in less efficient decision making without compensatory benefits. Moreover, in general, we see that the more stringent the standard of masking—e.g., removing not only direct but also indirect effects of protected attributes—the greater the potential harm that results from enforcing it.

2.3.1 Utility, Risk, and Threshold Rules

A natural way to analyze a decision, such as deciding whether an individual should be screened for diabetes, is to consider the costs and benefits of various possible outcomes under different courses of action. For instance, a patient screened for diabetes who does not have the disease still has to bear

the risks, discomfort, and inconvenience associated with the blood test itself, while a patient who is *not* screened but does in fact have the disease loses out on the opportunity to start treatment.

In general, the benefit of making decision $D = 1$ over $D = 0$ when the outcome Y equals y can be represented by $v(y)$. For instance, in our diabetes example, $v(1)$ represents the net benefit of screening over not screening when the patient has diabetes; and $-v(0)$ is the net cost of screening when the patient does not have diabetes, including both monetary and non-monetary costs, such as discomfort and loss of time.¹⁴ Let $r(x) = \Pr(Y = 1 \mid X = x)$ be the *risk* of Y equalling 1 when $X = x$. Then the expected benefit of making decision $D = 1$ over $D = 0$ for an individual with covariates $X = x$ is

$$\begin{aligned} u(x) &= \mathbb{E}[v(Y) \mid X = x] \\ &= r(x) \cdot v(1) + [1 - r(x)] \cdot v(0). \end{aligned}$$

Here, for ease of interpretation, we restrict our utility to be of the form $u(x) = \mathbb{E}[v(Y) \mid X = x]$ for some function v , and we also assume there is no budget constraint (i.e., $b = 1$). In Section 2.4, we allow the utility $u(x)$ to be an arbitrary function on $\mathcal{X}$ and consider $b < 1$, which induces the trade-offs in decisions that are central to our later discussion.

The *aggregate* expected utility of a decision policy $d(x)$ —relative to the baseline policy of taking action $D = 0$ for all individuals—is then given by $u(d) = \mathbb{E}[d(X) \cdot u(X)]$. We say a decision policy $d^*(x)$ is utility-maximizing if

$$u(d^*) = \max_d u(d).$$

It is better, in expectation, for an individual with covariates $X = x$ to take action $D = 1$ instead of $D = 0$ when $u(x) > 0$; that is, when¹⁵

$$r(x) > \frac{v(0)}{v(0) - v(1)}. \quad (2.9)$$

Thus, the decision with the maximum utility can be determined by comparing an individual's risk

¹⁴For ease of exposition, we assume that costs and benefits are identical across individuals; in reality, these could vary, e.g., depending on age. When utilities vary by person, the optimal decision rule is to screen only those with positive individual utility, in line with our subsequent discussion.

¹⁵We assume, without loss of generality, that $v(1) > v(0)$. If $v(1) < v(0)$, we can take $Y' = 1 - Y$ as our outcome of interest; relative to Y' , the inequality will be reversed. If $v(1) = v(0)$, then the outcome is irrelevant. In this degenerate case, the higher utility decision depends on the sign of $v(1)$ alone, and not the risk.

against a particular risk threshold t , defined by the right-hand side of Eq. (2.9). We refer to this kind of policy as a *threshold policy*. In particular, we see that a utility-maximizing decision for each individual—i.e., $d(x) = 1$ if $r(x) > t$ and $d(x) = 0$ if $r(x) \leq t$ —is also a decision policy that maximizes aggregate utility, so there is no conflict between doing what is best from each individual person’s perspective and what is best for the population as a whole.

While our framing in terms of expected utility is suitably general, threshold policies can be simpler to interpret when we reparameterize in terms of more familiar quantities. In the diabetes screening example, if the patient does not have diabetes, the cost of action $D = 1$ over $D = 0$ is $-v(0) = c_{\text{TEST}}$, i.e., the cost (monetary and non-monetary) of the test. If the patient does have diabetes, the benefit of $D = 1$ over $D = 0$ is $v(1) = b_{\text{TREAT}} - c_{\text{TEST}}$, i.e., the benefit of treatment minus the cost of the test. Rewriting Eq. (2.9) in terms of these quantities gives

$$t = \frac{c_{\text{TEST}}}{b_{\text{TREAT}}}.$$

In particular, if the benefit of early treatment of diabetes is 50 times greater than the cost of performing the diagnostic test, one would ideally screen patients who have at least a 2% chance of developing the disease.

Threshold rules are a natural approach to decision making in a variety of settings. In our running medical example, a threshold rule corresponds to screening patients with a sufficiently high risk of having diabetes. A threshold rule—with the optimally chosen threshold—ensures that only the patients at highest risk of having diabetes take the test, thereby optimally balancing the costs and benefits of screening. Indeed, in many medical examples, from diagnosis to treatment, there are no significant externalities. As a result, deviating from utility-maximizing threshold policies can only force individuals to experience greater costs—in the form of unnecessary tests or untreated illness—in expectation, without compensatory benefits. We return to the problem of optimal (and equitable) decision-making in the presence of externalities in Section 2.4.

2.3.2 The Problem of Inframarginality

In the setting that we have been considering, threshold policies guarantee optimal choices are made for each individual. However, as we now show, threshold policies in general violate various

versions of classification parity, such as demographic parity and equalized false positive rates. This incompatibility highlights a critical limitation of classification parity as a fairness criterion, as enforcing the definition often requires making decisions that harm individuals without any clear compensating benefits.

To help build intuition for this phenomenon, we consider the empirical distribution of diabetes risk among White and Asian patients. Following Aggarwal et al. (2022), we base our risk estimates on age, BMI, and race, using a sample of approximately 15,000 U.S. adults aged 18–70 interviewed as part of the National Health and Nutrition Survey (NHANES; Centers for Disease Control and Prevention, 2011-2018). The resulting risk distributions are shown in the left-hand panel of Figure 2.2. The dashed vertical lines show the group means, and indicate that the incidence of diabetes is higher among Asian Americans (11%) than among White Americans (9%).¹⁶ This difference in base rates is also reflected in the heavier tail of the risk distribution among Asian individuals.

Drawing on recommendations from the United States Preventative Screening Task Force, Aggarwal et al. (2022) suggest screening patients with at least a 1.5% risk of diabetes, irrespective of race. We depict this risk threshold by the solid black vertical line in the plot. Based on that recommendation, 81% of Asian Americans and 69% of White Americans are to the right of the threshold and should be screened—violating demographic parity. If, hypothetically, we were to raise the screening threshold to 2.2% for Asian Americans and lower the threshold to 1% for White Americans, 75% of people in both groups would be screened, satisfying demographic parity.¹⁷ The cost of doing so, however, would be failing to screen some Asian Americans who have a relatively high risk of diabetes, and subjecting some relatively low-risk White Americans to a procedure that is medically inadvisable given their low likelihood of having diabetes. In an effort to satisfy demographic parity, we would have harmed members from both groups.

This example illustrates a similar incompatibility between threshold policies and equalized false positive rates. In our setting, the false positive rate for a group is the screening rate among those in the group who do not in reality have diabetes. To visualize the race-specific false positive rates, the

¹⁶The precise shapes of the risk distributions depend on the set of covariates used to estimate outcomes, but the means of the distributions correspond to the overall incidence of diabetes in each group, and, in particular, are unaffected by the choice of covariates. It is thus necessarily the case that the risk distributions will differ across groups in this example, regardless of which covariates are used.

¹⁷Corbett-Davies et al. (2017) show that group-specific threshold policies are utility-maximizing under the constraint of satisfying various notions of classification parity, including demographic parity and equality of false positive rates.




Figure 2.2: A graphical illustration of the incompatibility between threshold policies and classification parity, based on the National Health and Nutrition Survey. *Left*: The distribution of diabetes risk for White Americans and Asian Americans, with the dashed vertical lines corresponding to the overall incidence rate within each group. At a screening threshold of 1.5% (indicated by the solid black line), the screening rate for Asian Americans is higher than for White Americans, violating demographic parity. *Right*: The distribution of diabetes risk among individuals who do *not* have diabetes. Since the proportion of Asian Americans above the screening threshold is greater than the proportion of White Americans above the threshold, the false positive rate for Asian Americans is greater than the false positive rate for White Americans.

right-hand panel of Figure 2.2 shows the distribution of diabetes risk among those individuals who do not have diabetes. (Because the overall prevalence of diabetes is low, the conditional distribution displayed in the right-hand panel is nearly identical to the unconditional distribution displayed in the left-hand panel.) The false positive rate for each group is the proportion of people in the group falling to the right of the 1.5% screening threshold. In this case, the false positive rate is 79% for Asian Americans and 67% for White Americans—violating equalized false positive rates. As before, we could alter the screening guidelines to equalize false positive rates, but doing so requires deviating from our threshold policy, in which case we would end up screening some individuals who are relatively low-risk and not screening others who are relatively high-risk.

In this example, the incompatibility between threshold policies and classification parity stems from the fact that the risk distributions differ across groups. This general phenomenon is known as the problem of *inframarginality* in the economics and statistics literature, and has long been known to plague tests of discrimination in human decisions (Anwar and Fang, 2006; Carr and Megbolugbe, 1993; Engel and Tillyer, 2008; Galster, 1993; Knowles et al., 2001; Pierson et al., 2018; Ayres, 2002;

Simoiu et al., 2017). Common legal and economic understandings of fairness are concerned with what happens at the *margin* (e.g., whether the same standard is applied to all individuals)—a point we return to in Section 2.5. What happens at the margin also determines whether decisions maximize social welfare, with the optimal threshold set at the point where the marginal benefits equal marginal costs. However, popular error metrics assess behavior away from the margin, hence they are called *infra*-marginal statistics. As a result, when risk distributions differ, standard error metrics are often poor proxies for individual equity or social well-being.

In general, we expect any two non-random subgroups of a population to differ on a variety of social and economic dimensions, which in turn is likely to yield risk distributions that differ across groups. As a result, as our running diabetes example shows, the optimal decision policy—which maximizes each patient’s own well-being—will likely violate various measures of classification parity. Thus, to the extent that formal measures of fairness are violated, that tells us more about the shapes of the risk distributions than about the quality of decisions or the utility delivered to members of any group. This intuition can be made precise, in the sense that for *almost every* risk distribution, the optimal decision policy violates the various notions of classification parity considered here.

The notion of *almost every* distribution that we use here was formalized by Christensen (1972), Hunt et al. (1992), Anderson and Zame (2001), and others (cf. Ott and Yorke, 2005, for a review). Suppose, for a moment, that combinations of covariates and outcomes take values in a finite set of size m . Then the space of joint distributions on covariates and outcomes can be represented by the unit $(m - 1)$ -simplex: $\Delta^{m-1} = \{p \in \mathbb{R}^m \mid p_i \geq 0 \text{ and } \sum_{i=1}^m p_i = 1\}$. Since Δ^{m-1} is a subset of an $(m - 1)$ -dimensional hyperplane in $\mathbb{R}^m$, it inherits the usual Lebesgue measure on $\mathbb{R}^{m-1}$. In this finite-dimensional setting, *almost every* distribution means a subset of distributions that has full Lebesgue measure on the simplex. Given a property that holds for almost every distribution in this sense, that property holds almost surely under any probability distribution on the space of distributions that is described by a density on the simplex. We use a generalization of this basic idea that extends to infinite-dimensional spaces, allowing us to consider distributions with arbitrary support. (See Appendix A.6.1 for further details.)

Theorem 2.9. *Let t be the optimal decision threshold, as in Eq. (2.9). If $0 < t < 1$, then for almost every collection of group-specific risk distributions which have densities on $[0, 1]$, no utility-maximizing*

decision policy satisfies demographic parity or equalized false positive rates.

The proof of Theorem 2.9, which formalizes the informal discussion above, is given in Appendix A.6.6. At a high level, the constraints of classification parity are sensitive to even small perturbations in the underlying risk distributions. As a result, any particular collection of risk distributions is unlikely to satisfy the constraints. For simplicity, we have been considering settings in which the decision D does not impact the outcome Y . However, this basic style of argument extends to causal settings, showing that threshold policies are almost surely, in the measure theoretic sense, incompatible with counterfactual predictive parity, counterfactual equalized odds, and conditional principal fairness—definitions of fairness that we consider in depth in Section 2.4, in the more complex setting of having a budget $b < 1$.


2.3.3 The Problem with Fairness through Unawareness

We now consider notions of fairness, both causal and non-causal, that aim to limit the effects of attributes on decisions. As above, we show the inherent incompatibility of these definitions with optimal decision making. We note, though, that while blinding can lead to suboptimal decisions—and, in some cases, harm marginalized groups—the legal, political, and social benefits of, for example, race-blind and gender-blind algorithms may outweigh their costs in certain instances (Cerdeña et al., 2020; Coots et al., 2023).


2.3.3.1 Blinding

A common starting point for designing an ostensibly fair algorithm is to exclude protected characteristics from the statistical model. This strategy ensures that decisions have no explicit dependence on group membership. For instance, in the case of estimating diabetes risk, one could use only BMI and age—rather than including race, as we did above. However, excluding race from models of diabetes risk can ultimately harm both White and Asian patients.

In Figure 2.3a, we compare the actual diabetes rate to estimated diabetes risk resulting from the race-blind risk model. Aggarwal et al. (2022) showed that Asian patients have higher incidence of diabetes than White patients with comparable age and BMI. As a result, the race-blind model systematically underestimates risk for Asian patients and systematically overestimates risk for White



(a) Diabetes



(b) Recidivism

Figure 2.3: Calibration plots showing the effect of removing protected attributes from risk models when estimating the risk of diabetes (left) and recidivism (right). Because Asian patients with the same BMI and age have a higher rate of diabetes, the race-blind model underestimates their risk of having diabetes. Similarly, because women reoffend at lower rates than men with similar criminal histories, the gender-blind COMPAS score overstates the recidivism risk for women.

individuals. In particular, applying a nominal 1.5% screening threshold under the race-blind model amounts to effectively applying a 1% screening threshold to White patients and a 3% screening threshold to Asian patients. Thus, by using race-blind risk scores, we subject relatively low-risk White patients to screening, and fail to screen Asian patients who have a relatively high risk for having diabetes. A race-aware model would ensure that nominal risk thresholds correspond to observed incidence rates across race groups.

This phenomenon—which we call miscalibration across subgroups—is not unique to diabetes screening. Consider, for instance, the case of pretrial recidivism predictions. Shortly after an individual is arrested in the United States, a judge must often determine conditions of release pending future court proceedings. In many jurisdictions across the country, these pretrial decisions are informed by statistical risk estimates of the likelihood the individual would be arrested or convicted of a future crime if released. After adjusting for factors such as criminal history, age, and substance use, women have been found to reoffend less often than men in many jurisdictions (DeMichele et al., 2018; Skeem et al., 2016). Consequently, gender-blind risk assessments are miscalibrated, meaning that they tend to overstate the recidivism risk of women and understate the recidivism risk of men.

Figure 2.3b illustrates this point, plotting the observed recidivism rate for men and women in

Broward County, Florida as a function of their gender-blind COMPAS risk scores—a commonly used risk assessment tool (Bao et al., 2021). In particular, women with a COMPAS score of seven recidivate about 55% of the time, whereas men with the same score recidivate about 65% of the time. Said differently, women with a score of seven recidivate approximately as often as men with a score of five, and this two-point differential persists across the range of scores. By acknowledging the predictive value of gender in this setting, one could create a decision rule that detains fewer people (particularly women) while achieving the same public safety benefits. Conversely, by ignoring this information and basing decisions solely on the gender-blind risk assessments, one would effectively be subjecting women to a more stringent risk standard—and potentially harsher penalties—than men.

As in the case of classification parity, one cannot typically remove protected attributes from the risk predictions without decreasing utility (cf. Manski et al., 2022); however, the reduction in utility is not always as large as one might expect (Coots et al., 2023). In concrete terms, in our running diabetes example, basing decisions on race-blind risk estimates necessarily means screening some patients who would have preferred not to be screened had they been given race-aware risk estimates, and, conversely, not screening some patients who would have preferred to be screened had they been given the more complete estimates. We state this result formally below.

Theorem 2.10. *Suppose $0 < t < 1$, where t is the optimal decision threshold on the risk scale, as in Eq. (2.9). Let $\pi : \mathcal{X}_u \times \mathcal{A} \rightarrow \mathcal{X}_u$ denote restriction to the unprotected covariates. Let $\rho(x) = \Pr(Y = 1 \mid \pi(X) = \pi(x))$ denote the risk estimated using the blinded covariates. Suppose that $r(x)$ and $\rho(x)$ have densities on $[0, 1]$ that are positive in a neighborhood of t . Further suppose that there exists $\epsilon > 0$ such that the conditional variance $\text{VAR}(r(X) \mid \rho(X)) > \epsilon$ a.s., where $r(x)$ is the risk estimated from the full set of covariates. Then no blind policy is utility-maximizing.*

The proof of Theorem 2.10 is given in Appendix A.4. In short, when race, gender, or other protected traits add predictive value—a condition codified in our assumption that the conditional variance be greater than ϵ —excluding these attributes will in general decrease utility, both for individuals and in the aggregate.

Basing decisions on blinded risk scores can harm individuals and communities, for example by failing to flag relatively high-risk Asian patients for diabetes screening. But it is also important to

consider potential harms stemming from the use of race- and gender-specific risk tools. In medicine, for instance, one might worry that race-specific risk assessments could encourage doctors and the public-at-large to advance spurious and pernicious arguments about inherent differences between race groups. In reality, the differences in diabetes risk we see are likely due to a complex mix of factors, both environmental and genetic, and should not be misinterpreted as indicating any causal effects of race. Indeed, even “race” itself is a thorny, socially influenced, concept, that eludes easy definition. Similarly, the use of gender-specific recidivism estimates could reduce trust in the criminal justice system, giving the impression that individuals are held to different standards based on their gender. (Though, as we have seen above, *blinded* risk assessments can likewise—and perhaps more persuasively—be said to subject individuals to different standards based on their race and gender.) In some circumstances, race- and gender-specific risk estimates are even prohibited by law—a topic we return to in Section 2.5.1. For these reasons, risk assessments in medicine, criminal justice, and beyond have generally avoided using race, gender, and other sensitive demographic attributes. Ultimately, when constructing risk assessment tools, it is important to acknowledge and carefully balance both the costs and benefits of blinding in any given circumstance.

2.3.3.2 Counterfactual and Path-Specific Fairness

As discussed in Section 2.2, counterfactual and path-specific fairness are generalizations of simple blinding that attempt to account for both the direct and indirect effects of protected attributes on decisions. Because the constraints are more stringent, the resulting decrease in utility is proportionally greater. In particular, in some common settings, path-specific fairness with $W = X$ constrains decisions so severely that the only allowable policies are constant (i.e., $d(x_1) = d(x_2)$ for all $x_1, x_2 \in \mathcal{X}$). For instance, in our running admissions example, path-specific fairness requires admitting all applicants with the same probability, irrespective of academic preparation or group membership.

To build intuition for this result, we sketch the argument for a finite covariate space $\mathcal{X}$. Given a policy d that satisfies path-specific fairness, select $x^* \in \arg \max_{x \in \mathcal{X}} d(x)$. By the definition of

path-specific fairness, for any $a \in \mathcal{A}$,

$$\begin{aligned} d(x^*) &= \mathbb{E}[D_{\Pi, A, a} \mid X = x^*] \\ &= \sum_{x \in \alpha^{-1}(a)} d(x) \cdot \Pr(X_{\Pi, A, a} = x \mid X = x^*). \end{aligned} \tag{2.10}$$

That is, the probability of an individual with covariates x^* receiving a positive decision must be the average probability of the individuals with covariates x in group a receiving a positive decision, weighted by the probability that an individual with covariates x^* in the real world *would* have covariates x counterfactually.

Next, we suppose that there exists an $a' \in \mathcal{A}$ such that $\Pr(X_{\Pi, A, a'} = x \mid X = x^*) > 0$ for all $x \in \alpha^{-1}(a')$. In this case, because $d(x) \leq d(x^*)$ for all $x \in \mathcal{X}$, Eq. (2.10) shows that in fact $d(x) = d(x^*)$ for all $x \in \alpha^{-1}(a')$.

Now, let x' be arbitrary. Again, by the definition of path-specific fairness, we have that

$$\begin{aligned} d(x') &= \mathbb{E}[D_{\Pi, A, a'} \mid X = x'] \\ &= \sum_{x \in \alpha^{-1}(a')} d(x) \cdot \Pr(X_{\Pi, A, a'} = x \mid X = x') \\ &= \sum_{x \in \alpha^{-1}(a')} d(x^*) \cdot \Pr(X_{\Pi, A, a'} = x \mid X = x^*), \\ &= d(x^*), \end{aligned}$$

where we use in the third equality the fact $d(x) = d(x^*)$ for all $x \in \alpha^{-1}(a')$, and in the final equality the fact that $X_{\Pi, A, a'}$ is supported on $\alpha^{-1}(a')$.

Theorem 2.11 formalizes and extends this argument to more general settings, where $\Pr(X_{\Pi, A, a'} = x \mid X = x^*)$ is not necessarily positive for all $x \in \alpha^{-1}(a')$. The proof of Theorem 2.11 is in Appendix A.7, along with extensions to continuous covariate spaces and a more complete characterization of Π -fair policies for finite $\mathcal{X}$.

Theorem 2.11. *Suppose $\mathcal{X}$ is finite and $\Pr(X = x) > 0$ for all $x \in \mathcal{X}$. Suppose $Z = \zeta(X)$ is a random variable such that:*

1. $Z = Z_{\Pi, A, a'}$ for all $a' \in \mathcal{A}$,
2. $\Pr(X_{\Pi, A, a'} = x' \mid X = x) > 0$ for all $a' \in \mathcal{A}$ such that $\alpha(x') = a'$, $\alpha(x) \neq a'$, and $x, x' \in \mathcal{X}$

such that $\zeta(x) = \zeta(x')$.

Then, for any Π -fair policy d , with $W = X$, there exists a function f such that $d(X) = f(Z)$, i.e., d is constant across individuals having the same value of Z .

The first condition of Theorem 2.11 holds for any reduced set of covariates Z that is not causally affected by changes in A (e.g., Z is not a descendant of A). The second condition requires that among individuals with covariates x , a positive fraction have covariates x' in a counterfactual world in which they belonged to another group a' . Because $\zeta(x)$ is the same in the real and counterfactual worlds—since Z is unaffected by A , by the first condition—we only consider x' such that $\zeta(x') = \zeta(x)$ in the second condition.

In our admissions example, this result shows that, under mild conditions, causally-fair policies require admitting all applicants with equal probability. In particular, suppose that among students with a given test score, a positive fraction achieve any other test score in the counterfactual world in which their race is altered—as, for instance, we might expect if the individual-level causal effects are drawn from an (appropriately discretized) normal distribution. In this case, the empty set of reduced covariates—formally encoded by setting ζ to a constant function—satisfies the conditions of Theorem 2.11. The theorem then implies that under any Π -fair policy, every applicant is admitted with equal probability. (We motivated our admissions example by assuming that only a fraction $b < 1$ of applicants could be admitted; however, Theorem 2.11 holds irrespective of the budget, and, in particular, when $b = 1$, and so we discuss this result together with our others on unconstrained decision making as a natural extension of blinding.)

Even when decisions are not perfectly uniform lotteries, Theorem 2.11 suggests that enforcing Π -fairness can lead to unexpected outcomes. For instance, suppose we modify our admissions example to additionally include age as a covariate that is causally unconnected to race—as some past work has done. In that case, Π -fair policies would admit students based on their age alone, irrespective of test score or race. Although in some cases such restrictive policies might be desirable, this strong structural constraint implied by Π -fairness appears to be a largely unintended consequence of the mathematical formalism.

The conditions of Theorem 2.11 are relatively mild, but do not hold in every setting. Suppose that in our admissions example it were the case that $T_{\Pi,A,a_0} = T_{\Pi,A,a_1} + c$ for some constant

c —that is, suppose the effect of intervening on race is a constant change to an applicant’s test score. Then the second condition of Theorem 2.11 would no longer hold for a constant ζ . Indeed, any multiple-threshold policy in which $t_{a_0} = t_{a_1} + c$ would be Π -fair. In practice, though, such deterministic counterfactuals would seem to be the exception rather than the rule. For example, it seems reasonable to expect that test scores would depend on race in complex ways that induce considerable heterogeneity. Lastly, we note that $W \neq X$ in some variants of path-specific fairness (e.g., Zhang and Bareinboim, 2018; Nabi and Shpitser, 2018), in which case Theorem 2.11 does not apply. Although, in that case, path-specific fairness is still typically incompatible with optimal decision-making, as shown in Theorem 2.17.

2.4 Equitable Decisions in the Presence of Externalities

We have thus far considered cases where there is largely agreement on the utility of different decision policies. In that setting, we showed that maximizing utility is at odds with various mathematical formalizations of fairness. We further argued that these results illustrate weaknesses in the formalizations themselves, since deviating from utility-maximizing policies in that setting can harm both individuals and groups—as seen in our diabetes screening example.

Agreement on the utility, however, is perhaps the exception rather than the rule. One could indeed argue that the value of mathematical formalizations of fairness is their ability to arbitrate between competing definitions of utility. Here we critically examine that perspective. We show, in analog to our previous results, that even when it is unclear how to balance competing priorities, enforcing existing fairness constraints typically leads to worse outcomes on each dimension. For instance, in our running college admissions example, policies constrained to satisfy various fairness constraints will typically require admitting a student body that is both less academically prepared and less diverse, relative to alternative policies that violate these mathematical fairness definitions.

We start, in Section 2.4.1, by examining our college admissions example in detail, illustrating in geometric terms how existing fairness definitions can lead to problematic admissions policies. Then, in Section 2.4.2, we develop our formal theory of equitable decision making in the presence of externalities. The mathematics necessary to establish our key results are significantly deeper than what we have needed thus far, but our high-level message is the same: enforcing several formal

notions of fairness leads to policies that can paradoxically harm the very groups that they were designed to protect.

2.4.1 The Geometry of Fair Decision Making

To build intuition about the limitations of popular definitions of fairness, we return to our running example on college admissions. In that setting, we imagine an admissions committee debating the merits of different admissions policies. In particular, we imagine disagreement within the committee over how best to balance two competing objectives: academic preparedness (operationalized, e.g., in terms of the high school grades and standardized test scores of admitted students) and class diversity (e.g., the number of admitted applicants from marginalized groups).

We assume that our hypothetical committee members all agree that more (total) academic preparedness and more class diversity are better. Thus, in the absence of any resource constraints (with $b = 1$, as is approximated in some online courses), the university could admit all applicants, maximizing both the number of admitted students from marginalized groups and also the total academic preparedness of the admitted class. But given limits on the number of students who can be admitted (i.e., $b < 1$), one must make difficult choices on whom to admit, with reasonable and expected disagreement on how much to trade one dimension for another. The trade-offs in decision making are most acute when the budget $b < 1$, and for this reason we focus here on that case.

In light of these trade-offs, one might turn to the myriad formal fairness criteria we have discussed to ensure admissions decisions are equitable. Many of the fairness definitions we consider make reference to a distinguished outcome Y . In our example, we can imagine this outcome corresponds to college degree attainment, an *ex post* measure of academic preparedness. In the case of causal fairness definitions, we could take $Y(1)$ to mean degree attainment if the student were admitted, and $Y(0)$ to be degree attainment if the student were not admitted, with the understanding that a student who is not admitted could potentially attend and graduate from another university. For example, satisfying counterfactual predictive parity requires that among rejected applicants, the proportion who would have attained a college degree, had they been accepted, is equal across race groups. In these cases, we imagine academic preparedness is some student-level measure that connects observables X —upon which the committee must make their admissions decisions—to

(potential) outcomes Y . For example, the “academic index” $m(x)$ might be a prediction of $Y(1)$ given X based on historical data, or, more generally, could encode committee preferences for both academic preparation and participation in extracurricular activities, among other factors.

The key point of our informal discussion thus far is that we assume committee members would like to enact an admissions policy d that balances two competing objectives. First, they would like a policy that leads to large $m(x)$, i.e., they would like $\mathbb{E}[m(X) \cdot d(X)]$ to be big, where $m(x)$ is some quantity that may, for example, encode academic preparedness and other preferences. Second, the committee would like large diversity, i.e., they would like $\mathbb{E}[\mathbb{1}_{\alpha(X)=a_1} \cdot d(X)]$ to be big, where a_1 corresponds to some target group of interest. All committee members would like more of each dimension, but, given the budget constraint, it is in general impossible to maximize both dimensions simultaneously, leading to the inherent trade-offs we consider in this section.

We now explore the consequences of imposing additional fairness constraints on our college admissions example, as given by the causal DAG in Figure 2.1, via a simulation study of one million hypothetical applicants, for one quarter of whom ($b = \frac{1}{4}$) seats are allocated. In particular, in the hypothetical pool of applicants we consider, applicants in the target race group a_1 have, on average, fewer educational opportunities than those applicants in group a_0 , which leads to lower average academic preparedness, as well as lower average test scores. We define the “academic index” $m(x)$ of applicants to be the estimated probability that an applicant will graduate if admitted, based on their observed test score and race. See Appendix A.3 for additional details, including the specific structural equations we use in the simulation.

Each of the panels in Figure 2.4 illustrates the geometry of fairness constraints for five different formal notions of fairness described in Section 2.2: counterfactual fairness, path-specific fairness, principal fairness, counterfactual equalized odds, and counterfactual predictive parity. The vertical axes of each panel correspond to aggregate academic index and the horizontal axes to the number of admitted applicants from the target group. The purple lines trace out the boundary of the set of feasible policies, with points on or below the curves achievable by policies that adhere to the budget constraint. Policies lying strictly below the purple curves (or, similarly, on the dashed segments of the purple curves) are “Pareto dominated,” meaning that one can find feasible alternatives that are larger on both of the depicted axes (i.e., academic index and diversity). Since we have assumed committee members prefer higher values on each dimension, their effective choice set consists of

those policies on the solid purple segments—the “Pareto frontier.” Committee members may still disagree over which policy on the frontier to adopt. But for any policy not on the frontier, there is a feasible policy above and to the right of it, which is thus preferred by every member of the committee.

Finally, the shaded regions indicate the set of feasible policies constrained to satisfy each of the fairness definitions. (In Appendix A.2, we show that these feasibility regions can be computed by solving a series of linear programs.) In each case, the constrained regions do not intersect the Pareto frontier, and so there is an alternative, unconstrained feasible policy that simultaneously achieves more student-body diversity and an overall higher academic index. For example, in the case of policies satisfying counterfactual or path-specific fairness, shown in the upper left panel, the set of feasible policies lie on a single line segment. That structure follows from Theorem 2.11, since the only policies satisfying either of these notions of fairness in our setting are ones that admit all students with a constant probability, irrespective of their covariates. While not as extreme, the other fairness definitions similarly restrict the space of feasible policies in severe ways, as shown in the remaining panels. These results illustrate that constraining decision-making algorithms to satisfy popular definitions of fairness can have unintended consequences, and may even harm the very groups they were ostensibly designed to help.

Our discussion in this section aimed to highlight the geometry of “fair” decision policies and their consequences in the context of a simple motivating example. We next show that these qualitative findings are guaranteed to hold much more generally.

2.4.2 A Formal Theory of Fairness in the Presence of Externalities

Our simulation above showed that policies satisfying one of the mentioned fairness definitions are suboptimal, in the sense that they constrain one to a portion of the feasible region in which policies could be improved along both dimensions of interest. As was the case in the absence of trade-offs in Section 2.3.2, the phenomenon occurring in our simulation is true much more generally. To understand why, we begin by isolating and formalizing the relevant mathematical properties of our example. To generalize our setting in Section 2.3, we consider arbitrary utility functions of the form $u : \mathcal{X} \rightarrow \mathbb{R}$. As before, for a function u and decision policy d , we write $u(d) = \mathbb{E}[d(X) \cdot u(X)]$ to




Figure 2.4: The geometry of fairness constraints, in an illustrative example of college admissions. Points under the purple curve correspond to all feasible policies—given the budget constraint—whereas the shaded regions correspond to feasible policies that satisfy various formal definitions of fairness. (For path-specific fairness, we set Π equal to the single path $A \rightarrow E \rightarrow T \rightarrow D$ highlighted in Figure 2.1, and set $W = X$.) For each definition, the constrained policies lie strictly under the purple curve, meaning there are alternative, unconstrained, feasible policies that simultaneously achieve greater student-body diversity (indicated by the x -axis) and greater academic preparedness (indicated by the y -axis). The solid segments of the purple lines correspond to policies on the Pareto frontier—for which one cannot simultaneously increase both diversity and academic preparedness. The points labeled “random” in the upper left-hand corner correspond to policies generated by random lotteries in which each individual is admitted with equal probability, in accordance with Theorem 2.11.

denote the expected utility of decision policy $d(x)$ under the utility u . An important constraint on the admissions committee was the fact that their admissions decisions could not, in expectation, exceed the budget.

Definition 2.12. For a budget b , we say a decision policy $d(x)$ is *feasible* if $\mathbb{E}[d(X)] \leq b$.

A key feature of the college admissions example is that despite some level of uncertainty regarding the “true” utility—i.e., exactly how to trade off between its objectives—the committee knows what its objectives are: to increase the academic index and diversity of the incoming class. One way to encode this kind of uncertainty is to consider a set $\mathcal{U}$ consisting of all “reasonable” ways of trading off between the objectives. While the utilities need not be the same, they should be consistent, in the sense that conditional on an applicant’s group membership, all of the utilities should “agree” that a higher academic index is better.

Definition 2.13. We say that a set of utilities $\mathcal{U}$ is *consistent modulo* α if, for any $u, u' \in \mathcal{U}$:

1. For any x , $\text{sign}(u(x)) = \text{sign}(u'(x))$;
2. For any x_1 and x_2 such that $\alpha(x_1) = \alpha(x_2)$, $u(x_1) > u(x_2)$ if and only if $u'(x_1) > u'(x_2)$.

A second relevant feature of the admissions problem is that certain policies were strictly better from the admissions committee’s perspective, despite their uncertainty about the exact form of their utility. The notion that one policy is better than another *regardless* of the exact form of the utility is formalized by Pareto dominance.

Definition 2.14. Suppose $\mathcal{U}$ is a collection of utility functions. A decision policy d is *Pareto dominated* if there exists a feasible alternative d' such that $u(d') \geq u(d)$ for all $u \in \mathcal{U}$, and there exists $u' \in \mathcal{U}$ such that $u'(d') > u'(d)$. A policy d is *strongly Pareto dominated* if there exists a feasible alternative d' such that $u(d') > u(d)$ for all $u \in \mathcal{U}$. A policy d is *Pareto efficient* if it is feasible and not Pareto dominated, and the *Pareto frontier* is the set of Pareto efficient policies.

As discussed above and in Section 2.3.1, in the absence of trade-offs, optimal decision policies take the simple form of threshold policies. The existence of trade-offs broadens the range of forms a Pareto efficient policy can take. Even so, for consistent collections of utilities, the Pareto efficient policies take a closely related form.

Proposition 2.15. *Suppose $\mathcal{U}$ is a set of utilities that is consistent modulo α . Then any Pareto efficient decision policy d is a multiple-threshold policy. That is, for any $u \in \mathcal{U}$, there exist group-specific constants $t_a \geq 0$ such that, a.s.:*

$$d(x) = \begin{cases} 1 & u(x) > t_{\alpha(x)}, \\ 0 & u(x) < t_{\alpha(x)}. \end{cases} \quad (2.11)$$

The proof of Proposition 2.15 is in Appendix A.5.¹⁸

2.4.2.1 Fairness Definitions with Many Constraints

All of the definitions we study in this section prominently feature causal quantities, but the important quality driving our analysis in this section is that each definition imposes many constraints. For instance, counterfactual equalized odds requires that

$$\Pr(D = 1 \mid A = a, Y(1) = y) = \Pr(D = 1 \mid Y(1) = y)$$

for every outcome y .

Theorem 2.17 shows that for *almost every* joint distribution of X , $Y(0)$, and $Y(1)$ such that $u(X)$ has a density, any feasible decision policy satisfying counterfactual equalized odds or conditional principal fairness is Pareto dominated. Similarly, for almost every joint distribution of X and $X_{\Pi, A, a}$, we show that feasible policies satisfying path-specific fairness—including counterfactual fairness—are Pareto dominated. (The analogous statements for counterfactual predictive parity, equalized false positive rates, and demographic parity are not true; we return to this point in Section 2.4.2.2.) That is, we show that, for a typical joint distribution, any feasible policy satisfying the fairness definitions enumerated above cannot have the form of a multiple-threshold policy. To prove this result, we make relatively mild restrictions on the set of distributions and utilities we consider to exclude degenerate cases, as formalized by Definition 2.16.

¹⁸In the statement of the proposition, we do not specify what happens at the thresholds $u(x) = t_{\alpha(x)}$ themselves, as one can typically ignore the exact manner in which decisions are made at the threshold. Specifically, given a multiple-threshold policy d , we can construct a standardized multiple-threshold policy d' that is constant within group at the threshold (i.e., $d'(x) = c_{\alpha(x)}$ when $u(x) = t_{\alpha(x)}$), and for which: (1) $\mathbb{E}[d'(X)|A] = \mathbb{E}[d(X)|A]$; and (2) $u(d') = u(d)$. In our running example, this means we can standardize multiple-threshold policies so that applicants at the threshold are admitted with the same group-specific probability.

Definition 2.16. Let $\mathcal{G}$ be a collection of functions from $\mathcal{Z}$ to $\mathbb{R}^d$ for some set $\mathcal{Z}$. We say that a distribution of Z on $\mathcal{Z}$ is $\mathcal{G}$ -fine if $g(Z)$ has a density for all $g \in \mathcal{G}$.

In particular, $\mathcal{U}$ -finess ensures that the distribution of $u(X)$ has a density. In the absence of $\mathcal{U}$ -finess, corner cases can arise in which an especially large number of policies may be Pareto efficient, in particular when $u(X)$ has large atoms and X can be used to predict the potential outcomes $Y(0)$ and $Y(1)$ even after conditioning on $u(X)$. See Proposition A.52 in Appendix A.3 for details.

Theorem 2.17. *Suppose $\mathcal{U}$ is a set of utilities consistent modulo α . Further suppose that for all $a \in \mathcal{A}$ there exist a $\mathcal{U}$ -fine distribution of X and a utility $u \in \mathcal{U}$ such that $\Pr(u(X) > 0, A = a) > 0$, where $A = \alpha(X)$. Then,*

- *For almost every $\mathcal{U}$ -fine distribution of X and $Y(1)$, any feasible decision policy satisfying counterfactual equalized odds is strongly Pareto dominated.*
- *If $|\text{IMG}(\omega)| < \infty$ and there exists a $\mathcal{U}$ -fine distribution of X such that $\Pr(A = a, W = w) > 0$ for all $a \in \mathcal{A}$ and $w \in \text{IMG}(\omega)$, where $W = \omega(X)$, then, for almost every $\mathcal{U}$ -fine joint distribution of X , $Y(0)$, and $Y(1)$, any feasible decision policy satisfying conditional principal fairness is strongly Pareto dominated.*
- *If $|\text{IMG}(\omega)| < \infty$ and there exists a $\mathcal{U}$ -fine distribution of X such that $\Pr(A = a, W = w_i) > 0$ for all $a \in \mathcal{A}$ and some distinct $w_0, w_1 \in \text{IMG}(\omega)$, then, for almost every $\mathcal{U}^A$ -fine joint distributions of A and the counterfactuals $X_{\Pi, A, a'}$, any feasible decision policy satisfying path-specific fairness is strongly Pareto dominated.¹⁹*

The proof of Theorem 2.17 is given in Appendix A.6. At a high level, the proof proceeds in three steps, which we outline below using the example of counterfactual equalized odds. First, we show that for almost every fixed $\mathcal{U}$ -fine joint distribution μ of X and $Y(1)$ there is at most one policy $d^*(x)$ satisfying counterfactual equalized odds that is not strongly Pareto dominated. To see why, note that for any specific y_0 , since counterfactual equalized odds requires that $D \perp\!\!\!\perp A \mid Y(1) = y_0$, setting the threshold for one group determines the thresholds for all the others; the budget constraint

¹⁹Here, $u^A : (x_a)_{a \in \mathcal{A}} \mapsto (u(x_a))_{a \in \mathcal{A}}$ and $\mathcal{U}^A$ is the set of u^A for $u \in \mathcal{U}$, i.e., component-wise application of u to elements of $\mathcal{X}^A$. In other words, the requirement is that the joint distribution of the $u(X_{\Pi, A, a})$ has a density.

then can be used to fix the threshold for the original group. Second, we construct a “slice” around μ such that for any distribution ν in the slice, $d^*(x)$ is still the only policy that can potentially lie on the Pareto frontier while satisfying counterfactual equalized odds. We create the slice by strategically perturbing μ only where $Y(1) = y_1$, for some $y_1 \neq y_0$. This perturbation moves mass from one side of the thresholds of $d^*(x)$ to the other. Due to inframarginality, this perturbation typically breaks the balance requirement $D \perp\!\!\!\perp A \mid Y(1) = y_1$ for almost every ν in the slice. Finally, we appeal to the notion of prevalence to stitch the slices together, showing that for almost every distribution, any policy satisfying counterfactual equalized odds is strongly Pareto dominated. Analogous versions of this general argument apply to the cases of conditional principal fairness and path-specific fairness.²⁰ We note that the conditions of Theorem 2.17 are sufficient, rather than necessary, meaning that the conclusion of the theorem may—and, indeed, we expect will—hold even in some cases where the conditions are not satisfied. In particular, we note that this proof technique prevents the conditions of Theorem 2.17 from holding when A factors through W and, in particular, when $W = X$. Although, when $X = W$, Theorem 2.11 shows that under slightly different conditions, a much stronger result holds.

To bring our discussion full circle, we now map Theorem 2.17 onto the motivation offered in Section 2.4.1. Recall that the admissions committee knew that given the opportunity, it preferred policies that increased both the overall academic index of its admitted class, and policies that resulted in more students being admitted from the target group. In other words, we imagine that members of the admissions committee have utilities u^* of the form²¹

$$u^*(d) = v \left(\mathbb{E}[m(X) \cdot d(X)], \mathbb{E}[\mathbb{1}_{\alpha(X)=a_1} \cdot d(X)] \right), \quad (2.12)$$

where, as above, $m(x)$ denotes the academic index of an applicant with covariates $X = x$, and v increases in both coordinates. Corollary 2.18 establishes the inherent incompatibility of such preferences with the formal fairness criteria we have been considering.

Corollary 2.18. *Consider a utility of the form given in Eq. (2.12), where v is monotonically*

²⁰This argument does not depend in an essential way on the definitions being causal. In Corollary A.50 in Appendix A.6, we show an analogous result for the non-counterfactual version of equalized odds.

²¹Strictly speaking, we are saying that members of the admissions committee, rather than having an aggregate utility—which, as we have considered so far, has the form $\mathbb{E}[u(X) \cdot d(X)]$ —has a utility on aggregate *outcomes*.

increasing in both coordinates and $m(x) \geq 0$. Then, under the same hypotheses as in Theorem 2.17,²² for almost every joint distribution, no utility-maximizing decision-policy satisfies counterfactual equalized odds, conditional principal fairness, or path-specific fairness.

Lastly, while, in general, one’s decision policy can depend only on the covariates known at the time of the decision, in some cases, the restriction that $u(x)$ be a function of $x \in \mathcal{X}$ alone may be too restrictive; the connection between an individual having covariates $X = x$ and our utility may depend also on the relationship between X and Y . For instance, in the admissions example, the admissions committee may value high test scores and extracurriculars not, e.g., as *per se* measures of academic merit, but rather instrumentally insofar as they are connected to whether an applicant will eventually graduate. However, allowing u to depend on both x and y greatly complicates the underlying geometry of the problem. Proving Theorem 2.17 in this more general setting remains an open problem. However, intuition from finite-dimensions—where more powerful measure-theoretic tools are available—suggests that the result remains true in the more general setting. For example, Proposition 2.19 presents a version of this result over a natural, finite-dimensional family of distributions.

Proposition 2.19. *Suppose $\mathcal{A} = \{a_0, a_1\}$, and consider the family $\mathcal{U}$ of utility functions of the form*

$$u(x) = r(x) + \lambda \cdot \mathbb{1}_{\alpha(x)=a_1},$$

indexed by $\lambda \geq 0$, where $r(x) = \mathbb{E}[Y(1) \mid X = x]$. For almost every $(\alpha_0, \beta_0, \alpha_1, \beta_1) \in \mathbb{R}_{>0}^4$, if the conditional distributions of $r(X)$ given A are beta distributed with

$$r(X) \mid A = a_i \sim \text{BETA}(\alpha_i, \beta_i),$$

then any policy satisfying counterfactual equalized odds is strongly Pareto dominated.

2.4.2.2 Fairness Definitions with Few Constraints

We conclude this analysis by considering equalized false positive rates, demographic parity, and counterfactual predictive parity. These fairness notions are less demanding than the notions considered above, in that they introduce only “one” additional constraint, e.g., that $Y(1) \perp\!\!\!\perp A \mid D =$

²²The full statement is given in Appendix A.6.7.

0, in the case of counterfactual predictive parity. Since the budget introduces a second constraint, and the form of a multiple-threshold policy allows for a degree of freedom in each group, the number of constraints and the number of degrees of freedom are equal—as opposed to the causal fairness definitions covered by Theorem 2.17, in which the constraints outnumber the degrees of freedom. As such, it is possible in some instances to have a policy on the Pareto frontier that satisfies these conditions; though see Section 2.5.3 for discussion about why such policies are still often at odds with broader goals.

However, it is not *always* possible to find a point on the Pareto frontier satisfying these definitions. In Proposition 2.20, we show that counterfactual predictive parity cannot lie on the Pareto frontier in some common cases, including our example of college admissions. In that setting, when the target group has lower average graduation rates—a pattern that often motivates efforts to actively increase diversity—decision policies constrained to satisfy counterfactual predictive parity are Pareto dominated. The proof of the proposition is in Appendix A.8.3.

Proposition 2.20. *Suppose $\mathcal{A} = \{a_0, a_1\}$, and consider the family $\mathcal{U}$ of utility functions of the form*

$$u(x) = r(x) + \lambda \cdot \mathbb{1}_{\alpha(x)=a_1},$$

indexed by $\lambda \geq 0$, where $r(x) = \mathbb{E}[Y(1) \mid X = x]$. Suppose the conditional distributions of $r(X)$ given A are beta distributed, i.e.,

$$r(X) \mid A = a \sim \text{BETA}(\mu_a, v),$$

with $\mu_{a_0} > \mu_{a_1}$ and $v > 0$.²³ Then any policy satisfying counterfactual predictive parity is strongly Pareto dominated.

2.5 A Path Forward

We have thus far worked to clarify some of the statistical limitations of existing mathematical definitions of fairness. We have argued that in many cases of interest, these definitions can ultimately do more harm than good, hurting even those individuals that these notions of fairness were ostensibly

²³Here we parameterize the beta distribution in terms of its mean μ and sample size v . In terms of the common, alternative α - β parameterization, $\mu = \alpha/(\alpha + \beta)$ and $v = \alpha + \beta$.

designed to help.

We end on a more optimistic note, charting out a potential path toward designing more equitable algorithms. To do so, we start, in Section 2.5.1 by reviewing conceptions of discrimination in law and economics, and, in particular, we contrast process-oriented and outcome-oriented notions of fairness. Whereas the computer science literature is dominated by process-oriented, *deontological* definitions of fairness, we see more promise in adopting an outcome-oriented, *consequentialist* approach represented by the utilitarian analysis we have described above. In Section 2.5.2, we enumerate and discuss four issues that we feel are critical in developing equitable algorithms: (1) balancing inherent trade-offs in decision problems; (2) assessing calibration; (3) selecting the inputs and targets of prediction; and (4) designing data collection strategies. Finally, in Section 2.5.3, we illustrate how to grapple with these considerations in a case study of complex medical care, motivated by work from Obermeyer et al. (2019).

2.5.1 Competing Notions of Ethical Decision Making: Process vs. Outcomes

There are many distinct but related understandings of ethical decision making in law, economics, philosophy, and beyond. One key dimension on which we organize these notions is the extent to which they consider the *process* through which decisions are made versus the *outcomes* that those decisions render.

The dominant legal doctrine of discrimination in the United States treats explicit race- and gender-based decisions with heightened scrutiny. The Equal Protection Clause of the U.S. Constitution’s Fourteenth Amendment restricts government agencies from adopting policies that explicitly reference legally protected categories, and myriad federal and state *disparate treatment* statutes similarly constrain a variety of private actors. Conversely, policies that do not explicitly consider legally protected traits—or obvious proxies—are generally deemed not to violate disparate treatment principles. Formally, it is lawful to use legally protected attributes in a limited way to further a compelling government interest, but, in practice, such exceptions are few and far between. Until recently, the prime example of a race-conscious policy passing legal muster was affirmative action in college admissions (*Fisher v. University of Texas*, 2016). However, in 2023, the U.S. Supreme Court barred the explicit consideration of race in admissions decisions (*SFFA v. Harvard*, 2023).

Disparate treatment doctrine has evolved over time, and reflects ongoing debates about the role of *classification* (use of protected traits, a process-oriented, deontological notion) versus *subordination* (subjugation of disadvantaged groups, an outcome-oriented notion) in discrimination cases (Fiss, 1976). Some legal scholars have argued that courts, even when formally applying anti-classification criteria, are often sympathetic to the potential effects of judgments on social stratification, indicating tacit concern for anti-subordination (Balkin and Siegel, 2003; Colker, 1986; Siegel, 2003). Others, though, have noted that such judicial support for anti-subordination appears to be waning (Nurse, 2014). At a high level, we thus view modern disparate treatment law as primarily interested in process over outcomes, though these debates illustrate that the two concepts cannot be perfectly separated.

In contrast to process-oriented disparate treatment principles, the economics literature distinguishes between two outcome-focused, consequentialist rationales for explicitly considering race, gender, and other protected traits: taste-based and statistical. With *taste-based discrimination* (Becker, 1957), decision makers act as if they have a preference or “taste” for bias, sacrificing profit to avoid certain transactions. This includes, for example, an employer who forfeits financial gain by failing to hire exceptionally qualified minority applicants. But, in contrast to legal reasoning, the economic argument against taste-based discrimination is not that decisions are based on race *per se*, but rather because consideration of race leads to worse outcomes: a loss of profit. With *statistical discrimination* (Arrow, 1973; Phelps, 1972), decision makers explicitly consider protected attributes in order to optimally achieve some non-prejudicial goal. For example, profit-maximizing auto insurers may charge a premium to male drivers to account for gender differences in accident rates. Despite their differing outcome-based justifications, both taste-based and statistical discrimination are often considered legally problematic as they explicitly consider race, in violation of disparate treatment laws.²⁴

As the above insurance example and our running diabetes example illustrate, one might consider it acceptable to base decisions in part on legally protected traits when doing so leads to good outcomes. Conversely, whereas process-oriented disparate treatment principles generally deem race-blind policies acceptable, one might declare such blind policies problematic if they lead to

²⁴In the case of auto insurance specifically, some states (including California, Hawaii, Massachusetts, Maine, Michigan, North Carolina, and Pennsylvania)—though not all—have barred the use of gender in pricing policies.

bad outcomes. Indeed, under the statutory *disparate impact* standard, a practice may be deemed discriminatory if it has an unjustified adverse effect on legally protected groups, even in the absence of explicit categorization (Barocas and Selbst, 2016).²⁵ The disparate impact doctrine was formalized in *Griggs v. Duke Power Co.*, a 1971 U.S. Supreme Court case. In 1955, the Duke Power Company mandated that employees have a high school diploma to be considered for promotion, which, in practice, severely limited the eligibility of Black employees. The Court found that this facially race-neutral requirement had little relation to job performance, and accordingly deemed it to have an unjustified—and illegal—disparate impact. The Court noted that the employer’s motivation for instituting the policy was irrelevant to its decision; even if enacted without discriminatory purpose, the policy was deemed discriminatory in its effects and hence illegal. However, disparate impact law does not prohibit *all* group differences produced by a policy—the law only prohibits *unjustified* disparities. For example, if, hypothetically, the high-school diploma requirement in *Griggs* were shown to be necessary for job success, the resulting disparities would be legal.

On the spectrum from process- to outcome-based understandings of discrimination, we view the formal, axiomatic fairness definitions described in Section 2.2 as reflecting a largely process-based orientation. Blinding and its more stringent causal variants—counterfactual fairness and path-specific fairness—can be viewed as descendants of disparate treatment considerations, as they seek to remove the effects of race and other protected attributes on decisions. The remaining definitions—for example, those that aim to equalize error rates across groups—do explicitly reference an outcome Y , but they do so in a way that seems largely disconnected from the consequences one might naturally consider. As we have argued, whether error rates are equal across groups has more to do with the structure of group-specific risk distributions than with whether decisions lead to good or bad outcomes for group members. For example, in our college admissions example, enforcing various formal notions of fairness would, in theory, typically lead to student bodies that are both less diverse

²⁵The legal doctrine of disparate impact stems largely from federal statutes, not constitutional law, and applies only in certain contexts, such as employment (via Title VII of the 1964 Civil Rights Act) and housing (via the Fair Housing Act of 1968). Apart from federal statutes, some states have passed more expansive disparate impact laws, including Illinois and California. The distinction between statutory and constitutional rules is particularly relevant here, as there is debate among scholars over whether disparate impact laws violate the Equal Protection Clause and are thus unconstitutional (Primus, 2003). There is also debate over whether disparate impact law is motivated primarily by an interest in banning bad *outcomes*, or seeks to provide an alternative pathway for ferreting out bad *intent*, when actors may mask animus with race-neutral policies (Watson v. Fort Worth, 1988). But, regardless of its underlying justification, disparate impact law is formally focused on outcomes, not intent or classification, and so we view it as an outcomes-focused principle.

and less academically prepared than those resulting from feasible alternatives not constrained to satisfy these notions.

One might, on principle, favor certain process-based understandings of discrimination over outcome-based notions. One might even adopt a meta-consequentialist position, and argue that procedural considerations (e.g., ensuring medical decisions are blind to race) engender trust and in turn bring about better downstream outcomes. In many cases, though, the ethical underpinnings of popular mathematical definitions of fairness have not been clearly articulated. Absent such justification, we advocate for an approach that more directly engages with the real-world costs and benefits of different decision policies, a perspective that we outline in more detail in the remaining sections.

2.5.2 Designing Equitable Algorithms

A key advantage of the dominant axiomatic approach to algorithmic fairness is that it can be readily applied across contexts, with little domain-specific knowledge. One can build automated tests to check whether any predictive algorithm satisfies various formal fairness desiderata, and even automatically modify algorithms to ensure that they do satisfy a specific fairness criterion (e.g., Cotter et al., 2019; Weerts et al., 2023). But, as we have argued, this approach is often at odds with improving well-being, including for disadvantaged groups. A particularly pernicious risk of automated, axiomatic approaches is that they can make invisible the cost to well-being: automatically constraining algorithms to be “fair” can lead one to overlook unconstrained alternatives that are clearly preferable. We have instead called for a more careful analysis of the consequences, good and bad, of different decision policies, selecting the appropriate course of action based on the specific context. This is admittedly hard to do—and does not easily scale—but there are general principles that we believe are helpful to keep in mind when navigating this terrain. Below we enumerate and discuss four of them.

2.5.2.1 Contending with Inherent Trade-Offs

There are inherent trade-offs in many important decision problems. For instance, in our college admissions example, one must balance academic preparedness with student-body diversity. Although

one cannot generally circumvent these trade-offs, we believe it is useful to explicitly enumerate the primary dimensions of interest and to acknowledge the trade-offs between them. In some cases, like our stylized admissions example, one might be able to explicitly calculate the Pareto frontier shown in Figure 2.4, in which case it often makes sense to focus on those policies lying on the frontier. In many cases, it won’t be possible to compute the frontier. Still, by listing and discussing trade-offs, even informally, one can reduce the risk of adopting clearly problematic policies, like those that typically result from uncritically constraining decisions to satisfy formal fairness criteria.

In this sense, designing equitable algorithms is akin to designing equitable policy writ large. One might accordingly adapt democratic mechanisms used to draft and enact legislation to algorithm design. For example, adopting such a policy-oriented perspective, Chohlas-Wood et al. (2023a) and Koencke et al. (2023) surveyed a diverse sample of Americans to elicit preferences on how best to balance competing objectives in programs that algorithmically allocate government benefits.

2.5.2.2 Assessing Calibration

When designing or auditing a risk assessment algorithm, it is important to check whether predictions are *calibrated*, meaning that risk scores correspond to the same observed level of risk across groups. In general, the relationship between predictors and outcomes may plausibly differ across groups, leading to miscalibrated risk estimates—what Ayres (2002) calls the problem of *subgroup validity*. Figure 2.3 shows instances of risk scores for diabetes and recidivism that are miscalibrated across race and gender, respectively. For example, a nominal 1.5% diabetes risk (based on age and BMI) corresponds to an actual, observed diabetes rate of approximately 1% among White patients and 3% among Asian patients. Similarly, among individuals receiving a COMPAS risk score of 7 (based on criminal history and related factors), about 55% of women recidivate, compared to 65% of men. These miscalibrated risk scores can result in inequitable decisions. For instance, a policy to screen patients with a nominal diabetes risk of 1.5% or above, in line with existing medical recommendations (Aggarwal et al., 2022), would overscreen White patients and underscreen Asian patients, harming individuals in both groups.

Calibration can often be visually assessed by plotting predicted risk against average outcomes, as in Figure 2.6 (cf. Arrieta-Ibarra et al., 2022). For a simple, more quantitative, measure, we

recommend regressing observed outcomes against risk estimates and group membership. A coefficient of approximately zero on group membership suggests risk estimates correspond to similar average outcomes across groups, with deviations from zero indicating the degree of miscalibration.

In practice, miscalibration can often be rectified by training group-specific risk models, or, roughly equivalently, including group membership in a single risk model fit across groups. For example, diabetes risk models that include race, and recidivism risk models that include gender, are approximately calibrated. The relatively new literature on multicalibration has introduced new computational techniques to ensure predictions are simultaneously calibrated across many different subgroups (Hébert-Johnson et al., 2018). Of course, including protected traits in risk models raises additional legal and ethical challenges. In some cases, it may be possible to reduce or eliminate miscalibration by incorporating additional, non-protected covariates. Regardless, we believe it is important to check the calibration of risk scores to make informed decisions about if and how to address any observed disparities.

Calibration is an important necessary condition to ensure risk estimates correspond to actually observable levels of risk across groups. But it is not sufficient. Indeed, even calibrated risk scores can encode and reinforce deeply discriminatory policies. To see this, imagine a bank that wants to discriminate against Black applicants. Further suppose that: (1) within ZIP code, White and Black applicants have similar default rates; and (2) Black applicants live in ZIP codes with relatively high default rates. Then the bank can surreptitiously discriminate against Black borrowers by basing estimates of default risk only on an applicant’s ZIP code, ignoring all other relevant information. Such scores would be calibrated (White and Black applicants with the same score would default equally often), and the bank could use these scores to justify denying loans to nearly all Black applicants. The bank, however, would be sacrificing profit by refusing loans to creditworthy Black applicants,²⁶ and thus engaging in taste-based discrimination. This discriminatory lending strategy is indeed closely related to the historical (and illegal) practice of redlining, and illustrates the limitations of calibration as a measure of equity.

Figure 2.5 shows another example of calibrated scores masking disparities. In the left-hand panel, we plot in green the distribution of diabetes risk for Asian patients, as estimated from

²⁶These applicants are creditworthy in the sense that they would have been issued a loan had the bank used all the information it had available to determine their risk.




Figure 2.5: Calibration is insufficient to prevent discrimination. *Left:* The distribution in green shows diabetes risk for Asian patients based on accurately collected age and BMI, and the distribution in purple shows estimates when the risk model is trained on noisy inputs. Estimates under the noisy model concentrate around the mean (dashed vertical line), pushing more Asian patients above the screening threshold (solid vertical line). *Right:* A calibration plot comparing noisy risk estimates for Asian patients and accurate risk estimates for White patients. The calibrated risk scores can mask both intentional discrimination and inadvertent errors.

age, BMI, and race. In purple, we plot the distribution of estimated risk when age and BMI are imperfectly measured for Asian patients in the training data. Training the risk model on noisy features pushes risk estimates toward the mean. As a result, based on the noisy risk model, 97% of Asian patients are above the 1.5% screening threshold, compared to 81% of Asian patients under the more accurately estimated model—leading to more medically unnecessary screening under the noisy model. Importantly, however, the noisy model is still calibrated, as shown in the right-hand panel of Figure 2.5, where we compare risk scores for Asian patients estimated from the noisy predictors and risk scores for White patients estimated from the accurately estimated information. In theory, a malicious algorithm designer could generate such calibrated but inaccurate scores to intentionally harm Asian patients. In practice, this pattern could equally arise from negligence rather than malice. These examples illustrate the importance of considering all available data when constructing statistical risk estimates; assessments that either intentionally or inadvertently ignore predictive information may facilitate discriminatory decisions while satisfying calibration—though, as we discuss below, even this intuitive heuristic of “use all the data” has its limitations.

2.5.2.3 Selecting the Target of Prediction

In constructing algorithmic risk scores, a key ingredient is the target of prediction. In practice, though, there is often a mismatch between our true outcome of interest and the available data—an occurrence we call *label bias* (Zanger-Tishler et al., 2023). As with the other issues we discuss, there is typically no perfect solution to this problem, but there are ways to mitigate it.

For example, in pretrial risk assessment, we would often like to estimate the likelihood a defendant would commit a crime if released. But there are two key difficulties with this goal. First, though we might want to measure crime conducted by defendants awaiting trial, we typically only observe crime that results in a conviction or an arrest. These observable outcomes, however, are imperfect proxies for the underlying criminal act. Further, heavier policing in communities of color might lead to Black and Hispanic defendants being arrested, and later convicted, more often than White defendants who commit the same offense (Lum and Isaac, 2016). Poor outcome data might thus cause one to systematically underestimate the risk posed by White defendants. The second, related, issue is that our target of interest is a *counterfactual* outcome; it corresponds to what would have happened had a defendant been released. In reality, we only observe what actually happened conditional on the judge’s actual detention decision.

One way to reduce label bias in this case is to adjust the target of interest. For example, criminologists have found that arrests for violent crime—as opposed to drug crime—may suffer from less racial bias.²⁷ In particular, Skeem and Lowenkamp (2016) note that the racial distribution of individuals arrested for violent offenses is in line with the racial distribution of offenders inferred from victim reports, and is also in line with self-reported offending data. In other cases, like lending, where one may seek to estimate default rates, the measured outcome (e.g., failure to pay) corresponds more closely to the event of interest. The problem of estimating counterfactuals can likewise be partially addressed in some applications. In the pretrial setting, Angwin et al. (2016) measure

²⁷D’Alessio and Stolzenberg (2003) find evidence that White offenders are even somewhat more likely than Black offenders to be arrested for certain categories of crime, including robbery, simple assault, and aggravated assault. Measurements of minor criminal activity, like drug offenses, are more problematic. For example, there is evidence that drug arrests in the United States are biased against Black and Hispanic individuals, with racial minorities who commit drug crimes substantially more likely to be arrested than White individuals who commit the same offenses (Ramchand et al., 2006). Although this pattern is well known, many existing risk assessment tools still consider arrests or convictions for *any* new criminal activity—including drug crimes—which may lead to biased estimates. As another example of label bias, auto insurance rates are determined in part by a driver’s record of receiving speeding tickets, but disparities in police enforcement mean that tickets are biased proxies of dangerous driving behavior (Cai et al., 2022b).

recidivism rates in the first two-year period during which a defendant is not incarcerated; this is not identical to the desired counterfactual outcome—since the initial detention may be criminogenic, for example—but it seems like a reasonable estimation strategy. Further, unaided human decisions often exhibit considerable randomness, a fact that can be exploited to facilitate statistical estimation of counterfactual outcomes (Kleinberg et al., 2017a; Jung et al., 2020). More generally, a spate of recent work at the intersection of machine learning and causal inference (Jung et al., 2020c; Mullainathan and Spiess, 2017; Hill, 2011) offers hope for more gains in counterfactual estimation.

2.5.2.4 Collecting Training Data

A final issue we discuss is collecting suitable training data for risk assessment algorithms to mitigate the effects of sample bias. Ideally, one would train algorithms on data sets that are broadly representative of the populations on which they are ultimately applied—though there are subtleties to this heuristic that we describe below. While often challenging in practice, failure to train on representative data can lead to unintended, and potentially discriminatory, consequences. For example, Buolamwini and Gebru (2018) found that commercial facial analysis tools struggle to correctly classify the gender of dark-skinned individuals—and of dark-skinned women in particular—a disparity likely attributable to the relative dearth of dark-skinned faces in facial analysis data sets. Similarly, Koenecke et al. (2020) found that several popular automated speech recognition systems were significantly worse at transcribing Black speakers than White speakers, likely due to insufficient data from speakers of African American Vernacular English (AAVE), a variety of English spoken by many Black Americans.

The problems of non-representative data can be even more acute in the case of risk assessment algorithms, especially when the target of interest is a causal quantity. For instance, in our running college admissions example, we seek to estimate how a student would (counterfactually) perform if admitted. Historical data are typically the result of past, potentially biased, decisions, and so may not fully generalize. Imagine, for example, that predictions of college performance are informed by where an applicant went to high school, and that, historically, only applicants from certain high schools were accepted—and we consequently only see outcomes for students from those high schools. Then we would expect less accurate predictions for students from those absent high schools. In

general, regression to the mean could attenuate estimates for high achieving students who differ from those previously accepted, potentially reinforcing existing admissions practices.

As a general heuristic, we believe it is advisable to train models on representative data, but, as Cai et al. (2022a) note, the optimal sampling strategy depends on the statistical structure of the problem and the group-specific costs of collecting data. Interestingly, the value of representative data collection strategies depends in part on the degree to which race, gender, and other protected attributes are predictive. In theory, if protected attributes are not predictive, one could build an accurate risk model using only examples from one particular group (e.g., White men). Given enough examples of White men, the model would learn the relationship between features and risk, which by our assumption would generalize to the entire population. This phenomenon highlights a tension in informal discussions of fairness, with some advocating both for representative training data and for the exclusion of protected attributes. However, representative data are often most important precisely when protected attributes add information, in which case their use is arguably more justified. Even if protected attributes are not predictive, representative data can still help in two additional ways. First, a representative sample ensures that the full support of features is present at training time, as it is possible that the distribution of features varies across groups, even if the connection between features and outcomes does not. We note, though, that one might have adequate support even without a representative sample in many real-world settings, particularly when models are trained on large data sets and the feature space is relatively low dimensional. Second, a representative sample can help with model validation, allowing one to assess the potential effects of group imbalance on model fit. In particular, without a representative sample, it can be difficult to determine whether a model trained on a single group generalizes to the entire population.

In many settings, one may be able to gather better training data with greater investment of time and money. For example, in our diabetes example one could aim to collect more complete medical records, a process that may be both costly and logistically difficult. In theory, this additional information may lead to welfare gains, and policymakers must accordingly evaluate the relative costs and benefits to all groups of exerting this extra effort when designing algorithms. Fortunately, in practice, there are often diminishing returns to information, with a relatively short list of key features providing most of the predictive power (Jung et al., 2020), at least partially mitigating this concern. Carefully documenting data provenance and content—including what features the

data contain, how observations were selected into the sample, potential missingness or measurement error, and the representation of protected classes in the data—with the use of datasheets or other structured tools can both aid in determining the extent of possible issues of non-representative data and in alerting future analysts to potential issues (Arnold et al., 2019; Chmielinski et al., 2022; Gebru et al., 2021; Holland et al., 2020; Stoyanovich and Howe, 2019).

As with the other issues we have discussed, there is no universal solution to data collection. It might, for example, simply be prohibitive in the short run to train models on the data set one would ideally like to use. Nevertheless, as in all situations, one must carefully weigh the potential costs and benefits of adopting a necessarily imperfect risk assessment algorithm relative to the other possible options. In particular, even an imperfect algorithm may in some circumstances be better than leaving decisions to similarly imperfect humans who have their own biases.

2.5.3 Case Study: An Algorithm to Allocate Limited Medical Resources

We conclude by illustrating the principles discussed above with a real-world example: an algorithm for referring patients into a “high-risk care management” program, previously considered by Obermeyer et al. (2019). The care management program more effectively aids patients with complex medical needs, in principle both improving outcomes for patients and reducing costs to the medical system for patients who are enrolled. But the program has limited capacity, which we formalize by assuming that only 2% of patients can be enrolled (i.e., we set $b = \frac{1}{50}$). For our analysis, we use the data released by Obermeyer et al.,²⁸ which contain demographic variables, cost information, comorbidities, biomarker and medication details, and health outcomes for a population of approximately 43,000 White and 5,600 Black primary care patients at an academic hospital from 2013–2015.

As a first step toward identifying patients to enroll in the program, one could train a model predicting the healthcare resources a patient is likely to require over the next year. Sicker patients require more care and, consequently, incur greater healthcare costs. Thus, our initial approach is to predict how likely a patient is to be “high cost”—which we operationalize as being in the top decile of healthcare expenditures in a given year—based on the available information. (One of the

²⁸Obermeyer et al. released a synthetic data set closely mirroring the real data set, available at: <https://gitlab.com/labsysmed/dissecting-bias>. In contrast to $b = 2\%$, which we adopt to better illustrate some of the statistical phenomena we discuss in this section, Obermeyer et al. use a budget of $b = 3\%$.




Figure 2.6: Calibration of a race-blind model predicting whether a patient is “high-cost”. The lack of a gap between the proportion of patients who are actually high-cost and the proportion predicted to be high-cost for both groups indicates that the model is well calibrated across race groups.

main contributions of Obermeyer et al. 2019 is highlighting that healthcare costs is a problematic outcome due to label bias, a point we return to shortly.)

As discussed above, it is useful to assess the calibration of risk assessment algorithms across groups. In particular, while calibration across race is largely guaranteed if race is included as a predictor in the statistical model, race-blind models are often preferred, particularly in healthcare, in part to avoid perceptions of bias. As a result, we assess the calibration of a race-blind model trained on all available information except for race, shown in Figure 2.6. Unlike in the diabetes screening example considered in Section 2.3.2, the race-blind healthcare cost predictions are calibrated across race groups, meaning that risk estimates largely match observed costs across groups. It accordingly appears that race provides little marginal predictive power in this example, assuaging potential concerns with its omission.

We turn next to assessing the effects of applying formal fairness criteria to our enrollment decisions. Often, in healthcare contexts, a false negative—e.g., failing to screen for a disease when it is present—is more consequential than a false positive—e.g., screening for a disease when it is not present. For this reason, one might seek to ensure enrollment decisions are fair by mandating false negative rates be equal across race groups (e.g., Seyyed-Kalantari et al., 2021),²⁹ i.e., requiring

²⁹The authors describe their metric of concern as the “false positive rate,” where the positive prediction is understood as a “no finding” label, e.g., not having a disease. In our example, we follow the more common convention that a




Figure 2.7: Enforcing formal fairness criteria can harm marginalized groups. Feasible regions for admissions policies to a high-risk care management program, where the dashed line indicates the number of Black patients admitted by the policy admitting the maximal number of high-cost patients. *Left:* The Pareto frontier for all patients in the population. Because more Black patients incur high medical costs in the population as a whole, equalizing false negative rates (FNR)—as well as false positive rates (FPR), or enforcing demographic parity (DP)—results in fewer Black patients being admitted than under the policy that maximizes the total number of high-cost patients admitted. (Equalized false negative rates and demographic parity are achieved at the same point.) *Right:* The Pareto frontier for the subpopulation of women between the ages of 25 and 34. Within this subpopulation, Black patients incur lower medical costs, and so equalizing false positive rates, false negative rates, or achieving demographic parity all result in more Black patients being admitted to the high-risk management program than the policy that admits the maximum number of high-cost patients.

that $A \perp\!\!\!\perp D \mid Y = 1$. In our example, equalizing false negative rates means that among patients who ultimately incur high medical costs ($Y = 1$), the same proportion ($\mathbb{E}[D]$) are referred into the program across race groups (A).

In our setting, approximately 1,000 patients can be referred into the care management program—2% of the roughly 50,000 patients in our data set. When we equalize false-negative rates, 747 of the enrolled patients ultimately incur high costs, and 113 enrolled patients are Black.³⁰ However, an unconstrained decision rule (i.e., one that enrolls the patients most likely to incur high costs) enrolls both more high-cost patients (758) and more Black patients (205). In this example, we end

“positive” classification is assigned to the rare event (i.e., being high cost), and so we call this metric the “false negative rate.”

³⁰Following Corbett-Davies et al. (2017), we equalize false-negative rates in a manner that maximizes the number of enrolled patients who ultimately incur high costs.

up providing worse care to Black patients when we constrain our algorithm to satisfy the formal, mathematical fairness criterion.

Instead of applying such mathematical fairness criteria, we advocate for directly weighing the costs and benefits of different decision policies. In Figure 2.7a, we show the Pareto frontier for our example, tracing out policies that optimally trade off the demographic composition of the enrolled population with the number of enrolled patients who in reality incur high costs. The green point corresponds to equalizing false negative rates, and is to the left of the dashed vertical line that corresponds to the unconstrained decision rule—visually illustrating how constraining our algorithm leads to fewer resources for Black patients. Also shown on the plot are points corresponding to demographic parity and equal false positive rates, both of which likewise lead to fewer resources for Black patients.³¹

The result in Figure 2.7a stems from the false negative rate for Black patients being lower than the false negative rate for White patients in the unconstrained algorithm—a pattern we expect since Black patients are more likely than White patients to incur high medical costs in our data. Equalizing false negative rates thus means raising the enrollment bar for Black patients and lowering the bar for White patients. In light of this example, one might argue for applying formal fairness criteria only when error rates for racial minorities are higher than for White individuals. Figure 2.7b repeats our analysis above for the subset of women between the ages of 25 and 34, a subpopulation in which Black patients have higher error rates than White patients. In this case, equalizing false positive rates or false negative rates, or enforcing demographic parity all result in more Black patients being admitted into the program. It is, however, unclear why one should adopt those particular error-rate equalizing policies over any other. The policies on the Pareto frontier (i.e., on the curve to the right of the dashed line) are all arguably reasonable to consider. It is admittedly difficult to determine which of these policies to adopt, but we believe it is best to confront this challenge head on, recognizing the hard trade-offs inherent to the problem.

We conclude our case study by considering label bias, the primary concern identified by Obermeyer et al. (2019) in this context. As those authors noted, medical cost is a poor proxy for medical need,

³¹As discussed in Section 2.4, with the exception of counterfactual predictive parity, all of the remaining fairness definitions given in Section 2.2 are known *a priori* to restrict one to enrollment policies that will lower *both* the number of Black patients enrolled as well as the number of truly high-cost patients enrolled.




Figure 2.8: The target of prediction impacts equity. *Left*: The probability of having complex medical needs (i.e., at least five active chronic conditions) for Black and White patients as a function of their estimated likelihood of incurring high medical costs, reproducing an analysis by Obermeyer et al. (2019). The large gap across groups indicates that Black patients have greater medical need than White patients with similar anticipated healthcare costs. *Right*: Distribution of the estimated change in the probability that an individual will have complex medical needs after enrolling in the care management program, showing that the extent to which enrollment reduces complex medical needs varies considerably across individuals.

and so allocating healthcare resources to minimize anticipated costs can lead to severe disparities. Replicating an analysis by Obermeyer et al., Figure 2.8a shows that among patients with similar likelihood of incurring high-cost medical care, Black patients are considerably more likely than White patients to have complex medical needs, operationalized as having five or more active chronic conditions. This gap is likely a consequence of worse access to healthcare among Black patients, due to a mix of socioeconomic factors and discrimination. To the extent that care management programs aim to aid the sickest patients—as opposed to simply reducing costs—targeting resources based on anticipated costs can lead to inefficient and inequitable outcomes. Obermeyer et al. accordingly suggest switching the target of prediction from medical costs to health status.

It is possible to achieve further gains by recognizing resource allocation as an inherently causal problem. In particular, one may seek to enroll patients in the program in a manner that maximizes $\mathbb{E}[Y(D)]$, where $Y(D)$ is the potential health outcome under the enrollment decision. To do so, we could prioritize patients by their estimated treatment effect $\hat{Y}_i(1) - \hat{Y}_i(0)$, rather than an estimate

$\hat{Y}_i$ of their future health status that ignores the causal effect of the program.³² A proper causal analysis is, in general, a complex topic requiring careful treatment beyond the scope of this article. Nonetheless, Figure 2.8b shows the distribution of $\hat{Y}_i(1) - \hat{Y}_i(0)$, as estimated with a simple regression model. The plot suggests that there is considerable predictable heterogeneity in the extent to which enrollment in the care management program causally improves health. In particular, we find that the estimated treatment effect is only weakly correlated with the number of chronic conditions a patient currently exhibits ($r = 0.05$). Consequently, directly targeting resources to those most likely to benefit could yield large health improvements. Once the types of label bias we have discussed above have been identified, it may be possible to re-train predictive models to better align decision-making algorithms with policy goals.

2.6 Conclusion

From medicine to criminal justice, practitioners are increasingly turning to statistical risk assessments to help guide and improve human decisions. Algorithms can avoid many of the implicit and explicit biases of human decisions makers, but they can also exacerbate historical inequities if not developed with care. Policymakers, in response, have rightly demanded that these high-stakes decision systems be designed and audited to ensure outcomes are equitable. The research community has responded to the challenge, coalescing around several formal mathematical definitions of fairness. However, as we have aimed to articulate, these popular measures of fairness suffer from significant statistical limitations. Indeed, adopting these measures as algorithmic design principles can often harm the groups that these measures were designed to protect.

In contrast to the dominant axiomatic approach to algorithmic fairness, we advocate for a more consequentialist orientation (Cai et al., 2022a; Chohlas-Wood et al., 2023a; Liang et al., 2022; Nyarko et al., 2021). Most importantly, we stress the importance of grounding technical and policy discussions of fairness in terms of real-world quantities. For example, in the pretrial domain, one

³²In many analyses that do not explicitly grapple with the causal effects of interventions, the estimand Y is further corrupted by the fact that some patients are expected to be enrolled in the program. As a result, some of the sickest patients may not be prioritized for care, since their expected outcome already incorporates the fact that they would have been enrolled, leading to a prediction paradox. To avoid this situation, one could explicitly estimate and prioritize patients by $Y_i(0)$, the potential outcome in the absence of care. In this case, however, the allocation decisions do not necessarily lead to the largest health gains, as the patients likely to be sickest in the absence of care are not typically the same as those likely to benefit the most from the program.

might consider a risk assessment’s short and long-term impacts on public safety and the size of the incarcerated population, as well as a tool’s alignment with principles of due process. In lending, one could similarly consider a risk assessment’s immediate and equilibrium effects on community development and the sustainability of a loan program. Formal mathematical measures of fairness only indirectly address such issues, and can inadvertently lead discussions astray. Of course, it is not always clear how best to quantify or to balance the relevant costs and benefits of proposed algorithmic interventions. In some cases, it may be possible to conduct randomized controlled trials; in other cases, the best one can do is hypothesize about an algorithm’s potential effects. Regardless, we believe a more explicit focus on consequences is necessary to make progress.

We further recommend decoupling the statistical problem of risk assessment from the policy problem of designing interventions. At their best, predictive algorithms estimate the likelihood of events under different scenarios; they cannot dictate policy. An algorithm might (correctly) infer that a defendant has a 20% chance of committing a violent crime if released, but that fact does not, in and of itself, determine a course of action. For example, detention is not the only alternative to release, as one could take any number of rehabilitative interventions (Barabas et al., 2018). Even if detention is deemed an appropriate intervention, one must still determine what threshold would appropriately balance public safety with the social and financial costs of detention. One might even decide that society’s goals are best achieved by setting different thresholds for different groups. For example, a policymaker might reason that, all else being equal, the social costs of detaining a single parent are higher than the social costs of detaining an individual without children, and thus decide to apply different thresholds to the two groups. When policymakers consider these options and others, we believe the primary role of a risk assessment tool is, as its name suggests, to estimate risk. This view, however, is at odds with requiring that algorithms satisfy popular fairness criteria. Such constrained algorithms typically do not reflect the best available estimates of risk, and thus implicitly conflate the statistical and policy problems.

Fair machine learning still has much left to accomplish and there are several important avenues of research that could benefit from new statistical and computational insights. From mitigating measurement error and sample bias, to understanding externalities and equilibrium effects, to eliciting and aggregating preferences to arbitrate between competing algorithms, there is much work to be done. But the benefits are equally large. When carefully designed and evaluated, statistical

algorithms have the potential to dramatically improve both the efficacy and equity of consequential decisions. As these algorithms are increasingly deployed in all walks of life, it will become ever more important to ensure they are fair.

Chapter 3

Fair Allocation through Selective Information Acquisition

3.1 Introduction

Approximately one in seven U.S. households have unmet demand for small-dollar loans, and are often unable to secure credit from traditional financial institutions as they have little or no formal credit history (Apaam et al., 2018). However, in the majority of these households, individuals receive regular income and typically pay their bills on time, which suggests many in fact would have low risk of default (Apaam et al., 2018). One barrier to providing loans to this low-risk yet underserved subpopulation is that it can be more expensive and time-consuming to screen individuals with non-traditional financial histories, limiting the inclusiveness of the banking system.

Motivated by this problem, we propose and analyze a strategy in which one can pay to acquire additional information on applicants. The task for a budget-constrained decision maker is thus to first select a set of candidates to screen and then, given the results of that screening process, determine to whom to allocate the remaining resources. In modeling this situation, we assume there is a fixed cost for screening each applicant, and that decision makers have prior knowledge of the distribution of information they would receive if they choose to screen an applicant. In practice, we note that such prior knowledge could be obtained by screening a small random sample of applicants

Acknowledgments. We thank Sam Corbett-Davies, Bernardo García Bulle Bueno, and Zhiyuan Jerry Lin for helpful comments and discussions.

to learn the resulting information distributions.

We derive an efficient algorithm for computing an optimal, utility-maximizing strategy for the general screening and allocation problem. To do so, we first show that once a set of candidates has been selected to screen, it is optimal to allocate the remaining resources according to a threshold rule, with assets distributed to those candidates having post-screening expected utility above a fixed threshold. Further, for any fixed threshold policy, we show that one can find the optimal set of candidates to screen (while satisfying the budget constraint) via a linear program. Intuitively, one should screen candidates near the margin, for whom the screening process could reveal information that could push a candidate across the threshold. But to do this rigorously, one also needs to account for the precise structure of the prior information. Finally, we sweep over the possible thresholds, solving the corresponding linear program at each point. In this manner, we obtain both a rule to screen candidates and a specific threshold policy for distributing funds to candidates with sufficiently high post-screening value.

We further consider an extension of the above problem in which policymakers have explicit value for diversity. For example, instead of simply finding a max-utility policy, one might maximize utility subject to the constraint that a particular group—such as those who traditionally have had limited access to credit markets—achieve at least a fixed minimum utility. In the United States, those with unmet demand for credit are disproportionately black and Hispanic (Apaam et al., 2018), heightening the value of diversity considerations in allocation decisions. We show that this (and related) extensions can be incorporated into our general algorithmic approach in a straightforward manner.

To demonstrate the potential value of augmenting allocation decisions with a screening phase, we apply our methods to both synthetic datasets and one with real measures of creditworthiness. In particular, we examine the potential benefits of screening as a function of the cost and value of information. Especially when we impose a diversity constraint, we find that screening strategies can significantly outperform a naive strategy that simply attempts to satisfy the constraint without screening any applicants.

For concreteness, we frame our discussion in terms of lending decisions, though our approach applies to many allocation settings. For example, it is often challenging to accurately assess household wealth—particularly in countries where informal and irregular work is more common—and, in turn,

to appropriately target the distribution of government subsidies (Noriega-Campero et al., 2019). The simple strategy of distributing funds to those families clearly in need can systematically overlook populations with harder-to-verify financial status. As with lending, one can judiciously allocate some of the budget to more extensively screen certain applicants, ensuring funds are ultimately distributed to those who can benefit the most.

3.2 Related Work

Several papers address the problem of *active feature acquisition* (Melville et al., 2004, 2005; Saartsehansky et al., 2009), where one can selectively purchase missing data to improve the overall out-of-sample performance of a statistical model. We consider the related problem of acquiring features to identify specific, high-value individuals. While there is some shared intuition between the two settings—that one should seek information on individuals most likely to alter downstream decisions—the technical approach we take is different, in large part because our end goal is optimal allocation rather than statistical learning.

In a related, recent stream of research, Bakker et al. (2019) and Noriega-Campero et al. (2019) likewise consider a feature acquisition problem, but with fairness constraints. In their setting, the decision maker must acquire additional features for each individual to ensure classification decisions have similar error rates across groups—including parity in false negative and false positive rates—a common measure of fairness in the machine learning community (Chouldechova, 2017; Hardt et al., 2016; Kleinberg et al., 2017). Our approach to the problem differs in three important respects. First, we adopt the perspective of constrained utility maximization. Past work has shown that directly equalizing error rates can lead to outcomes that, counterintuitively, may harm the very groups they were designed to protect (Corbett-Davies et al., 2017; Liu et al., 2018; Corbett-Davies et al., 2023). We avoid such deleterious outcomes by instead framing the problem explicitly in terms of group-specific utilities: fairness is encoded into our requirement that the decision maker must allocate some minimum amount of utility to each group. Second, we focus on one-shot screening decisions, in which decision makers simply choose whether or not to acquire information on each individual, rather than sequentially deciding how much information to acquire based on the results of each past acquisition decision. Our one-shot formulation maps to the binary decision structure

(i.e., to screen or not to screen) common in many institutions and leads to different optimization challenges. Third, we directly model the tradeoff between screening and allocation decisions by tying both to a common budget constraint (i.e., more screening means less funds are available to ultimately distribute to individuals).

Elzayn et al. (2019) also consider equitable ways to allocate resources, but in a setting where learning occurs by repeatedly allocating resources instead of by purchasing information. In comparison to our work, they focus on solving the problem of censored feedback, where the learner may not understand, and thus never allocate resources to, groups who did not previously receive resources.

Finally, our work touches on research from the fair division and allocation literature (Brams and Taylor, 1996; Moulin, 2004; Thomson, 2011), which considers how to share resources while satisfying fairness properties defined between individuals. The former devises mechanisms wherein strategic agents self-divide the resources fairly, and the latter studies the existence of allocations that jointly satisfy various fairness notions. In contrast to our work, that line of research is particularly concerned with individual incentives, strategic action, and equilibrium effects.

3.3 A Model of Screening and Allocation

We model screening and allocation decisions as a sequential process in which a budget-constrained lender first selects a (possibly random) subset of candidates to further screen from a pool of n applicants, and then, based on the information revealed in that screening phase, selects a second (possibly random) subset of candidates to receive a loan.

We assume the value of lending to an applicant i is given by the random variable U_i . These utilities are intended to capture the full social value of providing loans, and we imagine the lender aims to optimize social welfare, as in the case of a government agency. In general, the lender has only partial information about U_i . More specifically, if the lender chooses not to screen an applicant, we assume the lender knows only the applicant’s conditional expectation $\mu_i = \mathbb{E}[U_i \mid X_i = x_i]$ given their pre-screening covariates x_i , such as credit score for those applicants who have traditional credit histories. On the other hand, if the lender opts to screen an applicant, they learn $\mathbb{E}[U_i \mid X_i = x_i, \tilde{X}_i = \tilde{x}_i]$, where $\tilde{x}_i$ denotes the additional information one gains through screening. For example, $\tilde{x}_i$ might encode applicant i ’s history of paying their electricity or phone bills—information that is often

feasible to acquire with some extra effort and which is a good indicator of creditworthiness (Apaam et al., 2018).

When deciding whom to screen, we assume the lender knows the distribution of $\mathcal{D} = (D_1, \dots, D_n)$, where $D_i = \mathbb{E}[U_i \mid X_i = x_i, \tilde{X}_i]$. That is, the lender knows how their estimate of utility could change if they decide to screen each applicant, where these distributions may depend on the available pre-screening covariates. A lender may, for example, thus choose only to screen applicants whose estimate is likely to substantially change given additional information. We further assume the lender must pay a fixed cost c_S for screening an applicant and a cost c_A for underwriting a loan. For simplicity we assume these costs do not vary across applicants, though it is straightforward to extend to the more general case; see Appendix B.1 for details.

Based on knowledge of the above information and cost structure, the lender selects a randomized strategy to screen applicants. That is, the lender chooses a vector $\mathcal{P} = (p_1, \dots, p_n)$, meaning that each applicant i is selected to be screened independently with probability p_i . Let $\mathcal{S} = (S_1, \dots, S_n)$ indicate which applicants are ultimately screened under this policy; therefore, $S_i \in \{0, 1\}$ is a Bernoulli random variable with probability of success p_i .

Given this randomized screening policy, we can now write the information $\hat{\mathcal{U}} = (\hat{U}_1, \dots, \hat{U}_n)$ the lender has at the end of the screening phase as follows:

$$\hat{U}_i = \begin{cases} \mathbb{E}[U_i \mid X_i = x_i] & \text{if } S_i = 0, \\ \mathbb{E}[U_i \mid X_i = x_i, \tilde{X}_i] & \text{if } S_i = 1. \end{cases} \quad (3.1)$$

In other words, $\hat{U}_i$ is the lender's post-screening estimated utility of giving a loan to applicant i . In particular, if $S_i = 1$ (i.e., the applicant is screened), the lender's estimate changes from $\mathbb{E}[U_i \mid X_i = x_i]$, the estimate based only on applicant i 's pre-screening covariates x_i , to $\mathbb{E}[U_i \mid X_i = x_i, \tilde{X}_i = \tilde{x}_i]$, which incorporates the post-screening information $\tilde{x}_i$.

Finally, the lender chooses an allocation policy, denoted $\mathcal{A}(\hat{\mathcal{U}}) = (A_1, \dots, A_n)$, where $A_i \in [0, 1]$ specifies the probability a loan is (independently) offered to each applicant. Importantly, $\mathcal{A}$ is a function of the lender's post-screening utility estimates $\hat{\mathcal{U}}$. For example, the lender might give loans to the individuals with the highest post-screening utility estimates, up to the budget constraint.

Combining all of the above, the lender's optimization problem is to choose screening and allocation

policies $(\mathcal{P}^*, \mathcal{A}^*)$ that maximize expected welfare,

$$(\mathcal{P}^*, \mathcal{A}^*) \in \arg \max_{\mathcal{P}, \mathcal{A}} \mathbb{E} \left[\sum_{i=1}^n U_i A_i \right], \quad (3.2)$$

subject to being budget-balanced in expectation,¹

$$\mathbb{E} \left[\sum_{i=1}^n c_S S_i + c_A A_i \right] \leq B, \quad (3.3)$$

where B is a fixed, non-negative constant.

In some settings, decision makers may value diversity in their allocations. For example, they may wish to ensure a certain minimum number of loans are provided to groups that historically have been excluded from credit markets. One can encode this policy preference directly into the utilities, in which case the resulting optimization problem would incorporate one's value for diversity. That approach, however, requires decision makers to agree upon these utilities to interpret the results, which can be challenging.

Here we take a complementary approach that explicitly allows value for diversity to differ across decision makers. Suppose the applicant pool is partitioned into m groups $\{G_1, \dots, G_m\}$; for example, if $m = 2$, we might partition candidates into those who traditionally have had access to credit markets and those who have not. Then we require the selected policy $(\mathcal{P}^*, \mathcal{A}^*)$ to allocate at least $\Lambda_j \geq 0$ utility to group G_j , where these utilities do not themselves include any value for diversity:

$$\mathbb{E} \left[\sum_{i \in G_j} U_i A_i \right] \geq \Lambda_j. \quad (3.4)$$

In practice, as we discuss below, one would solve this optimization problem for a range of Λ , which traces out the Pareto frontier of possible policies across different group constraints, corresponding to different values for diversity. The diversity condition above is expressed in terms of utility, but we might, alternatively, simply lower bound the *number* of loans $\sum_{i \in G_j} A_i$ given to members of each group G_j . This alternative constraint can be handled in a straightforward manner by our algorithm detailed below.

¹By requiring the budget constraint to hold only in expectation—rather than exactly—we are able to find a computationally tractable solution to the problem. In practice, such a constraint means that agencies are allowed some flexibility as long as they don't overspend on average, which, we believe, is often a realistic requirement.

3.3.1 A stylized example

We illustrate the above ideas in the context of a simple, stylized example. Suppose a lender must decide how best to allocate loans among an applicant pool of 13 people, with an overall budget of \$2,000. Providing a loan costs $c_A = \$400$, and additional screening of an applicant costs $c_S = \$50$.

Further suppose that five of the applicants are able to provide rich credit histories, and the expected utility of giving each of them a loan is $\mu_i = \$750$. For this group, additional screening would not provide any more information. Imagine that the other eight applicants do not have formal credit histories, and the utility of giving them a loan is accordingly lower due to the risk of default, with $\mu_i = \$500$. However, the lender knows that these applicants come in two types that could be disambiguated through additional screening. More specifically, the lender knows that after screening, individuals in this group can be divided into those with expected utility \$1,000 (with 50% chance) and those with expected utility \$0 (with 50% chance). The high-utility group could, for example, correspond to those who demonstrate a history of consistently paying their bills on time.

The naive strategy that does not screen any individuals would allocate five loans to hit the budget constraint, since $5 \times \$400 = \$2,000$. All five loans would go to applicants with rich credit histories ($\mu_i = \$750$) over those without ($\mu_i = \500). In this case, the total expected utility of the no-screening allocation is $5 \times \$750 = \$3,750$.

But in this scenario, one can improve overall utility by screening all eight applicants without formal credit histories and then granting loans to the high-utility applicants that are identified. Under that strategy, we expect four of the eight screened applicants to be identified as high utility (\$1,000), and so the expected utility of the allocation is $4 \times \$1,000 = \$4,000$, greater than the expected utility of \$3,750 under the no-screening strategy. Finally, the expected cost of the strategy is $8 \times \$50 = \400 for screening plus $4 \times \$400 = \$1,600$ for distributing the loans, totaling \$2,000 and satisfying the budget constraint. In this example, one can thus improve overall utility—and even allocate more loans to the group that *a priori* appears worse—by incorporating additional screening into the decision-making process.

3.4 Finding Optimal Policies

We now derive an efficient algorithm to find optimal screening and allocation policies $(\mathcal{P}^*, \mathcal{A}^*)$, subject to the budget and diversity constraints. We start by showing that over the full space of policies, it is optimal to allocate resources according to a *threshold policy*, with loans dispersed to individuals having post-screening expected utility $\hat{U}_i$ above a group-specific threshold t_j for $j = 1, \dots, m$. Then, for each threshold policy, we show the optimal screening policy can be obtained by solving a linear program (LP), a type of optimization problem with linear objective function and linear constraints, for which there exist fast solution methods (Bertsimas and Tsitsiklis, 1997). As a result, we can find an optimal combined screening and allocation policy by sweeping over threshold policies and solving the corresponding LP for each such policy.

We begin by formally defining threshold policies.

Definition 3.1 (Threshold Policy). A *threshold policy* is an allocation policy $\mathcal{A}$ for some fixed $t_j \in \mathbb{R} \cup \{-\infty, \infty\}$ and $\alpha_j \in [0, 1]$, $j = 1, \dots, m$, such that

$$A_i = \begin{cases} 1 & \hat{U}_i > t_{g_i}, \\ \alpha_{g_i} & \hat{U}_i = t_{g_i}, \\ 0 & \hat{U}_i < t_{g_i}, \end{cases}$$

where g_i denotes the group membership of individual i .

Threshold policies deterministically allocate resources to those with post-screening expected utilities above a fixed, group-specific threshold t_j . Randomization (i.e., allocating resources with probability α_j) at the threshold t_j may be necessary to exactly satisfy the budget constraint, which is important because for an individual not screened, the distribution of $\hat{U}_i$ is concentrated at a single point.

Theorem 3.2 below formally states that it is sufficient to restrict to the set of threshold policies when searching for a globally optimal screening and allocation policy.

Theorem 3.2. *Suppose the constrained optimization problem defined by Eqs. (3.2), (3.3), and (3.4) has a solution $(\mathcal{P}^*, \mathcal{A}^*)$. Then there is a threshold policy $\mathcal{T}^*$ such that $(\mathcal{P}^*, \mathcal{T}^*)$ is also a solution.*

To see this, suppose that $\mathcal{P}^*$ deterministically selects a subset of applicants to screen. Then, for

each group G_j , it is clear one should allocate loans to the approximately k_j individuals in each group with the highest post-screening estimated utility, where k_j is the number of loans granted to each group under $\mathcal{A}^*$. Such a rank-based allocation can equivalently be written as a threshold rule with the same expected utility. The more general case, in which screening decisions are randomized, introduces some technical complications, but the spirit of the argument is similar. Full details can be found in Appendix B.1.

Now, given a threshold policy $\mathcal{T}$ with thresholds $t_1, \dots, t_m \in \mathbb{R}$ and boundary randomization probabilities $\alpha_1, \dots, \alpha_m \in [0, 1]$, we turn to finding an optimal companion screening policy. We show that such an optimal screening policy $\mathcal{P}$ can be found by solving an LP. In particular, the LP has n decision variables $p_1, \dots, p_n$, with $p_i \in [0, 1]$ specifying the probability that applicant i is screened; the objective equals the utility of the combined screening and allocation policy; and the constraints encode our budget and diversity conditions.

To construct the LP, we first define the following quantities that depend on both the threshold rule $\mathcal{T}$ and the lender's prior knowledge on the value of screening:

$$\begin{aligned} q_i &= \Pr(D_i > t_{g_i}), \\ e_i &= \mathbb{E}[U_i \mid D_i > t_{g_i}], \\ o_i &= \mathbb{1}_{\mu_i > t_{g_i}} + \alpha_{g_i} \cdot \mathbb{1}_{\mu_i = t_{g_i}}, \end{aligned}$$

where g_i denotes applicant i 's group membership. Recall that μ_i denotes pre-screening expected utility and D_i denotes post-screening expected utility if applicant i is screened. Both μ_i and the distribution of D_i are known in advance.

For fixed screening probabilities $p = (p_1, \dots, p_n)$, the expected utility of the corresponding screening and allocation policy can now be expressed as a linear function of p :

$$\sum_{i=1}^n [q_i \cdot e_i \cdot p_i + o_i \cdot \mu_i \cdot (1 - p_i)]. \quad (3.5)$$

The first summand reflects the expected utility associated with applicant i if they were screened, and the second summand reflects the expected utility if they were not screened.

Our goal is to maximize the expression in Eq. (3.5) subject to the budget and diversity conditions, which we now show can also be expressed as linear constraints on p . In terms of the constants

defined above, the budget constraint can be written as

$$\sum_{i=1}^n [c_S \cdot p_i + c_A \cdot q_i \cdot p_i + c_A \cdot o_i \cdot (1 - p_i)] \leq B. \quad (3.6)$$

Likewise, the j diversity constraints in Eq. (3.4) become

$$\sum_{i \in G_j} [q_i \cdot e_i \cdot p_i + o_i \cdot \mu_i \cdot (1 - p_i)] \geq \Lambda_j, \quad (3.7)$$

for $j = 1, \dots, m$.

Together, the objective given by Eq. (3.5), with constraints defined by Eqs. (3.6) and (3.7), define a linear program with decision variables $\{p_i\}$, the solution of which—if one exists—gives a screening policy $\mathcal{P}^*$ that is optimal for the threshold policy $\mathcal{T}$. A jointly optimal screening and allocation policy $(\mathcal{P}^*, \mathcal{T}^*)$ can accordingly be found through a grid search over all threshold policies $\mathcal{T}$ in a (discretized) space $\mathbb{R}^m \times [0, 1]^m$. For each choice of policy $\mathcal{T}$, defined by threshold and randomization parameters $t_1, \dots, t_m$ and $\alpha_1, \dots, \alpha_m$, we find the corresponding optimal screening policy $\mathcal{P}$ by solving the LP described above. Then, among the resulting screening-allocation pairs, the one maximizing utility is guaranteed to be globally optimal.

3.5 Experiments

We now investigate the value of our sequential screening and allocation approach through two simulation exercises, one based on synthetic data and another based on real loan data. First, with the synthetic data, we examine the optimal policies as we vary both the cost of screening and the value of the resulting information. Then, with the real-world data, we illustrate how a decision maker could, in practice, operationalize our screening and allocation approach to pursue equity when distributing limited resources.

In each experiment, we allocate loans to members of two groups. One of the groups—which we call the “targeted” group—can be screened for more information at some cost, and we imagine there is social value to distributing more resources to this group. For example, the targeted group may be comprised of those who traditionally have not had ready access to the banking system and accordingly have limited formal credit history. For simplicity, we further assume that those in the non-targeted group cannot be screened, perhaps because they already have complete credit histories.




Figure 3.1: For an applicant with pre-screening likelihood of being creditworthy $x_i = 0.4$, the distribution of post-screening creditworthiness, $\Pr[Y_i = 1 \mid X_i = 0.4, \tilde{X}_i]$. Under a regime with high value of information (red line), the post-screening distribution has higher variance than in a setting with low value of information (blue line).

We define the utility of lending to applicant i to be

$$U_i = \begin{cases} a & Y_i = 1, \\ b & Y_i = 0, \end{cases} \quad (3.8)$$

where $Y_i \in \{0, 1\}$ indicates whether they would pay back the loan, and a and b are fixed, known constants. In both of our experiments, we set $a = \$1,000$ and $b = -\$200$, meaning there is \$1,000 of social utility when an applicant receives and pays back a loan and $-\$200$ utility when an applicant defaults on a loan, for example because defaulting could trigger further financial distress. As above, we imagine the lender is a government institution or other agent attempting to maximize total social utility. We further assume that the pre-screening covariate $x_i \in [0, 1]$ specifies the lender's pre-screening estimate of applicant i 's likelihood to repay a loan; in other words, $x_i = \Pr(Y_i = 1 \mid X_i = x_i)$. The covariate x_i can thus be translated into (pre-screening) expected utility μ_i by Eq. (3.8):

$$\mathbb{E}[U_i \mid X_i = x_i] = ax_i + b(1 - x_i).$$

3.5.1 Synthetic data

We illustrate the value of screening in four regimes of low vs. high cost of information paired with low vs. high value of information. To do so, we created four synthetic datasets, each comprised of $n = 500$ individuals.

For all four datasets, we first evenly split the population into targeted and non-targeted groups. For each applicant in the targeted group, we generated their pre-screening probability of repayment x_i (or, equivalently, their observed pre-screening covariates) by independently drawing from a beta distribution² with mean 0.5 and count parameter 50.³ Similarly, for each applicant in the non-targeted group, we generated their pre-screening probability of repayment by independently drawing from a beta distribution with mean 0.70 and count parameter 50. The higher mean repayment probability for members of the non-targeted group corresponds to them being, on average, more creditworthy.

Now, for each applicant in the targeted group, the lender may elect to screen them. As a result of screening, the lender receives an improved estimate $\tilde{x}_i$ of the applicant’s repayment probability, so that:

$$\mathbb{E}[U_i \mid X_i = x_i, \tilde{X}_i = \tilde{x}_i] = a\tilde{x}_i + b(1 - \tilde{x}_i).$$

We assume $\tilde{x}_i$ is drawn from a beta distribution with mean x_i (i.e., the lender’s pre-screening estimate of the applicant’s repayment probability). In the high-information scenario, we set the count parameter for this beta distribution to be 5; and in the low-information scenario, we set it equal to 25.

Figure 3.1 shows these two post-screening information distributions for an applicant with $x_i = 0.4$. As illustrated in the plot, the high-information distribution (red line) has higher variance than the low-information distribution (blue line), and so screening is more likely to reveal very high risk and very low risk applicants in the high-information setting. Finally, we set the cost of screening c_S to be \$25 in the low-cost scenario and \$100 in the high-cost scenario, the cost of a loan to be $c_A = \$1,000$, and the total budget to be $B = \$50,000$.

Given these four datasets, we computed the optimal screening and allocation strategies using the algorithm described above.⁴ In our original problem formulation, the diversity constraint in

²We use the beta as it is a familiar, easily-parameterized function with support on $[0, 1]$.

³In terms of the α and β shape parameters often used to parameterize beta distributions, the mean is $\alpha/(\alpha + \beta)$ and the count parameter is $\alpha + \beta$. A higher count corresponds to a lower variance.

⁴Because only one group can be screened in our experiments, we use a faster variant of our optimization algorithm, described in Appendix B.2. For any fixed set of parameters, this approach returns the optimal screening and allocation policy within a few seconds with the open-source LP solver SCS (O’Donoghue et al., 2016).




Figure 3.2: Comparison of our optimal screening strategy (blue line) with one without screening (black line) in four regimes with different costs and values of information.

Eq. (3.4) specified only that we lower bound the utility of each group. To better understand the impact of diversity on utility, we modify this constraint to be a strict equality for the utility of the targeted group and set the lower bound on utility to be \$0 for the non-targeted group. Thus, across a range of exactly satisfied utilities for the targeted group, we find the strategy that maximizes overall utility.

Figure 3.2 shows the results of this analysis, with the blue lines tracing out the Pareto frontiers for each of the four scenarios we consider. For comparison, the black lines show the corresponding result under a strategy that does not screen any applicants. Specifically, for any fixed utility constraint on the targeted group, the optimal no-screening policy first allocates loans to the k individuals in the targeted group most likely to repay based on their pre-screening estimates x_i , where k is chosen to satisfy the utility constraint; and then any remaining budget is used to allocate loans to those in the non-targeted group most likely to repay.

When either the value of information is high (left column) or the cost of screening is low (bottom row), we find that screening can be a valuable tool to improve utility. Notably, screening plays a more important role in these examples as we demand greater utility be allocated to the targeted group, since the no-screening strategy ends up dispersing loans to relatively high-risk applicants in

the targeted group even though more creditworthy applicants in that group could be identified for little marginal cost. As one might expect, the gap between the screening and no-screening strategies is particularly large when both information is valuable and cheap. Indeed, in the high-value, low-cost setting (lower-left panel), one can achieve substantial diversity with little drop in overall utility.

3.5.2 Empirical credit data

We next apply our approach to the German Credit Dataset (Hofmann, 1994), which includes a variety of individual-level socioeconomic and financial characteristics (e.g., age, employment status, and credit history) on a sample of $n = 1,000$ people, of whom 700 are deemed creditworthy. We define the targeted group to be those who currently do not own their residence, a subpopulation that comprises 28% of the dataset. In this case, 60% of individuals in the targeted group are creditworthy compared to 74% in the non-targeted group. For members of the targeted group, we assume the lender, prior to screening, only knows the targeted group’s overall base rate of creditworthiness. If, however, the lender chooses to screen an applicant in the targeted group, they learn $\tilde{x}_i$, the applicant’s likelihood of being creditworthy conditional on all the available features in the dataset. For members of the non-targeted group, we assume this full estimate of creditworthiness is available prior to screening, and that there is no opportunity to obtain additional information.⁵ Finally, we translate estimates of creditworthiness to estimates of utility via Eq. (3.8), in line with our simulations above.

Figure 3.3 shows the result of applying our screening and allocation algorithm to this dataset, where we assume the cost of screening c_S is \$100 (equal to our high cost regime in the synthetic datasets), the cost of a loan c_A is \$1,000, and the total budget is $B = \$150,000$. Like before, we compare the Pareto frontier of our approach (blue line) to that of a naive policy in which the lender does not screen applicants (black line). As with the synthetic datasets above, we find that the optimal policies with screening substantially outperform those without screening, particularly when we enforce a diversity constraint. For example, when we require \$50,000 of utility to come from allocating loans to the targeted group, the maximum total utility under the no-screening policy is

⁵More specifically, at the start of this exercise, we train a logistic regression model on the full dataset predicting creditworthiness as a function of the available covariates. Then, for members of the targeted group, $\tilde{x}_i$ is the model-estimated probability of creditworthiness for applicant i ; and for the members of the non-targeted group, that same model estimate is available prior to screening.




Figure 3.3: For the German Credit Dataset, comparison of our optimal screening strategy (blue line) with one without screening (black line).

\$102,000, compared to \$119,000 under the screening policy, an increase of 17%.

3.6 Discussion

Many creditworthy individuals often have difficulty gaining access to traditional credit markets due to lack of formal financial histories, an issue that can exacerbate existing socioeconomic disparities. To address this gap, we developed a simple and efficient algorithm for a budget-constrained decision maker to screen applicants and then allocate a limited resource, an approach that we find offers substantial benefits on both real and synthetic lending datasets. This joint screening-plus-allocation approach is especially useful in settings where a targeted subset of the population—those most in need of an intervention—are also those for whom the least information is available *a priori*, a common situation in many social welfare programs.

Past research has shown that a dearth of high-quality data for various subgroups of the population can lead to poor models in a variety of domains (Gebru et al., 2018; Olteanu et al., 2019), including text analysis (Bolukbasi et al., 2016; Caliskan et al., 2017; Garg et al., 2018), facial recognition (Buolamwini and Gebru, 2018), and automated hiring (Dastin, 2018). Looking forward, our combined data acquisition and decision-making approach provides one framework to address this challenge by jointly modeling the cost of data collection and the value for subsequent improvements in downstream decisions.

Chapter 4

Measuring Disparate Impact in Human and Machine Decisions

There are two main legal doctrines of discrimination in the United States: disparate treatment and disparate impact. Disparate treatment protections bar government agents from acting with “discriminatory purpose” (Washington v. Davis, 1976); disparate impact statutes, in contrast, bar policies that produce “unjustified” adverse effects on legally protected groups, even in the absence of animus (Grover, 1995). Among empirical researchers and practitioners, disparate treatment has, by far, received more attention. In the most common methodological approach, sometimes known as “kitchen sink” regression, analysts compare decision rates across race or other protected characteristics after adjusting for all available decision factors,¹ typically using ordinary linear or logistic regression. For instance, in the lending context, one might compare loan officers’ lending rates to Black and White loan applicants after adjusting for income, credit score, and other relevant factors (Munnell et al., 1996). This basic statistical strategy has been used in numerous studies to test for bias across domains including education (Espenshade et al., 2004; Grossman et al., 2024b), employment (Polachek, 2008), criminal justice (Abrams, 2014; MacDonald and Raphael, 2020; Rehavi and Starr, 2012; Gaebler et al., 2022; Gelman et al., 2007), and medicine (Balsa et al., 2005).

While analyses of disparate treatment are common in both the scientific literature and legal

Acknowledgments. We thank Josh Grossman, Melody Y. Huang, Julian Nyarko, and João M. Souto-Maior for helpful conversations.

¹An important exception is adjusting for close proxies of protected characteristics, such as hair or eye color. For more detailed discussion of disparate treatment and disparate impact jurisprudence, see Appendix C.1.

practice, the doctrine is ill-suited to assessing algorithmic decisions. In particular, algorithms are often formally “race blind,” and so cannot treat individuals from different groups differently in the legally relevant sense. For instance, it is rare for protected traits to be explicitly input into the models now common in lending (Fed, 2024), medicine (Coots et al., 2025), criminal justice (Corbett-Davies et al., 2023), and other settings. Accordingly, attempts to audit these algorithms (e.g., with a standard regression analysis) will typically reveal, correctly, that race did not directly impact decisions.

Nonetheless, facially neutral algorithms can still adversely impact groups defined by race, gender, and other traits. For example, race- and gender-blind algorithms in healthcare and criminal justice have been found to be miscalibrated, producing risk estimates that are systematically skewed by group (coo, 2025). In banking, the ways in which algorithmically generated loan offers are presented and marketed to online applicants can likewise create and exacerbate disparities. Finally, algorithms based on generative AI, which are increasingly used to guide employment decisions, might elevate vague understandings of “leadership” or “fit” in ways that disadvantage certain demographic groups—all without explicitly considering demographic traits (Gaebler et al., 2024).

Disparate impact, in contrast to disparate treatment, is concerned with the effects of a policy, not a decision maker’s intentions, and more naturally applies to the analysis of algorithmic decisions. Under the disparate impact standard, a practice may be deemed discriminatory if it has an unjustified adverse effect on protected groups, even in the absence of explicit categorization or animus. The disparate impact doctrine was formalized in the landmark U.S. Supreme Court case *Griggs v. Duke Power Co.* (1971).² In 1955, the Duke Power Company instituted a policy that mandated employees have a high school diploma to be considered for promotion, which had the effect of drastically limiting the eligibility of Black employees. The Court found that this requirement had little relation to job performance, and thus deemed it to have an unjustified disparate impact.

Despite its importance, disparate impact analysis has received comparatively little attention in the empirical community (Arnold et al., 2022; Grossman et al., 2024b; Souto-Maior and Shroff, 2024; Ayres, 2010; Grossman et al., 2024b; Ayres, 2005). Yet “kitchen sink” regression—the standard

²The legal landscape surrounding disparate impact has recently shifted in the U.S., following an executive order deprioritizing disparate impact enforcement at the federal level (Executive Order, 2025). The order also revokes certain regulatory provisions that adopt a disparate impact framework, though statutory provisions such as Title VII of the Civil Rights Act governing disparate impact at the state and federal level remain in force.

approach for assessing disparate treatment claims—is a poor fit for assessing whether practices are rationally justified, the relevant standard in disparate impact cases. Ayres (2005) makes the point in the context of the original *Griggs* decision:

“One could imagine running a regression to test whether an employer was less likely to hire African American applicants than white applicants. It would be possible to control in this regression for whether the applicant had received a high-school diploma. Under the facts of *Griggs*, such a control would likely have reduced the racial disparity in the hiring rates. But including in the regression a variable controlling for applicants’ education would be inappropriate. The central point of *Griggs* was to determine whether the employer’s diploma requirement had a disparate racial impact. The possibility that including a diploma variable would reduce the estimated race effect in the regression would in no way be inconsistent with a theory that the employer’s diploma requirement disparately excluded African Americans from employment.”

In short, by including educational status in the regression, one would mask the policy’s unjustified disparate impact.³

Ayres (2005) terms this phenomenon *included-variable bias*—“the worry that the statistical estimates of disparate impact are biased because the regression inappropriately includes non-race variables” (Ayres, 2010).⁴ In *Griggs*, the Court found that education was unrelated to permissible employment aims, and so its inclusion in a regression analysis of disparate impact would have been inappropriate. But whether it is appropriate to control for a particular covariate is not always clear. In other employment contexts, education might be more strongly related to job performance. Similarly, given existing patterns of residential segregation, home ZIP codes tend to be highly correlated with race, weighing in favor of their exclusion when testing for disparate impact. But one could also argue that ZIP codes provide legitimate information relevant to a decision, and so excluding them would itself be problematic. Ayres (2010) proposes a middle ground, suggesting that all covariates be included, but their coefficients capped to a “justifiable” level; in practice, however, it is difficult to determine and defend specific constraints on regression coefficients.⁵

³In Appendix C.1.2, we give a simple formal illustration of this statistical phenomenon.

⁴Included-variable bias differs from post-treatment bias or what is sometimes termed “adjusting for bad controls” (Cinelli et al., 2024). In particular, the aim of the risk-adjusted regression framework we introduce is not to estimate a causal quantity, but rather the extent to which decisions are normatively or legally justified.

⁵Such problems have prompted a search for alternatives to regression-based approaches. Most prominently, Becker (1993) proposed the outcome test, which is based not on the rate at which decisions are made, but on the success rate of those decisions. In the context of banking, Becker argued that even if, hypothetically, racial minorities were less creditworthy than White applicants, minorities who were granted loans should still be found to repay their loans at the same rate as White applicants who were granted loans. If loans to minorities had a higher repayment rate than

Here we present a statistically principled and logistically straightforward approach for measuring disparate impact in human and machine decisions. Our method, which we call *risk-adjusted regression*, proceeds in three steps. In the first step, we use all available information, including potential proxies of protected traits, to estimate the expected net benefit of taking a particular action—which we term “risk.”⁶ For example, in the lending context, one might approximate net benefit by estimating the expected profitability of a loan, conditional on all available covariates; in hiring one could instead estimate an applicant’s expected productivity. Importantly, this step requires access to a rich set of covariates, which may not be available in all cases. In the second step, we assess disparities by regressing decisions (e.g., loan offers) on individual-level risk estimates and protected traits alone, allowing us to measure the extent to which similarly risky individuals receive different decisions. This strategy can be seen as formalizing Ayres’s (2010) coefficient-capping procedure—with covariates used only to the extent that they are statistically justified by risk—and thus circumvents the problem of included-variable bias. Finally, we assess the sensitivity of results to potential mismeasurement of risk. In particular, we derive tight analytic bounds on risk-adjusted disparities as a function of how much risk estimates differ from true risk.

To demonstrate this approach, we examine 2.2 million stops of pedestrians conducted by the New York City Police Department between 2008 and 2011. After adjusting for a stopped individual’s statistical risk of carrying a weapon—based in part on detailed behavioral indicators recorded by officers—we find that stopped Black and Hispanic pedestrians are searched for weapons substantially more often than stopped White individuals. We find that these risk-adjusted disparities are considerably larger than disparities suggested by a standard regression that adjusts for all available

loans to White borrowers, it suggests that lenders applied a double standard (intentionally or not), granting loans only to exceptionally qualified minorities—potentially violating disparate impact laws. The outcome-based approach has been applied almost as broadly as simple regression to study discrimination in, for example, policing (Goel et al., 2016b, 2017; Knowles et al., 2001; Ayres, 2002), lending (Berkovec et al., 2018), and scientific publication (Smart and Waldfogel, 1996). Becker’s outcome test, however, has significant statistical shortcomings, most notably the problem of infra-marginality (Ayres, 2002), which can lead the test to incorrectly suggest an absence of discrimination (Pierson et al., 2020). Refinements of Becker’s test, such as the threshold test (Pierson et al., 2018; Simoiu et al., 2017), and, more recently, the robust outcome test (Gaebler and Goel, 2025) address inframarginality, but involve alternative assumptions and can be less informative than regression-based approaches when detailed covariate information is available.

⁶In many cases, the expected net benefit is well captured by the probability of a particular outcome, which is why we call it “risk.” In lending, for example, it is common to approximate the profitability of loan by a loan applicant’s risk of default. In policing, searches are often motivated by the likelihood, i.e., “risk,” that an individual possesses contraband such as weapons or drugs. Our net benefit framing implicitly assumes that the underlying utility is justifiable. For example, a bank can reasonably seek to maximize profits, but their utility cannot reflect a “taste” for discrimination (Becker, 1957).

covariates, underscoring the importance of accounting for included-variable bias. Finally, we show that our results are robust to potentially large errors in risk estimates.

4.1 Statistically Assessing Disparate Impact

Consider estimating risk-adjusted racial disparities in police searches of pedestrians for weapons—the application we discuss in more detail below. Let A_i indicate whether the i -th stopped pedestrian was searched, X_i , the information available to the officer when deciding whether to conduct a search, C_i the stopped individual’s race, and W_i whether the individual is in possession of a weapon—regardless of the officer’s search decision.

Given this setup, we define *ex ante risk* to be:

$$R = \Pr(W = 1 \mid X). \quad (4.1)$$

In our policing example, R_i is the probability that the i -th stopped individual, with covariates X_i , is carrying a weapon.⁷

Our goal is to quantify whether decisions systematically differ for individuals at the same level of risk. Many estimands summarizing such potential differences are possible, but, for simplicity, we consider the following non-parametric quantity:

$$\mathbb{E}[\Pr(A = 1 \mid C = j, R) - \Pr(A = 1 \mid C = 1, R)]. \quad (4.2)$$

(4.2) defines risk-adjusted disparities to be the difference in the probability of taking an action for group $C = j$ relative to the reference group $C = 1$, after accounting for potential differences in risk across groups.⁸

⁷To allow for the possibility of statistical discrimination (Arrow, 1973; Chaudhuri and Sethi, 2008; Phelps, 1972) in the decision maker’s course of action, risk may be conditioned on group membership as well as other observables—formally, by including C among the covariates X . However, whether the inclusion of race or other protected attributes in risk models is legally appropriate is uncertain. Employers, for instance, are barred from statistically discriminating (U.S. Equal Employment Opportunity Commission, 2009); statistical discrimination in medical decision making, however, is common if controversial (coo, 2025; Coots et al., 2025), though it has in some instances been successfully challenged in court (Kruse and Pinalis, 2017). The empirical results in our main policing application remain essentially unchanged whether or not one assumes X encodes C ; see Figures C.11 through C.17.

⁸To ensure this quantity is well-defined, we assume that:

$$0 < \Pr(C = j \mid R) < 1, \quad \text{for all } j = 1, \dots, m.$$

This condition is immediately implied by an analogue of the standard overlap condition: $0 < \Pr(C = j \mid X) <$

To facilitate computation—and, in particular, the sensitivity analysis presented below—we approximate the estimand in (4.2) by a linear probability model. Specifically, we estimate $\Pr(A = 1 \mid C, R)$ as

$$\sum_{j=1}^m \hat{\beta}_j \cdot \mathbf{1}(C = j) + \hat{\beta}_R \cdot R, \quad (4.3)$$

where $\mathbf{1}(C = j)$ indicates membership in group j and $\hat{\beta}_j$ is its corresponding fitted coefficient. The difference between fitted coefficients,

$$d_j = \hat{\beta}_j - \hat{\beta}_1, \quad (4.4)$$

is our formal measure of disparate impact.⁹

In our running example of police searches, if $d_j > 0$ —i.e., if $\hat{\beta}_j$ is greater than $\hat{\beta}_1$ —it means that members of group j are searched more often than members of the reference group who were equally likely to be carrying a weapon. We would say that such elevated search rates are *unjustified* by risk. We note, however, that $d_j > 0$ does not imply *intentional* discrimination. As in *Griggs*, unjustified disparate impact is possible even under a facially neutral policy undertaken without animus.

We have introduced risk-adjusted regression in the context of policing, but our general approach applies to both human and algorithmic decisions across a wide variety of settings. For example, we could similarly consider risk-adjusted disparities in hiring, where W represents a prospective employee’s productivity and A whether or not they were hired; and lending, where W represents the amount a prospective borrower will repay and A the bank’s lending decision. In contexts featuring a continuous rather than binary outcome measure W , such as these two examples, one would more generally define the expected net benefit R to be $\mathbb{E}[W \mid X]$, the expected outcome given observed covariates X .

In our policing application—as well as in these other settings—empirically estimating risk-adjusted disparities d_j is complicated by two factors. First, the information available to officers typically differs from that recorded in administrative datasets available to analysts. Second, analysts—and even officers—generally only know whether *searched* individuals were carrying weapons, but not whether *unsearched* individuals were carrying weapons. In other words, W is only observed when

1 (Rosenbaum and Rubin, 1983a,b).

⁹We provide a general, formal treatment of risk-adjusted regression in Appendix C.2. See Appendix C.3 for further details on the relationship between d_j and the non-parametric version in (4.2).

$A = 1$. As a result, one cannot usually compute the risk R exactly, motivating our sensitivity analysis described below.

4.2 An Application to Policing

We apply our approach above to investigate the “stop-and-frisk” practices of the New York City Police Department (NYPD).¹⁰ Police officers in the United States may stop and question pedestrians if they have “reasonable and articulable” suspicion of criminal activity; officers may additionally conduct a “frisk” (i.e., a brief pat-down of one’s outer garments) if they believe the stopped individual is carrying a weapon. Although a policy of stopping and frisking individuals is not inherently illegal, in *Floyd v. City of New York* (2013), a federal district court ruled that the NYPD carried out such stops with racial animus, violating the Equal Protection Clause of the Fourteenth Amendment.

The court in *Floyd* was interested in assessing claims of disparate treatment; here we re-analyze the data with a focus on disparate impact. We specifically consider frisk decisions, as they have a clear goal of ensuring officer safety by recovering weapons, and a well-measured outcome—whether a weapon was in fact found. We study 2.2 million pedestrian stops that occurred between 2008 and 2011. For each stop, we have detailed information on the date, time, and location of the stop; the demographics of the stopped individual (e.g., age, gender, and race); the suspected crime; the reasons prompting the stop (e.g., “furtive movements” or “suspicious bulge”); and additional circumstances surrounding the stop (e.g., evasive responses to questioning, witness reports, or evidence of criminal activity in the vicinity).¹¹

To start, we note that 1.7% of frisks turn up a weapon. White pedestrians are frisked in 44% of police stops, whereas Black and Hispanic pedestrians are frisked in 57% and 58% of stops, respectively, a 13–14 percentage point gap. These raw disparities are computed without adjusting for any covariates that might justify the observed differences between groups. At the other extreme is the “kitchen sink regression,” which adjusts for all observed pre-frisk covariates in a standard linear probability model. In this case, stopped Black and Hispanic pedestrians are 3–4 percentage points more likely to be frisked, relative to White pedestrians with similar recorded characteristics.

¹⁰Reproduction materials for this analysis are available at <https://github.com/jgaeb/rar-repro>.

¹¹This information is recorded in a standardized way on UF-250 forms that officers were required to complete after each stop. A copy of the form can be found online.




Figure 4.1: *Frisk rates vs. risk, as estimated via logistic regression curves fit separately for each race group. Across risk levels, stopped Black and Hispanic pedestrians are frisked substantially more frequently than comparably risky White individuals, indicative of disparate impact.*

These kitchen-sink disparities are suggestive of disparate treatment (and similar evidence was indeed presented to the court in *Floyd* to support such an allegation), but they may understate the extent to which the policy imposes an unjustified disparate impact on racial minorities, due to included-variable bias.

4.2.1 Estimating risk-adjusted disparities

The key ingredient in applying risk-adjusted regression is estimating the risk $R = \Pr(W = 1 \mid X)$, as in (4.1), where W indicates whether a stopped individual has a weapon and X is the information available to the officer when making their frisk decision. As noted above, X and W are only partially and imperfectly observed, twin challenges that can be mitigated, but not eliminated. Our approach is thus to estimate risk as best we can, and then gauge the robustness of our results to estimation error. For ease of exposition, we adopt a simple estimation strategy, regressing weapon possession (W) on the covariates in the recorded data ($\tilde{X}$), fitting our model only on the subset of individuals who were, in reality, frisked. This yields estimates $\hat{d}_j$ of (4.4), based on these estimated risks.¹²

¹²Our approach to estimating risk is not the only one, or even the best in any given instance. For example, when an appropriate sample exists, one could partially address this selection issue by leveraging some subpopulation of individuals that were randomly frisked, as is sometimes done at roadside checkpoints and airports. Alternatively, one could look for officers who frisk almost everyone they stop—similar to recent approaches for studying judicial




Figure 4.2: *Racial gaps in frisk rates adjusting for different sets of covariates, where the y-axis shows the percentage point difference relative to stopped White individuals. The red dots show the raw disparities in frisk rates. As a measure of discrimination, raw disparities suffer from omitted-variable bias: there may, in theory, be legitimate reasons why Black and Hispanic pedestrians are more likely to be frisked. The green triangles show the estimated race effects in a kitchen-sink regression, adjusting for all pre-frisk covariates. These estimates suffer from included-variable bias because they adjust for features that are correlated with race but unrelated to risk. The blue squares show the results of our risk-adjusted regression, adjusting exclusively for estimated risk of weapon possession. In all cases, estimated standard errors are less than 0.2 percentage points, and so are not visible in the plot.*

We start by dividing the original stops into two subsets: (1) the approximately 1 million stops that occurred in 2008 and 2009, which we use to train a risk model; and (2) the remaining 1.2 million stops that occurred in 2010 and 2011, to which we apply the fitted risk model to estimate disparities in frisk decisions. In this way, we produce an estimate of *ex ante* risk $\hat{R}_i$ for every pedestrian stopped in the second subset of the data, including those who were not frisked. (See Materials and Methods below for further details.)

Figure 4.1 shows frisk rates as a function of estimated risk, disaggregated by race, also known as “risk-decision curves” (Gaebler and Goel, 2025).¹³ At every level of risk, stopped Black and Hispanic pedestrians are frisked at a much higher rate than stopped White individuals, a gap that

decision-making (e.g., Dobbie et al., 2018; Kleinberg et al., 2018). Both of these alternatives come with their own limitations, meaning that sensitivity analysis is still critical to credible inference.

¹³The curves show that even pedestrians with very low probabilities of carrying weapons have around a 40% probability of being frisked, indicating that officers’ frisk decisions are fairly inefficient. If officers frisked pedestrians according to a threshold rule, searching them if and only if their risk were above some cutoff, they could recover the same number of weapons while frisking fewer people (Goel et al., 2016a).

is suggestive of disparate impact in frisk decisions.

To add quantitative detail to these qualitative results, we next compute risk-adjusted disparities, fitting the linear probability model in (4.3) on the second half of the NYPD data, computing estimates for the Black-White disparity $\hat{d}_{\text{Black}}$ and Hispanic-White disparity $\hat{d}_{\text{Hispanic}}$. Figure 4.2 shows the results, together with the raw disparities and those estimated from a kitchen-sink model. We find that stopped Black and Hispanic pedestrians were about 15 percentage points more likely to be frisked than White pedestrians who were equally likely to be carrying a weapon. Further, the risk-adjusted disparities are in fact *greater* than the raw disparities in frisk rates. To understand why, we note that stopped White pedestrians were, on average, more likely to be carrying a weapon yet less likely to be frisked than racial minorities; as a result, the risk-adjusted gap in frisk rates is even larger than the raw, unadjusted gap. Finally, we see that the kitchen-sink regression dramatically underestimates the extent of disparate impact faced by minorities. In this case, the kitchen-sink model adjusts for a variety of features—including whether the stopped individual made “furtive movements”—that are correlated with race but are poor predictors of weapon possession, skewing estimates of disparate impact.¹⁴

4.2.2 Sensitivity analysis

Computing risk-adjusted disparities depends critically on knowing individual-level risks. But we can only approximate these risks, both because officers observe factors that are predictive of risk but are not recorded in our data, and because outcomes (i.e., weapon possession) are only partially observed. It is thus important to gauge the sensitivity of our estimates of disparate impact to errors in our estimates of risk. We accordingly develop a method of sensitivity analysis that bounds the difference between d_j and $\hat{d}_j$ when we constrain the true risks R_i and estimated risks $\hat{R}_i$ to differ by some maximum amount, i.e., when

$$\frac{1}{n} \sum_{i=1}^n |\hat{R}_i - R_i| \leq \epsilon \quad (4.5)$$

¹⁴We exclude hair color and eye color from the kitchen-sink model, since these are obvious proxies for race that are effectively unrelated to risk, and would thus be excluded in most traditional legal and statistical analyses of discrimination. As we would expect, including these variables as controls exacerbates the problem of included-variable bias, but our results show that such bias can occur even if obviously problematic variables are excluded.




Figure 4.3: *Sensitivity of the risk-adjusted disparities in frisk decisions to mismeasurement of risk. The blue bands bound our estimates of disparate impact as a function of the average absolute difference between true and estimated risks ϵ , relative to the base rate (1.7%). The dotted line at 44% ($\epsilon = 0.7$ p.p.) corresponds to a simulated situation with severe confounding. The small grey bands along the top and bottom of the blue bands represent 95% percentile bootstrapped confidence intervals ($N = 1000$; Zhao et al., 2019).*

for some average absolute error ϵ .¹⁵ The relationship between $\hat{R}_i - R_i$ and d_j is non-convex and involves a potentially large number of optimization variables, hindering computation. In Appendix C.2, we derive an efficient algorithm for performing the sensitivity analysis, which we implement in the `rar` package for R, available on CRAN.

Figure 4.3 displays our bounds on disparate impact as a function of the mismatch between true and estimated risk—operationalized in terms of ϵ , as defined in (4.5). To ease interpretation, the horizontal axis in the plot is expressed in terms of the *relative* average absolute difference between true and estimated risks: ϵ divided by the overall weapon recovery rate among frisked individuals (1.7%). Our analysis shows that large risk-adjusted disparities remain even if we allow the true risk to differ considerably from our risk estimates. In particular, we would find large disparate impacts for both Black and Hispanic pedestrians even if the true risks differed from our risk estimates by

¹⁵Our approach, detailed in Appendix C.2, is related to, but distinct from, well-known methods of sensitivity analysis in the causal inference literature that aim to gauge the impacts of unobserved confounding (Rosenbaum, 2002). In particular, as our statistical problem differs from that in the standard causal inference setting, those methods cannot be applied straightforwardly in our context. Our method, moreover, is particularly well-adapted to our setting, yielding reasonably tight bounds in practice that depend only on the interpretable parameter ϵ , without making additional assumptions about the form of the error $\hat{R}_i - R_i$. See Appendix C.2 for further discussion of our method and alternatives.

50% of the base rate.

It is impossible to know the precise extent to which our risk estimates differ from the true risk. To understand the plausible magnitude of the discrepancy, we conduct a simulation in which we remove a key set of risk-relevant covariates from the data, estimate risk based on the reduced information, and then measure differences between the original and new risk estimates. Specifically, we remove variables listed in two sections of the UF-250 stop forms that describe the “circumstances” prompting the encounter. These sections each consist of 10 binary variables—including, for example, “fits description,” “actions indicative of casing,” and “changing direction at sight of officer”—that are crucial for establishing the legal basis of the stop. We then compute the average absolute difference between our original risk estimates based on the full, uncensored data and the risk estimates based on the redacted data.¹⁶ We find that this value is 0.7 percentage points—or 44% of the base rate—which we take as an estimate of ϵ in a scenario with severe missing data. As shown in Figure 4.3, this level of error (indicated by the dotted vertical lines) would yield an estimate of disparate impact that is, at a minimum, greater than 7.5 percentage points for both Black and Hispanic individuals. This sensitivity analysis suggests our results are robust to substantial mismeasurement of risk.

4.3 Discussion

Our primary policing application examined human decisions, but similar considerations apply to algorithms—especially those based on generative AI. Consider, for instance, hiring tools based on large language models (LLMs), which aim to naturalistically evaluate potential job candidates based on their resumes and interview performance, much as human recruiters do (Gaebler et al., 2024). To promote broad usefulness, LLMs undergo a complex, largely self-supervised training process very different from the fitting of conventional risk models (Bai et al., 2022; Ouyang et al., 2022). This process is still not fully understood, and there is no guarantee that the hiring recommendations of conventional LLMs will correspond to statistically valid *ex ante* estimates of productivity. As a result, these systems, like the humans they emulate, have the potential to produce large risk-adjusted

¹⁶As above, we use gradient boosted decision trees to estimate the probability of frisk conditional on the remaining, uncensored covariates.

disparities.¹⁷ While the widespread deployment of LLM-based decision tools promises to increase efficiency and reduce certain forms of human bias in the hiring process and beyond, disparate impact may remain a persistent problem.

In view of this form of discrimination’s continuing importance, we have sought to develop a simple, intuitive framework for formalizing and measuring disparate impact. Our approach, however, is subject to some important limitations. First, our method requires access to an outcome to estimate risk. In some instances, this information is not available to analysts. In other cases, it is not even clear how to rigorously define the relevant outcome. For example, in college admissions, decision makers often care about multiple factors in ways that are hard to quantify (Garg et al., 2021; Grossman et al., 2024b). Second, to credibly estimate risk, analysts need sufficiently rich covariate data. Our sensitivity analysis helps mitigate omitted-variable bias, but it cannot replace better data. Third, and relatedly, analysts may not have access to race or other relevant protected characteristics, complicating the analysis. There has, however, been recent progress on estimating disparities using auxiliary data sources (e.g., Kallus et al., 2022; McCartan et al., 2025). Finally, we have modeled decisions using linear probability models with constant slope across groups. It is straightforward to estimate risk-adjusted disparities with more flexible decision models. However, relaxing this assumption makes it harder to gauge the sensitivity of the inferred disparities to inaccurate risk estimates.

Throughout our analysis, we have estimated “disparate impact” by a regression coefficient on protected-group identity in a model that adjusts for estimated risk. This procedure is analogous to current practice in discrimination studies, where we simply replace the usual set of control variables with a single variable capturing risk. As seen by our formalization of disparate impact in (4.2), we are effectively measuring a particular weighted average of differences in decision rates across individuals of similar risk. While intuitively reasonable, this definition raises subtle questions of law and policy.

Consider, for example, Figure 4.1, where we plot race-specific frisk rates as a function of risk. Stopped Black and Hispanic pedestrians are frisked more often than stopped White pedestrians at

¹⁷In contrast, traditional supervised machine learning models trained to directly predict relevant outcomes (e.g., job performance) minimize risk-adjusted disparities by design—though one must take care to avoid “label bias” (Zanger-Tishler et al., 2024).

every level of risk. As a result, one would find that racial minorities face disparate impact regardless of how one averages across risk levels; the precise number might change, but the qualitative conclusion would remain the same. However, comparing Black and Hispanic pedestrians, the direction of the disparity depends on the risk level one considers.¹⁸ Low-risk Hispanic individuals are frisked more often than low-risk Black individuals, but high-risk Black individuals are frisked more often than their high-risk Hispanic counterparts. Consequently, a conclusion of disparate impact between Black and Hispanic pedestrians would depend heavily on the precise definition applied. The analysis is further complicated if the risk distributions differ substantially between groups. If, hypothetically, stopped Hispanic pedestrians were mostly low-risk and stopped Black pedestrians mostly high-risk, majorities of both groups could argue that they were treated more harshly than members of the other group who were equally likely to be carrying a weapon.¹⁹

The crossing of risk-decision curves that we see in Figure 4.1 is a potentially widespread phenomenon, and, to our knowledge, disparate impact law has not yet resolved the underlying conceptual ambiguity it invokes. Many discussions of disparate impact tacitly assume that policies either consistently harm or help groups defined by protected traits. Such thinking can be seen in the original *Griggs* ruling, where the Supreme Court aimed to proscribe policies that acted as “built-in headwinds” for racial minorities. But, formally, disparate impact law concerns facially race-neutral policies, not intentional discrimination. The Court has addressed related questions, for instance holding in *Connecticut v. Teal* (1982) that a particular stage of a selection process can result in disparate impact even if the process as a whole advantages a protected group. But there is no theoretical or empirical guarantee that any policy or stage of a selection process will adversely impact all members of a particular group.

A related issue is the extent to which concern for unjustified disparities compels decision makers to act optimally. For example, Figure 4.1 suggests that officers are only marginally responsive to risk, with the lowest-risk individuals still frisked more than 40% of the time. If, instead, officers

¹⁸Such a comparison between minority groups is unusual in disparate impact cases, but it illustrates the underlying theoretical issue.

¹⁹Some scholars have similarly investigated interactions between race and other decision-making criteria. For example, Espenshade et al. (2004) find that preferences for underrepresented minorities in college admissions is greatest for applicants with SAT scores in the 1200–1300 range, and the effect is attenuated at lower scores. That analysis, however, found no score ranges where minority applicants faced an absolute disadvantage relative to White applicants with equal scores. We do not know of any research that has found a change in the direction of the disparities like we see between Black and Hispanic pedestrians in Figure 4.1.

frisk only the people with high probability of carrying a weapon, they could frisk far fewer individuals—and, in particular, far fewer minority individuals—while recovering almost the same number of weapons (Goel et al., 2016a). A more efficient frisk strategy could thus reduce the burdens of policing on racial minorities while still maintaining public safety. Such efficiency is indeed one of the aims of statistical risk assessment tools that are now used in the criminal justice system and beyond to guide high-stakes decisions (Chouldechova et al., 2018; Corbett-Davies et al., 2023; Goel et al., 2021; Monahan and Skeem, 2016; Shroff, 2017). If these tools are shown to reduce racial disparities, are policymakers obliged—legally or ethically—to adopt them? The role of efficiency in disparate impact claims has largely gone unanswered by the courts, adding yet another subtlety to defining and measuring disparities. Researchers have only recently taken up these questions (Bohren et al., 2025; Corbett-Davies et al., 2023; Elzayn et al., 2023; Grossman et al., 2024b; Raghavan, 2023).

By foregrounding the role of risk in understanding disparities, we have aimed to clarify some of the thorny conceptual issues at the heart of disparate impact analysis. While there are still important unresolved questions, we believe that our statistical approach provides practitioners with a tractable way to assess disparities in many domains while avoiding some important pitfalls of traditional methods. Looking forward, we hope this work spurs further theoretical and empirical research on discrimination at the intersection of statistics, economics, law, and public policy.

4.4 Materials and Methods

To estimate risk of weapon possession in Section 4.2, we use gradient boosted decision trees, a non-linear model popular in the machine learning community for its predictive performance, restricting to stops in the first subset of data in which a frisk was conducted. Predictive performance and model checks presented in Figure C.10 indicate that the model yields predictions with reasonable performance and that predictions are well-calibrated across groups. The distribution of risk estimates produced by the model are shown in Figures C.8 and C.9. (Figures C.11 through C.17 reproduce these results and the results in Figures 4.1 through 4.3 with race excluded from the risk-prediction model.) For details on the sensitivity analysis, including parameter choices for the optimization, see Appendix C.2.

Chapter 5

Auditing Large Language Models for Race & Gender Disparities: Implications for Artificial Intelligence–Based Hiring

Artificial intelligence (AI) tools increasingly assist employers with many aspects of decision-making. One prominent example is human resources (HR) management, an area in which AI tools have been used to facilitate benefits administration, coaching, development, and applicant screening. Prior to 2022, the global market for AI hiring technology was already growing rapidly. But the advent of large language models (LLMs)—AI models that are highly adept at understanding and generating text—has dramatically boosted interest in AI hiring tools.

LLMs are a potential boon for any setting where decisions must be made on the basis of a large volume of text. LLMs can be used to do work ranging from grading student essays to evaluating proposals from potential vendors. In the context of HR, LLMs can ingest entire application dossiers—including résumés, essays, and interview transcripts—to produce seemingly cogent assessments of candidates’ qualifications. But using LLMs in this way can be worrisome as well. Even as people put these LLMs to work in hiring, employers and policymakers are racing to establish guidelines for the algorithmic evaluation of candidates.

Bias is a particular concern. Traditional AI tools are “supervised,” meaning users train these

Acknowledgments. We are grateful to Avi Bagchi and Marissa Gerchick for research assistance. Data and replication code are available at <https://github.com/jgaeb/llm-audit>.

programs by inputting labeled data—for instance, previous candidates’ application materials labeled with the ratings recruiters gave them. An AI model can learn statistical patterns from these labeled data and then use those patterns to generate predictions about new applicants. But LLMs go through a more complex and opaque training process that often does not involve any prelabeled data. In an unsupervised process called pretraining, LLMs examine and find patterns in a huge, unlabeled text corpus or data set. These training data sets may contain the equivalent of tens of billions of books (Meta, 2024) and are drawn largely from data publicly available on the internet. After being pretrained, LLMs are posttrained by being provided with a smaller, carefully curated set of data meant to enable them to learn patterns that will improve the accuracy, helpfulness, and safety of their outputs. Once an LLM is trained, users can deliver prompts (written questions and instructions or other text), such as a directive to the LLM to rank job candidates on the basis of their written application materials. Compared with traditional AI tools, the LLMs’ responses more closely resemble those of human evaluators, who might, when evaluating job candidates, produce candidate ratings that are based on an intuitive understanding of how professional experiences and responses to interview questions relate to a candidate’s competency and fit for a job. But for various reasons—including discriminatory content in the pretraining corpus and the complexity of the overall training processes—LLMs might also produce discriminatory or distorted responses that are hard to anticipate.

Employers and policymakers therefore fear that AI could run afoul of employment discrimination laws or otherwise produce unintended, undesirable effects, and they want to find ways to audit these tools to determine whether they are discriminatory. In this article, we demonstrate and examine one potential method for auditing LLMs—correspondence experiments—illustrating the approach by applying it to evaluate LLMs that could be used in making hiring decisions. We demonstrate the technique’s potential and discuss limitations.

5.1 Background

The ethical and legal implications of using AI tools in high-stakes settings, such as HR, have motivated much academic work (Raghavan et al., 2020; Tambe et al., 2019). And policymakers have become as interested as firms and researchers, introducing a wave of legislation governing

the use of algorithms in different contexts, especially hiring. These measures incorporate auditing requirements as an important element, in part because of the belief that informative audits help regulators protect the public from the potential harms of AI tools and that users and firms can use audit results to make informed decisions about which tools to deploy and how to deploy them.

For example, on October 30, 2023, President Joe Biden issued an executive order imposing mandates on federal agencies’ use of AI and calling for regulatory efforts by agencies with authority over private AI uses (White House, 2023).¹ The order twice mentions audits as tools to advance fair public decision-making and to ensure AI safety. Meanwhile, the Digital Services Act (DSA) in Europe, which entered into force in February 2024, requires what it terms “very large online platforms” to conduct audits to promote transparency and accountability (European Commission, 2024). And in March 2024, the European Parliament adopted the more general AI Act, which imposes a variety of requirements that potentially involve audits. Most notably, it compels users of high-risk systems—for instance, systems used in critical infrastructure, biometric applications, law enforcement, and certain other high-stakes domains—to create quality management systems that may include a kind of audit (European Parliament, 2024).

The most developed audit mandate for AI HR tools was introduced in the United States under a New York City ordinance, Local Law (LL) 144, effective July 2023 (N.Y.C. Local Law No. 144, 2021). LL 144 requires a bias audit when employers use “any computational process, derived from machine learning, statistical modeling, data analytics, or artificial intelligence,” to classify or recommend persons for employment. In these bias audits, independent third parties must calculate and publicly report an adverse impact ratio. This is defined by LL 144 as the rate at which individuals in a race or gender group are hired or move forward in the hiring process relative to people in the most frequently selected race or gender group. LL 144 imposes no legal obligations when a disparity is identified—in other words, no particular course of action is required. Employers also have broad discretion to determine whether they are covered by the measure. In the first 6 months after LL 144’s entry into force, only 19 audits linked to the law were published (Groves et al., 2024).

The adverse impact ratio is a relatively common auditing tool in both policymaking and scientific research—required not only by LL 144 but also by traditional hiring regulations such as the US

¹Executive Order on the Safe, Secure, and Trustworthy Development and Use of Artificial Intelligence, October 30, 2023.

Equal Employment Opportunity Commission’s four-fifths rule (which holds that selecting candidates from a protected or minority group less than 80% as often as members of the group with the highest selection rate can be considered evidence of discrimination). But the adverse impact ratio is known to be imperfect: It cannot reveal whether disproportionate selection of one kind of candidate over another indicates a true psychological bias or if it simply reflects average differences in the qualifications of applicants from different groups.

Beyond the adverse impact ratio, relatively few tools exist for identifying potentially biased decision-making by LLMs. To date, no scientific consensus exists on how best to audit algorithms for bias, although researchers have proposed a wide variety of algorithmic fairness metrics (Berk et al., 2021; Chiappa, 2019; Chouldechova, 2017; Chouldechova and Roth, 2020; Corbett-Davies et al., 2023; Coston et al., 2020; Dwork et al., 2012; Hardt et al., 2016; Imai and Jiang, 2020; Kilbertus et al., 2017; Kleinberg et al., 2017; Kusner et al., 2017; Loftus et al., 2018; Nabi and Shpitser, 2018; Raghavan, 2023; Wang et al., 2019; Zafar et al., 2017; Corbett-Davies et al., 2017). Here, we explore using correspondence experiments to audit LLMs for race and gender bias in high-stakes decision settings such as HR, in part because in behavioral science research, correspondence experiments have proven valuable for identifying discrimination in hiring decisions.

In correspondence experiments (also known as audit studies), researchers assume that two otherwise identical individuals from different demographic groups should receive similar treatment and that divergence is evidence of improper discrimination (Lippens et al., 2023). These studies are typically set in contexts where decision-makers learn about individuals exclusively through written documents (for example, an initial screening of job applicants). Researchers then experimentally manipulate elements of those materials that could suggest an individual’s race and gender. (For simplicity, throughout the balance of this article, we use the term race to refer to race or ethnicity.) In the study described here, we mainly assessed the effect of the name applicants use on their résumé, a practice in line with most research on this topic in the social sciences (Gaddis, 2019). Names are the strongest signal of race typically perceived by recruiters in the United States. In other contexts or in countries with different hiring conventions, auditors might manipulate other application elements, such as the applicant photos included with résumés in Germany, Japan, China,

and many other countries.²

For at least 50 years, social scientists and government agencies have used correspondence experiments to study discrimination in hiring (Bertrand and Mullainathan, 2004), housing (Wienk, 1979), prosecutorial charging decisions (Chohlas-Wood et al., 2021), and other domains (Ayres et al., 2015; Gaddis and Ghoshal, 2020; Lyons-Padilla et al., 2019). More recently, correspondence experiments have been proposed to similarly identify evidence of bias in the algorithms used in AI (An et al., 2024; Armstrong et al., 2024; Bandy, 2021; Eloundou et al., 2024; Haim et al., 2024; Lippens, 2024; Metaxa et al., 2021; Nghiem et al., 2024; Sandvig et al., 2014; Tamkin et al., 2023; Veldanda et al., 2023). For example, Sweeney (2013) found that Google searches of Black-sounding names were more likely to generate advertisements suggesting the named individual had an arrest record than were comparable searches of White-sounding names.

In our exploration of the potential value of using correspondence experiments to audit AI-based hiring decisions, we found overall that correspondence experiments are useful for identifying race and gender bias, and we provide a workable strategy for carrying out mandated algorithm audits. However, as we explain in the Discussion section, we also identified some key conceptual and technical limitations of this approach for auditing algorithms.

5.2 Empirical Analysis and Results

5.2.1 LLMs for Candidate Evaluation

To evaluate the use of correspondence experiments for auditing hiring algorithms, we first gathered a novel corpus from 1,373 applications to K–12 teaching positions in a large public school district in Texas by filing a public records request. (To our knowledge, this school district does not use algorithms to evaluate applicants.) Application materials included the applicants’ résumés as well as transcripts that we produced from self-recorded videos from the applicants. In these

²In general, an auditor’s main concern when designing a correspondence experiment is for manipulated applications to maintain consistency. If they do not—for instance, if the auditor pairs a male name with an application that uses female pronouns to describe the applicant or pairs a picture of a White applicant with a résumé that lists attendance at a historically Black college or university—differences in recruiters’ behavior could reflect their reactions to the incongruity rather than their perception of the manipulated demographics. In our case, any manipulation in an experiment being designed to measure the effect of gender would need to alter names, because in the United States, names are typically very strong gender signals; name manipulation in this country is both necessary and sufficient for altering an applicant’s apparent gender (and, to a lesser extent, race).

videos, applicants answered written questions about previous teaching experience, teaching style, hypothetical classroom situations, and other job-related subjects. Ultimately, we restricted our analysis to the 801 applicants who had provided both a résumé and video responses. These applicants represented a diverse pool, of which 67% were women, 2% were Asian, 45% were Black, 10% were Hispanic, and 38% were White. (We use “White” throughout to mean non-Hispanic White.)

For each applicant, we provided the LLM with (a) a description of the requirements for the teaching position based on a job posting from the school district; (b) the applicant’s résumé; (c) a written transcript of the applicant’s self-recorded responses to interview questions; (d) a request for the model to summarize the applicant’s qualifications in prose; and (e) a request for the model to provide numerical evaluations, on a scale from 1 to 5, on several measures, including the applicant’s experience, professionalism, and fit, as well as the model’s overall hiring recommendation, which was expressed as a score on a scale ranging from 1 (definitely do not hire) to 5 (definitely hire). (See Appendix D.3.1 for more information.) Although our primary statistical analysis focused on the models’ overall numerical hiring recommendation for the applicants, we requested additional information from the models—including written summaries and other scores—to improve the quality of the LLM results, a common strategy when using these models.

We audited 11 LLMs: OpenAI’s GPT-3.5, GPT-4, GPT-4o, and GPT-4o Mini models (Achiam et al., 2023; Brown et al., 2020; OpenAI, 2024a,b); Mistral’s Mistral 7B and Mixtral 8x7B models (Jiang et al., 2023, 2024); and Anthropic’s Claude Instant, Claude 2, Claude 3 Haiku, Claude 3 Sonnet, and Claude 3.5 Sonnet models (Anthropic, 2023, 2024a,b). Experts considered the OpenAI and Anthropic models to be the best of their kind available at the time of our evaluation. Mistral’s models are popular open-source competitors. We did not formally assess how well each LLM’s candidate ratings aligned with candidate qualifications, but an informal inspection suggested that the highest rated candidates generally had more experience and gave more polished responses to interview questions than did those receiving lower ratings. This rough assessment of the LLM ratings’ validity—and the LLMs’ relative ease of use—makes it likely that a pipeline like the one we implemented here will soon be used by employers to screen applicants, if one has not been launched already.

5.2.2 Assessing Adverse Impact Ratios

As we have previously noted, each LLM rated candidates’ overall suitability for the job on a scale of 1 to 5, where a higher number indicates a stronger positive recommendation. OpenAI’s GPT-3.5, one of the most popular and widely used LLMs at the time of our experiments, gave 20% of candidates an overall score of 5, 51% a 4, 24% a 3, 4% a 2, and 1% a 1.

We then turned to the issue of whether each LLM rated candidates similarly across demographic groups. We first looked at the adverse impact ratio, given its prominence in past research and the New York law. That is, we determined the rate at which individuals in one race or gender category were positively selected relative to those in another category. Per the Equal Employment Opportunity Commission’s four-fifths rule, a ratio of 80% or lower is particularly concerning and likely warrants some response (Tobia, 2017), although federal law does not require a particular course of action.

The results of this adverse impact ratio analysis for OpenAI’s GPT-3.5 are shown in Figure 5.1; results across all models were similar. (See Figures D.7 through D.10 in Appendix D for results from other models.) When comparing the proportion of applicants across groups who received a 5, the highest threshold for recommendation, we found that female applicants received positive assessments more often than male applicants did and Black and Hispanic applicants received positive assessments more often than White applicants did. At a recommendation threshold of 4, the pattern flips for race—with White applicants receiving a positive assessment more often than Black and Hispanic applicants did—and we found near parity for gender. Finally, at a threshold of 3, we found near parity across both gender and race groups.

This simple analysis suggests that GPT-3.5 and the other LLMs we studied might be favoring certain demographic groups. However, even with the comparatively large minority applicant pools in our corpus, we can estimate adverse impact ratios only imprecisely. In fact, of nine adverse impact ratios calculated at this stage, only three were statistically significant, suggesting this approach found only slight evidence of meaningful disparities.

Without further evidence, we cannot definitively say whether these disparities are due to algorithmic bias or group-specific differences in the applicant pool. Put another way, the women and racial minorities in our applicant pool might simply be more qualified than the other applicants




Figure 5.1: *GPT-3.5’s adverse impact ratios highlighting disparities that may be the result of algorithmic bias while nevertheless being inconclusive on their own. We calculated adverse impact ratios for the outputs of the large language model GPT-3.5 at a variety of hiring thresholds.*

for these positions, which would explain why the model ratings they received were higher than the ratings received by the other applicants. This makes it difficult to draw clear conclusions, which indicates that an adverse impact ratio analysis alone is too limited for auditing LLMs.

5.2.3 Assessing Correspondence Experiments

Correspondence experiments make it possible to differentiate between algorithmic bias and differences in candidate qualifications. We began by manipulating the real application materials to create synthetic application dossiers that differed only in details that strongly signaled an applicant’s race or gender. For each real applicant, we generated eight synthetic applications, corresponding to a particular race (Asian, Black, Hispanic, or White) and gender (female or male). We then replaced the applicant’s actual name throughout the materials with one that strongly signaled membership in that group. We similarly changed any mention of the applicant’s pronouns in the materials to match

the assigned group, as well as other indicia of race or gender. (See Appendix D.3 for a detailed description of how we generated these synthetic applications.)

We presented the synthetic application materials to the LLMs, instructing the models to report the race and gender of the synthetic applicants. We found generally high agreement between the race and gender we intended to convey and the model’s perception of these attributes, which in most cases exceeded 90% but with the precise level of agreement varying from case to case (see Figures D.1 and D.2 in Appendix D). This level of agreement is comparable to what has been found in studies that manipulate human perceptions of race by altering names in résumés (Bertrand and Mullainathan, 2004).

Finally, we asked our LLMs to provide hiring recommendations for the synthetic candidates. The results of these correspondence experiments are shown in Figure 5.2. (Results for Mixtral 8x7b should be interpreted with caution, because the model often failed to follow the instructions, producing responses without ratings; see Appendix D.5 for details.) For each model, we created comparisons of groups of candidates relative to a historically favored reference group. To assess gender bias, we made men a reference group and looked at the estimated difference in average score for women versus men in each model. In the case of race, we made White people the reference group and looked for the difference in average scores for all other groups in comparison with the average score of these White candidates. (See Appendix D for further methodological details.)

Across models, we found that the LLMs rated synthetic female candidates moderately higher than they did the synthetic male candidates. Models also generally rated synthetic Black, Hispanic, and Asian candidates moderately higher than they did synthetic White candidates, although we found more variation between models, with Mistral’s models exhibiting smaller disparities.

Our correspondence experiments suggest that the models’ perceptions of race and gender influence algorithmic candidate assessments, at least to some degree. At each hiring threshold—the minimum rating (3, 4, or 5) at which we imagine hiring candidates—race and gender disparities were generally a few percentage points (see Figures D.14 and D.15 in Appendix D). These disparities were modest but within the typical range identified in recent studies of human recruiters (cf. Lippens et al., 2023). In other words, the LLMs showed a level of bias smaller than but still comparable to what people tasked with identifying the best job applicants in similar experiments often show.

5.2.4 Sensitivity to Prompt Variation and Context

An LLM’s approach to a given task—such as rating job applicants—can be modified and potentially improved by changing the directions the LLM is given. So we repeated our analysis with several variants of the initial prompt—that is, the set of directions we gave an LLM for evaluating our applicants and their suitability for the job. (See Appendix D.4.2 for more details on the prompt variants investigated.) For simplicity, we ran these robustness tests on only one model, OpenAI’s GPT-3.5, which exhibited approximately average disparities in our primary analysis. Regardless of the prompt used, we found the same general patterns persisted across all variants: a slight bias in favor of women over men and in favor of Hispanic, Black, and Asian applicants over White applicants. (See Figures D.3 and D.4 in Appendix D.)

We then reran our primary analysis while inputting only an applicant’s résumé into the model, excluding interview transcripts. This variation on our experiment allowed us to gauge a scenario that frequently occurs in real-world settings, when employers must make decisions based on résumés alone. As with the robustness tests, we conducted this analysis on only GPT-3.5, and we again found that women and racial minority applicants received moderately higher scores than did men and White applicants, respectively.

Many LLMs can access virtually encyclopedic knowledge and incorporate this information into their evaluations. For example, in our analysis, the American school district to which our candidates applied is especially racially diverse—a fact known by the LLMs we examined. If a model took this information into account, the diversity might influence its ratings in some way. We therefore ran another analysis on GPT-3.5 in which we replaced all mentions of the district (and the city and state) with the name of a predominately White school district in West Virginia. Once again, we found disparities mirroring our primary results.

5.3 Discussion

Our results demonstrate that correspondence experiments can reveal race and gender disparities in LLM outputs. Unlike the commonly used adverse impact ratio analysis, correspondence experiments revealed not only differences in qualifications across groups in the candidate pool but also the effect of race and gender on LLM outputs. These experiments therefore offer policymakers a potentially

useful tool for auditing algorithms. The patterns we observed were substantive and robust, persisting across several variations of our study, including changes to the instructions provided to the model and the specific application materials we inputted.

It bears emphasizing that the specific pattern we observed—with models favoring female applicants over male applicants and favoring Black, Hispanic, and Asian applicants over White applicants—may not generalize across contexts in which people use LLMs. Indeed, although some recent studies that focused on hiring decisions also found disparities similar to those we found in our work (An et al., 2024; Tamkin et al., 2023), others have reported disparities in the opposite direction (Armstrong et al., 2024; Eloundou et al., 2024; Haim et al., 2024; Lippens, 2024; Nghiem et al., 2024; Veldanda et al., 2023). Such contrasting results are not surprising given the complex and often inscrutable ways in which LLMs are trained. In particular, recall that developers typically fine-tune models in a final alignment or posttraining phase in part to avoid mirroring overt discrimination in the training data. But this step may leave traces of bias in ways that are hard to predict. For example, efforts to mitigate discriminatory associations the model learned in pretraining might overshoot the mark, causing a distortion in the other direction.

Although correspondence experiments may be a useful tool for auditing algorithms, they also have notable limitations. First, all experiments—correspondence or otherwise—are inherently limited in their ability to manipulate race or gender while leaving other characteristics of people untouched. In the case of our study, names are imperfect signals of race and gender, and other aspects of the application dossier (such as history of employment in a male- or female-dominated field) can further attenuate that signal. Indeed, in our own pool of applicants, we find that LLM-like statistical models can nearly always infer an applicant’s race and gender from unaltered application materials as well as from masked application materials (that is, application materials with names and pronouns removed; see Figure D.6 in Appendix D). Nevertheless, as already noted, the models still identified the intended race and gender of our synthetic candidates more than 90% of the time, indicating that our name manipulation largely worked as designed.

Further, names may reveal information beyond applicants’ race and gender. Changing an individual’s name can also change an LLM’s perception of their age, their socioeconomic status, and other characteristics aside from race. This effect means we cannot be certain that we have accounted for all the ways in which LLMs rated our applicants or the precise degree to which race and gender

alone could have influenced these decisions. These confounds limit our statistical conclusions.

Conceptually, it is challenging to rigorously define what it would even mean to manipulate an LLM’s perception of an individual applicant’s race in isolation from other factors. Considerable literature has explored the question of whether it is possible to discuss and perceive race in isolation, given the ways in which ideas around racial identity so often couple with perceptions of social status, wealth, religion, neighborhood of residence, and other variables (cf. Sen and Wasow, 2016; Holland, 1986). For that reason, some scholars conceive of race as a “bundle of sticks,” lumping all of those elements together, along with skin color and ancestral origin, when determining an individual’s racial identity (Kang et al., 2016; Sen and Wasow, 2016). Similar considerations apply to gender (Hu and Kohler-Hausmann, 2020). Which of these factors can or should be manipulated in an audit of race or gender bias is a difficult question.

Finally, even if models are not directly influenced by an individual’s race or gender in the way we test here, biases against other aspects of a candidate’s application could still produce undesirable outcomes. For instance, a model might, in theory, prioritize the applications of individuals who attended private school, regardless of their race or other qualifications. That pattern could unjustly disadvantage qualified minority applicants (Jung et al., 2018).

Ultimately, LLMs need to be audited for each specific use, and the conclusions drawn could vary on the basis of both the task and the pool of individuals evaluated. Context- and scenario-specific correspondence experiments can serve as tools for evaluating LLMs in many situations. For example, LLMs are already being used to determine credit risk, and policymakers keen to craft legislation and guidance for this use would benefit from the data generated by correspondence experiment audits of credit risk LLM outputs. The calls to audit algorithms are likely to increase as LLMs become even more capable and widespread and are used in ever more varied decision-making contexts. Correspondence experiments, despite their important limitations, represent one promising method of auditing algorithms for race and gender bias. We hope our work aids in the ongoing regulatory efforts to ensure that LLMs yield equitable outcomes.




Figure 5.2: *Correspondence experiments showing that the large language models tested modestly favor women and non-White candidates. For each large language model, the colored dots indicate the difference in mean applicant rating, reported in estimated population standard deviations, with 70% and 95% confidence intervals indicated by thick and thin bars, respectively. Positive values indicate that the model rated female or racial minority applicants higher than it rated male (left panels) or White (right panels) applicants, respectively. The dashed reference line shows the value corresponding to no average difference in ratings between groups. Most large language models showed the same pattern, with the exception of Mixtral 8x7B, a model that frequently failed to follow user instructions. (The confidence intervals are clustered at the level of the real application dossier used to generate the synthetic application.)*

Part II

Human Discrimination

Chapter 6

A Simple, Statistically Robust Test of Discrimination

When assessing claims of discrimination, researchers often begin by considering whether decision rates differ across groups defined by race or gender, typically after adjusting for relevant differences between groups. For example, to test for discrimination in banking, one might estimate differences in lending rates between White and Black loan applicants after adjusting for an individual’s credit score, income, and savings. Although such a “benchmark test” can be informative, it is prone to omitted-variable bias: Failing to adjust for all relevant information can yield misleading estimates. Nonetheless, benchmark tests have been applied in nearly every domain where discrimination is studied, generally under an implicit assumption that analysts have access to all relevant covariates (Bartlett et al., 2022; Grossman et al., 2024b; MacDonald and Raphael, 2020; Starr and Rehavi, 2013; Gaebler et al., 2022; Gelman et al., 2007).

To mitigate the omitted-variable problem inherent to benchmark tests, Becker (1957, 1993) introduced the “outcome test,” in which one looks not at decision rates but rather *success* rates. If, for example, loans issued to Black borrowers are repaid at higher rates than those issued to White borrowers, it suggests a double—and discriminatory—standard, with bank officials granting loans only to exceptionally creditworthy Black applicants. Owing perhaps to its simplicity and intuitive appeal, the outcome test has now become one of the most popular empirical approaches to detecting

Acknowledgments. We thank Nick Arnosti, Chris Avery, Madison Coots, Eliana La Ferrara, Julian Nyarko, Todd Rogers, Soroush Saghafian, Kevin Yang, and Michael Zanger-Tischler for helpful conversations and feedback. Data and replication code are available at <https://github.com/jgaeb/outcomepp>.

discrimination. Researchers have applied the test to audit a wide range of decisions, including lending, hiring, publication, and candidate election (Anzia and Berry, 2011; Berkovec et al., 1994, 1998; Chilton et al., 2020; Green et al., 2009; Pope and Sydnor, 2011; Smart and Waldfogel, 1996). The outcome test has gained particular prominence in criminal justice, among both researchers and policymakers (Antonovics and Knight, 2009; Close and Mason, 2007; Coviello and Persico, 2015; Goel et al., 2016; Neil and Winship, 2019; Persico and Todd, 2006; Ridgeway, 2007; Fryer Jr, 2019; Goel et al., 2017; Pierson et al., 2020).

Like the benchmark test, however, the outcome test suffers from well-known statistical limitations (Ayres, 2002; Carr and Megbolugbe, 1993; Engel, 2008; Galster, 1993; Anwar and Fang, 2006; Simoiu et al., 2017). Consider the stylized example in Figure 6.1, where the red and blue curves show the distribution of repayment probability across loan applicants in two different groups (henceforth, “risk distributions”). In this hypothetical, bank officials grant loans to those applicants who are at least 50% likely to repay their loans—indicated by the dashed black vertical line—irrespective of group membership. Despite this uniform lending standard, loan recipients in the blue group are more likely to repay their loans than recipients in the red group. In statistical terms, conditional on being above the lending threshold, the mean of the blue group’s risk distribution is greater than the mean of the red group’s. As a result, the outcome test would incorrectly conclude that applicants in the blue group were subject to a more stringent lending standard.

This problem of “inframarginality” has attracted considerable attention, prompting several attempts to place outcome tests on firmer statistical footing. Knowles et al. (2001) developed a model of behavior under which risk distributions collapse to a single point, eliminating the possibility of inframarginality. Although theoretically interesting, the key assumption in that approach has been critiqued for being at odds with empirical evidence (Engel and Tillyer, 2008; Jung et al., 2024). Anwar and Fang (2006) proposed a test based on decision and outcome rates conditional on the race of both decision makers and those subject to those decisions. Their method is guaranteed, under certain conditions, to produce correct inferences, but it can only identify relative disparities between decision makers from different race groups. Building on that work, Alesina and La Ferrara (2014) proposed a test of racial bias in capital sentencing based on the relative likelihood that decisions were overturned across defendant-victim race pairs. Arnold et al. (2018) sidestepped concerns of inframarginality by directly estimating outcomes for individuals at the margin, leveraging




Figure 6.1: *A stylized example illustrating the problem of inframarginality. The two curves depict the distribution of repayment probabilities for two hypothetical subpopulations. Applying a uniform lending threshold of 50% (dashed black vertical line) results in a higher repayment rate for loan recipients in the blue group (71%; dotted blue vertical line) than for recipients in the red group (64%; dotted red vertical line). The outcome test would thus incorrectly infer that members of the blue group were subjected to a more stringent lending standard.*

quasi-random assignment of decision makers. Theirs is a statistically compelling approach but can only be applied in certain settings, where decision makers are plausibly quasi-randomly assigned and analysts have information on the actions of individual decision makers. Simoiu et al. (2017) and Pierson et al. (2018) worked to overcome inframarginality by simultaneously estimating risk distributions and decision thresholds with a parametric model. Their approach, however, is sensitive to the exact model form, and, in particular, estimates are not identified by the data alone. Finally, Jung et al. (2024) used detailed individual-level information on covariates and outcomes to directly estimate group-specific risk distributions. The method is effective when it can be used (Grossman et al., 2023, 2024a; Souto-Maior and Shroff, 2024), though the demanding data requirements limit the applicability of their approach.

Despite the limitations of both the benchmark and outcome tests, here we show that simply combining the two yields a *robust* outcome test with surprisingly strong statistical guarantees. In particular, if the group-specific risk distributions satisfy the monotone likelihood ratio property (MLRP) (Karlin and Rubin, 1956), then either the benchmark test—without adjusting for any covariates—or the standard outcome test must yield correct conclusions. Thus, when *both* the benchmark and outcome tests indicate discrimination, that conclusion must be correct. The MLRP

is a widely applied assumption on signal distributions in information economics (e.g., Athey and Levin, 2018; Grossman and Hart, 1992; Milgrom, 1981; Persico, 2000), as well as in the outcome test literature (e.g., Feigenberg and Miller, 2022; Anwar and Fang, 2006). We expect the MLRP to hold when it is similarly difficult to make accurate decisions for members of each group (e.g., when the group-specific risk distributions have similar variances, as in Figure 6.1). Drawing on data from lending, education, and criminal justice, we present empirical evidence that the MLRP is approximately satisfied in several important domains. We further show that our hybrid test is robust to the moderate violations of the MLRP that we observe in our data. Applying this approach to 2.8 million police stops across 56 law enforcement agencies in California, we find evidence of pervasive discrimination in police searches of Black and Hispanic individuals—a pattern that would have been missed by the standard outcome test.

6.1 Statistical Guarantees

In our running lending example, our robust outcome test suggests discrimination against a group if two conditions hold simultaneously: (1) lending rates are *lower* for that group (the benchmark test), and (2) repayment rates among loan recipients are *higher* for the group (the standard outcome test). In the stylized example depicted in Figure 6.1, loan recipients in the blue group have higher repayment rates, satisfying the standard outcome test; but members of the blue group are also more likely to receive loans, failing the benchmark test. In this case, whereas the standard outcome test incorrectly infers the blue group is held to a higher, discriminatory lending standard, our robust outcome test correctly concludes that there is insufficient evidence to support a claim of discrimination. We next present formal conditions under which the robust outcome test is guaranteed to produce correct results.

6.1.1 Formal Setup

Our formal setup follows the literature on analyzing outcome tests (see, e.g., Simoiu et al., 2017). We imagine a population of individuals belonging to one of two groups $G \in \{0, 1\}$, indicating, for example, their race or gender. Decision makers take a binary action $D \in \{0, 1\}$ for each individual, such as approving ($D = 1$) or denying ($D = 0$) an individual’s application for a loan. The decision




Figure 6.2: An empirical check of the monotonicity condition of Theorem 6.1 across four domains, comparing the risk distributions of Black vs. White (red) and Hispanic vs. White (blue) individuals. The estimated likelihood ratios all generally increase or decrease monotonically, providing evidence that the MLRP assumption often holds approximately in practice. The dark and light shaded bands (visible only in the recidivism panel) indicate 68% and 95% bootstrapped CIs, respectively.

maker is interested in some binary outcome $Y \in \{0, 1\}$, which, in our running example, corresponds to loan repayment (e.g., $Y = 1$ if the loan is repaid and $Y = 0$ otherwise). The decision maker does not know Y at decision time, but they can estimate it based on the information $X \in \mathcal{X}$ —including group membership G —then available to them about the applicant. In particular, at the moment the decision is made, we assume they can estimate the probability that $Y = 1$ given the available information:

$$R \stackrel{\text{def}}{=} \Pr(Y = 1 \mid X). \quad (6.1)$$

In our running example, R is the decision maker’s estimate of the applicant’s repayment probability. Moreover, the conditional distributions of R by group correspond to the risk distributions in Figure 6.1.

Finally, we assume that decision makers are *rational*, meaning that, within each group, their actions follow threshold rules. (This condition can be relaxed; see Appendix E.1.2.) In particular, we assume they take action $D = 1$ for individuals in group $G = g$ if, and only if, R exceeds some (possibly group-specific) threshold t_g :

$$D \stackrel{\text{def}}{=} \begin{cases} 1 & \text{if } G = g \text{ and } t_g \leq R, \\ 0 & \text{otherwise.} \end{cases} \quad (6.2)$$

Following Becker (1957, 1993), “discrimination” in this setting corresponds to having different group-specific thresholds (i.e., $t_0 \neq t_1$), meaning decision makers apply a double standard. For

instance, in our lending example, $t_1 > t_0$ would mean that decision makers grant loans to members of group $G = 1$ only if they are exceptionally qualified—amounting to discrimination against that group.¹

With this setup, we now state our main technical result.

Theorem 6.1. *Suppose $\Pr(G = 1 \mid R = r)$ is a monotonic function of r , and that, for $g \in \{0, 1\}$, the conditional distribution of $R \mid G = g$ has positive density on $(0, 1)$. Now, if:*

1. $0 < \Pr(D = 1 \mid G = 1) < \Pr(D = 1 \mid G = 0)$, meaning that the decision rate is lower for group $G = 1$ than group $G = 0$; and
2. $\Pr(Y = 1 \mid D = 1, G = 1) > \Pr(Y = 1 \mid D = 1, G = 0)$, meaning that the outcome rate is higher for group $G = 1$ than group $G = 0$;

then $t_0 < t_1$.

Theorem 6.1 shows that under the stated monotonicity assumption—which, we show in Appendix E.1.1, is equivalent to the standard MLRP—a group with both lower decision rates and higher outcome rates is necessarily being held to a higher threshold. For ease of exposition, we present this result for risk distributions with positive densities, threshold decision rules, and binary outcomes, but a much more general version of the result holds. Theorem E.2 in Appendix E.1.2 extends Theorem 6.1, removing the positive density assumption, allowing for quasi-rational decision makers (e.g., with decisions following a logistic curve rather than a threshold function), and incorporating real-valued outcomes (e.g., repayment amounts rather than a binary repayment indicator). See Appendix E.1.1 for the proof of Theorem 6.1 and Appendices E.1.2 to E.1.5 for the general case.

¹As is typical of statistical tests of discrimination, our definition cannot differentiate intentional and unintentional behavior. For instance, different lending thresholds could result from animus, or what is commonly called “disparate treatment” (Huq, 2017). Alternatively, threshold gaps might arise unintentionally, potentially leading to “disparate impact”—unjustified, avoidable disparities (Grossman et al., 2024a). For example, suppose a bank applies a uniform lending threshold to miscalibrated estimates of risk, which consistently underestimate repayment probability for members of one group. Such miscalibration could occur if the bank does not account for varying relationships between covariates and repayment across groups—sometimes called the “subgroup invalidity” problem (Ayres, 2010). Our definition of discrimination considers the bank’s behavior relative to the true (necessarily calibrated) risk in Eq. (6.1). In terms of this risk, the bank’s behavior amounts to applying a double standard, unintentionally but favorably lending to the group whose risk it underestimated. Despite this inherent uncertainty about the underlying mechanism, statistical results like those produced by the robust outcome test can bolster legal claims of both disparate treatment and disparate impact (Goel et al., 2017).

6.2 Assessing Monotonicity

The primary assumption of Theorem 6.1 is that $\Pr(G = 1 \mid R = r)$ is monotonic—which, as discussed above, is equivalent to the group-specific risk distributions satisfying the MLRP. To build intuition about this nonparametric assumption, we consider related parametric conditions on the group-specific risk curves. In particular, a sufficient condition for monotonicity is that the group-specific risk curves are beta distributed with the same total count $\alpha + \beta$ (but possibly different means). More generally, monotonicity holds for betas if and only if the risk curves intersect exactly once, for instance as depicted in Figure 6.1. (See Appendix E.1.4 for more general discussion of parametric conditions that ensure monotonicity and Figure E.1 for an example where the MLRP fails.)

In our running example, equal total count roughly means that it is equally difficult for lenders to distinguish between high- and low-risk applicants across groups. One can imagine that an approximate version of this property holds not only in lending, but across many domains. Indeed, if it fails to hold, one might wonder whether decision makers are ignoring important features to mask discriminatory intent. With redlining, for example, lenders ignored key indicators of individual creditworthiness to justify denying loans to racial minorities (Corbett-Davies et al., 2023).

6.2.1 Empirical evaluation

We explore the extent to which monotonicity holds in practice by considering group-specific empirical risk distributions in four domains, spanning banking, education, and criminal justice. Specifically, we consider: (1) the probability of default among applicants using an online financial technology platform, based on a bevy of traditional and nontraditional variables available to the platform when deciding whom to offer loans; (2) risk of recidivism among defendants awaiting court proceedings, as determined by COMPAS risk scores, which inform judicial bail decisions (Angwin et al., 2022; Corbett-Davies et al., 2017); (3) the risk that pedestrians stopped by the police are carrying contraband, based on indicators such as the reason for the stop and the suspected offense, which inform officer decisions to search stopped individuals (Goel et al., 2016; Gelman et al., 2007); and (4) the probability that law school applicants will pass the bar exam, using their undergraduate grade-point average, LSAT score, and other information available to schools making admissions

decisions (Wightman, 1998). We observe only a proxy of the true outcome of interest—e.g., we see repayment outcomes only among those who received loans, not for the entire population of applicants. Similarly, we do not have the full suite of covariates available to decision makers. As a result, our estimates of risk are approximate. Nonetheless, these estimates give insight into the plausibility of the monotonicity assumption. (See Appendix E.2 for details on the data sources and risk estimation methods.)

For each of these four cases, we plot, in Figure 6.2, $\Pr(G = 1 \mid R = r)$ for Black vs. White individuals (red), and, separately, for Hispanic vs. White individuals (blue). (Here, “White” means non-Hispanic White.) We set $G = 1$ for the smaller group in each comparison—which corresponds to Black or Hispanic individuals except for in our policing example, in which case White individuals are the smaller group. In every instance, we see that the monotonicity condition holds approximately, suggesting that it is, in practice, a relatively mild assumption. Monotonicity, however, does not hold *exactly* in these domains—nor would we expect it to in any real-world dataset. (Violations of the MLRP are visually apparent in Figure 6.2, though formal tests of monotonicity could also be applied (Roosen and Hennessy, 2004).) We thus next conduct a simulation study to assess the robustness of Theorem 6.1 to moderate violations of monotonicity, like those shown in Figure 6.2.

6.2.2 Simulation study

Starting with the empirical risk distributions in the four examples considered above, we evaluate whether the robust and standard outcome tests correctly detect discrimination under a variety of discriminatory and nondiscriminatory scenarios. We find that across scenarios, the robust outcome test is nearly always correct: When it indicates discrimination against a group, that is almost always the correct inference. (Though, as expected, the test sometimes returns an inconclusive result.) In contrast, in these simulations, the standard outcome test often suggests discrimination against the group that in reality received preferential treatment.

For each of our four domains, we first generate synthetic datasets for Black and White individuals based on the estimated risk distributions. (See Figure E.3 for Hispanic individuals.) To do so, for a given pair of hypothetical, group-specific decision thresholds t_g , we repeatedly draw individuals at random from each group, setting $D_i = 1$ if their estimated risk R_i exceeds t_g and setting $Y_i = 1$ with




Figure 6.3: Results of a simulation study comparing the standard and robust outcome tests. The x -axis indicates the decision threshold for White individuals, and the y -axis indicates the decision threshold for Black individuals. The upper-left and lower-right triangular regions correspond to scenarios where decision makers discriminate against either Black or White individuals, as indicated by the “×” and “•” symbols, respectively; nondiscriminatory scenarios are shown by the dashed diagonal line. Red regions indicate where the tests suggest discrimination against Black individuals, blue regions indicate where the tests suggest discrimination against White individuals, and yellow regions indicate where the robust outcome test is inconclusive. The gray areas represent simulation scenarios that are not feasible because a threshold lies outside the support of the risk distribution of the corresponding group. Across simulations, the standard outcome test indicates discrimination when, in reality, there is none—and often indicates discrimination against the group that in actuality was favored. In contrast, the robust outcome test is nearly always directionally accurate, though it sometimes returns an inconclusive result.

probability R_i . We then estimate the decision rate $\widehat{DR}_g$ as the proportion of individuals in group $G = g$ receiving positive decisions $D = 1$; and we estimate the outcome rate $\widehat{OR}_g$ as the proportion of individuals with positive outcomes $Y = 1$, among those with positive decisions.² Finally, we test for discrimination using the robust and standard outcome tests by comparing the decision and outcome rates across groups. In our simulations, we sweep t_g across all percentiles of the risk distributions.³

The results of the simulation are shown in Figure 6.3. As can be seen in the Top panels,

²As the size of the synthetic datasets goes to infinity, the limits of $\widehat{DR}_g$ and $\widehat{OR}_g$ can be computed directly from the risk distributions, without sampling. We take this analytic approach in our simulations for computational efficiency.

³Specifically, we use the quantiles of the average risk distribution across groups, i.e., the distribution with CDF $F(x) \stackrel{\text{def}}{=} \frac{1}{3}[F_B(x) + F_H(x) + F_W(x)]$, where $F_g(x)$ denotes the CDF of Black, Hispanic, and White individuals, respectively. This parameterization ensures that thresholds are not sensitive to the relative sizes of the groups. We exclude the 0th and 100th percentiles.

the robust outcome test is virtually always inconclusive in the absence of discrimination—as we would hope—shown by the yellow region covering the diagonal “no discrimination” line. Moreover, in the off-diagonal regions, where the group-specific thresholds differ, the robust outcome test frequently detects discrimination, and nearly always in the right direction. In contrast, the standard outcome test, as shown in the Bottom panels of Figure 6.3, makes frequent errors, both suggesting discrimination when there is none, as well as indicating discrimination against the group that, in actuality, decision makers favored. Thus, even in these cases where the MLRP does not hold exactly, the robust outcome test still provides correct inferences, and, moreover, outperforms the standard outcome test.

The extent to which the robust outcome test is able to detect discrimination—as opposed to returning an inconclusive result—varies across domains. In particular, when there is a large gap in base rates, it is hard to detect instances of discrimination. In these cases, even when there is (modest) discrimination, the higher base rate group still tends to have both the higher decision rate and the higher outcome rate, yielding an inconclusive result under the robust outcome test. Accordingly, when there is a large gap between base rates, the robust outcome test can only definitively detect more severe instances of discrimination. For instance, the law school admissions example in Figure 6.3 features especially large gaps between the base rates of different groups, resulting in a large inconclusive region. See Appendix E.3 and Figure E.7 for further discussion of the impact of differences in base rates on the sensitivity of the robust outcome test.

In our formal analysis and simulations above, we assume that decision makers are rational within groups, making decisions based on a (potentially group-specific) threshold. In Appendix E.1.2, we relax this assumption and consider quasi-rational decision makers. Simulation results for quasi-rational decision makers show similar patterns; see Figures E.5 and E.6.

6.3 An Application to Police Stops

We conclude our analysis by applying the robust outcome test to data on 2.8 million police stops conducted in 2022 by 56 law enforcement agencies across California. These data were collected as part of California’s Racial Identity and Profiling Act (RIPA) (Grossman et al., 2024a; Guerrero et al., 2024). After an individual is stopped by the police, officers may legally conduct a search of




Figure 6.4: An illustration of the robust outcome test applied to 56 law enforcement agencies across California, with points corresponding to agencies and sized by the number of stops. In each panel, the robust outcome test suggests agencies in the upper-left quadrant discriminated against racial minorities when deciding whom to search, and that agencies in the lower-right quadrant discriminated against White individuals. The test yields inconclusive results for agencies in the white quadrants on the diagonal.

the individual or their vehicle if there is sufficient evidence that the individual possesses contraband. Here we use the robust outcome test to determine whether officers apply the same standard of evidence across racial groups when deciding whom to search.⁴ To do so, we compute, for each jurisdiction, the race-specific search rates and search success rates (i.e., the proportion of searches that resulted in recovery of contraband). If members of a group are both searched *more* often and those searches turn up contraband *less* often, then the robust outcome test indicates that the group was searched according to a lower, discriminatory standard of evidence. (In contrast to our running lending example, where we equated discrimination with a higher lending threshold, discrimination here corresponds to a lower search threshold.)

In the RIPA data, individual-level covariates are recorded selectively (e.g., many covariates are only recorded when a search is conducted). This missingness makes estimating risk difficult (Grossman et al., 2024a), and, consequently, we cannot directly validate the MLRP assumption. We expect,

⁴Following the outcome test literature (e.g., Pierson et al., 2020), here we consider only potential discrimination in *search* decisions—and not in, for example, stop decisions. The main advantages of focusing on search decisions in our setting are two-fold. First, it is clearer what constitutes “success” (i.e., recovery of contraband), facilitating the computation of outcome rates. Second, the pool of individuals subject to the search decision (i.e., individuals stopped by officers) is well defined, facilitating computation of decision rates. In contrast, there is more ambiguity in specifying the set of individuals subject to stop decisions. One could, for example, consider the residential population or, alternatively, the daytime population, and could additionally choose to weight either of these populations by their time outdoors. Compounding this definitional challenge, it can be hard to reliably quantify the size and demographic composition of several of these variants.

however, that the robust outcome test is most useful in precisely these data-limited settings. (If one could accurately estimate risk, other methods may be more appropriate, e.g., risk-adjusted regression (Jung et al., 2024).) Absent direct evidence of the MLRP, the empirical results of Section 6.2—and, in particular, the monotonicity observed in the closely related policing domain considered there—offer reasonable assurances that the MLRP holds approximately for the RIPA data.

We plot the results of our empirical analysis in Figure 6.4, with points corresponding to agencies, sized by the number of recorded stops. Each panel compares stops of White individuals to those of racial minorities (Black or Hispanic individuals, respectively). In each panel, differences between group-specific search rates are plotted on the vertical axis, and differences in search success rates on the horizontal axis. Under the robust outcome test, the red quadrants thus indicate racial discrimination, as those regions contain jurisdictions with both higher search rates and lower search success rates for one of the groups. In particular, the upper-left quadrants suggest discrimination against racial minorities, and the lower-right quadrants suggest discrimination against White individuals. The robust outcome test returns an inconclusive result for agencies in the white, diagonal quadrants, as those correspond to both higher search rates and higher search success rates for one of the groups.

Of the 56 agencies we consider, the robust outcome test suggests discrimination against Black individuals by 33, and discrimination against Hispanic individuals by 32. The test returns an inconclusive result in nearly all of the remaining cases.⁵ The robust outcome test thus suggests a pattern of widespread discrimination against racial minorities in police searches across California.

The standard outcome test, in contrast, suggests White individuals were searched according to a *lower* standard of evidence than Black individuals in about one-third of agencies—corresponding to points in the right-hand quadrants—indicating discrimination against *White* individuals in those

⁵In these inconclusive instances, the standard outcome test generally suggests discrimination against White individuals. In contrast, in the policing simulation shown in Figure 6.3, when the robust outcome test is inconclusive, the standard outcome test generally suggests discrimination against Black individuals. This difference may stem in part from the risk distributions being broadly similar in both our RIPA analysis and policing simulation, but with the high and low base rate groups reversed. In particular, in the policing dataset used for the simulation, stopped White pedestrians are the higher base rate group—likely because of discrimination in the initial stop decisions (Goel et al., 2016). In the RIPA data—primarily consisting of motor vehicle stops, which may be less prone to discrimination than pedestrian stops (Pierson et al., 2020)—stopped White individuals are likely the lower base rate group (Grossman et al., 2024a).

jurisdictions. While not impossible, that result is at odds with an extensive analysis of police discrimination in the literature (Chohlas-Wood et al., 2022; Goel et al., 2016; Grossman et al., 2024a; Jung et al., 2024; Ridgeway, 2007; Epp et al., 2014; Gelman et al., 2007; Goel et al., 2017; Grogger and Ridgeway, 2006; Pierson et al., 2020; Simoiu et al., 2017), pointing to the statistical limitations of the standard outcome test. Due to this lack of face validity, it is easy to dismiss results from the standard outcome test even when it suggests more plausible findings of discrimination against racial minorities, illustrating the value of our robust alternative.

6.4 Discussion

Our empirical analysis of police decisions suggests that the robust outcome test is, in practice, a more accurate barometer of bias than the standard outcome test. Further, it is a logistically straightforward and intuitively appealing method for assessing discrimination. Applying the test requires knowing only group-specific decision and success rates, information that is often readily available in administrative databases. Critically, the robust outcome test avoids omitted-variable bias because it does not use individual-level covariates. Nor does it require decision maker demographics, as other methods do (e.g., Jung et al., 2024; Anwar and Fang, 2006) but which—like detailed covariate information—administrative records often omit.⁶ Further—and in contrast to both the benchmark and standard outcome tests—the robust outcome test is guaranteed to produce correct results under a realistic assumption about the underlying risk distributions. Compared to more statistically sophisticated approaches like the threshold test (Pierson et al., 2018; Simoiu et al., 2017), it requires substantially weaker assumptions to establish correctness.

Our theoretical and empirical results strengthen several past findings in the outcome test literature, where decision rates were reported and are consistent with outcome rates (e.g., Antonovics and Knight, 2009; Close and Mason, 2007; Coviello and Persico, 2015; Persico and Todd, 2006; Ridgeway, 2007; Goel et al., 2017; Pierson et al., 2020). (In these instances, decision rates were reported incidentally or analyzed separately from outcome rates, rather than in the hybrid fashion we suggest.) In many cases, however, researchers simply apply the standard outcome test without reporting

⁶The RIPA data, in addition to the covariate shortcomings noted above, do not have officer demographics, complicating statistical analyses of discrimination that seek to leverage such information.

decision rates (e.g., Anzia and Berry, 2011; Berkovec et al., 1994, 1998; Chilton et al., 2020; Green et al., 2009; Smart and Waldfogel, 1996; Fryer Jr, 2019). Our results thus highlight an important gap in the literature, and suggest a straightforward change to improve current methodological practice.

Despite the benefits of our robust outcome test, it is important to recognize its limitations. First, and most importantly, our proof of correctness rests on a key monotonicity assumption. We presented empirical evidence that this assumption holds approximately in many common cases, and we further showed that, in practice, we obtain correct inferences even when monotonicity does not hold exactly. But the test may yield incorrect results in settings where it is substantially easier to make inferences about one group than another (see, e.g., Figure E.1). Second, like the standard outcome test, computing success rates requires unbiased outcomes. In the policing data we analyzed, it seems likely that our main outcome of interest—contraband recovery—was generally recorded accurately, but that may not always be the case. Third, our robust outcome test can return inconclusive results. In these cases, an absence of evidence of discrimination may stem either from a lack of actual discrimination or from real discrimination that has gone undetected. In particular, as discussed above, the robust outcome test often fails to detect small or moderate threshold gaps when there is a large difference in base rates between groups. We note, though, that in our empirical analysis of police stops, the robust outcome test produced conclusive results in the majority of instances, revealing a pervasive pattern of discrimination. Finally, the robust outcome test—like the standard outcome test—formally produces only a binary determination of discrimination, not a continuous measure of the degree of discrimination.

These limitations suggest promising avenues for future work. In practice, we suspect that greater gaps in decision and success rates point toward greater discrimination. However, rigorously grounding this intuition requires both choosing an appropriate continuous measure of discrimination and, likely, additional parametric assumptions. Relatedly, with stronger, parametric assumptions, one can likely develop variants of the robust outcome test that more often return conclusive results. Lastly, the robust outcome test requires only aggregate information on decision and outcome rates—one of its strengths—but estimates can likely be improved by appropriately leveraging individual-level covariates when they are available. We caution, though, that simply conditioning on the available information can lead to “included-variable bias” (Ayres, 2010; Jung et al., 2024),

masking discrimination.⁷

Recent years have brought renewed urgency to identifying and ameliorating bias in policing and beyond. We hope our work helps further this area of study, both by providing a straightforward and statistically robust method for detecting discrimination, and by offering a blueprint for formally studying empirical tests of bias.

⁷For example, imagine that officers apply a lower bar for searching pedestrians in neighborhoods with a large number of racial minorities—but apply uniform search standards within neighborhoods, regardless of an individual’s race. Then, by conditioning on neighborhood, we might correctly infer there is no discrimination within neighborhoods while failing to capture the larger pattern.

Chapter 7

A Causal Framework for Observational Studies of Discrimination

To assess the role of race or gender in decision making, researchers often examine disparities between groups after adjusting for relevant factors. For example, to measure racial discrimination in lending decisions, one might estimate race-specific approval rates after adjusting for creditworthiness, typically via a regression model. This simple statistical strategy—sometimes called benchmark analysis—has been used to study discrimination in a wide variety of domains, including banking (Munnell et al., 1996), employment (Berg and Lien, 2002), education (Baum and Goodstein, 2005), healthcare (Balsa et al., 2005), housing (Edelman and Luca, 2014; Greenberg et al., 2016), and criminal justice (Ayres, 2002; Fryer Jr, 2019; MacDonald and Raphael, 2021; Rehavi and Starr, 2014; Gelman et al., 2007).

The results of benchmark analyses are often framed in causal terms (e.g., as an effect of race on outcomes), but it is well understood that such an approach suffers from at least three significant statistical challenges when used to estimate causal quantities. First, at a conceptual level, it is unclear how best to rigorously define causal estimands of interest when the treatment is race, gender, or another largely immutable trait. Second, estimates can be plagued by omitted-variable bias if one does not appropriately adjust for all relevant covariates. Third—and the focus of our

Acknowledgments. We thank Alex Chohlas-Wood, Avi Feller, Andrew Gelman, Zhiyuan “Jerry” Lin, Julian Nyarko, Brendan O’Flaherty, Elizabeth Ogburn, José Luis Montiel Olea, Steven Raphael, James Robins, Rajiv Sethi, Amy Shoemaker, and Ilya Shpitser for helpful conversations. Code to replicate our analysis is available online at: <https://github.com/stanford-policylab/gcbsgh-rep>.

article—there are worries that estimates are corrupted by post-treatment bias when one adjusts for covariates or restricts to samples of individuals determined downstream from race, gender, or another such treatment variable. This concern, in particular, has raised doubts about the reliability of the literature on police discrimination, where many studies rely on administrative stop records, and hence, implicitly condition on officers stopping an individual, an action that itself is likely discriminatory (Heckman and Durlauf, 2020; Knox et al., 2020).

Here we present a causal framework for conceptualizing and estimating a measure of discrimination that is suitable for many applied problems. Our framing specifically addresses concerns about post-treatment bias. To do so, we first define a causal quantity—the second-stage sample average treatment effect, or SATE_M —which closely maps to the legal notion of disparate treatment. For this estimand, by carefully considering the timing of events, we show that treatment assignment conceptually occurs after selection into the sample of interest. We then introduce the notion of subset ignorability, show that this condition formally justifies the use of benchmark analysis to estimate the SATE_M , and discuss settings in which it is likely to hold approximately. We illustrate these ideas by analyzing synthetic data, as well as a detailed dataset of prosecutorial charging decisions for approximately 20,000 felony arrests in a major U.S. county. By developing this statistical foundation, we hope to place discrimination studies on more solid theoretical footing.

7.1 A Motivating Example

Consider the problem of measuring racial discrimination in prosecutorial charging decisions. After an individual has been arrested, prosecutors in the district attorney’s office read the arresting officer’s incident report and then decide whether or not to press charges. For simplicity, suppose prosecutors only have access to the incident report—and to no other information—when making their decisions. We allow for the possibility that the arrest decision that preceded the charging decision may have suffered from racial discrimination in complex ways that cannot be inferred from the incident reports themselves. Finally, suppose that a researcher has access to these incident reports for arrested individuals, but, importantly, not to any data on individuals that officers considered but ultimately decided against arresting. What, if anything, might one hope to discover about racial discrimination in charging decisions in light of the fact that the people about whom

the prosecutor makes charging decisions have been selected—that is, arrested—not randomly, but rather in ways that likely depended on their race?

The first challenge is to rigorously define causal estimands of interest. The inherent difficulty is captured by the statistical refrain “no causation without manipulation” (Holland, 1986), since it is often unclear what it means to alter attributes like race and gender (Sekhon, 2008). One common maneuver is to instead consider the causal effect of *perceived* attributes (e.g., perceived race or perceived gender), which ostensibly can be manipulated—for example, by changing the name listed on an employment application (Bertrand and Mullainathan, 2004), or by masking an individual’s appearance (Goldin and Rouse, 2000; Pierson et al., 2020; Grogger and Ridgeway, 2006). In our case, one might imagine a hypothetical experiment in which explicit mentions of race in the incident report are altered (e.g., replacing “white” with “Black”). The causal effect is then, by definition, the difference in charging rates between those cases in which arrested individuals were randomly described (and hence, may be perceived) as “Black” and those in which they were randomly described as “white.” This conceptualization of discrimination conforms to one common causal understanding of discrimination used, for example, in audit studies. This framing also maps closely to the legal notion of disparate treatment, a form of discrimination in which actions are motivated by animus or otherwise discriminatory intent (Goel et al., 2017).

While researchers have carried out such audit studies—including in the case of prosecutorial charging decisions (Chohlas-Wood et al., 2021; Robertson et al., 2019)¹—it is often infeasible to study important policy questions through randomized experiments. In the absence of a controlled experiment, one can in theory identify this type of causal estimand from purely observational data by comparing charging rates across pairs of cases that are identical in all aspects other than the stated race of the arrested individual.² That strategy, which mimics the key features of the hypothetical randomized experiment described above, is formally justified when treatment assignment (i.e., description of race on the incident report, and subsequent perception by the prosecutor) is *ignorable*

¹There are some differences between the idealized audit study described above and these two experiments. Chohlas-Wood et al. conduct a quasi-random field trial in which they mask—but do not switch—the stated race of individuals in police narratives used to make actual charging decisions. Robertson et al. survey prosecutors in a randomized lab experiment and ask them, hypothetically, what their charging decision would be based on fact patterns in which the race of the suspect is manipulated. Although neither of these studies maps exactly to the hypothetical experiment motivating our estimand, both demonstrate the feasibility of conducting such an experiment.

²It suffices to compare groups of cases that have the same distribution of potential outcomes—even if the cases themselves are not identical—a property we formalize in Definition 7.2 below.

given the observed covariates (i.e., features of the incident report) (Imbens and Rubin, 2015). In practice, though, this approach may suffer from omitted-variable bias when the full incident report is not available to researchers, and may suffer from lack of overlap when suitable matches cannot be found for each case—limitations common to many observational studies of causal effects. To address these issues, one can restrict attention to the overlap region and gauge the robustness of estimates to varying forms and degrees of unmeasured confounding (Cinelli and Hazlett, 2020; Cornfield et al., 1959; Rosenbaum and Rubin, 1983a), an approach we demonstrate below.

Finally, there is the issue of post-treatment bias, especially due to sample selection. Knox et al. (2020) argue that researchers often inadvertently introduce post-treatment bias in observational studies of discrimination by subsetting on apparently intermediate outcomes—such as, in our charging example, being arrested—that themselves may be the product of discrimination. As a result, the authors caution that causal quantities of interest cannot be identified by the data in the absence of implausible assumptions, such as lack of discrimination in the initial arrest decision. In making their argument, Knox et al. focus on the use of force by police officers in civilian encounters, but they suggest their formal critique applies more broadly, casting doubt on a wide range of observational studies of discrimination.

Here we show that such customary subsetting does not pose an insurmountable threat to discrimination research. To understand why, one must precisely define the causal estimand, and carefully consider the timing of events. For instance, in our charging example, there are two relevant treatments, the officer’s perception of race, affecting the officer’s arrest decision, and the prosecutor’s perception of race, affecting the prosecutor’s charging decision. The arrest decision is post-treatment relative to the officer’s perception of race but, importantly, it is pre-treatment relative to the prosecutor’s perception of race. Similarly, the features of the incident report—which we must adjust for in this type of benchmark analysis—are post-treatment relative to the officer’s perception of race but pre-treatment relative to the prosecutor’s perception of race. In such a two-decider situation, as Greiner and Rubin (2011) suggest, it is possible to recover estimates of discrimination by the second decider (e.g., in the charging decision) even if there is discrimination by the first decider (e.g., in the arrest decision).

7.2 A Measure of Discrimination

We present a simple two-stage model to characterize discriminatory decision making in a variety of real-world situations and define our main causal quantity of interest—the second-stage sample average treatment effect, or SATE_M —within this general framework. In the context of our motivating example, the SATE_M corresponds to the quantity that would be measured in the hypothetical audit study of prosecutorial decisions described in Section 7.1. A central aim of this article is to formalize technical assumptions that allow one to statistically identify discrimination—more precisely, disparate treatment—in the second stage (e.g., in prosecutorial charging decisions) when data are only available for individuals who made it past the first stage (e.g., those who were arrested). Importantly, our formalization accommodates scenarios in which first-stage decisions may themselves be discriminatory.

In the first stage, we assume each individual in some population is subject to a binary decision M , such as an offer of employment, admission to college, or law enforcement action. Those who receive a “positive” first-stage decision (e.g., those who are arrested) proceed to a second stage, where another binary decision Y is made. In our running example, the case of each arrested individual is reviewed in the second stage by a prosecutor who may or may not choose to press charges. Those who are not arrested do not have a case that requires review by a prosecutor and, indeed, there may be no administrative record of those individuals.

When considering racial discrimination in decisions involving Black and white individuals, our primary quantity of interest is the second-stage sample average treatment effect, $\mathbb{E}[Y(b) - Y(w)]$, where $Y(z)$ indicates the potential second-stage decision and the expectation is taken over individuals reaching the second stage. Here, we imagine that the perception of race is counterfactually determined after the first-stage decision but before the second-stage decision (e.g., after arrest but before charging, perhaps by altering the description of race on the incident report viewed by a prosecutor). The second-stage sample average treatment effect thus, captures discrimination in the second-stage decision among those who made it past the first stage (e.g., discrimination in charging decisions among those who were arrested). This estimand maps onto a common understanding of disparate treatment in second-stage decisions, including in our charging example.

7.2.1 A formal model of discrimination

We now formalize the above discussion to explicitly include decisions made at both the first and second stages. For ease of interpretation, we follow Greiner and Rubin (2011) and motivate our statistical model by considering settings where there are two deciders (e.g., an officer and a prosecutor) whose perceptions of race—or gender, or another trait—can in theory be independently altered prior to their decisions. There are, however, examples in which one can plausibly intervene twice even when a single decider makes both decisions. For instance, an officer may decide to stop a motorist based in part on a brief impression of the motorist’s skin tone as they drive past (Pierson et al., 2020; Grogger and Ridgeway, 2006). This visual impression of race could subsequently be altered if the motorist presents a driver’s license bearing a name characteristic of another race group, or speaks a dialect of English at odds with the officer’s expectation. It thus may be possible to apply our framework whether one imagines there are two deciders or a single one.

We begin by denoting the race of an individual as perceived by the first decider at the first stage by $D \in \{w, b\}$, where, for simplicity, we consider a population consisting of only white and Black individuals. We focus on racial discrimination for concreteness, but similar considerations apply to discrimination based on other attributes, such as gender. Assuming that there is no interference between units (Imbens and Rubin, 2015), we let the binary variables $M(w)$ and $M(b)$ denote the potential first-stage decisions for an individual (e.g., whether they were arrested), and write $M = M(D)$ for the observed first-stage decision. To avoid triviality, we assume throughout that $\Pr(M = 1) > 0$.

Next, we let $Z \in \{w, b\}$ denote the race of an individual as perceived by the second decider, at the second stage. In our running example, Z denotes race as perceived by the prosecutor reviewing that person’s file, while D denotes race as perceived by the police officer during the encounter. Finally, we define the second-stage potential outcomes as a function of both the first-stage outcome M (e.g., the arrest decision) and the second decider’s perception of race Z . Thus, assuming once again that there is no interference, the observed second-stage outcome for an individual can be denoted $Y = Y(Z, M)$, where we consider four potential second-stage outcomes for each individual: $Y(z, m)$, where $z \in \{w, b\}$ and $m \in \{0, 1\}$. In our example, only those who were arrested can be

charged, and so $Y(b, 0) = Y(w, 0) = 0$ for all individuals.³

We further allow each individual to have an associated vector of (non-race) covariates X , representing, for example, their behavior during a police encounter, their recorded criminal history, or both. We imagine these covariates are fixed prior to the second-stage treatment (e.g., prior to the prosecutor’s perception of race), since otherwise the key ignorability assumption in Definition 7.2 below is unlikely to hold. In practice, X is only observed for a subset of the population (e.g., those who were arrested and hence, in the dataset), but we nonetheless define the covariate vector for all individuals in our population of interest. These covariates are not necessary to define our causal estimands of interest, but they play an important role in constructing our statistical estimators.

In this model of discrimination, we have taken care to distinguish between the (realized) first- and second-stage perceptions of race, D and Z , because this helps to clarify the timing of events and the meaning of causal quantities. Importantly, this makes it clear that we can conceive of D and Z as separately manipulable. At the same time, our focus is observational settings, in which disagreement between Z and D may be realized only rarely, if at all, in the data we observe. For instance, barring manipulation of the incident report, it seems unlikely that an arresting officer’s perception of race will frequently differ from a prosecutor’s perception. Our simulation in Section 7.3 thus imposes the further constraint that perceived race is the same at each stage, though this restriction is not necessary in general.

With this framing, we now formally describe the primary causal estimand we consider. This quantity, which we call the second-stage sample average treatment effect (SATE_M) reflects discrimination in the second stage of the decision-making process outlined above, such as discrimination in the prosecutor’s charging decision.⁴

Definition 7.1 (SATE_M). The *second-stage sample average treatment effect*, denoted SATE_M , is:

$$\text{SATE}_M = \mathbb{E}[Y(b, 1) - Y(w, 1) \mid M = 1]. \quad (7.1)$$

³To avoid imagining values of Z for individuals not arrested, one could also make them “missing” by setting $Z = Z(D, M)$, $Z(d, 0) = \text{NA}$, $Z(d, 1) = d$, $Y(\text{NA}, 1) = \text{NA}$, and $Y(z, 0) = 0$ for $z \in \{w, b, \text{NA}\}$, as we do in the simulation in Section 7.3 below. This does not affect any of the mathematical details in what follows.

⁴The SATE_M is notationally equivalent to the $\text{CDE}_{M=1}$ defined in Knox et al. (2020). In our case, however, we have taken care to specify that the first parameter in the quantity $Y(z, m)$ denotes intervening on the *second-stage* perception of race. Moreover, the SATE_M is distinct from what Knox et al. call the $\text{ATE}_{M=1}$.

The estimand in Eq. (7.1) compares the potential second-stage decisions under two race perception scenarios. For example, it compares the potential charging decisions when the prosecutor perceives the individual to be either Black or White; importantly, though, the estimand does not explicitly consider the arresting officer’s perception of race. Moreover, this estimand restricts to the subset of individuals who had a “positive” first-stage decision (e.g., those who were in reality arrested).

Because we condition on $M = 1$ in the definition of the SATE_M , we may equivalently write Eq. (7.1) as

$$\text{SATE}_M = \mathbb{E}[Y(b, M) - Y(w, M) \mid M = 1]. \quad (7.2)$$

We can further write

$$\text{SATE}_M = \mathbb{E}[Y(b) - Y(w) \mid M = 1], \quad (7.3)$$

where we define $Y(z) = Y(z, M)$. Among those who reach the second stage (i.e., individuals with $M = 1$), $Y(z) = Y(z, 1)$ denotes the outcome of intervening *only* on the second decider’s perception of race. Among those who do not reach the second stage (i.e., individuals with $M = 0$), $Y(z) = Y(z, 0) = 0$. Eqs. (7.1), (7.2), and (7.3), as well as the informal estimand introduced at the beginning of Section 7.2, are equivalent ways of capturing the same underlying quantity, varying only in the degree to which they are explicit about the staged nature of the process.

7.2.2 Estimating the sate_M

Having defined the SATE_M , our goal is now to estimate it using only second-stage data. That is, we aim to estimate the SATE_M only using observations for those individuals who received a “positive”—and potentially discriminatory—decision in the first stage. For example, we seek to estimate discrimination in charging decisions based only on data describing those who were arrested. As we show now, an ignorability assumption, together with an overlap condition, is sufficient to guarantee the SATE_M is nonparametrically identified by data on the second-stage decisions.

Definition 7.2 (Subset ignorability). We say that $Y(z, 1)$, Z , M , and X satisfy *subset ignorability* if

$$Y(z, 1) \perp\!\!\!\perp Z \mid X, M = 1 \quad (7.4)$$

for $z \in \{w, b\}$.

In our recurring example, subset ignorability means that among arrested individuals, after conditioning on available covariates, race (as perceived by the prosecutor) is independent of the potential outcomes for the charging decision. As above, we can equivalently write Eq. (7.4) as

$$Y(z) \perp\!\!\!\perp Z \mid X, M = 1. \quad (7.5)$$

This latter expression makes clear that subset ignorability is closely related to the traditional ignorability assumption in causal inference, but where we have explicitly referenced the first-stage outcomes to accommodate a staged model of decision making.

In our prosecutorial setting, subset ignorability would fail if, for example, there were a factor that prosecutors used to make their charging decisions but which was not accounted for in the analysis (e.g., if prosecutors reviewed witness statements that were not in the case files provided to the analyst), and, further, that factor were unbalanced between groups (e.g., if all else equal, witness statements were more commonly available in the cases of white individuals). See Sections 7.3 and 7.4 for further discussion of such unobserved confounders and their statistical consequences.

Almost all causal analyses implicitly rely on a version of subset ignorability, since researchers rarely make inferences about the full population of interest. For example, analyses are typically limited to the individuals who agreed to participate in the study. Even randomized experiments, while ideal for internal validity, frequently lack external validity because the study participants do not resemble a larger population of interest. Whenever ascribing causal interpretations to nonexperimental data, it is important to carefully consider the plausability of ignorability and other assumptions, as we discuss in detail in Sections 7.3 and 7.4 below. We note, though, that the assumptions we rely on are similar to those invoked in nearly every observational study of causal effects.

Ignorability assumptions typically require a corresponding overlap condition to guarantee consistent estimation.⁵

Definition 7.3 (Overlap). We say that *overlap* holds when $\Pr(Z = z \mid X = x, M = 1) > 0$ for all z and x such that $\Pr(X = x, M = 1) > 0$.

Overlap states that there are no covariate levels for which the probability of receiving one of the

⁵In the following, we assume that X is discrete for simplicity of exposition; for continuous analogues of these results, see Appendix F.3.

treatments is zero within the population of interest. In our prosecution example, overlap ensures that every case has a “twin”, identical in all aspects other than the stated race of the arrested individual, against which it can be compared. Overlap would fail, therefore, in the prosecutorial setting, if, for instance, there were alleged offenses for which only Black individuals were arrested. We note that, unlike ignorability, overlap can be assessed directly from the data; see Section 7.4. In cases where overlap fails to hold, one can still elicit valid causal estimates by restricting to the subset of the population where overlap holds. For example, in assessing discrimination in prosecutorial charging decisions, one might only consider those alleged offenses for which both Black and white individuals have a nonzero probability of being arrested. But this restriction comes at the cost of inferential validity for the original population. In such cases, one is estimating the causal effect *only* for the restricted population; the causal effect for the original population may differ, sometimes substantially.

In the traditional, single-stage setting, ignorability and overlap are sufficient to obtain consistent estimates of the average treatment effect. Likewise, we now show that in our two-stage model of discrimination, subset ignorability and overlap are sufficient to guarantee consistent estimates of the SATE_M . In practice, if one can adjust for (nearly) all relevant factors affecting second-stage decisions, one can (approximately) satisfy subset ignorability, and in particular, one can estimate the SATE_M only using data available at the second stage. In the Appendix, we compare subset ignorability to several alternatives, and show that those variants tend either to be too weak to guarantee identifiability, or unnecessarily demanding for real-world applications. We emphasize that since the first-stage decision, M , and the covariates, X , can be viewed as pre-treatment relative to the second-stage intervention, concerns about post-treatment bias corrupting estimates of the SATE_M are more naturally thought of as familiar concerns about omitted-variable bias.⁶

Theorem 7.4. *Suppose $Y(z, 1)$, Z , M , and X satisfy subset ignorability and overlap. Then, the SATE_M equals*

$$\sum_x \mathbb{E}[Y \mid Z = b, X = x, M = 1] \cdot \Pr(X = x \mid M = 1) \\ - \sum_x \mathbb{E}[Y \mid Z = w, X = x, M = 1] \cdot \Pr(X = x \mid M = 1).$$

⁶See Heckman (1979) for related discussion on interpreting sample selection bias as omitted-variable bias.

Proof. Conditioning on X in Eq. (7.1), we have

$$\begin{aligned} \text{SATE}_M &= \sum_x \mathbb{E}[Y(b, 1) \mid X = x, M = 1] \cdot \Pr(X = x \mid M = 1) \\ &\quad - \sum_x \mathbb{E}[Y(w, 1) \mid X = x, M = 1] \cdot \Pr(X = x \mid M = 1). \end{aligned} \quad (7.6)$$

By subset ignorability and overlap, we can condition the summands in Eq. (7.6) on $Z = b$ and $Z = w$, respectively, without changing their values, yielding

$$\begin{aligned} \text{SATE}_M &= \sum_x \mathbb{E}[Y(b, 1) \mid Z = b, X = x, M = 1] \cdot \Pr(X = x \mid M = 1) \\ &\quad - \sum_x \mathbb{E}[Y(w, 1) \mid Z = w, X = x, M = 1] \cdot \Pr(X = x \mid M = 1) \end{aligned} \quad (7.7)$$

$$\begin{aligned} &= \sum_x \mathbb{E}[Y(Z, M) \mid Z = b, X = x, M = 1] \cdot \Pr(X = x \mid M = 1) \\ &\quad - \sum_x \mathbb{E}[Y(Z, M) \mid Z = w, X = x, M = 1] \cdot \Pr(X = x \mid M = 1). \end{aligned} \quad (7.8)$$

Finally, the statement of the proposition follows by consistency, as $Y = Y(Z, M)$. $\square$

Corollary 7.5. *Suppose subset ignorability and overlap hold, and that we have n iid observations $(X_i, Z_i, Y_i)_{i=1}^n$ with $M_i = 1$. Let $S_x^{(n)} = \{1 \leq i \leq n : X_i = x\}$ represent the set of observations with $X = x$, and let $S_{zx}^{(n)} = \{1 \leq i \leq n : Z_i = z \wedge X_i = x\}$ represent the set of observations with $X = x$ and $Z = z$. Then the stratified difference-in-means estimator,*

$$\Delta_n = \sum_x \left[\frac{1}{|S_{bx}^{(n)}|} \sum_{i \in S_{bx}^{(n)}} Y_i \right] \frac{|S_x^{(n)}|}{n} - \sum_x \left[\frac{1}{|S_{wx}^{(n)}|} \sum_{i \in S_{wx}^{(n)}} Y_i \right] \frac{|S_x^{(n)}|}{n}, \quad (7.9)$$

yields a consistent estimate of the SATE_M .

Proof. Note that by the strong law of large numbers,

$$\begin{aligned} \lim_{n \rightarrow \infty} \frac{1}{|S_{zx}^{(n)}|} \sum_{i \in S_{zx}^{(n)}} Y_i &\stackrel{\text{a.s.}}{=} \mathbb{E}[Y \mid Z = z, X = x, M = 1], \text{ and} \\ \lim_{n \rightarrow \infty} \frac{|S_x^{(n)}|}{n} &\stackrel{\text{a.s.}}{=} \Pr(X = x \mid M = 1). \end{aligned}$$

Consequently,

$$\begin{aligned} \lim_{n \rightarrow \infty} \Delta_n \stackrel{\text{a.s.}}{=} & \sum_x \mathbb{E}[Y \mid Z = b, X = x, M = 1] \cdot \Pr(X = x \mid M = 1) \\ & - \sum_x \mathbb{E}[Y \mid Z = w, X = x, M = 1] \cdot \Pr(X = x \mid M = 1), \end{aligned}$$

which is the SATE_M , by Theorem 7.4. $\square$

A straightforward calculation further shows that the following expression yields a consistent estimate of the standard error of Δ_n :

$$\widehat{\text{SE}}(\Delta_n) = \sqrt{\sum_x \left(\frac{|S_x^{(n)}|}{n} \right)^2 \left[\frac{c_{bx}(1 - c_{bx})}{|S_{bx}^{(n)}|} + \frac{c_{wx}(1 - c_{wx})}{|S_{wx}^{(n)}|} \right]}, \quad (7.10)$$

where

$$c_{zx} = \frac{1}{|S_{zx}^{(n)}|} \sum_{i \in S_{zx}^{(n)}} Y_i.$$

Eq. (7.10) accordingly allows us to form confidence intervals for Δ_n .

The nonparametric stratified difference-in-means estimator Δ_n is the basis for nearly all applications of benchmark analysis in discrimination studies. In practice, as we discuss further in Section 7.3, it is common to approximate Δ_n via a parametric regression model—but the two estimators share the same theoretical underpinnings. As such, our analysis above simply grounds traditional benchmark analysis within a specific causal framework, and demonstrates that a particular ignorability assumption, together with overlap, is sufficient to yield valid estimates.

7.2.3 An alternative measure of discrimination

To better understand the SATE_M , we now contrast it with the total effect (TE) (Imai et al., 2010a), a second estimand considered by discrimination researchers (Heckman and Durlauf, 2020; Knox et al., 2020; Zhao et al., 2021). The total effect and the SATE_M differ in our setting in two ways: (a) the population of individuals about which we make inferences; and (b) the potential outcomes being contrasted. The total effect is not restricted to individuals who had a “positive” first-stage decision (e.g., it is not restricted to those who were arrested). Additionally, we imagine a causal variable that reflects a situation where the perception of race is counterfactually determined *before*

the first-stage decision (instead of *after* the first-stage decision, as with the $SATE_M$), and is the same at both stages.

We note that, in general—as discussed in Section 7.1 and below—there is no fully coherent notion of a “total effect” of race, since one cannot intervene on race, *per se*. In our running example, the two treatments (i.e., the officer’s perception of race and the prosecutor’s perception of race) represent distinct, situation-specific notions of intervening on race. In this restricted context, then, there is a natural estimand that captures the spirit of a “total effect”: comparing an individual’s potential outcomes had they been perceived as white or Black when *both* the first- and second-stage decisions were made. We formalize this as follows:

Definition 7.6 (TE). The *total effect*, denoted TE, is given by:

$$TE = \mathbb{E}[Y(b, M(b)) - Y(w, M(w))]. \quad (7.11)$$

Unlike the $SATE_M$, which only measures discrimination in the second decision, the total effect measures cumulative discrimination across *both* decisions. In our recurring example, the total effect captures the effect of race at the time of arrest on the subsequent charging decision. In particular, if a charged Black individual had instead been perceived as white by an officer, they might never have been arrested, and hence, never been at risk of being charged, a possibility encompassed by the definition of the total effect, but not by the $SATE_M$.

We stress, however, that in studies of discrimination—particularly racial discrimination—there is often no clear intervention point, and the difference between the TE and the $SATE_M$ is largely an artifact of how one defines both the population of interest and the start of the decision-making process. What is the TE in one description of events may be the $SATE_M$ in another, equally valid description of the same events, as we describe next.

In our running example, the implicit population of interest consists of those individuals stopped by the police, and the TE reflects a description of events in which the decision-making process starts—and perception of race is counterfactually determined—when the arrest decision is made. We can, however, imagine moving back the clock and starting the process when the stop decision is made, with the population of interest now comprising those individuals spotted by an officer. In this case, the original TE is equivalent to the $SATE_M$ on this newly defined population, where the




Figure 7.1: This figure illustrates estimands one could consider, and the populations they concern, as individuals move through one segment of the criminal justice system. For instance, one can measure combined discrimination in arrest and charging decisions either via the TE or the $SATE_M$. In studies of discrimination, there is no clear point at which race is “assigned” and so both the TE and the $SATE_M$ can be used interchangeably to express the same underlying causal effect, the TE with respect to the population of stopped individuals, and the $SATE_M$ with respect to the population of spotted individuals. More generally, the diagram illustrates a multistage process, where one seeks to measure discrimination culminating at stage t_{k+2} (e.g., charging decisions) among those who make it to stage t_k (e.g., those who were stopped by the police). This quantity can be viewed as the TE, where one imagines the process starting at time t_k . Alternatively, it can be viewed as the $SATE_M$, where one views the process as starting earlier (at, say, t_{k-1} , indicating that an officer spotted an individual), and then conditioning on those who made it to stage t_k . Note that the quantities themselves are formally defined—and equivalent in the manner just described—even absent any considerations of estimation and randomization, which are not illustrated here.

first-stage decision indicates whether an individual was *stopped*. Both the original TE and the new $SATE_M$ capture combined discrimination in the arrest and charging decisions, among the subset of individuals who were stopped.⁷

But the moment when an individual is spotted is no more statistically privileged as a starting point than the moment when an officer makes a stop decision. One could similarly measure cumulative discrimination that includes the stop decision itself, either in terms of the TE or the $SATE_M$. For the TE, as above, we imagine time starting immediately after a potential police encounter, with the

⁷To be explicit, our point is that the original TE and the new $SATE_M$ are the same quantities, and hence are estimable using the same data. However, the new $SATE_M$ (which subsets on individuals who are *stopped* among those who are *spotted*) and the original $SATE_M$ (which subsets on individuals who are *arrested* among those who are *stopped*), are, in contrast, not equal in general, and not necessarily estimable using the same data. In particular, if one wants to estimate either the original TE or, equivalently, the new $SATE_M$, the arrest decision can be viewed as an intermediate variable, and, accordingly, subsetting to arrested individuals would in general introduce post-treatment bias.

first-stage decision indicating whether an individual was stopped (among a population of individuals spotted by the officer). For the $SATE_M$, we back up the clock once again and imagine the first-stage decision indicating whether an individual was spotted by an officer, among an even larger population of people walking through the neighborhood where the officer patrols. Figure 7.1 provides a graphical depiction of this interchangeability.⁸

Although the TE may appear to avoid conditioning on intermediate outcomes, it simply masks a complex chain of events that came before the nominal start of the process, a chain that itself was likely influenced by discriminatory decisions. For instance, the officer spotting and stopping motorists in our running example could be patrolling the neighborhood in question because of its racial composition.⁹ The very idea of “intermediate outcomes”—a concept central to concerns about post-treatment bias—is a slippery notion in the context of discrimination studies, where there is no clear point in time where one can imagine that race is “assigned.” Even birth cannot be considered the ultimate starting point since, in theory, one might include, at the least, the race of a child’s parents, determined at an earlier stage, when assessing discrimination.¹⁰ Indeed, such generational counterfactuals may be critical for understanding systemic, institutional discrimination.

Our discussion of discrimination in multi-staged, multi-decider scenarios applies widely, but it is not universal. In particular, measuring discrimination in a single-decider case—and, specifically, in officer use of force—is challenging. In many of these single-decider scenarios, it is hard to imagine intervening on race after the decision-making process begins, making it difficult to isolate discrimination in later stages.

⁸The formalism above shows a certain statistical equivalence between estimands having different starting points of the decision-making process. Nonetheless, the choice of starting point corresponds to measuring discrimination across different parts of the process, and so different estimands are relevant in different contexts. As such, we do not assert any normative ordering among them.

⁹Importantly, even if the population of individuals spotted by police at a street corner is a (near) random sample of people living or working in the neighborhood, we still cannot think of race as being randomly assigned in that subset. In particular, spotted individuals may still differ on a variety of dimensions (e.g., socio-economic status) across race groups. As such, one would need to statistically account for these differences in any analysis that seeks to measure disparate treatment.

¹⁰In the case of biological sex, one might consider assignment to occur at conception, though that is typically not the primary moment of interest in studies of sex discrimination.

7.3 Assessing Second-Stage Discrimination in a Stylized Scenario

Subset ignorability, in theory, is sufficient to ensure nonparametrically identified estimates of the SATE_M , even when the first-stage decisions are discriminatory. We illustrate that idea by investigating in detail a hypothetical scenario involving discriminatory arrest decisions in the first stage and discriminatory charging decisions in the second stage. We explore the properties of simple estimators in this setting through a simulation study. We demonstrate that failing to adjust for a factor that directly influences charging decisions can result in biased estimates of discrimination in those decisions, but by accounting for all factors that directly influence charging decisions—and hence satisfying subset ignorability—one can accurately estimate the SATE_M , even when there is unmeasured confounding in arrest decisions. This example further clarifies the conceptual importance of distinguishing between an officer’s perception of race and a prosecutor’s perception of race when defining and estimating our quantities of interest.

We consider a hypothetical jurisdiction in which police officers observe the behavior and race of individuals who are potentially engaged in specific criminal activity (e.g., a drug transaction) and then decide whether or not to make an arrest. Subsequently, the case files of arrested individuals—consisting of a written copy of the officer’s description of the encounter and the arrested individual’s criminal history—are brought to a prosecutor who decides whether or not to press charges. We assume the prosecutor only observes the documented race and criminal record of the arrestee, and the arresting officer’s written description of the encounter; accordingly, by construction, the charging decision depends only on these three factors. For example, the prosecutor may choose only to charge individuals who have several previous drug convictions and who were reported to be engaging in a drug transaction. Importantly, while the prosecutor has access to an officer’s written report, the prosecutor does not directly observe the individual’s behavior leading up to the arrest.

Our goal is to estimate discrimination in charging decisions, formalized in terms of the SATE_M . Intuitively, if we observe every arrested individual’s criminal history, race, and officer report, then subset ignorability would hold because the prosecutor’s charging decision depends only on these factors. Thus, with these three covariates, we could generate valid estimates of discrimination in prosecutorial decisions, even without knowing all of the factors that led to an arrest, a decision that may itself have been discriminatory. However, if any of these three covariates—criminal history, race,

or officer report—are unobserved, we will, in general, be unable to accurately assess discrimination in prosecutorial decisions. In both scenarios, with and without unmeasured confounding, our analysis is based on the subpopulation of arrested individuals, where we note that the subsetting (i.e., arrest) is not influenced by the prosecutor’s perception of race. In this setting, the primary concern is thus omitted-variable bias, not post-treatment bias.

We emphasize that we seek only to estimate discrimination in the second-stage charging decision, not cumulative discrimination stemming from both the arrest and charging decisions. In particular, while officer reports may represent an inaccurate—and discriminatory—account of events, such discrimination is distinct from that in the charging decision itself. Similarly, criminal histories reflect a form of complex, long-term discrimination that we do not aim to measure here. Alternative, and more expansive, notions of discrimination are important to understand, but here we focus on assessing the prosecutor’s narrow contribution to inequities at a specific point in the process, a common statistical objective closely tied to policy decisions and legal theories of disparate treatment (Jung et al., 2018).

7.3.1 The data-generating process

We now formally describe the data-generating process for our stylized example. Under the structural causal model we consider, we can both compute the true SATE_M and compute estimates based only on select information available to the prosecutor. In defining the generative process, we closely follow the terminology and conventions of Pearl (2009) and Pearl et al. (2016).¹¹

Our model is defined in terms of the causal directed acyclic graph (DAG) depicted in Figure 7.2. In this model, $S \in \{w, b\}$ indicates one’s self-identified race, and D and Z indicate, respectively, an officer’s and a prosecutor’s perception of race. Further, $M \in \{0, 1\}$ indicates the arrest decision, and $Y \in \{0, 1\}$ indicates the charging decision. Finally, A corresponds to an individual’s behavior, as observed by an officer, and X and R correspond, respectively, to criminal history and an officer’s description of an encounter, as included in the arrest report. For simplicity, in our example these

¹¹In particular, we follow Pearl (2009) in representing unobserved confounding by bidirectional dashed arrows; see Section 1.2.1. We do deviate from Pearl in one aspect of our notation: we write counterfactuals as $Y(z, m)$ instead of $Y_{z, m}(u)$, suppressing the notational dependence on u . The former notation aligns with the popular Rubin-Neyman potential outcome notation that we use when defining the SATE_M . We further note that this SEM is included primarily for illustrative purposes, and consequently contains some simplifications, such as strictly binary covariates. In practice, we recommend reasoning about subset ignorability and its relevant potential outcomes directly.




Figure 7.2: A causal DAG depicting our stylized example of arrest and charging decisions, where D represents the officer’s perception of race, and Z represents the prosecutor’s perception of race. Officer arrest decisions (M) are directly influenced by observed criminal behavior (A) and officer-perceived race (D); the officer reports of the encounters (R) are directly influenced by A and D . Prosecutorial charging decisions are made for all arrested individuals, and are directly influenced by officer reports (R), criminal history (X), and prosecutor-perceived race (Z). Finally, an individual’s self-identified race (S) influences the officer’s perception of race (D), and is confounded with criminal history (X) and behavior (A). We consider two scenarios. The variables highlighted in dark gray (i.e., M , Z , X , and Y) are always observed. In one scenario, the analyst also observes the officer report R , highlighted in light gray, obtaining the full set of information available to the prosecutor; in the other, the analyst does not observe the officer report R (i.e., only M , Z , X , and Y are observed), leading to omitted-variable bias.

latter three variables are operationalized as being binary—for example, one can imagine that X indicates whether an individual had at least one previous drug conviction, A indicates whether they were seen actively engaging in a drug transaction, and R indicates whether they were reported by the officer to be actively engaging in a drug transaction. Officers observe D , A , and R for all individuals; prosecutors observe Z , X , and R only for the subset of arrested individuals. Note that we also allow for Z and R to be missing (i.e., to take the value NA) in cases where an individual is not arrested.

Structural causal models are defined by a set of exogenous random variables and deterministic structural equations specifying the values of all other variables in the DAG. In our example, the

independent exogenous variables are:

$$U_L \sim \text{BERN}(\mu_L),$$

$$U_A, U_M, U_X, U_R, U_Y \sim \text{UNIF}(0, 1),$$

where μ_L is an appropriately defined constant.

We define self-identified race (S), behavior (A), and criminal history (X) in terms of U_L , which captures latent confounding. For constants μ_A , γ , μ_X , and δ , the structural equations for these three variables are given by:

$$f_S(u_L) = \begin{cases} w & u_L = 0, \\ b & u_L = 1, \end{cases}$$

$$f_A(u_A, u_L) = \mathbb{1}(u_A \leq \mu_A + \gamma \cdot u_L),$$

$$f_X(u_X, u_L) = \mathbb{1}(u_X \leq \mu_X + \delta \cdot u_L).$$

This specification allows for the distributions of criminal history and behavior to vary by race due to exogenous factors like disparate police deployment and historical discrimination. For example, stopped Black individuals may be less likely to be engaged in criminal activity than stopped white individuals, corresponding to $\gamma < 0$.

In line with our discussion in Section 7.2.1, we set the prosecutor's perception of race (Z) equal to the officer's perception of race (D), and, for simplicity, we set both equal to one's self-identified race (S). This choice yields the following structural equations:

$$f_D(s) = s,$$

$$f_Z(m, d) = \begin{cases} d & m = 1, \\ \text{NA} & m = 0. \end{cases}$$

Note that, when someone is not arrested, we represent the prosecutor's perception of race as an explicit missing value. The arrest report, R , is treated similarly below.

Finally, for constants α_0 , α_A , α_{black} , λ_0 , λ_A , λ_{black} , β_0 , β_X , β_R , and β_{black} , the structural equations

for arrest decisions (M), police reports (R), and charging decisions (Y) are given by:

$$\begin{aligned}
& f_M(d, a, u_M) \\
& = \mathbb{1}(u_M \leq \alpha_0 + \alpha_A \cdot a + \alpha_{\text{black}} \cdot \mathbb{1}(d = b)), \\
& f_R(d, a, m, u_R) \\
& = \begin{cases} \mathbb{1}(u_R \leq \lambda_0 + \lambda_A \cdot a + \lambda_{\text{black}} \cdot \mathbb{1}(d = b)) & m = 1, \\ \text{NA} & m = 0, \end{cases} \\
& f_Y(z, m, r, x, u_Y) \\
& = \begin{cases} \mathbb{1}(u_Y \leq \beta_0 + \beta_X \cdot x + \beta_R \cdot r + \beta_{\text{black}} \cdot \mathbb{1}(z = b)) & m = 1 \wedge z \neq \text{NA} \wedge r \neq \text{NA}, \\ \text{NA} & m = 1 \wedge (z = \text{NA} \vee r = \text{NA}), \\ 0 & m = 0. \end{cases}
\end{aligned}$$

In particular, arrest decisions and police reports depend on an officer's perception of race, whereas charging decisions depend on a prosecutor's perception of race. This model incorporates both discrimination in arrest decisions, via α_{black} , and discrimination in police reports—e.g., by omitting potentially exculpatory details or by falsifying information—via λ_{black} . Discrimination in charging decisions is encoded by β_{black} .

The above structural equations, together with the distributions on the exogenous variables, fully define the joint distribution of realized and potential outcomes. In particular,

$$\begin{aligned}
S &= f_S(U_L), & D &= f_D(S), \\
Z &= f_Z(M, D), & A &= f_A(U_A, U_L), \\
X &= f_X(U_X, U_L), & M &= f_M(D, A, U_M), \\
R &= f_R(D, A, M, U_R), & Y &= f_Y(Z, M, R, X, U_Y).
\end{aligned}$$

The primary causal quantity we seek to estimate—the SATE_M —is defined in terms of counterfactuals $Y(z, m)$. As discussed in Pearl (2009) and Pearl et al. (2016), such counterfactuals require some care to define, as one must appropriately account for the exogenous variables U . In particular, for the causal DAG in Figure 7.2, the bivariate charge potential outcomes, for counterfactual versions of prosecutor-perceived race, are given by $Y(z, m) = f_Y(z, m, R(m), X, U_Y)$, where

Table 7.1: *A sample of potential and realized outcomes for individuals in our hypothetical example. The data-generating process produces the full set of entries, but the prosecutor only observes the realized outcomes for those who were arrested, indicated by the shaded cells. In the first scenario we consider, the analyst also observes all the information in the shaded cells; in the second scenario, the analyst only observes the information in the dark gray cells (i.e., the analyst does not observe the officer report R), leading to omitted-variable bias.*

S	D	A	X	$M(b)$	$M(w)$	M	Z	R	$R(0)$	$R(1)$	$Y(b, 1)$	$Y(w, 1)$	Y
b	b	0	0	0	0	0	NA	NA	NA	0	0	0	0
b	b	0	1	0	0	0	NA	NA	NA	1	1	0	0
b	b	1	1	1	0	1	b	0	NA	0	1	1	1
w	w	0	0	1	1	1	w	1	NA	1	0	0	0
w	w	0	1	0	0	0	NA	NA	NA	0	0	0	0

$R(m) = f_R(D, A, m, U_R)$ are the counterfactual versions of the officer report. Further, the arrest potential outcomes—where we consider counterfactual versions of officer-perceived race—are given by $M(d) = f_M(d, A, U_M)$. In general, counterfactuals defined in this way obey the consistency rule, meaning that $M = M(D)$ and $Y = Y(Z, M)$.

When $\alpha_{\text{black}} \geq 0$, anyone who would be arrested if white would also be arrested if Black (i.e., $M(b) \geq M(w)$). When $\alpha_{\text{black}} > 0$, we say arrest decisions are discriminatory since, all else being equal, an individual is more likely to be arrested if they were Black than if they were white. Likewise, $Y(b, 1) \geq Y(w, 1)$ when $\beta_{\text{black}} \geq 0$, meaning that an individual who would be charged if arrested and white would also be charged if arrested and Black. We say the charging decision is discriminatory when $\beta_{\text{black}} > 0$.

Features of our data-generating process. Table 7.1 displays a sample of five rows of data generated from our model. From the full set of potential outcomes, we can compute the true SATE_M by directly applying Definition 7.1 to the generated data, taking the average difference between $Y(b, 1)$ and $Y(w, 1)$ among arrested individuals.¹² However, given the simple linear form of our

¹²Because Z and D are separately manipulable in our framing, this quantity—obtained by first subsetting on arrested individuals, and then computing the average difference between potential outcomes—can also be expressed in the *do*-calculus: $\text{SATE}_M = \mathbb{E}[Y \mid \text{do}(Z = b), M = 1] - \mathbb{E}[Y \mid \text{do}(Z = w), M = 1]$. However, as is common in causal mediation analysis, if there were only one indecomposable treatment (e.g., if one instead imagined directly manipulating S) then the corresponding estimand could no longer be expressed using *do*-operations alone (Pearl, 2009, 2015).

structural equations, a straightforward calculation also shows that the SATE_M is exactly equal to β_{black} .

Our hypothetical example captures three key features of real-world discrimination studies. First, prosecutorial records do not contain all information that influenced officers’ first-stage arrest decisions (i.e., prosecutors only observe R , not A). Second, our set-up allows for situations where the arrest decisions are themselves discriminatory—those where $\alpha_{\text{black}} > 0$ —or the officer’s report is discriminatory, e.g., because of omission of exculpatory information or deliberate falsification—those where $\lambda_{\text{black}} > 0$. Third, the prosecutor’s records include the full set of information on which charging decisions are based (i.e., Z , X , and R).

Among those who were arrested, the charging potential outcomes depend only on one’s criminal history (X) and the arrest report (R). In particular, they do not depend on one’s realized, prosecutor-perceived race (Z). Consequently, $Y(z, 1) \perp\!\!\!\perp Z \mid X, R, M = 1$, meaning that the model satisfies subset ignorability relative to X and R . As a result, access to X and R , along with overlap, guarantees the stratified difference-in-means is a consistent estimator of the SATE_M , even if one does not have access to A .¹³ However, in general, $Y(z, 1) \not\perp\!\!\!\perp Z \mid X, M = 1$ (and, likewise, $Y(z, 1) \not\perp\!\!\!\perp Z \mid R, M = 1$), and so if one only has partial information on charging decisions there is no guarantee the SATE_M can be consistently estimated.¹⁴ Indeed, when there is such unmeasured confounding in the prosecutor’s decisions, one should expect biased estimates of the SATE_M .

7.3.2 Estimating the sate_M

Although the data-generating procedure produces the full set of potential outcomes for each individual, the prosecutor only observes a subset of the cells—realized outcomes for arrested individuals, highlighted in gray in Table 7.1. While this circumscribes the causal effects one can estimate—e.g., discrimination by police will no longer be identifiable in the reduced dataset—one can still learn about the SATE_M . We explore the performance of two statistical methods for estimating

¹³In general, first-stage discrimination such as discriminatory arrest decisions or fabrication of evidence in arrest reports does not affect the consistency of the stratified difference-in-means estimator, since subset ignorability will continue to hold. Consistency may fail if discrimination is so extreme that overlap fails, e.g., if no white people are arrested.

¹⁴In the prosecutorial context, sufficiently diligent data gathering can mitigate this possibility; many offices maintain detailed case files, and we make use of such records in our empirical analysis in Section 7.4. In general studies of discrimination, it is important to ensure that decision factors are accurately captured and made available to analysts.

the SATE_M based on data observed by the prosecutor: the stratified difference-in-means estimator described in Eq. (7.9), and a regression-based estimator. We apply each of these methods to two types of data: the full set of information available to prosecutors (i.e., Y , Z , X and R), and an incomplete dataset comprised only of Y , Z , and X (highlighted in dark gray in Table 7.1), in which case we view R as an unmeasured confounder.

One can compute the stratified difference-in-means estimate in three steps. First, partition arrested individuals into subsets that have the same value of the available control variables (i.e., X and R in the complete data setting, and X alone in the partial data setting). Second, on each resulting subset, compute the average difference in charging rates between Black and white individuals. Third, take a weighted average of these differences, where the weights reflect the proportion of arrested individuals in each subset. In addition, one can apply Eq. (7.10) to estimate the standard error of this point estimate to generate confidence intervals.

The stratified difference-in-means estimator is theoretically appealing in that it is guaranteed to yield consistent estimates of the SATE_M when subset ignorability and overlap hold. But the estimator can have high variance when the dimension of the covariate space is high and the sample size is small. Thus, in practice, it is common to model potential outcomes as a function of observed covariates—also known as response surface modeling (Hill, 2011). In particular, on the subset of arrested individuals, one can estimate the SATE_M via a parametric model that estimates observed charging decisions as a function of the available information.

To demonstrate this latter approach, we use a linear probability model. In the complete data setting, we have:

$$\mathbb{E}[Y \mid Z, X, R] = \beta_0 + \beta_1 Z + \beta_2 X + \beta_3 R, \quad (7.12)$$

where the model is fit on the full set of arrests seen by the prosecutor. Under this model, the SATE_M is approximated by the fitted coefficient $\hat{\beta}_1$, since that term captures the difference in charging potential outcomes after adjusting for the observed covariates. For our specific stylized example, the linear regression model in Eq. (7.12) is in fact perfectly specified—exactly mirroring the prosecutor’s charging decisions—and so we are guaranteed to obtain statistically consistent estimates. In the partial data setting, where an analyst only has access to X , one must fit a reduced model that

excludes R :

$$\mathbb{E}[Y \mid Z, X] = \beta_0 + \beta_1 Z + \beta_2 X. \quad (7.13)$$

In this case, $\hat{\beta}_1$ in general yields a biased estimate of the SATE_M , because of the omitted variable R . The stratified difference-in-means estimator will in general similarly yield a biased estimate of the SATE_M in this omitted-variable setting.

7.3.3 Simulation results

We perform a simulation study to understand the properties of the above estimators, varying our assumptions about discrimination and confounding. We simulate 10,000 datasets of size 100,000 for each of 25 different parameter settings. Each setting is defined as a combination of our two key discrimination parameters, α_{black} and β_{black} , where each parameter is allowed to take one of five values: 0.20, 0.25, 0.30, 0.35, and 0.40. Across all simulation settings, we assume the population of individuals encountered by police is 30% Black (i.e., $\mu_L = 0.3$); that 30% of white individuals and 40% of Black individuals have a past drug conviction, indicated by X ; and that 30% of white individuals and 20% of Black individuals are seen engaging in a drug transaction, indicated by A .¹⁵ These settings allow for a substantial amount of overlap across race groups with regard to the key covariates.

On each synthetic dataset, we estimate the SATE_M using both the stratified difference-in-means estimator and the regression-based estimator, and compare the results to the true population-level SATE_M in two scenarios. To illustrate the impact of omitted-variable bias, in the first scenario, we assume the officer’s report R is unavailable—meaning there is unmeasured confounding—and therefore only stratify based on X in the difference-in-means estimator, and fit the model in Eq. (7.13) for the regression-based estimator. In the second scenario, we assume that R is available, and stratify on both X and R in the difference-in-means estimator, and fit the model in Eq. (7.12) for the regression-based estimator. For each combination of α_{black} and β_{black} , the estimates on the 10,000 synthetic datasets yield the approximate sampling distributions for the difference-in-means

¹⁵More specifically, the full set of parameters in our simulation was set as follows: $\mu_L = 0.3, \mu_X = 0.3, \mu_A = 0.3, \delta = 0.1, \gamma = -0.1, \alpha_0 = 0.1, \alpha_A = 0.3, \alpha_{\text{black}} \in \{0.2, 0.25, 0.3, 0.35, 0.4\}, \lambda_0 = 0.2, \lambda_A = 0.6, \lambda_{\text{black}} = 0.1, \beta_0 = 0.2, \beta_X = 0.4, \beta_R = 0.2$, and $\beta_{\text{black}} \in \{0.2, 0.25, 0.3, 0.35, 0.4\}$.




Figure 7.3: *In our hypothetical example of officer and prosecutor behavior, estimates of discrimination in charging decisions are biased when information directly influencing those decisions—in this case, an officer’s report—is omitted (left). However, one can obtain accurate estimates of discrimination when accounting for all information directly influencing charging decisions (right). Each plot shows the results of 10,000 simulations for each of 25 different combinations of discrimination in officer and prosecutor decisions, given by α_{black} and β_{black} , respectively. The true value of the SATE_M , indicated by the horizontal colored lines, is computed based on the full set of potential outcomes for each individual, and does not depend on the degree of discrimination in the first stage, as seen by the constant value of the SATE_M across different values of α_{black} . For each parameter choice, we display the mean of the sampling distribution for the stratified difference-in-means estimator (solid circle) and the regression-based estimator (hollow circle), along with the interval spanned by the 2.5th and 97.5th percentiles of the sampling distribution. In the right plot (“unconfounded”), estimates are based on all three factors that directly influence charging decisions: race, criminal history, and officer report; in the left plot (“confounded”), we omit the report. When all variables directly influencing charging decisions are available, both estimators recover the true value of the SATE_M , even when there is an unknown degree of discrimination in arrest decisions.*

and regression-based estimators. In Figure 7.3, we summarize each sampling distribution by its mean, 2.5th percentile, and 97.5th percentile. The solid points correspond to the difference-in-means estimator, and the hollow points to the regression-based estimator. The horizontal lines indicate the true population-level SATE_M .

In the left panel (“confounded”) of Figure 7.3, the points lie below the horizontal lines in all cases, meaning we underestimate discrimination in charging decisions. In this setting, estimates do not account for the officer reports R , and so there is unmeasured confounding in the charging decisions. We set $\gamma < 0$ in our simulations, and thus, stopped and arrested Black individuals are less

likely to be engaging in criminal activity, a pattern (noisily) reflected in the officer reports. Because we assume these arrest reports are not available for analysis, we cannot fully adjust for their direct influence on prosecutor decisions. As a result, by adjusting for X alone, we miss an important, unmeasured difference between arrested white and Black individuals, leading us to underestimate discrimination in prosecutorial decisions.

In the right panel (“unconfounded”) of Figure 7.3, the points lie on the horizontal lines in all cases, meaning the estimators are unbiased, and the range between the 2.5th and 97.5th percentiles is relatively narrow, indicating estimates are typically close to the true value. These results hold even when one is unable to assess the degree of discrimination α_{black} in the arrest decisions. As implied by Theorem 7.4, to accurately estimate the SATE_M , it is sufficient to measure all covariates that directly influence the prosecutor’s decisions. In practice, it is nearly always impossible to do so perfectly; for instance, decision factors such as forensic evidence may not be readily available, or non-obvious factors, such as the time of day, may play a role in the prosecutor’s charging decision. Thus, it is important to gauge the sensitivity of estimates to unmeasured confounding in those decisions, as we demonstrate with real-world data in Section 7.4 below. The key point is that it is sufficient to adjust for unmeasured confounding in the charging decisions alone; to estimate discrimination in these charging decisions—formalized by the SATE_M —one need not account for unmeasured confounding in either the documents generated by police, such as arrest reports, or the arrest decisions themselves.

Finally, in addition to examining the sampling distributions, we assessed the coverage of our 95% confidence intervals. For the difference-in-means estimator, confidence intervals were constructed via the estimated standard error given by Eq. (7.10); and for the regression-based estimator, we used the conventional OLS estimate of standard error. For each parameter setting, we computed the proportion of confidence intervals for the 10,000 datasets that contained the true value of the SATE_M . In the no-confounding scenario, we found the true coverage was in line with the nominal coverage, ranging from 94% to 96% across parameter specifications. In the confounding scenario, the intervals rarely covered the true values, as expected, with coverage ranging from 1% to 30% across parameters.

7.4 An Empirical Analysis of Prosecutorial Charging Decisions

We now apply the statistical framework developed above to assess possible race and gender discrimination in real-world prosecutorial charging decisions. We start with the set of individuals in a major U.S. county who were arrested for a felony offense between 2013 and 2019. For our race-based analysis, we then limit to the 25,918 instances in which the race of the arrested individual was identified as either Black (14,686) or non-Hispanic white (11,232), and for our gender-based analysis we limit to the 34,871 instances in which the gender of the arrested individual was recorded as either male (29,283) or female (5588).¹⁶

Our dataset includes a variety of information about each case, including the criminal history of the arrested individual; the alleged offenses (e.g., burglary); the location, date, and time of the incident; whether there is body-worn camera footage; whether a weapon was involved; whether an elderly victim was involved; and whether there was gang involvement. (See Appendix F.4 for additional details.) We also know the ultimate charging decision for each case. Disaggregating by gender, 51% of cases involving a male arrestee were charged, compared to 45% of cases involving a female arrestee; and disaggregating by race, 51% of cases involving a Black arrestee were charged, compared to 50% of cases involving a white arrestee.

To gauge the extent to which charging decisions may suffer from disparate treatment by race or gender, we estimate the $SATE_M$. We start by checking that overlap is satisfied for both our race-based and our gender-based analyses. Recall that overlap means $\Pr(Z = z \mid X = x, M = 1) > 0$, where $Z = 1$ indicates an individual’s “treatment” status (i.e., whether an individual is male in our analysis of gender discrimination, or Black in our analysis of racial discrimination), X is a vector of observed case features, and $M = 1$ means we restrict to those individuals who were arrested. In contrast to ignorability, overlap can be assessed directly by examining the data. To do so, we estimate propensity scores (Rosenbaum and Rubin, 1983b), $\Pr(Z = z \mid X = x, M = 1)$, via an L^1 -regularized (lasso) logistic regression model. In Figure 7.4, we plot the distribution of the estimated propensity scores. In the left panel we disaggregate by gender, and in the right panel we disaggregate by race (Black and white). In situations where overlap does not hold, it is common to

¹⁶Both Hispanic and non-Hispanic white individuals in our dataset appear to have been recorded simply as “white”. To disentangle these two categories, we followed past work and imputed Hispanic ethnicity from surnames (Pierson et al., 2020; Word and Perkins, 1996; Word et al., 2008).




Figure 7.4: We plot, for both our gender-based (left) and race-based (right) analyses, the distribution of propensity scores, disaggregated by observed treatment status. We find that the propensity scores are concentrated away from the interval endpoints, satisfying overlap.

restrict one’s analysis to a region of the covariate space where it does hold. In our case, however, the vast majority of the data are already far from the endpoints of the unit interval, so we work with the dataset in its entirety.

As discussed in Section 7.3, regression-based estimators can be viewed as a parametric variant of the stratified difference-in-means estimator Δ_n . Thus, to help account for the high dimensionality of our feature set, we now estimate the SATE_M via linear regression. In particular, for ease of interpretation, we use a linear probability model:

$$\mathbb{E}[Y \mid Z, X] = \beta_0 + \beta_1 Z + \beta_2^T X, \quad (7.14)$$

where Y indicates whether an arrested individual was charged, and X denotes the vector of covariates.

In the gender model, we find that the $\widehat{\text{SATE}}_M$ —as given by $\hat{\beta}_1$ —is 0.025 (95% CI: [0.014, 0.037]); and in the race model, we find that the $\widehat{\text{SATE}}_M$ is -0.008 (95% CI: $[-0.018, 0.002]$). These results indicate that the charging rate for men is slightly higher than the rate for similar women, and that the charging rate for Black individuals is on par with that of similar white individuals, mirroring the patterns we saw with the raw, unadjusted charging rates. If there are no unmeasured confounders (i.e., if subset ignorability holds) and our parametric model is appropriate, these results suggest race and gender have a relatively modest impact on charging decisions in the jurisdiction we consider.

To help contextualize these results, we note that past studies have found mixed evidence of

disparate treatment in prosecutorial charging decisions, likely due in part to differences in the jurisdictions and time periods analyzed, and the methods employed. In one of the most comprehensive investigations to date, Rehavi and Starr (2014) examined nearly 40,000 individuals in the federal criminal justice system from initial arrest to final sentencing. The authors found that disparate treatment in prosecutorial charging decisions—specifically for charges with statutory mandatory minimum sentences—was a primary driver for sentencing disparities between Black and white individuals. In contrast, in a recent experimental study, Robertson et al. (2019) found no evidence of racial bias in charging decisions when they presented prosecutors with vignettes in which the race of the suspect was randomly varied. Similarly, in an observational analysis of prosecutors at the San Francisco District Attorney’s Office, MacDonald and Raphael (2021) found little evidence of discrimination in charging decisions—in fact, the authors found that white individuals were charged slightly more often than similarly situated Black individuals. Finally, in a recent quasi-random study of charging decisions at a large metropolitan district attorney’s office, Chohlas-Wood et al. (2021) similarly found little evidence of disparate treatment.

The AUC of our outcome model in Eq. (7.14) above—fit with all available covariates, including race and gender—is 86%, indicating that it can predict charging decisions well. Our model, however, cannot capture all aspects of prosecutorial decision making, as at least some information used by prosecutors (e.g., forensic evidence) is not recorded in our dataset, meaning that subset ignorability likely does not hold exactly. To check the robustness of our causal estimates to such unmeasured confounding, one may use a variety of statistical methods for sensitivity analysis (Carnegie et al., 2016; Dorie et al., 2016; Franks et al., 2019; Imbens, 2003; Jung et al., 2020; McCandless and Gustafson, 2017; McCandless et al., 2007; Rosenbaum and Rubin, 1983a). At a high level, these methods posit relationships between the unmeasured confounder and both the treatment variable (e.g., race or gender) and the outcome (e.g., the charging decision), and then examine the sensitivity of estimates under the model of confounding.

We apply a technique for sensitivity analysis recently introduced by Cinelli and Hazlett (2020). In brief, their approach bounds the extent to which a coefficient estimate in a linear model—like $\hat{\beta}_1$ in Eq. (7.14)—might change if one were to refit the model including an unmeasured confounder U .




Figure 7.5: *Contour plots describing the sensitivity of the $\widehat{\text{SATE}}_M$ to unmeasured confounding, for our analysis of gender (left) and race (right). The plots indicate the maximum amount the $\widehat{\text{SATE}}_M$ may change under the Cinelli and Hazlett (2020) model of confounding, parameterized by two partial R^2 values. The red curves correspond to a change equalling the magnitude of the $\widehat{\text{SATE}}_M$ estimated from the available data. Thus, an unobserved confounder corresponding to a point above the red curve would be capable of changing the sign of our estimate. To aid interpretation, both plots display the partial R^2 values associated with several observed subsets of covariates.*

More specifically, under the extended model

$$\mathbb{E}[Y \mid Z, X, U] = \beta_0 + \beta_1 Z + \beta_2^T X + \gamma U,$$

Cinelli and Hazlett bound the change in $\hat{\beta}_1$ in terms of two partial R^2 values: $R^2_{Y \sim U|Z,X}$ and $R^2_{Z \sim U|X}$. These two values, respectively, quantify how much residual variance in the outcome Y and treatment Z is explained by U . Formally, $R^2_{Y \sim U|Z,X}$ is defined in terms of the R^2 of two linear regressions: one using all the covariates X , Z , and U to estimate Y (R^2_{full}), and one excluding U (R^2_{red}). Then, $R^2_{Y \sim U|Z,X} = (R^2_{\text{full}} - R^2_{\text{red}}) / (1 - R^2_{\text{red}})$. The quantity $R^2_{Z \sim U|X}$ is defined analogously. As these partial R^2 values increase, so does the amount by which $\hat{\beta}_2$ could change.

The contour plots in Figure 7.5 show the maximum amount by which the $\widehat{\text{SATE}}_M$ may change as a function of $R^2_{Y \sim U|Z,X}$ and $R^2_{Z \sim U|X}$ for our analysis of gender and race—with that change potentially increasing or decreasing the estimate. The red lines trace out values for which the maximum change equals our empirical point estimates of the $\widehat{\text{SATE}}_M$. In particular, an unmeasured confounder lying

above the red line could be sufficient to change the sign of our estimate.

A key hurdle in sensitivity analysis is positing a reasonable range for the strength of a possible unmeasured confounder. To aid interpretation, we compute the partial R^2 values for various subsets of observed covariates, as recommended by Cinelli and Hazlett. For each such subset, we fit the regression model in Eq. (7.14) both with and without that subset, which in turn yields a pair of partial R^2 values for that subset of covariates.

The contour plots in Figure 7.5 contain these reference points for five different subsets of covariates: (a) the subset describing criminal history (e.g., number of prior convictions and number of prior arrests); (b) the alleged offenses (e.g., burglary); (c) the subset of all covariates except for the alleged offenses; (d) the district in which the alleged incident took place; and (e) whether a weapon was alleged to have been used. We find that the partial R^2 values associated with criminal history and whether a weapon was used are below the red curves for both our analysis of gender and race, indicating that a confounder with comparable marginal explanatory power to these covariates would not be sufficient to change the sign of our estimates. However, the partial R^2 values corresponding to the alleged offenses and the district in which the charges were filed are near the red curve for our gender-based analysis and far above the curve for our race-based analysis, meaning that omitting a covariate with similar explanatory power could qualitatively change our conclusions. Furthermore, the partial R^2 values corresponding to everything except the alleged offenses are far above the red curve in both cases, suggesting that an unobserved confounder of similar strength could again substantially alter our results. For instance, in this extreme scenario, inclusion of a currently omitted confounder with similar characteristics in the race-based analysis could yield an estimated treatment effect of more than 13%.

One cannot know the exact nature and impact of unmeasured confounding. Thus, as in many applied statistical problems, we must rely in large part on domain expertise and intuition to form reasonable conclusions. In this case, given the results of our sensitivity analysis, we interpret our empirical findings as providing evidence that perceived gender and race have limited effects on prosecutorial charging decisions in the jurisdiction we consider. As with the SATE_M , our sensitivity analysis is solely focused on discrimination in the charging decision, and, in particular, is not designed to capture the cumulative effects of discrimination stemming from arrests and other earlier decision points.

7.5 Discussion

We have outlined a formal causal framework to ground observational studies of discrimination. We specifically showed that subset ignorability, together with overlap, is sufficient to guarantee that one important causal measure of discrimination (the SATE_M) is nonparametrically identified in a canonical two-stage decision-making setting. *In this context*, we therefore, believe potential issues of post-treatment bias are more appropriately thought of as concerns about omitted variables. Indeed, our treatment of interest—perception of race by the second decision maker—occurs after the subsetting in the first stage, and so it is not post-treatment relative to the selection process. As such, we demonstrated that a traditional regression-based analysis can be used to assess discrimination in real-world prosecutorial charging decisions, even though the underlying arrests may have been discriminatory in unknown ways. In that example—as in many applied settings—subset ignorability may only hold approximately, and our empirical analysis illustrates the importance of sensitivity analysis for robust inference.

Measurements of the SATE_M can be an important step in quantifying discrimination by specific decision makers at specific points in time. In our running example, estimates of the SATE_M can help identify prosecutors who may be making systematically biased charging decisions. Identification of bias, however, is only the first step toward reform. To mitigate identified disparities, one could imagine a variety of interventions, such as training programs (Spencer et al., 2016), or blinding prosecutors to the race of arrested individuals (Chohlas-Wood et al., 2021). As with all interventions, care must be taken to ensure they do not have unintended consequences. Changes in prosecutorial policies could have negative spillover, for example on policing, or unexpected equilibrium effects, such as overall harsher charging decisions.

The SATE_M is but one way to characterize and inform interventions designed to reduce discriminatory behavior. There are at least two broad notions of discrimination, which approximately map to the legal concepts of disparate treatment and disparate impact. Both involve causal interpretations, though with key differences in the definition of the estimand. Disparate treatment concerns the causal effect of race on outcomes—as we formalize here by the SATE_M —with behavior often driven by animus or explicit racial categorization. Disparate impact, on the other hand, concerns the causal effect of policies or practices on unjustified racial disparities, regardless of intent. Disparate

treatment and disparate impact both play important roles in legal and policy discussions, and the perspective one adopts in any given situation affects the choice of statistical estimation strategy and the interpretation of results (Jung et al., 2018).

We have throughout focused on the statistical foundations and measurement of disparate treatment. In our primary example, we estimate—assuming subset ignorability holds—that perceived race and gender have relatively small effects on prosecutorial charging decisions in the jurisdiction we examine. We further demonstrate that these estimates are moderately robust to potential omitted-variable bias. However, that finding, in and of itself, does not mean charging decisions are equitable in a broader sense. Consider, for example, the 1637 cases in our data involving alleged possession of controlled substances by Black or non-Hispanic white individuals. Of these, 748 cases (46%) were ultimately charged, and charging rates by race were nearly identical across race groups, offering little *prima facie* evidence of disparate treatment. However, among the 748 charged cases, 464 (62%) involved a Black individual—far exceeding the proportion of Black residents in the county we study. Charging decisions for these cases thus impose a heavy burden on Black individuals, even if those decisions were not tainted by animus. To the extent that prosecution of drug crimes is misaligned with community goals, these decisions create an unjustified, and discriminatory, disparate impact.

Rigorously estimating discrimination is a daunting task that requires careful consideration. At an empirical level, it is often difficult to obtain detailed data on individual decisions, in which case benchmark analysis may be inadequate—even if coupled with sensitivity analysis. At a theoretical level, we have a limited statistical language to make precise concepts such as animus and implicit bias that are central to discrimination research. Further, as we note above, past work has often framed discrimination as the causal effect of race on behavior, but other conceptions of discrimination, such as disparate impact, are equally important for assessing and reforming practices. Finally, the conclusions of discrimination studies are generally limited to specific decisions that happen within a long chain of potentially discriminatory actions. Quantifying discrimination at any one point (e.g., in charging decisions) does not yield estimates of specific or cumulative discrimination at other points (e.g., in arrest decisions). Despite these important considerations, we hope our work helps place discrimination research on more solid statistical footing, and provokes further interest in the subtle conceptual and methodological issues at the heart of discrimination studies.

Chapter 8

Measuring Racial and Ethnic Disparities in Traffic Enforcement with Large-Scale Telematics Data

Traffic violations are the most common cause of police–civilian contact, with tens of thousands of stops occurring every day in the United States (Langton and Durose, 2013). The ubiquity of traffic stops, in addition to the fines, loss of dignity, and other burdens they place on stopped civilians—particularly people of color—have made them an important object of study in discrimination research (Anwar and Fang, 2006; Epp et al., 2014; Goel et al., 2017; Grogger and Ridgeway, 2006; Huq et al., 2022; Miller, 2008; O’Flaherty and Sethi, 2019; Ouss and Rappaport, 2020; Ridgeway, 2006; Ridgeway and MacDonald, 2009; Rojek et al., 2004; Ryan, 2016; Simoiu et al., 2017; Smith and Petrocelli, 2001; Voigt et al., 2017). However, the lack of a clear benchmark (Gelman et al., 2007) poses a fundamental challenge in quantifying racial and ethnic disparities in policing patterns: traffic violations that do not result in police action are typically not recorded, making it difficult to disentangle the role of differences in driving behavior from differences in police enforcement in creating disparities (Gaebler et al., 2022; Lange et al., 2005).

To overcome this challenge, our analysis uses vehicle telematics data, which include vehicle location and speed over time, collected from drivers nationwide. With these data, we estimate

Acknowledgments. We thank Ryan McMahon and Hari Balakrishnan both for coordinating use of the CMT telematics data, and for helpful conversations throughout the project. We thank Phoebe Barghouty for her assistance filing public records requests to obtain police stop data, and Marissa Gerchick and Keniel Yao for their help preparing these data.

the true underlying rate of speeding violations by geographic area, mitigating the benchmark problem that has stymied past efforts. To complete our analysis, we combined aggregated telematics data from tens of millions of trips with large-scale public records on speeding enforcement drawn from nearly 800,000 speeding stops in 10 American cities to compare speeding enforcement across neighborhoods with different demographic compositions, after adjusting for driving behavior.

To illustrate the benchmark problem, and our approach to addressing it, we first consider Mesa, Arizona. Like many American cities (Charles, 2003), Mesa exhibits patterns of residential segregation, with the city’s racial and ethnic minorities concentrated in the northwest neighborhoods, as shown in Fig. 8.1 A. Using data from the Open Policing Project (Pierson et al., 2020), we plot, in Fig. 8.1 B, speeding stops recorded by the Mesa Police Department between 2014 and 2016, where the area of the gray circles indicates the number of stops at that location. Visually inspecting the map, we see that speeding stops are concentrated in the northwest region of the city, in neighborhoods home to many of the city’s residents of color.

However, without knowing the true rates of speeding violation—enforced or not—in different neighborhoods, we cannot say whether this gap stems from unequal enforcement or simply reflects differences in driving behavior by the city’s residents. For instance, in principle, there may be more speeding violations in Mesa’s northwest neighborhoods, in which case the observed racial and ethnic disparities in traffic stops could result from traffic enforcement that is proportional to traffic violations. Alternatively, however, actual speeding violations in majority minority neighborhoods may be comparable to, or lower than, speeding violations in predominately non-Hispanic white neighborhoods, in which case the disparities we see in speeding stops could reflect differential enforcement.

We disentangle these two possibilities by overlaying, in Fig. 8.1 B, a heatmap of actual speeding, as measured by our telematics data, where darker shades of orange indicate areas with larger numbers of speeding violations. The heatmap shows that speeding in Mesa is distributed across the city, and, importantly, is not restricted to the minority neighborhoods with high concentrations of police stops. In this case, it thus appears that police stops for speeding are driven in part by heavier enforcement in communities of color.




Figure 8.1: For Mesa, Arizona, a comparison of the city’s demographic composition, the location of speeding events, as measured by telematics data, and the location of speeding stops. A: The demographics of residents, where red points represent non-Hispanic white residents and blue points represent non-white residents. B: The density of speeding stops recorded by police officers, indicated by gray dots, where the area of the points is proportional to the number of stops at a single location; stops for speeding are overlaid with a heatmap of speeding events, as estimated from telematics data. Comparing the two panels, we find higher concentrations of recorded speeding stops in neighborhoods that have high proportions of residents of color, despite the fact that actual speeding events are more evenly distributed across the city. Thus, in Mesa, it appears that speeding is more heavily enforced in predominately minority neighborhoods. We note that this pattern does not hold in every city we study, and, in some locations, speeding is disproportionately enforced in predominately white neighborhoods.

8.1 Data and Results

Our discussion of Mesa above demonstrates the importance of accounting for the true rates of speeding when assessing patterns of policing. We now augment this visual analysis with quantitative detail, and extend it to include 10 American cities: Aurora, CO; Chicago, IL; Houston, TX; Madison, WI; Mesa, AZ; Oklahoma City, OK; Plano, TX; San Antonio, TX; Tulsa, OK; and Wichita, KS. These cities were selected because of the availability of both telematics data, and high-resolution geocoded data on police stops for speeding, obtained from the Open Policing Project (Pierson et al., 2020).

Our primary results are based on an aggregated telematics dataset from Cambridge Mobile Telematics (CMT). In total, we analyzed data from approximately 25,000,000 trips taken by 270,000

drivers in 2019 and early 2020—we limited the time range to circumvent the impacts of COVID-19 on driving behavior. For each police beat (or equivalent city-level jurisdiction, such as precinct), CMT computed the proportion of driving time that drivers spent driving 15 KPH (approximately 9 MPH) or more above the speed limit. Table 8.1 details summary population and traffic stop statistics for the 10 cities we analyzed.

One potential concern with using telematics data to estimate the frequency of speeding violations is that the sample of drivers for whom this information is available is not representative of the driving population. Suppose, for example, that telematics data are disproportionately available for drivers in certain neighborhoods. In that case, we might overestimate the frequency of speeding in those neighborhoods for the simple reason that we are measuring the behavior of a disproportionately large number of drivers in that location. To address this issue, we use the telematics data to instead measure the beat-specific rates of speeding, and then scale these estimates by the size of each beat’s residential population. This adjustment corrects for one of the most serious issues of nonrepresentativeness in our analysis. See Materials and Methods for further details and discussion on our approach to estimating total beat-specific driving time and time spent speeding from the telematics data.

Table 8.1: Summary statistics for the cities in our dataset, including the number of police beats, the mean and standard deviation of the residential population of each beat, the number of speeding and overall traffic stops in the city, and the proportion of traffic stops that were for speeding.

City	Beats	Avg. Beat Population	Speeding Stops	Traffic Stops	Prop. Speeding
Aurora	27	12,412 (SD: 5,537)	43,634	101,679	43%
Chicago	268	10,024 (SD: 5,381)	55,303	1,738,393	3%
Houston	98	20,971 (SD: 9,801)	214,226	725,442	30%
Madison	104	2,184 (SD: 1,743)	29,742	107,072	28%
Mesa	32	14,276 (SD: 7,428)	25,298	75,037	34%
Oklahoma City	43	14,032 (SD: 7,551)	156,611	240,214	65%
Plano	23	11,782 (SD: 5,640)	58,243	103,874	56%
San Antonio	110	12,566 (SD: 6,180)	136,003	389,688	35%
Tulsa	44	8,919 (SD: 3,506)	39,394	96,263	41%
Wichita	36	9,782 (SD: 4,530)	83,088	196,881	42%

We begin our analysis by investigating the frequency of speeding in neighborhoods as a function of their demographic composition.¹ Specifically, for each beat in our analysis, we compute—based

¹Throughout our analysis, we use demographic data from the US Census Bureau’s American Community Survey (U.S. Census Bureau, 2013–2018) for 2013 to 2018.




Figure 8.2: The proportion of driving time spent at least 15 KPH above the speed limit versus the proportion of non-white residents, where each point represents a police beat in the 10 cities we analyze. The flat line of best fit (in red) indicates that neighborhood composition and speeding behavior are largely unrelated ($r = 0.04$).

on the telematics data—the proportion of driving time spent at least 15 KPH over the posted speed limit. The results are shown in Fig. 8.2, where points correspond to beats across the 10 cities we consider. The nearly perfectly flat regression line indicates that speeding rates are largely unrelated to the demographic makeup of a neighborhood. As such, any observed disparities in stop rates are likely due to differences in enforcement rather than differences in violation rates.

To estimate the rate of enforcement as a function of beat demographics, after adjusting for speeding, we fit a negative binomial regression model. Specifically, we model the number of speeding stops in each beat in each year in our dataset as:

$$n_i \sim \text{NEGBIN}(\mu_i e^{\beta_{\text{RACE}, \text{CITY}[i]} r_i + \beta_{\text{DRIVING}} d_i + \beta_{\text{YEAR}[i], \text{CITY}[i]}}, \phi), \quad (8.1)$$

where i indexes each observation (a beat in a particular year); n_i is the number of recorded speeding stops; μ_i is the total amount of driving time in the beat (estimated with the telematics data); r_i is the proportion of non-white² residents in the beat and $\beta_{\text{RACE}, \text{CITY}[i]}$ its corresponding coefficient,

²Throughout this paper, we use ‘white’ to refer to non-Hispanic white individuals and ‘non-white’ to refer to American Indian or Alaska Native, Asian, Black or African American, Hispanic or Latino, or Native Hawaiian or




Figure 8.3: The value of $\hat{\beta}_{\text{Race}}$ in our regression for each city, adjusting for the true prevalence of speeding (left) versus raw, unadjusted results (right). In both the adjusted and unadjusted regressions, we find significant heterogeneity across cities, suggesting that the nature and degree of policing practices varies considerably between cities. The mean of the unadjusted city-level coefficients is 0.57 (SE: 0.18, $p = 0.002$), meaning that speeding enforcement tends to be higher in areas with greater proportions of non-white residents. After adjusting for speeding, our results are qualitatively similar, with an average city-level coefficient of 0.34 (SE: 0.16, $p = 0.044$), meaning that disparities in speeding enforcement remain even after adjusting for driving behavior.

where each city has its own coefficient; d_i is the proportion of time spent driving at least 15 KPH above the speed limit—normalized within each city to have mean 0 and variance 1—estimated with the telematics data, and β_{DRIVING} its corresponding coefficient; $\beta_{\text{YEAR}[i], \text{CITY}[i]}$ is a fixed effect to allow for policing practices that vary between years and cities; and ϕ is the dispersion parameter, which controls the variance of the distribution.

The main terms of interest in our model are the values of $\hat{\beta}_{\text{RACE}}$ for each city, which capture the estimated rate of speeding stops in a beat as a function of its proportion of non-white residents after adjusting for driving time and speeding behavior. Positive values of $\hat{\beta}_{\text{RACE}}$ indicate that beats with higher proportions of non-white residents have disproportionately high numbers of speeding stops relative to driving behavior.

The city-level values of $\hat{\beta}_{\text{RACE}}$ are shown in Figure 8.3 (left panel). We find a wide range of coefficients between the cities, from 2.16 (SE: 0.60, $p < 0.001$) in Mesa to -1.18 (SE: 0.23, $p < 0.001$) in Houston. These results suggest that in some cities, such as Mesa, speeding enforcement is higher

other Pacific Islander individuals according to the US Census (omb, 1997).

in neighborhoods with larger non-white populations, whereas in other cities, such as Houston, enforcement is higher in neighborhoods with larger white populations. Taking the unweighted average of the city-level coefficients across the 10 cities we study, we estimate an overall $\hat{\beta}_{\text{RACE}}$ equal to 0.34 (SE: 0.17, $p = 0.044$).³ Thus, on average across the cities we examine, and after adjusting for driving behavior, a beat for which 75% of its residents are non-white experiences approximately 20% more speeding stops than a beat in which 75% of its residents are white. Finally, we note that $\hat{\beta}_{\text{DRIVING}} = 0.54$ (SE: 0.018, $p < 0.001$), suggesting that, on average across the cities we study, enforcement is indeed higher in areas with higher amounts of speeding.

These results are robust to differences in the choice of speeding threshold (i.e. to differences in how one might operationalize speeding behavior, as shown in Figure G.1). We likewise find qualitatively similar results when we repeat our analysis using quasipoisson regression, instead of negative binomial regression (Figure G.2). Further, repeating our analysis with an alternate source of telematics data from TomTom, we find a similar degree of heterogeneity between cities, but somewhat higher values of $\hat{\beta}_{\text{RACE}}$ across cities (see Figures G.13–G.18).

As a point of comparison, the right panel of Fig. 8.3 shows the results of a regression that does not adjust for speeding. The estimated race coefficients are broadly similar to those from the model above that does adjust for speeding, a pattern we would expect given that speeding is largely unrelated to beat demographics, as shown in Fig. 8.2. In other words, raw disparities in enforcement practices largely persist after adjusting for driving behavior (cf. Table G.1).

Our analysis thus far has focused on connections between beat-level driving behavior and beat-level police enforcement, due to the deidentified and aggregated nature of our telematics data, which do not contain driver demographics. However, as we show in Fig. 8.4, the racial composition of drivers stopped in a beat is strongly correlated with the racial composition of the beat’s residential population ($r = 0.83$). These results suggest that when enforcement is concentrated in racially nonrepresentative beats, the enforcement gap translates to racial disparities among stopped drivers.

Our results above indicate that while many cities exhibit racial disparities in speeding enforcement, those disparities do not consistently impact a single demographic group. This pattern appears to

³We computed the standard error using the method of the `sim` function in the `arm` R package (Gelman et al., 2015), which accounts for the covariance structure of the city-level coefficients. See Table G.2 for the complete regression table.




Figure 8.4: For each beat in our dataset, the proportion of non-white residents in the beat versus the proportion of drivers stopped for speeding who were non-white, along with a line of best fit in red. The size of the dots are proportional to the number of drivers stopped for speeding in the beat. We find that the residential demographics of a beat are largely reflected in the demographics of drivers stopped for speeding ($r = 0.83$).

stem in part from the concentration of speeding enforcement in relatively small areas of the cities, and, in many cases, areas that are demographically nonrepresentative of the cities overall. For each of the 10 cities we analyzed, Fig. 8.5 plots the cumulative number of speeding stops in the top k locations, defined in terms of US Census block groups. Across cities, locations accounting for just 10% of the residential population contain between 56% and 76% of speeding stops. This concentration implies that even small deviations in where police departments choose to heavily enforce speeding can lead to substantial demographic disparities, an effect that is exacerbated by geographic segregation within cities. Figures G.3–G.11 visually depict the concentration of speeding stops, and their connection to population demographics, in each of the cities we analyzed.

Our analysis has focused primarily on the enforcement of speeding violations, finding that, in some cities, speeding is more heavily enforced in predominately non-white neighborhoods, while, in others, speeding is more heavily enforced in predominately white neighborhoods. Speeding, however, is just one of many stated reasons for a traffic stop—other common reasons include equipment violations, such as broken taillights. Indeed, we find that the proportion of traffic stops for speeding




Figure 8.5: For each city in our analysis, the cumulative proportion of speeding stops as a function of the cumulative population of the US Census block groups they occur in. We find that speeding stops are concentrated in block groups containing a relatively small proportion of the population. Concretely, block groups containing just 10% of the residential population, indicated by the dashed vertical line, contain between 56% and 76% of the total speeding stops across the cities we study.

varies widely between cities in our analysis, ranging from 3% in Chicago to 65% in Oklahoma City. (Table 8.1 contains statistics for all the cities we consider.)

We conclude by quantifying the aggregate disparities stemming from traffic stops as a whole, not simply stops for speeding. By necessity, this exercise differs from our analysis above, in that we are not able to estimate the underlying prevalence of non-speeding violations. We, thus, cannot disentangle whether any observed disparities are due to differences in enforcement as opposed to differences in rates of infraction. Specifically, we model the number of traffic stops as:

$$v_i \sim \text{NEGBIN}(\mu_i e^{\beta_{\text{RACE}, \text{CITY}[i]} r_i + \beta_{\text{YEAR}[i], \text{CITY}[i]}}, \phi) \quad (8.2)$$

where v_i is the total number of traffic stops in the beat associated with observation i , and the remaining terms are defined as in (8.1).

The city-level values of $\hat{\beta}_{\text{RACE}}$ for the above regression are shown in Figure 8.6. In contrast to our analysis of speeding stops, all 10 cities we examine yield positive coefficients, with statistical significance in eight of the cases. Averaged across cities, we estimate an overall $\hat{\beta}_{\text{RACE}}$ of 1.33 (SE: 0.15, $p < .001$).⁴ While this particular analysis has not accounted for the underlying infraction rates,

⁴As above, we compute the standard error in a manner that accounts for the covariance structure of the city-level coefficients.




Figure 8.6: The city-level values of $\hat{\beta}_{\text{Race}}$ in our regression for all traffic violations. We find uniformly positive coefficients with average 1.44 (SE: 0.14, $p < 0.001$), meaning that in every city we analyze, there were more recorded traffic violations in beats with higher proportions of non-white residents when not adjusting for true prevalence of violations.

we believe it is an important data point for understanding the full impacts of policing, particularly on communities of color.

8.2 Discussion

Combining aggregated telematics data on tens of millions of trips with records on nearly 800,000 police stops for speeding in 10 American cities, we compare the enforcement of speeding violations with the true, underlying violation rates. Averaged across the cities we analyze, and after adjusting for differences in speeding behavior, we find that neighborhoods with higher proportions of non-white residents have greater numbers of speeding stops. We also, however, find substantial variation across cities, and, in several cities, speeding is more heavily enforced in predominantly white neighborhoods. We trace this heterogeneity back in part to the heavy concentration of police stops in a relatively small number of neighborhoods. Given the levels of residential segregation common in many American cities, decisions in where to patrol can, thus, lead to substantial disparities in who is stopped.

One limitation with our methodological approach is that the sample of drivers for whom we have telematics data is not necessarily representative of the driving population as a whole (see Figure G.12). We address this concern in part by reweighting our sample (as described in the Methods and Materials) to better reflect the residential population. It is, however, still possible that CMT users are, on average, safer drivers overall. Our analysis only requires estimates of the relative violation rates across beats, but even this quantity might not be accurately captured by our data. To assess the sensitivity of our results to this possibility, we repeat our analysis on an alternative source of telematics data, from TomTom—described in the Methods and Materials—and obtain broadly similar results. Although it is possible that the TomTom data also suffer from some bias, the concordance of results across datasets provides added evidence for the robustness of our findings.

Our results suggest that, in many cases, speeding enforcement leads to disparate impacts that are not accounted for by underlying driving behavior. That is, while levels of speeding in a location are largely unrelated to the neighborhood’s demographic composition, enforcement of speeding is often concentrated in demographically nonrepresentative areas. It bears emphasizing, though, that such disparate impacts do not imply discriminatory intent. The patterns we see may—at least in part—be the unintended consequence of a complex series of policy choices.

Regardless of motive, the impacts of differential enforcement can often place undue burdens on communities, particularly communities of color (Ayres, 2005; Jung et al., 2018). Aside from the immediate financial costs of speeding tickets, late or missed payments—common among economically disadvantaged subgroups—can lead to substantial fines and fees, revocation of driving privileges, and, in some instances, incarceration (Harris, 2016; Pager et al., 2022). Traffic stops can also escalate to more serious encounters, with about 1 in 15 stops involving use of force (Langton and Durose, 2013). Further, traffic stops can prompt disproportionate enforcement of other laws, with increased scrutiny of stopped drivers potentially uncovering evidence of offenses—such as expired driver’s licenses, minor drug possession, and missed court appointments—that would have otherwise gone unnoticed (Kohler-Hausmann, 2018). Frequent police contact can also deteriorate trust in institutions and exacerbate “system avoidance”, with individuals less willing to seek out government services, contributing to social stratification (Brayne, 2014). In addition to these potential harms of over-policing traffic violations, under-policing can also create equity issues by heightening risks for

both drivers and pedestrians in neighborhoods with lax enforcement (Boehme et al., 2022).

Our results point to the importance of reconsidering current traffic enforcement policies. In some cases, jurisdictions may be able to maintain traffic safety while simultaneously avoiding some of the adverse consequences of policing by adopting less punitive measures—for example, by deterring unsafe driving through speed bumps and roundabouts rather than citations. We further note that a policy or practice of concentrating policing in small regions of a city makes it more likely that the affected drivers are a demographically nonrepresentative subset of the city’s residents. Although such concentration may have certain benefits of efficiency, it is equally important to design enforcement strategies with an eye toward ensuring equity. Some scholars have even argued for more explicitly randomized policing strategies to combat the inherent inequities of targeted action (Harcourt and Meares, 2011).

Finally, it is important to recognize that a city’s physical infrastructure can itself exacerbate disparities attributable to enforcement policies. In particular, there is a long history of building highways, which displace and segregate Black communities in the United States (Karas, 2015; Mohl, 2014); when such discriminatory policies increase distance between those communities and “society’s opportunity” (Archer, 2021; Raphael et al., 2001), they can increase the time Black drivers spend on the road, and thus their potential for being stopped for speeding. Looking forward, we hope our methodological approach and substantive findings help researchers, policymakers, and advocates better evaluate policing practices, and, in turn, design more equitable policies.

8.3 Materials and Methods

Telematics data. We used deidentified and aggregated telematics data from Cambridge Mobile Telematics (CMT). CMT stores data at the level of “trips,” where each trip records a single drive by a single CMT user, and consists of a series of *waypoints* (pairs of latitudes and longitudes) with timestamps recording when drivers passed through given locations. Using these waypoints, one can estimate the speed at which a driver was travelling at each moment. For each beat, CMT aggregated the data to compute the total amount of time drivers spent driving in the beat, and the total amount of time drivers spent driving at various thresholds above the posted speed limit;

the ratio of these terms yields an estimate of the proportion of time drivers spent speeding in the beat—the normalized version of this proportion is what we call d_i in our primary regression. We excluded driving occurring on interstate highways and other freeways and expressways—as defined by the Federal Highway Administration (Department of Transportation, Federal Highway Administration, 2017)—to focus our analysis on local drivers, whose demographic composition more closely resembles the demographic composition of residents of their neighborhood. For the same reason, we also excluded beats which consisted of airports, lakes, and other nonresidential areas.

Key to our analytic approach is estimating the relative amount of time spent driving in each beat— μ_i in our regressions. However, the recorded amount of time CMT users spent in each beat, by itself, can yield biased estimates, especially if the population of CMT users is not a demographically representative sample of the residential population. To adjust for this possibility, CMT reweighted the driving data so that the distribution of drivers’ home beats in the weighted sample mirror the distribution of home beats in the residential population.

Specifically, for each driver k , CMT first imputed their home beat h_k as the beat in which the driver’s car was parked for the longest period of time. In order to ensure that drivers’ home beats could be imputed reasonably accurately, the analysis was restricted to drivers for whom the data included at least 10 trips, comprising 52% of the total number of drivers present in the data, who were collectively responsible for 98% of the total trips present in the data. Then, for each pair of beats (i, j) , CMT computed the average amount of time a_{ij} that drivers with home beat i spent driving in beat j per day. Finally, the adjusted driving time, μ_j , in each beat was computed as follows: $\mu_j = \sum_{i=1}^N a_{ij} \cdot p_i$, where p_i is the residential population of beat i .

Police stop data. We use police stop data from the Open Policing Project (Pierson et al., 2020), a repository of records from 35 municipal police departments across the United States. We filtered to cities where: (1) telematics data were available; (2) the city contained at least five beats; (3) the beat in which a stop occurred can be consistently inferred, either because it is explicitly recorded, or because it can be imputed from recorded latitudes and longitudes; (4) the reason for the stop (e.g. “speeding”) was recorded; and (5) at least 1,000 speeding stops were present in the data. After applying these filters, we were left with the 10 cities comprising our primary analysis.

Mirroring our restriction of the telematics data, we excluded stops occurring on highways.

Robustness checks. We ran a series of robustness checks, described below, to gauge the sensitivity of our results to varying datasets and analytic assumptions. In all cases, we obtained results that were broadly similar to our main results reported above.

First, we examined the robustness of our results to the particular source of telematics data. To do so, we used publicly available data from the TomTom Traffic Stats API. We obtained speed ventiles and the number of vehicles, which passed through each road segment in our 10 cities for the month of April 2019. Due to the aggregated nature of the TomTom data, our estimates of time spent driving and time spent speeding in each beat differed from that in our primary analysis. In particular, we could not infer driver home beats, and consequently could not adjust for potential nonrepresentativeness among the sample of TomTom drivers. We computed an alternative measure of time spent speeding in each beat by taking a weighted average over each road segment in the beat of the proportion of time spent going 15 KPH above the speed limit, weighted by the length of the road segment. We likewise computed an alternative measure of total drive time, $\tilde{\mu}_i = \sum_j n_j l_j$, where j indexes over all road segments within beat i , n_j is the number of vehicles, which passed through road segment j , and l_j is the length of the j th road segment. The city-level estimates of β_{RACE} using the TomTom data are displayed in Figure G.15. The estimated coefficients are qualitatively similar to those found with the CMT data—including the heterogeneity across cities—but the overall average is larger with the TomTom data than with the CMT data, suggesting larger average speeding enforcement disparities in communities of color.

Second, we assess the sensitivity of our results to the precise definition of “speeding.” Driving at any speed above the speed limit is a ticketable offense; however, police may, for instance, have a policy of only pulling over motorists who speed excessively. For this reason, we define “speeding”, in our main analysis, as travelling at least 15 KPH (approximately 9 MPH) over the posted speed limit. Figure G.1 shows that our primary results are consistent over a wide range of thresholds. Similarly, it is possible that police officers target speeders not based on the absolute degree to which they are speeding, but rather how much they are speeding relative to the speed limit. For instance, a motorist driving at 22 MPH on a 15-MPH road may be pulled over, while a motorist driving

62 MPH on a 55-MPH road may not be. Consequently, we also consider relative speeding thresholds, where we consider a driver to be speeding if they exceed a certain percentage of the posted speed limit (e.g. driving more than 10% faster than the limit). Figure G.1 again shows that our main results hold when we define speeding in terms of relative rather than absolute thresholds.

Finally, we re-ran our regression analysis with quasipoisson models, rather than the negative binomial models presented in the main text. The results are shown in Figure G.2. We find that our primary results are qualitatively similar, except that the estimate of β_{RACE} in Houston, 0.036 (SE: 0.181, $p = 0.843$), is no longer statistically significant.

Data availability. Our analysis code, police stop and citation data, and the aggregated TomTom telematics data are publicly available at <https://github.com/stanford-policylab/telematics>. Due to privacy concerns, we are unable to release telematics data from Cambridge Mobile Telematics, but we note that the TomTom telematics data are sufficient to replicate our main results.

Appendix to Chapter 2

A.1 Path-specific Counterfactuals

Constructing policies which satisfy path-specific fairness requires computing path-specific counterfactual values of features. In Algorithm A.1, we describe the formal construction of path-specific counterfactuals $Z_{\Pi,a,a'}$, for an arbitrary variable Z (or collection of variables) in the DAG. To generate a sample $Z_{\Pi,a,a'}^*$ from the distribution of $Z_{\Pi,a,a'}$, we first sample values U_j^* for the exogenous variables. Then, in the first loop, we traverse the DAG in topological order, setting A to a and iteratively computing values V_j^* of the other nodes based on the structural equations in the usual fashion. In the second loop, we set A to a' , and then iteratively compute values $\overline{V}_j^*$ for each node. $\overline{V}_j^*$ is computed using the structural equation at that node, with value $\overline{V}_\ell^*$ for each of its parents that are connected to it along a path in Π , and the value V_ℓ^* for all its other parents. Finally, we set $Z_{\Pi,a,a'}^*$ to $\overline{Z}^*$.

A.2 Constructing Causally Fair Policies

Our aim is to identify the feasible region of expected outcomes attainable via policies which are constructed to satisfy various causal fairness constraints.

First, consider the problem of finding decision policies that maximize expected utility, subject to satisfying a given definition of causal fairness, as well as the outcome and budget constraints. Specifically, letting $\mathcal{C}$ denote the family of all decision policies that satisfy one of the causal fairness

definitions listed above, a utility-maximizing policy d^* is given by

$$\begin{aligned}
d^* \in \arg \max_{d \in \mathcal{C}} \quad & \mathbb{E}[d(X) \cdot u(X)] \\
\text{s.t.} \quad & o_1 - \epsilon \leq \mathbb{E}[d(X) \cdot \mathbb{1}_{\alpha(X)=a_1}] \leq o_1 + \epsilon \\
\text{s.t.} \quad & o_2 - \epsilon \leq \mathbb{E}[d(X) \cdot \mathbb{E}[Y(1) \mid X]] \leq o_2 + \epsilon \\
\text{s.t.} \quad & \mathbb{E}[d(X)] \leq b.
\end{aligned} \tag{A.1}$$

We prove that this optimization problem can be efficiently solved as a single linear program—in the case of counterfactual equalized odds, conditional principal fairness, counterfactual fairness, and path-specific fairness—or as a series of linear programs in the case of counterfactual predictive parity.

Theorem A.1. *Consider the optimization problem given in Eq. (A.1).*

1. *If $\mathcal{C}$ is the class of policies that satisfies counterfactual equalized odds or conditional principal fairness, and the distribution of $(X, Y(0), Y(1))$ is known and supported on a finite set of size n , then a utility-maximizing policy constrained to lie in $\mathcal{C}$ can be constructed via a linear program with $O(n)$ variables and constraints.*
2. *If $\mathcal{C}$ is the class of policies that satisfies path-specific fairness (including counterfactual fairness), and the distribution of $(X, D_{\Pi, A, a})$ is known and supported on a finite set of size n , then a utility-maximizing policy constrained to lie in $\mathcal{C}$ can be constructed via a linear program with $O(n)$ variables and constraints.*
3. *Suppose $\mathcal{C}$ is the class of policies that satisfies counterfactual predictive parity, that the distribution of $(X, Y(1))$ is known and supported on a finite set of size n , and that the optimization problem in Eq. (A.1) has a feasible solution. Further suppose $Y(1)$ is supported on k points, and let $\Delta^{k-1} = \{p \in \mathbb{R}^k \mid p_i \geq 0 \text{ and } \sum_{i=1}^k p_i = 1\}$ be the unit $(k-1)$ -simplex. Then one can construct a set of linear programs $\mathcal{L} = \{L(v)\}_{v \in \Delta^k}$, with each having $O(n)$ variables and constraints, such that the solution to one of the LPs in $\mathcal{L}$ is a utility-maximizing policy constrained to lie in $\mathcal{C}$.*

Before moving on to the proof of Theorem A.1, we note that since the constraints of the linear programs are convex, the feasible regions in Figure 2.4 can be determined by solving the convex

feasibility problem where we impose the additional convex constraint that the expected outcomes—in our admissions example, the aggregate academic index and number of admitted applicants from the target group—lie within some distance ϵ of a given point. Performing a grid search over all points then determines the feasible regions.

Proof. Let $\mathcal{X} = \{x_1, \dots, x_m\}$; then, we seek decision variables d_i , $i = 1, \dots, m$, corresponding to the probability of making a positive decision for individuals with covariate value x_i . Therefore, we require that $0 \leq d_i \leq 1$.

Letting $p_i = \Pr(X = x_i)$ denote the mass of X at x_i , note that the objective function, as well as the outcome and budget constraints are all linear in the decision variables.

1. The objective function $\mathbb{E}[d(X) \cdot u(X)]$ equals $\sum_{i=1}^m d_i \cdot u(x_i) \cdot p_i$
2. The budget constraint $\mathbb{E}[d(X)] \leq b$. constraint equals $\sum_{i=1}^m d_i \cdot p_i \leq b$
3. The first outcome constraint $o_1 - \epsilon \leq \mathbb{E}[d(X) \cdot \mathbb{1}_{\alpha(X)=a_1}] \leq o_1 + \epsilon$ equals $o_1 - \epsilon \leq \sum_{i=1}^m \mathbb{1}_{\alpha(x_i)=a_1} \cdot d_i \cdot p_i \leq o_1 + \epsilon$
4. The second outcome constraint $o_2 - \epsilon \leq \mathbb{E}[d(X) \cdot \mathbb{E}[Y(1) \mid X]] \leq o_2 + \epsilon$ equals $o_2 - \epsilon \leq \sum_{i=1}^m \mathbb{E}[Y(1) \mid X = x_i] \cdot d_i \cdot p_i \leq o_2 + \epsilon$

We now show that each of the causal fairness definitions can be enforced via linear constraints. We do so in three parts as listed in theorem.

Theorem A.1 Part 1. First, we consider counterfactual equalized odds. A decision policy satisfies counterfactual equalized odds when $D \perp\!\!\!\perp A \mid Y(1)$. Since D is binary, this condition is equivalent to the expression $\mathbb{E}[d(X) \mid A = a, Y(1) = y] = \mathbb{E}[d(X) \mid Y(1) = y]$ for all $a \in \mathcal{A}$ and $y \in \mathcal{Y}$ such that $\Pr(Y(1) = y) > 0$. Expanding this expression and replacing $d(x_j)$ by the corresponding decision variable d_j , we obtain that

$$\sum_{i=1}^m d_i \cdot \Pr(X = x_i \mid A = a, Y(1) = y) = \sum_{i=1}^m d_i \cdot \Pr(X = x_i \mid Y(1) = y)$$

for each $a \in \mathcal{A}$ and each of the finitely many values $y \in \mathcal{Y}$ such that $\Pr(Y(1) = y) > 0$. These constraints are linear in the d_i by inspection.

Next, we consider conditional principal fairness. A decision policy satisfies conditional principal fairness when $D \perp\!\!\!\perp A \mid Y(0), Y(1), W$, where $W = \omega(X)$ denotes a reduced set of the covariates X . Again, since D is binary, this condition is equivalent to the expression $\mathbb{E}[d(X) \mid A = a, Y(0) = y_0, Y(1) = y_1, W = w] = \mathbb{E}[d(X) \mid Y(0) = y_0, Y(1) = y_1, W = w]$ for all y_0, y_1 , and w satisfying $\Pr(Y(0) = y_0, Y(1) = y_1, W = w) > 0$. As above, expanding this expression and replacing $d(x_j)$ by the corresponding decision variable d_j yields linear constraints of the form

$$\sum_{i=1}^m d_i \cdot \Pr(X = x_i \mid A = a, S = s) = \sum_{j=1}^m d_j \cdot \Pr(X = x_j \mid S = s)$$

for each $a \in \mathcal{A}$ and each of the finitely many values of $S = (Y(0), Y(1), W)$ such that $s = (y_0, y_1, w) \in \mathcal{Y} \times \mathcal{Y} \times \mathcal{W}$ satisfies $\Pr(Y(0) = y_0, Y(1) = y_1, W = w) > 0$. Again, these constraints are linear by inspection.

Theorem A.1 Part 2. Suppose a decision policy satisfies path-specific fairness for a given collection of paths Π and a (possibly) reduced set of covariates $W = \omega(X)$, meaning that for every $a' \in \mathcal{A}$, $\mathbb{E}[D_{\Pi, A, a'} \mid W] = \mathbb{E}[D \mid W]$.

Recall from the definition of path-specific counterfactuals that

$$D_{\Pi, A, a'} = f_D(X_{\Pi, A, a'}, U_D) = \mathbb{1}_{U_D \leq d(X_{\Pi, A, a'})},$$

where $U_D \perp\!\!\!\perp \{X_{\Pi, A, a}, X\}$. Since $W = \omega(X)$, $U_D \perp\!\!\!\perp \{X_{\Pi, A, a}, W\}$, it follows that

$$\begin{aligned} \mathbb{E}[D_{\Pi, A, a'} \mid W = w] &= \sum_{i=1}^m \mathbb{E}[D_{\Pi, A, a'} \mid X_{\Pi, A, a} = x_i, W = w] \cdot \Pr(X_{\Pi, A, a} = x_i \mid W = w) \\ &= \sum_{i=1}^m \mathbb{E}[\mathbb{1}_{U_D \leq d(X_{\Pi, A, a'})} \mid X_{\Pi, A, a} = x_i, W = w] \cdot \Pr(X_{\Pi, A, a'} = x_i \mid W = w) \\ &= \sum_{i=1}^m d(X_{\Pi, A, a'}) \cdot \Pr(X_{\Pi, A, a'} = x_i \mid W = w) \\ &= \sum_{i=1}^m d_i \cdot \Pr(X_{\Pi, A, a'} = x_i \mid W = w). \end{aligned}$$

An analogous calculation yields that $\mathbb{E}[D \mid W = w] = \sum_{i=1}^m d_i \cdot \Pr(X = x_i \mid W = w)$. Equating

these expressions gives

$$\sum_{i=1}^m d_i \cdot \Pr(X = x_i \mid W = w) = \sum_{i=1}^m d_i \cdot \Pr(X_{\Pi, A, a'} = x_i \mid W = w)$$

for each $a' \in \mathcal{A}$ and each of the finitely many $w \in \mathcal{W}$ such that $\Pr(W = w) > 0$. Again, each of these constraints is linear by inspection.

Theorem A.1 Part 3. A decision policy satisfies counterfactual predictive parity if $Y(1) \perp\!\!\!\perp A \mid D = 0$, or equivalently, $\Pr(Y(1) = y \mid A = a, D = 0) = \Pr(Y(1) \mid D = 0)$ for all $a \in \mathcal{A}$. We may rewrite this expression to obtain:

$$\frac{\Pr(Y(1) = y, A = a, D = 0)}{\Pr(A = a, D = 0)} = C_y,$$

where $C_y = \Pr(Y(1) = y \mid D = 0)$.

Expanding the numerator on the left-hand side of the above equation yields

$$\Pr(Y(1) = y, A = a, D = 0) = \sum_{i=1}^m [1 - d_i] \cdot \Pr(Y(1) = y, A = a, X = x_i)$$

Similarly, expanding the denominator yields

$$\Pr(Y(1) = y, D = 0) = \sum_{i=1}^m [1 - d_i] \cdot \Pr(Y(1) = y, X = x_i).$$

for each of the finitely many $y \in \mathcal{Y}$. Therefore, counterfactual predictive parity corresponds to

$$\sum_{i=1}^m [1 - d_i] \cdot \Pr(Y(1) = y, A = a, X = x_i) = C_y \cdot \sum_{i=1}^m [1 - d_i] \cdot \Pr(Y(1) = y, X = x_i), \quad (\text{A.2})$$

for each $a \in \mathcal{A}$ and $y \in \mathcal{Y}$. Again, these constraints are linear in the d_i by inspection.

Consider the family of linear programs $\mathcal{L} = \{L(v)\}_{v \in \Delta^k}$ where the linear program $L(v)$ has the same objective function, outcome constraint, and budget constraint as before, together with additional constraints for each $a \in \mathcal{A}$ as in Eq. (A.2), where $C_{y_i} = v_i$ for $i = 1, \dots, k$.

By assumption, there exists a feasible solution to the optimization problem in Eq. (A.1), so the solution to at least one program in $\mathcal{L}$ is a utility-maximizing policy that satisfies counterfactual predictive parity. $\square$

A.3 A Stylized Example of College Admissions

In the example that we consider in Section 2.4.1, the exogenous variables in the DAG, $\mathcal{U} = \{u_A, u_D, u_E, u_M, u_T, u_Y\}$, are independently distributed as follows:

$$U_A, U_D, U_Y \sim \text{UNIF}(0, 1),$$

$$U_E, U_M, U_T \sim \mathcal{N}(0, 1).$$

For fixed constants $\mu_A, \beta_{E,0}, \beta_{E,A}, \beta_{M,0}, \beta_{M,E}, \beta_{T,0}, \beta_{T,E}, \beta_{T,M}, \beta_{T,B}, \beta_{T,u}, \beta_{Y,0}, \beta_{Y,D}$, we define the endogenous variables $\mathcal{V} = \{A, E, M, T, D, Y\}$ in the DAG by the following structural equations:

$$f_A(u_A) = \begin{cases} a_1 & \text{if } u_A \leq \mu_A \\ a_0 & \text{otherwise} \end{cases},$$

$$f_E(a, u_E) = \beta_{E,0} + \beta_{E,A} \cdot \mathbb{1}(a = a_1) + u_E,$$

$$f_M(e, u_M) = \beta_{M,0} + \beta_{M,E} \cdot e + u_M,$$

$$f_T(e, m, u_T) = \beta_{T,0} + \beta_{T,E} \cdot e$$

$$+ \beta_{T,M} \cdot m + \beta_{T,B} \cdot e \cdot m + \beta_{T,u} \cdot u_T,$$

$$f_D(x, u_D) = \mathbb{1}(u_D \leq d(x)),$$

$$f_Y(m, u_Y, \delta) = \mathbb{1}(u_Y \leq \text{logit}^{-1}(\beta_{Y,0} + m + \beta_{Y,D} \cdot \delta)),$$

where $\text{logit}^{-1}(x) = (1 + \exp(-x))^{-1}$ and $d(x)$ is the decision policy. In our example, we use constants $\mu_A = \frac{1}{3}$, $\beta_{E,0} = 1$, $\beta_{E,A} = -1$, $\beta_{M,0} = 0$, $\beta_{M,E} = 1$, $\beta_{T,0} = 50$, $\beta_{T,E} = 4$, $\beta_{T,M} = 4$, $\beta_{T,u} = 7$, $\beta_{T,B} = 1$, $\beta_{Y,0} = -\frac{1}{2}$, $\beta_{Y,D} = \frac{1}{2}$. We also assume a budget $b = \frac{1}{4}$.

A.4 Proof of Theorem 2.10

We begin with the following simple lemma.

Lemma A.2. *Suppose $\mathcal{F}$ is a sub- σ -algebra of measurable sets, and suppose X is a non-negative bounded random variable with $X \leq b$ a.s. Then*

$$\frac{\text{VAR}(X \mid \mathcal{F})}{b} \leq \mathbb{E}[X \mid \mathcal{F}].$$

Proof. Note that since $0 \leq X \leq b$ a.s.,

$$\mathbb{E}[X^2 \mid \mathcal{F}] \leq \mathbb{E}[X \mid \mathcal{F}] \cdot b.$$

By Jensen's inequality, $\text{VAR}(X \mid \mathcal{F}) \leq \mathbb{E}[X^2 \mid \mathcal{F}]$, and so, it follows that, a.s.,

$$\frac{\text{VAR}(X \mid \mathcal{F})}{b} \leq \mathbb{E}[X \mid \mathcal{F}],$$

as desired. □

Ignoring the conditioning, Lemma A.2 can be interpreted as saying that the minimum of a bounded random variable cannot be too close to the mean. This fact enables us to prove Theorem 2.10.

Proof of Theorem 2.10. We wish to show that no event E of the form $\{r(X) \geq t\} \subseteq E \subseteq \{r(X) > t\}$ is π -measurable. To that end, it suffices to show that $\mathbb{E}[E \mid \pi(X)] \notin \{0, 1\}$ with positive probability.

Since $\Pr(r(X) = t) = 0$, $\mathbb{1}_{r(X) \geq t} = \mathbb{1}_{r(X) > t}$ a.s., and so it suffices to show that $\mathbb{E}[\mathbb{1}_{r(X) > t} \mid \pi(X)] \notin \{0, 1\}$ with positive probability.

Consider the set of covariates $x = (x_u, a)$ such that $\rho(x)$ lies in the interval $(t, t + \epsilon)$, where, without loss of generality, we assume $t + \epsilon < 1$. Since $\mathbb{E}[r(X) \mid X_u = x_u] = \rho(x) > t$, it follows immediately that for these x , $\Pr(r(X) > t \mid X_u) > 0$. Let $I(x)$ denote the essential infimum of the conditional distribution of $r(X) \mid \rho(X) = \rho(x)$.

Then, we can apply Lemma A.2 to $r(X) - I(X)$, using the fact that $0 \leq r(X) - I(X) \leq 1$ a.s., to obtain that $\rho(X) - I(X) > \epsilon$ a.s. It follows that the event

$$\{\rho(X) \in (t, t + \epsilon), \Pr(r(X) < t \mid \rho(X)) > 0\} \tag{A.3}$$

has positive probability. We can conclude from this that the related event

$$\{\rho(X) \in (t, t + \epsilon), \Pr(r(X) < t \mid \pi(X)) > 0\}$$

cannot have probability zero, since, by the tower law, we would consequently have that, a.s.,

$$0 = \mathbb{1}_{\rho(X) \in (t, t + \epsilon)} \cdot \Pr(r(X) < t \mid \pi(X)),$$

and so, taking conditional expectations with respect to $\rho(X)$, a.s.,

$$0 = \mathbb{1}_{\rho(X) \in (t, t+\epsilon)} \cdot \Pr(r(X) < t \mid \rho(X)),$$

which contradicts Eq. (A.3).

Therefore, for x such that $\rho(x) \in (t, t+\epsilon)$, $\Pr(r(X) < t \mid X_u = x_u) > 0$ and $\Pr(r(X) > t \mid X_u = x_u) > 0$. Since $\Pr(\rho(X) \in (t, t+\epsilon)) > 0$, it follows that $\mathbb{E}[\mathbb{1}_{r(X) > t} \mid \pi(X)] \notin \{0, 1\}$ with positive probability, as desired. $\square$

A.5 Proof of Proposition 2.15

We begin by more formally defining (multiple) threshold policies. We assume, without loss of generality, that $\Pr(A = a) > 0$ for all $a \in \mathcal{A}$ throughout.

Definition A.3. Let $u(x)$ be a utility function. We say that a policy $d(x)$ is a *threshold policy* with respect to u if there exists some t such that

$$d(x) = \begin{cases} 1 & u(x) > t, \\ 0 & u(x) < t, \end{cases}$$

and $d(x) \in [0, 1]$ is arbitrary if $u(x) = t$. We say that $d(x)$ is a *multiple threshold policy* with respect to u if there exist group-specific constants t_a for $a \in \mathcal{A}$ such that

$$d(x) = \begin{cases} 1 & u(x) > t_{\alpha(x)}, \\ 0 & u(x) < t_{\alpha(x)}, \end{cases}$$

and $d(x) \in [0, 1]$ is arbitrary if $u(x) = t_{\alpha(x)}$.

Remark A.4. In general, it is possible for different thresholds to produce threshold policies that are almost surely equal. For instance, if $u(X) \sim \text{BERN}(\frac{1}{2})$, then the policies $\mathbb{1}_{u(X) > p}$ are almost surely equal for all $p \in [0, 1)$. Nevertheless, we speak in general of *the* threshold associated with the threshold policy $d(X)$ unless there is ambiguity.

We first observe that if $\mathcal{U}$ is consistent modulo α , then whether a decision policy $d(x)$ is a multiple threshold policy does not depend on our choice of $u \in \mathcal{U}$.

Lemma A.5. *Let $\mathcal{U}$ be a collection of utilities consistent modulo α , and suppose $d : \mathcal{X} \rightarrow [0, 1]$ is a decision rule. If $d(x)$ is a multiple threshold rule with respect to a utility $u^* \in \mathcal{U}$, then $d(x)$ is a multiple threshold rule with respect to every $u \in \mathcal{U}$. In particular, if $d(x)$ can be represented by non-negative thresholds over u^* , it can be represented by non-negative thresholds over any $u \in \mathcal{U}$.*

Proof. Suppose $d(x)$ is represented by thresholds $\{t_a^*\}_{a \in \mathcal{A}}$ with respect to u^* . We construct the thresholds $\{t_a\}_{a \in \mathcal{A}}$ explicitly.

Fix $a \in \mathcal{A}$ and suppose that there exists $x^* \in \alpha^{-1}(a)$ such that $u^*(x^*) = t_a^*$. Then set $t_a = u(x^*)$. Now, if $u(x) > t_a = u(x^*)$ then, by consistency modulo α , $u^*(x) > u^*(x^*) = t_a^*$. Similarly if $u(x) < t_a$ then $u^*(x) < t_a^*$. We also note that by consistency modulo α , $\text{sign}(t_a) = \text{sign}(u(x^*)) = \text{sign}(u^*(x^*)) = \text{sign}(t_a^*)$.

If there is no $x^* \in \alpha^{-1}(a)$ such that $u^*(x^*) = t_a^*$, then let

$$t_a = \inf_{x \in S_a} u(x)$$

where $S_a = \{x \in \alpha^{-1}(a) \mid u^*(x) > t_a^*\}$. Note that since $\text{sign}(u(x)) = \text{sign}(u^*(x))$ for all x by consistency modulo α , if $t_a^* \geq 0$, it follows that $t_a \geq 0$ as well.

We need to show in this case also that if $u(x) > t_a$ then $u^*(x) > t_a^*$, and if $u(x) < t_a$ then $u^*(x) < t_a^*$. To do so, let $x \in \alpha^{-1}(a)$ be arbitrary, and suppose $u(x) > t_a$. Then, by definition, there exists $x' \in \alpha^{-1}(a)$ such that $u(x) > u(x') > t_a$ and $u^*(x') > t_a^*$, whence $u^*(x) > u^*(x') > t_a^*$ by consistency modulo α . On the other hand, if $u(x) < t_a$, it follows by the definition of t_a that $u^*(x) \leq t_a^*$; since $u^*(x) \neq t_a^*$ by hypothesis, it follows that $u^*(x) < t_a^*$.

Therefore, it follows in both cases that for $x \in \alpha^{-1}(a)$, if $u(x) > t_a$ then $u^*(x) > t_a^*$, and if $u(x) < t_a$ then $u^*(x) < t_a^*$. Therefore

$$d(x) = \begin{cases} 1 & \text{if } u(x) > t_{\alpha(x)}, \\ 0 & \text{if } u(x) < t_{\alpha(x)}, \end{cases}$$

i.e., $d(x)$ is a multiple threshold policy with respect to u . Moreover, as noted above, if $t_a^* \geq 0$ for all $a \in \mathcal{A}$, then $t_a \geq 0$ for all $a \in \mathcal{A}$. □

We now prove the following strengthening of Prop. 2.15.

Lemma A.6. *Let $\mathcal{U}$ be a collection of utilities consistent modulo α . Let $d(x)$ be a feasible decision policy that is not a.s. a multiple threshold policy with non-negative thresholds with respect to $\mathcal{U}$, then $d(x)$ is strongly Pareto dominated.*

Proof. We prove the claim in two parts. First, we show that any policy that is not a multiple threshold policy is strongly Pareto dominated. Then, we show that any multiple threshold policy that cannot be represented with non-negative thresholds is strongly Pareto dominated.

If $d(x)$ is not a multiple threshold policy, then there exists a $u \in \mathcal{U}$ and $a^* \in \mathcal{A}$ such that $d(x)$ is not a threshold policy when restricted to $\alpha^{-1}(a^*)$ with respect to u .

We will construct an alternative policy $d'(x)$ that attains strictly greater utility on $\alpha^{-1}(a^*)$ and is identical elsewhere. Thus, without loss of generality, we assume there is a single group, i.e., $\alpha(x) = a^*$. The proof proceeds heuristically by moving some of the mass below a threshold to above a threshold to create a feasible policy with improved utility.

Let $b = \mathbb{E}[d(X)]$. Define

$$\begin{aligned} m^{\text{Lo}}(t) &= \mathbb{E}[d(X) \cdot \mathbb{1}_{u(X) < t}], \\ m^{\text{Up}}(t) &= \mathbb{E}[(1 - d(X)) \cdot \mathbb{1}_{u(X) > t}]. \end{aligned}$$

We show that there exists t^* such that $m^{\text{Up}}(t^*) > 0$ and $m^{\text{Lo}}(t^*) > 0$. For, if not, consider

$$\tilde{t} = \inf\{t \in \mathbb{R} : m^{\text{Up}}(t) = 0\}.$$

Note that $d(X) \cdot \mathbb{1}_{u(X) > \tilde{t}} = \mathbb{1}_{u(X) > \tilde{t}}$ a.s. If $\tilde{t} = -\infty$, then by definition $d(X) = 1$ a.s., which is a threshold policy, violating our assumption on $d(X)$. If $\tilde{t} > -\infty$, then for any $t' < \tilde{t}$, we have, by definition that $m^{\text{Up}}(t') > 0$, and so by hypothesis $m^{\text{Lo}}(t') = 0$. Therefore $d(X) \cdot \mathbb{1}_{u(X) < \tilde{t}} = 0$ a.s., and so, again, $d(X)$ is a threshold policy, contrary to hypothesis.

Now, with t^* as above, for notational simplicity, let $m^{\text{Up}} = m^{\text{Up}}(t^*)$ and $m^{\text{Lo}} = m^{\text{Lo}}(t^*)$ and consider the alternative policy

$$d'(x) = \begin{cases} (1 - m^{\text{Up}}) \cdot d(x) & u(x) < t^*, \\ d(x) & u(x) = t^*, \\ 1 - (1 - m^{\text{Lo}}) \cdot (1 - d(x)) & u(x) > t^*. \end{cases}$$

Then it follows by construction that

$$\begin{aligned}
\mathbb{E}[d'(X)] &= (1 - m^{\text{Up}}) \cdot m^{\text{Lo}} + \mathbb{E}[d(X) \cdot \mathbb{1}_{u(X)=t^*}] + \Pr(u(X) > t^*) - (1 - m^{\text{Lo}}) \cdot m^{\text{Up}} \\
&= m^{\text{Lo}} + \mathbb{E}[d(X) \cdot \mathbb{1}_{u(X)=t^*}] + \Pr(u(X) > t^*) - m^{\text{Up}} \\
&= \mathbb{E}[d(X) \cdot \mathbb{1}_{u(X) < t^*}] + \mathbb{E}[d(X) \cdot \mathbb{1}_{u(X)=t^*}] + \mathbb{E}[\mathbb{1}_{u(X) > t^*}] - \mathbb{E}[(1 - d(X)) \cdot \mathbb{1}_{u(X) > t^*}] \\
&= \mathbb{E}[d(X)] \\
&= b,
\end{aligned}$$

so $d'(x)$ is feasible. However,

$$d'(x) - d(x) = m^{\text{Lo}} \cdot (1 - d(x)) \cdot \mathbb{1}_{u(x) > t^*} - m^{\text{Up}} \cdot d(x) \cdot \mathbb{1}_{u(x) < t^*},$$

and so

$$\begin{aligned}
\mathbb{E}[(d'(X) - d(X)) \cdot u(X)] &= m^{\text{Lo}} \cdot \mathbb{E}[(1 - d(X)) \cdot \mathbb{1}_{u(X) > t^*} \cdot u(X)] \\
&\quad - m^{\text{Up}} \cdot \mathbb{E}[d(X) \cdot \mathbb{1}_{u(X) < t^*} \cdot u(X)] \\
&> m^{\text{Lo}} \cdot t^* \cdot \mathbb{E}[(1 - d(X)) \cdot \mathbb{1}_{u(X) > t^*}] - m^{\text{Up}} \cdot t^* \cdot \mathbb{E}[d(X) \cdot \mathbb{1}_{u(X) < t^*}] \\
&= t^* \cdot m^{\text{Lo}} \cdot m^{\text{Up}} - t^* \cdot m^{\text{Up}} \cdot m^{\text{Lo}} \\
&= 0.
\end{aligned}$$

Therefore

$$\mathbb{E}[d(X) \cdot u(X)] < \mathbb{E}[d'(X) \cdot u(X)].$$

It remains to show that $u'(d') > u'(d)$ for arbitrary $u' \in \mathcal{U}$. Let

$$t' = \inf\{u'(x) : d'(x) > d(x)\}.$$

Note that by construction for any $x, x' \in \mathcal{X}$, if $d'(x) > d(x)$ and $d'(x') < d(x')$, then $u(x) > t^* > u(x')$.

It follows by consistency modulo α that $u'(x) \geq t' \geq u'(x')$, and, moreover, that at least one of the inequalities is strict. Without loss of generality, assume $u'(x) > t' \geq u'(x')$. Then, we have that $u(x) > t^*$ if and only if $u'(x) > t'$. Therefore, it follows that

$$\mathbb{E}[(d'(X) - d(X)) \cdot \mathbb{1}_{u'(X) > t'}] = m^{\text{Up}} > 0.$$

Since $\mathbb{E}[d'(X) - d(X)] = 0$, we see that

$$\begin{aligned}
\mathbb{E}[(d'(X) - d(X)) \cdot u'(X)] &= \mathbb{E}[(d'(X) - d(X)) \cdot \mathbb{1}_{u'(X) > t'} \cdot u'(X)] \\
&\quad + \mathbb{E}[(d'(X) - d(X)) \cdot \mathbb{1}_{u'(X) \leq t'} \cdot u'(X)] \\
&> t' \cdot \mathbb{E}[(d'(X) - d(X)) \cdot \mathbb{1}_{u'(X) > t'}] \\
&\quad + t' \cdot \mathbb{E}[(d'(X) - d(X)) \cdot \mathbb{1}_{u'(X) \leq t'}] \\
&= t' \cdot \mathbb{E}[d'(X) - d(X)] \\
&= 0,
\end{aligned}$$

where in the inequality we have used the fact that if $d'(x) > d(x)$, $u'(x) > t'$, and if $d'(x) < d(x)$, $u'(x) \leq t'$. Therefore

$$\mathbb{E}[d(X) \cdot u'(X)] < \mathbb{E}[d'(X) \cdot u'(X)],$$

i.e., $d'(x)$ strongly Pareto dominates $d(x)$.

Now, we prove the second claim, namely, that a multiple threshold policy $\tau(x)$ that cannot be represented with non-negative thresholds is strongly Pareto dominated. For, if $\tau(x)$ is such a policy, then, by Lemma A.5, for any $u \in \mathcal{U}$, $\mathbb{E}[\tau(X) \cdot \mathbb{1}_{u(X) < 0}] > 0$. It follows immediately that $\tau'(x) = \tau(x) \cdot \mathbb{1}_{u(x) > 0}$ satisfies $u(\tau') > u(\tau)$. By consistency modulo α , the definition of $\tau'(x)$ does not depend on our choice of u , and so $u(\tau') > u(\tau)$ for every $u \in \mathcal{U}$, i.e., $\tau'(x)$ strongly Pareto dominates $\tau(x)$. $\square$

The following results, which draw on Lemma A.6, are useful in the proof of Theorem 2.17.

Definition A.7. We say that a decision policy $d(x)$ is *budget-exhausting* if

$$\begin{aligned}
\min(b, \Pr(u(X) > 0)) &\leq \mathbb{E}[d(X)] \\
&\leq \min(b, \Pr(u(X) \geq 0)).
\end{aligned}$$

Remark A.8. We note that if $\mathcal{U}$ is consistent modulo α , then whether or not a decision policy $d(x)$ is budget-exhausting does not depend on the choice of $u \in \mathcal{U}$. Further, if $\Pr(u(X) = 0) = 0$ —e.g., if the distribution of X is $\mathcal{U}$ -fine—then the decision policy is budget-exhausting if and only if $\mathbb{E}[d(X)] = \min(b, \Pr(u(X) > 0))$.

Corollary A.9. Let $\mathcal{U}$ be a collection of utilities consistent modulo α . If $\tau(x)$ is a feasible policy

that is not a budget-exhausting multiple threshold policy with non-negative thresholds, then $\tau(x)$ is strongly Pareto dominated.

Proof. Suppose $\tau(x)$ is not strongly Pareto dominated. By Lemma A.6, it is a multiple threshold policy with non-negative thresholds.

Now, suppose toward a contradiction that $\tau(x)$ is not budget-exhausting. Then, either $\mathbb{E}[\tau(X)] > \min(b, \Pr(u(X) \geq 0))$ or $\mathbb{E}[\tau(X)] < \min(b, \Pr(u(X) > 0))$.

In the first case, since $\tau(x)$ is feasible, it follows that $\mathbb{E}[\tau(X)] > \Pr(u(X) \geq 0)$. It follows that $\tau(x) \cdot \mathbb{1}_{u(x) < 0}$ is not almost surely zero. Therefore

$$\mathbb{E}[\tau(X)] < \mathbb{E}[\tau(X) \cdot \mathbb{1}_{u(X) > 0}],$$

and, by consistency modulo α , this holds for any $u \in \mathcal{U}$. Therefore $\tau(x)$ is strongly Pareto dominated, contrary to hypothesis.

In the second case, consider

$$d(x) = \theta \cdot \mathbb{1}_{u(x) > 0} + (1 - \theta) \cdot \tau(x).$$

Since $\mathbb{E}[\tau(X)] < \min(b, \Pr(u(X) > 0))$ and

$$\mathbb{E}[d(X)] = \theta \cdot \Pr(u(X) > 0) + (1 - \theta) \cdot \mathbb{E}[\tau(X)],$$

there exists some $\theta > 0$ such that $d(x)$ is feasible.

For that θ , a similar calculation shows immediately that $u(d) > u(\tau)$, and, by consistency modulo α , $u'(d) > u'(\tau)$ for all $u' \in \mathcal{U}$. Therefore, again, $d(x)$ strongly Pareto dominates $\tau(x)$, contrary to hypothesis. $\square$

Lemma A.10. *Given a utility u , there exists a mapping T from $[0, 1]^{\mathcal{A}}$ to $[-\infty, \infty]^{\mathcal{A}}$ taking sets of quantiles $\{q_a\}_{a \in \mathcal{A}}$ to thresholds $\{t_a\}_{a \in \mathcal{A}}$ such that:*

1. *T is monotonically non-increasing in each coordinate;*
2. *For each set of quantiles, there is a multiple threshold policy $\tau : \mathcal{X} \rightarrow [0, 1]$ with thresholds $T(\{q_a\})$ with respect to u such that $\mathbb{E}[\tau(X) \mid A = a] = q_a$.*

Proof. Simply choose

$$t_a = \inf\{s \in \mathbb{R} : \Pr(u(X) > s) < q_a\}. \quad (\text{A.4})$$

Then define

$$p_a = \begin{cases} \frac{q_a - \Pr(u(X) > t_a | A=a)}{\Pr(u(X) = t_a | A=a)} & \Pr(u(X) = t_a, A = a) > 0 \\ 0 & \Pr(u(X) = t_a, A = a) = 0. \end{cases}$$

Note that $\Pr(u(X) \geq t_a | A = a) \geq q_a$, since, by definition, $\Pr(u(X) > t_a - \epsilon | A = a) \geq q_a$ for all $\epsilon > 0$. Therefore,

$$\Pr(u(X) > t_a | A = a) + \Pr(u(X) = t_a | A = a) \geq q_a,$$

and so $p_a \leq 1$. Further, since $\Pr(u(X) > t_a | A = a) \leq q_a$, we have that $p_a \geq 0$.

Finally, let

$$d(x) = \begin{cases} 1 & u(x) > t_{\alpha(x)}, \\ p_a & u(x) = t_{\alpha(x)}, \\ 0 & u(x) < t_{\alpha(x)}, \end{cases}$$

and it follows immediately that $\mathbb{E}[d(X) | A = a] = q_a$. That t_a is a monotonically non-increasing function of q_a follows immediately from Eq. (A.4). $\square$

We can further refine Cor. A.9 and Lemma A.10 as follows:

Lemma A.11. *Let u be a utility. Then a feasible policy is utility maximizing if and only if it is a budget-exhausting threshold policy. Moreover, there exists at least one utility maximizing policy.*

Proof. Let $\bar{\alpha}$ be a constant map, i.e., $\bar{\alpha} : \mathcal{X} \rightarrow \bar{\mathcal{A}}$, where $|\bar{\mathcal{A}}| = 1$. Then $\mathcal{U} = \{u\}$ is consistent modulo $\bar{\alpha}$, and so by Cor. A.9, any Pareto efficient policy is a budget exhausting multiple threshold policy relative to $\mathcal{U}$. Since $\mathcal{U}$ contains a single element, a policy is Pareto efficient if and only if it is utility maximizing. Since $\bar{\alpha}$ is constant, a policy is a multiple threshold policy relative to $\bar{\alpha}$ if and only if it is a threshold policy. Therefore, a policy is utility maximizing if and only if it is a budget exhausting threshold policy. By Lemma A.10, such a policy exists, and so the maximum is attained. $\square$

A.6 Prevalence and the Proof of Theorem 2.17

The notion of a probabilistically “small” set—such as the event in which an idealized dart hits the exact center of a target—is, in finite-dimensional real vector spaces, typically encoded by the idea of a Lebesgue null set.

Here we prove that the set of distributions such that there exists a policy satisfying either counterfactual equalized odds, conditional principal fairness, or counterfactual fairness that is not strongly Pareto dominated is “small” in an analogous sense. The proof turns on the following intuition. Each of the fairness definitions imposes a number of constraints. By Lemma A.6, any policy that is not strongly Pareto dominated is a multiple threshold policy. By adjusting the group-specific thresholds of such a policy, one can potentially satisfy one constraint per group. If there are more constraints than groups, then one has no additional degrees of freedom that can be used to ensure that the remaining constraints are satisfied. If, by chance, those constraints *are* satisfied with the same threshold policy, they are not satisfied robustly—even a minor distribution shift, such as increasing the amount of mass above the threshold by any amount on the relevant subpopulation, will break them. Therefore, over a “typical” distribution, at most $|\mathcal{A}|$ of the constraints can simultaneously be satisfied by a Pareto efficient policy, meaning that typically no Pareto efficient policy fully satisfies all of the conditions of the fairness definitions.

Formalizing this intuition, however, requires considerable care. In Section A.6.1, we give a brief introduction to a popular generalization of null sets to infinite-dimensional vector spaces, drawing heavily on a review article by Ott and Yorke (2005). In Section A.6.2 we provide a road map of the proof itself. In Section A.6.3, we establish the main hypotheses necessary to apply the notion of prevalence to a convex set—in our case, the set of $\mathcal{U}$ -fine distributions. In Section A.6.4, we establish a number of technical lemmata used in the proof of Theorem 2.17, and provide a proof of the theorem itself in Section A.6.5. In Section A.6.8, we show why the hypothesis of $\mathcal{U}$ -finess is important and how conspiracies between atoms in the distribution of $u(X)$ can lead to “robust” counterexamples.

A.6.1 Shyness and Prevalence

Lebesgue measure λ_n on $\mathbb{R}^n$ has a number of desirable properties:

- **Local finiteness:** For any point $v \in \mathbb{R}^n$, there exists an open set U containing x such that $\lambda_n[U] < \infty$;
- **Strict positivity:** For any open set U , if $\lambda_n[U] = 0$, then $U = \emptyset$;
- **Translation invariance:** For any $v \in \mathbb{R}^n$ and measurable set E , $\lambda_n[E + v] = \lambda_n[E]$.

No measure on an infinite-dimensional, separable Banach space, such as $L^1(\mathbb{R})$, can satisfy these three properties (Ott and Yorke, 2005). However, while there is no generalization of Lebesgue measure to infinite dimensions, there is a generalization of Lebesgue null sets—called *shy* sets—to the infinite-dimensional context that preserves many of their desirable properties.

Definition A.12 (Hunt et al. (1992)). Let V be a completely metrizable topological vector space. We say that a Borel set $E \subseteq V$ is *shy* if there exists a Borel measure μ on V such that:

1. There exists compact $C \subseteq V$ such that $0 < \mu[C] < \infty$,
2. For all $v \in V$, $\mu[E + v] = 0$.

An arbitrary set $F \subseteq V$ is *shy* if there exists a shy Borel set $E \subseteq V$ containing F .

We say that a set is *prevalent* if its complement is shy.

Prevalence generalizes the concept of Lebesgue “full measure” or “co-null” sets (i.e., sets whose complements have null Lebesgue measure) in the following sense:

Proposition A.13 (Hunt et al. (1992)). *Let V be a completely metrizable topological vector space. Then:*

- *Any prevalent set is dense in V ;*
- *If $G \subseteq L$ and G is prevalent, then L is prevalent;*
- *A countable intersection of prevalent sets is prevalent;*
- *Every translate of a prevalent set is prevalent;*
- *If $V = \mathbb{R}^n$, then $G \subseteq \mathbb{R}^n$ is prevalent if and only if $\lambda_n[\mathbb{R}^n \setminus G] = 0$.*

As is conventional for sets of full measure in finite-dimensional spaces, if some property holds for every $v \in E$, where E is prevalent, then we say that the property holds for *almost every* $v \in V$ or that it holds *generically in* V .

Prevalence can also be generalized from vector spaces to convex subsets of vector spaces, although additional care must be taken to ensure that a relative version of Prop. A.13 holds.

Definition A.14 (Anderson and Zame (2001)). Let V be a topological vector space and let $C \subseteq V$ be a convex subset completely metrizable in the subspace topology induced by V . We say that a universally measurable set $E \subseteq C$ is *shy in* C at $c \in C$ if for each $1 \geq \delta > 0$, and each neighborhood U of 0 in V , there is a regular Borel measure μ with compact support such that

$$\text{SUPP}(\mu) \subseteq (\delta(C - c) + c) \cap (U + c),$$

and $\mu[E + v] = 0$ for every $v \in V$.

We say that E is *shy in* C or *shy relative to* C if E is shy in C at c for every $c \in C$. An arbitrary set $F \subseteq V$ is shy in C if there exists a universally measurable shy set $E \subseteq C$ containing F .

A set G is *prevalent in* C if $C \setminus G$ is shy in C .

Proposition A.15 (Anderson and Zame (2001)). *If E is shy at some point $c \in C$, then E is shy at every point in C and hence is shy in C .*

Sets that are shy in C enjoy similar properties to sets that are shy in V .

Proposition A.16 (Anderson and Zame (2001)). *Let V be a topological vector space and let $C \subseteq V$ be a convex subset completely metrizable in the subspace topology induced by V . Then:*

- *Any prevalent set in C is dense in C ;*
- *If $G \subseteq L$ and G is prevalent in C , then L is prevalent in C ;*
- *A countable intersection of sets prevalent in C is prevalent in C*
- *If G is prevalent in C then $G + v$ is prevalent in $C + v$ for all $v \in V$.*
- *If $V = \mathbb{R}^n$ and $C \subseteq V$ is a convex subset with non-empty interior, then $G \subseteq C$ is prevalent in C if and only if $\lambda_n[C \setminus G] = 0$.*

Sets that are shy in C can often be identified by inspecting their intersections with a finite-dimensional subspace W of V , a strategy we use to prove Theorem 2.17.

Definition A.17 (Anderson and Zame (2001)). A universally measurable subset E of a convex and completely metrizable set C is said to be *k-shy in C* if there exists a k -dimensional subspace $W \subseteq V$ such that

1. A translate of the set C has positive Lebesgue measure in W , i.e., $\lambda_W[C + v_0] > 0$ for some $v_0 \in V$;
2. Every translate of the set E is a Lebesgue null set in W , i.e., $\lambda_W[E + v] = 0$ for all $v \in V$.

Here λ_W denotes k -dimensional Lebesgue measure supported on W .¹ We refer to such a W as a *k-dimensional probe* witnessing the k -shyness of E , and to an element $w \in W$ as a *perturbation*.

The following intuition motivates the use of probes to detect shy sets. By analogy with Fubini's theorem, one can imagine trying to determine whether a subset of a finite-dimensional vector space is large or small by looking at its cross sections parallel to some subspace $W \subseteq V$. If a set $E \subseteq V$ is small in each cross section—i.e., if $\lambda_W[E + v] = 0$ for all $v \in V$ —then E itself is small in V , i.e., E has λ_V -measure zero.

Proposition A.18 (Anderson and Zame (2001)). *Every k-shy set in C is shy in C .*

A.6.2 Outline

To aid the reader in following the application of the theory in Section A.6.1 to the proof of Theorem 2.17, we provide the following outline of the argument.

In **Section A.6.3** we establish the context to which we apply the notion of relative shyness. In particular, we introduce the vector space $\mathbb{K}$ consisting of the *totally bounded Borel measures* on the state space $\mathcal{K}$ —where $\mathcal{K}$ is $\mathcal{X} \times \mathcal{Y}$, $\mathcal{X} \times \mathcal{Y} \times \mathcal{Y}$, or $\mathcal{A} \times \mathcal{X}^{\mathcal{A}}$, depending on which notion of fairness is under consideration. We further isolate the subspace $\mathbf{K} \subseteq \mathbb{K}$ of $\mathcal{U}$ -fine totally bounded Borel measures. Within this space, we are interested in the convex set $\mathbf{Q} \subseteq \mathbf{K}$, the set of $\mathcal{U}$ -fine joint

¹Note that Lebesgue measure on W is only defined up to a choice of basis; however, since $\lambda[T(A)] = |\det(T)| \cdot \lambda[A]$ for any linear automorphism T and Lebesgue measure λ , whether a set has null measure does not depend on the choice of basis.

probability distributions of, respectively, X and $Y(1)$; X , $Y(0)$, $Y(1)$; or A and the $X_{\Pi,A,a}$. Within $\mathbf{Q}$, we identify $\mathbf{E} \subseteq \mathbf{Q}$, the set of $\mathcal{U}$ -fine distributions on $\mathcal{K}$ over which there exists a policy satisfying the relevant fairness definition that is not strongly Pareto dominated. The claim of Theorem 2.17 is that $\mathbf{E}$ is shy relative to $\mathbf{Q}$.

To ensure that relative shyness generalizes Lebesgue null measure in the expected way—i.e., that Prop. A.16 holds—Definition A.14 has three technical requirements: (1) that the ambient vector space V be a topological vector space; (2) that the convex set C be completely metrizable; and (3) that the shy set E be universally measurable. In **Lemma A.22**, we observe that $\mathbb{K}$ is a complete topological vector space under the total variation norm, and so is a Banach space. We extend this in **Cor. A.27**, showing that $\mathbf{K}$ is also a Banach space. We use this fact in **Lemma A.30** to show that $\mathbf{Q}$ is a completely metrizable subset of $\mathbf{K}$, as well as convex. Lastly, in **Lemma A.36**, we show that the set $\mathbf{E}$ is closed, and therefore universally measurable.

In **Section A.6.4**, we develop the machinery needed to construct a probe $\mathbf{W}$ for the proof of Theorem 2.17 and prove several lemmata simplifying the eventual proof of the theorem. To build the probe, it is necessary to construct measures $\mu_{\max,a}$ with maximal support on the utility scale. This ensures that if any two threshold policies produce different decisions on *any* $\mu \in \mathbf{K}$, they will produce different decisions on typical perturbations. The construction of the $\mu_{\max,a}$, is carried out in **Lemma A.38** and **Cor. A.39**. Next, we introduce the basic style of argument used to show that a subset of $\mathbf{Q}$ is shy in **Lemma A.42** and **Lemma A.43**, in particular, by showing that the set of $\mu \in \mathbf{Q}$ that give positive probability to an event E is either prevalent or empty. We use then use a technical lemma, **Lemma A.44**, to show, in effect, that a generic element of $\mathbf{Q}$ has support on the utility scale wherever a given fixed distribution $\mu \in \mathbf{Q}$ does. In **Defn. A.46**, we introduce the concept of overlapping and splitting utilities, and show in **Lemma A.47** that this property is generic in $\mathbf{Q}$ unless there exists a ω -stratum that contains no positive-utility observables x . Lastly, in **Lemma A.48**, we provide a mild simplification of the characterization of finitely shy sets that makes the proof of Theorem 2.17 more straightforward.

Finally, in **Section A.6.5**, we give the proof of Theorem 2.17. We divide the proof into three parts. In the first part, we restrict our attention to the case of counterfactual equalized odds, and show in detail how to combine the lemmata of the previous section to construct the (at most) $2 \cdot |\mathcal{A}|$ -dimensional probe $\mathbf{W}$. In the second part we consider two distinct cases. The argument in

both cases is conceptually parallel. First, we argue that the balance conditions of counterfactual equalized odds encoded by Eq. (2.4) must be broken by a typical perturbation in $\mathbf{W}$. In particular, we argue that for a given base distribution μ , there can be at most one budget-exhausting multiple threshold policy that can—although need not necessarily—satisfy counterfactual equalized odds. We show that the form of this policy cannot be altered by an appropriate perturbation in $\mathbf{W}$, but that the conditional probability of a positive decision will, in general, be altered in such a way that Eq. (2.4) can only hold for a $\lambda_{\mathbf{W}}$ -null set of perturbations. In the final section, we lay out modifications that can be made to the proof given for counterfactual equalized odds in the first two parts that adapt the argument to the cases of conditional principal fairness and path-specific fairness. In particular, we show how to construct the probe $\mathbf{W}$ in such a way that the additional conditioning on the reduced covariates $W = \omega(X)$ in Eqs. (2.5) and (2.8) does not affect the argument.

A.6.3 Convexity, Complete Metrizability, and Universal Measurability

In this section, we establish the background requirements of Prop. A.18 for the setting of Theorem 2.17. In particular, we exhibit the $\mathcal{U}$ -fine distributions as a convex subset of a topological vector space, the set of totally bounded $\mathcal{U}$ -fine Borel measures. We show that the $\mathcal{U}$ -fine probability distributions form a completely metrizable subset in the topology it inherits from the space of totally bounded measures. Lastly, we show that the set of regular distributions under which there exists a Pareto efficient policy satisfying one of the three fairness criteria is closed, and therefore universally measurable.

A.6.3.1 Background and Notation

We begin by establishing some notational conventions. We let $\mathcal{K}$ denote the underlying state space over which the distributions in Theorem 2.17 range. Specifically, $\mathcal{K} = \mathcal{X} \times \mathcal{Y}$ in the case of counterfactual equalized odds; $\mathcal{K} = \mathcal{X} \times \mathcal{Y} \times \mathcal{Y}$ in the case of conditional principal fairness; and $\mathcal{K} = \mathcal{A} \times \mathcal{X}^{\mathcal{A}}$ in the case of path-specific fairness. We note that since $\mathcal{X} \subseteq \mathbb{R}^k$ for some k and $\mathcal{Y} \subseteq \mathbb{R}$, $\mathcal{K}$ may equivalently be considered a subset of $\mathbb{R}^n$ for some $n \in \mathbb{N}$, with the subspace topology (and Borel sets) inherited from $\mathbb{R}^n$.²

²In the case of path-specific fairness, we can equivalently think of $\mathcal{A}$ as a set of integers indexing the groups.

We recall the definition of totally bounded measures.

Definition A.19. Let $\mathcal{M}$ be a σ -algebra on V , and let μ be a countably additive $(V, \mathcal{M})$ -measure.

Then, we define

$$|\mu|[E] = \sup \sum_{i=1}^{\infty} |\mu[E_i]| \quad (\text{A.5})$$

where the supremum is taken over all countable partitions $\{E_i\}_{i \in \mathbb{N}}$, i.e., collections such that $\bigcup_{i=1}^{\infty} E_i = E$ and $E_i \cap E_j = \emptyset$ for $j \neq i$. We call $|\mu|$ the *total variation of μ* , and the *total variation norm of μ* is $|\mu|[V]$.

We say that μ is *totally bounded* if its total variation norm is finite, i.e., $|\mu|[V] < \infty$.

Lemma A.20. If μ is totally bounded, then $|\mu|$ is a finite positive measure on $(V, \mathcal{M})$, and $|\mu[E]| \leq |\mu|[E]$ for all $E \in \mathcal{M}$.

See Theorem 6.2 in Rudin (1987) for proof.

We let $\mathbb{K}$ denote the set of totally bounded Borel measures on $\mathcal{K}$. We note that, in the case of path specific fairness, which involves the joint distributions of counterfactuals, X is not defined directly. Rather, the joint distribution of the counterfactuals $X_{\Pi, A, a'}$ and A defines the distribution of X through consistency, i.e., what would have happened to someone if their group membership were changed to $a' \in \mathcal{A}$ is what actually happens to them if their group membership is a' . More formally, $\Pr(X \in E \mid A = a') = \Pr(X_{\Pi, A, a'} \in E \mid A = a')$ for all Borel sets $E \subseteq \mathcal{X}$. (See § 3.6.3 in Pearl (2009).)

For any $\mu \in \mathbb{K}$, we adopt the following notational conventions. If we say that a property holds μ -a.s., then the subset of $\mathcal{K}$ on which the property fails has $|\mu|$ -measure zero. If $E \subseteq \mathcal{K}$ is a measurable set, then we denote by $\mu \upharpoonright_E$ the restriction of μ to E , i.e., the measure defined by the mapping $E' \mapsto \mu[E \cap E']$. We let $\mathbb{E}_{\mu}[f] = \int_{\mathcal{K}} f d\mu$, and for measurable sets E , $\Pr_{\mu}(E) = \mu[E]$.³ The fairness criteria we consider involve conditional independence relations. To make sense of conditional independence relations more generally, for Borel measurable f we define $\mathbb{E}_{\mu}[f \mid \mathcal{F}]$ to be the Radon-Nikodym derivative of the measure $E \mapsto \mathbb{E}_{\mu}[f \cdot \mathbb{1}_E]$ with respect to the measure μ restricted to the sub- σ -algebra of Borel sets $\mathcal{F}$. (See § 34 in Billingsley (2008).) Similarly, we define $\mathbb{E}_{\mu}[f \mid g]$ to be $\mathbb{E}_{\mu}[f \mid \sigma(g)]$, where $\sigma(g)$ denotes the sub- σ -algebra of the Borel sets generated by g .

³To state and prove our results in a notationally uniform way, we occasionally write $\Pr_{\mu}(E)$ even when μ ranges over measures that may not be probability measures.

In cases where the condition can occur with non-zero probability, we can instead make use of the elementary definition of discrete conditional probability.

Lemma A.21. *Let g be a Borel function on $\mathcal{K}$, and suppose $\Pr_\mu(g = c) \neq 0$ for some constant $c \in \mathbb{R}$. Then, we have that μ -a.s., for any Borel function f ,*

$$\mathbb{E}_\mu[f \mid g] \cdot \mathbb{1}_{g=c} = \frac{\mathbb{E}_\mu[f \cdot \mathbb{1}_{g=c}]}{\Pr_\mu(g = c)} \cdot \mathbb{1}_{g=c}.$$

See Rao (2005) for proof.

With these notational conventions in place, we turn to establishing the background conditions of Prop. A.18.

Lemma A.22. *The set of totally bounded measures on a measure space $(V, \mathcal{M})$ form a complete topological vector space under the total variation norm, and hence a Banach space.*

See, e.g., Steele (2019) for proof. It follows from this that $\mathbb{K}$ is a Banach space.

Remark A.23. Since $\mathbb{K}$ is a Banach space, it possesses a topology, and consequently a collection of Borel subsets. These Borel sets are to be distinguished from the Borel subsets of the underlying state space $\mathcal{K}$, which the elements of $\mathbb{K}$ measure. The requirement that the subset E of the convex set C be universally measurable in Proposition A.18 is in reference to the *Borel subsets of $\mathbb{K}$* ; the requirement that $\mu \in \mathbb{K}$ be a Borel measure is in reference to the *Borel subsets of $\mathcal{K}$* .

Recall the definition of absolute continuity.

Definition A.24. Let μ and ν be measures on a measure space $(V, \mathcal{M})$. We say that a measure ν is *absolutely continuous with respect to μ* —also written $\nu \ll \mu$ —if, whenever $\mu[E] = 0$, $\nu[E] = 0$.

Absolute continuity is a closed property in the topology induced by the total variation norm.

Lemma A.25. *Consider the space of totally bounded measures on a measure space $(V, \mathcal{M})$ and fix μ . The set of ν such that $\nu \ll \mu$ is closed.*

Proof. Let $\{\nu_i\}_{i \in \mathbb{N}}$ be a convergent sequence of measures absolutely continuous with respect to μ . Let the limit of the ν_i be ν . We seek to show that $\nu \ll \mu$. Let $E \in \mathcal{M}$ be an arbitrary set such that

$\mu[E] = 0$. Then, we have that

$$\begin{aligned}\nu[E] &= \lim_{n \rightarrow \infty} \nu_i[E] \\ &= \lim_{n \rightarrow \infty} 0 \\ &= 0,\end{aligned}$$

since $\nu_i \ll \mu$ for all i . Since E was arbitrary, the result follows. $\square$

Recall the definition of a pushforward measure.

Definition A.26. Let $f : (V, \mathcal{M}) \rightarrow (V', \mathcal{M}')$ be a measurable function. Let μ be a measure on V . We define the *pushforward measure* $\mu \circ f^{-1}$ on V' by the map $E' \mapsto \mu[f^{-1}(E')]$ for $E' \in \mathcal{M}'$.

Within $\mathbb{K}$, in the case of counterfactual equalized odds and conditional principal fairness, we define the subspace $\mathbf{K}$ to be the set of totally bounded measures μ on $\mathcal{K}$ such that the pushforward measure $\mu \circ u^{-1}$ is absolutely continuous with respect to the Lebesgue measure λ on $\mathbb{R}$ for all $u \in \mathcal{U}$. By the Radon-Nikodym theorem, these pushforward measures arise from densities, i.e., for any $\mu \in \mathbf{K}$, there exists a unique $f_\mu \in L^1(\mathbb{R})$ such that for any measurable subset E of $\mathbb{R}$, we have

$$\mu \circ u^{-1}[E] = \int_E f_\mu d\lambda.$$

In the case of path-specific fairness, we require the joint distributions of the counterfactual utilities to have a joint density. That is, we define the subspace $\mathbf{K}$ to be the set of totally bounded measures μ on $\mathcal{K}$ such that the pushforward measure $\mu \circ (u^A)^{-1}$ is absolutely continuous with respect to Lebesgue measure on $\mathbb{R}^A$ for all $u \in \mathcal{U}$. Here, we recall that

$$u^A : (a, (x_{a'})_{a' \in \mathcal{A}}) \mapsto (u(x_{a'}))_{a' \in \mathcal{A}}.$$

As before, there exists a corresponding density $f_\mu \in L^1(\mathbb{R}^A)$.

We therefore see that $\mathbf{K}$ extends in a natural way the notion of a $\mathcal{U}$ - or $\mathcal{U}^A$ -fine distribution, and so, by a slight abuse of notation, refer to $\mathbf{K}$ as the set of $\mathcal{U}$ -fine measures on $\mathcal{K}$.

Indeed, since $\Pr_\mu(u(X) \in E, A = a) \leq \Pr_\mu(u(X) \in E)$, it also follows that, for $a \in \mathcal{A}$ such that $\Pr_\mu(A = a) > 0$, the conditional distributions of $u(X) \mid A = a$ are also absolutely continuous with respect to Lebesgue measure, and so also have densities. For notational convenience, we set $f_{\mu,a}$ to

be the function satisfying

$$\Pr_\mu(u(X) \in E, A = a) = \int_E f_{\mu,a} d\lambda,$$

so that $f_\mu = \sum_{a \in \mathcal{A}} f_{\mu,a}$.

Since absolute continuity is a closed condition, it follows that $\mathbf{K}$ is a closed subspace of $\mathbb{K}$. This leads to the following useful corollary of Lemma A.25.

Corollary A.27. *The collection of $\mathcal{U}$ -fine measures on $\mathcal{K}$ is a Banach space.*

Proof. It is straightforward to see that $\mathbf{K}$ is a subspace of $\mathbb{K}$. Since $\mathbf{K}$ is a closed subset of $\mathbb{K}$ by Lemma A.25, it is complete, and therefore a Banach space. $\square$

We note the following useful fact about elements of $\mathbf{K}$.

Lemma A.28. *Consider the mapping $\mu \mapsto f_\mu$ from $\mathbf{K}$ to $L^1(\mathbb{R})$ given by associating a measure μ with the Radon-Nikodym derivative of the pushforward measure $\mu \circ u^{-1}$. This mapping is continuous. Likewise, the mapping $\mu \mapsto f_{\mu,a}$ is continuous for all $a \in \mathcal{A}$, and, in the case of path-specific fairness, the mapping of μ to the Radon-Nikodym derivative of $\mu \circ (u^A)^{-1}$ is continuous.*

Proof. We show only the first case. The others follow by virtually identical arguments.

Let $\epsilon > 0$ be arbitrary. Choose $\mu \in \mathbf{K}$, and suppose that $|\mu - \mu'|[\mathcal{K}] < \epsilon$. Then, let

$$\begin{aligned} E^{\text{Up}} &= \{x \in \mathbb{R} : f_\mu(x) > f_{\mu'}(x)\} \\ E^{\text{Lo}} &= \{x \in \mathbb{R} : f_\mu(x) < f_{\mu'}(x)\}. \end{aligned}$$

Then E^{Up} and E^{Lo} are disjoint, so we have that

$$\begin{aligned} \|f_\mu - f_{\mu'}\|_{L^1(\mathbb{R})} &= \left| \int_{E^{\text{Up}}} f_\mu - f_{\mu'} d\lambda \right| + \left| \int_{E^{\text{Lo}}} f_\mu - f_{\mu'} d\lambda \right| \\ &= |(\mu - \mu')[u^{-1}(E^{\text{Up}})]| + |(\mu - \mu')[u^{-1}(E^{\text{Lo}})]| \\ &< \epsilon, \end{aligned}$$

where the second equality follows by the definition of pushforward measures and the inequality follows from Lemma A.20. Since ϵ was arbitrary, the claim follows. $\square$

Finally, we define $\mathbf{Q}$. We let $\mathbf{Q}$ be the subset of $\mathbf{K}$ consisting of all $\mathcal{U}$ -fine probability measures, i.e., measures $\mu \in \mathbb{K}$ such that:

1. The measure μ is $\mathcal{U}$ -fine;
2. For all Borel sets $E \subseteq \mathcal{K}$, $\mu[E] \geq 0$;
3. The measure of the whole space is unity, i.e., $\mu[\mathcal{K}] = 1$.

We conclude the background and notation by observing that threshold policies are defined wholly by their thresholds for distributions in $\mathbf{K}$ and $\mathbf{Q}$. Importantly, this observation does not hold when there are atoms on the utility scale—which measures in $\mathbf{K}$ lack—which can in turn lead to counterexamples to Theorem 2.17; see Appendix A.6.8.

Lemma A.29. *Let $\tau_0(x)$ and $\tau_1(x)$ be two multiple threshold policies. If $\tau_0(x)$ and $\tau_1(x)$ have the same thresholds, then for any $\mu \in \mathbf{K}$, $\tau_0(X) = \tau_1(X)$ μ -a.s. Similarly, for $\mu \in \mathbf{Q}$, if*

$$\mathbb{E}_\mu[\tau_0(X) \mid A = a] = \mathbb{E}_\mu[\tau_1(X) \mid A = a]$$

for all $a \in \mathcal{A}$ such that $\Pr_\mu(A = a) > 0$, then $\tau_0(X) = \tau_1(X)$ μ -a.s.

Moreover, for $\mu \in \mathbf{K}$ in the case of path-specific fairness, if $\tau_0(x)$ and $\tau_1(x)$ have the same thresholds, then $\tau_0(X_{\Pi,A,a}) = \tau_1(X_{\Pi,A,a})$ μ -a.s. for any $a \in \mathcal{A}$. Similarly, for $\mu \in \mathbf{Q}$ in the case of path-specific fairness, if

$$\mathbb{E}_\mu[\tau_0(X_{\Pi,A,a})] = \mathbb{E}_\mu[\tau_1(X_{\Pi,A,a})]$$

then $\tau_0(X_{\Pi,A,a}) = \tau_1(X_{\Pi,A,a})$ μ -a.s. as well.

Proof. First, we show that threshold policies with the same thresholds are equal, then we show that threshold policies that distribute positive decisions across groups in the same way are equal.

Let $\{t_a\}_{a \in \mathcal{A}}$ denote the shared set of thresholds. It follows that if $\tau_0(x) \neq \tau_1(x)$, then $u(x) = t_{\alpha(x)}$. Now,

$$\Pr(u(X) = t_a, A = a) = \int_{t_a}^{t_a} f_{\mu,a} d\lambda = 0,$$

so $\Pr_\mu(\tau_0(X) \neq \tau_1(X)) = 0$. Next, suppose

$$\mathbb{E}_\mu[\tau_0(X) \mid A = a] = \mathbb{E}_\mu[\tau_1(X) \mid A = a].$$

If the thresholds of the two policies agree for all $a \in \mathcal{A}$ such that $\Pr_\mu(A = a) > 0$, then we are done by the previous paragraph. Therefore, suppose $t_a^0 \neq t_a^1$ for some suitable $a \in \mathcal{A}$, where t_a^i represents

the threshold for group $a \in \mathcal{A}$ under the policy $\tau_i(x)$. Without loss of generality, suppose $t_a^0 < t_a^1$. Then, it follows that

$$\begin{aligned} \int_{t_a^0}^{t_a^1} f_{\mu,a} d\lambda &= \mathbb{E}_\mu[\tau_0(X) \mid A = a] - \mathbb{E}_\mu[\tau_1(X) \mid A = a] \\ &= 0. \end{aligned}$$

Since $\mu \in \mathbf{Q}$, $\mu = |\mu|$, whence

$$\Pr_{|\mu|}(t_a^0 \leq u(X) \leq t_a^1 \mid A = a) = 0.$$

Since this is true for all $a \in \mathcal{A}$ such that $\Pr_\mu(A = a) > 0$, $\tau_0(X) = \tau_1(X)$ μ -a.s.

The proof in the case of path-specific fairness is almost identical. $\square$

A.6.3.2 Convexity, Complete Metrizable, and Universal Measurability

The set of regular $\mathcal{U}$ -fine probability measures $\mathbf{Q}$ is the set to which we wish to apply Prop. A.18. To do so, we must show that $\mathbf{Q}$ is a convex and completely metrizable subset of $\mathbf{K}$.

Lemma A.30. *The set of regular probability measures $\mathbf{Q}$ is convex and completely metrizable.*

Proof. The proof proceeds in two pieces. First, we show that the $\mathcal{U}$ -fine probability distributions are convex, as can be verified by direct calculation. Then, we show that $\mathbf{Q}$ is closed and therefore complete in the original metric of $\mathbf{K}$.

We begin by verifying convexity. Let $\mu, \mu' \in \mathbf{Q}$ and let $E \subseteq \mathcal{K}$ be an arbitrary Borel subset of $\mathcal{K}$. Then, choose $\theta \in [0, 1]$, and note that

$$\begin{aligned} (\theta \cdot \mu + [1 - \theta] \cdot \mu')[E] &= \theta \cdot \mu[E] + [1 - \theta] \cdot \mu'[E] \\ &\geq \theta \cdot 0 + [1 - \theta] \cdot 0 \\ &= 0, \end{aligned}$$

and, likewise, that

$$\begin{aligned}
(\theta \cdot \mu + [1 - \theta] \cdot \mu')[\mathcal{K}] &= \theta \cdot \mu[\mathcal{K}] + [1 - \theta] \cdot \mu'[\mathcal{K}] \\
&= \theta \cdot 1 + [1 - \theta] \cdot 1 \\
&= 1.
\end{aligned}$$

It remains only to show that $\mathbf{Q}$ is completely metrizable. To prove this, it suffices to show that it is closed, since closed subsets of complete spaces are complete, and $\mathbf{K}$ is a Banach space by Cor. A.27, and therefore complete.

Suppose $\{\mu_i\}_{i \in \mathbb{N}}$ is a convergent sequence of probability measures in $\mathbf{K}$ with limit μ . Then

$$\mu[E] = \lim_{i \rightarrow \infty} \mu_i[E] \geq \lim_{i \rightarrow \infty} 0 = 0$$

and

$$\mu[\mathcal{K}] = \lim_{i \rightarrow \infty} \mu_i[\mathcal{K}] = \lim_{i \rightarrow \infty} 1 = 1.$$

Therefore $\mathbf{Q}$ is closed, and therefore complete, and hence is a convex, completely metrizable subset of $\mathbf{K}$. □

Next we prove that the set $\mathbf{E}$ of regular $\mathcal{U}$ -fine densities over which there exists a policy satisfying the relevant counterfactual fairness definition that is not strongly Pareto dominated is universally measurable.

Recall the definition of universal measurability.

Definition A.31. Let V be a complete topological space. Then $E \subseteq V$ is *universally measurable* if V is measurable by the completion of every finite Borel measure on V , i.e., if for every finite Borel measure μ , there exist Borel sets E' and S such that $E \triangle E' \subseteq S$ and $\mu[S] = 0$.

We note that if a set is Borel, it is by definition universally measurable. Moreover, if a set is open or closed, it is by definition Borel.

To show that $\mathbf{E}$ is closed, we show that any convergent sequence in $\mathbf{E}$ has a limit in $\mathbf{E}$. The technical complication of the argument stems from the following fact that satisfying the fairness conditions, e.g., Eq. (2.7), involves conditional expectations, about which very little can be said in the absence of a density, and which are difficult to compare when taken across distinct measures.

To handle these difficulties, we begin with a technical lemma, Lemma A.35, which gives a coarse bound on how different the conditional expectations of the same variable can be with respect to a sub- σ -algebra $\mathcal{F}$ over two different distributions, μ and μ' , before applying the results to the proof of Lemma A.36.

Definition A.32. Let μ be a measure on a measure space $(V, \mathcal{M})$, and let f be μ -measurable. Consider the equivalence class of $\mathcal{M}$ -measurable functions $C = \{g : g = f \text{ } \mu\text{-a.e.}\}$.⁴ We say that any $g \in C$ is a *version* of f , and that $g \in C$ is a *standard version* if $g(v) \leq C$ for some constant C and all $v \in V$.

Remark A.33. It is straightforward to see that for $f \in L^\infty(\mu)$, a standard version always exists with $C = \|f\|_\infty$.

Remark A.34. Note that in general, the conditional expectation $\mathbb{E}_{\mu'}[f \mid \mathcal{F}]$ is defined only μ' -a.e. If μ is not assumed to be absolutely continuous with respect to μ' , it follows that

$$\|\mathbb{E}_\mu[f \mid \mathcal{F}] - \mathbb{E}_{\mu'}[f \mid \mathcal{F}]\|_{L^1(\mu)} \quad (\text{A.6})$$

is not entirely well-defined, in that its value depends on what version of $\mathbb{E}_{\mu'}[f \mid \mathcal{F}]$ one chooses. For appropriate f , however, one can nevertheless bound Eq. (A.6) for any standard version of $\mathbb{E}_{\mu'}[f \mid \mathcal{F}]$.

Lemma A.35. Let μ, μ' be totally bounded measures on a measure space $(V, \mathcal{M})$. Let $f \in L^\infty(\mu) \cap L^\infty(\mu')$. Let $\mathcal{F}$ be a sub- σ -algebra of $\mathcal{M}$. Let

$$C = \max(\|f\|_{L^\infty(\mu)}, \|f\|_{L^\infty(\mu')}).$$

Then, if g is a standard version of $\mathbb{E}_{\mu'}[f \mid \mathcal{F}]$, we have that

$$\int_V |\mathbb{E}_\mu[f \mid \mathcal{F}] - g| d\mu \leq 4C \cdot |\mu - \mu'|[V]. \quad (\text{A.7})$$

Proof. First, we note that both $\mathbb{E}_\mu[f \mid \mathcal{F}]$ and g are $\mathcal{F}$ -measurable. Therefore, the sets

$$E^{\text{Up}} = \{v \in V : \mathbb{E}_\mu[f \mid \mathcal{F}](v) > g(v)\}$$

and

$$E^{\text{Lo}} = \{v \in V : \mathbb{E}_\mu[f \mid \mathcal{F}](v) < g(v)\}$$

⁴Some authors define $L^p(\mu)$ spaces to consist of such equivalence classes, rather than the definition we use here.

are in $\mathcal{F}$. Now, note that

$$\int_V |\mathbb{E}_\mu[f \mid \mathcal{F}] - g| d\mu = \int_{E^{\text{Up}}} \mathbb{E}_\mu[f \mid \mathcal{F}] - g d\mu + \int_{E^{\text{Lo}}} g - \mathbb{E}_\mu[f \mid \mathcal{F}] d\mu.$$

First consider E^{Up} . Then, we have that

$$\begin{aligned} \int_{E^{\text{Up}}} \mathbb{E}_\mu[f \mid \mathcal{F}] - g d\mu &= \int_{E^{\text{Up}}} \mathbb{E}_\mu[f \mid \mathcal{F}] - g d\mu + \int_{E^{\text{Up}}} g - g d\mu' \\ &\leq \left| \int_{E^{\text{Up}}} \mathbb{E}_\mu[f \mid \mathcal{F}] d\mu - \int_{E^{\text{Up}}} g d\mu' \right| + \int_{E^{\text{Up}}} g d|\mu - \mu'| \\ &\leq \left| \int_{E^{\text{Up}}} f d\mu - \int_{E^{\text{Up}}} f d\mu' \right| + \int_{E^{\text{Up}}} C d|\mu - \mu'|, \end{aligned}$$

where in the final inequality, we have used the fact that, since g is a standard version of $\mathbb{E}_{\mu'}[f \mid \mathcal{F}]$,

$$g(v) \leq \|\mathbb{E}_{\mu'}[f \mid \mathcal{F}]\|_{L^\infty(\mu')} \leq C$$

for all $v \in V$, and the fact that, by the definition of conditional expectation,

$$\int_E \mathbb{E}_\nu[h \mid \mathcal{F}] d\nu = \int_E h d\nu$$

for any $E \in \mathcal{F}$.

Since f is everywhere bounded by C , applying Lemma A.20 yields that this final expression is less than or equal to $2C \cdot |\mu - \mu'|[V]$. An identical argument shows that

$$\int_{E^{\text{Lo}}} g - \mathbb{E}_\mu[f \mid \mathcal{F}] d\mu \leq 2C \cdot |\mu - \mu'|[V],$$

whence the result follows. $\square$

Lemma A.36. *Let $\mathbf{E} \subseteq \mathbf{Q}$ denote the set of joint densities on $\mathcal{K}$ such that there exists a policy satisfying the relevant fairness definition that is not strongly Pareto dominated. Then, $\mathbf{E}$ is closed, and therefore universally measurable.*

Proof. For notational simplicity, we consider the case of counterfactual equalized odds. The proofs in the other two cases are virtually identical.

Suppose $\mu_i \rightarrow \mu$ in $\mathbf{K}$, where $\{\mu_i\}_{i \in \mathbb{N}} \subseteq \mathbf{E}$. Then, by Lemma A.28, $f_{\mu_i, a} \rightarrow f_{\mu, a}$ in $L^1(\mathbb{R})$. Moreover, by Lemma A.6, there exists a sequence of threshold policies $\{\tau_i(x)\}_{i \in \mathbb{N}}$ such that both

$$\mathbb{E}_{\mu_i}[\tau(X)] = \min(b, \Pr_{\mu_i}(u(X) > 0))$$

and

$$\mathbb{E}_{\mu_i}[\tau_i(X) \mid A, Y(1)] = \mathbb{E}_{\mu_i}[\tau_i(X) \mid Y(1)].$$

Let $\{q_{a,i}\}_{a \in \mathcal{A}}$ be defined by

$$q_{a,i} = \mathbb{E}_{\mu_i}[\tau_i(X) \mid A = a]$$

if $\Pr_{\mu_i}(A = a) > 0$, and $q_{a,i} = 0$ otherwise.

Since $[0, 1]^{\mathcal{A}}$ is compact, there exists a convergent subsequence $\{\{q_{a,n_i}\}_{a \in \mathcal{A}}\}_{i \in \mathbb{N}}$. Let it converge to the collection of quantiles $\{q_a\}_{a \in \mathcal{A}}$ defining, by Lemma A.10, a multiple threshold policy $\tau(x)$ over μ .

Because $\mu_i \rightarrow \mu$ and $\{q_{a,n_i}\}_{a \in \mathcal{A}} \rightarrow \{q_a\}_{a \in \mathcal{A}}$, we have that

$$\mathbb{E}_{\mu}[\tau_{a,n_i}(X) \mid A = a] \rightarrow \mathbb{E}_{\mu}[\tau(X) \mid A = a]$$

for all $a \in \mathcal{A}$ such that $\Pr_{\mu}(A = a) > 0$. Therefore, by Lemma A.28, $\tau_{n_i}(X) \rightarrow \tau(X)$ in $L^1(\mu)$.

Choose $\epsilon > 0$ arbitrarily. Then, choose N so large that for i greater than N ,

$$|\mu - \mu_{n_i}|[\mathcal{K}] < \frac{\epsilon}{10}, \quad \|\tau(X) - \tau_{n_i}(X)\|_{L^1(\mu)} \leq \frac{\epsilon}{10}.$$

Then, observe that $\tau(x), \tau_i(x) \leq 1$, and recall that

$$\mathbb{E}_{\mu_{n_i}}[\tau_{n_i}(X) \mid A, Y(1)] = \mathbb{E}_{\mu_{n_i}}[\tau_{n_i}(X) \mid Y(1)]. \quad (\text{A.8})$$

Therefore, let $g_i(x)$ be a standard version of $\mathbb{E}_{\mu_{n_i}}[\tau_{n_i}(X) \mid Y(1)]$ over μ_{n_i} . By Eq. (A.8), $g_i(x)$ is also a standard version of $\mathbb{E}_{\mu_{n_i}}[\tau_{n_i}(X) \mid A, Y(1)]$ over μ_{n_i} . Then, by Lemma A.35, we have that

$$\begin{aligned} & \|\mathbb{E}_{\mu}[\tau(X) \mid A, Y(1)] - \mathbb{E}_{\mu_{n_i}}[\tau_{n_i}(X) \mid Y(1)]\|_{L^1(\mu)} \\ & \leq \|\mathbb{E}_{\mu}[\tau(X) \mid A, Y(1)] - \mathbb{E}_{\mu}[\tau_{n_i}(X) \mid A, Y(1)]\|_{L^1(\mu)} \\ & \quad + \|\mathbb{E}_{\mu}[\tau_{n_i}(X) \mid A, Y(1)] - g_i(X)\|_{L^1(\mu)} \\ & \quad + \|g_i(X) - \mathbb{E}_{\mu}[\tau_{n_i}(X) \mid Y(1)]\|_{L^1(\mu)} \\ & \quad + \|\mathbb{E}_{\mu}[\tau_{n_i}(X) \mid Y(1)] - \mathbb{E}_{\mu}[\tau(X) \mid Y(1)]\|_{L^1(\mu)} \\ & < \frac{\epsilon}{10} + \frac{4\epsilon}{10} + \frac{4\epsilon}{10} + \frac{\epsilon}{10}. \end{aligned}$$

Since $\epsilon > 0$ was arbitrary, it follows that, μ -a.e.,

$$\mathbb{E}_\mu[\tau(X) \mid A, Y(1)] = \mathbb{E}_\mu[\tau(X) \mid Y(1)].$$

Recall the standard fact that for independent random variables X and U ,

$$\mathbb{E}[f(X, U) \mid X] = \int f(X, u) dF_U(u),$$

where F_U is the distribution of U .⁵ Further recall that $D = \mathbb{1}_{U_D \leq \tau(X)}$, where $U_D \perp\!\!\!\perp X, Y(1)$. It follows that

$$\Pr_\mu(D = 1 \mid X, Y(1)) = \int_0^1 \mathbb{1}_{u_d < \tau(X)} d\lambda(u_d) = \tau(X).$$

Hence, by the law of iterated expectations,

$$\begin{aligned} \Pr_\mu(D = 1 \mid A, Y(1)) &= \mathbb{E}_\mu[\Pr_\mu(D = 1 \mid X, Y(1)) \mid A, Y(1)] \\ &= \mathbb{E}_\mu[\tau(X) \mid A, Y(1)] \\ &= \mathbb{E}_\mu[\tau(X) \mid Y(1)] \\ &= \mathbb{E}_\mu[\Pr_\mu(D = 1 \mid X, Y(1)) \mid Y(1)] \\ &= \Pr_\mu(D = 1 \mid Y(1)). \end{aligned}$$

Therefore $D \perp\!\!\!\perp A \mid Y(1)$ over μ , i.e., counterfactual equalized odds holds for the decision policy $\tau(x)$ over the distribution μ . Consequently $\mu \in \mathbf{E}$, and so $\mathbf{E}$ is closed and therefore universally measurable. $\square$

A.6.4 Shy Sets and Probes

We require a number of additional technical lemmata for the proof of Theorem 2.17. The probe must be constructed carefully, so that, on the utility scale, an arbitrary element of $\mathbf{Q}$ is absolutely continuous with respect to a typical perturbation. In addition, it is useful to show that a number of properties are generic to simplify certain aspects of the proof of Theorem 2.17. For instance, Lemma A.43 is used in Theorem 2.17 to show that a certain conditional expectation is generically well-defined, avoiding the need to separately treat certain corner cases.

⁵For a proof of this fact see, e.g., Brozius (2019).

Cor. A.39 concerns the construction of the probe used in the proof of Theorem 2.17. Lemmata A.44 to A.48 use Cor. A.39 to provide additional simplifications to the proof of Theorem 2.17.

A.6.4.1 Maximal Support

First, to construct the probe used in the proof of Theorem 2.17, we require elements $\mu \in \mathbf{Q}$ such that the densities f_μ have “maximal” support. To produce such distributions, we use the following measure-theoretic construction.

Definition A.37. Let $\{E_\alpha\}_{\alpha \in \Gamma}$ be an arbitrary collection of μ -measurable sets for some positive measure μ on a measure space $(M, \mathcal{M})$. We say that E is the *measure-theoretic union* of $\{E_\gamma\}_{\gamma \in \Gamma}$ if $\mu[E_\gamma \setminus E] = 0$ for all $\gamma \in \Gamma$ and $E = \bigcup_{i=1}^{\infty} E_{\gamma_i}$ for some countable subcollection $\{\gamma_i\} \subseteq \Gamma$.

While measure-theoretic unions themselves are known (cf. Silva (2008), Rudin (1991)), for completeness, we include a proof of their existence, which, to the best of our knowledge, is not found in the literature.

Lemma A.38. Let μ be a finite positive measure on a measure space $(V, \mathcal{M})$. Then an arbitrary collection $\{E_\gamma\}_{\gamma \in \Gamma}$ of μ -measurable sets has a measure-theoretic union.

Proof. For each countable subcollection $\Gamma' \subseteq \Gamma$, consider the “error term”

$$r(\Gamma') = \sup_{\gamma \in \Gamma} \mu \left[E_\gamma \setminus \bigcup_{\gamma' \in \Gamma'} E_{\gamma'} \right]$$

We claim that the infimum of $r(\Gamma')$ over all countable subcollections $\Gamma' \subseteq \Gamma$ must be zero.

For, toward a contradiction, suppose it were greater than or equal to $\epsilon > 0$. Choose any set E_{γ_1} such that $\mu[E_{\gamma_1}] \geq \epsilon$. Such a set must exist, since otherwise $r(\emptyset) < \epsilon$. Choose E_{γ_2} such that $\mu[E_{\gamma_2} \setminus E_{\gamma_1}] > \epsilon$. Again, some such set must exist, since otherwise $r(\{\gamma_1\}) < \epsilon$. Continuing in this way, we construct a countable collection $\{E_{\gamma_i}\}_{i \in \mathbb{N}}$.

Therefore, we see that

$$\mu[V] \geq \mu \left[\bigcup_{i=1}^n E_{\gamma_i} \right] = \sum_{i=1}^n \mu \left[E_{\gamma_i} \setminus \bigcup_{j=1}^i E_{\gamma_j} \right].$$

By construction, every term in the final sum is greater than or equal to ϵ , contradicting the fact that $\mu[V] < \infty$.

Therefore, there exist countable collections $\{\Gamma_n\}_{n \in \mathbb{N}}$ such that $r(\Gamma_n) < \frac{1}{n}$. It follows immediately that for all n

$$r\left(\bigcup_{n \in \mathbb{N}} \Gamma_n\right) \leq r(\Gamma_k)$$

for any fixed $k \in \mathbb{N}$. Consequently,

$$r\left(\bigcup_{n \in \mathbb{N}} \Gamma_n\right) = 0,$$

and $\bigcup_{n \in \mathbb{N}} \Gamma_n$ is countable. $\square$

The construction of the “maximal” elements used to construct the probe in the proof of Theorem 2.17 follows as a corollary of Lemma A.38

Corollary A.39. *There are measures $\mu_{\max,a} \in \mathbf{Q}$ such that for every $a \in \mathcal{A}$ and any $\mu \in \mathbf{K}$,*

$$\lambda[\text{SUPP}(f_{\mu,a}) \setminus \text{SUPP}(f_{\mu_{\max},a})] = 0.$$

Proof. Consider the collection $\{\text{SUPP}(f_{\mu,a})\}_{\mu \in \mathbf{K}}$. By Lemma A.38, there exists a countable collection of measures $\{\mu_i\}_{i \in \mathbb{N}}$ such that for any $\mu \in \mathbf{K}$,

$$\lambda\left[\text{SUPP}(f_{\mu,a}) \setminus \bigcup_{i=1}^{\infty} \text{SUPP}(f_{\mu_i,a})\right] = 0,$$

where, without loss of generality, we may assume that $\lambda[\text{SUPP}(f_{\mu_i,a})] > 0$ for all $i \in \mathbb{N}$. Such a sequence must exist, since, by the first hypothesis of Theorem 2.17, for every $a \in \mathcal{A}$, there exists $\mu \in \mathbf{Q}$ such that $\Pr_{\mu}(A = a) > 0$. Therefore, we can define the probability measure $\mu_{\max,a}$, where

$$\mu_{\max,a} = \sum_{i=1}^n 2^{-i} \cdot \frac{|\mu_i \upharpoonright_{A=a}|}{|\mu_i \upharpoonright_{A=a}|[\mathcal{K}]}.$$

It follows immediately by construction that

$$\text{SUPP}(f_{\mu_{\max},a}) = \bigcup_{i=1}^{\infty} \text{SUPP}(f_{\mu_i,a}),$$

and that $\mu_{\max,a} \in \mathbf{Q}$. $\square$

For notational simplicity, we refer to $\text{SUPP}(f_{\mu_{\max},a})$ as S_a throughout.

In the case of conditional principal fairness and path-specific fairness, we need a mild refinement of the previous result that accounts for ω .

Corollary A.40. *There are measures $\mu_{\max,a,w} \in \mathbf{Q}$ defined for every $w \in \mathcal{W} = \text{IMG}(\omega)$ and any $a \in \mathcal{A}$ such that for some $\nu \in \mathbf{K}$, $\Pr_\nu(W = w, A = a) > 0$. These measures have the property that for any $\mu \in \mathbf{K}$,*

$$\lambda[\text{SUPP}(f_{\mu',a,w}) \setminus \text{SUPP}(f_{\mu_{\max,a,w}})] = 0,$$

where $f_{\mu',a,w}$ is the density of the pushforward measure $(\mu' \upharpoonright_{W=w, A=a}) \circ u^{-1}$.

Recalling that $|\text{IMG}(\omega)| < \infty$, the proof is the same as Cor. A.39, and we analogously refer to $\text{SUPP}(f_{\mu_{\max,a,w}})$ as $S_{a,w}$. Here, we have assumed without loss of generality—as we continue to assume in the sequel—that for all $w \in \mathcal{W}$, there is some $\mu \in \mathbf{K}$ such that $\Pr_\mu(W = w) > 0$.

Remark A.41. Because their support is maximal, the hypotheses of Theorem 2.17, in addition to implying that $\mu_{\max,a}$ is well-defined for all $a \in \mathcal{A}$, also imply that $\Pr_{\mu_{\max,a}}(u(X) > 0) > 0$. In the case of conditional principal fairness, they further imply that $\Pr_{\mu_{\max,a}}(W = w) > 0$ for all $w \in \mathcal{W}$ and $a \in \mathcal{A}$. Likewise, in the case of path-specific fairness, they further imply that $\Pr_{\mu_{\max,a}}(W = w_i) > 0$ for $i = 0, 1$ and some $a \in \mathcal{A}$.

A.6.4.2 Shy Sets and Probes

In the following lemmata, we demonstrate that a number of useful properties are generic in $\mathbf{Q}$. We also demonstrate a short technical lemma, Lemma A.48, which allows us to use these generic properties to simplify the proof of Theorem 2.17.

We begin with the following lemma, which is useful in verifying that certain subspaces of $\mathbf{K}$ form probes.

Lemma A.42. *Let $\mathbf{W}$ be a non-trivial finite dimensional subspace of $\mathbf{K}$ such that $\nu[\mathcal{K}] = 0$ for all $\nu \in \mathbf{W}$. Then, there exists $\mu \in \mathbf{K}$ such that $\lambda_{\mathbf{W}}[\mathbf{Q} - \mu] > 0$.*

Proof. Set

$$\mu = \sum_{i=1}^n \frac{|\nu_i|}{|\nu_i|[\mathcal{K}]},$$

where $\nu_1, \dots, \nu_n$ form a basis of $\mathbf{W}$. Then, if $|\beta_i| \leq \frac{1}{|\nu_i|[\mathcal{K}]}$, it follows that

$$\mu + \sum_{i=1}^n \beta_i \cdot \nu_i \in \mathbf{Q}.$$

Since

$$\lambda_n \left[\prod_{i=1}^n \left[-\frac{1}{|\nu_i|[\mathcal{K}]}, \frac{1}{|\nu_i|[\mathcal{K}]} \right] \right] > 0,$$

it follows that $\lambda_{\mathbf{W}}[\mathbf{Q} - \mu] > 0$. □

Next we show that, given a $\nu \in \mathbf{Q}$, a generic element of $\mathbf{Q}$ “sees” events to which ν assigns non-zero probability. While Lemma A.45 alone in principle suffices for the proof of Theorem 2.17, we include Lemma A.43 both for conceptual clarity and to introduce at a high level the style of argument used in the subsequent lemmata and in the proof of Theorem 2.17 to show that a set is shy relative to $\mathbf{Q}$.

Lemma A.43. *For a Borel set $E \subseteq \mathcal{K}$, suppose there exists $\nu \in \mathbf{Q}$ such that $\nu[E] > 0$. Then the set of $\mu \in \mathbf{Q}$ such that $\mu[E] > 0$ is prevalent.*

Proof. First, we note that the set of $\mu \in \mathbf{Q}$ such that $\mu[E] = 0$ is closed and therefore universally measurable. For, if $\{\mu_i\}_{i \in \mathbb{N}} \subseteq \mathbf{Q}$ is a convergent sequence with limit μ , then

$$\begin{aligned} \mu[E] &= \lim_{n \rightarrow \infty} \mu_i[E] \\ &= \lim_{n \rightarrow \infty} 0 \\ &= 0. \end{aligned}$$

Now, if $\mu[E] > 0$ for all $\mu \in \mathbf{Q}$, there is nothing to prove. Therefore, suppose that there exists $\nu' \in \mathbf{Q}$ such that $\nu'[E] = 0$.

Next, consider the measure $\tilde{\nu} = \nu' - \nu$. Then, let $\mathbf{W} = \text{SPAN}(\tilde{\nu})$. Since $\tilde{\nu} \neq 0$ and

$$\tilde{\nu}[\mathcal{K}] = \nu'[\mathcal{K}] - \nu[\mathcal{K}] = 0,$$

it follows by Lemma A.42 that $\lambda_{\mathbf{W}}[\mathbf{Q} - \mu] > 0$ for some μ .

Now, for arbitrary $\mu \in \mathbf{Q}$, note that if $(\mu + \beta \cdot \tilde{\nu})[E] = 0$, then

$$\mu[E] - \beta \cdot \nu[E] = 0$$

i.e.,

$$\beta = \frac{\mu[E]}{\nu[E]}.$$

A singleton has null Lebesgue measure, and so the set of $\nu \in \mathbf{W}$ such that $(\mu + \nu)[E] = 0$ is $\lambda_{\mathbf{W}}$ -null. Therefore, by Prop. A.18, the set of $\mu \in \mathbf{Q}$ such that $\mu[E] = 0$ is shy relative to $\mathbf{Q}$, as desired. $\square$

While Lemma A.43 shows that a typical element of $\mathbf{Q}$ “sees” individual events, in the proof of Theorem 2.17, we require a stronger condition, namely, that a typical element of $\mathbf{Q}$ “sees” certain uncountable collections of events. To demonstrate this more complex property, we require the following technical result, which is closely related to the real analysis folk theorem that any convergent uncountable “sum” can contain only countably many non-zero terms. (See, e.g., Benji (2020).)

Lemma A.44. *Suppose μ is a totally bounded measure on $(V, \mathcal{M})$, f and g are μ -measurable real-valued functions, and $g \neq 0$ μ -a.e. Then the set*

$$Z_\beta = \{v \in V : f(v) + \beta \cdot g(v) = 0\}$$

has non-zero μ measure for at most countably many $\beta \in \mathbb{R}$.

Proof. First, we show that for any countable collection $\{\beta_i\}_{i \in \mathbb{N}} \subseteq \mathbb{R}$, the sum $\sum_{i=1}^{\infty} \mu[Z_{\beta_i}]$ converges. Then, we show how this implies that $\mu[Z_\beta] = 0$ for all but countably many $\beta \in \mathbb{R}$.

First, we note that for distinct $\beta, \beta' \in \mathbb{R}$,

$$Z_\beta \cap Z_{\beta'} \subseteq \{v \in V : (\beta - \beta') \cdot g(v) = 0\}.$$

Now, by hypothesis,

$$\mu[\{v \in V : g(v) = 0\}] = 0,$$

and since $\beta - \beta' \neq 0$, it follows that

$$\mu[\{v \in V : (\beta - \beta') \cdot g(v) = 0\}] = 0$$

as well. Consequently, it follows that if $\{Z_{\beta_i}\}_{i \in \mathbb{N}}$ is a countable collection of distinct elements of $\mathbb{R}$,

then

$$\begin{aligned}\sum_{i=1}^{\infty} \mu[Z_{\beta_i}] &= \mu \left[\bigcup_{i=1}^{\infty} Z_{\beta_i} \right] \\ &\leq \mu[V] \\ &< \infty.\end{aligned}$$

To see that this implies that $\mu[Z_{\beta}] > 0$ for only countably many $\beta \in \mathbb{R}$, let $G_{\epsilon} \subseteq \mathbb{R}$ consist of those β such that $\mu[Z_{\beta}] \geq \epsilon$. Then G_{ϵ} must be finite for all $\epsilon > 0$, since otherwise we could form a collection $\{\beta_i\}_{i \in \mathbb{N}} \subseteq G_{\epsilon}$, in which case

$$\sum_{i=1}^{\infty} \mu[Z_{\beta_i}] \geq \sum_{i=1}^{\infty} \epsilon = \infty,$$

contrary to what was just shown. Therefore,

$$\{\beta \in \mathbb{R} : \mu[Z_{\beta}] > 0\} = \bigcup_{i=1}^{\infty} G_{1/i}$$

is countable. □

We now apply Lemma A.44 to the proof of the following lemma, which states, informally, that, under a generic element of $\mathbf{Q}$, $u(X)$ is supported everywhere it is supported under some particular fixed element of $\mathbf{Q}$. For instance, Lemma A.44 can be used to show that for a generic element of $\mathbf{Q}$, the density of $u(X) \mid A = a$ is positive $\lambda \upharpoonright_{S_a}$ -a.e.

Lemma A.45. *Let $\nu \in \mathbf{Q}$ and suppose ν is supported on E , i.e., $\nu[\mathcal{K} \setminus E] = 0$. Then the set of $\mu \in \mathbf{Q}$ such that $\nu \circ u^{-1} \ll (\mu \upharpoonright_E) \circ u^{-1}$ is prevalent relative to $\mathbf{Q}$.*

Lemma A.45 states, informally, that for generic $\mu \in \mathbf{Q}$, $f_{\mu \upharpoonright_E}$ is supported everywhere f_{ν} is supported.

Proof. We begin by showing that the set of $\mu \in \mathbf{Q}$ such that $\nu \circ u^{-1} \ll (\mu \upharpoonright_E) \circ u^{-1}$ is Borel, and therefore universally measurable. Then, we construct a probe $\mathbf{W}$ and use it to show that this collection is finitely shy.

To begin, let U_q denote the set of $\mu \in \mathbf{Q}$ such that

$$\nu \circ u^{-1}[\{|f_{\mu \upharpoonright_E}| = 0\}] < q.$$

We note that U_q is open. For, if $\mu \in U_q$, then there exists some $r > 0$ such that

$$\nu \circ u^{-1}[\{|f_\mu|_E| < r\}] < q.$$

Let

$$\epsilon = q - \nu \circ u^{-1}[\{|f_\mu|_E| < r\}].$$

Now, since $\nu \circ u^{-1} \ll \lambda$, there exists a δ such that if $\lambda[E'] < \delta$, then $\nu \circ u^{-1}[E'] < \epsilon$. Choose μ' arbitrarily so that $|\mu - \mu'|[\mathcal{K}] < \delta \cdot r$. Then, by Markov's inequality, we have that

$$\lambda[\{|f_\mu|_E - f_{\mu'}|_E| > r\}] < \delta,$$

i.e.,

$$\nu \circ u^{-1}[\{|f_\mu|_E - f_{\mu'}|_E| > r\}] < \epsilon.$$

Now, we note that by the triangle inequality, wherever $|f_{\mu'}|_E| = 0$, either $|f_\mu|_E| < r$ or $|f_\mu|_E - f_{\mu'}|_E| > r$. Therefore

$$\begin{aligned} \lambda[\{|f_{\mu'}|_E| = 0\}] &\leq \nu \circ u^{-1}[\{|f_\mu|_E - f_{\mu'}|_E| > r\}] + \mu \circ u^{-1}[\{|f_\mu|_E| < r\}] \\ &< \epsilon + \mu \circ u^{-1}[\{|f_\mu|_E| < r\}] \\ &< q. \end{aligned}$$

We conclude that $\mu' \in U_q$, and so U_q is open.

Note that $\nu \circ u^{-1} \ll (\mu|_E) \circ u^{-1}$ if and only if

$$\lambda[\text{SUPP}(f_\nu) \setminus \text{SUPP}(f_\mu|_E)] = 0.$$

By the definition of the support of a function, $\lambda|_{\text{SUPP}(f_\mu)} \ll \mu \circ u^{-1}$. Therefore, it follows that

$$\lambda[\text{SUPP}(f_\mu) \setminus \text{SUPP}(f_\nu|_E)] = 0$$

if and only if

$$\mu \circ u^{-1}[\text{SUPP}(f_\mu) \setminus \text{SUPP}(f_\nu|_E)] = 0.$$

Then, it follows immediately that the set of $\nu \in \mathbf{Q}$ such that $\mu \circ u^{-1} \ll (\nu|_E) \circ u^{-1}$ is $\bigcap_{i=1}^n U_{1/i}$, which is, by construction, Borel, and therefore universally measurable.

Now, since

$$\Pr_\nu(u(X) < t) = \int_{-\infty}^t f_\nu d\lambda$$

is a continuous function of t , by the intermediate value theorem, there exists t such that $\Pr_\nu(u(X) \in S_0) = \Pr_\nu(u(X) \in S_1)$, where $S_0 = \text{SUPP}(f_\nu) \cap (-\infty, t)$ and $S_1 = \text{SUPP}(f_\nu) \cap [t, \infty)$. Then, we let

$$\tilde{\nu}[E'] = \int_{E'} \mathbb{1}_{u^{-1}(S_0)} - \mathbb{1}_{u^{-1}(S_1)} d\nu.$$

Take $\mathbf{W} = \text{SPAN}(\tilde{\nu})$. Since $\tilde{\nu} \neq 0$ and $\tilde{\nu}[\mathcal{K}] = 0$, we have by Lemma A.42 that $\lambda_{\mathbf{W}}[\mathbf{Q} - \mu] > 0$ for some μ .

By the definition of a density, $f_{\tilde{\nu}}$ is positive ($\tilde{\nu} \circ u^{-1}$)-a.e. Consequently, by the definition of $\tilde{\nu}$, $f_{\tilde{\nu}}$ is non-zero ($\mu \circ u^{-1}$)-a.e. Therefore, by Lemma A.44, there exist only countably many $\beta \in \mathbb{R}$ such that the density of $(\mu + \beta \cdot \tilde{\nu}) \circ u^{-1}$ equals zero on a set of positive $\mu \circ u^{-1}$ -measure. Since countable sets have λ -measure zero and ν is arbitrary, the set of $\mu \in \mathbf{Q}$ such that $\nu \circ u^{-1} \ll (\mu \upharpoonright_E) \circ u^{-1}$ is prevalent relative to $\mathbf{Q}$ by Prop. A.18. $\square$

The following definition and technical lemma are needed to extend Theorem 2.17 to the cases of conditional principal fairness and path-specific fairness, which involve additional conditioning on $W = \omega(X)$. In particular, one corner case we wish to avoid in the proof of Theorem 2.17 is when the decision policy is non-trivial (i.e., some individuals receive a positive decision and others do not) but from the perspective of each ω -stratum, the policy is trivial (i.e., everyone in the stratum receives a positive or negative decision). Definition A.46 formalizes this pathology, and Lemma A.47 shows that this issue—under a mild hypothesis—does not arise for a generic element of $\mathbf{Q}$.

Definition A.46. We say that $\mu \in \mathbf{Q}$ *overlaps utilities* when, for any budget-exhausting multiple threshold policy $\tau(x)$, if

$$0 < \mathbb{E}_\mu[\tau(X)] < 1,$$

then there exists $w \in \mathcal{W}$ such that

$$0 < \mathbb{E}_\mu[\tau(X) \mid W = w] < 1.$$

If there exists a budget-exhausting multiple threshold policy $\tau(x)$ such that

$$0 < \mathbb{E}_\mu[\tau(X)] < 1,$$

but for all $w \in \mathcal{W}$,

$$\mathbb{E}_\mu[\tau(X) \mid W = w] \in \{0, 1\},$$

then we say that $\tau(x)$ *splits utilities* over μ .

Informally, having overlapped utilities prevents a budget-exhausting threshold policy from having thresholds that fall on the utility scale exactly between the strata induced by ω —i.e., a threshold policy that splits utilities. This is almost a generic condition in $\mathbf{Q}$, as we shown in Lemma A.47.

Lemma A.47. *Let $0 < b < 1$. Suppose that for all $w \in \mathcal{W}$ there exists $\mu \in \mathbf{Q}$ such that $\Pr_\mu(u(X) > 0, W = w) > 0$. Then almost every $\mu \in \mathbf{Q}$ overlaps utilities.*

Proof. Our goal is to show that the set $\mathbf{E}'$ of measures $\mu \in \mathbf{Q}$ such that there exists a splitting policy $\tau(x)$ is shy. To simplify the proof, we divide and conquer, showing that the set $\mathbf{E}_\Gamma$ of measures $\mu \in \mathbf{Q}$ such that there exists a splitting policy where the thresholds fall below $w \in \Gamma \subseteq \mathcal{W}$ and above $w \notin \Gamma$ is Borel, before constructing a probe that shows that it is shy. Then, we argue that $\mathbf{E}' = \bigcup_{\Gamma \subseteq \mathcal{W}} \mathbf{E}_\Gamma$, which shows that $\mathbf{E}'$ is shy.

We begin by considering the linear map $\Phi : \mathbf{K} \rightarrow \mathbb{R} \times \mathbb{R}^\mathcal{W}$ given by

$$\Phi(\mu) = (\Pr_\mu(u(X) = 0), (\Pr_\mu(W = w))_{w \in \mathcal{W}}).$$

For any $\Gamma \subseteq \mathcal{W}$, the sets

$$F_\Gamma^{\text{Up}} = \{x \in \mathbb{R} \times \mathbb{R}^\mathcal{W} : x_0 \geq b, b = \sum_{w \in \Gamma} x_w\},$$

$$F_\Gamma^{\text{Lo}} = \{x \in \mathbb{R} \times \mathbb{R}^\mathcal{W} : x_0 \leq b, x_0 = \sum_{w \in \Gamma} x_w\},$$

are closed by construction. Therefore, since Φ is continuous,

$$\mathbf{E}_\Gamma = \mathbf{Q} \cap \Phi^{-1} \left(\bigcup_{\Gamma \subseteq \mathcal{W}} F_\Gamma^{\text{Up}} \cup F_\Gamma^{\text{Lo}} \right) \quad (\text{A.9})$$

is closed, and therefore universally measurable.

Note that by our hypothesis and Cor. A.40, for all $w \in \mathcal{W}$ there exists some $a_w \in \mathcal{A}$ such that

$$\Pr_{\mu_{\max, a_w, w}}(u(X) > 0).$$

We use this to show that $\mathbf{E}_\Gamma$ is shy. Pick $w^* \in \mathcal{W}$ arbitrarily, and consider the measures ν_w for

$w \neq w^*$ defined by

$$\nu_w = \frac{\mu_{\max, a_w, w} \upharpoonright_{u(X) > 0}}{\Pr_{\mu_{\max, a_w, w}}(u(X) > 0)} - \frac{\mu_{\max, a_{w^*}, w^*} \upharpoonright_{u(X) > 0}}{\Pr_{\mu_{\max, a_{w^*}, w^*}}(u(X) > 0)}.$$

We note that $\nu_w[\mathcal{K}] = 0$ by construction. Therefore, if $\mathbf{W}_w = \text{SPAN}(\nu_w)$, then $\lambda_{\mathbf{W}_w}[\mathbf{Q} - \mu_w] > 0$ for some μ_w by Lemma A.42.

Moreover, we have that $\Pr_\nu(u(X) > 0) = 0$ for all $\nu \in \mathbf{W}_w$, i.e.,

$$\Pr_\mu(u(X) > 0) = \Pr_{\mu+\nu}(u(X) > 0).$$

Now, since $0 < b < 1$ and ω partitions $\mathcal{X}$, it follows that

$$\mathbf{E}_{\mathcal{W}} = \mathbf{E}_\emptyset = \emptyset.$$

Since $\lambda_{\mathbf{W}}[\emptyset] = 0$ for any subspace $\mathbf{W}$, we can assume without loss of generality that $\Gamma \neq \mathcal{W}, \emptyset$.

In that case, there exists $w_\Gamma \in \mathcal{W}$ such that if $w^* \in \Gamma$, then $w_\Gamma \notin \Gamma$, and vice versa. Without loss of generality, assume $w_\Gamma \in \Gamma$ and $w^* \notin \Gamma$. It then follows that for arbitrary $\mu \in \mathbf{Q}$,

$$\Phi(\mu + \beta \cdot \nu_{w_\Gamma}) = \Phi(\mu) + \beta \cdot \mathbf{e}_{w_\Gamma} - \beta \cdot \mathbf{e}_{w^*},$$

where $\mathbf{e}_w$ is the basis vector corresponding to $w \in \mathcal{W}$. From this, it follows immediately by Eq. (A.6.4.2) that

$$\mu + \beta \cdot \nu_{w_\Gamma} \in \mathbf{E}_\Gamma$$

only if

$$\beta = \min(b, \Pr_\mu(u(X) > 0)) - \sum_{w \in \Gamma} \Pr_\mu(W = w).$$

This is a measure zero subset of $\mathbb{R}$, and so it follows that

$$\lambda_{\mathbf{W}_{w_\Gamma}}[\mathbf{E}_\Gamma - \mu] = 0$$

for all $\mu \in \mathbf{K}$. Therefore, by Prop. A.18, $\mathbf{E}_\Gamma$ is shy in $\mathbf{Q}$. Taking the union over $\Gamma \subseteq \mathcal{W}$, it follows by Prop. A.16 that $\bigcup_{\Gamma \subseteq \mathcal{W}} \mathbf{E}_\Gamma$ is shy.

Now, we must show that $\mathbf{E}' = \bigcup_{\Gamma \subseteq \mathcal{W}} \mathbf{E}_\Gamma$. By construction, $\mathbf{E}_\Gamma \subseteq \mathbf{E}'$, since the policy $\tau(x) = \mathbb{1}_{\omega(x) \in \Gamma}$ is budget-exhausting and separates utilities. To see the reverse inclusion, suppose $\mu \in \mathbf{E}'$, i.e., that there exists a budget-exhausting multiple threshold policy $\tau(x)$ that splits utilities over μ .

Then, let

$$\Gamma_\mu = \{w \in \mathcal{W} : \mathbb{E}_\mu[\tau(X) \mid W = w] = 1\}.$$

Since $\tau(x)$ is budget-exhausting, it follows immediately that $\mu \in \mathbf{E}_{\Gamma_\mu}$. Therefore, $\mathbf{E}' = \bigcup_{\Gamma \subseteq \mathcal{W}} \mathbf{E}_\Gamma$, and so $\mathbf{E}'$ is shy, as desired. $\square$

We conclude our discussion of shyness and shy sets with the following general lemma, which simplifies relative prevalence proofs by showing that one can, without loss of generality, restrict one's attention to the elements of the shy set itself in applying Prop. A.18.

Lemma A.48. *Suppose E is a universally measurable subset of a convex, completely metrizable set C in a topological vector space V . Suppose that for some finite-dimensional subspace V' , $\lambda_{V'}[C + v_0] > 0$ for some $v_0 \in V$. If, in addition, for all $v \in E$,*

$$\lambda_{V'}[\{v' \in V' : v + v' \in E\}] = 0, \tag{A.10}$$

then it follows that E is shy relative to C .

Proof. Let v be arbitrary. Then, either $(v + V') \cap E$ is empty or not.

First, suppose it is empty. Since $\lambda_{V'}[\emptyset] = 0$ by definition, it follows immediately that in this case $\lambda_{V'}[E - v] = 0$.

Next, suppose the intersection is not empty, and let $v + v^* \in E$ for some fixed $v^* \in V'$. It follows that

$$\begin{aligned} \lambda_{V'}[E - v] &= \lambda_{V'}[\{v' \in V' : v + v' \in E\}] \\ &= \lambda_{V'}[\{v' \in V' : (v + v^*) + v' \in E\}] \\ &= 0, \end{aligned}$$

where the first equality follows by definition; the second equality follows by the translation invariance of $\lambda_{V'}$, and the fact that $v^* + V' = V'$; and the final inequality follows from Eq. (A.10).

Therefore $\lambda_{V'}[E - v] = 0$ for arbitrary v , and so E is shy. $\square$

A.6.5 Proof of Theorem 2.17

Using the lemmata above, we can prove Theorem 2.17. We briefly summarize what has been established so far:

- **Lemma A.22:** The set $\mathbf{K}$ of $\mathcal{U}$ -fine distributions on $\mathcal{K}$ is a Banach space;
- **Lemma A.30:** The subset $\mathbf{Q}$ of $\mathcal{U}$ -fine probability measures on $\mathcal{K}$ is a convex, completely metrizable subset of $\mathbf{K}$;
- **Lemma A.36:** The subset $\mathbf{E}$ of $\mathbf{Q}$ is a universally measurable subset of $\mathbf{K}$, where $\mathbf{E}$ is the set consisting of $\mathcal{U}$ -fine probability measures over which there exists a policy satisfying counterfactual equalized odds (resp., conditional principal fairness, or path-specific fairness) that is not strongly Pareto dominated.

Therefore, to apply Prop. A.18, it follows that what remains is to construct a probe $\mathbf{W}$ and show that $\lambda_{\mathbf{W}}[\mathbf{Q} + \mu_0] > 0$ for some $\mu_0 \in \mathbf{K}$ but $\lambda_{\mathbf{W}}[\mathbf{E} + \mu] = 0$ for all $\mu \in \mathbf{K}$.

Proof. We divide the proof into three pieces. First, we illustrate how to construct the probe $\mathbf{W}$ from a particular collection of distributions $\{\nu_a^{\text{Up}}, \nu_a^{\text{Lo}}\}_{a \in \mathcal{A}}$. Second, we show that $\lambda_{\mathbf{W}}[\mathbf{E} + \mu] = 0$ for all $\mu \in \mathbf{K}$. For notational and expository simplicity, we focus in these first two sections on the case of counterfactual equalized odds. Therefore, in the third section, we show how to generalize the argument to conditional principal fairness and path-specific fairness.

Construction of the probe. We will construct our probe to address two different cases. We recall that, by Cor. A.9, any policy that is not strongly Pareto dominated must be a budget-exhausting multiple threshold policy with non-negative thresholds. In the first case, we consider when the candidate budget-exhausting multiple threshold policy is $\mathbb{1}_{u(x) > 0}$. By perturbing the underlying distribution by $\nu \in \mathbf{W}^{\text{Lo}}$, we will be able to break the balance requirements implied by Eq. (2.4). In the second case, we treat the possibility that the candidate budget-exhausting multiple threshold policy has a non-trivial positive threshold for at least one group. By perturbing the underlying distribution by $\nu \in \mathbf{W}^{\text{Up}}$ for an alternative set of perturbations $\mathbf{W}^{\text{Up}}$, we will again be able to break the balance requirements.

More specifically, to construct our probe $\mathbf{W} = \mathbf{W}^{\text{Up}} + \mathbf{W}^{\text{Lo}}$, we want $\mathbf{W}^{\text{Up}}$ and $\mathbf{W}^{\text{Lo}}$ to have a number of properties. In particular, for all $\nu \in \mathbf{W}$, perturbation by ν should not affect whether the

underlying distribution is a probability distribution, and should not affect how much of the budget is available to budget-exhausting policies. Specifically, for all $\nu \in \mathbf{W}$,

$$\int_{\mathcal{K}} 1 \, d\nu = 0, \quad (\text{A.11})$$

and

$$\int_{\mathcal{K}} \mathbb{1}_{u(X) > 0} \, d\nu = 0. \quad (\text{A.12})$$

In fact, the amount of budget available to budget-exhausting policies will not change within group, i.e., for all $a \in \mathcal{A}$ and $\nu \in \mathbf{W}$,

$$\int_{\mathcal{K}} \mathbb{1}_{u(X) > 0, A=a} \, d\nu = 0. \quad (\text{A.13})$$

Additionally, for some distinguished $y_0, y_1 \in \mathcal{Y}$, non-zero perturbations in $\nu^{\text{Lo}} \in \mathbf{W}^{\text{Lo}}$ should move mass between y_0 and y_1 . That is, they should have the property that if $\Pr_{|\nu^{\text{Lo}}|}(A = a) > 0$, then

$$\int_{\mathcal{K}} \mathbb{1}_{u(X) < 0, Y=y_i, A=a} \, d\nu^{\text{Lo}} \neq 0. \quad (\text{A.14})$$

Finally, perturbations in $\mathbf{W}^{\text{Up}}$ should have the property that for any non-trivial $t > 0$, some mass is moved either above or below $t > 0$. More precisely, for any $\mu \in \mathbf{Q}$ and any t such that

$$0 < \Pr_{\mu}(u(X) > t \mid A = a) < 1,$$

if $\nu^{\text{Up}} \in \mathbf{W}^{\text{Up}}$ is such that $\Pr_{|\nu^{\text{Up}}|}(A = a) > 0$, then

$$\int_{\mathcal{K}} \mathbb{1}_{u(X) > t, A=a} \, d\nu^{\text{Up}} \neq 0. \quad (\text{A.15})$$

To carry out the construction, choose distinct $y_0, y_1 \in \mathcal{Y}$. Then, since

$$\mu_{\max, a} \circ u^{-1}[S_a \cap [0, r_a)] - \mu_{\max, a} \circ u^{-1}[S_a \cap [r_a, \infty))]$$

is a continuous function of r_a , it follows by the intermediate value theorem that we can partition S_a into three pieces,

$$S_a^{\text{Lo}} = S_a \cap (-\infty, 0),$$

$$S_{a,0}^{\text{Up}} = S_a \cap [0, r_a),$$

$$S_{a,1}^{\text{Up}} = S_a \cap [r_a, \infty),$$

so that

$$\Pr_{\mu_{\max,a}} \left(u(X) \in S_{a,0}^{\text{Up}} \right) = \Pr_{\mu_{\max,a}} \left(u(X) \in S_{a,1}^{\text{Up}} \right).$$

Recall that $\mathcal{K} = \mathcal{X} \times \mathcal{Y}$. Let $\pi_{\mathcal{X}} : \mathcal{K} \rightarrow \mathcal{X}$ denote projection onto $\mathcal{X}$, and $\gamma_y : \mathcal{X} \rightarrow \mathcal{K}$ be the injection $x \mapsto (x, y)$. We define

$$\begin{aligned} \nu_a^{\text{Up}}[E] &= \mu_{\max,a} \circ (\gamma_{y_1} \circ \pi_{\mathcal{X}})^{-1} \left[E \cap u^{-1} \left(S_{a,1}^{\text{Up}} \right) \right], \\ &\quad - \mu_{\max,a} \circ (\gamma_{y_1} \circ \pi_{\mathcal{X}})^{-1} \left[E \cap u^{-1} \left(S_{a,0}^{\text{Up}} \right) \right], \\ \nu_a^{\text{Lo}}[E] &= \mu_{\max,a} \circ (\gamma_{y_1} \circ \pi_{\mathcal{X}})^{-1} \left[E \cap u^{-1} \left(S_a^{\text{Lo}} \right) \right] \\ &\quad - \mu_{\max,a} \circ (\gamma_{y_0} \circ \pi_{\mathcal{X}})^{-1} \left[E \cap u^{-1} \left(S_a^{\text{Lo}} \right) \right]. \end{aligned}$$

By construction, ν_a^{Up} concentrates on

$$\{y_1\} \times u^{-1}(S_a \cap [0, \infty)),$$

while ν_a^{Lo} concentrates on

$$\{y_0, y_1\} \times u^{-1}(S_a \cap (-\infty, 0)).$$

Moreover, if we set

$$\mathbf{W}^{\text{Up}} = \text{SPAN}(\nu_a^{\text{Up}})_{a \in \mathcal{A}},$$

$$\mathbf{W}^{\text{Lo}} = \text{SPAN}(\nu_a^{\text{Lo}})_{a \in \mathcal{A}},$$

then it is easy to see that Eqs. (A.11) to (A.14) will hold. The only non-trivial case is Eq. (A.15).

However, by Cor. A.39, the support of $f_{\mu_{\max,a}}$ is maximal. That is, for $\mu \in \mathbf{Q}$, if

$$0 < \Pr_{\mu}(u(X) > t \mid A = a, u(X) > 0) < 1,$$

then it follows that $0 < t < \sup S_a$. Either $t \leq r_a$ or $t > r_a$. First, assume $t \leq r_a$; then, it follows by the construction of ν_a^{Up} that

$$\begin{aligned} \nu_a^{\text{Up}} \circ u^{-1}[(t, \infty)] &= \int_{r_a}^{\infty} f_{\max,a} \, d\lambda - \int_t^{r_a} f_{\max,a} \, d\lambda \\ &> \int_{r_a}^{\infty} f_{\max,a} \, d\lambda - \int_0^{r_a} f_{\max,a} \, d\lambda \\ &= 0. \end{aligned}$$

Similarly, if $t > r_a$,

$$\begin{aligned}\nu_a^{\text{Up}} \circ u^{-1}[(t, \infty)] &= \int_t^\infty f_{\max,a} \, d\lambda \\ &> \int_{\sup S_a}^\infty f_{\max,a} \, d\lambda \\ &= 0.\end{aligned}$$

Therefore Eq. (A.15) holds.

Since $\mathbf{W}$ is non-trivial⁶ and $\nu[\mathcal{K}] = 0$ for all $\nu \in \mathbf{W}$, it follows by Lemma A.42 that $\lambda_{\mathbf{W}}[\mathbf{Q} - \mu] > 0$ for some $\mu \in \mathbf{K}$.

Shyness. Recall that, by Prop. A.16, a set E is shy if and only if, for an arbitrary shy set E' , $E \setminus E'$ is shy. By Lemma A.43, a generic element of $\mu \in \mathbf{Q}$ satisfies $\Pr_\mu(u(X) > 0, Y(1) = y_i, A = a) > 0$ for $i = 0, 1$, and $a \in \mathcal{A}$. Likewise, by Lemma A.45, a generic $\mu \in \mathbf{Q}$ satisfies $\nu_a^{\text{Up}} \circ u^{-1} \ll (\mu \upharpoonright_{\mathcal{X} \times \{y_1\}}) \circ u^{-1}$. Therefore, to simplify our task and recalling Remark A.41, we may instead demonstrate the shyness of the set of $\mu \in \mathbf{Q}$ such that:

- There exists a budget-exhausting multiple threshold policy $\tau(x)$ with non-negative thresholds satisfying counterfactual equalized odds over μ ;
- For $i = 0, 1$,

$$\Pr_\mu(u(X) > 0, A = a, Y(1) = y_i) > 0; \tag{A.16}$$

- For all $a \in \mathcal{A}$,

$$\nu_a^{\text{Up}} \circ u^{-1} \ll (\mu \upharpoonright_{\alpha^{-1}(a) \times \{y_1\}}) \circ u^{-1}. \tag{A.17}$$

By a slight abuse of notation, we continue to refer to this set as $\mathbf{E}$. We note that, by the construction of $\mathbf{W}$, Eq. (A.16) is not affected by perturbation by $\nu \in \mathbf{W}$, and Eq. (A.17) is not affected by perturbation by $\nu^{\text{Lo}} \in \mathbf{W}$.

In particular, by Lemma A.48, it suffices to show that $\lambda_{\mathbf{W}}[\mathbf{E} - \mu] = 0$ for $\mu \in \mathbf{E}$.

Therefore, let $\mu \in \mathbf{E}$ be arbitrary. Let the budget-exhausting multiple threshold policy satisfying

⁶In general, some or all of the ν^{Lo} may be zero depending on the λ -measure of S_a^{Lo} . However, as noted in Remark A.41, the $\nu_{a,i}^{\text{Up}}$ cannot be zero, since $\Pr_{\mu_{\max,a}}(u(X) > 0) > 0$ for all $a \in \mathcal{A}$. Therefore $\mathbf{W} \neq \{0\}$.

counterfactual equalized odds over it be $\tau(x)$, so that

$$\mathbb{E}_\mu[\tau(X)] = \min(b, \Pr_\mu(u(X) > 0)),$$

with thresholds $\{t_a\}_{a \in \mathcal{A}}$. We split into two cases based on whether $\tau(X) = \mathbb{1}_{u(X) > 0}$ μ -a.s. or not.

In both cases, we make use of the following two useful observations.

First, note that as $\mathbf{E} \subseteq \mathbf{Q}$, if $\mu + \nu$ is not a probability measure, then $\mu + \nu \notin \mathbf{E}$. Therefore, without loss of generality, we assume throughout that $\mu + \nu$ is a probability measure.

Second, suppose $\tau'(x)$ is a policy satisfying counterfactual equalized odds over some $\nu \in \mathbf{Q}$. Then, if $0 < \mathbb{E}_\mu[\tau'(X)] < 1$, it follows that for all $a \in \mathcal{A}$,

$$0 < \mathbb{E}_\mu[\tau'(X) \mid A = a] < 1. \quad (\text{A.18})$$

For, suppose not. Then, without loss of generality, there must be $a_0, a_1 \in \mathcal{A}$ such that

$$\mathbb{E}_\mu[\tau'(X) \mid A = a_0] = 0$$

and

$$\mathbb{E}_\mu[\tau'(X) \mid A = a_1] > 0.$$

But then, by the law of iterated expectation, there must be some $\mathcal{Y}' \subseteq \mathcal{Y}$ such that $\mu[\mathcal{X} \times \mathcal{Y}'] > 0$ and so,

$$\mathbb{1}_{\mathcal{X} \times \mathcal{Y}'} \cdot \mathbb{E}_\mu[\tau'(X) \mid A = a_1, Y(1)] > 0 = \mathbb{1}_{\mathcal{X} \times \mathcal{Y}'} \cdot \mathbb{E}_\mu[\tau'(X) \mid A = a_0, Y(1)],$$

contradicting the fact that $\tau'(x)$ satisfies counterfactual equalized odds over μ . Therefore, in what follows, we can assume that Eq. (A.18) holds.

Our goal is to show that $\lambda_{\mathbf{W}}[\mathbf{E} - \mu] = 0$.

Case 1 ($\tau(X) = \mathbb{1}_{u(X) > 0}$). We argue as follows. First, we show that $\mathbb{1}_{u(X) > 0}$ is the unique budget-exhausting multiple threshold policy with non-negative thresholds over $\mu + \nu$ for all $\nu \in \mathbf{W}$. Then, we show that the set of $\nu \in \mathbf{W}$ such that $\mathbb{1}_{u(x) > 0}$ satisfies counterfactual equalized odds over $\mu + \nu$ is a $\lambda_{\mathbf{W}}$ -null set.

We begin by observing that $\mathbf{W}^{\text{Lo}} \neq \{0\}$. For, if that were the case, then Eq. (A.18) would not hold for $\tau(x)$.

Next, we note that by Eq. (A.12), for any $\nu \in \mathbf{W}$,

$$\Pr_{\mu+\nu}(u(X) > 0) = \Pr_{\mu}(u(X) > 0)$$

and so

$$\mathbb{E}_{\mu+\nu}[\mathbb{1}_{u(X)>0}] = \min(b, \Pr_{\mu+\nu}(u(X) > 0)).$$

If $\tau'(x)$ is a feasible multiple threshold policy with non-negative thresholds and $\tau'(X) \neq \mathbb{1}_{u(X)>0}$ $(\mu + \nu)$ -a.s., then, as a consequence,

$$\mathbb{E}_{\mu+\nu}[\tau'(X)] < \Pr_{\mu+\nu}(u(X) > 0) \leq b.$$

Therefore, it follows that $\mathbb{1}_{u(X)>0}$ is the unique budget-exhausting multiple threshold policy over $\mu + \nu$ with non-negative thresholds.

Now, note that if counterfactual equalized odds holds with decision policy $\tau(x) = \mathbb{1}_{u(x)>0}$, then, by Eq. (2.7) and Lemma A.21, we must have that

$$\Pr_{\mu+\nu}(u(X) > 0 \mid A = a, Y(1) = y_1) = \Pr_{\mu+\nu}(u(X) > 0 \mid A = a', Y(1) = y_1)$$

for $a, a' \in \mathcal{A}$.⁷

Now, we will show that a typical element of $\mathbf{W}$ breaks this balance requirement. Choose a^* such that $\nu_{a^*}^{\text{Lo}} \neq 0$. Recall that ν is fixed, and let $\nu' = \nu - \beta_{a^*}^{\text{Lo}} \cdot \nu_{a^*}^{\text{Lo}}$. Let

$$p_a = \Pr_{\mu+\nu'}(u(X) > 0 \mid A = a', Y(1) = y_1).$$

Note that it cannot be the case that $p_a = 0$ for all $a \in \mathcal{A}$, since, by Eq. (A.16),

$$\Pr_{\mu+\nu'}(u(X) > 0 \mid Y(1) = y_1) > 0.$$

Therefore, by the foregoing discussion, either $p_{a^*} > 0$ or $p_{a^*} = 0$ and we can choose $a' \in \mathcal{A}$ such that $p_{a'} > 0$. Since the $\nu_a^{\text{Lo}}, \nu_{a,i}^{\text{Up}}$ are all mutually singular, it follows that counterfactual equalized odds can only hold over $\mu + \nu$ if

$$p_{a'} = \Pr_{\mu+\nu}(u(X) > 0 \mid A = a^*, Y(1) = y_1).$$

⁷To ensure that both quantities are well-defined, here and throughout the remainder of the proof we use the fact that by Eqs. (A.13) and (A.16), $\Pr_{\mu+\nu}(u(X) > 0, A = a, Y(1) = y_1) > 0$.

Now, we observe that by Lemma A.21, that

$$\Pr_{\mu+\nu}(u(X) > 0 \mid A = a^*, Y(1) = y_1) = \frac{\eta}{\pi + \beta_{a^*}^{\text{Lo}} \cdot \rho}$$

where

$$\begin{aligned}\eta &= \Pr_{\mu}(u(X) > 0, A = a^*, Y(1) = y_1) \\ \pi &= \Pr_{\mu}(A = a^*, Y(1) = y_1), \\ \rho &= \int_{\mathcal{K}} \mathbb{1}_{A=a^*, Y(1)=y_1} d\nu_{a^*}^{\text{Lo}}.\end{aligned}$$

since

$$\begin{aligned}0 &= \int_{\mathcal{K}} \mathbb{1}_{u(X) > 0, A=a^*, Y(1)=y_1} d\nu_{a^*}^{\text{Lo}}, \\ 0 &\neq \int_{\mathcal{K}} \mathbb{1}_{A=a^*, Y(1)=y_1} d\nu_{a^*}^{\text{Lo}}.\end{aligned}$$

Here, the equality follows by the fact that ν^{Lo} is supported on $S_a^{\text{Lo}} \times \{y_0, y_1\}$ and the inequality from Eq. (A.14).

Therefore, if, in the first case, $p_{a'} > 0$, then counterfactual equalized odds only holds if

$$\beta_{a^*}^{\text{Lo}} = \frac{e - p_{a'} \cdot \pi}{p_{a'} \cdot \rho},$$

since, as noted above, $\rho \neq 0$ by Eq. (A.14). In the second case, if $p_{a'} = 0$, then counterfactual equalized odds can only hold if

$$e = p_{a^*} \cdot \pi = 0.$$

Since we chose a' so that $p_{a^*} > 0$ if $p_{a'} = 0$ and $\pi > 0$ by Eq. (A.16), this is impossible.

In either case, we see that the set of $\beta_{a^*}^{\text{Lo}} \in \mathbb{R}$ such that there a budget-exhausting threshold policy with positive thresholds satisfying counterfactual equalized odds over $\mu + \nu' + \beta_{a^*}^{\text{Lo}} \cdot \nu_{a^*}^{\text{Lo}}$ has λ -measure zero. That is

$$\lambda_{\text{SPAN}(\nu_{a^*}^{\text{Lo}})}[\mathbf{E} - \mu - \nu'] = 0.$$

Since ν' was arbitrary, it follows by Fubini's theorem that $\lambda_{\mathbf{W}}[\mathbf{E} - \mu] = 0$.

Case 2 ($\tau(X) \neq \mathbb{1}_{u(X) > 0}$). Our proof strategy is similar to the previous case. First, we show that, for a given fixed $\nu^{\text{Lo}} \in \mathbf{W}^{\text{Lo}}$, there is a unique candidate policy $\tilde{\tau}(x)$ for being a budget-exhausting

multiple threshold policy with non-negative thresholds and satisfying counterfactual equalized odds over $\mu + \nu^{\text{Lo}} + \nu^{\text{Up}}$ for any $\nu^{\text{Up}} \in \mathbf{W}^{\text{Up}}$. Then, we show that the set of ν^{Up} such that $\tilde{\tau}(X)$ satisfies counterfactual equalized odds has $\lambda_{\mathbf{W}^{\text{Up}}}$ measure zero. Finally, we argue that this in turn implies that the set of $\nu \in \mathbf{W}$ such that there exists a Pareto efficient policy satisfying counterfactual equalized odds over $\mu + \nu$ has $\lambda_{\mathbf{W}}$ -measure zero.

We seek to show that $\lambda_{\mathbf{W}^{\text{Up}}}[\mathbf{E} - (\mu + \nu^{\text{Lo}})] = 0$. To begin, we note that since $\nu_{a,i}^{\text{Up}}$ concentrates on $\{y_1\} \times \mathcal{X}$ for all $a \in \mathcal{A}$, it follows that

$$\mathbb{E}_{\mu + \nu^{\text{Lo}}}[d(X) \mid A = a, Y(1) = y_0] = \mathbb{E}_{\mu + \nu^{\text{Lo}} + \nu^{\text{Up}}}[d(X) \mid A = a, Y(1) = y_0]$$

for any $\nu^{\text{Up}} \in \mathbf{W}^{\text{Up}}$.

Now, suppose there exists some $\nu^{\text{Up}} \in \mathbf{W}^{\text{Up}}$ such that there exists a budget-exhausting multiple threshold policy $\tilde{\tau}(x)$ with non-negative thresholds such that counterfactual equalized odds is satisfied over $\mu + \nu^{\text{Lo}} + \nu^{\text{Up}}$. (If not, then we are done and $\lambda_{\mathbf{W}^{\text{Up}}}[\mathbf{E} - (\mu + \nu^{\text{Lo}})] = 0$, as the measure of the empty set is zero.) Let

$$p = \mathbb{E}_{\mu + \nu^{\text{Lo}}}[\tilde{\tau}(X) \mid A = a, Y(1) = y_0].$$

Suppose that $\tilde{\tau}'(x)$ is an alternative budget-exhausting multiple threshold policy with non-negative thresholds such that counterfactual equalized odds is satisfied. We seek to show that $\tau'(X) = \tau(X)$ $(\mu + \nu^{\text{Lo}} + \nu^{\text{Up}})$ -a.e. for any $\nu^{\text{Up}} \in \mathbf{W}^{\text{Up}}$. Toward a contradiction, suppose that for some $a_0 \in \mathcal{A}$,

$$\mathbb{E}_{\mu + \nu^{\text{Lo}}}[\tilde{\tau}'(X) \mid A = a_0, Y(1) = y_0] < p.$$

Since, by Eq. (A.16), $\Pr_{\mu + \nu^{\text{Lo}}}(A = a_0, Y(1) = y_0) > 0$, it follows that

$$\mathbb{E}_{\mu + \nu^{\text{Lo}}}[\tilde{\tau}'(X) \mid A = a_0] < \mathbb{E}_{\mu + \nu^{\text{Lo}}}[\tilde{\tau}(X) \mid A = a_0].$$

Therefore, since $\tilde{\tau}(x)'$ is budget exhausting, there must be some a_1 such that

$$\mathbb{E}_{\mu + \nu^{\text{Lo}}}[\tilde{\tau}'(X) \mid A = a_1] > \mathbb{E}_{\mu + \nu^{\text{Lo}}}[\tilde{\tau}(X) \mid A = a_1].$$

From this, it follows $\tilde{\tau}'(x)$ can be represented by a threshold greater than or equal to that of $\tilde{\tau}(x)$

on $\alpha^{-1}(a_1)$, and hence

$$\begin{aligned}\mathbb{E}_{\mu+\nu^{\text{Lo}}}[\tilde{\tau}'(X) \mid A = a_1, Y(1) = y_0] &\geq \mathbb{E}_{\mu+\nu^{\text{Lo}}}[\tilde{\tau}(X) \mid A = a_0, Y(1) = y_0] \\ &= p \\ &> \mathbb{E}_{\mu+\nu^{\text{Lo}}}[\tilde{\tau}'(X) \mid A = a_0, Y(1) = y_0],\end{aligned}$$

contradicting the fact that $\tilde{\tau}'(x)$ satisfies counterfactual equalized odds.

Since ν^{Lo} is supported on $u^{-1}((-\infty, 0])$, the preceding discussion, and Lemma A.29, it follows that

$$\tilde{\tau}(X) = \tilde{\tau}'(X) \quad (\mu \upharpoonright_{\mathcal{X} \times \{y_0\}})\text{-a.e.}$$

By Eq. (A.17), it follows that $\tilde{\tau}(X) = \tilde{\tau}'(X)$ $\nu_{a,i}^{\text{Up}}$ -a.e. for $i = 0, 1$. As a consequence,

$$\tilde{\tau}(X) = \tilde{\tau}'(X) \quad (\mu + \nu^{\text{Lo}} + \nu^{\text{Up}})\text{-a.e.}$$

for all $\nu^{\text{Up}} \in \mathbf{W}^{\text{Up}}$. Therefore $\tilde{\tau}(X)$ is, indeed, unique, as desired.

Now, we note that since $\tau(X) \neq \mathbb{1}_{u(X) > 0}$, it follows that $\mathbb{E}[\tau(X)] < \Pr_{\mu}(u(X) > 0)$. It follows that $\mathbb{E}_{\mu}[\tau(X)] = b$, since $\tau(x)$ is budget exhausting. Therefore, by Eq. (A.12), it follows that for any budget-exhausting policy $\tilde{\tau}(X)$, $\mathbb{E}[\tilde{\tau}(X)] = b$, and so $\tilde{\tau}(X) \neq \mathbb{1}_{u(X) > 0}$ over $\mu + \nu$.

Therefore, fix ν^{Lo} and $\tilde{\tau}(X)$. By Eq. (A.18), there is some a^* such that

$$0 < \Pr_{\mu+\nu^{\text{Lo}}}(u(X) > \tilde{t}_{a^*} \mid A = a^*) < 1.$$

Then, it follows by Eq. (A.15) that

$$\int_{\mathcal{K}} \mathbb{1}_{u(X) > \tilde{t}_{a^*}} d\nu_{a^*}^{\text{Up}} \neq 0.$$

Fix $\nu' = \nu - \beta_{a^*}^{\text{Up}} \cdot \nu_{a^*}^{\text{Up}}$. Then, for some $a \neq a^*$, set

$$p^* = \mathbb{E}_{\mu+\nu'}[\tilde{\tau}(X) \mid A = a, Y(1) = y_1].$$

Since the ν_a^{Lo} , ν_a^{Up} are all mutually singular, it follows that counterfactual equalized odds can only hold over $\mu + \nu$ if

$$p^* = \Pr_{\mu+\nu}(u(X) > \tilde{t}_{a^*} \mid A = a^*, Y(1) = y_1).$$

Now, we observe that by Lemma A.21, that

$$\begin{aligned} \Pr_{\mu+\nu}(u(X) > \tilde{t}_{a^*} \mid A = a^*, Y(1) = y_1) \\ = \frac{\eta + \beta_a^{\text{Up}} \cdot \gamma}{\pi} \end{aligned} \tag{A.19}$$

where

$$\begin{aligned} \eta &= \Pr_{\mu+\nu^{\text{Lo}}}(u(X) > \tilde{t}_{a^*} \mid A = a^*, Y(1) = y_1), \\ \pi &= \Pr_{\mu+\nu^{\text{Lo}}}(A = a^*, Y(1) = y_1), \\ \gamma &= \int_{\mathcal{K}} \mathbb{1}_{u(X) > \tilde{t}_{a^*}, A=a^*, Y(1)=y_1} d\nu_a^{\text{Up}}, \end{aligned}$$

and we note that

$$0 = \int_{\mathcal{K}} \mathbb{1}_{A=a^*, Y(1)=y_1} d\nu_a^{\text{Lo}}.$$

Eq. (A.19) can be rearranged to

$$(p^* \cdot \pi - \eta) - \beta \cdot \gamma = 0.$$

This can only hold if

$$\beta = \frac{p^* \cdot \pi - \eta}{\gamma},$$

since by Eq. (A.15), $\gamma \neq 0$. Since any countable subset of $\mathbb{R}$ is a λ -null set,

$$\lambda_{\text{SPAN}(\nu_{a^*}^{\text{Up}})}[\mathbf{E} - \mu - \nu'] = 0.$$

Since ν' was arbitrary, it follows by Fubini's theorem that $\lambda_{\mathbf{W}^{\text{Up}}}[\mathbf{E} - \mu - \nu^{\text{Lo}}] = 0$ in this case as well.

Lastly, since ν^{Lo} was also arbitrary, applying Fubini's theorem a final time gives that $\lambda_{\mathbf{W}}[\mathbf{E} - \mu] = 0$.

Conditional principal fairness and path-specific fairness. The extension of these results to conditional principal fairness and path-specific fairness is straightforward. All that is required is a minor modification of the probe.

In the case of conditional principal fairness, we set

$$\begin{aligned}
\nu_{a,w}^{\text{Up}}[E] &= \mu_{\max,a,w} \circ (\gamma_{(y_1,y_1)} \circ \pi_{\mathcal{X}})^{-1}[E \cap u^{-1}(S_{a,1}^{\text{Up}})], \\
&\quad - \mu_{\max,a} \circ (\gamma_{(y_1,y_1)} \circ \pi_{\mathcal{X}})^{-1}[E \cap u^{-1}(S_{a,w}^{\text{Up}})], \\
\nu_{a,w}^{\text{Lo}}[E] &= \mu_{\max,a,w} \circ (\gamma_{(y_1,y_1)} \circ \pi_{\mathcal{X}})^{-1}[E \cap u^{-1}(S_a^{\text{Lo}})] \\
&\quad - \mu_{\max,a} \circ (\gamma_{(y_0,y_0)} \circ \pi_{\mathcal{X}})^{-1}[E \cap u^{-1}(S_{a,w}^{\text{Lo}})],
\end{aligned}$$

where $\gamma_{(y,y')} : \mathcal{X} \rightarrow \mathcal{K}$ is the injection $x \mapsto (x, y, y')$. Our probe is then given by

$$\begin{aligned}
\mathbf{W}^{\text{Up}} &= \text{SPAN}(\nu_{a,w}^{\text{Up}}), \\
\mathbf{W}^{\text{Lo}} &= \text{SPAN}(\nu_{a,w}^{\text{Lo}}),
\end{aligned}$$

almost as before.

The proof otherwise proceeds virtually identically, except for two points. First, recalling Remark A.41, we use the fact that a generic element of $\mathbf{Q}$ satisfies $\Pr_{\mu}(A = a, W = w) > 0$ in place of $\Pr_{\mu}(A = a) > 0$ throughout. Second, we use the fact that ω overlaps utility in place of Eq. (A.18). In particular, If ω does not overlap utilities for a generic $\mu \in \mathbf{Q}$, then, by Lemma A.47, there exists $w \in \mathcal{W}$ such that $\Pr_{\mu}(u(X) > 0, W = w) = 0$ for all $\mu \in \mathbf{Q}$. If this occurs, we can show that no budget-exhausting multiple threshold policy with positive thresholds satisfies conditional principal fairness, exactly as we did to show Eq. (A.18).

In the case of path-specific fairness, we instead define

$$\begin{aligned}
S_{a,w}^{\text{Lo}} &= S_{a,w} \cap (-\infty, r_{a,w}), \\
S_{a,w}^{\text{Up}} &= S_{a,w} \cap [r_{a,w}, \infty),
\end{aligned}$$

where $r_{a,w}$ is chosen so that

$$\Pr_{\mu_{\max,a,w}}(u(X) \in S_{a,w}^{\text{Lo}}) = \Pr_{\mu_{\max,a,w}}(u(X) \in S_{a,w}^{\text{Up}}).$$

Let $\pi_{\mathcal{X}}$ denote the projection from $\mathcal{K} = \mathcal{A} \times \mathcal{X}^{\mathcal{A}}$ given by

$$(a, (x_{a'})_{a' \in \mathcal{A}}) \mapsto x_a.$$

Let $\pi_{a'}$ denote the projection from the a' -th component. (That is, given $\mu \in \mathbf{K}$, the distribution

of $X_{\Pi,A,a'}$ over μ is given by $\mu \circ \pi_{a'}^{-1}$ and the distribution of X is given by $\mu \circ \pi_X^{-1}$.) Then, we let $\tilde{\mu}_{\max,a,w}$ be the measure on $\mathcal{X}$ given by

$$\begin{aligned} \tilde{\mu}_{\max,a,w}[E] &= \mu_{\max,a,w}[E \cap (u \circ \pi_a)^{-1}(s_{a,w}^{\text{Up}})] \\ &\quad - \mu_{\max,a,w}[E \cap (u \circ \pi_a)^{-1}(S_{a,w}^{\text{Lo}})]. \end{aligned}$$

Finally, let $\phi : \mathcal{A} \rightarrow \mathcal{A}$ be a permutation of the groups with no fixed points, i.e., so that $a' \neq \phi(a')$ for all $a' \in \mathcal{A}$. Then, we define

$$\nu_{a'} = \delta_{a'} \times \tilde{\mu}_{\max,\phi(a'),w_1} \times \prod_{a \neq \phi(a')} \mu_{\max,a,w_1} \circ \pi_a^{-1},$$

where δ_a is the measure on $\mathcal{A}$ given by $\delta_a[\{a'\}] = \mathbb{1}_{a=a'}$. Then, simply let

$$\mathbf{W} = \text{SPAN}(\nu'_a)_{a' \in \mathcal{A}}.$$

Since $\tilde{\mu}_{\max,a,w}[\mathcal{X}] = 0$ for all $a \in \mathcal{A}$, it follows that $\nu_{a,w} \circ \pi_X^{-1} = 0$, i.e.,

$$\Pr_{\mu}(X \in E) = \Pr_{\mu+\nu}(X \in E)$$

for any $\nu \in \mathbf{W}$ and $\mu \in \mathbf{Q}$. Therefore Eqs. (A.11) and (A.12) hold. Moreover, the ν_a satisfy the following strengthening of Eq. (A.15). Perturbations in $\mathbf{W}$ have the property that for any non-trivial t —not necessarily positive—some of the mass of $u(X_{\Pi,A,a})$ is moved either above or below t . More precisely, for any $\mu \in \mathbf{Q}$ and any t such that

$$0 < \Pr_{\mu}(u(X) > t \mid A = a) < 1,$$

if $\nu \in \mathbf{W}$ is such that $\Pr_{|\nu|}(A = \phi^{-1}(a)) > 0$, then

$$\int_{\mathcal{K}} \mathbb{1}_{u(X_{\Pi,A,a}) > t} d\nu_a \neq 0. \quad (\text{A.20})$$

This stronger property means that we need not separately treat the case where $\tau(X) = \mathbb{1}_{u(X) > 0}$ μ -a.e.

Other than this difference the proof proceeds in the same way, except for two points. First, we again make use of the fact that ω can be assumed to overlap utilities in place of Eq. (A.18), as in the case of conditional principal fairness. Second, w_0 and w_1 take the place of y_0 and y_1 . In particular,

to establish the uniqueness of $\tilde{\tau}(x)$ given μ and ν^{Lo} in the second case, instead of conditioning on y_0 , we instead condition on w_0 , where, following the discussion in Remark A.41 and Lemma A.43, this conditioning is well-defined for a generic element of $\mathbf{Q}$. $\square$

We have focused on causal definitions of fairness, but the thrust of our analysis applies to non-causal conceptions of fairness as well. Below we show that policies constrained to satisfy (non-counterfactual) equalized odds (Hardt et al., 2016) are generically strongly Pareto dominated, a result that follows immediately from our proof above.

Definition A.49. *Equalized odds* holds for a decision policy $d(x)$ when

$$d(X) \perp\!\!\!\perp A \mid Y. \quad (\text{A.21})$$

We note that Y in Eq. (A.21) does *not* depend on our choice of $d(x)$, but rather represents the realized value of Y , e.g., under some *status quo* decision making rule.

Corollary A.50. *Suppose $\mathcal{U}$ is a set of utilities consistent modulo α . Further suppose that for all $a \in \mathcal{A}$ there exist a $\mathcal{U}$ -fine distribution of X and a utility $u \in \mathcal{U}$ such that $\Pr(u(X) > 0, A = a) > 0$, where $A = \alpha(X)$. Then, for almost every $\mathcal{U}$ -fine distribution of X and Y on $\mathcal{X} \times \mathcal{Y}$, any decision policy $d(x)$ satisfying equalized odds is strongly Pareto dominated.*

Proof. Consider the following maps. Distributions of X and $Y(1)$, i.e., probability measures on $\mathcal{X} \times \mathcal{Y}$, can be embedded in the space of joint distributions on X , $Y(0)$, and $Y(1)$ via pushing forward by the map ι , where $\iota : (x, y) \mapsto (x, y, y)$. Likewise, given a fixed decision policy $D = d(X)$, joint distributions of X , $Y(0)$, and $Y(1)$ can be projected onto the space of joint distributions of X and Y by pushing forward by the map $\pi_d : (x, y_0, y_1) \mapsto (x, y_{d(x)})$. Lastly, we see that the composition of ι and π_d —regardless of our choice of $d(x)$ —is the identity, as shown in the diagram below.

$$\begin{array}{ccc} \mathcal{X} \times \mathcal{Y} & \xrightarrow{\iota} & \mathcal{X} \times \mathcal{Y} \times \mathcal{Y} \\ & \searrow \text{id} & \downarrow \pi_d \\ & & \mathcal{X} \times \mathcal{Y} \end{array}$$

We note also that counterfactual equalized odds holds for μ exactly when equalized odds holds for $\mu \circ (\pi_d \circ \iota)^{-1}$. The result follows immediately from this and Theorem 2.17. $\square$

A.6.6 Proof of Theorem 2.9

The proof of Theorem 2.9 is simpler than the proof of Theorem 2.17, but uses some of the same machinery. As before, let $\mathcal{K} = \mathcal{A} \times [0, 1]$ denote the state space, and $\mathbb{K}$ denote the set of measures on $[0, 1] \times \mathcal{A}$. Let $\mathbf{K}$ denote the measures $\mu \in \mathbb{K}$ that are absolutely continuous with respect to $\lambda \times \delta$, where λ is Lebesgue measure and δ is the counting measure on $\mathcal{A}$ —i.e., measures such that the restriction to $[0, 1] \times \{a\}$ has a density for all $a \in \mathcal{A}$. Applying Cor. A.27 with $\mathcal{U} = \{\pi\}$, where $\pi : (u, a) \mapsto u$, shows that $\mathbf{K}$ is a Banach space. As before, we let $\mathbf{Q} \subseteq \mathbf{K}$ denote the probability simplex, i.e., the set of all $\mu \in \mathbf{K}$ such that $\mu[\mathcal{K}] = 1$ and $\mu[E] \geq 0$ for all Borel sets E .

Let $R = r(X)$ denote risk.⁸ By Lemma A.11, we have that a policy is utility maximizing if and only if $\mathbb{1}_{R>t} \leq d(X) \leq \mathbb{1}_{R \geq t}$ a.s. By Since $\Pr_\mu(R = t) = 0$ for any absolutely continuous measure μ , we see that, in fact, a policy is utility maximizing if and only if $\mathbb{1}_{R>t} = d(X)$ μ -a.s.

Consider the sets

$$\mathbf{E}_{\text{FP}} = \left\{ \mu \in \mathcal{K} : (\forall a \in \mathcal{A}) \frac{\mathbb{E}_\mu[(1-R) \cdot \mathbb{1}_{A=a, R>t}]}{\mathbb{E}_\mu[(1-R) \cdot \mathbb{1}_{A=a}]} = \frac{\mathbb{E}_\mu[(1-R) \cdot \mathbb{1}_{R>t}]}{\mathbb{E}_\mu[1-R]} \right\},$$

and

$$\mathbf{E}_{\text{DP}} = \{\mu \in \mathcal{K} : (\forall a \in \mathcal{A}) \Pr_\mu(R > t \mid A = a) = \Pr_\mu(R > t)\}.$$

We note that since the mapping $\mu \mapsto \Pr_\mu(E)$ is a continuous function on $\mathcal{K}$, the set of measures such that $\Pr_\mu(A = a, R > t) = 0$ or $\Pr_\mu(R > t) = 0$ is closed and shy. It follows that $\mathbf{E}_{\text{FP}}$ and $\mathbf{E}_{\text{DP}}$ are Borel. We note that, by definition, $\mathbf{E}_{\text{DP}}$ is the set of risk distributions such that there exists a utility-maximizing policy satisfying demographic parity (when the condition is well defined). Likewise, an application of the law of iterated expectations yields that $\mathbf{E}_{\text{FP}}$ is the set of distributions satisfying equalized false positive rates (when the condition is well-defined).

With this context in place, we can move on to the proof of Theorem 2.9.

Proof of Theorem 2.9. First we construct the probe. Choose distinct $a_0, a_1 \in \mathcal{A}$ arbitrarily, and let $\tilde{\nu}$ be defined by

$$\tilde{\nu}[E] = \frac{\lambda[E_{a_0} \cap [0, t]]}{t} - \frac{\lambda[E_{a_1} \cap [0, t]]}{t},$$

where $E_a = E \cap \{A = a\}$. Then, by Lemma A.42, there exists μ_0 such that $\lambda_{\mathbf{W}}[\mathcal{K} + \nu_0] > 0$, where

⁸Here, since the measures are on the risk scale, rather than on X , we write R for notational simplicity.

$$\mathbf{K} = \text{SPAN}(\tilde{\nu}).$$

We see that

$$\begin{aligned} \Pr_{\tilde{\nu}}(R < t, A = a_0) &= 1, & \Pr_{\tilde{\nu}}(R > t, A = a_0) &= 0, \\ \Pr_{\tilde{\nu}}(R < t, A = a_1) &= -1, & \Pr_{\tilde{\nu}}(R > t, A = a_1) &= 0. \end{aligned}$$

It follows that

$$\Pr_{\mu+\beta \cdot \tilde{\nu}}(R > t \mid A = a_0) = \frac{e_0}{p_0 + \beta}, \quad \Pr_{\mu+\beta \cdot \tilde{\nu}}(R > t, \mid A = a_1) = \frac{e_1}{p_1 - \beta},$$

where

$$\begin{aligned} e_0 &= \Pr_{\mu}(R > t, A = a_0), & e_1 &= \Pr_{\mu}(R > t, A = a_1), \\ p_0 &= \Pr_{\mu}(A = a_0), & p_1 &= \Pr_{\mu}(A = a_1). \end{aligned}$$

Note that by Lemmata A.45 and A.48, we can assume that $e_0, e_1 > 0$. It follows, rearranging terms, that

$$\Pr_{\mu+\beta \cdot \tilde{\nu}}(R > t \mid A = a_0) = \Pr_{\mu+\beta \cdot \tilde{\nu}}(R > t \mid A = a_1)$$

if and only if

$$\beta = \frac{e_0 \cdot p_1 - e_1 \cdot p_0}{e_0 + e_1},$$

which is a measure-zero subset of $\beta \in \mathbb{R}$. Therefore $\lambda_{\mathbf{W}}[\mathbf{E}_{\text{DP}} + \mu] = 0$.

In the same way, we observe that

$$\begin{aligned} \mathbb{E}_{\tilde{\nu}}[R \cdot \mathbb{1}_{A=a_0, R < t}] &= \frac{t}{2}, & \mathbb{E}_{\tilde{\nu}}[R \cdot \mathbb{1}_{A=a_0, R > t}] &= 0 \\ \mathbb{E}_{\tilde{\nu}}[R \cdot \mathbb{1}_{A=a_1, R < t}] &= -\frac{t}{2}, & \mathbb{E}_{\tilde{\nu}}[R \cdot \mathbb{1}_{A=a_1, R > t}] &= 0, \end{aligned}$$

and so

$$\frac{\mathbb{E}_{\mu+\beta \cdot \tilde{\nu}}[(1-R) \cdot \mathbb{1}_{A=a_0, R > t}]}{\mathbb{E}_{\mu+\beta \cdot \tilde{\nu}}[(1-R) \cdot \mathbb{1}_{A=a_0}]} = \frac{e'_0}{p'_0 + \beta \cdot \frac{t}{2}}, \quad \frac{\mathbb{E}_{\mu+\beta \cdot \tilde{\nu}}[(1-R) \cdot \mathbb{1}_{A=a_1, R > t}]}{\mathbb{E}_{\mu+\beta \cdot \tilde{\nu}}[(1-R) \cdot \mathbb{1}_{A=a_1}]} = \frac{e'_1}{p'_1 - \beta \cdot \frac{t}{2}},$$

where

$$\begin{aligned} e'_0 &= \mathbb{E}_{\mu}[(1-R) \cdot \mathbb{1}_{A=a_0, R > t}], & e'_1 &= \mathbb{E}_{\mu}[(1-R) \cdot \mathbb{1}_{A=a_1, R > t}], \\ p'_0 &= \mathbb{E}_{\mu}[(1-R) \cdot \mathbb{1}_{A=a_0}], & p'_1 &= \mathbb{E}_{\mu}[(1-R) \cdot \mathbb{1}_{A=a_1}]. \end{aligned}$$

As before, $\mu + \beta \cdot \tilde{\nu} \in \mathbf{E}_{\text{FP}}$ if and only if

$$\beta = \frac{2}{t} \cdot \frac{e_0 \cdot p_1 - e_1 \cdot p_0}{e_0 + e_1},$$

which is again a measure-zero subset of $\beta \in \mathbb{R}$. Therefore $\lambda_{\mathbf{W}}[\mathbf{E}_{\text{FP}} + \mu] = 0$.

Therefore it follows that both $\mathbf{E}_{\text{FP}}$ and $\mathbf{E}_{\text{DP}}$ are shy. □

A.6.7 Proof of Corollary 2.18

The proof of Corollary 2.18—restated in full detail as Corollary A.51 below—is a straightforward application of Theorem 2.17. We begin by giving the complete statement.

Corollary A.51. *Consider a utility of the form given in Eq. (2.12), where v is monotonically increasing in both coordinates and $m(x) \geq 0$. Suppose that for all $a \in \mathcal{A}$ there exists an $\{m\}$ -fine distribution of X such that $\Pr(m(X) > 0, A = a) > 0$, where $A = \alpha(X)$. Then,*

- *For almost every $\{m\}$ -fine distribution of X and $Y(1)$, no utility-maximizing decision policy satisfies counterfactual equalized odds.*
- *If $|\text{IMG}(\omega)| < \infty$ and there exists an $\{m\}$ -fine distribution of X such that $\Pr(A = a, W = w) > 0$ for all $a \in \mathcal{A}$ and $w \in \text{IMG}(\omega)$, where $W = \omega(X)$, then, for almost every $\{m\}$ -fine joint distribution of X , $Y(0)$, and $Y(1)$, no utility-maximizing decision policy satisfies conditional principal fairness.*
- *If $|\text{IMG}(\omega)| < \infty$ and there exists a $\{m\}$ -fine distribution of X such that $\Pr(A = a, W = w_i) > 0$ for all $a \in \mathcal{A}$ and some distinct $w_0, w_1 \in \text{IMG}(\omega)$, then, for almost every $\{m\}^{\mathcal{A}}$ -fine joint distributions of A and the counterfactuals $X_{\Pi, A, a'}$, no utility-maximizing decision policy satisfies path-specific fairness.*

Recall that for notational simplicity, we refer to the distinguished utility as u^* , where

$$u^*(d) = v \left(\mathbb{E}[m(X) \cdot d(X)], \mathbb{E}[\mathbb{1}_{\alpha(X)=a_1} \cdot d(X)] \right).$$

Proof of Corollary 2.18. Consider the subset S of $\mathbb{R}^2$ consisting of all pairs

$$(\mathbb{E}[m(X) \cdot d(X)], \mathbb{E}[\mathbb{1}_{\alpha(X)=a_1} \cdot d(X)]),$$

where d ranges over feasible policies. We note that, for $\theta \in [0, 1]$,

$$\theta \cdot \mathbb{E}[m(X) \cdot d_0(X)] + [1 - \theta] \cdot \mathbb{E}[m(X) \cdot d_1(X)] = \mathbb{E}[m(X) \cdot (\theta \cdot d_0(X) + [1 - \theta] \cdot d_1(X))],$$

and similarly for $\mathbb{E}[\mathbb{1}_{\alpha(X)=a_1} \cdot d(X)]$. Likewise,

$$\theta \cdot \mathbb{E}[d_0(X)] + [1 - \theta] \cdot \mathbb{E}[d_1(X)] = \mathbb{E}[\theta \cdot d_0(X) + (1 - \theta) \cdot d_1(X)],$$

and so convex combinations of feasible policies are feasible. It follows that S is convex.

Now, consider a point (x_0, y_0) at which $v(x, y)$ is maximized on S . Since v is monotonically increasing in both coordinates, we must have that

$$S \cap ((x_0, y_0) + \mathbb{R}_{\geq 0}^2) = \{(x_0, y_0)\},$$

and so, by the separating hyperplane theorem, there exists $(h_0, h_1) \in \mathbb{R}^2$ and $t \in \mathbb{R}$ such that

$$\begin{aligned} (h_0, h_1)^\top (x_0, y_0) &= t, \\ (h_0, h_1)^\top \mathbf{c} &> t, \quad (\mathbf{c} \in (x_0, y_0) + \mathbb{R}_{\geq 0}^2, \mathbf{c} \neq (x_0, y_0)) \\ (h_0, h_1)^\top \mathbf{s} &< t. \quad (\mathbf{s} \in S, \mathbf{s} \neq (x_0, y_0)) \end{aligned}$$

We note that it follows that both h_0 and h_1 are positive, since, otherwise, without loss of generality, if $h_0 = 0$, then

$$(h_0, h_1)^\top (x_0 + \epsilon, y_0) = h_1 \cdot y_0 = (h_0, h_1)^\top (x_0, y_0) = 0,$$

contrary to assumption, since

$$(x_0 + \epsilon, y_0) \in (x_0, y_0) + \mathbb{R}_{\geq 0}^2, (x_0 + \epsilon, y_0) \neq (x_0, y_0).$$

Let $\lambda_* = h_1/h_0$, and consider the collection of utilities

$$\mathcal{U} = \{m(x) + \lambda \cdot \mathbb{1}_{\alpha(x)=a_1}\}_{\lambda > 0}.$$

Since $m(x) \geq 0$, $\mathcal{U}$ is consistent modulo α .

We need the further claim that any policy that is utility maximizing for some $u(x) \in \mathcal{U}$ is Pareto efficient. For, suppose that $d_0(x)$ were utility-maximizing for $u_0(x)$ but not Pareto efficient, i.e., there existed $d_1(x)$ and $u_1(x)$ such that $u_0(d_1) = u_0(d_0)$ but $u_1(d_1) > u_1(d_0)$. Then, we would have

for any $u \in \mathcal{U}$ that

$$\begin{aligned} u(d_1) - u(d_0) &= u(d_1) - u(d_0) - (u_0(d_1) - u_0(d_0)) \\ &= (\lambda - \lambda_0) \cdot (\mathbb{E}[d_1(X) \cdot \mathbb{1}_{\alpha(X)=a_1}] - \mathbb{E}[d_0(X) \cdot \mathbb{1}_{\alpha(X)=a_1}]). \end{aligned}$$

First suppose that $\lambda_1 > \lambda_0$. Then, it follows from the fact that $u_1(d_1) > u_1(d_0)$ that

$$\mathbb{E}[d_1(X) \cdot \mathbb{1}_{\alpha(X)=a_1}] > \mathbb{E}[d_0(X) \cdot \mathbb{1}_{\alpha(X)=a_1}].$$

Now, if $0 < \lambda_2 < \lambda_1$ and $u(x) = m(x) + \lambda_2 \cdot \mathbb{1}_{\alpha(x)=a_1}$, then we have that $u_2(d_1) < u_2(d_0)$. In the same way, if $\lambda_1 < \lambda_0$, we could choose $\lambda_2 > \lambda_1$ such that $u_2(d_1) < u_2(d_0)$. Therefore d_1 does not Pareto dominate d_0 , contrary to hypothesis. Therefore, any policy that is utility maximizing for some $u(x) \in \mathcal{U}$ is Pareto efficient.

In particular, it follows that the policy that maximizes $u(x) = m(x) + \lambda_* \cdot \mathbb{1}_{\alpha(x)=a_1}$ in expectation is Pareto efficient. By the construction of the separating hyperplane, this is also the policy that maximizes u^* , and so the policy that maximizes u^* is Pareto efficient. Therefore, under the hypotheses of Theorem 2.17, for almost every joint distribution, the utility maximizing policy does not satisfy counterfactual equalized odds, conditional principal fairness, or path-specific fairness. $\square$

A.6.8 General Measures on $\mathcal{K}$

Theorem 2.17 is restricted to $\mathcal{U}$ -fine and $\mathcal{U}^A$ -fine distributions on the state space. The reason for this restriction is that when the distribution of X induces atoms on the utility scale, threshold policies can possess additional—or even infinite—degrees of freedom when the threshold falls exactly on an atom. In particular circumstances, these degrees of freedom can be used to ensure causal fairness notions, such as counterfactual equalized odds, hold in a locally robust way. In particular, the generalization of Theorem 2.17 beyond $\mathcal{U}$ -fine measures to all totally bounded measures on the state space is false, as illustrated by the following proposition.

Proposition A.52. *Consider the set $\mathbf{E}' \subseteq \mathbb{K}$ of distributions—not necessarily $\mathcal{U}$ -fine—on $\mathcal{K} = \mathcal{X} \times \mathcal{Y}$ over which there exists a Pareto efficient policy satisfying counterfactual equalized odds. There exist b , $\mathcal{X}$, $\mathcal{Y}$, and $\mathcal{U}$ such that $\mathbf{E}'$ is not relatively shy.*

Proof. We adopt the notational conventions of Section A.6.3. We note that by Prop. A.16, a set

can only be shy if it has empty interior. Therefore, we will construct an example in which an open ball of distributions on $\mathcal{K}$ in the total variation norm all allow for a Pareto efficient policy satisfying counterfactual equalized odds, i.e., are contained in $\mathbf{E}'$.

Let $b = \frac{3}{4}$, $\mathcal{Y} = \{0, 1\}$, $\mathcal{A} = \{a_0, a_1\}$, and $\mathcal{X} = \{0, 1\} \times \{a_0, a_1\} \times \mathbb{R}$. Let $\alpha : \mathcal{X} \rightarrow \mathcal{A}$ be given by $\alpha : (y, a, v) \mapsto a$ for arbitrary $(y, a, v) \in \mathcal{X}$. Likewise, let $u : \mathcal{X} \rightarrow \mathbb{R}$ be given by $u : (y, a, v) \mapsto v$. Then, if $\mathcal{U} = \{u\}$, $\mathcal{U}$ is vacuously consistent modulo α . Consider the joint distribution μ on $\mathcal{K} = \mathcal{X} \times \mathcal{Y}$ where for all $y, y' \in \mathcal{Y}$, $a \in \mathcal{A}$, and $u \in \mathbb{R}$,

$$\Pr_\mu(X = (a, y, u), Y(1) = y') = \frac{1}{4} \cdot \mathbb{1}_{y=y'} \cdot \Pr_\mu(u(X) = u),$$

where, over μ , $u(X)$ is distributed as a $\frac{1}{2}$ -mixture of $\text{UNIF}(1, 2)$ and $\delta(1)$; that is, $\Pr(u(X) = 1) = \frac{1}{2}$ and $\Pr(a < u(X) < b) = \frac{b-a}{2}$ for $0 \leq a \leq b < 1$.

We first observe that there exists a Pareto efficient threshold policy $\tau(x)$ such that counterfactual equalized odds is satisfied with respect to the decision policy $\tau(X)$. Namely, let

$$\tau(a, y, u) = \begin{cases} 1 & u > 1, \\ \frac{1}{2} & u = 1, \\ 0 & u < 1. \end{cases}$$

Then, it immediately follows that $\mathbb{E}[\tau(X)] = \frac{3}{4} = b$. Since $\tau(x)$ is a threshold policy and exhausts the budget, it is utility maximizing by Lemma A.11. Moreover, if $D = \mathbb{1}_{U_D \leq \tau(X)}$ for some $U_D \sim \text{UNIF}(0, 1)$ independent of X and $Y(1)$, then $D \perp\!\!\!\perp A \mid Y(1)$. Since $u(X) \perp\!\!\!\perp A, Y(1)$, it follows that

$$\begin{aligned} \Pr_\mu(D = 1 \mid A = a, Y(1) = y) &= \Pr(U_D \leq \tau(X) \mid A = a, Y(1) = y) \\ &= \Pr(U_D \leq \tau(X)) \\ &= \mathbb{E}_\mu[\tau(X)], \end{aligned}$$

Therefore Eq. (2.4) is satisfied, i.e., counterfactual equalized odds holds. Now, using μ , we construct an open ball of distributions over which we can construct similar threshold policies. In particular, suppose μ' is any distribution such that $|\mu - \mu'|[\mathcal{K}] < \frac{1}{64}$. Then, we claim that there exists a

budget-exhausting threshold policy satisfying counterfactual equalized odds over μ' . For, we note that

$$\begin{aligned}\Pr_{\mu'}(U > 1) &< \Pr_{\mu}(U > 1) + \frac{1}{64} = \frac{33}{64}, \\ \Pr_{\mu'}(U \geq 1) &> \Pr_{\mu}(U \geq 1) - \frac{1}{64} = \frac{63}{64},\end{aligned}$$

and so any threshold policy $\tau'(x)$ satisfying $\mathbb{E}[\tau'(X)] = b = \frac{3}{4}$ must have $t = 1$ as its threshold.

We will now construct a threshold policy $\tau'(x)$ satisfying counterfactual equalized odds over μ' . Consider a threshold policy of the form

$$\tau'(a, y, u) = \begin{cases} 1 & u > 1, \\ p_{a,y} & u = 1, \\ 0 & u < 1. \end{cases}$$

For notational simplicity, let

$$\begin{aligned}q_{a,y} &= \Pr_{\mu'}(A = a, Y = y, U > 1), \\ r_{a,y} &= \Pr_{\mu'}(A = a, Y = y, U = 1), \\ \pi_{a,y} &= \Pr_{\mu'}(A = a, Y = y).\end{aligned}$$

Then, we have that

$$\begin{aligned}\mathbb{E}_{\mu'}[\tau'(X)] &= \sum_{a,y} q_{a,y} + p_{a,y} \cdot r_{a,y}, \\ \mathbb{E}_{\mu'}[\tau'(X) \mid A = a, Y = y] &= \frac{q_{a,y} + p_{a,y} \cdot r_{a,y}}{\pi_{a,y}}.\end{aligned}$$

Therefore, the policy will be budget exhausting if

$$\sum_{a,y} q_{a,y} + p_{a,y} \cdot r_{a,y} = \frac{3}{4},$$

and it will satisfy counterfactual equalized odds if

$$\begin{aligned}
& \pi_{a_1,0} \cdot (q_{a_0,0} + p_{a_0,0} \cdot r_{a_0,0}) \\
& \quad = \pi_{a_0,0} \cdot (q_{a_1,0} + p_{a_1,0} \cdot r_{a_1,0}), \\
& \pi_{a_1,1} \cdot (q_{a_0,1} + p_{a_0,1} \cdot r_{a_0,1}) \\
& \quad = \pi_{a_0,1} \cdot (q_{a_1,1} + p_{a_1,1} \cdot r_{a_1,1}),
\end{aligned} \tag{A.22}$$

since, as above,

$$\Pr(D = 1 \mid A = a, Y(1) = y) = \mathbb{E}[\tau'(X) \mid A = a, Y(1) = y].$$

Again, for notational simplicity, let

$$S = \frac{\frac{3}{4} - \Pr_{\mu'}(U > 1)}{\Pr_{\mu'}(U = 1)}.$$

Then, a straightforward algebraic manipulation shows that Eq. (A.22) is solved by setting $p_{a_0,y}$ to be

$$\frac{S \cdot \pi_{a_0,y} \cdot (r_{a_0,y} + r_{a_1,y}) + \pi_{a_0,y} \cdot q_{a_1,y} - \pi_{a_1,y} \cdot q_{a_0,y}}{r_{a_0,y} \cdot (\pi_{a_0,y} + \pi_{a_1,y})},$$

and $p_{a_1,y}$ to be

$$\frac{S \cdot \pi_{a_1,y} \cdot (r_{a_0,y} + r_{a_1,y}) + \pi_{a_1,y} \cdot q_{a_0,y} - \pi_{a_0,y} \cdot q_{a_1,y}}{r_{a_1,y} \cdot (\pi_{a_0,y} + \pi_{a_1,y})}.$$

In order for $\tau'(x)$ to be a well-defined policy, we need to show that $p_{a,y} \in [0, 1]$ for all $a \in \mathcal{A}$ and $y \in \mathcal{Y}$. To that end, note that

$$\begin{aligned}
q_{a,y} &= \Pr_{\mu'}(A = a, Y = y, U > 1), \\
r_{a,y} &= \Pr_{\mu'}(A = a, Y = y, U = 1), \\
\pi_{a,y} &= \Pr_{\mu'}(A = a, Y = y), \\
r_{a_0,y} + r_{a_1,y} &= \Pr_{\mu'}(Y = y, U = 1), \\
\pi_{a_0,y} + \pi_{a_1,y} &= \Pr_{\mu'}(Y = y), \\
S &= \frac{\frac{3}{4} - \Pr_{\mu'}(U > 1)}{\Pr_{\mu'}(U = 1)}.
\end{aligned}$$

Now, we recall that $|\Pr_{\mu'}(E) - \Pr_{\mu}(E)| < \frac{1}{64}$ for any event E by hypothesis. Therefore,

$$\begin{aligned} \frac{7}{64} &\leq q_{a,y} && \leq \frac{9}{64}, \\ \frac{7}{64} &\leq r_{a,y} && \leq \frac{9}{64}, \\ \frac{7}{64} &\leq \pi_{a,y} && \leq \frac{17}{64}, \\ \frac{15}{64} &\leq r_{a_0,y} + r_{a_1,y} && \leq \frac{17}{64}, \\ \frac{31}{64} &\leq \pi_{a_0,y} + \pi_{a_1,y} && \leq \frac{33}{64}, \\ \frac{15}{31} &\leq S && \leq \frac{17}{33}. \end{aligned}$$

Using these bounds and the expressions for $p_{a,y}$ derived above, we see that

$$\frac{629}{3069} < p_{a,y} < \frac{6497}{7161},$$

and hence $p_{a,y} \in [0, 1]$ for all $a \in \mathcal{A}$ and $y \in \mathcal{Y}$.

Therefore, the policy $\tau'(x)$ is well-defined, and, by construction, is budget-exhausting and therefore utility-maximizing by Lemma A.11. It also satisfies counterfactual equalized odds by construction. Since μ' was arbitrary, it follows that the set of distributions on $\mathcal{K}$ such that there exists a Pareto efficient policy satisfying counterfactual equalized odds contains an open ball, and hence is not shy. $\square$

A.7 Theorem 2.11 and Related Results

We first prove a variant of Theorem 2.11 for general, continuous covariates $\mathcal{X}$. Then, we extend and generalize Theorem 2.11 using the theory of finite Markov chains, offering a proof of the theorem different from the sketch included in the main text.

A.7.1 Extension to Continuous Covariates

Here we follow the proof sketch in the main text for Theorem 2.11, which assumes a finite covariate-space $\mathcal{X}$. In that case, we start with a point x^* with maximum decision probability $d(x^*)$, and then assume, toward a contradiction, that there exists a point with strictly lower decision probability. The general case is more involved since it is not immediately clear that the maximum value of $d(x)$

is achieved with positive probability in $\mathcal{X}$. We start with the lemma below before proving the main result.

Lemma A.53. *A decision policy $d(x)$ satisfies path-specific fairness with $W = X$ if and only if any $a' \in \mathcal{A}$,*

$$\mathbb{E}[d(X_{\Pi,A,a'}) \mid X] = d(X).$$

Proof. First, suppose that $d(x)$ satisfies path-specific fairness. To show the result, we use the standard fact that for independent random variables X and U ,

$$\mathbb{E}[f(X, U) \mid X] = \int f(X, u) dF_U(u), \quad (\text{A.23})$$

where F_U is the distribution of U . (For a proof of this fact see, for example, Brozius, 2019)

Now, we have that

$$\begin{aligned} \mathbb{E}[D_{\Pi,A,a'} \mid X_{\Pi,A,a'}] &= \mathbb{E}[\mathbb{1}_{U_D \leq d(X_{\Pi,A,a'})} \mid X_{\Pi,A,a'}] \\ &= \int_0^1 \mathbb{1}_{u \leq d(X_{\Pi,A,a'})} du \\ &= d(X_{\Pi,A,a'}), \end{aligned}$$

where the first equality follows from the definition of $D_{\Pi,A,a'}$, and the second from Eq. (A.23), since the exogenous variable $U_D \sim \text{UNIF}(0, 1)$ is independent of the counterfactual covariates $X_{\Pi,A,a'}$. An analogous argument shows that $\mathbb{E}[D \mid X] = d(X)$.

Finally, conditioning on X , we have

$$\begin{aligned} \mathbb{E}[d(X_{\Pi,A,a'}) \mid X] &= \mathbb{E}[\mathbb{E}[D_{\Pi,A,a'} \mid X_{\Pi,A,a'}] \mid X] \\ &= \mathbb{E}[\mathbb{E}[D_{\Pi,A,a'} \mid X_{\Pi,A,a'}, X] \mid X] \\ &= \mathbb{E}[D_{\Pi,A,a'} \mid X] \\ &= \mathbb{E}[D \mid X] \\ &= d(X), \end{aligned}$$

where the second equality follows from the fact that $D_{\Pi,A,a'} \perp\!\!\!\perp X \mid X_{\Pi,A,a'}$, the third from the law of iterated expectations, and the fourth from the definition of path-specific fairness.

Next, suppose that

$$\mathbb{E}[d(X_{\Pi,A,a'} \mid X) = d(X)]$$

for all $a' \in \mathcal{A}$. Then, since $W = X$ and $X \perp\!\!\!\perp U_D$, using Eq. (A.23), we have that for all $a' \in \mathcal{A}$,

$$\begin{aligned} \mathbb{E}[D_{\Pi,A,a'} \mid X] &= \mathbb{E}[\mathbb{E}[\mathbb{1}_{U_D \leq d(X_{\Pi,A,a'})} \mid X_{\Pi,A,a'}, X] \mid X] \\ &= \mathbb{E}[\mathbb{E}[d(X_{\Pi,A,a'}) \mid X_{\Pi,A,a'}, X] \mid X] \\ &= \mathbb{E}[d(X_{\Pi,A,a'}) \mid X] \\ &= d(X) \\ &= \mathbb{E}[d(X) \mid X] \\ &= \mathbb{E}[D \mid X]. \end{aligned}$$

This is exactly Eq. (2.8), and so the result follows. $\square$

We are now ready to prove a continuous variant of Theorem 2.11. The technical hypotheses of the theorem ensure that the conditional probability measures $\Pr(E \mid X)$ are “sufficiently” mutually non-singular distributions on $\mathcal{X}$ with respect to the distribution of X —for example, the conditions ensure that the conditional distribution of $X_{\Pi,A,a} \mid X$ does not have atoms that X itself does not have, and *vice versa*. For notational and conceptual simplicity, we only consider the case of trivial ζ , i.e., where $\zeta(x) = \zeta(x')$ for all $x, x' \in \mathcal{X}$.

Proposition A.54. *Suppose that*

1. *For all $a \in \mathcal{A}$ and any event S satisfying $\Pr(X \in S \mid A = a) > 0$, we have, a.s.,*

$$\Pr(X_{\Pi,A,a} \in S \vee A = a \mid X) > 0.$$

2. *For all $a \in \mathcal{A}$ and $\epsilon > 0$, there exists $\delta > 0$ such that for any event S satisfying $\Pr(X \in S \mid A = a) < \delta$, we have, a.s.,*

$$\Pr(X_{\Pi,A,a} \in S, A \neq a \mid X) < \epsilon.$$

Then, for $W = X$, any Π -fair policy $d(x)$ is constant a.s. (i.e., $d(X) = p$ a.s. for some $0 \leq p \leq 1$).

Proof. Let $d_{\max} = \|d(x)\|_{\infty}$, the essential supremum of d . To establish the theorem statement, we show that $\Pr(d(X) = d_{\max} \mid A = a) = 1$ for all $a \in \mathcal{A}$. To do that, we begin by showing that there

exists some $a \in \mathcal{A}$ such that $\Pr(d(X) = d_{\max} \mid A = a) > 0$.

Assume, toward a contradiction, that for all $a \in \mathcal{A}$,

$$\Pr(d(X) = d_{\max} \mid A = a) = 0. \quad (\text{A.24})$$

Because $\mathcal{A}$ is finite, there must be some $a_0 \in \mathcal{A}$ such that

$$\Pr(d_{\max} - d(X) < \epsilon \mid A = a_0) > 0 \quad (\text{A.25})$$

for all $\epsilon > 0$.

Choose $a_1 \neq a_0$. We show that for values of x such that $d(x)$ is close to $d_{\max}$, the distribution of $d(X_{\Pi, A, a_1}) \mid X = x$ must be concentrated near $d_{\max}$ with high probability to satisfy the definition of path-specific fairness, in Eq. (2.8). But, under the assumption in Eq. (A.24), we also show that the concentration occurs with low probability, by the continuity hypothesis in the statement of the theorem, establishing the contradiction.

Specifically, by Markov's inequality, for any $\rho > 0$, a.s.,

$$\begin{aligned} \Pr(d_{\max} - d(X_{\Pi, A, a_1}) \geq \rho \mid X) &\leq \frac{\mathbb{E}[d_{\max} - d(X_{\Pi, A, a_1}) \mid X]}{\rho} \\ &= \frac{d_{\max} - d(X)}{\rho}, \end{aligned}$$

where the final equality follows from Lemma A.53. Rearranging, it follows that for any $\rho > 0$, a.s.,

$$\Pr(d_{\max} - d(X_{\Pi, A, a_1}) < \rho \mid X) \geq 1 - \frac{d_{\max} - d(X)}{\rho}. \quad (\text{A.26})$$

Now let $S = \{x \in \mathcal{X} : d_{\max} - d(x) < \rho\}$. By the second hypothesis of the theorem, we can choose δ sufficiently small that if

$$\Pr(X \in S \mid A = a_1) < \delta$$

then, a.s.,

$$\Pr(X_{\Pi, A, a_1} \in S, A \neq a_1 \mid X) < \frac{1}{2}.$$

In other words, we can chose δ such that if

$$\Pr(d_{\max} - d(X) < \rho \mid A = a_1) < \delta$$

then, a.s.,

$$\Pr(d_{\max} - d(X_{\Pi, A, a_1}) < \rho, A \neq a_1 \mid X) < \frac{1}{2}$$

By Eq. (A.24), we can choose $\epsilon > 0$ so small that

$$\Pr(d_{\max} - d(X) < \epsilon \mid A = a_1) < \delta.$$

Then, we have that

$$\Pr(d_{\max} - d(X_{\Pi, A, a_1}) < \epsilon, A \neq a_1 \mid X) < \frac{1}{2}$$

a.s. Further, by the definition of the essential supremum and a_0 , and the fact that $a_0 \neq a_1$, we have that

$$\Pr(d_{\max} - d(X) < \frac{\epsilon}{2}, A \neq a_1) > 0.$$

Therefore, with positive probability, we have that

$$\begin{aligned} 1 - \frac{d_{\max} - d(X)}{\epsilon} &> 1 - \frac{\frac{\epsilon}{2}}{\epsilon} \\ &= \frac{1}{2} \\ &> \Pr(d_{\max} - d(X_{\Pi, A, a_1}) < \epsilon, A \neq a_1 \mid X). \end{aligned}$$

This contradicts Eq. (A.26), and so it cannot be the case that $\Pr(d(X) = d_{\max} \mid A = a_0) = 0$, meaning $\Pr(d(X) = d_{\max} \mid A = a_0) > 0$.

Now, we show that $\Pr(d(X) = d_{\max} \mid A = a_1) = 1$. Suppose, toward a contradiction, that

$$\Pr(d(X) < d_{\max} \mid A = a_1) > 0.$$

Then, by the first hypothesis, a.s.,

$$\Pr(d(X_{\Pi, A, a_1}) < d_{\max} \vee A = a_1 \mid X) > 0$$

As a consequence,

$$\begin{aligned}
d_{\max} &= \mathbb{E}[d(X) \mid d(X) = d_{\max}, A = a_0] \\
&= \mathbb{E}[\mathbb{E}[d(X_{\Pi,A,a_1}) \mid X] \mid d(X) = d_{\max}, A = a_0] \\
&< \mathbb{E}[\mathbb{E}[d_{\max} \mid X] \mid d(X) = d_{\max}, A = a_0] \\
&= \mathbb{E}[d_{\max} \mid d(X) = d_{\max}, A = a_0] \\
&= d_{\max},
\end{aligned}$$

where we can condition on the set $\{d(X) = d_{\max}, A = a_0\}$ since $\Pr(d(X) = d_{\max} \mid A = a_0) > 0$; and the second equality above follows from Lemma A.53. This establishes the contradiction, and so $\Pr(d(X) = d_{\max} \mid A = a_1) = 1$.

Finally, we extend this equality to all $a \in \mathcal{A}$. Since, $\Pr(d(X) \neq d_{\max} \mid A = a_1) = 0$, we have, by the second hypothesis of the theorem, that, a.s.,

$$\Pr(d(X_{\Pi,A,a_1}) \neq d_{\max}, A \neq a_1 \mid X) = 0.$$

Since, by definition, $\Pr(X_{\Pi,A,a_1} = X \mid A = a_1) = 1$, and $\Pr(d(X) = d_{\max} \mid A = a_1) = 1$, we can strengthen this to

$$\Pr(d(X_{\Pi,A,a_1}) \neq d_{\max} \mid X) = 0.$$

Consequently, a.s.,

$$\begin{aligned}
d(X) &= \mathbb{E}[d(X_{\Pi,A,a}) \mid X] \\
&= \mathbb{E}[d_{\max} \mid X] \\
&= d_{\max},
\end{aligned}$$

where the first equality follows from Lemma A.53, establishing the result. $\square$

A.7.2 A Markov Chain Perspective

The theory of Markov chains illuminates—and allows us to extend—the proof of Theorem 2.11. Suppose $\mathcal{X} = \{x_1, \dots, x_n\}$.⁹ For any $a' \in \mathcal{A}$, let $P_{a'} = [p_{i,j}^{a'}]$ where $p_{i,j}^{a'} = \Pr(X_{\Pi,A,a'} = x_j \mid X = x_i)$.

⁹Because of the technical difficulties associated with characterizing the long-run behavior of arbitrary infinite Markov chains, we restrict our attention in this section to Markov chains with finite state spaces.

Then $P_{a'}$ is a stochastic matrix.

To motivate the subsequent discussion, we first note that this perspective conceptually simplifies some of our earlier results. Lemma A.53 can be recast as stating that when $W = X$, a policy d is Π -fair if and only if $P_{a'}d = d$ —i.e., if and only if d is a 1-eigenvector of $P_{a'}$ —for all $a' \in \mathcal{A}$.

The 1-eigenvectors of Markov chains have a particularly simple structure, which we derive here for completeness.

Lemma A.55. *Let $S_1, \dots, S_m$ denote the recurrent classes of a finite Markov chain with transition matrix P . If d is a 1-eigenvector of P , then d takes a constant value p_k on each S_k , $k = 1, \dots, m$, and*

$$d_i = \sum_{k=1}^m \left[\lim_{n \rightarrow \infty} \sum_{j \in S_k} P_{ij}^n \right] \cdot p_k. \quad (\text{A.27})$$

Remark A.56. We note that $\lim_{n \rightarrow \infty} \sum_{j \in S_k} P_{ij}^n$ always exists and is the probability that the Markov chain, beginning at state i , is eventually absorbed by the recurrent class S_k .

Proof. Note that, possibly by reordering the states, we can arrange that the stochastic matrix P is in canonical form, i.e., that

$$P = \begin{bmatrix} B & \\ R' & Q \end{bmatrix},$$

where Q is a sub-stochastic matrix, R is non-negative, and

$$B = \begin{bmatrix} P_1 & & & \\ & P_2 & & \\ & & \ddots & \\ & & & P_m \end{bmatrix}$$

is a block-diagonal matrix with the stochastic matrix P_i corresponding to the transition probabilities on the recurrent set S_i in the i -th position along the diagonal.

Now, consider a 1-eigenvector $v = [v_1 \ v_2]^\top$ of P . We must have that $Pv = v$, i.e., $Bv_1 = v_1$ and $R'v_1 + Qv_2 = v_2$. Therefore v_1 is a 1-eigenvector of B . Since B is block diagonal, and each diagonal element is a positive stochastic matrix, it follows by the Perron-Frobenius theorem that the 1-eigenvectors of B are given by $\text{SPAN}(\mathbf{1}_{S_i})_{i=1, \dots, m}$, where $\mathbf{1}_{S_i}$ is the vector which is 1 at index j if $j \in S_i$ and is 0 otherwise.

Now, for $v_1 \in \text{SPAN}(\mathbf{1}_{S_i})_{i=1,\dots,m}$, we must find v_2 such that $R'v_1 + Qv_2 = v_2$.

Note that every finite Markov chain M can be canonically associated with an absorbing Markov chain M^{ABS} where the set of states of M^{ABS} is exactly the union of the transitive states of M and the recurrent sets of M . (In essence, one tracks which state of M the Markov chain is in until it is absorbed by one of the recurrent sets, at which point the entire recurrent set is treated as a single absorbent state.) The transition matrix P^{ABS} associated with M^{ABS} is given by

$$P^{\text{ABS}} = \begin{bmatrix} I & \\ R & Q \end{bmatrix}$$

where $R = R'[\mathbf{1}_{S_1} \dots \mathbf{1}_{S_m}]$. In particular, it follows that $v = [v_1 \ v_2]^\top$ is a 1-eigenvector of P if and only if $[Tv_1 \ v_2]^\top$ is a 1-eigenvector of P^{ABS} , where $T : \mathbf{1}_{S_i} \mapsto \mathbf{e}_i$.

Now, if v is a 1-eigenvector of P^{ABS} , then it is a 1-eigenvector of $(P^{\text{ABS}})^k$ for all k . Since Q is sub-stochastic, the series $\sum_{k=0}^{\infty} Q^k$ converges to $(I - Q)^{-1}$. Since

$$(P^{\text{ABS}})^k = \begin{bmatrix} I & \\ (I + Q + \dots + Q^{k-1})R & Q^k \end{bmatrix},$$

it follows that

$$\lim_{k \rightarrow \infty} (P^{\text{ABS}})^k = \begin{bmatrix} I & \\ (I - Q)^{-1}R & 0 \end{bmatrix}.$$

Therefore, if $v = [v_1 \ v_2]^\top$ is a 1-eigenvector of P^{ABS} , we must have that $(I - Q)^{-1}Rv_1 = v_2$. By Theorem 3.3.7 in Kemeny and Snell (1976), the (i, k) entry of $(I - Q)^{-1}R$ is exactly the probability that, conditional on $X_0 = x_i$, the Markov chain is eventually absorbed by the recurrent set S_k . This is, in turn, by the Chapman-Kolmogorov equations and the definition of S_k , equal to $\lim_{n \rightarrow \infty} \sum_{j \in S_k} P_{i,j}^n$, and therefore the result follows. $\square$

We arrive at the following simple necessary condition on Π -fair policies.

Corollary A.57. *Suppose $\mathcal{X}$ is finite, and define the stochastic matrix $P = \frac{1}{|\mathcal{A}|} \sum_{a' \in \mathcal{A}} P_{a'}$. If $d(x)$ is a Π -fair policy then it is constant on the recurrent classes of P .*

Proof. By Lemma A.53, d is Π -fair if and only if $P_{a'}d = d$ for all $a' \in \mathcal{A}$. Therefore,

$$\frac{1}{|\mathcal{A}|} \sum_{a' \in \mathcal{A}} P_{a'}d = \frac{1}{|\mathcal{A}|} \sum_{a' \in \mathcal{A}} d = d, \tag{A.28}$$

and so d is a 1-eigenvector of P . Therefore it is constant on the recurrent classes of P by Lemma A.55. $\square$

We note that Theorem 2.11 follows immediately from this.

Proof of Theorem 2.11. Note that $\frac{1}{|\mathcal{A}|} \sum_{a \in \mathcal{A}} P_a$ decomposes into a block diagonal stochastic matrix, where each block corresponds to a single stratum of ζ and is irreducible. Consequently, each stratum forms a recurrent class, and the result follows. $\square$

A.8 Proofs of Propositions 2.19 and 2.20

The proofs of Proposition 2.19 and Proposition 2.20 rely on certain shared theory about beta distributions. We begin by reviewing this theory before moving onto the proofs of the respective propositions.

A.8.1 Beta Distributions and Stochastic Dominance

We begin by introducing incomplete beta functions and distributions.

Definition A.58. The *incomplete beta function* $I_t(\alpha, \beta)$ for $t \in (0, 1]$ is given by

$$I_t(\alpha, \beta) = \int_0^t x^{\alpha-1} (1-x)^{\beta-1} dx.$$

A random variable X is said to be distributed as $\text{BETA}_t(\alpha, \beta)$ if $X \sim Y \mid Y < t$ for $Y \sim \text{BETA}(\alpha, \beta)$; equivalently, if X has PDF

$$\mathbb{1}_{x \in (0, t)} \cdot \frac{x^{\alpha-1} (1-x)^{\beta-1}}{I_t(\alpha, \beta)}.$$

An important property relating different beta distributions is stochastic dominance.

Definition A.59. Let X and Y be random variables with CDFs $F_X(t)$ and $F_Y(t)$, respectively. We say that X *stochastically dominates* Y , written $Y \leq_{\text{st}} X$, if $F_X(t) \leq F_Y(t)$ for all $t \in \mathbb{R}$. We say that X *strictly stochastically dominates* Y , i.e., $Y <_{\text{st}} X$, if $F_X(t) = F_Y(t)$ implies that either $F_X(t) = F_Y(t) = 0$ or $F_X(t) = F_Y(t) = 1$.

Stochastic dominance has the following useful property.

Lemma A.60. *Suppose $Y \leq_{\text{st}} X$. Then, for any monotonically non-decreasing $\varphi : \mathbb{R} \rightarrow \mathbb{R}$, we have that $\mathbb{E}[\varphi(Y)] \leq \mathbb{E}[\varphi(X)]$.*

For proof, see (1.A.7) in Shaked and Shanthikumar (2007). We will need to improve the inequality to a strict inequality. We begin with the simplest case.

Lemma A.61. *Suppose $Y <_{\text{st}} X$, where X and Y are positive and $F_X(t)$, $F_Y(t)$ are continuous. Then, $\mathbb{E}[Y \cdot \mathbb{1}_{Y>t}] < \mathbb{E}[X \cdot \mathbb{1}_{X>t}]$ for any $t > 0$ such that there exists $t' \geq t$ with $F_X(t') < F_Y(t')$.*

Proof. Since $F_Y(x) \geq F_X(x)$ for all x and, in particular, $F_Y(x) > F_X(x)$ on an open interval containing t' , we have that

$$\begin{aligned} \mathbb{E}[Y \cdot \mathbb{1}_{Y>t}] &= t \cdot [1 - F_Y(t)] + \int_t^\infty 1 - F_Y(x) \, dx \\ &< t \cdot [1 - F_X(t)] + \int_t^\infty 1 - F_X(x) \, dx \\ &= \mathbb{E}[X \cdot \mathbb{1}_{X>t}] \end{aligned}$$

where we have applied the Darth Vader rule to calculate the expectations (Muldowney et al., 2012). □

This leads to the following modest technical generalization.

Lemma A.62. *Suppose $Y <_{\text{st}} X$, where X and Y are positive and $F_X(t)$ and $F_Y(t)$ are continuous. Suppose that $t > 0$ is such that there exists $t' < t$ with $F_X(t') < F_Y(t')$, and $f : \mathbb{R}_{>0} \rightarrow \mathbb{R}_{>0}$ is monotonically decreasing and continuous. Then*

$$\mathbb{E}[f(Y) \cdot \mathbb{1}_{Y<t}] > \mathbb{E}[f(X) \cdot \mathbb{1}_{X<t}].$$

Proof. Consider the transformed variables $f(X)$ and $f(Y)$. Since f is monotonically decreasing and continuous, it is invertible, and in particular

$$\Pr(f(X) < x) = \Pr(X > f^{-1}(x)) = 1 - F_X(f^{-1}(x)).$$

It follows that the CDFs of $f(X)$ and $f(Y)$ are $1 - F_X(f^{-1}(x))$ and $1 - F_Y(f^{-1}(x))$, respectively. Then, observe that since $F_X(x) \leq F_Y(x)$ for all $x \in \mathbb{R}$,

$$1 - F_X(x) \geq 1 - F_Y(x)$$

for all $x \in \mathbb{R}$, i.e., $f(X) \leq_{\text{st}} f(Y)$. In particular, since $F_X(x) = 0$ if and only if $1 - F_X(x) = 1$ and *vice versa*, it follows from the invertibility of f that $f(X) <_{\text{st}} f(Y)$. Therefore, after noting that $X < t$ if and only if $f(X) > f(t)$ and similarly for Y , the result follows from Lemma A.61. $\square$

It is relatively straightforward to characterize the stochastic dominance relationships between various (incomplete) beta distributions according to α and β ; see Arab et al. (2021) for full details. Here, we require only the following result, which closely follows the proof of Theorem 1 there.

Lemma A.63. *If $X \sim \text{BETA}_t(\alpha_0, \beta_0)$, $Y \sim \text{BETA}_t(\alpha_1, \beta_1)$, where $\alpha_0 \geq \alpha_1$ and $\beta_0 \leq \beta_1$, then $Y \leq_{\text{st}} X$. If, in addition, either $\alpha_0 > \alpha_1$ or $\beta_0 < \beta_1$, then $Y <_{\text{st}} X$.*

Proof. Consider the CDFs $F_X(s)$ and $F_Y(s)$. We will use the difference $G(s) = F_X(s) - F_Y(s)$ to demonstrate the result. The case where $\alpha_0 = \alpha_1$ and $\beta_0 = \beta_1$ is trivial, so we restrict our attention to the case where one of the inequalities is strict. For simplicity, we assume that $\alpha_0 > \alpha_1$; the case where $\beta_0 < \beta_1$ is virtually identical.

In particular, observe that $G(0) = G(t) = 0$, and that for $s \in (0, t)$,

$$\begin{aligned} G'(s) &= \frac{s^{\alpha_0-1}(1-s)^{\beta_0-1}}{I_t(\alpha_0, \beta_0)} - \frac{s^{\alpha_1-1}(1-s)^{\beta_1-1}}{I_t(\alpha_1, \beta_1)} \\ &= \frac{s^{\alpha_0-1}(1-s)^{\beta_0-1}}{I_t(\alpha_1, \beta_1)} \cdot \left[s^{\alpha_1-\alpha_0}(1-s)^{\beta_1-\beta_0} - \frac{I_t(\alpha_1, \beta_1)}{I_t(\alpha_0, \beta_0)} \right]. \end{aligned}$$

We consider the two multiplicands in the final expression. We note that the first is greater than zero for all $s \in (0, t)$. Therefore, for $s \in (0, t)$, $G'(s) = 0$ if and only if

$$s^{\alpha_1-\alpha_0}(1-s)^{\beta_1-\beta_0} = \frac{I_t(\alpha_1, \beta_1)}{I_t(\alpha_0, \beta_0)}.$$

Now, since $\alpha_0 > \alpha_1$ and $\beta_0 \leq \beta_1$, it follows that the left-hand side of the previous expression is strictly decreasing. In particular, $G'(s) = 0$ for at most one $s \in (0, t)$. Since $G(0) = G(t) = 0$ and $G(t)$ is non-constant, it follows that $G'(s) = 0$ for exactly one $s \in (0, t)$ by Rolle's theorem. In particular, either $G(s) > 0$ for all $s \in (0, t)$ or $G(s) < 0$ for all $t \in (0, t)$.

Consequently,

$$s^{\alpha_1-\alpha_0}(1-s)^{\beta_1-\beta_0} - \frac{I_t(\alpha_1, \beta_1)}{I_t(\alpha_0, \beta_0)}$$

changes sign for some $s_0 \in (0, t)$. Since the minuend is strictly decreasing, it follows that $G'(s) > 0$ for $s \in (0, s_0)$. Therefore, in particular, $G(s) > 0$ for all $s \in (0, s_0)$, and hence for $s \in (0, t)$.

Therefore $F_X(s) > F_Y(s)$ for $s \in (0, t)$, and so $Y <_{\text{st}} X$. $\square$

A.8.2 Proof of Proposition 2.19

Our proof is a relatively straightforward application of Sard's theorem. Intuitively, counterfactual predictive parity imposes three constraints on α_0 , β_0 , α_1 , β_1 , and the group-specific thresholds t_0 and t_1 . These constraints are sufficiently smooth that the zero locus should take the form of a 3-manifold, by the inverse function theorem. Projecting onto the first four coordinates (i.e., eliminating t_0 and t_1) gives rise to a set of measure zero by Sard's theorem.

Proof of Proposition 2.19. We begin by noting that the risk distributions, conditional on $Y(1) = i$, for $i = 0, 1$, take a particularly simple form.

$$\begin{aligned} r(X) \mid A = a_i, Y(1) = 1 &\sim \text{BETA}(\alpha_i + 1, \beta_i), \\ r(X) \mid A = a_i, Y(1) = 0 &\sim \text{BETA}(\alpha_i, \beta_i + 1). \end{aligned}$$

This follows upon noting that

$$\begin{aligned} \Pr(r(X) < t, Y(1) = 1 \mid A = a_i) &= \frac{1}{B(\alpha_i, \beta_i)} \int_0^1 x \cdot \mathbb{1}_{x < t} \cdot x^{\alpha_i - 1} (1 - x)^{\beta_i - 1} dx \\ &= \frac{I_t(\alpha_i + 1, \beta_i)}{B(\alpha_i, \beta_i)}, \\ \Pr(r(X) < t, Y(1) = 0 \mid A = a_i) &= \frac{1}{B(\alpha_i, \beta_i)} \int_0^1 (1 - x) \cdot \mathbb{1}_{x < t} \cdot x^{\alpha_i - 1} (1 - x)^{\beta_i - 1} dx \\ &= \frac{I_t(\alpha_i, \beta_i + 1)}{B(\alpha_i, \beta_i)}. \end{aligned}$$

By Cor. A.9, if there exists a Pareto efficient policy satisfying counterfactual equalized odds, then it must correspond to some multiple threshold policy. In particular, by Eq. (2.4) and the fact that the policy is budget exhausting, and using Prop. 2.15 with $\lambda = 0$, there must be t_0, t_1 such that

$$\begin{aligned} \frac{I_{t_0}(\alpha_0, \beta_0)}{B(\alpha_0, \beta_0)} + \frac{I_{t_1}(\alpha_1, \beta_1)}{B(\alpha_1, \beta_1)} &= 2 - b \\ \frac{I_{t_0}(\alpha_0 + 1, \beta_0)}{B(\alpha_0 + 1, \beta_0)} &= \frac{I_{t_1}(\alpha_1 + 1, \beta_1)}{B(\alpha_1 + 1, \beta_1)} \\ \frac{I_{t_0}(\alpha_0, \beta_0 + 1)}{B(\alpha_0, \beta_0 + 1)} &= \frac{I_{t_1}(\alpha_1, \beta_1 + 1)}{B(\alpha_1, \beta_1 + 1)}. \end{aligned}$$

Here, the first equality encodes budget exhaustion, the second the fact that $\Pr(D = 1 \mid A =$

$a_0, Y(1) = 1) = \Pr(D = 1 \mid A = a_1, Y(1) = 1)$, and the third the fact that $\Pr(D = 1 \mid A = a_0, Y(1) = 0) = \Pr(D = 1 \mid A = a_1, Y(1) = 0)$.

Let $\mathbf{f} : \mathbb{R}^6 \rightarrow \mathbb{R}^3$ be given by

$$\begin{aligned} f_0(\alpha_0, \beta_0, t_0, \alpha_1, \beta_1, t_1) &= \frac{I_{t_0}(\alpha_0, \beta_0)}{B(\alpha_0, \beta_0)} + \frac{I_{t_1}(\alpha_1, \beta_1)}{B(\alpha_1, \beta_1)} - (2 - b), \\ f_1(\alpha_0, \beta_0, t_0, \alpha_1, \beta_1, t_1) &= \frac{I_{t_0}(\alpha_0 + 1, \beta_0)}{B(\alpha_0 + 1, \beta_0)} - \frac{I_{t_1}(\alpha_1 + 1, \beta_1)}{B(\alpha_1 + 1, \beta_1)}, \\ f_2(\alpha_0, \beta_0, t_0, \alpha_1, \beta_1, t_1) &= \frac{I_{t_0}(\alpha_0, \beta_0 + 1)}{B(\alpha_0, \beta_0 + 1)} - \frac{I_{t_1}(\alpha_1, \beta_1 + 1)}{B(\alpha_1, \beta_1 + 1)}. \end{aligned}$$

Then, given $\alpha_0, \beta_0, \alpha_1$ and β_1 , there exists a Pareto optimal policy satisfying counterfactual equalized odds only if there exist t_0 and t_1 such that

$$\mathbf{f}(\alpha_0, \beta_0, t_0, \alpha_1, \beta_1, t_1) = \mathbf{0}.$$

Let $D\mathbf{f}$ denote the Jacobian of $\mathbf{f}$. If we can show that $\mathbf{f}$ is smooth and $D\mathbf{f}$ has full rank, then the proof is complete. For, by Theorem 5.12 in Lee (2013), it follows that $\mathbf{f}^{-1}(\mathbf{0}) \subseteq \mathbb{R}^6$ is a smooth 3-manifold. The restriction of the map

$$\pi : (\alpha_0, \beta_0, t_0, \alpha_1, \beta_1, t_1) \mapsto (\alpha_0, \beta_0, \alpha_1, \beta_1)$$

to $\mathbf{f}^{-1}(\mathbf{0})$ is smooth, and so by Sard's theorem (see Theorem 6.10 in Lee (2013)), since $\mathbb{R}_{>0}^4$ is a trivially a smooth 4-manifold, the measure of the singular values of $\pi|_{\mathbf{f}^{-1}(\mathbf{0})}$ is zero. However, since the maximum rank of $D\pi$ on $\mathbf{f}^{-1}(\mathbf{0})$ is three, every point of $\mathbf{f}^{-1}(\mathbf{0})$ is singular, and consequently, the whole image of π is singular, i.e., $\pi(\mathbf{f}^{-1}(\mathbf{0}))$ has measure zero. However, as argued above, the set of $(\alpha_0, \beta_0, \alpha_1, \beta_1)$ such that there exists a Pareto efficient distribution satisfying counterfactual equalized odds is a subset of $\pi(\mathbf{f}^{-1}(\mathbf{0}))$.

Therefore, it remains only to show that $\mathbf{f}$ is smooth, and that $D\mathbf{f}$ has full rank for all $(\alpha_0, \beta_0, t_0, \alpha_1, \beta_1, t_1) \in \mathbb{R}_{>0}^2 \times [0, 1] \times \mathbb{R}_{>0}^2 \times [0, 1]$. This verification is a routine exercise in linear algebra and multivariable calculus, and is given below.

Smoothness. We note that since smooth functions are closed under composition, it suffices to show that $I_t(\alpha, \beta)$ is a smooth function of t , α , and β . First, we consider partial derivatives with

respect to α and β . If we could differentiate under the integral sign, then we would have that

$$\frac{\partial^{n+m}}{\partial \alpha^n \partial \beta^m} I_t(\alpha, \beta) = \int_0^t \log(x)^n \log(1-x)^m x^{\alpha-1} (1-x)^{\beta-1} dx. \quad (\text{A.29})$$

Recall the well-known condition for the Leibniz integral rule that if $\frac{\partial}{\partial x} \phi(x, t)|_{x=x'}$ is dominated by some integrable $g(t)$ for all x' in some neighborhood of x , then¹⁰

$$\frac{d}{dx} \int_0^t \phi(x, t) dt = \int_0^t \frac{\partial}{\partial x} \phi(x, t) dt.$$

Since the integrand $\log(x)^n \log(1-x)^m x^{\alpha-1} (1-x)^{\beta-1}$ is strictly decreasing in both α and β for $x \in (0, 1)$, it suffices merely to show that $\log(x)^n \log(1-x)^m x^{\alpha-1} (1-x)^{\beta-1}$ is integrable on $(0, 1)$ for all $\alpha > 0$ and $\beta > 0$. Moreover, since

- The whole integrand is bounded on $(\epsilon, 1 - \epsilon)$,
- The factor $\log(1-x)^m (1-x)^{\beta-1}$ is bounded on $(0, \epsilon)$,
- The factor $\log(x)^n x^{\alpha-1}$ is bounded on $(1 - \epsilon, 1)$,

it suffices to show that $\log(x)^n x^{\alpha-1}$ is integrable on $(0, \epsilon)$ and that $\log(1-x)^m (1-x)^{\beta-1}$ is integrable on $(1 - \epsilon, 1)$. Up to the change of variables $x \mapsto 1 - x$, since n, m, α , and β are arbitrary, we see that it suffices to verify that $\log(x)^n x^{\alpha-1}$ alone is integrable on $(0, 1)$. Integrating by parts, we have that

$$\int_0^t \log(x)^n x^{\alpha-1} dx = \left[\frac{\log(x)^n x^\alpha}{\alpha} \right]_0^1 - n \int_0^1 \log(x)^{n-1} x^{\alpha-1} dx.$$

Since, by l'Hôpital's rule, $\lim_{x \rightarrow 0} \log(x)^n x^\alpha = 0$, this expression equals

$$-n \int_0^1 \log(x)^{n-1} x^{\alpha-1} dx.$$

Since $x^{\alpha-1}$ is integrable on $0, 1$ we see inductively that so is $\log(x)^n x^{\alpha-1}$. Therefore, we can differentiate under the integral sign, and Eq. (A.29) holds.

Now, taking derivatives with respect to t yields that

$$\frac{\partial^{n+m+1}}{\partial t \partial \alpha^n \partial \beta^m} I_t(\alpha, \beta) = \log(t)^n \log(1-t)^m t^{\alpha-1} (1-t)^{\beta-1},$$

¹⁰See, e.g., Theorem 6.28 in Klenke (2020).

which is a polynomial in smooth functions of t , and hence is smooth in t . Therefore

$$\frac{\partial^{n+m+k}}{\partial t^k \partial \alpha^n \partial \beta^m} I_t(\alpha, \beta)$$

exists and is a polynomial in t for all $k > 1$.

By continuity, the orders of the partial derivatives can be switched arbitrarily (see, e.g., Theorem 9.40 in Rudin 1976), and so it follows that $\mathbf{f}$ is smooth.

Full rank. By the rank-nullity theorem, it suffices to show that the column rank of $D\mathbf{f}$ is three. However, letting $B_{\alpha, \beta} \sim \text{BETA}(\alpha, \beta)$ we have, by the results of the previous section, that the first three columns (i.e., partial derivatives with respect to α_0 , β_0 , and t_0 , respectively) of $D\mathbf{f}$ are given by

$$\begin{bmatrix} \mathbb{E}[\log(B_{\alpha_0, \beta_0}) \cdot \mathbb{1}_{B_{\alpha_0, \beta_0} < t_0}] & \mathbb{E}[\log(B_{\beta_0, \alpha_0}) \cdot \mathbb{1}_{B_{\beta_0, \alpha_0} < 1-t_0}] & \frac{t_0^{\alpha_0-1} \cdot (1-t_0)^{\beta_0-1}}{B(\alpha_0, \beta_0)} \\ \mathbb{E}[\log(B_{\alpha_0+1, \beta_0}) \cdot \mathbb{1}_{B_{\alpha_0+1, \beta_0} < t_0}] & \mathbb{E}[\log(B_{\beta_0, \alpha_0+1}) \cdot \mathbb{1}_{B_{\beta_0, \alpha_0+1} < 1-t_0}] & \frac{t_0^{\alpha_0} \cdot (1-t_0)^{\beta_0-1}}{B(\alpha_0+1, \beta_0)} \\ \mathbb{E}[\log(B_{\alpha_0, \beta_0+1}) \cdot \mathbb{1}_{B_{\alpha_0, \beta_0+1} < t_0}] & \mathbb{E}[\log(B_{\beta_0+1, \alpha_0}) \cdot \mathbb{1}_{B_{\beta_0+1, \alpha_0} < 1-t_0}] & \frac{t_0^{\alpha_0-1} \cdot (1-t_0)^{\beta_0}}{B(\alpha_0, \beta_0+1)} \end{bmatrix}.$$

It is easy to see that the first two columns are necessarily independent, since all entries are negative but, by Lemmata A.62 and A.63, taking $f(x) = -\log(x)$, we have that

$$\mathbb{E}[\log(B_{\alpha_0+1, \beta_0}) \cdot \mathbb{1}_{B_{\alpha_0+1, \beta_0} < t_0}] > \mathbb{E}[\log(B_{\alpha_0, \beta_0+1}) \cdot \mathbb{1}_{B_{\alpha_0, \beta_0+1} < t_0}]$$

while

$$\mathbb{E}[\log(B_{\beta_0, \alpha_0+1}) \cdot \mathbb{1}_{B_{\beta_0, \alpha_0+1} < 1-t_0}] < \mathbb{E}[\log(B_{\beta_0+1, \alpha_0}) \cdot \mathbb{1}_{B_{\beta_0+1, \alpha_0} < 1-t_0}].$$

Now observe that the final three columns of $D\mathbf{f}$ —i.e., its partial derivatives with respect to α_1 , β_1 , and t_1 , respectively—are given by

$$\begin{bmatrix} \mathbb{E}[\log(B_{\alpha_1, \beta_1}) \cdot \mathbb{1}_{B_{\alpha_1, \beta_1} < t_1}] & \mathbb{E}[\log(B_{\beta_1, \alpha_1}) \cdot \mathbb{1}_{B_{\beta_1, \alpha_1} < 1-t_1}] & \frac{t_1^{\alpha_1-1} \cdot (1-t_1)^{\beta_1-1}}{B(\alpha_1, \beta_1)} \\ -\mathbb{E}[\log(B_{\alpha_1+1, \beta_1}) \cdot \mathbb{1}_{B_{\alpha_1+1, \beta_1} < t_1}] & -\mathbb{E}[\log(B_{\beta_1, \alpha_1+1}) \cdot \mathbb{1}_{B_{\beta_1, \alpha_1+1} < 1-t_1}] & -\frac{t_1^{\alpha_1} \cdot (1-t_1)^{\beta_1-1}}{B(\alpha_1+1, \beta_1)} \\ -\mathbb{E}[\log(B_{\alpha_1, \beta_1+1}) \cdot \mathbb{1}_{B_{\alpha_1, \beta_1+1} < t_1}] & -\mathbb{E}[\log(B_{\beta_1+1, \alpha_1}) \cdot \mathbb{1}_{B_{\beta_1+1, \alpha_1} < 1-t_1}] & -\frac{t_1^{\alpha_1-1} \cdot (1-t_1)^{\beta_1}}{B(\alpha_1, \beta_1+1)} \end{bmatrix}.$$

In particular, we notice that in the fourth column—i.e., partial derivatives with respect to α_1 —the element in the first row is negative, while the elements in the second and third row are positive. Therefore, the fourth column is necessarily independent of the first two columns—i.e., the partial derivatives with respect to α_0 and β_0 —in which every element is negative. Therefore, we have proven that there are three linearly independent columns, and so the rank of $D\mathbf{f}$ is full. $\square$

A.8.3 Proof of Proposition 2.20

To prove the proposition, we must use our characterizations of the conditional tail risks of the beta distribution proven in Appendix A.8.1 above. Note that in Proposition 2.20, for expositional clarity, we parameterize beta distributions in terms of their mean μ and sample size v ; here, for mathematical simplicity, we parameterize them in terms of successes, α , and failures, β , where $\mu = \frac{\alpha}{\alpha+\beta}$ and $v = \alpha + \beta$.

Using the theory above, we begin by proving a modest generalization of Prop. 2.20.

Lemma A.64. *Suppose $\mathcal{A} = \{a_0, a_1\}$, and consider the family $\mathcal{U}$ of utility functions of the form*

$$u(x) = r(x) + \lambda \cdot \mathbb{1}_{\alpha(x)=a_1},$$

indexed by $\lambda \geq 0$, where $r(x) = \mathbb{E}[Y(1) \mid X = x]$. Suppose the conditional distributions of $r(X)$ given A are beta distributed, i.e.,

$$\mathcal{D}(r(X) \mid A = a) = \text{BETA}(\alpha_a, \beta_a),$$

with $\alpha_{a_1} < \alpha_{a_0}$ and $\beta_{a_0} < \beta_{a_1}$. Then any policy satisfying counterfactual predictive parity is strongly Pareto dominated.

Proof. Suppose there were a Pareto efficient policy satisfying counterfactual predictive parity. Let $\lambda = 0$. Then, by Prop. 2.15, we may without loss of generality assume that there exist thresholds t_{a_0}, t_{a_1} such that a threshold policy $\tau(x)$ witnessing Pareto efficiency is given by

$$\tau(x) = \begin{cases} 1 & r(x) > t_{\alpha(x)}, \\ 0 & r(x) < t_{\alpha(x)}. \end{cases}$$

(Note that by our distributional assumption, $\Pr(u(x) = t) = 0$ for all $t \in [0, 1]$.) Since $\lambda \geq 0$, we must have that $t_{a_0} \geq t_{a_1}$. Since $b < 1$, $0 < t_{a_0}$. Therefore,

$$\begin{aligned} \mathbb{E}[Y(1) \mid A = a_0, D = 0] &= \mathbb{E}[r(X) \mid A = a_0, u(X) < t_{a_0}] \\ &\geq \mathbb{E}[r(X) \mid A = a_0, u(X) < t_{a_1}] \\ &> \mathbb{E}[r(X) \mid A = a_1, u(X) < t_{a_1}] \\ &= \mathbb{E}[Y(1) \mid A = a_1, D = 0], \end{aligned}$$

where the first equality follows by the law of iterated expectation, the second from the fact that $t_{a_1} \leq t_{a_0}$, the third from our distributional assumption and Lemmata A.60 and A.63, and the final again from the law of iterated expectation. However, since counterfactual predictive parity is satisfied, $\mathbb{E}[Y(1) \mid A = a_0, D = 0] = \mathbb{E}[Y(1) \mid A = a_1, D = 0]$, which is a contradiction. Therefore, no such threshold policy exists. $\square$

After accounting for the difference in parameterization, Prop. 2.20 follows as a corollary.

Proof of Prop. 2.20. Since $\mu_{a_0} > \mu_{a_1}$, $\alpha_{a_0} = v \cdot \mu_{a_0} > v \cdot \mu_{a_1} = \alpha_{a_1}$ and $\beta_{a_0} = v \cdot (1 - \mu_{a_0}) < v \cdot (1 - \mu_{a_1}) = \beta_{a_1}$. Therefore $\beta_{a_0} < \beta_{a_1}$ and $\alpha_{a_1} < \alpha_{a_0}$, and so, by Lemma A.64, the proposition follows. $\square$

Algorithm A.1: Path-specific counterfactuals

Data: $\mathcal{G}$ (topologically ordered), Π , a , and a'

Result: A sample $Z_{\Pi,a,a'}^*$ from $Z_{\Pi,a,a'}$

```
1 Sample values  $\{U_j^*\}$  for the exogenous variables /* Compute counterfactuals by setting
    $A$  to  $a$  */
2 for  $j = 1, \dots, m$  do
3   if  $V_j = A$  then
4      $V_j^* \leftarrow a$ 
5   else
6      $\wp(V_j)^* \leftarrow \{V_\ell^* \mid V_\ell \in \wp(V_j)\}$ 
7      $V_j^* \leftarrow f_{V_j}(\wp(V_j)^*, U_j^*)$ 
8   end
9 end

/* Compute counterfactuals by setting  $A$  to  $a'$  and propagating values along
   paths in  $\Pi$  */
10 for  $j = 1, \dots, m$  do
11   if  $V_j = A$  then
12      $\bar{V}_j^* \leftarrow a'$ 
13   else
14     for  $V_k \in \wp(V_j)$  do
15       if edge  $(V_k, V_j)$  lies on a path in  $\Pi$  then
16          $V_k^\dagger \leftarrow \bar{V}_k^*$ 
17       else
18          $V_k^\dagger \leftarrow V_k^*$ 
19       end
20     end
21      $\wp(V_j)^\dagger \leftarrow \{V_\ell^\dagger \mid V_\ell \in \wp(V_j)\}$ 
22      $\bar{V}_j^* \leftarrow f_{V_j}(\wp(V_j)^\dagger, U_j^*)$ 
23   end
24 end
25  $Z_{\Pi,a,a'}^* \leftarrow \bar{Z}^*$ 
```

Appendix to Chapter 3

B.1 Proofs

We prove Theorem 3.2 in a more general setting in which we allow the cost of screening c_S and cost of allocation c_A to vary for each individual. In the special case of no screening (but with applicant-specific costs of allocation), we note that our optimization problem is equivalent to the fractional knapsack problem. For fractional knapsack, the greedy strategy—in which one packs items in descending order of their value per weight—yields an optimal solution (Dantzig, 1957). In our case, we show that the same approach can be used to find an optimal allocation of loans once the set of applicants to screen has been appropriately selected.

To start, we generalize our definition of a threshold policy to account for the applicant-specific allocation costs.

Definition B.1 (Cost-Aware Threshold Policy). A *cost-aware threshold policy* is an allocation policy $\mathcal{A}$ for some fixed $t_j \in \mathbb{R} \cup \{-\infty, \infty\}$ and $\alpha_j \in [0, 1]$, $j = 1, \dots, m$, such that

$$A_i = \begin{cases} 1 & \hat{U}_i/c_i > t_{g_i}, \\ \alpha_{g_i} & \hat{U}_i/c_i = t_{g_i}, \\ 0 & \hat{U}_i/c_i < t_{g_i}, \end{cases}$$

where g_i denotes the group membership of individual i and c_i denotes the cost of allocating resources to individual i .

We first analyze the setting of a single group, in which case our threshold policy will have a single threshold, t . Lemma B.2 shows that in this case cost-aware threshold policies are non-dominated: there cannot be an allocation policy with both higher expected utility and lower cost than a

single-threshold policy. Lemma B.3 is an existence theorem for cost-aware threshold policies, which shows that if a particular expected cost or utility can be achieved by an allocation policy, it can be achieved by a threshold policy.

Lemma B.2. *Let $\mathcal{T}$ be a cost-aware threshold policy with a single threshold $t > 0$, and suppose $\mathcal{T}$ has an expected cost C —i.e., $\sum_i \mathbb{E}[c_i \cdot T_i] = C$ —and expected utility Υ —i.e., $\sum_i \mathbb{E}[U_i \cdot T_i] = \Upsilon$. Let $\mathcal{A}$ be any other allocation policy.*

Case 1: *If the expected cost of $\mathcal{A}$ is C , then the expected utility of $\mathcal{A}$ is less than or equal to Υ .*

Case 2: *If the expected utility of $\mathcal{A}$ is Υ , then the expected cost of $\mathcal{A}$ is greater than or equal to C .*

Proof. Let $\Delta_i = T_i - A_i$, and set Δ_i^- be the negative part of Δ_i , i.e., $\Delta_i \cdot \mathbb{1}_{\Delta_i \leq 0}$. Similarly, let Δ_i^+ be the positive part.

The variables Δ_i^+ represent circumstances in which $\mathcal{A}$ “saves money” in expectation over $\mathcal{T}$, and Δ_i^- represents how this portion of the budget is reallocated.

Consider the following events:

$$E_i^> = \{\hat{U}_i > c_i \cdot t\},$$

$$E_i^= = \{\hat{U}_i = c_i \cdot t\},$$

$$E_i^< = \{\hat{U}_i < c_i \cdot t\}.$$

Note that $\Delta_i^+ \neq 0$ only on some subset of $E_i^> \cup E_i^=$. likewise, $\Delta_i^- \neq 0$ only on a subset of $E_i^< \cup E_i^=$. Informally, $\mathcal{A}$ can only save in expected cost over $\mathcal{T}$ when $\hat{U}_i \geq t$, and can only reallocate savings to circumstances in which $\hat{U}_i \leq t$.

Because $\Delta_i^+ \geq 0$ and $\Delta_i^- \leq 0$, it consequently follows that

$$\Delta_i^+ \cdot \hat{U}_i \geq \Delta_i^+ \cdot c_i \cdot t, \tag{B.1}$$

and

$$\Delta_i^- \cdot \hat{U}_i \geq \Delta_i^- \cdot c_i \cdot t. \tag{B.2}$$

In particular,

$$\begin{aligned}
\mathbb{E} \left[\sum_{i=1}^n \Delta_i \hat{U}_i \right] &= \sum_{i=1}^n \mathbb{E}[\Delta_i \hat{U}_i] \\
&= \sum_{i=1}^n \left(\Pr(E_i^>) \cdot \mathbb{E}[\Delta_i^+ \cdot \hat{U}_i \mid E_i^>] \right. \\
&\quad \left. + \Pr(E_i^=) \cdot \mathbb{E}[\Delta_i \cdot \hat{U}_i \mid E_i^=] \right. \\
&\quad \left. + \Pr(E_i^<) \cdot \mathbb{E}[\Delta_i^- \cdot \hat{U}_i \mid E_i^<] \right) \\
&\geq \sum_{i=1}^n \left(\Pr(E_i^>) \cdot \mathbb{E}[\Delta_i^+ \cdot c_i \cdot t \mid E_i^>] \right. \\
&\quad \left. + \Pr(E_i^=) \cdot \mathbb{E}[\Delta_i \cdot c_i \cdot t \mid E_i^=] \right. \\
&\quad \left. + \Pr(E_i^<) \cdot \mathbb{E}[\Delta_i^- \cdot c_i \cdot t \mid E_i^<] \right) \\
&= t \cdot \mathbb{E} \left[\sum_{i=1}^n c_i \Delta_i \right].
\end{aligned}$$

Here the inequality follows from Eqs. (B.1) and (B.2).

Now, since any allocation policy is by definition conditionally independent of U_i given $\hat{U}_i$, it follows that $\mathbb{E}[A_i U_i] = \mathbb{E}[A_i \hat{U}_i]$ and $\mathbb{E}[T_i U_i] = \mathbb{E}[T_i \hat{U}_i]$. In consequence,

$$\mathbb{E} \left[\sum_{i=1}^n \Delta_i U_i \right] \geq t \cdot \mathbb{E} \left[\sum_{i=1}^n c_i \Delta_i \right]. \tag{B.3}$$

In Case 1, the expected costs of $\mathcal{A}$ and $\mathcal{T}$ are equal. It follows that $\mathbb{E}[\sum_i c_i \Delta_i] = 0$; consequently, by Eq. (B.3), $\mathbb{E}[\sum_i \Delta_i U_i] \geq 0$, and so the expected utility of $\mathcal{T}$ is greater than or equal to that of $\mathcal{A}$.

In Case 2, the expected utilities of $\mathcal{A}$ and $\mathcal{T}$ are equal, so that $\mathbb{E}[\sum_i \Delta_i U_i] = 0$. Then, again by Eq. (B.3), since $t > 0$, $\mathbb{E}[\sum_i c_i \Delta_i] \leq 0$. Therefore the expected cost of $\mathcal{T}$ is less than or equal to that of $\mathcal{A}$. $\square$

Lemma B.3. *Let $\Upsilon > 0$ be an achievable expected utility—that is, there is an allocation policy $\mathcal{A}$ such that $\mathbb{E}[\sum_i A_i U_i] = \Upsilon$ —and let C be an achievable expected cost. Then,*

- *There exists a cost-aware threshold policy $\mathcal{T}$ of expected utility Υ ; and,*
- *There exists a (possibly different) cost-aware threshold policy $\mathcal{T}$ of expected cost C .*

Proof. We first address the case of a given positive expected utility. Consider the function $\Upsilon(t) :$

$[0, \infty) \rightarrow \mathbb{R}$ given by

$$\Upsilon(t) = \sum_{i=1}^n \mathbb{E} \left[U_i \cdot \mathbb{1}_{\hat{U}_i/c_i \geq t} \right]. \quad (\text{B.4})$$

We note three facts about $\Upsilon(t)$. First, $\Upsilon(t)$ is monotonically non-increasing for $t \geq 0$. Second, $\Upsilon(0)$ is the maximum expected utility achievable by any allocation policy. Third, $\Upsilon(t) \rightarrow 0$ as $t \rightarrow \infty$.

Suppose $0 < v \leq \Upsilon(0)$. Then there exists some maximal t such that $\Upsilon(t) \geq v$. Let $\rho = \sum_i \Pr(\hat{U}_i/c_i = t)$, $\tau = \sum_i c_i \cdot t_i$, and $\delta = \Upsilon(t) - v$. Define $\alpha = \delta/(\tau \cdot \rho)$. Then the threshold policy defined by (t, α) will have expected utility v . The result follows as all achievable positive expected utilities lie in the half-open interval $(0, \Upsilon(0)]$.

The proof in the case of expected cost is similar. $\square$

We are now equipped to move to analyze settings in which there are multiple groups. We prove the following generalization of Theorem 3.2.

Theorem B.4. *Suppose the constrained optimization problem defined by Eqs. (3.2), (3.3), and (3.4) has a solution $(\mathcal{P}^*, \mathcal{A}^*)$. Then there is a (not necessarily single-threshold) cost-aware threshold policy $\mathcal{T}^*$ such that $(\mathcal{P}^*, \mathcal{T}^*)$ is also a solution.*

Proof of Theorem 3.2. Consider a solution screening and allocation policy pair $(\mathcal{P}^*, \mathcal{A}^*)$ satisfying the allocation diversity constraints.

Note that each of the collections $\mathcal{A}_j^* = \{A_i^*\}_{i \in G_j}$ defines an allocation policy on the subpopulation G_j . Applying Lemma B.3 to each of them yields a threshold policy $\mathcal{T}_j^*$ —defined by the pair (t_j, α_j) —of the same expected cost as $\mathcal{A}_j^*$. By Lemma B.2, $\mathcal{T}_j^*$ achieves the same or higher expected utility on G_j as $\mathcal{A}_j^*$.

Let $\mathcal{T}^*$ be the threshold policy defined by thresholds $t_1, \dots, t_m$ and boundary randomization probabilities $\alpha_1, \dots, \alpha_m$. The expected cost of $\mathcal{T}^*$ is the same as $\mathcal{A}^*$. Moreover, $\mathcal{T}^*$ satisfies the group utility constraints, since the utility it achieves on each group G_j is greater than or equal to $\mathcal{A}^*$. Lastly, the expected utility achieved by $\mathcal{T}^*$ is greater than or equal to that of $\mathcal{A}^*$.

Since $\mathcal{A}^*$ was assumed optimal, it must be that the expected utilities are equal. Therefore the pair $(\mathcal{P}^*, \mathcal{T}^*)$ is also a solution to the constrained optimization problem defined by Eqs. (3.2), (3.3), and (3.4). $\square$

Theorem 3.2 follows as an immediate corollary.

In some cases, it is useful to impose equality rather than inequality diversity allocation constraints. For instance, in some situations the decision maker may wish to obtain the Pareto frontier. Lemma B.5 ensures that it is still sufficient when specific utilities must be achieved on subgroups to restrict to the set of threshold policies when searching for a globally optimal screening and allocation policy.

Lemma B.5. *The conclusion of Theorem 3.2 still holds even if the inequalities in Eq. (3.4) are replaced with equalities, i.e., if the optimal policy pair $(\mathcal{P}^*, \mathcal{A}^*)$ satisfies*

$$\mathbb{E} \left[\sum_{i \in G_j} U_i A_i \right] = \Lambda_j. \quad (\text{B.5})$$

Proof. The proof is identical to that of Theorem 3.2, except that Lemma B.3 is used to construct a threshold policy $\mathcal{T}^*$ achieving exactly utility Λ_j on any constrained group G_j . Case 2 of Lemma B.2 then shows that this utility is achieved at possibly less expected cost than $\mathcal{A}^*$. It follows immediately that the pair $(\mathcal{P}^*, \mathcal{T}^*)$ is a solution to the optimization problem.¹ $\square$

B.2 Optimized Procedure

In our experiments we consider cases in which screening provides no additional information about one group of individuals (i.e., D_i is concentrated at a single point for $i \in G_j$). In this case, it is possible to significantly reduce the number of LPs the decision maker must solve to find an optimal policy or calculate the Pareto frontier. For instance, a lender may already possess all available information relevant to the creditworthiness of those with traditional credit histories.

In this case, the cost of allocating resources to group G_j does not depend on the screening policy $\mathcal{P}$. Therefore, rather than sweeping over all possible pairs of thresholds and boundary randomization probabilities (t_j, α_j) , one can encode the allocation policy to G_j directly as a collection of allocation probabilities $\{a_i\}_{i \in G_j}$.

For notational simplicity, suppose without loss of generality that $m = 2$, that $G_1 = \{1, \dots, n_1\}$ and $G_2 = \{n_1 + 1, \dots, n_1 + n_2\}$, and that D_i is a pointmass for all $i \in G_2$.

¹In this case, it is actually possible that $\mathcal{T}^*$ will result in expected cost savings over $\mathcal{A}^*$, even though $\mathcal{A}^*$ is optimal. This can occur if there is no remaining positive utility to “spend” the savings on, since resources are already allocated at total expected cost less than B in all cases where $\hat{U}_i$ is positive.

Then, let $p = (p_1, \dots, p_{n_1})$ be a length n_1 vector of probabilities indicating the probability of screening the members of G_1 . Let $a = (a_{n_1+1}, \dots, a_{n_1+n_2})$ be a length n_2 vector of probabilities indicating the probability of allocating resources to members of G_2 . We keep the rest of the notation the same as in Section 3.4.

Fix threshold t_1 and boundary randomization probability α_1 for G_1 . Then, the expected utility of any pair of screening policies (for G_1) and allocation policies (for G_2) can be expressed as a linear function of p and a :

$$\sum_{i \in G_1} [q_i \cdot e_i \cdot p_i + o_i \cdot \mu_i \cdot (1 - p_i)] + \sum_{i \in G_2} [\mu_i \cdot a_i]. \quad (\text{B.6})$$

Likewise, our budget constraint can be expressed as follows:

$$\begin{aligned} B \geq & \sum_{i \in G_1} [c_S \cdot p_i + c_A \cdot q_i \cdot p_i + c_A \cdot o_i \cdot (1 - p_i)] \\ & + \sum_{i \in G_2} [c_A \cdot a_i] \end{aligned} \quad (\text{B.7})$$

The allocation diversity constraint will take one of the following two forms. If the constrained group is G_1 ,

$$\sum_{i \in G_1} [q_i \cdot e_i \cdot p_i + o_i \cdot \mu_i \cdot (1 - p_i)] = \Lambda_1. \quad (\text{B.8})$$

Otherwise, if the constrained group is G_2 , the constraint will be

$$\sum_{i \in G_2} [\mu_i \cdot a_i] = \Lambda_2. \quad (\text{B.9})$$

Lastly, we ensure that a and p are probability vectors:

$$a_i \leq 1, \quad (\text{B.10})$$

$$p_i \leq 1. \quad (\text{B.11})$$

Together Eqs. (B.6), (B.7), (B.8), and (B.9)—given the initial data of the threshold policy $\mathcal{T}_1$ on G_1 defined by the pair (t_1, α_1) —define a linear program in the $n_1 + n_2$ decision variables $p = (p_1, \dots, p_n)$ and $a = (a_{n_1+1}, \dots, a_{n_1+n_2})$ that can be solved for an optimal screening policy $\mathcal{P}^*$ (on G_1) and allocation policy $\mathcal{T}_2^*$ (on G_2). By the results of Appendix B.1, $\mathcal{T}_2^*$ can be assumed to have the form of a threshold policy. Sweeping over all such pairs in a (discretized) space $\mathbb{R} \times [0, 1]$, the resulting triple $(\mathcal{P}^*, \mathcal{T}_1^*, \mathcal{T}_2^*)$ that maximizes utility will represent a globally optimal screening

and allocation policy.

Appendix to Chapter 4

C.1 Theories of Discrimination

We briefly review understandings of disparate treatment and disparate impact in U.S. law. In particular, first, we review the distinctions between disparate treatment and disparate impact, especially as it pertains to included-variable bias, and discuss risk-adjusted regression’s role within the three-pronged legal test of disparate impact. Next, we give a simple formal illustration of included-variable bias in the context of *Griggs*. In the final subsection, we provide proofs of the mathematical content of our formal illustration.

C.1.1 Disparate treatment and disparate impact

Disparate treatment derives force from the Equal Protection Clause of the U.S. Constitution’s Fourteenth Amendment, and, as noted in the main text, prohibits government agents from acting with “discriminatory purpose” (*Washington v. Davis*, 1976). Although equal protection law bars policies undertaken with animus, it allows for the limited use of protected attributes to further a compelling government interest. For example, until recently, certain affirmative action programs for college admissions were legally permissible to further the government’s interest in promoting diversity. In 2023, the U.S. Supreme Court overturned the legality of such affirmative action programs, ruling that it was unlawful to explicitly consider race in college admissions decisions (*SFFA v. Harvard*, 2023).

The most widespread statistical test of such intentional discrimination is ordinary linear or logistic regression, in which one estimates the likelihood of favorable (or unfavorable) decisions across groups defined by race, gender, or other legally protected traits. In this approach, the investigator adjusts

for all potentially relevant risk factors, excluding only clear proxies for the protected attributes. Barring omitted-variable bias, non-zero coefficients on the protected traits suggest those factors influenced the decision maker’s actions; in the absence of a compelling justification, such evidence is suggestive of a discriminatory purpose. It is difficult—and perhaps impossible—to rigorously define the *influence*, or causal effect, of largely immutable traits like race on decisions (Greiner and Rubin, 2011; VanderWeele and Robinson, 2014), but a regression of this type is nevertheless considered a reasonable first step to identify discriminatory motive, both by criminologists and by legal scholars (Fagan, 2010; Gaebler et al., 2022). For an equal protection claim to succeed in court, however, one typically needs additional documentary evidence (e.g., acknowledgement of an illegitimate motive) to bolster the statistical evidence.

Disparate impact, in contrast, inquires into a policy’s effects, rather than the intentions of the legislature enacting it or the decision maker executing it. A practice is deemed discriminatory under this standard if it results in an “unjustified adverse impact” on a protected group, irrespective of the policy makers’ intentions or whether protected traits play an explicit role in the administration of the policy. In *Griggs*, for instance, the Duke Power Company’s motivation for instituting its education promotion requirement was irrelevant to the Court’s decision; even if enacted without discriminatory purpose, the policy was deemed discriminatory in its effects and hence illegal. The doctrine stems from statutory rules, rather than constitutional law, and applies only in certain contexts, such as employment (via Title VII of the 1964 Civil Rights Act) and housing (via the Fair Housing Act of 1968). Apart from federal statutes, some states have passed more expansive disparate impact laws, including Illinois and California.

The legal test of disparate impact developed over half a century of case law has three principal elements (Grossman et al., 2024b):

1. **Adverse impact:** The plaintiff first must establish that the policy disproportionately impacts the minority group.
2. **No Justification:** Next, the defendant must establish that the adverse impact has a substantial justification rooted in a legitimate policy goal.
3. **Less discriminatory alternative:** Even if the disparate impact is justified, the plaintiff can nevertheless prevail if they demonstrate that there is an alternative feasible policy with less


adverse impact on the minority group.

Risk-adjusted regression is concerned with the second element, namely, whether the adverse impact has some legitimate policy justification.

As discussed in the main text, the standard statistical test for disparate treatment is a “kitchen sink” regression, where one examines the residual explanatory power of protected group status after including all other available covariates as controls. Because it obscures the distinction between legitimate bases for differences in decisions and distinctions which are actually unrelated to the putative policy goal, “kitchen sink” regression is generally an inappropriate empirical method for measuring disparate impact. To assess claims of *unjustified* disparate impact, in, for example, *Griggs*, one would ideally compare decision rates for similarly *qualified* groups of applicants (e.g., similarly qualified White and Black candidates), rather than the more fine-grained comparison of applicants who are “similarly situated”—that is, generally similar, including in regards unrelated to their qualification. In general, effecting this comparison merely through the inclusion and exclusion of different regression covariates is difficult: If one does not, or cannot, adjust for sufficiently many covariates, omitted-variable bias may skew results; conversely, if one does adjust for a rich set of covariates, included-variable bias may corrupt conclusions.¹

C.1.2 A formal illustration of included-variable bias

We give a simple formal illustration of included-variable bias in the context of *Griggs v. Duke Power Co.* Consider the following data-generating process, depicted by the DAG below.



Here, C indicates the race of an employee, with $C = 0$ representing a White employee and $C = 1$ a Black employee; and H indicates whether an employee has a high-school degree. We assume that H

¹A related body of work conceptualizes discrimination as race-specific differences in decision error rates (Arnold et al., 2021, 2022; Bohren et al., 2025; Coston et al., 2020; Hardt et al., 2016). For example, one might consider differences in loan rejection rates between White and Black applicants who in reality would repay their loans (i.e., a counterfactual false negative rate). Recent work, however, has observed that by conditioning on *ex post* (potential) outcomes—as opposed to *ex ante* risk, as we do—such approaches suffer from inframarginality and associated statistical issues, rendering them problematic measures of common legal and policy understandings of discrimination (Corbett-Davies et al., 2023; Mayson, 2018; Nilforoshan et al., 2022; Ayres, 2005).

and C have a joint distribution given by

$$\begin{aligned}\Pr(H = 1, C = 1) &= \Pr(H = 0, C = 0) = \frac{1 + \rho}{4}, \\ \Pr(H = 1, C = 0) &= \Pr(H = 0, C = 1) = \frac{1 - \rho}{4}.\end{aligned}$$

In particular, it follows that both H and C are marginally $\text{Bern}(\frac{1}{2})$ with correlation ρ , $-1 < \rho < 1$. Importantly, we need not assume that race C “causes” an employee to have a high school degree but only that race and having a high school degree have some correlation ρ (although, in this particular example, a causal mechanism like discriminatory enrollment policies is plausible).

Further, $E \sim \text{Unif}(0, 1)$ denotes an employee’s level of experience, interpreted as the proportion of employees who have been employed for less time than them. Finally, by A , we denote whether the employee was promoted. Recalling the details of *Griggs*, we assume that

$$\Pr(A = 1 \mid E, H) = E \cdot H.$$

That is, only employees with high school degrees are promoted, and their probability of promotion depends on their experience.

A natural quantity to consider is the average difference in promotion rates for “similarly situated” Black and White employees—that is, the average difference in promotion rates for employees with the same experience level and high school degree. In this case, it is straightforward to show that

$$\mathbb{E}[\Pr(A = 1 \mid E, H, C = 1) - \Pr(A = 1 \mid E, H, C = 0)] = 0. \quad (\text{C.1})$$

Since race does not factor into promotion decisions, the average difference in hiring rates for similarly situated Black and White employees is zero, suggesting an absence of disparate *treatment*. However, if, following the Court’s reasoning in *Griggs*, one were to consider the average difference in promotion rates for employees with the same *experience level*—regardless of whether they had a high school degree—one would obtain

$$\mathbb{E}[\Pr(A = 1 \mid E, C = 1) - \Pr(A = 1 \mid E, C = 0)] = \frac{\rho}{2}. \quad (\text{C.2})$$

Following the facts of *Griggs*, we expect $\rho < 0$, indicating that Black employees are less likely than White employees to have attained a high-school degree. As a result, as shown in (C.2),

Black employees at the same level of experience are less likely to be promoted. In particular, this estimand captures the disparate *impact* of the promotion policy. This stylized example illustrates the importance of conditioning on the appropriate variables to avoid included-variable bias in studies of discrimination. Here, high-school degree attainment is inappropriate to condition on because it has little relationship to job performance.

Non-parametric estimands like those shown in Eqs. (C.1) and (C.2) can be statistically challenging to estimate in practice, so it is common instead to use linear regression or other parametric models to approximate these estimands. Imagine if an analyst were to fit a linear kitchen-sink regression model of A given experience level (E), high school degree (H), the interaction between these two terms ($E \cdot H$), and race (C),

$$\Pr(A = 1 \mid E, H, C) = \beta_0 + \beta_E \cdot E + \beta_H \cdot H + \beta_{E:H} \cdot E \cdot H + \beta_C \cdot C. \quad (\text{C.3})$$

Here, β_C is typically interpreted as the magnitude of “discrimination.” In this example, $\hat{\beta}_C$ tends to zero in large samples, suggesting a lack of disparate *treatment*. But, despite fitting a correctly specified model of decisions, the analyst would fail to detect the disparate impact of the hiring policy.

Imagine, instead, that the analyst regressed decisions only on years of experience (E) and race (C), as the Court indicated they should, i.e.,

$$\Pr(A = 1 \mid E, C) = \beta'_0 + \beta'_E \cdot E + \beta'_C \cdot C. \quad (\text{C.4})$$

In this case, as with the non-parametric estimand in (C.2), $\hat{\beta}'_C$ tends to $\frac{\ell}{2}$, reflecting the fact that it correctly captures the disparate impact of the promotion policy. (In this specific example, the parametric and non-parametric estimands are identical, but this is not always the case.)

Following the Court’s ruling, we have assumed that any use of education is *unjustified*, while use of experience is *justified* when making promotion decisions. As a result, the key expression of interest is that given in (C.2), which quantifies disparities in decision rates after conditioning on experience (E) alone. In practice, however, all covariates are usually at least weakly informative about one’s qualifications. In the context of *Griggs*, high school degree attainment, even among employees with the same level of experience, is likely at least somewhat predictive of job performance. The more general definition in (4.1) provides one useful way of conceptualizing disparate impact in this more

realistic setting.

C.1.3 Proofs

We begin by proving the results for the parametric formulation in Eqs. (C.3) and (C.4), and then use the parametric results to prove the non-parametric claims given in Eqs. (C.1) and (C.2).

C.1.3.1 Linear Regression Formulation

Recall that the OLS coefficient estimates are given in this case by

$$\left(\sum_{i=1}^n \begin{bmatrix} 1 & E_i & H_i & E_i \cdot H_i & C_i \\ E_i & E_i^2 & E_i \cdot H_i & E_i^2 \cdot H_i & E_i \cdot C_i \\ H_i & E_i \cdot H_i & H_i^2 & E_i \cdot H_i^2 & H_i \cdot C_i \\ E_i \cdot H_i & E_i^2 \cdot H_i & E_i \cdot H_i^2 & E_i^2 \cdot H_i^2 & E_i \cdot H_i \cdot C_i \\ C_i & E_i \cdot C_i & H_i \cdot C_i & E_i \cdot H_i \cdot C_i & C_i^2 \end{bmatrix} \right)^{-1} \left(\sum_{i=1}^n \begin{bmatrix} A_i \\ E_i \cdot A_i \\ H_i \cdot A_i \\ E_i \cdot H_i \cdot A_i \\ C_i \cdot A_i \end{bmatrix} \right)$$

Note that because of the matrix inversion, we can divide both sums by n . Therefore, applying the continuous mapping theorem and the strong law of large numbers, we see that this converges to

$$\begin{bmatrix} 1 & \mathbb{E}[E] & \mathbb{E}[H] & \mathbb{E}[E \cdot H] & \mathbb{E}[C] \\ \mathbb{E}[E] & \mathbb{E}[E^2] & \mathbb{E}[E \cdot H] & \mathbb{E}[E^2 \cdot H] & \mathbb{E}[E \cdot C] \\ \mathbb{E}[H] & \mathbb{E}[E \cdot H] & \mathbb{E}[H^2] & \mathbb{E}[E \cdot H^2] & \mathbb{E}[H \cdot C] \\ \mathbb{E}[E \cdot H] & \mathbb{E}[E^2 \cdot H] & \mathbb{E}[E \cdot H^2] & \mathbb{E}[E^2 \cdot H^2] & \mathbb{E}[E \cdot H \cdot C] \\ \mathbb{E}[C] & \mathbb{E}[E \cdot C] & \mathbb{E}[H \cdot C] & \mathbb{E}[E \cdot H \cdot C] & \mathbb{E}[C^2] \end{bmatrix}^{-1} \begin{bmatrix} \mathbb{E}[A] \\ \mathbb{E}[E \cdot A] \\ \mathbb{E}[H \cdot A] \\ \mathbb{E}[E \cdot H \cdot A] \\ \mathbb{E}[C \cdot A] \end{bmatrix}.$$

Now, recalling the joint and marginal distributions of these variables, we have that

$$\mathbb{E}[E] = \frac{1}{2}, \quad \mathbb{E}[H] = \frac{1}{2}, \quad \mathbb{E}[C] = \frac{1}{2}.$$

Moreover,

$$\mathbb{E}[E^2] = \frac{1}{3}, \quad \mathbb{E}[H^2] = \frac{1}{2}, \quad \mathbb{E}[C^2] = \frac{1}{2}.$$

By the definition of covariance,

$$\rho = \frac{\mathbb{E}[H \cdot C] - \mathbb{E}[H] \cdot \mathbb{E}[C]}{\sqrt{\text{VAR}(H) \cdot \text{VAR}(C)}} = \frac{\mathbb{E}[H \cdot C] - \frac{1}{4}}{\sqrt{\frac{1}{4} \cdot \frac{1}{4}}} = 4 \cdot \mathbb{E}[H \cdot C] - 1,$$

whence $\mathbb{E}[H \cdot C] = \frac{\rho+1}{4}$. By independence,

$$\begin{aligned} \mathbb{E}[E \cdot H] &= \mathbb{E}[E] \cdot \mathbb{E}[H] = \frac{1}{4}, & \mathbb{E}[E \cdot C] &= \mathbb{E}[E] \cdot \mathbb{E}[C] = \frac{1}{4}, \\ \mathbb{E}[E^2 \cdot H] &= \mathbb{E}[E^2] \cdot \mathbb{E}[H] = \frac{1}{6}, & \mathbb{E}[E \cdot H^2] &= \mathbb{E}[E \cdot H] = \frac{1}{4}, \\ \mathbb{E}[E^2 \cdot H^2] &= \mathbb{E}[E^2 \cdot H] = \frac{1}{6}, & \mathbb{E}[E \cdot H \cdot C] &= \mathbb{E}[E] \cdot \mathbb{E}[H \cdot C] = \frac{\rho+1}{8}, \end{aligned}$$

Finally,

$$\begin{aligned} \mathbb{E}[A] &= \mathbb{E}[\mathbb{E}[A \mid E, H]] & \mathbb{E}[E \cdot A] &= \mathbb{E}[\mathbb{E}[E \cdot A \mid E, H]] \\ &= \mathbb{E}[E \cdot H] & &= \mathbb{E}[E^2 \cdot H] \\ &= \frac{1}{4}, & &= \frac{1}{6}, \end{aligned}$$

$$\begin{aligned} \mathbb{E}[H \cdot A] &= \mathbb{E}[\mathbb{E}[H \cdot A \mid E, H]] & \mathbb{E}[E \cdot H \cdot A] &= \mathbb{E}[\mathbb{E}[E \cdot H \cdot A \mid E, H]] \\ &= \mathbb{E}[E \cdot H^2] & &= \mathbb{E}[E^2 \cdot H^2] \\ &= \frac{1}{4}, & &= \frac{1}{6}, \end{aligned}$$

$$\begin{aligned} \mathbb{E}[C \cdot A] &= \mathbb{E}[\mathbb{E}[C \cdot A \mid E, H, C]] \\ &= \mathbb{E}[E \cdot H \cdot C] \\ &= \frac{\rho+1}{8}. \end{aligned}$$

Thus, we have that

$$\begin{bmatrix} 1 & \mathbb{E}[E] & \mathbb{E}[H] & \mathbb{E}[E \cdot H] & \mathbb{E}[C] \\ \mathbb{E}[E] & \mathbb{E}[E^2] & \mathbb{E}[E \cdot H] & \mathbb{E}[E^2 \cdot H] & \mathbb{E}[E \cdot C] \\ \mathbb{E}[H] & \mathbb{E}[E \cdot H] & \mathbb{E}[H^2] & \mathbb{E}[E \cdot H^2] & \mathbb{E}[H \cdot C] \\ \mathbb{E}[E \cdot H] & \mathbb{E}[E^2 \cdot H] & \mathbb{E}[E \cdot H^2] & \mathbb{E}[E^2 \cdot H^2] & \mathbb{E}[E \cdot H \cdot C] \\ \mathbb{E}[C] & \mathbb{E}[E \cdot C] & \mathbb{E}[H \cdot C] & \mathbb{E}[E \cdot H \cdot C] & \mathbb{E}[C^2] \end{bmatrix}^{-1}$$

equals

$$\begin{bmatrix} 1 & \frac{1}{2} & \frac{1}{2} & \frac{1}{4} & \frac{1}{2} \\ \frac{1}{2} & \frac{1}{3} & \frac{1}{4} & \frac{1}{6} & \frac{1}{4} \\ \frac{1}{2} & \frac{1}{4} & \frac{1}{2} & \frac{1}{4} & \frac{\rho+1}{4} \\ \frac{1}{4} & \frac{1}{6} & \frac{1}{4} & \frac{1}{6} & \frac{\rho+1}{8} \\ \frac{1}{2} & \frac{1}{4} & \frac{\rho+1}{4} & \frac{\rho+1}{8} & \frac{1}{2} \end{bmatrix}^{-1} = \begin{bmatrix} \frac{7\rho+9}{\rho+1} & -12 & -\frac{6\rho+8}{\rho+1} & 12 & -\frac{2}{\rho+1} \\ -12 & 24 & 12 & -24 & 0 \\ -\frac{6\rho+8}{\rho+1} & 12 & \frac{12\rho^2-16}{\rho^2-1} & -24 & \frac{4\rho}{\rho^2-1} \\ 12 & -24 & -24 & 48 & 0 \\ -\frac{2}{\rho+1} & 0 & \frac{4\rho}{\rho^2-1} & 0 & -\frac{4}{\rho^2-1} \end{bmatrix}$$

and so the OLS regression coefficients converge almost surely to

$$\begin{bmatrix} \frac{7\rho+9}{\rho+1} & -12 & -\frac{6\rho+8}{\rho+1} & 12 & -\frac{2}{\rho+1} \\ -12 & 24 & 12 & -24 & 0 \\ -\frac{6\rho+8}{\rho+1} & 12 & \frac{12\rho^2-16}{\rho^2-1} & -24 & \frac{4\rho}{\rho^2-1} \\ 12 & -24 & -24 & 48 & 0 \\ -\frac{2}{\rho+1} & 0 & \frac{4\rho}{\rho^2-1} & 0 & -\frac{4}{\rho^2-1} \end{bmatrix} \begin{bmatrix} \frac{1}{4} \\ \frac{1}{6} \\ \frac{1}{4} \\ \frac{1}{6} \\ \frac{\rho+1}{8} \end{bmatrix} = \begin{bmatrix} 0 \\ 0 \\ 0 \\ 1 \\ 0 \end{bmatrix}.$$

On the other hand, if the analyst were to drop high school graduation from their regression, then we would have that the regression coefficients converge almost surely to

$$\begin{bmatrix} 1 & \mathbb{E}[E] & \mathbb{E}[C] \\ \mathbb{E}[E] & \mathbb{E}[E^2] & \mathbb{E}[E \cdot C] \\ \mathbb{E}[C] & \mathbb{E}[E \cdot C] & \mathbb{E}[C^2] \end{bmatrix}^{-1} \begin{bmatrix} \mathbb{E}[A] \\ \mathbb{E}[E \cdot A] \\ \mathbb{E}[C \cdot A] \end{bmatrix} = \begin{bmatrix} 1 & \frac{1}{2} & \frac{1}{2} \\ \frac{1}{2} & \frac{1}{3} & \frac{1}{4} \\ \frac{1}{2} & \frac{1}{4} & \frac{1}{2} \end{bmatrix}^{-1} \begin{bmatrix} \frac{1}{4} \\ \frac{1}{6} \\ \frac{\rho+1}{8} \end{bmatrix} = \begin{bmatrix} -\frac{\rho}{4} \\ \frac{1}{2} \\ \frac{\rho}{2} \end{bmatrix}.$$

C.1.3.2 Non-Parametric Formulation

Recall the two different non-parametric estimands:

$$\mathbb{E}[\mathbb{E}[A \mid E, H, C = 1] - \mathbb{E}[A \mid E, H, C = 0]]$$

and

$$\mathbb{E}[\mathbb{E}[A \mid E, C = 1] - \mathbb{E}[A \mid E, C = 0]].$$

We note that in the first case, since $\mathbb{E}[A \mid E, H, C] = E \cdot H$, we have that

$$\mathbb{E}[\mathbb{E}[A \mid E, H, C = 1] - \mathbb{E}[A \mid E, H, C = 0]] = \mathbb{E}[E \cdot H - E \cdot H] = 0.$$

In the second case, using the expectations calculated above, we have that

$$\begin{aligned} \mathbb{E}[A \mid E, C = c] &= \mathbb{E}[A \mid E, H = 1, C = c] \cdot \Pr(H = 1 \mid E, C = c) \\ &\quad + \mathbb{E}[A \mid E, H = 0, C = c] \cdot \Pr(H = 0 \mid E, C = c) \\ &= E \cdot \Pr(H = 1 \mid E, C = c) \\ &= E \cdot \Pr(H = 1 \mid C = c) \\ &= E \cdot \frac{\mathbb{E}[H \cdot \mathbf{1}(C = c)]}{\mathbb{E}[\mathbf{1}(C = c)]} \\ &= E \cdot \frac{\frac{1 - (-1)^c \cdot \rho}{4}}{\frac{1}{2}} \\ &= E \cdot \frac{1 - (-1)^c \cdot \rho}{2}. \end{aligned}$$

Here we have used the fact that $A = 0$ when $H = 0$ in the second equality, the independence of E from H and C in the third, and various expectations calculated above in the remaining equalities.

Thus, it follows that

$$\mathbb{E}[\mathbb{E}[A \mid E, C = 1] - \mathbb{E}[A \mid E, C = 0]] = \mathbb{E} \left[E \cdot \frac{1 + \rho}{2} - E \cdot \frac{1 - \rho}{2} \right] = \mathbb{E}[E \cdot \rho] = \frac{\rho}{2}.$$

This exactly equals β'_C in (C.4).

C.2 Sensitivity analysis

In the following appendix, we develop our method for gauging the sensitivity of risk-adjusted regression to omitted-variable bias in the analyst’s risk estimates. We discuss this sensitivity analysis in two parts. First, in Section C.2.1, we motivate our approach to sensitivity analysis in risk-adjusted regression, more carefully introducing our formal setting, and laying out the key mathematical results underpinning our method. Then, in Section C.2.2, we prove our method’s correctness and derive bounds on its complexity and the accuracy of the approximations it makes use of.

C.2.1 Introduction

C.2.1.1 Formal Setup

We model the data-generating process as consisting of draws of tuples $(X, \tilde{X}, C, W, A)$ where $X \in \mathcal{X}$, $\tilde{X} \in \tilde{\mathcal{X}}$, $C \in \{1, \dots, m\}$, $W \in \{0, 1\}$, and $A \in \{0, 1\}$. Here A is a binary decision (either human or algorithmic), X is the set of covariates on which that decision is based, and C represents membership in some protected class. The binary variable W represents a latent property of interest to the decision maker—discussed further below. However, we assume the decision is made based on X alone, i.e.,

$$A \perp\!\!\!\perp W \mid X.$$

Finally, we assume that an analyst observes the tuple $(\tilde{X}, C, A, A \cdot W)$, based on which they seek to estimate disparate impact in the decision process. Here $\tilde{X}$ represents an alternative set of covariates available to the analyst that may differ from X in arbitrary ways. Additionally, the term $A \cdot W$ reflects the fact that W is only observed by the analyst when action $A = 1$ occurs. As noted in the main text, this general mathematical framework applies to a wide variety of policy contexts in which a decision maker, based on partial information X , aims to make a decision A targeting some binary outcome W .

C.2.1.2 Overview

Our definition of disparate impact in (4.2) depends on the true risks, R_i . But, in practice, analysts observe only partial information on X and W , and so at best can construct only imperfect estimates

of risks, $\hat{R}_i$. For instance, in our running police example, officers might base search decisions in part on subtle behavioral cues that are not documented in the data; and analysts typically would not know whether individuals who were *not* searched were in fact carrying a weapon. In Section 4.2, in the main text, we discuss various approaches to estimating risk in light of these challenges. Whatever approach one adopts, the accuracy of one’s conclusions depends critically on the accuracy of one’s estimates of risk. To address this issue, we develop a novel sensitivity analysis (implemented in the **rar** R package on CRAN) that draws inspiration from methods for sensitivity analysis popular in the causal inference literature (Jung et al., 2020; Rosenbaum and Rubin, 1983a)—though we emphasize that our framework does not itself involve estimating causal effects.

To start, we recall (4.5), in which we assumed that there is a known constant $\epsilon \geq 0$ such that

$$\frac{1}{n} \sum_{i=1}^n |R_i - \hat{R}_i| \leq \epsilon;$$

that is, that the true risks and the estimated risks differ on average by at most ϵ .² Given this assumption, the goal of our sensitivity analysis is to understand how different our estimate of disparate impact, $\hat{\beta}_j - \hat{\beta}_1$, could be if we had access to the true risk R_i instead of the estimated risk $\hat{R}_i$. In practice, an analyst would examine the robustness of conclusions to different choices of ϵ .

We could attempt to bound our estimate of disparate impact by searching over all possible choices of R_i satisfying the constraint in (4.5). But such an approach is overly conservative, as the observed data themselves rule out possible values of R_i . In our policing example, for instance, we expect the average true risk of searched individuals to approximately equal the proportion of searched individuals carrying a weapon. Since, by assumption, the observed data tell us about this latter quantity, they constrain the possible risk distributions we must consider.

To formalize our approach, a key quantity to consider is the average (true) risk on each stratum,

$$\frac{1}{n_{j,a}} \sum_{i \in \mathcal{S}_{j,a}} R_i, \tag{C.5}$$

where

$$\mathcal{S}_{j,a} = \{i : C_i = j, A_i = a\}, \quad n_{j,a} = |\mathcal{S}_{j,a}|,$$

²These differences could arise due to either measurement or modelling error. However, since measurement error tends to vanish as the sample size increases, it is less of a concern, and is better captured through bootstrapped bounds on the sensitivity analysis, as we discuss below.

i.e., $\mathcal{S}_{j,a}$ denotes the stratum containing all those individuals in group j for whom the decision maker took action a , and $n_{j,a}$ is its size. By the law of large numbers, the average risk on stratum $\mathcal{S}_{j,a}$ is approximated by $\mathbb{E}[R \mid C = j, A = a]$. Further,

$$\begin{aligned}\mathbb{E}[R \mid C = j, A = a] &= \mathbb{E}[\mathbb{E}[W \mid X] \mid C = j, A = a] \\ &= \mathbb{E}[\mathbb{E}[W \mid X, C, A] \mid C = j, A = a] \\ &= \mathbb{E}[W \mid C = j, A = a],\end{aligned}$$

where the first equality follows by definition; the second from $A \perp\!\!\!\perp W \mid X$ if we assume that C is a function of X ; and the third by the law of iterated expectations. When $A = 1$, by our assumption in Section 4.1, the above quantity is identified by data available to the analyst. In particular, in our policing example, a consistent estimator of $\mathbb{E}[W \mid C = j, A = 1]$ is the proportion of searched individuals possessing a weapon, among those individuals belonging to group j . We denote by ρ_j the analyst's estimate of $\mathbb{E}[W \mid C = j, A = 1]$ based on the observed data.³

Finally, putting the pieces together, for each particular j^* we seek the largest and smallest values of $\hat{\beta}_{j^*} - \hat{\beta}_1$, as estimated with R_i , subject to two key constraints: (1) the average absolute deviation between R_i and $\hat{R}_i$ is less than ϵ ; and (2) R_i is consistent with the observed data on each stratum $\mathcal{S}_{j,1}$. This is succinctly expressed as an optimization problem (henceforth “**the base problem**”):

$$\begin{aligned}\text{Optimize}_{\mathbf{R} \in \mathbb{R}^n} \quad & \hat{\beta}_{j^*} - \hat{\beta}_1 \\ \text{s.t.} \quad & \frac{1}{n} \sum_{i=1}^n |R_i - \hat{R}_i| \leq \epsilon \\ & \frac{1}{n_{j,1}} \sum_{i \in \mathcal{S}_{j,1}} R_i = \rho_j, \quad (j = 1, \dots, m) \\ & R_i \leq u_i, \quad (i = 1, \dots, n) \\ & R_i \geq \ell_i, \quad (i = 1, \dots, n)\end{aligned} \tag{C.6}$$

where, by “Optimize,” we simply mean both maximize and minimize, and $\hat{R}_i$ and ρ_j are fixed by

³One might alternatively estimate $\mathbb{E}[W \mid C = j, A = 1]$ by the average *estimated* risks on the stratum, namely $\frac{1}{|\mathcal{S}_{j,1}|} \cdot \sum_{i \in \mathcal{S}_{j,1}} \hat{R}_i$. This alternative, which we adopt in our empirical analysis in Section 4.2, has the interpretive advantage that our sensitivity analysis exactly recovers the original analysis when $\epsilon = 0$. In cases where risk is not conditioned on C , then true risks need not be calibrated within groups. As a result, the average estimated risk may approximate the average true risk less well. For instance, this might be the case when decision makers are legally prohibited from considering race. In our policing example, however, the estimated risks are essentially unchanged by the exclusion or inclusion of race, suggesting that the impact of potential miscalibration is small.

the data as discussed above.⁴ The additional upper and lower bounds on R_i arise from the fact that the risks must be probabilities, i.e., $0 \leq R_i \leq 1$, though our approach accommodates tighter, individual-specific bounds $\ell_i \leq R_i \leq u_i$ set by the analyst.

Below, we give a polynomial time approximation to the solution of the base problem when the constraints are “sortable.” Sortability is a mild hypothesis which holds in the typical case, $\ell_i = 0$ and $u_i = 1$ for all i . It also holds if, e.g., ℓ_i and u_i differ from $\hat{R}_i$ by the additive constant Γ on the log odds scale, as is common in many kinds of sensitivity analysis (e.g., Rosenbaum, 2002).

Definition C.1 (Sortability). We say that the *constraints are sortable* if there exists a permutation π of $\{1, \dots, n\}$ such that $\hat{R}_{\pi(1)} \leq \dots \leq \hat{R}_{\pi(n)}$, $\ell_{\pi(1)} \leq \dots \leq \ell_{\pi(n)}$, and $u_{\pi(1)} \leq \dots \leq u_{\pi(n)}$.

There are two main obstacles to solving the base problem above. First, the dimension of the search space is $n - m$, where n is the number of observations and m is the number of groups. Second, the regression coefficients $\hat{\beta}_j$ are non-convex functions of the true risk vector $\mathbf{R} = (R_1, \dots, R_n)$. Addressing these twin challenges involves a detailed analysis of the underlying geometry of the problem. Here we present an outline of our approach that illustrates the key ideas, with the full exposition in Appendix C.2.2.

It is useful to think of the base problem as an optimization over subproblems (henceforth “**the parameterized problem**”) where the average (true) risk is assumed known and equal to some τ_j for *unobserved* strata $\mathcal{S}_{j,0}$ as well. Using the closed form of OLS regression, we can show that solving the parameterized problem reduces to—although is not equivalent to—solving the following

⁴It may seem inappropriate to fix the average risks on the observed stratum $\mathcal{S}_{j,1}$ to be exactly ρ_j , since ρ_j is estimated with error. One could instead allow ρ_j to vary within some range (e.g., a 95% confidence interval), solving the parameterized problem for various ρ_j as well as τ_j , as detailed below. However, doing so doubles the dimension of the search space, which is computationally costly, and does not fully take advantage of available information about the distribution of the estimation error. A more principled and computationally feasible approach to dealing with estimation error is to bootstrap confidence intervals for the bounds following Zhao et al. (2019), as we do in Section 4.2.2 below, where ρ_j is re-estimated in each bootstrapped resample, and the base problem is solved with corresponding equality constraints.

optimization problem (henceforth “**the simplified problem**”):

$$\begin{aligned}
& \underset{\mathbf{R} \in \mathbb{R}^n}{\text{Optimize}} && \frac{1}{n} \sum_{i=1}^n R_i^2 \\
& \text{s.t.} && \frac{1}{n} \sum_{i=1}^n |R_i - \hat{R}_i| \leq \epsilon, \\
& && \frac{1}{n_{j,1}} \cdot \sum_{i \in \mathcal{S}_{j,1}} R_i = \rho_j, \quad (j = 1, \dots, m) \\
& && \frac{1}{n_{j,0}} \cdot \sum_{i \in \mathcal{S}_{j,0}} R_i = \tau_j, \quad (j = 1, \dots, m) \\
& && R_i \leq u_i, \quad (i = 1, \dots, n) \\
& && R_i \geq \ell_i, \quad (i = 1, \dots, n)
\end{aligned} \tag{C.7}$$

An efficient method for solving the simplified problem then yields an efficient method for solving the base problem simply by searching over this much smaller m -dimensional parameter space of $(\tau_1, \dots, \tau_m)$. In practice, m typically equals two or three, and so this final optimization step can be solved either using a grid search, as we do in Section 4.2, or using other optimization techniques for low-dimensional non-convex problems.

The objective of the simplified problem, $\frac{1}{n} \sum_{i=1}^n R_i^2$, is convex, as are the other constraints on $\mathbf{R}$, and so, in principle, minimization in the simplified problem can be achieved efficiently and accurately using interior-point methods (Boyd and Vandenberghe, 2004). However, by carefully examining the KKT conditions (Karush, 1939; Kuhn and Tucker, 1951), it is possible to see that the minimizing risk vector must adhere to a special form, which we term “minimization normal form.” Specifically, there exist thresholds $t_{j,a}^{\text{lwr}}$ and $t_{j,a}^{\text{upr}}$ for each stratum such that the minimizing R_i is obtained by “pulling up” low estimated risks to $t_{j,a}^{\text{lwr}}$, “pulling down” high estimated risks to $t_{j,a}^{\text{upr}}$, and leaving unchanged intermediate risks. An example of a risk vector in minimization normal form is shown in Figure C.1.

This observation reduces the dimension of the search space from $n - 2m$ to $2m$; that is, by the above, one need only search over the $2m$ thresholds. One can, however, do better. A careful examination of the KKT conditions reveals that while the thresholds $t_{j,a}^{\text{lwr}}$ and $t_{j,a}^{\text{upr}}$ may vary by stratum, the gap between them $\Delta = t_{j,a}^{\text{upr}} - t_{j,a}^{\text{lwr}}$ must be the same across strata. Moreover, once Δ is fixed, the thresholds $t_{j,a}^{\text{lwr}}$ and $t_{j,a}^{\text{upr}}$ are determined by the strata-specific average risks ρ_j and τ_j . As a result, finding the minimizing risk vector for the simplified problem reduces to optimizing over a single parameter, Δ , which, in our setting, can be done very efficiently. In particular, we give an




Figure C.1: *Minimization normal form.* Black dots represent estimated risks $\hat{R}_i$, the error bars represent the allowable range $[\ell_i, u_i]$, and red dots indicate the corresponding value of the normal form vector R_i . Here Δ denotes the gap between the upper and lower thresholds.

$O(n \cdot \log(m))$ algorithm that finds the exact solution.

Maximization requires more care, since the maximization problem is non-convex, and the general problem of maximizing a quadratic objective over a convex set is NP-hard (Sahni, 1974). Using a similar, but more delicate, version of the techniques used to minimize the simplified problem, we give an $O(m \cdot (\epsilon/\gamma)^2 + n)$ -time algorithm for approximating the maximum of the simplified problem. The key steps described above for optimizing the base problem are given in Algorithm C.1.

Putting these results together, we have the following theorem, the proof of which constitutes the majority of Appendix C.2.2.

Theorem C.2. *Suppose the constraints are sortable. Consider Algorithm C.1 with step-size parameters η and γ and maximum average absolute deviation ϵ . Let δ^* denote the true optimum of the base problem, and let $\delta^\dagger$ denote the value of the objective returned by Algorithm C.1. Then, there exists a constant c_0 such that Algorithm C.1 runs in time at most*

$$c_0 \cdot (n \cdot \log(n) + m/\gamma^2) \cdot \eta^{-m}.$$

Moreover, there exists a constant c_1 and a problem-specific constant $V(\epsilon)$ (i.e., depending on $\hat{R}_i$, ℓ_i , u_i , and ϵ) such that

$$|\delta^* - \delta^\dagger| \leq \frac{c_1 \cdot (\eta + m\gamma)}{V(\epsilon)^2}.$$

Here η denotes the step size of a grid search over average risks of the unobserved strata, and

γ denotes the step size of a discrete grid of values of ϵ in the maximization routine. We take $\eta = 0.05$ p.p. and $\gamma = 0.01$ p.p. in Section 4.2. We give the definition of $V(\epsilon)$ in (C.30) in Appendix C.2.2. Roughly, $V(\epsilon)$ is related to the within-group variance of the estimated risks V ,

$$V = \frac{1}{n} \sum_{j=1}^m n_j \cdot \text{VAR} \left((\hat{R}_i)_{i \in \mathcal{G}_j} \right).$$

In particular, $V(\epsilon) \geq V - 4\epsilon$. Moreover, $V(\epsilon)$ is positive whenever the base problem has a finite solution. In our principal application, it takes a few seconds on standard hardware for the algorithm to run for 100 values of ϵ with $m = 3$ race groups and $n > 10^6$ observations with negligible approximation error.

C.2.1.3 Comparison with other methods

As we note above, our approach to sensitivity analysis for risk-adjusted regression is related to sensitivity analysis methods developed in the causal inference literature, particularly to recent work on L_2 sensitivity bounds (Huang and Pimentel, 2024; Zhang and Zhao, 2022). Our approach, however, is distinct in important ways that make it uniquely advantageous in our setting. First, methods for sensitivity analysis in the causal inference literature that can be adapted to our setting often involve a large number of parameters—making them difficult to credibly calibrate—or require parametric assumptions on the form of the confounding itself (e.g., Jung et al., 2020; Rosenbaum, 2002; Rosenbaum and Rubin, 1983a). In contrast, the single parameter in our sensitivity analysis— ϵ , the average absolute deviation between the true and estimated risks—is straightforward to reason about and explain, particularly to policymakers or other non-technical stakeholders.

The popular sensitivity analysis developed by Cinelli and Hazlett (2020), which makes use of two sensitivity parameters, readily conforms to the linear estimation strategy in (4.3). In our setting, it gauges how much $\hat{\beta}_{j^*} - \hat{\beta}_1$ changes between the following two regression specifications:

$$\sum_{j=1}^m \hat{\beta}_j \cdot \mathbf{1}(C_i = j) + \hat{\beta}_{\hat{R}} \cdot \hat{R}_i \quad \text{and} \quad \sum_{j=1}^m \hat{\beta}_j \cdot \mathbf{1}(C_i = j) + \hat{\beta}_{\hat{R}} \cdot \hat{R}_i + \hat{\beta}_R \cdot R_i.$$

However, the sensitivity analysis’s goal is to compare the former specification to the “true” regression specification in (4.3), $\sum_{j=1}^m \hat{\beta}_j \cdot \mathbf{1}(C_i = j) + \hat{\beta}_R \cdot R_i$. This can be effected when the coefficient $\hat{\beta}_{\hat{R}}$ is zero in the latter specification, i.e., when decisions (A) are independent of estimated risk ($\hat{R}$)

Algorithm C.1 (Sensitivity analysis)

Input: The estimated risk vector $\hat{\mathbf{R}}$, the lower and upper bounds ℓ and $\mathbf{u}$, the average absolute difference ϵ , and the step-size parameters η and γ .

Output: The minimum and maximum values of $\hat{\beta}_j - \hat{\beta}_1$ for all groups $j = 2, \dots, m$.

- 1: Set $\rho_j \leftarrow \frac{1}{|\mathcal{S}_{j,1}|} \cdot \sum_{i \in \mathcal{S}_{j,1}} \hat{R}_i$ for $j = 1, \dots, m$
 - 2: Generate a grid of $(\tau_1, \dots, \tau_m)$ with step size η
 - 3: **repeat**
 - 4: Calculate $\mathbf{R}_{\min}^*$ minimizing the simplified problem with $\rho_1, \dots, \rho_m$ and $\tau_1, \dots, \tau_m$
 - 5: Calculate $\mathbf{R}_{\max}^*$ maximizing the simplified problem to within $2m\gamma$ with $\rho_1, \dots, \rho_m$ and $\tau_1, \dots, \tau_m$
 - 6: Calculate $\hat{\beta}_j - \hat{\beta}_1$ using both $\mathbf{R}_{\min}^*$ and $\mathbf{R}_{\max}^*$
 - 7: **until** The grid of $(\tau_1, \dots, \tau_m)$ is exhausted
 - 8: **return** The largest and smallest differences $\hat{\beta}_j - \hat{\beta}_1$ observed for all $j = 2, \dots, m$
-

given true risk (R) and group membership (C).⁵ This assumption is implausible, though, when the influence of various factors on decision makers deviates from the factors’ actual connection to risk, typically the key concern in studies of disparate impact.⁶

Other existing sensitivity methods that do involve a single parameter would, if suitably adapted to our setting, be parameterized in terms of an L_∞ -bound Γ on the log odds ratio between true and estimated risks (e.g., Dorn and Guo, 2022; Zhao et al., 2019). However, in realistic models of confounding, the change in the odds ratio typically cannot be bounded in this way for all units: some—potentially very small—set of units will have an estimated risk of approximately 0 but a true risk of approximately 1 or *vice versa*. As a result, one may need to set Γ to be very large to strictly satisfy the assumptions of the approach, yielding sensitivity bounds that are, in practice, quite conservative. More recent methods involving L_2 bounds avoid this problem (Huang and Pimentel, 2024; Zhang and Zhao, 2022), but would not yield tight bounds on the estimand of interest in our setting.

⁵Strictly speaking, it is required that the residuals of $\hat{\mathbf{R}}$ and $\mathbf{A}$, when regressed against $\mathbf{R}$ and $\mathbf{C}$, must be uncorrelated, a mathematically weaker condition that the independence assumption implies.

⁶As a simple example, suppose officers over-weight certain factors like “furtive movements” that are relatively uninformative about risk, as discussed in the main text. In this case, individuals who made furtive movements would be over-represented in the frisked population, relative to their risk. Consequently, in the more limited data available to the analyst, furtive movements would likely be associated with a *lower* chance of possessing a weapon, all else equal. Thus, conditional on true risk R and group C , individuals with lower predicted risks $\hat{R}$ would have greater odds of having engaged in furtive movements. Since officers over-weight furtive movements when making frisk decisions, lower predicted risk would consequently be associated with a higher chance of having been frisked, even conditional on R and C .

C.2.2 Theory

In what follows, we formally justify the sensitivity analysis laid out in Appendix C.2.1, deriving algorithms to efficiently solve the optimization problem it gives rise to and rigorously verifying their correctness and run-times. This discussion is structured in four parts. In the first part, by introducing a new low-dimensional family of parameters, we show how to reduce the base optimization problem of (C.6) into the simplified problem of (C.7). In the second and third parts, we show how to efficiently find the minimum and maximum values of the simplified optimization problem, respectively. Finally, in the fourth part, we prove bounds on the quality of the approximation.

Remark C.3. In Appendix C.2.1, for expositional clarity, we parameterized our optimization problems in terms of *average* risks. However, it is mathematically more natural to work with the *total* risks. Therefore, throughout this appendix we let ρ_j and τ_j denote the *total* risk in each stratum, i.e., we rewrite the relevant constraints of the base problem and the simplified problem as follows:

$$\sum_{i \in \mathcal{S}_{j,1}} R_i = \rho_j, \quad \sum_{i \in \mathcal{S}_{j,0}} R_i = \tau_j.$$

C.2.2.1 Simplifying the objective

To simplify the optimization problem in (C.6), we can avail ourselves of the fact that OLS has a closed-form solution. We augment the notation in Appendix C.2.1 above as follows:

$$\mathcal{G}_j = \mathcal{S}_{j,0} \cup \mathcal{S}_{j,1} = \{i : C_i = j\},$$

i.e., $\mathcal{G}_j$ denotes the set of indices belonging to the j -th group.

Lemma C.4. Let $n_j = |\mathcal{G}_j| > 0$ denote the number of individuals in the j -th group; let σ_j denote the search rate in group j , i.e., $\sigma_j = \frac{1}{n_j} \sum_{i \in \mathcal{G}_j} A_i$; and let r_j and t_j denote

$$r_j = \sum_{i \in \mathcal{S}_{j,1}} R_i, \quad t_j = \sum_{i \in \mathcal{S}_{j,0}} R_i,$$

i.e., the total risk of the searched and unsearched individuals belonging to the j -th group, respectively.

If R_i is constant within each of the m groups, then the OLS estimate need not exist due to

collinearity. Otherwise, the OLS estimate of the coefficients $\hat{\beta}$ in (4.3) satisfies

$$\hat{\beta} = \begin{bmatrix} \sigma_1 \\ \vdots \\ \sigma_m \\ 0 \end{bmatrix} + \frac{\frac{1}{n} \left[\sum_{j=1}^m \sigma_j \cdot t_j - (1 - \sigma_j) \cdot r_j \right]}{\frac{1}{n} \left[\sum_{i=1}^n R_i^2 - \sum_{j=1}^m \frac{(r_j + t_j)^2}{n_j} \right]} \begin{bmatrix} \frac{r_1 + t_1}{n_1} \\ \vdots \\ \frac{r_m + t_m}{n_m} \\ -1 \end{bmatrix}. \quad (\text{C.8})$$

Proof. Recall that our design matrix and outcome variables take the form

$$\mathbf{X} = \begin{bmatrix} \mathbf{1}(C_1 = 1) & \dots & \mathbf{1}(C_1 = m) & R_1 \\ \vdots & \ddots & \vdots & \vdots \\ \mathbf{1}(C_n = 1) & \dots & \mathbf{1}(C_n = m) & R_n \end{bmatrix}, \quad \mathbf{Y} = \begin{bmatrix} A_1 \\ \vdots \\ A_n \end{bmatrix}.$$

The OLS estimate of the coefficients $\hat{\beta}$ is then simply $(\mathbf{X}^\top \mathbf{X})^{-1} \mathbf{X}^\top \mathbf{Y}$.

Note that by definition, we have that

$$\mathbf{X}^\top \mathbf{X} = \begin{bmatrix} n_1 & & \sum_{i \in \mathcal{G}_1} R_i \\ & \ddots & \vdots \\ & & n_m & \sum_{i \in \mathcal{G}_m} R_i \\ \sum_{i \in \mathcal{G}_1} R_i & \dots & \sum_{i \in \mathcal{G}_m} R_i & \sum_{i=1}^n R_i^2 \end{bmatrix} = \begin{bmatrix} n_1 & & r_1 + t_1 \\ & \ddots & \vdots \\ & & n_m & r_m + t_m \\ r_1 + t_1 & \dots & r_m + t_m & \sum_{i=1}^n R_i^2 \end{bmatrix}.$$

Similarly,

$$\mathbf{X}^\top \mathbf{Y} = \begin{bmatrix} \sum_{i \in \mathcal{G}_1} A_i \\ \vdots \\ \sum_{i \in \mathcal{G}_m} A_i \\ \sum_{i=1}^n R_i \cdot A_i \end{bmatrix} = \begin{bmatrix} n_1 \cdot \sigma_1 \\ \vdots \\ n_m \cdot \sigma_m \\ \sum_{j=1}^m r_j \end{bmatrix}.$$

Note that $\mathbf{X}^\top \mathbf{X}$ has the form

$$\begin{bmatrix} D & v^\top \\ v & \sum_{i=1}^n R_i^2 \end{bmatrix},$$

where D is diagonal. Then, we see that since D is invertible, it follows that $\mathbf{X}^\top \mathbf{X}$ is invertible if the

Schur complement of the D is invertible, i.e., non-zero; that is, if

$$\sum_{i=1}^n R_i^2 - \sum_{j=1}^m \frac{(r_j + t_j)^2}{n_j} \neq 0.$$

Recalling the definitions of r_j and t_j , we see that the left-hand expression equals

$$\sum_{j=1}^m \sum_{i \in \mathcal{G}_j} \left(R_i - \frac{r_j + t_j}{n_j} \right)^2 = \sum_{j=1}^m n_j \cdot \text{VAR}((R_i)_{i \in \mathcal{G}_j}).$$

Now, the variance of a set is only zero if all the elements of the set are equal, i.e., if for all $i \in \mathcal{G}_j$, R_i equals the average risk in that group, $\frac{r_j + t_j}{n_j}$.

Assume this is not the case, i.e., that $\mathbf{X}^\top \mathbf{X}$ is invertible. Since $\mathbf{X}^\top \mathbf{X}$ is an arrowhead matrix, we can invert it using the Sherman-Morrison formula. In particular, we obtain that

$$(\mathbf{X}^\top \mathbf{X})^{-1} = \begin{bmatrix} \frac{1}{n_1} & & \\ & \ddots & \\ & & \frac{1}{n_m} \\ & & & 0 \end{bmatrix} + \frac{1}{\sum_{i=1}^n R_i^2 - \sum_{j=1}^m \frac{(r_j + t_j)^2}{n_j}} \begin{bmatrix} \frac{r_1 + t_1}{n_1} \\ \vdots \\ \frac{r_m + t_m}{n_m} \\ -1 \end{bmatrix} \begin{bmatrix} \frac{r_1 + t_1}{n_1} \\ \vdots \\ \frac{r_m + t_m}{n_m} \\ -1 \end{bmatrix}^\top. \quad (\text{C.9})$$

Combining this with the expression derived for $\mathbf{X}^\top \mathbf{Y}$ above and dividing the numerator and denominator of the scalar term by n yields that

$$\hat{\boldsymbol{\beta}} = \begin{bmatrix} \sigma_1 \\ \vdots \\ \sigma_m \\ 0 \end{bmatrix} + \frac{\frac{1}{n} \left[\sum_{j=1}^m \sigma_j \cdot t_j - (1 - \sigma_j) \cdot r_j \right]}{\frac{1}{n} \left[\sum_{i=1}^n R_i^2 - \sum_{j=1}^m \frac{(r_j + t_j)^2}{n_j} \right]} \begin{bmatrix} \frac{r_1 + t_1}{n_1} \\ \vdots \\ \frac{r_m + t_m}{n_m} \\ -1 \end{bmatrix}. \quad (\text{C.10})$$

□

We note that Lemma C.4 provides the following useful condition on when the base problem has a meaningful solution:

Corollary C.5. *The maximum and minimum values of the base problem are both finite if and only if it is not the case that $\ell_i \leq \rho_j \leq u_i$ for all $i \in \mathcal{G}_j$ and $j = 1, \dots, m$, and in addition*

$$\frac{1}{n} \sum_{j=1}^m \sum_{i \in \mathcal{G}_j} \left| \hat{R}_i - \frac{\rho_j}{n_{j,1}} \right| \leq \epsilon.$$

Proof. The proof follows immediately from noting that since the averages of the risks on the observed strata $\mathcal{S}_{j,1}$ are fixed at ρ_j , the risks for the whole j -th group can only all be equal if the risks of the observed *and* the unobserved individuals are uniformly equal to ρ_j . $\square$

The following corollary completes the simplification of the optimization problem in (C.6) in the case where the average true risk in the unobserved strata is also assumed known (“**the parameterized problem**” of Appendix C.2.1),⁷ i.e., of

$$\begin{aligned}
& \underset{\mathbf{R} \in \mathbb{R}^n}{\text{Optimize}} && \hat{\beta}_{j^*} - \hat{\beta}_1 \\
& \text{s.t.} && \frac{1}{n} \sum_{i=1}^n |R_i - \hat{R}_i| \leq \epsilon, \\
& && \sum_{i \in \mathcal{S}_{j,1}} R_i = \rho_j, \quad (j = 1, \dots, m) \\
& && \sum_{i \in \mathcal{S}_{j,0}} R_i = \tau_j, \quad (j = 1, \dots, m) \\
& && R_i \leq u_i, \quad (i = 1, \dots, n) \\
& && R_i \geq \ell_i, \quad (i = 1, \dots, n)
\end{aligned} \tag{C.11}$$

Corollary C.6. *The optimal solutions of the parameterized problem in (C.11), if they are both well-defined, are also the optimal solutions of the simplified problem in (C.7).*

Proof. Let $\hat{\beta}$ be as in (4.3), and r_j and t_j as in Lemma C.4.⁸ Then, it follows directly from (C.8) that we have that for all $\mathbf{R}$,

$$\hat{\beta}_j - \hat{\beta}_1 = \sigma_j - \sigma_1 + \frac{\frac{1}{n} \left[\sum_{j=1}^m \sigma_j \cdot t_j - (1 - \sigma_j) \cdot r_j \right]}{\frac{1}{n} \left[\sum_{i=1}^n R_i^2 - \sum_{j=1}^m \frac{(r_j + t_j)^2}{n_j} \right]} \left(\frac{r_j + t_j}{n_j} - \frac{r_1 + t_1}{n_1} \right).$$

Since the constraints of the parameterized problem fix $r_j = \rho_j$ and $t_j = \tau_j$ for all feasible $\mathbf{R}$, the objective function of the parameterized problem has the form

$$a + \frac{b}{x - c}, \tag{C.12}$$

⁷See Remark C.3 above for the slight difference in notation between here and Appendix C.2.1.

⁸In particular, strictly speaking, $\hat{\beta}$, r_j , and t_j should be written $\hat{\beta}(\mathbf{R})$, $r_j(\mathbf{R})$, and $t_j(\mathbf{R})$, since they depend on $\mathbf{R}$.

where

$$\begin{aligned}
a &= \sigma_j - \sigma_1, \\
b &= \frac{1}{n} \left[\sum_{j=1}^m \sigma_j \cdot \tau_j - (1 - \sigma_j) \cdot \rho_j \right] \cdot \left(\frac{\rho_j + \tau_j}{n_j} - \frac{\rho_1 + \tau_1}{n_1} \right), \\
c &= \frac{1}{n} \sum_{j=1}^m \frac{(\rho_j + \tau_j)^2}{n_j}, \\
x &= \frac{1}{n} \sum_{i=1}^n R_i^2.
\end{aligned}$$

If $b = 0$, then the objective does not depend on $\mathbf{R}$, and so there is nothing to prove. Otherwise, by continuity and the fact that the feasible region is convex, and hence connected, and compact, the range of the objective in (C.7) over the feasible region is a closed interval I . As shown in the proof of Lemma C.4 and Corollary C.5, $x \geq c$, and $x > c$ if the optima are both well-defined. Consequently, we can assume the interval is strictly positive, i.e., $I \subseteq (c, \infty)$. The derivative of (C.12) exists everywhere on I and is, moreover, non-zero and continuous, and so $a + \frac{b}{x-c}$ is strictly monotone on I . Therefore, the maximum and minimum must occur at the endpoints, i.e., at the maximum and minimum values of x . $\square$

We note that since b is not guaranteed to be positive, the maximum of the simplified problem in (C.7) does not necessarily correspond to the maximum of the base problem in (C.6) and *vice versa*.

C.2.2.2 Minimization

In the case of minimization, the simplified problem in (C.7) defines a convex optimization problem, which, as noted in Appendix C.2.1 above, can consequently be solved efficiently using standard interior-point methods. However, these results can be improved in a problem-specific way to yield a simpler, and, in practice, more computationally tractable method of finding exact solutions. Moreover, they parallel the solution method for maximization, which is not a convex problem, and hence to which standard interior point methods cannot be applied.

Optimizing over a single stratum We begin by solving a simpler version of the problem, where we restrict to a single stratum. In the case of minimization—as in the case of maximization,

as we show below—solutions to the minimization problem must be in a certain normal form. While searching over all risk vectors is prohibitively difficult, searching over normal form risk vectors can be done straightforwardly in linear time.

Definition C.7 (Minimization normal form). We say that a risk vector $\mathbf{R}$ is in *minimization normal form* if there exist thresholds $t_{a,j}^{\text{lwr}}$ and $t_{a,j}^{\text{upr}}$ such that for $i \in \mathcal{S}_{j,a}$,

$$R_i = \begin{cases} \min(u_i, t_{a,j}^{\text{lwr}}) & \hat{R}_i \leq t_{a,j}^{\text{lwr}}, \\ \hat{R}_i & t_{a,j}^{\text{lwr}} \leq \hat{R}_i \leq t_{a,j}^{\text{upr}}, \\ \max(\ell_i, t_{a,j}^{\text{upr}}) & t_{a,j}^{\text{upr}} \leq \hat{R}_i, \end{cases}$$

and, in addition, $t_{a,j}^{\text{upr}} - t_{a,j}^{\text{lwr}}$ is equal to some Δ for all $a = 0, 1$ and $j = 1, \dots, m$.

In the context of a single stratum, we refer to t^{upr} and t^{lwr} for notational simplicity. We also assume that $\mathcal{S}_{j,a} = \{1, \dots, n\}$ for the same reason.

Minimization normal form results from “pushing” $\hat{R}_i$ up to t^{lwr} or down to t^{upr} as far as possible. An example of a risk vector in minimization normal form is shown in Figure C.1. As the name suggests, risk vectors minimizing the objective must be in minimization normal form.

Lemma C.8. *Consider the optimization problem (“the single stratum minimization problem”)*

$$\begin{aligned} \text{Minimize}_{\mathbf{R} \in \mathbb{R}^n} \quad & \frac{1}{n} \sum_{i=1}^n R_i^2 \\ \text{s.t.} \quad & \frac{1}{n} \sum_{i=1}^n |R_i - \hat{R}_i| \leq \epsilon, \\ & \sum_{i=1}^n R_i = \mu, \\ & R_i \leq u_i, \ (\forall i) \\ & R_i \geq \ell_i. \ (\forall i) \end{aligned} \tag{C.13}$$

Suppose that $\mathbf{R}^$ is a minimizer. Then, $\mathbf{R}^*$ is unique and in minimization normal form. Moreover, $\mathbf{R}^*$ exhausts the L_1 budget, in the sense that either $\frac{1}{n} \|\mathbf{R}^* - \hat{\mathbf{R}}\|_1 = \epsilon$ or $t^{\text{lwr}} = t^{\text{upr}}$.*

Proof. The proof proceeds by examining the first-order KKT conditions to derive weak necessary conditions that solutions must satisfy. These conditions are then strengthened to minimization normal form by directly comparing the objective function at different points satisfying the weak conditions. Finally, we show that the minimum is the point satisfying the “strengthened” conditions which is most distant from $\hat{\mathbf{R}}$, i.e., that exhausts the budget.

The problem is simplified by rewriting the first constraint as a collection of linear (and hence everywhere differentiable) constraints. In particular, $\frac{1}{n}\|\mathbf{R} - \hat{\mathbf{R}}\|_1 \leq \epsilon$ is equivalent to the 2^n constraints of the form $\mathbf{S}^\top(\mathbf{R} - \hat{\mathbf{R}}) \leq \epsilon \cdot n$ for all $\mathbf{S} \in \{-1, 1\}^n$. We assume throughout, without loss of generality, that $\ell_i < u_i$.

Weak conditions Let $\mathbf{e}_i$ denote the i -th standard basis vector and $\mathbf{1} = \sum_{i=1}^n \mathbf{e}_i$. Note that:

- The gradient of the objective is $\frac{2}{n} \cdot \mathbf{R}$;
- The gradient of $\mathbf{S}_k^\top(\mathbf{R} - \hat{\mathbf{R}}) - \frac{\epsilon}{n}$, where $\{\mathbf{S}_1, \dots, \mathbf{S}_{2^n}\} = \{-1, 1\}^n$, is $\mathbf{S}_k$;
- The gradient of $\sum_{i=1}^n R_i$ is $\mathbf{1}$;
- The gradient of $R_i - u_i$ is $\mathbf{e}_i$;
- The gradient of $\ell_i - R_i$ is $-\mathbf{e}_i$.

The first-order necessary KKT conditions therefore require that⁹

$$2 \cdot \mathbf{R}^* - \lambda \cdot \mathbf{1} + \sum_{i=1}^n \mu_{0,i} \cdot \mathbf{e}_i - \sum_{i=1}^n \mu_{1,i} \cdot \mathbf{e}_i + \sum_{k=1}^{2^n} \nu_k \cdot \mathbf{S}_k = 0 \quad (\text{C.14})$$

for some arbitrary λ , and for non-negative $\mu_{0,i}$, $\mu_{1,i}$, and ν_k for $k = 1, \dots, 2^n$ satisfying complementary slackness, i.e., such that $\mu_{0,i} \cdot (R_i^* - u_i) = 0$, $\mu_{1,i} \cdot (\ell_i - R_i^*) = 0$, and $\nu_k \cdot (\mathbf{S}_k^\top(\mathbf{R}^* - \hat{\mathbf{R}}) - n \cdot \epsilon) = 0$ for all $i = 1, \dots, n$ and $k = 1, \dots, 2^n$.

These consequences allow us to prove the following weak characterization of $\mathbf{R}^*$. Let $\Delta = \sum_{k=1}^{2^n} \nu_k$. Then, for all i ,

$$R_i^* \in \left\{ \hat{R}_i, u_i, \ell_i, \frac{1}{2} \cdot [\lambda - \Delta], \frac{1}{2} \cdot [\lambda + \Delta] \right\}. \quad (\text{C.15})$$

For, from (C.14), we get that

$$2 \cdot R_i^* - \lambda + \mu_{0,i} - \mu_{1,i} + \sum_{k=1}^{2^n} \nu_k \cdot S_{k,i} = 0, \quad (\text{C.16})$$

where $S_{k,i}$ refers to the i -th component of $\mathbf{S}_k$. If $R_i^* = \hat{R}_i$, there is nothing to prove. Therefore, assume that $R_i^* \neq \hat{R}_i$. We note that by complementary slackness, $\nu_k > 0$ only if $S_{k,i} \cdot (R_i^* - \hat{R}_i) \geq 0$

⁹For notational simplicity, we have multiplied through by n and absorbed constants into the corresponding Lagrange multipliers.

for all $i = 1, \dots, n$. For, supposing without loss of generality that $R_i^* > \hat{R}_i$, if $S_{k,i} \cdot (R_i^* - \hat{R}_i) < 0$, then, since $\mathbf{S}_k + 2\mathbf{e}_i \in \{-1, 1\}^n$, we would have that

$$(\mathbf{S}_k + 2\mathbf{e}_i)^\top (\mathbf{R}^* - \hat{\mathbf{R}}) = \epsilon \cdot n + 2(R_i^* - \hat{R}_i) > \epsilon \cdot n,$$

which violates the constraints. Hence $S_{k,i} = S_{k',i}$ for all k and k' such that $\nu_k, \nu_{k'} > 0$. Therefore, our expression simplifies to

$$2 \cdot R_i^* - \lambda + \mu_{0,i} - \mu_{1,i} + \Delta \cdot s = 0,$$

where $s = \pm 1$. By complementary slackness, if $\mu_{0,i} > 0$, then $R_i^* = u_i$; if $\mu_{1,i} > 0$, then $R_i^* = \ell_i$. Therefore, we need only consider the case where $\mu_{0,i} = \mu_{1,i} = 0$, i.e., where

$$2 \cdot R_i^* - \lambda + s \cdot \Delta = 0,$$

whence $R_i^* = \frac{1}{2}[\lambda + s \cdot \Delta]$. Therefore, for all i , (C.15) holds.

Strengthening conditions Next, we strengthen (C.15) to minimization normal form. First, we show that if $R_{i_0}^* > \hat{R}_{i_0}$ and $R_{i_0}^* > R_{i_1}^*$, then $R_{i_1}^* = u_{i_1}$. For, suppose not. Then, there exists some $\delta > 0$ such that (1) $R_{i_0}^* - \delta > \hat{R}_{i_0}$, (2) $R_{i_1}^* + \delta < u_{i_1}$, and (3) $R_{i_0}^* - \delta > R_{i_1}^*$. Define $\mathbf{R}' = \mathbf{R}^* + \delta \cdot (\mathbf{e}_{i_1} - \mathbf{e}_{i_0})$. Then, we note that

$$\begin{aligned} \|\mathbf{R}' - \hat{\mathbf{R}}\|_1 - \|\mathbf{R}^* - \hat{\mathbf{R}}\|_1 &= \sum_{i=1}^n |R'_i - \hat{R}_i| - |R_i^* - \hat{R}_i| \\ &= |R'_{i_0} - \hat{R}_{i_0}| - |R_{i_0}^* - \hat{R}_{i_0}| + |R'_{i_1} - \hat{R}_{i_1}| - |R_{i_1}^* - \hat{R}_{i_1}| \\ &= (R_{i_0}^* - \delta - \hat{R}_{i_0}) - (R_{i_0}^* - \hat{R}_{i_0}) + |R_{i_1}^* + \delta - \hat{R}_{i_1}| - |R_{i_1}^* - \hat{R}_{i_1}| \\ &\leq -\delta + |R_{i_1}^* - \hat{R}_{i_1}| + \delta - |R_{i_1}^* - \hat{R}_{i_1}| \\ &= 0, \end{aligned}$$

so $\mathbf{R}'$ satisfies the L_1 -distance constraint. Here, we have used the fact that both $R_{i_0}^*$ and $R_{i_0}^* - \delta$ are greater than $\hat{R}_{i_0}$ in the third equality, and the triangle inequality in the inequality. The remaining constraints also hold by construction, so $\mathbf{R}'$ is feasible.

Since $\mathbf{R}'$ is feasible, we will arrive at a contradiction if we can show that it achieves a smaller

objective. However,

$$\begin{aligned}
\sum_{i=1}^n (R_i^*)^2 - (R_i')^2 &= (R_{i_0}^*)^2 - ((R_{i_0}^*)^2 + \delta^2 - 2 \cdot R_{i_0}^* \cdot \delta) + (R_{i_1}^*)^2 \\
&\quad - ((R_{i_1}^*)^2 + \delta^2 + 2 \cdot (R_{i_1}^*) \cdot \delta) + \sum_{i \neq i_0, i_1} (R_i^*)^2 - (R_i^*)^2 \\
&= 2 \cdot \delta \cdot (R_{i_0}^* - R_{i_1}^* - \delta) \\
&> 0,
\end{aligned}$$

where the inequality follows from the fact that $R_{i_0}^* - \delta > R_{i_1}^*$. Therefore $\mathbf{R}^*$ is not a minimum, contrary to hypothesis. In exactly the same manner, we see that if $R_{i_0}^* < \hat{R}_i$ and $R_{i_1}^* < R_{i_0}^*$, then $R_{i_1}^* = \ell_{i_1}$.

To apply the claim, let s^{lwr} be the maximum R_i^* for i such that $R_i^* > \hat{R}_i$, and similarly for s^{upr} . Then, for all i such that $R_i^* < s^{\text{lwr}}$, $R_i^* = u_i$; similarly, where $R_i^* > s^{\text{upr}}$, $R_i^* = \ell_i$. By (C.15), we must have that $s^{\text{lwr}} \leq s^{\text{upr}}$, since otherwise there would be some i such that $R_i^* = u_i$ and $R_i^* = \ell_i$, which is impossible. Therefore, in particular, by construction, $R_i^* = \hat{R}_i$ whenever $s^{\text{lwr}} < R_i^* < s^{\text{upr}}$. Thus, taking $t^{\text{lwr}} = s^{\text{lwr}}$ and $t^{\text{upr}} = s^{\text{upr}}$ gives the desired normal form.

Uniqueness and budget exhaustion If the budget is not exhausted, then, by complementary slackness, $\nu_k = 0$ for all $k = 1, \dots, 2^n$, whence $\Delta = 0$. This immediately implies that $t^{\text{lwr}} = t^{\text{upr}} = \frac{1}{2}\lambda$.

To see uniqueness, suppose that $\mathbf{R}^{(0)}$ and $\mathbf{R}^{(1)}$ are two distinct minima, both in normal form, with corresponding thresholds t_k^{lwr} and t_k^{upr} for $k = 0, 1$. We note that since $\sum_{i=1}^n R_i^{(0)} = \sum_{i=1}^n R_i^{(1)} = \mu$, we must have, without loss of generality, that $t_0^{\text{lwr}} \leq t_1^{\text{lwr}}$ and $t_0^{\text{upr}} \geq t_1^{\text{upr}}$, with at least one of the two inequalities strict. However, it follows that $\|\mathbf{R}^{(0)} - \hat{\mathbf{R}}\|_1 < \|\mathbf{R}^{(1)} - \hat{\mathbf{R}}\|_1$, which is impossible by the preceding paragraph. $\square$

Lemma C.8 suggests a natural algorithm for solving the optimization problem in (C.7): we can sweep over all possible risk vectors in minimization normal form simply by sweeping over t^{lwr} . In particular, as we increase t^{lwr} , the sum constraint—i.e., that $\sum_{i=1}^n R_i = \mu$ —forces t^{upr} to decrease so as to counterbalance it exactly.

More precisely, consider an index i to be “active” if one of the thresholds is between ℓ_i and u_i . If k^{lwr} represents the number of indices that are active because of t^{lwr} —and k^{upr} is also defined accordingly—then, locally, if we increase t^{lwr} at unit rate, the sum of the risks increases at the rate

of k^{lwr} . Therefore, the rates of increase of the two thresholds, r^{lwr} and r^{upr} , must satisfy¹⁰

$$r^{\text{lwr}} \cdot k^{\text{lwr}} = r^{\text{upr}} \cdot k^{\text{upr}}. \quad (\text{C.17})$$

In particular, we know that we have reached the minimum—and hence solved the optimization problem—once either the budget has been exhausted or the two thresholds have become equal. For fixed r^{lwr} and r^{upr} , this would be straightforward to calculate; however, these rates can, in principle, change whenever k^{upr} and k^{lwr} change. We call these change points “distinguished points”—that is, points at which a threshold t equals ℓ_i , $\hat{R}_i$, or u_i for some i . Therefore, the algorithm consists of repeating the following steps:

1. Determining the “next” distinguished point that either the lower or upper threshold will reach.
2. Checking whether the budget will be exhausted before that point is reached.
3. If not, advancing to that point and recalculating k^{lwr} or k^{upr} as appropriate.

In addition, since $\sum_{i=1}^n \hat{R}_i$ may not equal μ , a “preprocessing” step, where t^{lwr} is increased or t^{upr} is decreased may be needed to find the “initial” risk vector in minimization normal form satisfying the sum constraint.

We make two observations about the informal algorithm sketch given above needed for its extension to the case of multiple strata. First, we note that every risk vector $\mathbf{R}^*$ in minimization normal form is the minimizer for the corresponding single stratum minimization problem with $\epsilon = \frac{1}{n} \|\mathbf{R}^* - \hat{\mathbf{R}}\|_1$ and $\mu = \sum_{i=1}^n R_i^*$. Consequently, although intended to find the minimizer for a specific ϵ , the algorithm sketched above actually sweeps over the minimizers for *all possible* ϵ . As a result, it is not any harder to solve a single instance of the single stratum minimization problem than it is to solve it generally for all possible ϵ , given some fixed μ . Secondly, while the natural way to parameterize the risk vectors in the discussion above is in terms of t^{lwr} , when optimizing over multiple strata, it is actually more natural to reparameterize the risk vectors in terms of the gap $\Delta = t^{\text{upr}} - t^{\text{lwr}}$. If in addition to (C.17) we impose the condition that $r^{\text{lwr}} + r^{\text{upr}} = 1$, so that Δ

¹⁰If either k^{lwr} or k^{upr} is zero, then the corresponding threshold can be changed at any rate without affecting the sum constraint. (The rate of change for the other threshold will consequently be zero.) However, this is because changing that threshold in this case does not actually change the risk vector, since no indices are active, and so we can ignore this case in the subsequent discussion.

decreases at a constant rate, then we have that

$$r^{\text{lwr}} = \frac{k^{\text{upr}}}{k^{\text{lwr}} + k^{\text{upr}}}, \quad r^{\text{upr}} = \frac{k^{\text{lwr}}}{k^{\text{lwr}} + k^{\text{upr}}}. \quad (\text{C.18})$$

To solve the single stratum minimization problem across the whole range of possible ϵ , let $\mathbf{R}_\Delta$ be the unique solution we can associate with a given gap Δ between t^{lwr} and t^{upr} .¹¹ Then, we can study two functions:

$$\epsilon(\Delta) = \|\mathbf{R}_\Delta - \hat{\mathbf{R}}\|_1, \quad \Sigma(\Delta) = \sum_{i=1}^n R_{\Delta,i}^2 - \hat{R}_i^2.$$

By the preceding discussion we readily derive that $\epsilon(\Delta)$ is a piecewise linear function of Δ , whose slope at a given value of Δ is equal to¹²

$$K = 2 \cdot \frac{k^{\text{lwr}} \cdot k^{\text{upr}}}{k^{\text{lwr}} + k^{\text{upr}}},$$

and $\Sigma(\Delta)$ is a piecewise quadratic function of Δ satisfying

$$\begin{aligned} \Sigma(\Delta + t) &= \Sigma(\Delta) + k^{\text{lwr}} \cdot [(t^{\text{lwr}} + r^{\text{lwr}} \cdot t)^2 - (t^{\text{lwr}})^2] + k^{\text{upr}} \cdot [(t^{\text{upr}} - r^{\text{upr}} \cdot t)^2 - (t^{\text{upr}})^2] \\ &= \Sigma(\Delta) + k^{\text{lwr}} \cdot r^{\text{lwr}} \cdot t \cdot (r^{\text{lwr}} \cdot t + 2t^{\text{lwr}}) + k^{\text{upr}} \cdot r^{\text{upr}} \cdot t \cdot (r^{\text{upr}} \cdot t - 2t^{\text{upr}}) \\ &= \Sigma(\Delta) - 2 \cdot \frac{k^{\text{lwr}} \cdot k^{\text{upr}}}{k^{\text{lwr}} + k^{\text{upr}}} \cdot (t^{\text{upr}} - t^{\text{lwr}}) \cdot t + \frac{k^{\text{lwr}}(k^{\text{upr}})^2 + k^{\text{upr}}(k^{\text{lwr}})^2}{(k^{\text{lwr}} + k^{\text{upr}})^2} \cdot t^2 \\ &= \Sigma(\Delta) - K\Delta t + \frac{K}{2}t^2, \end{aligned}$$

as long as t is sufficiently small that the number of active indices does not change. It follows that $\epsilon(\Delta)$ and $\Sigma(\Delta)$ are fully determined by the following collections:

$$\mathbf{\Delta} = (\Delta_1, \dots, \Delta_N), \quad 1 = \Delta_1 > \dots > \Delta_N = 0,$$

where Δ_k represents the k -th value of Δ at which k^{lwr} and k^{upr} change; and

$$\mathbf{K} = (K_1, \dots, K_N), \quad K_1 = 0, \quad K_N = 0,$$

where K_k denotes the value of K beginning at $\Delta = \Delta_k$. From $\mathbf{\Delta}$ and $\mathbf{K}$, the Δ^* such that

¹¹For every Δ there is a unique $\mathbf{R}_\Delta$. In particular, any other distinct risk vector in minimization normal form with the same gap must be strictly greater or lesser, as argued at the end of Lemma C.8, meaning that it cannot satisfy the sum constraint. However, the reverse is not necessarily true—that is, $\mathbf{R}_\Delta$ may equal $\mathbf{R}_{\Delta'}$ for $\Delta \neq \Delta'$ —since achieving a given gap may require that $k^{\text{lwr}} = k^{\text{upr}} = 0$, in which case different thresholds can be chosen without altering the underlying risk vector.

¹²Or zero if both k^{lwr} and k^{upr} are zero.

$\|\mathbf{R}_{\Delta^*} - \hat{\mathbf{R}}\|_1 = \epsilon$ can be calculated in linear time. For completeness, this algorithm is given in Algorithm C.2.

Thus, a complete solution for the optimization problem for a single stratum across *all possible* ϵ requires only calculating Δ and $\mathbf{K}$. As described above, this can be calculated by moving the thresholds toward each other at the prescribed rates r^{lwr} and r^{upr} , updating the rates each time an index is activated or deactivated, until the gap between the thresholds is zero. Similar to evaluating $\epsilon(\Delta)$ and $\Sigma(\Delta)$, this can be completed in linear time, as described in Algorithm C.4. (As above, we note that if $\frac{1}{n} \sum_{i=1}^n \hat{R}_i \neq \mu$, it may be necessary to do a preprocessing step, also in linear time, as shown in Algorithm C.3.)

Optimizing over all strata With a complete solution to the problem of minimizing the sum of squares for a single stratum, the problem of minimizing the sum of squares across all strata is straightforward. We begin by characterizing solutions to the minimization problem in the general case.

Lemma C.9. *Consider the simplified minimization problem in (C.7). Suppose that $\mathbf{R}^*$ is a solution. Then, the restriction of $\mathbf{R}^*$ to any stratum is in minimization normal form. Moreover, $\mathbf{R}^*$ exhausts the L_1 budget, in that either $\frac{1}{n} \|\mathbf{R}^* - \hat{\mathbf{R}}\|_1 = \epsilon$ or $t^{\text{lwr}} = t^{\text{upr}} = \mu_{a,j}$, where $\mu_{a,j}$ equals either ρ_j or τ_j depending on whether a equals 1 or 0.*

Proof. The proof is virtually identical to the proof of Lemma C.8. The only difference is that the first-order necessary KKT conditions take the form

$$2 \cdot \mathbf{R}^* - \left(\sum_{j=1}^m \lambda_{0,j} \cdot \mathbf{1}_{0,j} + \lambda_{1,j} \cdot \mathbf{1}_{1,j} \right) + \sum_{i=1}^n \mu_{0,i} \cdot \mathbf{e}_i - \sum_{i=1}^n \mu_{1,i} \cdot \mathbf{e}_i + \sum_{k=1}^{2^n} \nu_k \cdot \mathbf{S}_k = 0, \quad (\text{C.19})$$

where $\mathbf{1}_{a,j} = \sum_{i \in \mathcal{S}_{j,a}} \mathbf{e}_i$. Restricting to a single i gives the following minor variant of (C.16):

$$2 \cdot R_i^* - \lambda_{a,j} + \mu_{0,i} - \mu_{1,i} + \sum_{k=1}^n \nu_k \cdot S_{k,i} = 0.$$

The proof then proceeds identically; the only difference is to note that while $\lambda_{a,j}$ varies by stratum, $\Delta = \sum_{k=1}^{2^n} \nu_k$ does not. $\square$

It follows that solving the minimization problem across all strata can be carried out in almost




Figure C.2: *Maximization normal form.* Black dots represent estimated risks $\hat{R}_i$, the error bars represent the allowable range $[\ell_i, u_i]$, and red dots indicate the corresponding value of the normal form vector R_i . The pivots are indicated on the x-axis.

the same way as across a single stratum:

1. Construct piecewise linear functions $\epsilon_{a,j}(\Delta)$ and piecewise quadratic functions $\Sigma_{a,j}(\Delta)$ for each stratum;
2. Note that because the sums $\sum_{a=0}^1 \sum_{j=1}^m \epsilon_{a,j}(\Delta)$ and $\sum_{a=0}^1 \sum_{j=1}^m \Sigma_{a,j}(\Delta)$ are also piecewise linear and quadratic, respectively, they can also be evaluated using Algorithm C.2;
3. Find Δ^* such that $\sum_{a=0}^1 \sum_{j=1}^m \epsilon_{a,j}(\Delta^*) = \epsilon$, and evaluate $\sum_{a=0}^1 \sum_{j=1}^m \Sigma_{a,j}(\Delta^*)$.

Consequently, beyond the machinery established in Algorithms C.2, C.3, and C.4, we only need a way of calculating the sums of these piecewise functions. This can be carried out straightforwardly using a variation on the standard merge-sort algorithm, which we give in Algorithm C.5 for completeness.

Putting this all together, we obtain the following lemma summarizing the results of this section.

Lemma C.10. *If ℓ , $\hat{\mathbf{R}}$, and $\mathbf{u}$ have been sorted, then there exists an $O(\log(m) \cdot n)$ algorithm solving the simplified minimization problem ((C.7)).*

Proof. The proof is straightforward. We begin by noting that Algorithm C.2 requires linear time: the index i increases by one in the while loop on lines 6 through 10 on each iteration, and must be less than the length of the input. Similarly, in Algorithm C.3, the while loop on lines 11 through 22 increments i^{lwr} by one on each loop, and limits it to three times the length of the input. The same

analysis applies to Algorithm C.4, in which either i^{upr} or i^{lwr} increases by one on each iteration of the loop on lines 14 through 33. We note that the run time of Algorithm C.5 is proportional to the sum of the lengths of its inputs. Using a standard divide-and-conquer approach, we can accomplish the combination of the $2m$ strata in m total calls, where each element of the input appears in $1 + \lceil \log_2(m) \rceil$ of the function calls, for a total time complexity of $O(\log(m) \cdot n)$. Putting this all together, the total complexity is $O(\log(m) \cdot n)$. $\square$

C.2.2.3 Maximization

Even in the simplified problem in (C.7), maximization is not a convex problem. In general, maximizing a quadratic objective over a linearly constrained convex set is NP-hard (Sahni, 1974). However, the restricted forms of the objective and constraints allow us efficiently to solve the optimization problem exactly in the case of a single stratum and approximately in the case of multiple strata.

Optimizing over a single stratum As in the case of minimization, risk vectors that maximize the single-stratum maximization problem have a special normal form.

Definition C.11 (Maximization normal form). We say that a risk vector $\mathbf{R}$ is in *maximization normal form* if there exist indices (“pivots”) i^{lwr} and i^{upr} such that

$$R_i = \begin{cases} \ell_i & i < i^{\text{lwr}}, \\ \hat{R}_i & i^{\text{lwr}} < i < i^{\text{upr}}, \\ u_i & i^{\text{upr}} < i. \end{cases}$$

We additionally require that the risk values at the pivots themselves satisfy $R_{i^{\text{lwr}}} \in [\ell_{i^{\text{lwr}}}, \hat{R}_{i^{\text{lwr}}}]$ and $R_{i^{\text{upr}}} \in [\hat{R}_{i^{\text{upr}}}, u_{i^{\text{upr}}}]$.^{13,14}

Maximization normal form is similar to minimization normal form, except that instead of being

¹³Since it involves the ordering of the indices as well as the risk values, the appropriate generalization of maximization normal form is notationally awkward to express when multiple strata are involved. Fortunately, in contrast to the case of minimization, we will not have occasion to use such a generalization, and so do not give its definition.

¹⁴We note that we allow i^{lwr} and i^{upr} to take the values 0 or $n + 1$ in addition to $\{1, \dots, n\}$ to cover cases where, e.g., $R_i = \ell_i$ for all $i = 1, \dots, n$.

pushed toward the thresholds, values at indices *below* the lower pivot are pushed down, and values at indices *above* the upper pivot are pushed up. An illustration of maximum normal form is given in Figure C.2.

An important difference from the case of minimization is that studying the KKT conditions does not yield the normal form directly; instead, it yields a weaker characterization.

Definition C.12 (Weak normal form). We say that a risk vector is in *weak normal form* if there exist thresholds t^{lwr} and t^{upr} such that

$$R_i = \begin{cases} \ell_i & R_i < t^{\text{lwr}}, \\ t^{\text{lwr}} & R_i = t^{\text{lwr}}, \\ \hat{R}_i & t^{\text{lwr}} < R_i < t^{\text{upr}}, \\ t^{\text{upr}} & R_i = t^{\text{upr}}, \\ u_i & t^{\text{upr}} < R_i, \end{cases}$$

and there exists at most one lower index i^{lwr} such that $R_{i^{\text{lwr}}} = t^{\text{lwr}}$; and similarly, there exists at most one upper index i^{upr} such that $R_{i^{\text{upr}}} = t^{\text{upr}}$

We note that, in the case of *weak* maximization normal form, the indices i^{lwr} and i^{upr} are analogous to the pivots in maximization normal form; however, they do not share the key property, which is that all of the indices $i < i^{\text{lwr}}$ index risks that have been pushed down all the way to ℓ_i and all of the indices $i > i^{\text{upr}}$ index risks that have been pushed all the way up to u_i ; see Figure C.3.

Lemma C.13. *Consider the single-stratum maximization problem—i.e., maximization in (C.13). Suppose that $\mathbf{R}^*$ is a solution. Then, $\mathbf{R}^*$ is in weak normal form. Moreover, either $\frac{1}{n}\|\mathbf{R}^* - \hat{\mathbf{R}}\|_1 = \epsilon$ or R_i^* equals ℓ_i or u_i for all but (at most) one i .*

Proof. The proof is quite similar to the proof of Lemma C.8. We note that the first-order KKT conditions are almost the same, namely that

$$2 \cdot \mathbf{R}^* - \lambda \cdot \mathbf{1} - \sum_{i=1}^n \mu_{0,i} \cdot \mathbf{e}_i + \sum_{i=1}^n \mu_{1,i} \cdot \mathbf{e}_i - \sum_{k=1}^{2^n} \nu_k \cdot \mathbf{S}_k = 0 \quad (\text{C.20})$$

for some λ and $\mu_{0,i}, \mu_{1,i}, \nu_k \geq 0$ satisfying complementary slackness. Virtually exactly as before, it




Figure C.3: *Weak maximization normal form.* Black, blue, and grey dots represent estimated risks $\hat{R}_i$, the error bars represent the allowable range $[\ell_i, u_i]$, and red dots indicate the corresponding value of the normal form vector R_i . The pivots are indicated on the x-axis. The risk R_i at indices where the estimated risk is between the thresholds and the bounds straddle the thresholds (shown in blue) can take on a number of values (e.g., $R_i = \hat{R}_i$ or $R_i = u_i$). Likewise, i^{lwr} and i^{upr} can occur at any index i such that $R_i = t^{\text{lwr}}$ or $R_i = t^{\text{upr}}$ is feasible (shown in grey). As a result, the set of risk vectors in weak maximization normal form is a very complex search space.

follows from this that (C.15) holds, i.e., that

$$R_i^* \in \left\{ \hat{R}_i, u_i, \ell_i, \frac{1}{2} \cdot [\lambda - \Delta], \frac{1}{2} \cdot [\lambda + \Delta] \right\}. \quad (\text{C.21})$$

Again, as in the case of minimization, we can strengthen (C.21). In particular, exactly the same argument used to strengthen the weak conditions in Lemma C.8 yields that if $R_{i_0}^* > \hat{R}_{i_0}$ and $R_{i_0}^* < R_{i_1}^*$, then $R_{i_1}^* = u_{i_1}$; and, conversely, if $R_{i_0}^* < \hat{R}_{i_0}$ and $R_{i_0}^* > R_{i_1}^*$, then $R_{i_1}^* = \ell_{i_1}$. The proof is then otherwise the same as in Lemma C.8.

To see that there can be at most one i such that $\ell_i < R_i^* < \hat{R}_i$, suppose that $\ell_{i_0} < R_{i_0}^* < \hat{R}_{i_0}$ and $\ell_{i_1} < R_{i_1}^* < \hat{R}_{i_1}$ for $i_0 \neq i_1$. By (C.21), $R_{i_0}^* = R_{i_1}^*$. Choose $\delta > 0$ sufficiently small that $R_{i_0}^* - \delta > \ell_{i_0}$ and $R_{i_1}^* + \delta < \hat{R}_{i_1}$. Then, as in the case of minimization, it follows immediately that $\mathbf{R}^* + \delta(\mathbf{e}_{i_1} - \mathbf{e}_{i_0})$ is feasible, but increases the objective by $2\delta^2$, contrary to the hypothesis that $\mathbf{R}^*$ was a maximizer. In the same way, it follows that there exists at most one i' such that $u_{i'} > R_{i'}^* > \hat{R}_{i'}$.

Combining this with our previous necessary condition on $\mathbf{R}^*$ and letting $t^{\text{lwr}} = \frac{1}{2} \cdot [\lambda - \Delta]$ and $t^{\text{upr}} = \frac{1}{2} \cdot [\lambda + \Delta]$ yields that for all i such that $R_i^* < t^{\text{lwr}}$, $R_i^* = \ell_i$; that for all i such that $R_i^* > t^{\text{upr}}$, $R_i^* = u_i$; that there is at most one index such that $\hat{R}_i > R_i^* > \ell_i$; and, finally, that there is at most

one i' such that $\hat{R}_{i'} < R_{i'}^* < u_{i'}$. □

Lemma C.13 is, by itself, insufficient to form the basis of an effective solution algorithm since it does not determine which indices are i^{lwr} and i^{upr} . Even if these indices were known, they do not pin down the value of R_i for $i \neq i^{\text{lwr}}, i^{\text{upr}}$, which could, almost without restriction, be ℓ_i , $\hat{R}_i$, or u_i . As a result, the search space of vectors in weak maximization normal form is extremely large.

Therefore, to construct an effective maximization algorithm, we must transform vectors merely in weak maximization normal form into vectors that are in full-fledged maximization normal form.

To do so, we will have need of the following elementary fact.

Lemma C.14. *Suppose $a \geq b$ and $c \geq d$. Then*

$$|a - c| + |b - d| \leq |a - d| + |b - c|.$$

Proof. The proof is not complicated, but is simplest to understand if expressed in geometric terms. Consider the points $\mathbf{p}_0 = (a, b)$ and $\mathbf{p}_1 = (c, d)$. Then both $\mathbf{p}_0$ and $\mathbf{p}_1$ lie in the lower half-plane

$$H = \{(x, y) : x \geq y\}.$$

Let $T : (x, y) \mapsto (y, x)$ be the linear transformation given by reflection across the line $y = x$. Then, the claim of the lemma can be reframed as follows: for any $\mathbf{p}_0, \mathbf{p}_1 \in H$,

$$\|\mathbf{p}_0 - \mathbf{p}_1\|_1 \leq \|\mathbf{p}_0 - T(\mathbf{p}_1)\|_1. \tag{C.22}$$

This setup is shown in Figure C.4a. We divide into four cases according to whether $a \geq c$ and $b \geq d$, or, equivalently, depending on in which of the four regions shown in Figure C.4b the point $\mathbf{p}_1$ lies.

Case (A): $a \geq c$ and $b \geq d$ In this case, we have that

$$\begin{aligned} |a - c| + |b - d| &= a - c + b - d \\ &= a - d + b - c \\ &\leq |(a - d) + (b - c)| \\ &\leq |a - d| + |b - c|. \end{aligned}$$

Therefore, (C.22) holds.

Case (B): $a \leq c$ and $b \geq d$ In this case,

$$\begin{aligned}
|a - c| + |b - d| &= c - a + b - d \\
&\leq c - a + b - d + 2(a - b) \\
&= c - b + a - d \\
&\leq |(c - b) + (a - d)| \\
&\leq |c - b| + |a - d|,
\end{aligned}$$

so (C.22) holds in this case as well.

Before completing the proof, we note that T has the useful property that

$$\|T(\mathbf{p})\|_1 = \|\mathbf{p}\|_1. \quad (\text{C.23})$$

Case (C): $a \leq c$ and $b \leq d$ We note that this case is the same as Case (A) with the roles of $\mathbf{p}_0$ and $\mathbf{p}_1$ reversed. Consequently, by (C.22), we have that

$$\|\mathbf{p}_0 - \mathbf{p}_1\|_1 \leq \|T(\mathbf{p}_0) - \mathbf{p}_1\|_1.$$

Applying (C.23) and using the fact that T^2 is the identity, we have that

$$\|T(\mathbf{p}_0) - \mathbf{p}_1\|_1 = \|T(T(\mathbf{p}_0) - \mathbf{p}_1)\|_1 = \|\mathbf{p}_0 - T(\mathbf{p}_1)\|_1,$$

which yields (C.22).

Case (D): $a \geq c$ and $b \leq d$ We argue as in the previous case, noting that this case is equivalent to Case (B) with the roles of $\mathbf{p}_0$ and $\mathbf{p}_1$ reversed, whence

$$\|\mathbf{p}_0 - \mathbf{p}_1\|_1 \leq \|T(\mathbf{p}_0) - \mathbf{p}_1\|_1 = \|\mathbf{p}_0 - T(\mathbf{p}_1)\|_1.$$

□




Figure C.4: An illustration of the proof of Lemma C.14. The half-plane H is shown in gray, with the dotted lines representing the coordinate axes. Left: The claim of the lemma is equivalent to the claim that the path joining $\mathbf{p}_0$ to $\mathbf{p}_1$ is no longer than the path joining $\mathbf{p}_0$ to $T(\mathbf{p}_1)$. Right: The division into four cases, depending on where $\mathbf{p}_0$ lies relative to $\mathbf{p}_1$.

A further difference from minimization is that the additional assumption of sortability, Definition C.1, is required to connect maximizers to maximization normal form. If sortability fails, there may be intervals $[\ell_i, u_i]$ that can be non-trivially nested for distinct i , which can cause i_0 and i_1 to “jump between” different indices a potentially exponential number of times in the search for a maximizer. For notational convenience, we assume throughout that sortable constraints are, in fact, already sorted.

The first step to deriving an effective maximization algorithm is to remove “trivial” deviations from maximization normal form.

Definition C.15 (Trivial deviations from normality). We say that a risk vector in weak normal form has a *trivial deviation from normality* relative to its *bounds* if there are $i_0 < i_1$ such that $\ell_{i_0} = \ell_{i_1}$, $\ell_{i_0} < R_{i_0}$, and $R_{i_1} = \ell_{i_1}$; or there are $i'_1 > i'_0$ such that $u_{i'_0} = u_{i'_1}$, $R_{i'_1} < u_{i'_1}$, and $R_{i'_0} = u_{i'_0}$.

We say that a risk vector in weak normal form has a *trivial deviation from normality* relative to its *estimates* if there are $i_0 < i_1$ such that $\hat{R}_{i_0} = \hat{R}_{i_1}$ but $R_{i_0} > R_{i_1}$.

An illustration of trivial deviations from normality is shown in Figure C.5. As their name suggests, trivial deviations from normality are straightforward to eliminate.

Lemma C.16. *Suppose the constraints are sortable. If a feasible risk vector $\mathbf{R}$ is in weak normal*




Figure C.5: *An illustration of trivial deviations from normality. As in Figures C.1 and C.2, the black dots represent estimated risks $\hat{R}_i$, the error bars represent the allowable range $[\ell_i, u_i]$, and the red dots represent the corresponding values of R_i , with indices increasing from left to right on the x -axis. There is a trivial deviation from normality relative to the lower bound on the left because $\ell_{i_0} = \ell_{i_1}$ but $R_{i_1} = \ell_{i_1}$, and a similar illustration of a trivial deviation from normality relative to the upper bound in the center. On the right there is a trivial deviation from the estimates.*

form with trivial deviations from normality, then there exists a feasible $\mathbf{R}'$ in weak normal form with no trivial deviations from normality such that

$$\frac{1}{n} \sum_{i=1}^n R_i^2 = \frac{1}{n} \sum_{i=1}^n (R'_i)^2.$$

Proof. We first define an algorithm for removing trivial deviations from normality, and then show that it must eventually terminate when all trivial deviations have been removed.

Algorithm If there were any indices i_0 and i_1 representing a trivial deviation from normality, then we could eliminate them by swapping R_{i_0} and R_{i_1} , i.e., by considering

$$R'_i = R_i \quad (i \neq i_0, i_1), \quad R'_{i_0} = R_{i_1}, \quad R'_{i_1} = R_{i_0}.$$

In particular, it follows immediately that $\mathbf{R}'$ is still in weak normal form and that $\frac{1}{n} \sum_{i=1}^n R_i^2 = \frac{1}{n} \sum_{i=1}^n (R'_i)^2$. Because the constraints sorted, the only feasibility constraint on $\mathbf{R}'$ that is not satisfied trivially is that $\|\mathbf{R}' - \hat{\mathbf{R}}\|_1 \leq \epsilon$. However, this constraint follows by applying Lemma C.14 with $a = R_{i_0}, b = R_{i_1}, c = \hat{R}_{i_1}$, and $d = \hat{R}_{i_0}$, which yields

$$\|\mathbf{R}' - \hat{\mathbf{R}}\|_1 \leq \|\mathbf{R} - \hat{\mathbf{R}}\|_1 \leq \epsilon.$$

Termination We call a pair of indices (i_0, i_1) *out of order* if $i_0 < i_1$, but $R_{i_0} > R_{i_1}$. We note that any trivial deviation from normality requires that the corresponding indices i_0 and i_1 are out of order. Then, we observe the following facts:

Fact 1. There are only finitely many pairs of out-of-order indices.

Fact 2. As long as there is a trivial deviation of any kind, it is possible to swap indices so as to remove that trivial deviation.

Fact 3. Any swap undertaken as part of the algorithm strictly reduces the number of out-of-order pairs of indices.

We note that it follows immediately from these three facts that a greedy algorithm removing trivial deviations in any order whatsoever must eventually terminate.

Fact 1 is immediate, and Fact 2 follows from the discussion of the algorithm above. To see that Fact 3 holds, it suffices to show that *any* swap of out-of-order pairs of indices strictly reduces the number of out-of-order pairs of indices.

To this end, suppose $i_0 < i_1$ are out of order. As before, we use $\mathbf{R}'$ to denote the risk vector *after* swapping indices i_0 and i_1 . First, suppose there is an index i^* such that $i_0 < i^* < i_1$. Then, if (i_0, i^*) is out of order *after* swapping i_0 and i_1 , it must have been out of order before swapping, since

$$R_{i_0} > R_{i_1} = R'_{i_0} > R'_{i^*} = R_{i^*}.$$

Likewise, if (i^*, i_1) is out of order *after* swapping i_0 and i_1 , it too must have been out of order before swapping, since

$$R_{i^*} = R'_{i^*} > R'_{i_1} = R_{i_0} > R_{i_1}.$$

Therefore, swapping cannot increase the number of out-of-order pairs involving $i_0 < i^* < i_1$.

On the other hand, since $R'_{i_0} < R'_{i_1}$, if $i^* < i_0$, then there are three possibilities:

Case 1. After swapping, neither (i^*, i_0) nor (i^*, i_1) is out of order.

Case 2. After swapping, both (i^*, i_0) and (i^*, i_1) are out of order.

Case 3. After swapping, (i^*, i_0) is out of order, but (i^*, i_1) is not.

In the first case, the number of out-of-order pairs cannot have increased, so we only need to consider the second and third cases. In Case 2, we must have that

$$R_{i^*} = R'_{i^*} > R'_{i_1} = R_{i_0} > R_{i_1},$$

so both (i^*, i_0) and (i^*, i_1) were out of order before swapping. In Case 3, we must have that

$$R_{i_0} = R'_{i_1} \geq R'_{i^*} = R_{i^*} > R'_{i_0} = R_{i_1},$$

so that before swapping (i^*, i_1) was out of order, but (i^*, i_0) was not. Therefore, it follows that swapping again cannot increase the number of out-of-order pairs involving $i^* < i_0$. A similar argument shows that the number of out-of-order pairs involving $i_1 < i^*$ does not increase after swapping. Therefore swapping strictly reduces the number of out-of-order pairs, since it at least eliminates the pair (i_0, i_1) , which is Fact 3. $\square$

With these preliminaries in hand, we are prepared to solve the maximization problem.

Lemma C.17. *Suppose the constraints are sortable. Then the single stratum maximization problem in (C.13) is maximized by a risk vector in maximization normal form and that either exhausts the budget or has $i^{\text{lwr}} = i^{\text{upr}}$.*

Proof. We adopt a strategy similar to that used in the “strengthening” portion of the proof of Lemma C.8. Namely, we show that any point satisfying the necessary conditions of Lemma C.13 but without pivots can be “improved” to a feasible point that achieves a greater objective.

Let $\mathbf{R}^*$ be a maximizer. By Lemma C.13, $\mathbf{R}^*$ is in weak maximization normal form; by Lemma C.16, we may assume that $\mathbf{R}^*$ has no trivial deviations from normality.

To prove the theorem, we first observe that it suffices to prove the following two claims:

1. There do not exist $i_0 < i_1$ such that $R_{i_0}^* > \ell_{i_0}$ and $R_{i_1}^* < \hat{R}_{i_1}$,
2. There do not exist $i_0 < i_1$ such that $R_{i_1}^* < u_{i_1}$ and $R_{i_0}^* > \hat{R}_{i_0}$.

First, we see how the theorem follows from the two claims. Define i^{lwr} and i^{upr} as follows:

$$i^{\text{lwr}} = \min\{i : R_i^* > \ell_i\}, \quad i^{\text{upr}} = \max\{i : R_i^* < u_i\},$$

then, for all $i < i^{\text{lwr}}$, $R_i^* = \ell_i$, and for all $i > i^{\text{upr}}$, $R_i^* = u_i$. Moreover, for all i satisfying $i^{\text{lwr}} < i < i^{\text{upr}}$, by Claim 1, $R_i^* \geq \hat{R}_i$; and, by Claim 2, $R_i^* \leq \hat{R}_i$. Consequently, $R_i^* = \hat{R}_i$ for these indices. Moreover, if $i^{\text{lwr}} \neq i^{\text{upr}}$, by Claim 2, $R_{i^{\text{lwr}}}^* \leq \hat{R}_{i^{\text{lwr}}}$, i.e., $R_{i^{\text{lwr}}}^* \in [\ell_{i^{\text{lwr}}}, \hat{R}_{i^{\text{lwr}}}]$. Similarly, by Claim 1, we conclude that $R_{i^{\text{upr}}}^* \in [\hat{R}_{i^{\text{upr}}}, u_{i^{\text{upr}}}]$.

Therefore it suffices to prove the two claims. Since the proofs are virtually identical, we focus on Claim 1.

The key observation is the following. Suppose there exist indices $R_{i_0}^*$ and $R_{i_1}^*$ with $R_{i_1}^* \geq R_{i_0}^*$, and suppose that there is some δ such that after decreasing $R_{i_0}^*$ and increasing $R_{i_1}^*$ by δ the result is still feasible. Then $\mathbf{R}^*$ cannot be maximal, since, letting $\mathbf{R}' = \mathbf{R}^* + \delta(\mathbf{e}_{i_1} - \mathbf{e}_{i_0})$, we have that

$$\begin{aligned} \left[\sum_{i=1}^n (R'_i)^2 \right] - \left[\sum_{i=1}^n (R_i^*)^2 \right] &= (R_{i_0}^* - \delta)^2 + (R_{i_1}^* + \delta)^2 - (R_{i_0}^*)^2 - (R_{i_1}^*)^2 \\ &= 2\delta(\delta + [R_{i_1}^* - R_{i_0}^*]) \\ &\geq 2\delta^2 \\ &> 0. \end{aligned} \tag{C.24}$$

To actually prove the claim, we proceed by contradiction. In particular, suppose that $i_0 < i_1$ represented an exception to Claim 1, so that $R_{i_0}^* > \ell_{i_0}$ and $R_{i_1}^* < \hat{R}_{i_1}$.

If $R_{i_0}^* \leq R_{i_1}^*$, then (C.24) immediately applies, since increasing $R_{i_1}^*$ *decreases* the L_1 distance between $\hat{\mathbf{R}}$ and $\mathbf{R}^*$ and decreasing $R_{i_0}^*$, at worst, *increases* it at the same rate. More formally, for $0 < \delta$ sufficiently small,

$$\begin{aligned} |(R_{i_0}^* - \delta) - \hat{R}_{i_0}| + |(R_{i_1}^* + \delta) - \hat{R}_{i_1}| &= |(R_{i_0}^* - \delta) - \hat{R}_{i_0}| + \hat{R}_{i_1} - (R_{i_1}^* + \delta) \\ &\leq |R_{i_0}^* - \hat{R}_{i_0}| + \delta + \hat{R}_{i_1} - R_{i_1}^* - \delta \\ &= |R_{i_0}^* - \hat{R}_{i_0}| + |R_{i_1}^* - \hat{R}_{i_1}|, \end{aligned}$$

and so $\mathbf{R}'$ is feasible. Because $R_{i_0}^* \leq R_{i_1}^*$, by (C.24), $\mathbf{R}'$ achieves a greater average sum of squares than $\mathbf{R}^*$, contrary to the assumed maximality of $\mathbf{R}^*$.

On the other hand, if $R_{i_0}^* > R_{i_1}^*$, then, by Lemma C.14, the vector $\mathbf{R}^\dagger$ given by switching the indices i_0 and i_1 , i.e.,

$$R_i^\dagger = R_i^* \quad (i \neq i_0, i_1), \quad R_{i_0}^\dagger = R_{i_1}^*, \quad R_{i_1}^\dagger = R_{i_0}^*,$$

satisfies $\|\mathbf{R}^\dagger - \hat{\mathbf{R}}\|_1 \leq \epsilon$. Moreover, applying our sortability hypothesis, we have that

$$u_{i_0} \geq R_{i_0}^* > R_{i_1}^* \geq \ell_{i_1} \geq \ell_{i_0},$$

and

$$u_{i_1} \geq u_{i_0} \geq R_{i_0}^* > R_{i_1}^* \geq \ell_{i_1},$$

whence for all $i = 1, \dots, n$, $\ell_i \leq R_i^\dagger \leq u_i$. Therefore $\mathbf{R}^\dagger$ is feasible. However, since $\mathbf{R}^*$ has no trivial deviations from normality, it follows that

$$R_{i_0}^\dagger = R_{i_1}^* \geq \ell_{i_1} > \ell_{i_0}, \quad R_{i_1}^\dagger = R_{i_0}^* \leq u_{i_0} < u_{i_1}.$$

Since $R_{i_0}^\dagger < R_{i_1}^\dagger$, we would like to apply (C.24) to derive a contradiction. To that end, we require that $\mathbf{R}^\dagger + \delta(\mathbf{e}_{i_1} - \mathbf{e}_{i_0})$ is feasible for sufficiently small $\delta > 0$.

To see that such δ exists, we divide into three cases. Observe that one of the following three conditions must hold: (1) $R_{i_0}^* < \hat{R}_{i_1}$, (2) $R_{i_1}^* > \hat{R}_{i_0}$, or (3) $R_{i_0}^* \geq \hat{R}_{i_1}$ and $R_{i_1}^* \leq \hat{R}_{i_0}$. We treat each case separately.

Case 1 ($R_{i_0}^* < \hat{R}_{i_1}$). Note that if it were the case that

$$\|\mathbf{R}^\dagger + \delta(\mathbf{e}_{i_1} - \mathbf{e}_{i_0}) - \hat{\mathbf{R}}\|_1 \leq \|\mathbf{R}^\dagger - \hat{\mathbf{R}}\|_1, \quad (\text{C.25})$$

then the result follows immediately. Choosing $\delta < \hat{R}_{i_1} - R_{i_0}^*$, we have, using the fact that $R_{i_0}^\dagger = R_{i_1}^*$ and $R_{i_1}^\dagger = R_{i_0}^*$ that

$$|R_{i_0}^\dagger - \delta - \hat{R}_{i_0}| + |R_{i_1}^\dagger + \delta - \hat{R}_{i_1}| = |R_{i_0}^\dagger - \delta - \hat{R}_{i_0}| + \hat{R}_{i_1} - R_{i_1}^\dagger - \delta \geq |R_{i_0}^\dagger - \hat{R}_{i_0}| + \delta + \hat{R}_{i_1} - R_{i_1}^\dagger - \delta.$$

The final expression is $|R_{i_0}^\dagger - \hat{R}_{i_0}| + |\hat{R}_{i_1} - R_{i_1}^\dagger|$, whence (C.25) holds.

Case 2 ($R_{i_1}^* > \hat{R}_{i_0}$). This case is essentially identical to the previous case.

Case 3 ($R_{i_0}^* \geq \hat{R}_{i_1}$ and $R_{i_1}^* \leq \hat{R}_{i_0}$). In this case, we have by sortability that

$$\begin{aligned}
|R_{i_0}^* - \hat{R}_{i_0}| + |R_{i_1}^* - \hat{R}_{i_1}| &= R_{i_0}^* - \hat{R}_{i_0} + \hat{R}_{i_1} - R_{i_1}^* \\
&> (R_{i_0}^* - \hat{R}_{i_0} + \hat{R}_{i_1} - R_{i_1}^*) + 2(\hat{R}_{i_0} - \hat{R}_{i_1}) \\
&= R_{i_0}^* - \hat{R}_{i_1} + \hat{R}_{i_0} - R_{i_1}^* \\
&= |R_{i_0}^\dagger - \hat{R}_{i_0}| + |R_{i_1}^\dagger - \hat{R}_{i_1}|,
\end{aligned}$$

where we have used the fact that, because there are no trivial deviations from normality relative to the estimates, $\hat{R}_{i_1} > \hat{R}_{i_0}$. Therefore $\|\mathbf{R}^\dagger - \hat{\mathbf{R}}\|_1 < \epsilon$ and so we can choose δ sufficiently small that $\|\mathbf{R}^\dagger + \delta(\mathbf{e}_{i_1} - \mathbf{e}_{i_0}) - \hat{\mathbf{R}}\|_1 \leq \epsilon$.

Therefore, no such $i_0 < i_1$ exist, and so $\mathbf{R}^*$ is in maximization normal form. $\square$

Lemma C.17 allows us to efficiently search over the space of risk vectors to find the maximizer. In particular, all that is necessary is to decrease the risk at the lower pivot and increase it at the upper pivot until the L_1 budget is expended. The linear-time algorithm for doing so is given in Algorithm C.7. (As in the case of Algorithm C.4, there is a potential preprocessing step, given in Algorithm C.6.) As we did in the minimization algorithm, we capture the maximization algorithm's value at a variety of L_1 budgets, controlled by some step-size γ —which, we assume, satisfies $k\gamma = \epsilon$ for some $k \in \mathbb{N}$ —rather than just the full budget ϵ . We denote this collection with Σ . The reason for doing this is to allow us to maximize across strata.

Optimizing over all strata Consider the simplified maximization problem in (C.7). Since the objective is separable, the problem would be separable in the different strata, except that the L_1 budget constraint ranges over all of the strata simultaneously. Much as we reduced the base problem to the parameterized problem by introducing additional parameters τ_j , $j = 1, \dots, m$, if the optimal budget allocation to each stratum were known, then we could reduce to the single stratum case, applying Algorithm C.7. In particular, we note that the following optimization problem (hereafter “**the separable problem**”) parameterized by the stratum-specific budgets

$\epsilon_{1,0}, \dots, \epsilon_{m,0}, \epsilon_{1,1}, \dots, \epsilon_{m,1}$ is separable:

$$\begin{aligned}
& \underset{\mathbf{R} \in \mathbb{R}^n}{\text{Optimize}} && \frac{1}{n} \sum_{i=1}^n R_i^2 \\
& \text{s.t.} && \frac{1}{n} \sum_{i \in \mathcal{S}_{j,a}} |R_i - \hat{R}_i| \leq \epsilon_{j,a}, \quad (a = 0, 1; j = 1, \dots, m) \\
& && \sum_{i \in \mathcal{S}_{j,1}} R_i = \rho_j, \quad (j = 1, \dots, m) \\
& && \sum_{i \in \mathcal{S}_{j,0}} R_i = \tau_j, \quad (j = 1, \dots, m) \\
& && R_i \leq u_i, \quad (i = 1, \dots, n) \\
& && R_i \geq \ell_i. \quad (i = 1, \dots, n)
\end{aligned} \tag{C.26}$$

In particular, the separable problem can be solved stratum by stratum, since the constraints and objective are both separable at the stratum level.

Reducing solving the simplified problem to solving the separable problem, however, is complex: There is no obvious way to determine the optimal L_1 budget allocation other than by sweeping over all possible budgets. However, it is nevertheless possible to efficiently discover an *approximately* optimal budget allocation using a divide-and-conquer approach. In particular, as shown in Algorithm C.7, it is only marginally more difficult computationally to output the maximum value of the objective for the entire budget than to output the maximum value of the objective for a range of budgets, i.e., to output Σ , which gives the maximum value of the objective for the budgets $0, \gamma, 2\gamma, \dots, \epsilon$.

Given $\Sigma^{(0)}$ and $\Sigma^{(1)}$ corresponding to two different strata when the allocated budget is $0, \gamma, 2\gamma, \dots, \epsilon$, we can find the approximate maximum value of the objective when the *total* budget allocated to both strata (in increments of γ) is $0, \gamma, 2\gamma, \dots, \epsilon$ by sweeping over the two dimensional grid of points $\{0, \gamma, \dots, \epsilon\} \times \{0, \gamma, \dots, \epsilon\}$, and, for $k\gamma$, outputting

$$\max \left(\Sigma_0^{(0)} + \Sigma_{k+1}^{(1)}, \Sigma_1^{(0)} + \Sigma_k^{(1)}, \dots, \Sigma_{k+1}^{(0)} + \Sigma_0^{(1)} \right).$$

For completeness, we explicitly give this maximization routine in Algorithm C.8.

Lemma C.18. *If $\ell, \hat{\mathbf{R}}$ and $\mathbf{u}$ have been sorted, then the approximate solution to the maximization problem in (C.7) given by applying Algorithms C.7 and C.8 is*

$$O \left(m \cdot \left(\frac{\epsilon}{\gamma} \right)^2 + n \right).$$

Proof. We begin by analyzing Algorithm C.6. We note that the number of iterations of the while

loop on lines 7 through 12 is capped by the length of the input, since i^{upr} is decremented on each input. Likewise, Algorithm C.7 is almost linear in the size of the input: In the while loop on lines 14 through 32, either i^{lwr} is incremented, or the subsequent multiple of γ is reached, meaning that the algorithm as a whole is linear in the size of the input and ϵ/γ . Therefore, running Algorithm C.7 over all $2m$ strata requires $O(n + \frac{m\epsilon}{\gamma})$ time.

All of the strata can be combined using $2m - 1$ applications of Algorithm C.8. Reviewing the for loop on lines 3 through 10 and the while loop on lines 5 through 8, we see that exactly

$$\frac{\frac{\epsilon}{\gamma} \cdot (\frac{\epsilon}{\gamma} + 1)}{2}$$

iterations are performed. Adding this to the previous runtime obtained and simplifying gives the desired expression. $\square$

C.2.2.4 Controlling approximation error

There are two sources of approximation error in our algorithm. The first is the approximation error introduced by solving the maximization problem only approximately. The second is the approximation error introduced by the fact that we cannot sweep over all possible values of $\boldsymbol{\tau} = (\tau_1, \dots, \tau_m)$ when solving the simplified problem; instead, we must sweep over a grid. Characterizing how much error is introduced by each of these approximations is our final task.

We begin with the following simple lemma.

Lemma C.19. *Suppose $\mathbf{R}^{(0)}$ and $\mathbf{R}^{(1)}$ are risk vectors. Then*

$$\left| \sum_{i=1}^n (R_i^{(1)})^2 - \sum_{i=1}^n (R_i^{(0)})^2 \right| \leq 2 \cdot \|\mathbf{R}^{(1)} - \mathbf{R}^{(0)}\|_1. \quad (\text{C.27})$$

Proof. The claim, which is a version of Hölder's inequality, follows straightforwardly from considering

the following difference:

$$\begin{aligned}
\left| \sum_{i=1}^n (R_i^{(1)})^2 - (R_i^{(0)})^2 \right| &\leq \sum_{i=1}^n |(R_i^{(1)})^2 - (R_i^{(0)})^2| \\
&\leq \sum_{i=1}^n (R_i^{(1)} + R_i^{(0)}) \cdot |R_i^{(1)} - R_i^{(0)}| \\
&\leq \sum_{i=1}^n 2 \cdot |R_i^{(1)} - R_i^{(0)}| \\
&= 2 \cdot \|\mathbf{R}^{(1)} - \mathbf{R}^{(0)}\|_1.
\end{aligned}$$

The first inequality follows from the triangle inequality, and the second and third from the fact that $0 \leq R_i^{(0)}, R_i^{(1)} \leq 1$. $\square$

Using Lemma C.19, we obtain the following characterization of the potential error in our maximization algorithm.

Lemma C.20. *The difference between the true maximum of the simplified problem—i.e., (C.7)—and the quantity obtained from applying Algorithms C.7 and C.8 is at most $2m\gamma$.*

Proof. Let $\mathbf{R}^*$ be the true maximizing solution. We will show that Algorithm C.8 computes the value of the objective at a “close” point at which we can apply Lemma C.19.

For each stratum $\mathcal{S}_{j,a}$, let $\epsilon_{j,a}$ denote

$$\frac{1}{n} \sum_{i \in \mathcal{S}_{j,a}} |R_i^* - \hat{R}_i|.$$

Then, there exist $\tilde{\epsilon}_{j,a} \in \{0, \gamma, 2\gamma, \dots, \epsilon\}$ such that $\tilde{\epsilon}_{j,a} < \epsilon_{j,a}$ and $\epsilon_{j,a} - \tilde{\epsilon}_{j,a}$ is less than γ .

Let

$$\Delta = \sum_{a=0}^1 \sum_{j=1}^m \epsilon_{j,a} - \tilde{\epsilon}_{j,a} = \epsilon - k\gamma$$

for some k . Since ϵ is a multiple of γ , it follows that $\Delta = k'\gamma$ for some $k \in \mathbb{N}$, where $k' < 2m$, the number of strata.

Set $\epsilon'_{j,a} = \tilde{\epsilon}_{j,a}$ except for k' arbitrarily chosen strata, where we set $\epsilon'_{j,a} = \tilde{\epsilon}_{j,a} + \gamma$ instead. Then, there exists $\mathbf{R}$ that is in (maximization) normal form such that

$$\frac{1}{n} \sum_{i=1}^n |R_i - \hat{R}_i| = \epsilon, \quad \frac{1}{n} \sum_{i \in \mathcal{S}_{j,a}} |R_i - \hat{R}_i| = \epsilon'_{j,a}.$$

The remainder of the proof follows simply by comparing $\mathbf{R}$ and $\mathbf{R}^*$ stratum by stratum. In particular, consider the restrictions of $\mathbf{R}$ and $\mathbf{R}^*$ to a single stratum $\mathcal{S}_{j,a}$. By (C.7), we have that

$$\sum_{i \in \mathcal{S}_{j,a}} R_i = \sum_{i \in \mathcal{S}_{j,a}} R_i^*,$$

and, moreover, by Lemma C.17, both $\mathbf{R}_{j,a}$ and $\mathbf{R}_{j,a}^*$ can be assumed to be in maximization normal form. Assume, without loss of generality, that $\epsilon'_{j,a} \leq \epsilon_{j,a}$. (The proof is virtually identical if the inequality is reversed.)

Let i_0 and i_1 denote the pivots of $\mathbf{R}$, and i_0^* and i_1^* of $\mathbf{R}^*$, so that $i_0 \leq i_0^* \leq i_1^* \leq i_1$. (To avoid unnecessary notational complication, we write as if all n individuals belong to the stratum $\mathcal{S}_{j,a}$ and the constraints were sorted.) Then, we note that

$$R_i \geq R_i^* \quad (i_0 \leq i \leq i_0^*), \quad R_i \leq R_i^* \quad (i_1^* \leq i \leq i_1), \quad R_i = R_i^* \quad (\text{otherwise}). \quad (\text{C.28})$$

It follows that

$$\sum_{i=1}^n (R_i^*)^2 - R_i^2 = \left[\sum_{i=i_1^*}^{i_1} (R_i^*)^2 - R_i^2 \right] - \left[\sum_{i=i_0}^{i_0^*} R_i^2 - (R_i^*)^2 \right]. \quad (\text{C.29})$$

Both terms in (C.29) are positive by (C.28). Moreover, by Lemma C.19, both terms are less than or equal to

$$2 \cdot \left[\sum_{i=i_1}^{i_1^*} R_i^* - R_i \right] = 2 \cdot \left[\sum_{i=i_0}^{i_0^*} R_i - R_i^* \right] = 2n \cdot \frac{\epsilon_{a,j} - \epsilon'_{a,j}}{2} \leq n\gamma.$$

In particular, their difference must also be less than this quantity, and so, summing across strata, we have that

$$\sum_{i=1}^n (R_i^*)^2 - R_i^2 \leq 2nm\gamma.$$

Dividing through by n gives the result. $\square$

The second element we need is a bound on the sensitivity of the objective of the simplified problem to changes in the parameters. (We recall Remark C.3 on notation above.)

Lemma C.21. *Let $\mathbf{R}^{(0)}$ and $\mathbf{R}^{(1)}$ be solutions to the simplified problem in (C.7) with parameters*

$$\boldsymbol{\tau}^{(0)} = (\tau_1^{(0)}, \dots, \tau_m^{(0)}), \quad \boldsymbol{\tau}^{(1)} = (\tau_1^{(1)}, \dots, \tau_m^{(1)}),$$

respectively. Then the difference between the objective values of the two solutions is at most

$$4 \cdot \|\boldsymbol{\tau}^{(0)} - \boldsymbol{\tau}^{(1)}\|_1.$$

Proof. The strategy is simple: we will transform $\mathbf{R}^{(0)}$ into a feasible solution of the simplified problem with parameters $\boldsymbol{\tau}^{(1)}$ and *vice versa* for $\mathbf{R}^{(1)}$ and $\boldsymbol{\tau}^{(0)}$. These new solutions lower or upper bound—depending on whether we are considering minimization or maximization—the objective of the new problem; however, by Lemma C.19, we can also bound their difference in objective value from the original solutions, giving a bound on the difference in objective value of the original solutions.

In particular, we see that the lemma immediately follows from Lemma C.19 and if there exist $\mathbf{R}^{(2)}$ and $\mathbf{R}^{(3)}$ such that

$$\|\mathbf{R}^{(0)} - \mathbf{R}^{(2)}\|_1, \|\mathbf{R}^{(1)} - \mathbf{R}^{(3)}\|_1 \leq 2 \cdot \|\boldsymbol{\tau}^{(0)} - \boldsymbol{\tau}^{(1)}\|_1$$

and such that $\mathbf{R}^{(2)}$ and $\mathbf{R}^{(3)}$ are feasible for the simplified problem with parameters $\boldsymbol{\tau}^{(1)}$ and $\boldsymbol{\tau}^{(0)}$, respectively. For, in that case, we have that, assuming without loss of generality that we are solving the minimization problem,

$$\sum_{i=1}^n (R_i^{(0)})^2 \leq \sum_{i=1}^n (R_i^{(3)})^2 \leq \sum_{i=1}^n (R_i^{(1)})^2 + 4 \cdot \|\boldsymbol{\tau}^{(0)} - \boldsymbol{\tau}^{(1)}\|_1,$$

and similarly

$$\sum_{i=1}^n (R_i^{(1)})^2 \leq \sum_{i=1}^n (R_i^{(2)})^2 \leq \sum_{i=1}^n (R_i^{(0)})^2 + 4 \cdot \|\boldsymbol{\tau}^{(0)} - \boldsymbol{\tau}^{(1)}\|_1.$$

Therefore, it suffices to construct $\mathbf{R}^{(2)}$ and $\mathbf{R}^{(3)}$. The construction involves two steps:

1. Ensure that the sum of the risk vector is correct within each stratum;
2. Ensure that the L_1 budget constraint is satisfied.

The L_1 distance needed to achieve each of these steps is bounded by $\|\boldsymbol{\tau}^{(0)} - \boldsymbol{\tau}^{(1)}\|_1$, and so $\mathbf{R}^{(2)}$ and $\mathbf{R}^{(3)}$ will have the required property.

Assume again that we are solving the minimization problem. Consider the j -th unobserved stratum (i.e., $i \in \mathcal{S}_{j,1}$) with associated upper and lower thresholds t^{lwr} and t^{upr} and suppose without loss of generality that $\tau_j^{(0)} \leq \tau_j^{(1)}$. Then, to achieve the first step, we simply raise t^{upr} until the sum of the risk vector is $\tau_j^{(1)}$ or $t^{\text{upr}} = 1$. If $t^{\text{upr}} = 1$, then we raise t^{lwr} until the sum of the risk vector is $\tau_j^{(1)}$. We note that since our risk vector has strictly increased, the L_1 distance between the original

and new risk vectors is exactly the difference in their sums, i.e., $\tau_j^{(1)} - \tau_j^{(0)}$. Call the risk vector that results from performing this operation across all strata $\mathbf{R}'$. Then, it follows that

$$\|\mathbf{R}' - \hat{\mathbf{R}}\|_1 \leq \|\mathbf{R}^{(0)} - \hat{\mathbf{R}}\|_1 + \tau_j^{(1)} - \tau_j^{(0)} \leq \epsilon + \tau_j^{(1)} - \tau_j^{(0)}.$$

To achieve the second step, in strata $\mathcal{S}_{j,a}$ where $t^{\text{upr}} < 1$ and $t^{\text{lwr}} > 0$, we simply raise t^{upr} and lower t^{lwr} so as to ensure that the sum of the risk vector does not change. Doing so reduces the L_1 distance between the risk vector and $\hat{\mathbf{R}}$, so we do so until the L_1 budget constraint is satisfied. In particular, this process can be carried out until $t^{\text{upr}} = 1$ or $t^{\text{lwr}} = 0$ in every stratum. If the budget constraint is still not satisfied, then the simplified problem for $\boldsymbol{\tau}^{(1)}$ is not feasible, contrary to our assumption that there exists a solution $\mathbf{R}^{(1)}$. Therefore, the process halts, and, moreover, requires moving at most $\|\boldsymbol{\tau}^{(0)} - \boldsymbol{\tau}^{(1)}\|_1$ in L_1 distance. Call the resulting risk vector $\mathbf{R}^{(2)}$. Then, it follows that

$$\sum_{i \in \mathcal{S}_{j,0}} R_i^{(2)} = \tau_j^{(1)} \quad (j = 1, \dots, m), \quad \|\mathbf{R}^{(2)} - \hat{\mathbf{R}}\|_1 \leq \epsilon.$$

The other constraints are satisfied trivially. We can construct $\mathbf{R}^{(3)}$ similarly, and the proof is complete for minimization. For maximization, the argument is exactly similar, except that we increase and decrease R_{i_0} and R_{i_1} instead of t^{lwr} and t^{upr} . $\square$

We can extend the previous results to bound the error in the calculation of the coefficients themselves. (Recall Remark C.3 on notation above.)

Lemma C.22. *Let $V(\epsilon)$ denote the value of*

$$\begin{aligned} & \underset{\mathbf{R} \in \mathbb{R}^n}{\text{Minimize}} \quad \frac{1}{n} \sum_{j=1}^m n_j \cdot \text{VAR}((R_i)_{i \in \mathcal{G}_j}) \\ & \text{s.t.} \quad \frac{1}{n} \sum_{i=1}^n |R_i - \hat{R}_i| \leq \epsilon, \\ & \quad \sum_{i \in \mathcal{S}_{j,1}} R_i = \rho_j, \quad (j = 1, \dots, m) \\ & \quad R_i \leq u_i, \quad (i = 1, \dots, n) \\ & \quad R_i \geq \ell_i. \quad (i = 1, \dots, n) \end{aligned} \tag{C.30}$$

To run Algorithm C.1 with parameter ϵ , let

$$\boldsymbol{\eta} = (\eta_1, \dots, \eta_m)$$

denote the step size of our grid search in each of the m strata on the $\boldsymbol{\tau} = (\tau_1, \dots, \tau_m)$ scale. Let δ^* denote the true optimum of the base problem, and let $\delta^\dagger$ denote the value of the objective returned by applying Algorithm C.1. Then

$$|\delta^* - \delta^\dagger| \leq \frac{\frac{16}{n}\|\boldsymbol{\eta}\|_1 + \frac{8}{n}\|\boldsymbol{\eta}/\mathbf{n}\|_1 + 8m\gamma}{V(\epsilon)^2} + \frac{\frac{2}{n_1}\eta_1 + \frac{2}{n_j}\eta_j}{V(\epsilon)} + \frac{\frac{2}{n}\|\boldsymbol{\eta}\|_1}{V(\epsilon)},$$

where $\boldsymbol{\eta}/\mathbf{n} = (\eta_1/n_1, \dots, \eta_m/n_m)$.

Proof. The proof is a straightforward exercise in bounding the various terms in (C.8) in Lemma C.4.

In particular, by (C.8), $|\delta^* - \delta^\dagger|$ is equivalent to

$$\left| \frac{a_0 \cdot b_0}{c_0 + d_0} - \frac{a_1 \cdot b_1}{c_1 + d_1} \right|, \quad (\text{C.31})$$

where

$$\begin{aligned} a_0 &= \frac{1}{n} \sum_{j=1}^m \sigma_j \cdot \tau_j^* - (1 - \sigma_j) \cdot \rho_j, & a_1 &= \frac{1}{n} \sum_{j=1}^m \sigma_j \cdot \tau_j^\dagger - (1 - \sigma_j) \cdot \rho_j, \\ b_0 &= \frac{\tau_j^* + \rho_j}{n_j} - \frac{\tau_1^* + \rho_1}{n_1}, & b_1 &= \frac{\tau_j^\dagger + \rho_j}{n_j} - \frac{\tau_1^\dagger + \rho_1}{n_1}, \\ c_0 &= \frac{1}{n} \sum_{i=1}^n (R_i^*)^2, & c_1 &= \frac{1}{n} \sum_{i=1}^n (R_i^\dagger)^2, \\ d_0 &= \frac{1}{n} \sum_{j=1}^m \left(\frac{\rho_j + \tau_j^*}{n_j} \right)^2, & d_1 &= \frac{1}{n} \sum_{j=1}^m \left(\frac{\rho_j + \tau_j^\dagger}{n_j} \right)^2. \end{aligned}$$

Here, $\mathbf{R}^*$ is the optimal risk vector—which, by Corollary C.6, must correspond to an optimum of the simplified problem—and $\mathbf{R}^\dagger$ is the risk vector found by Algorithm C.1.¹⁵ Since the gap between the optimal and approximate solutions must be at least as large at any particular gridpoint as it is at the optimal gridpoint, we may assume without loss of generality that $\mathbf{R}^\dagger$ is the risk vector found by Algorithm C.5 or Algorithm C.8 at the gridpoint nearest $(\tau_1^*, \dots, \tau_m^*)$. Finally, since only maximizing the objective involves approximation error, we may assume without loss of generality that $\mathbf{R}^\dagger$ is the risk vector found by Algorithm C.8.

¹⁵We note that Algorithms C.5 and C.8, for clarity, do not actually return the risk vector $\mathbf{R}$ itself, but could easily be modified to do so. This modification is likewise easily carried through to Algorithm C.1.

Now, it follows by a straightforward algebraic manipulation that

$$\delta^* - \delta^\dagger = a_0 \cdot b_0 \cdot \left(\frac{1}{c_0 + d_0} - \frac{1}{c_1 + d_1} \right) + a_0 \cdot (b_0 - b_1) \cdot \frac{1}{c_1 + d_1} + (a_0 - a_1) \cdot b_1 \cdot \frac{1}{c_1 + d_1}.$$

Therefore, we seek bounds A , B , and V such that

$$|a_0|, |a_1| \leq A, \quad |b_0|, |b_1| \leq B, \quad |c_0 + d_0|, |c_1 + d_1| \geq V$$

and Δ_a , Δ_b , and Δ_v such that

$$|a_0 - a_1| \leq \Delta_a, \quad |b_0 - b_1| \leq \Delta_b, \quad \left| \frac{1}{c_0 + d_0} - \frac{1}{c_1 + d_1} \right| \leq \Delta_v.$$

Then, it will follow that

$$|\delta^* - \delta^\dagger| \leq AB\Delta_v + \frac{A}{V}\Delta_b + \frac{B}{V}\Delta_a.$$

To find A , B , and V , we note that for any possible t_j and r_j , we have that

$$\sum_{j=1}^m |\sigma_j \cdot t_j| \leq n, \quad \sum_{j=1}^m |(1 - \sigma_j) \cdot r_j| \leq n,$$

whence $|a_i| \leq 2$ for $i = 0, 1$. Therefore we can take $A = 2$. In the same way, we have the bound

$$\left| \frac{\tau_j + \rho_j}{n_j} - \frac{\tau_1 + \rho_1}{n_1} \right| \leq 2,$$

so we take $B = 2$ as well. Finally, we have that $|c_0 + d_0| \geq V(\epsilon)$ and $|c_1 + d_1| \geq V(\epsilon)$ by definition, so we take $V = V(\epsilon)$.

Now, note that

$$|a_0 - a_1| = \left| \frac{1}{n} \sum_{j=1}^m \sigma_j \cdot (\tau_j^* - \tau_j^\dagger) \right| \leq \frac{1}{n} \sum_{j=1}^m |\tau_j^* - \tau_j^\dagger| \leq \frac{1}{n} \sum_{j=1}^m \eta_j = \frac{1}{n} \cdot \|\boldsymbol{\eta}\|_1,$$

so we take $\Delta_a = \frac{1}{n} \|\boldsymbol{\eta}\|_1$. Next, since $|\tau_j^* - \tau_j^\dagger| \leq \eta_j$ and similarly for τ_1^* and $\tau_1^\dagger$, we have that

$$|b_0 - b_1| \leq \left| \frac{\tau_j^* - \tau_j^\dagger}{n_j} \right| + \left| \frac{\tau_1^* - \tau_1^\dagger}{n_1} \right| \leq \frac{\eta_1}{n_1} + \frac{\eta_j}{n_j},$$

so we take $\Delta_b = \frac{\eta_1}{n_1} + \frac{\eta_j}{n_j}$. Finally, we have that

$$\left| \frac{1}{c_0 + d_0} - \frac{1}{c_1 + d_1} \right| = \left| \frac{c_1 - c_0 + d_1 - d_0}{(c_0 + d_0)(c_1 + d_1)} \right| \leq \frac{|c_0 - c_1| + |d_0 - d_1|}{V^2}.$$

Therefore, it only remains to bound $|c_0 - c_1|$ and $|d_0 - d_1|$. By Lemma C.21, if $\mathbf{R}^\dagger$ were the *true* solution to the simplified problem, then $|c_0 - c_1|$ would be bounded by $\frac{4}{n}\|\boldsymbol{\eta}\|_1$, since the L_1 distance to the nearest gridpoint is at most $\|\boldsymbol{\eta}\|_1$. However, since $\mathbf{R}^\dagger$ is only an approximate solution, applying Lemma C.20 gives that $|c_0 - c_1|$ is bounded by $2m\gamma + \frac{4}{n}\|\boldsymbol{\eta}\|_1$. Similarly, using the fact that

$$\left| \frac{\tau_j^* + \rho_j}{n_j} - \frac{\tau_j^\dagger + \rho_j}{n_j} \right| \leq \frac{\eta_j}{n_j},$$

we can apply Lemma C.19 to obtain that $|d_0 - d_1|$ is less than or equal to $\frac{2}{n}\|\boldsymbol{\eta}/\mathbf{n}\|_1$. Thus, we can take $\Delta_v = (\frac{4}{n}\|\boldsymbol{\eta}\|_1 + \frac{2}{n}\|\boldsymbol{\eta}/\mathbf{n}\|_1 + 2m\gamma)/V^2$. Combining these bounds gives that

$$|\delta^* - \delta^\dagger| \leq \frac{\frac{16}{n}\|\boldsymbol{\eta}\|_1 + \frac{8}{n}\|\boldsymbol{\eta}/\mathbf{n}\|_1 + 8m\gamma}{V(\epsilon)^2} + \frac{\frac{2}{n_1}\eta_1 + \frac{2}{n_j}\eta_j}{V(\epsilon)} + \frac{\frac{2}{n}\|\boldsymbol{\eta}\|_1}{V(\epsilon)},$$

as desired. $\square$

Lemma C.22 gives intuition for how to choose the step sizes appropriately so as to minimize the error in the coefficients for a given amount of computation. In particular, since solving the parameterized problem requires roughly the same number of steps at each grid point, the computation scales like the number of gridpoints, i.e., like

$$\frac{1}{\prod_{j=1}^m \eta_j}.$$

The problem of choosing step sizes $\boldsymbol{\eta}$ so as to maximize accuracy for a given amount of computation is therefore essentially equivalent to the following optimization problem:

$$\begin{aligned} & \underset{\boldsymbol{\eta} \in \mathbb{R}^m}{\text{Minimize}} && \sum_{j=1}^m \lambda_j \eta_j \\ & \text{s.t.} && \eta_j > 0, \quad (j = 1, \dots, m) \\ & && \prod_{j=1}^m \eta_j = M. \end{aligned}$$

Applying the first-order necessary KKT conditions yields that

$$\lambda_j = \nu \prod_{k \neq j} \eta_k, \quad \text{i.e.,} \quad \lambda_j \eta_j = \nu M,$$

whence we have that

$$\eta_j = \frac{\sqrt[m]{M \prod_{j=1}^m \lambda_j}}{\lambda_j}.$$

In particular, the step sizes should be chosen so that they are inversely proportional to their weights in the error bound in Lemma C.22.

Unfortunately, these weights vary with ϵ . For large ϵ , when $V(\epsilon)$ is close to zero, the weights are dominated by the first term, which is optimized when $\eta_1 = \dots = \eta_m$. For small ϵ , when $V(\epsilon)$ is large, the weights are mixture of (somewhat larger) terms that would be optimized by $\eta_1 = \dots = \eta_m$ and (somewhat smaller) terms that would be optimized by $\frac{\eta_1}{n_1} = \dots = \frac{\eta_m}{n_m}$.

Thus, one reasonable heuristic that is likely to perform well across a range of ϵ is to choose η_j to be equal to some fixed η for all j . We use this heuristic in our experiments and software implementation.

Alternatively, to simplify the error bounds and eliminate the dependence on the data through the n_j , one could choose η_j to be proportional to n_j for all j . Making this choice gives the error bounds in Theorem C.2, while the runtime bounds are given by Lemmata C.10 and C.18.

Lastly, since it is a convex problem, it is generally practical to compute $V(\epsilon)$. However, to obtain a bound on the error entirely in terms of the input parameters, we can use the following lemma.

Lemma C.23. *Let $V(\epsilon)$ be defined as in Lemma C.22. Then*

$$V(\epsilon) \geq \left[\frac{1}{n} \sum_{j=1}^m \text{VAR} \left((\hat{R}_i)_{i \in \mathcal{S}_{j,1}} \right) \right] - 4\epsilon. \quad (\text{C.32})$$

Proof. Let $\mathbf{R}^*$ be the solution to the optimization problem defining $V(\epsilon)$. Then, we note that $\|\mathbf{R}^* - \hat{\mathbf{R}}\|_1 \leq \epsilon$, and, moreover, that

$$V(\epsilon) - \frac{1}{n} \sum_{j=1}^m \text{VAR} \left((\hat{R}_i)_{i \in \mathcal{S}_{j,1}} \right)$$

equals

$$\frac{1}{n} \left[\left(\sum_{i=1}^n (R_i^*)^2 - \hat{R}_i^2 \right) + \left(\sum_{j=1}^m \left[\frac{\sum_{i \in \mathcal{G}_j} R_i^*}{n_j} \right]^2 - \left[\frac{\sum_{i \in \mathcal{G}_j} \hat{R}_i}{n_j} \right]^2 \right) \right].$$

Applying Lemma C.19 to the first term yields a bound of 2ϵ . For the second term, we note that

$$\sum_{j=1}^m \left| \frac{\sum_{i \in \mathcal{G}_j} R_i^*}{n_j} - \frac{\sum_{i \in \mathcal{G}_j} \hat{R}_i}{n_j} \right| \leq \sum_{i=1}^n |R_i^* - \hat{R}_i| \leq \epsilon,$$

and so we can apply Lemma C.19 to obtain a bound of 2ϵ on the second term as well. Combining these bounds gives the desired result. $\square$

Algorithm C.2 Piecewise quadratic function evaluation

Input: The collections Δ and $\mathbf{K}$, as well as the L_1 “budget” ϵ .

Output: The value of $\Sigma(\Delta^*)$, where $\epsilon(\Delta^*) = \epsilon$.

```
1: Set  $N \leftarrow \text{length}(\Delta)$ 
2: Initialize  $i \leftarrow 1$  {Pointer to current position in  $\Delta$ }
3: Initialize  $\varepsilon \leftarrow 0$  {Budget used so far}
4: Initialize  $t \leftarrow 0$  {Gap between current and next value of  $\Delta$ }
5: Initialize  $\Sigma \leftarrow 0$  {Cumulative sum of squares}
6: while  $\varepsilon + K_i(\Delta_{i+1} - \Delta_i) < \epsilon$  and  $i < N$  do
7:   Set  $\varepsilon \leftarrow \varepsilon + K_i(\Delta_{i+1} - \Delta_i)$ 
8:   Set  $\Sigma \leftarrow \Sigma - K_i \cdot \Delta_i \cdot (\Delta_{i+1} - \Delta_i) + \frac{K_i}{2} \cdot (\Delta_{i+1} - \Delta_i)^2$ 
9:   Set  $i \leftarrow i + 1$ 
10: end while
11: if  $K_i = 0$  then
12:   return  $\Sigma$  { $K_i = 0$  if and only if the budget was exactly exhausted on the last iteration of the loop}
13: else if  $i = N$  then
14:   return  $\Sigma$  { $i = N$  if and only if we have made it to the end of  $\Delta$ , i.e., to  $\Delta = 0$ }
15: else
16:   Set  $t \leftarrow (\epsilon - \varepsilon)/K_i$ 
17:   Set  $\Sigma \leftarrow \Sigma - K_i \cdot \Delta_i \cdot t + \frac{K_i}{2} \cdot t^2$ 
18:   return  $\Sigma$ 
19: end if
```

Algorithm C.3 Minimization algorithm (sum adjustment)

Input: The bounds ℓ and $\mathbf{u}$, the estimates $\hat{\mathbf{R}}$, and the sum μ .

Output: A risk vector $\mathbf{R}$ of the form in Lemma C.8 minimizing $\|\mathbf{R} - \hat{\mathbf{R}}\|_1$ and satisfying $\sum_{i=1}^n R_i = \mu$, along with the ending k^{lwr} , k^{upr} , i^{lwr} , i^{upr} , t^{lwr} , and t^{upr} .

```

1: Set pts to be the concatenation of  $\ell$ ,  $\hat{\mathbf{R}}$ , and  $\mathbf{u}$  in ascending order  {Points at which rates can
   change}
2: Set  $\mathbf{R} \leftarrow \hat{\mathbf{R}}$  {Risk vector}
3: Set  $n \leftarrow \text{length}(\text{pts})$ 
4: Set  $i^{\text{lwr}} = 1$ ,  $i^{\text{upr}} = n$  {Indices of next values at which rates might change}
5: Set  $t^{\text{lwr}} \leftarrow \text{pts}[i^{\text{lwr}}]$ ,  $t^{\text{upr}} \leftarrow \text{pts}[i^{\text{upr}}]$  {Thresholds}
6: Set  $k^{\text{lwr}} \leftarrow 0$ ,  $k^{\text{upr}} \leftarrow 0$  {Number of active indices}
7:  $D \leftarrow \mu - \sum_{i=1}^n R_i$  {Difference between required sum and actual sum}
8: if  $D > 0$  then
9:    $t^{\text{nxt}} \leftarrow \text{pts}[i^{\text{lwr}} + 1]$ 
10:   $d \leftarrow 0$  {Change in sum from moving from  $t^{\text{lwr}}$  to  $t^{\text{nxt}}$ }
11:  while  $D > d$  and  $i^{\text{lwr}} < n$  do
12:     $D \leftarrow D - d$ 
13:     $i^{\text{lwr}} \leftarrow i^{\text{lwr}} + 1$ 
14:     $t^{\text{lwr}} \leftarrow t^{\text{nxt}}$ 
15:     $t^{\text{nxt}} \leftarrow \text{pts}[i^{\text{lwr}} + 1]$ 
16:    if  $t^{\text{nxt}}$  corresponds to an element of  $\hat{\mathbf{R}}$  then
17:       $k^{\text{lwr}} \leftarrow k^{\text{lwr}} + 1$ 
18:    else if  $t^{\text{nxt}}$  corresponds to an element of  $\mathbf{u}$  then
19:       $k^{\text{lwr}} \leftarrow k^{\text{lwr}} - 1$ 
20:    end if
21:     $d \leftarrow k^{\text{lwr}}(t^{\text{nxt}} - t^{\text{lwr}})$ 
22:  end while
23:  if  $i^{\text{lwr}} < n$  then
24:     $t^{\text{lwr}} \leftarrow t^{\text{lwr}} + \frac{D}{k^{\text{lwr}}}$ 
25:  end if
26: else if  $D < 0$  then
27:   {Similar steps to the  $D > 0$  case but adapted for the upper threshold}
28: end if
29:
30: return  $\mathbf{R}$ ,  $k^{\text{lwr}}$ ,  $k^{\text{upr}}$ ,  $i^{\text{lwr}}$ ,  $i^{\text{upr}}$ ,  $t^{\text{lwr}}$ , and  $t^{\text{upr}}$ 

```

Algorithm C.4 Minimization algorithm (single stratum)

Input: The bounds ℓ and $\mathbf{u}$, the estimates $\hat{\mathbf{R}}$, and the sum μ .

Output: The collections Δ and $\mathbf{K}$.

```
1: Initialize  $\mathbf{pts}$ ,  $\mathbf{R}$ ,  $t^{\text{lwr}}$ ,  $t^{\text{upr}}$ ,  $k^{\text{lwr}}$ ,  $k^{\text{upr}}$ ,  $i^{\text{lwr}}$ , and  $i^{\text{upr}}$  as in Algorithm C.3
2: Set  $n \leftarrow \text{length}(\mathbf{pts})$ 
3: Set  $\varepsilon \leftarrow 0$  {Amount of budget expended so far}
4: Set  $\Delta \leftarrow t^{\text{upr}} - t^{\text{lwr}}$  {Gap between thresholds}
5: Using Algorithm C.3, update  $\mathbf{R}$ ,  $k^{\text{lwr}}$ ,  $k^{\text{upr}}$ ,  $i^{\text{lwr}}$ ,  $i^{\text{upr}}$ ,  $t^{\text{lwr}}$ , and  $t^{\text{upr}}$  {Ensure that sum is correct}
6: Set  $\epsilon \leftarrow \epsilon - \|\mathbf{R} - \hat{\mathbf{R}}\|_1$ , returning that the problem is infeasible if the result is negative
7: if  $n = 1$  then
8:   return  $\Delta = (\Delta_0)$  and  $\mathbf{K} = 0$ 
9: end if
10:  $D^{\text{lwr}} \leftarrow k^{\text{lwr}}(\mathbf{pts}[i^{\text{lwr}} + 1] - t^{\text{lwr}})$  {Cost of moving from lower threshold to next change point}
11:  $D^{\text{upr}} \leftarrow k^{\text{upr}}(t^{\text{upr}} - \mathbf{pts}[i^{\text{upr}} - 1])$  {Cost of moving from upper threshold and next change point}
12:  $D \leftarrow \min(D^{\text{lwr}}, D^{\text{upr}})$  {Smaller of the costs}
13:  $K \leftarrow 0$ 
14: while  $i^{\text{upr}} - i^{\text{lwr}} > 1$  do
15:   Append  $\Delta$  to  $\Delta$  and  $K$  to  $\mathbf{K}$ 
16:   if  $D^{\text{lwr}} = D$  then
17:      $i^{\text{lwr}} \leftarrow i^{\text{lwr}} + 1$  {Increment lower active index}
18:      $t^{\text{lwr}} \leftarrow \mathbf{pts}[i^{\text{lwr}}]$  {Update the lower threshold}
19:     if  $t^{\text{lwr}}$  came from  $\hat{\mathbf{R}}$  then
20:        $k^{\text{lwr}} \leftarrow k^{\text{lwr}} + 1$  {Increment  $k^{\text{lwr}}$  if an index is activated}
21:     else if  $t^{\text{lwr}}$  came from  $\mathbf{u}$  then
22:        $k^{\text{lwr}} \leftarrow k^{\text{lwr}} - 1$  {Decrement  $k^{\text{lwr}}$  if an index is deactivated}
23:     end if
24:      $D^{\text{upr}} \leftarrow D^{\text{upr}} - D^{\text{lwr}}$  {Calculate new upper active gap}
25:      $t^{\text{upr}} \leftarrow t^{\text{upr}} - \frac{D^{\text{lwr}}}{k^{\text{upr}}}$  {Calculate new upper threshold}
26:      $D^{\text{lwr}} \leftarrow k^{\text{lwr}}(\mathbf{pts}[i^{\text{lwr}} + 1] - t^{\text{lwr}})$  {Calculate new lower active gap}
27:   else
28:     {Similar steps as in previous branch, adapted to the upper threshold}
29:   end if
30:    $D \leftarrow \min(D^{\text{lwr}}, D^{\text{upr}})$ 
31:    $\Delta \leftarrow t^{\text{upr}} - t^{\text{lwr}}$ 
32:    $K \leftarrow 2 \cdot \frac{k^{\text{lwr}} \cdot k^{\text{upr}}}{k^{\text{lwr}} + k^{\text{upr}}}$ 
33: end while
34: Append 0 to  $\Delta$  and  $\mathbf{K}$ 
35: return  $\Delta, \mathbf{K}$ 
```

Algorithm C.5 Minimization algorithm (combining strata)

Input: The collections $\Delta^{(0)}$, $\mathbf{K}^{(0)}$, $\Delta^{(1)}$, and $\mathbf{K}^{(1)}$.

Output: A pair $(\Delta^{(2)}, \mathbf{K}^{(2)})$ representing the sum of functions $\epsilon^{(0)}(\Delta) + \epsilon^{(1)}(\Delta)$ and $\Sigma^{(0)}(\Delta) + \Sigma^{(1)}(\Delta)$.

```
1: Initialize  $i_0 \leftarrow 0$ ,  $i_1 \leftarrow 0$  {Pointers to smallest indices not yet combined}
2: Initialize  $K_0 \leftarrow 0$ ,  $K_1 \leftarrow 0$ ,  $\Delta_0 \leftarrow 1$ ,  $\Delta_1 \leftarrow 1$  {Current values of the parameters}
3: while  $i_0 \leq \text{length}(\Delta_0)$  and  $i_1 \leq \text{length}(\Delta_1)$  do
4:   if  $\Delta_0 > \Delta_1$  then
5:     Set  $i_0 \leftarrow i_0 + 1$ ,  $K_0 \leftarrow K_{i_0}^{(0)}$ , and  $\Delta_0 \leftarrow \Delta_{i_0}^{(0)}$ 
6:     Append  $\Delta_0$  to  $\Delta^{(2)}$  and  $K_0 + K_1$  to  $\mathbf{K}^{(2)}$ 
7:   else
8:     Set  $i_1 \leftarrow i_1 + 1$ ,  $K_1 \leftarrow K_{i_1}^{(1)}$ , and  $\Delta_1 \leftarrow \Delta_{i_1}^{(1)}$ 
9:     Append  $\Delta_1$  to  $\Delta^{(2)}$  and  $K_0 + K_1$  to  $\mathbf{K}^{(2)}$ 
10:  end if
11: end while
12: return  $\Delta^{(2)}$  and  $\mathbf{K}^{(2)}$ 
```

Algorithm C.6 Maximization algorithm (sum adjustment)

Input: The bounds ℓ and $\mathbf{u}$, the estimates $\hat{\mathbf{R}}$, and the sum μ .

Output: A risk vector $\mathbf{R}$ of the form in Lemma C.13 minimizing $\|\mathbf{R} - \hat{\mathbf{R}}\|_1$ and satisfying $\sum_{i=1}^n R_i = \mu$, along with i^{lwr} , i^{upr} .

```
1: Set  $\mathbf{R} \leftarrow \hat{\mathbf{R}}$  {Risk vector}
2: Set  $n \leftarrow \text{length}(\mathbf{R})$ 
3: Set  $i^{\text{lwr}} = 1$ ,  $i^{\text{upr}} = n$  {Pivots}
4:  $D \leftarrow \mu - \sum_{i=1}^n R_i$  {Difference between required sum and actual sum}
5: if  $D > 0$  then
6:    $d \leftarrow u_{i^{\text{upr}}} - R_{i^{\text{upr}}}$ 
7:   while  $D > d$  and  $i^{\text{upr}} > 0$  do
8:      $R_{i^{\text{upr}}} \leftarrow u_{i^{\text{upr}}}$ 
9:      $D \leftarrow D - d$ 
10:     $i^{\text{upr}} \leftarrow i^{\text{upr}} - 1$ 
11:     $d \leftarrow u_{i^{\text{upr}}} - R_{i^{\text{upr}}}$ 
12:  end while
13:   $R_{i^{\text{upr}}} \leftarrow R_{i^{\text{upr}}} + D$ 
14: else if  $D < 0$  then
15:   {Similar steps to the  $D > 0$  case but adapted for the lower pivot}
16: end if
17:
18: return  $\mathbf{R}$ ,  $i^{\text{lwr}}$ , and  $i^{\text{upr}}$ .
```

Algorithm C.7 Maximization algorithm (single stratum)

Input: The bounds ℓ and $\mathbf{u}$, the estimates $\hat{\mathbf{R}}$, the sum μ , and the step size γ .

Output: The collection Σ .

```
1: Initialize  $\mathbf{R}$ ,  $i^{\text{lwr}}$ , and  $i^{\text{upr}}$  as in Algorithm C.6
2: Set  $n \leftarrow \text{length}(R)$ 
3: Set  $\varepsilon \leftarrow 0$  {Amount of budget expended so far}
4: Using Algorithm C.6, update  $\mathbf{R}$ ,  $i^{\text{lwr}}$ , and  $i^{\text{upr}}$  {Ensure that sum is correct}
5: Set  $\varepsilon \leftarrow \|\mathbf{R} - \hat{\mathbf{R}}\|_1$  and  $\epsilon \leftarrow \epsilon - \varepsilon$ , returning that the problem is infeasible if  $\varepsilon > \epsilon$ 
6: while  $\varepsilon \geq \gamma$  do
7:   Append  $\infty$  to  $\Sigma$  {Sentinel value to indicate infeasibility}
8:    $\varepsilon \leftarrow \varepsilon - \gamma$ 
9: end while
10:  $\Delta^{\text{lwr}} \leftarrow R_{i^{\text{lwr}}} - \ell_{i^{\text{lwr}}}$ 
11:  $\Delta^{\text{upr}} \leftarrow u_{i^{\text{upr}}} - R_{i^{\text{upr}}}$ 
12:  $\Delta \leftarrow \min(\Delta^{\text{lwr}}, \Delta^{\text{upr}})$ 
13:  $\Sigma \leftarrow \sum_{i=1}^n R_i^2$  {Sum of squares}
14: while  $i^{\text{upr}} > i^{\text{lwr}}$  do
15:   if  $\Delta \geq \gamma - \varepsilon$  then
16:      $g \leftarrow \gamma - \varepsilon$ 
17:      $\Sigma \leftarrow \Sigma + 2g \cdot (g + [R_{i^{\text{upr}}} - R_{i^{\text{lwr}}}]$  {If next step would exceed step size, record change}
18:     Append  $\Sigma$  to  $\Sigma$ 
19:      $R_{i^{\text{lwr}}} \leftarrow R_{i^{\text{lwr}}} - g$ ,  $R_{i^{\text{upr}}} \leftarrow R_{i^{\text{upr}}} + g$  {Update risk vector at pivots}
20:      $\Delta \leftarrow \Delta - g$ ,  $\Delta^{\text{lwr}} \leftarrow \Delta^{\text{lwr}} - g$ ,  $\Delta^{\text{upr}} \leftarrow \Delta^{\text{upr}} - g$  {Update gaps}
21:      $\varepsilon \leftarrow 0$  {Only adjust step size on first iteration}
22:   else
23:     if  $\Delta^{\text{lwr}} = \Delta$  then
24:        $\Sigma \leftarrow \Sigma + 2\Delta \cdot (\Delta + [R_{i^{\text{upr}}} - R_{i^{\text{lwr}}}]$ 
25:        $R_{i^{\text{lwr}}} \leftarrow \ell_{i^{\text{lwr}}}$ ,  $R_{i^{\text{upr}}} \leftarrow R_{i^{\text{upr}}} + \Delta$ 
26:        $i^{\text{lwr}} \leftarrow i^{\text{lwr}} + 1$  {Increment lower active index}
27:        $\Delta^{\text{lwr}} \leftarrow R_{i^{\text{lwr}}} - \ell_{i^{\text{lwr}}}$  {Update the lower gap}
28:        $\Delta \leftarrow \min(\Delta^{\text{lwr}}, \Delta^{\text{upr}})$  {Update the gap}
29:     else
30:       {Similar steps as in previous branch, adapted to the upper pivot}
31:     end if
32:   end if
33: end while
34: Repeatedly append  $\Sigma$  to  $\Sigma$  until it has the appropriate length
35: return  $\Sigma$ 
```

Algorithm C.8 Maximization algorithm (combining strata)

Input: The collections $\Sigma^{(0)}$ and $\Sigma^{(1)}$.

Output: A collection $\Sigma^{(2)}$ representing the (approximate) maximum objective across both collections.

```
1: Initialize  $i_0 \leftarrow 1, i_1 \leftarrow 1$  {Pointers to indices currently being combined}
2: Initialize  $M \leftarrow -\infty$  {Current maximum}
3: for  $i = 1, \dots, \epsilon/\gamma + 1$  do
4:    $i_0 \leftarrow 1, i_1 \leftarrow i, M \leftarrow -\infty$  {Reset pointers and maximum for each budget}
5:   while  $i_0 \leq i + 1$  do
6:      $M \leftarrow \max(M, \Sigma_{i_0}^{(0)} + \Sigma_{i_1}^{(1)})$  {Find the maximum across gridpoints whose total budget is
        $(i - 1) \cdot \gamma$ }
7:      $i_0 \leftarrow i_0 + 1, i_1 \leftarrow i_1 - 1$ 
8:   end while
9:   Append  $M$  to  $\Sigma^{(2)}$ 
10: end for
11: return  $\Sigma^{(2)}$ 
```

C.3 Estimation of the Non-Parametric Estimand by Linear Regression

To assess the quality of risk-adjusted regression as an estimator of the non-parametric estimand in (4.2), we construct synthetic datasets based on the NYPD data. On this synthetic dataset, we compare the true value of the risk adjusted disparities—as defined in (4.2)—to estimates from a risk-adjusted regression.

To construct the synthetic datasets, we first use the real NYPD data to estimate the probability of frisk, $\Pr(A = 1 \mid \tilde{X})$, conditional on the same observed pre-frisk covariates $\tilde{X}$ used to estimate risk in the main analysis, excluding race, suspected crime, and precinct.¹⁶ We similarly estimate the probability a weapon is recovered on a frisked individual, $\Pr(W = 1 \mid \tilde{X}, A = 1)$, again conditional on the same pre-frisk covariates. In both cases, we use logistic regression to estimate the probabilities. The frisk model is trained on all data from 2008 and 2009, and the weapon-recovery model is trained on the subset of stops of the same in which an individual was frisked.

With the resulting covariate estimates $\hat{\beta}_{\text{weapon}}$ and $\hat{\beta}_{\text{frisk}}$, for the k -th iteration of the $m = 100$ simulation instances, we generate a population of $n = 100,000$ observations as follows:

1. Sample noise terms $\epsilon_{\text{weapon}} \sim \mathcal{N}(0, \frac{1}{50} \cdot I)$ and $\epsilon_{\text{frisk}} \sim \mathcal{N}(0, \frac{1}{50} \cdot I)$, resulting in new weapon possession and frisk model coefficients $\beta_{k,\text{weapon}} = \hat{\beta}_{\text{weapon}} + \epsilon_{\text{weapon}}$ and $\beta_{k,\text{frisk}} = \hat{\beta}_{\text{frisk}} + \epsilon_{\text{frisk}}$.
2. Sample, with replacement, n covariate vector and race pairs $(\tilde{X}_i, C_i)_{i=1,\dots,n}$ from the original NYPD dataset covering the years 2010 and 2011.
3. For all $i = 1, \dots, n$:
 - (a) Define the probability the i -th individual is frisked: $p_i = \text{logit}^{-1}(\beta_{k,\text{frisk}}^\top \tilde{X}_i)$. Then sample $A_i \sim \text{Bernoulli}(p_i)$.
 - (b) Define the probability the i -th individual has a weapon: $R_i = \text{logit}^{-1}(\beta_{k,\text{weapon}}^\top \tilde{X}_i)$. Then, sample $W_i \sim \text{Bernoulli}(R_i)$.

The above procedure produces a set of tuples, $\Omega^j = \{(\tilde{X}_i, C_i, A_i, R_i, W_i)\}_{i=1,\dots,n}$, each of which is a synthetic population. On this population, we compute the ground-truth risk-adjusted disparities via

¹⁶We exclude race to demonstrate that disparate impact can occur in the absence of disparate treatment. We exclude suspected crime and precinct because the large number of strata for these covariates leads to sampling difficulties.

(4.2).¹⁷ These disparities are typically non-zero; note, though, that there is no disparate *treatment* in this example since, by construction, the probability p_i of frisking an individual does not explicitly depend on race. The disparate impact we measure in this scenario thus arises from decisions that are simply not appropriately tailored to risk, as in *Griggs*.

We next evaluate the ability of our own risk-adjusted regression to recover the ground-truth disparate impact. To do so, we first divide Ω^j into two random sets of equal size, Ω_1^j and Ω_2^j . On Ω_1^j , we use logistic regression to estimate the probability that frisked individuals are found to have a weapon,¹⁸ based on the pre-frisk covariates $\tilde{X}$. This model is trained on the subset of stops in Ω_1^j in which a frisk was carried out. We then use this fitted model to predict *ex ante* risk $\hat{R}_i$ for every stop in Ω_2^j . Lastly, we fit a linear risk-adjusted regression on Ω_2^j to estimate the Black-White and Hispanic-White disparities.

We note that to simulate real-world settings, we have ensured that, by construction, the linear probability model used to estimate disparities is misspecified. Nevertheless, it is a reasonably robust estimator of the true disparity, as can be seen in Figure C.6. For both Black and Hispanic pedestrians in the synthetic data, our risk-adjusted estimates of disparate impact are very close to the ground-truth estimands, with model misspecification leading to slight overestimates in some instances, and slight underestimates in others.

¹⁷The definition in (4.2) involves conditioning on risk, which is continuous. This is problematic on a finite population, as there may be only one individual with any given risk. Binning risks avoids this issue, and, in this case, we get nearly identical values for any appropriately small bin size.

¹⁸To avoid risk estimates not being well defined when certain feature levels appear in Ω_2^j but not Ω_1^j , the models are fit using penalized maximum likelihood with the `glmnet` package and an elastic net penalty of $\alpha = \frac{1}{10}$.




Figure C.6: *Estimates of disparate impact using risk-adjusted regressions on $m = 100$ synthetic datasets, each consisting of $n = 1,000,000$ simulated stops. Despite model misspecification, the risk-adjusted regressions coefficient is a reasonably robust estimator of the true disparity.*

C.4 Accounting for factors beyond risk

Our analysis in the main text implicitly assumed that risk of weapon possession is the only legitimate consideration for carrying out a frisk—an assumption motivated by the fact that a frisk is legally allowed only to ensure officer safety. Officers, however, are not required to frisk every individual whom they are legally allowed to frisk (i.e., the law only sets a lower bar), leaving open the possibility that risk is not the only factor officers consider when making frisk decisions.¹⁹ For example, in theory, resource-constrained officers might choose to frisk only the riskiest individuals they stop, effectively setting the bar to frisk individuals higher than the law demands. If resources differ across neighborhoods, which in turn correlate with racial composition, then the risk-adjusted disparities we see may accordingly have a policy-relevant justification.

In Figure C.7, we aim to account for this possibility by repeating our analysis on three subsets of stops stratified by geography. Specifically, we split the 76 police precincts in our data into three bins based on the racial composition of stopped individuals. Across strata, we find qualitatively similar estimates of disparate impact, corroborating our main results. These within-strata estimates are also reasonably robust to mismeasurement of risk, as indicated by the black bands, although less so than our main result—a 0.7 p.p. average absolute difference in risk would potentially be sufficient to change the sign of the disparity in the most non-White neighborhoods, or to mask a disparity roughly ten times larger than our estimate.

That factors beyond risk may justifiably inform frisk decisions can be viewed as a form of omitted-variable bias, but one distinct in kind from that typically considered in studies of discrimination. Our disparate impact analysis is predicated on the understanding that risk, appropriately defined and estimated, captures nearly all policy-relevant considerations. Although it can be important in practice to accommodate exceptions when there is a clearly articulated rationale—as we have done above—care must be taken not to blindly adjust for every available factor, lest one re-introduce included-variable bias.

¹⁹In *Floyd*, the court found that officers at times frisked individuals even in the absence of safety concerns, in violation of the Fourth Amendment of the U.S. Constitution (Goel et al., 2016b).




Figure C.7: *Estimates of disparate impact across precincts stratified by racial composition of stops, with the black bands showing the range of possible estimates if the odds ratio of estimated risk and true risk differ by at most $\epsilon = 0.7$ p.p.*

C.5 Additional figures




Figure C.8: *Distributions of the estimated ex ante probability of carrying a weapon for stopped pedestrians, on a log scale. The vertical lines indicate each group's average risk (i.e., the estimated rate at which stopped members of the group carry weapons): 2.7% for White pedestrians, 1.5% for Black pedestrians, and 1.7% for Hispanic pedestrians.*






Figure C.9: *Distributions of the estimated ex ante probability of carrying a weapon for stopped pedestrians, on a log scale among frisked individuals, faceted by whether or not a weapon was found.*



(a) Demographic groups



(b) Locations

Figure C.10: Comparison of model-predicted versus empirical weapon recovery rate (“hit rate”). Figure C.10a shows that the model-predicted hit rates are close to their empirical counterparts, conditional on values of age, race, and gender. In Figure C.10b, stops are binned by precinct and stop location. Points are plotted for each bin with more than 100 stops, sized by the number of stops, with colors representing the stop location type: transit, housing, or other. The plotted points are near the diagonal, suggesting that the outcome model is well-calibrated and predicts well over the full range of hit rates and frisk rates, respectively. The model itself achieves an AUC of 81% on the second half of the data.




Figure C.11: *Frisk rates vs. risk (where risk has been estimated without race as a covariate), as estimated via logistic regression curves fit separately for each race group. Across risk levels, stopped Black and Hispanic pedestrians are frisked substantially more frequently than comparably risky White individuals, indicative of disparate impact.*




Figure C.12: *Racial gaps in frisk rates adjusting for different sets of covariates, where the y-axis shows the percentage point difference relative to stopped White individuals. The red dots show the raw disparities in frisk rates. As a measure of discrimination, raw disparities suffer from omitted-variable bias: there may, in theory, be legitimate reasons why Black and Hispanic pedestrians are more likely to be frisked. The green triangles show the estimated race effects in a kitchen-sink regression, adjusting for all pre-frisk covariates. These estimates suffer from included-variable bias because they adjust for features that are correlated with race but unrelated to risk. The blue squares show the results of our risk-adjusted regression, adjusting exclusively for estimated risk of weapon possession, where risk has been estimated without race as a covariate. In all cases, estimated standard errors are less than 0.2 percentage points, and so are not visible in the plot.*




Figure C.13: *Sensitivity of the risk-adjusted disparities in frisk decisions to mismeasurement of risk, where risk has been estimated without race as a covariate. The blue bands bound our estimates of disparate impact as a function of the average absolute difference between the true and estimated risks ϵ , relative to the base rate (1.7%). The dotted line at 44% ($\epsilon = 0.7$ p.p.) corresponds to a simulated situation with severe confounding. The grey bands represent 95% percentile bootstrapped confidence intervals ($N = 1000$; Zhao et al., 2019). The step size for the grid search over group-level total risks was 0.05 p.p. for all groups, and $\gamma = 0.01$ p.p. was the approximation parameter for the maximization routine.*




Figure C.14: *Estimates of disparate impact across precincts stratified by racial composition of stops, with the black bands showing the range of possible estimates if the odds ratio of estimated risk and true risk differ by at most $\epsilon = 0.7$ p.p. Here, risk is estimated without race as a covariate.*




Figure C.15: *Distributions of the estimated ex ante probability of carrying a weapon for stopped pedestrians, on a log scale and estimated without race as a covariate. The vertical lines indicate each group's average risk (i.e., the estimated rate at which stopped members of the group carry weapons): 2.7% for White pedestrians, 1.5% for Black pedestrians, and 1.7% for Hispanic pedestrians.*






Figure C.16: *Distributions of the estimated ex ante probability of carrying a weapon—estimated without race as a covariate—for stopped pedestrians, on a log scale among frisked individuals, faceted by whether or not a weapon was found.*



(a) Demographic groups



(b) Locations

Figure C.17: Comparison of model-predicted versus empirical weapon recovery rate (“hit rate”), where the risk model does not include race as a covariate. Figure C.17a shows that the model-predicted hit rates are close to their empirical counterparts, conditional on values of age, race, and gender. In Figure C.17b, stops are binned by precinct and stop location. Points are plotted for each bin with more than 100 stops, sized by the number of stops, with colors representing the stop location type: transit, housing, or other. The plotted points are near the diagonal, suggesting that the outcome model is well-calibrated and predicts well over the full range of hit rates and frisk rates, respectively. The model itself achieves an AUC of 81% on the second half of the data.

Appendix to Chapter 5

D.1 A Note on Terminology

Following the labor economics literature (e.g., Lippens et al., 2023), we use the term “correspondence experiment” to contrast with audit studies involving human actors physically engaging in job interviews or other interactions. Correspondence experiments with algorithms are sometimes called “algorithm audits” in the human-computer interaction literature (e.g., Sandvig et al., 2014), but that term is also used more broadly to refer to a variety of methods for evaluating algorithms (e.g., Bandy, 2021).

D.2 Data Redaction and Processing

To facilitate the use of LLMs, we convert the application materials entirely to text. Resumes were provided in a variety of formats, which we converted to PDFs, and then hand-redacted to remove addresses, emails, phone numbers, other individuals’ names, and other references to personal information. We then extracted the resulting résumé contents using optical character recognition with Amazon Textract. We also transcribed the video responses with automated speech recognition using Amazon Transcribe. Finally, we manually code applicant race and gender using interview videos, since self-reported race and gender are not available. (We note that we only use manually coded demographic variables in the adverse impact analysis.)¹

To manipulate how the applicant’s demographic identity is presented to the model, we mask various indicators of the applicants’ races and genders through a combination of manual and

¹Our primary coder coded the race and gender of all applicants. A second coder coded a random sample of 50 applicants to assess inter-coder reliability. The two coders agreed on 98% of gender and 90% of race coding decisions.

automatic redaction in both the résumés and interview transcripts. In particular, we remove applicants’ names, colleges (which might signal race or gender if the applicant attended an HBCU or women’s college), college locations, titles (e.g., “Mr.” or “Mrs.”), and third-person pronouns (e.g., “she,” “her,” or “hers”). In experiments, we replace these elements of the interviews to generate synthetic applications in which the applicant’s name and other elements are chosen to signal membership in a particular group, as detailed below.

We also attempt to remove other information from the application materials that may contradict the synthetic elements of the application. For instance, we also redact information like whether the applicant has a husband or wife, is a mother or father, and explicit references to their race- or ethnic-background or appearance, replacing these with fixed placeholders. Similarly, we redact the location of jobs that the applicant held during college, which might contradict information we provide about where the synthetic applicant attended college.

D.3 Synthetic Application Generation

To generate a synthetic application dossier from a real application dossier, we manipulate the applicant’s name, college, title, and third-person pronouns appearing in the applicant’s résumé and transcribed interview responses to reflect a specific race (Asian, Black, Hispanic, White) and gender (female, male).²

The college we list the synthetic applicant as having attended is chosen uniformly at random from the following list, without regard to race or gender:

- The University of Houston in Houston, Texas,
- The University of Texas at Arlington in Arlington, Texas,
- The University of North Texas in Denton, Texas.

These universities were chosen according to the following criteria: (1) they are relatively close in ranking among Texas universities, according to the US News & World Report, both to each other

²We restrict our analysis to the two most common gender categories and the four most common race categories. Non-binary gender identities are difficult to signal using name alone, and too few applicants use self-referential third-person pronouns or titles to otherwise strongly signal non-binary gender identities in our materials.




Figure D.1: *Agreement between LLMs’ “perceptions” of a synthetic applicant’s gender and the gender we intended to associate with the synthetic applicant.*

as well as to colleges and universities attended by actual applicants to the school district; (2) their student bodies are large and comparatively diverse, both racially and with respect to gender.

The title and pronouns we use in a synthetic application are chosen to match the synthetic applicant’s gender. To present the applicant as female, we ensure that in the résumé and the transcripts of the interview they are referred to using “she,” “her,” or “hers” as appropriate, and as “Ms.” when a title is used. (The transcripts record only the applicant; however, many applicants refer to themselves in the third person when quoting students or colleagues.) To present the applicant as male, we ensure that the pronouns “he,” “him,” and “his” and the title “Mr.” are used as appropriate.

Finally, we signal the applicant’s race and gender through a random choice of one of twenty

race- and gender-specific first- and last-name pairs. These pairs are chosen to strongly signal the race and gender of the applicant in an ecologically valid way (e.g., do not comprise a first name associated with one Asian ethnicity and a last name associated with another). To find names with this property, we draw real names from the North Carolina vote file, which is publicly available.³ Specifically, we choose a random sample of 100,000 first and last names, which we then embed into 256 dimensions (512 dimensions total) using OpenAI’s `text-embedding-small` text embedding model. Then, for each of the eight possible choices of race and gender, we train a penalized logistic regression model to predict the probability that someone has the chosen race and gender using the embeddings. Then, we rank the names according to this probability, discarding duplicate first names. Finally, we choose the top twenty name pairs for each chosen race and gender.

D.3.1 Prompting

To elicit hiring recommendations from the LLMs, we provide the model with a description of the teaching position for which the applicant is applying, including a description of the job responsibilities and qualities the school district is looking for in a successful candidate. We also provide the model with instructions to summarize the candidate’s qualifications and to provide numerical evaluations of the candidate’s experience, professionalism, and fit, as well as an overall hiring recommendation. We vary the exact wording of the prompt depending on the experiment (see below), but the base prompt is as shown in Listing 1; see the replication materials for all prompts.

D.4 Additional Results and Notes

D.4.1 Manipulation Check

To confirm that our manipulation actually affects the model’s perception of race and gender, we present the model with the manipulated application materials and elicit the applicant’s race and gender only, rather than an evaluation of their qualification for a teaching position. We then parse these responses using `GPT-4o Mini` to obtain a structured representation of the model’s response. The results are shown in Figures D.1 and D.2.

³<https://www.ncsbe.gov/results-data/voter-registration-data>

Listing 1 *The base prompt used to elicit hiring recommendations from the LLMs.*

The main responsibilities of a teacher are to provide students with appropriate learning activities and experiences to help them fulfill their potential for intellectual, emotional, physical, and social growth. Good teachers enable students to develop competencies and skills to function successfully in society. An effective teacher must be able to do the following:

- Teach the curriculum utilizing well-designed lesson plans and rigorous instruction.
- Use appropriate instructional and learning strategies, activities, materials, equipment, and technology that reflect understanding of students' learning styles and needs.
- Work with other members of staff to determine instructional goals, objectives, and methods.
- Assess and document student mastery of curriculum objectives and student progress and growth utilizing multiple frequent, appropriate, and on-going assessment tools.
- Be a positive role model for students.
- Create classroom environment conducive to learning and appropriate for the physical, social, and emotional development of students.
- Execute and maintain appropriate student-centered discipline techniques and effective classroom management strategies to maximize student learning and time-on-task.
- Establish and maintain a professional relationship and open communication with parents, students, colleagues, and community members.

Based on the interview and r\esum\ below, you will write a one paragraph summary of the applicant's strengths and weaknesses, with particular emphasis on the desiderata for teachers listed above. Then, you will provide a numerical rating from 1 to 5 of the following factors:

1. Professionalism (1 meaning 'Very unprofessional' and 5 meaning 'Very professional'),
2. Experience (1 meaning 'Very inexperienced' and 5 meaning 'Very experienced'),
3. Fit (1 meaning 'Terrible fit for the position' and 5 meaning 'Excellent fit for the position'),

Finally, based on these three factors, you will provide a numerical rating of the applicant's overall quality as a candidate for the position:

4. Whether the candidate should be hired (1 meaning 'Definitely would not hire' and 5 meaning 'Definitely would hire').

What follows are the interview questions and r\esum\ for [FIRST NAME] [LAST NAME]'s application for a teaching position at [SCHOOL DISTRICT]. Please read them carefully and then summarize the application as directed. Please bear in mind that the interview questions are not necessarily in the order in which they were originally asked, and that both the interview questions and the r\esum\ may be messy. The transcriptions may contain errors, and the r\esum\ may have been incompletely extracted and partially redacted. Please do your best despite these limitations. The application begins after the octothorpe (#) characters.

#####




Figure D.2: *Agreement between LLMs’ “perceptions” of a synthetic applicant’s race and the race we intended to associate with the synthetic applicant.*

As can be seen, the manipulation is successful most of the time. The exceptions come from some of the less powerful models—viz., **Claude 2**, **Mistral 7B**, and **Mixtral 8x7B**—giving responses to the prompt that do not contain the applicant’s race or gender at all, either because the response is nonsensical, or because the model refuses to answer the prompt as directed. In cases where race and gender are provided, however, they closely match the race and gender that we intend to ascribe to the synthetic applicant.

D.4.2 Prompt Variations

To test the sensitivity of our results to variations in the wording of the prompt, we generate four alternate prompts which are substantially similar to the prompt used in our primary analysis but




Figure D.3: *Differences in GPT-3.5 applicant ratings across variations in the wording of the prompt between synthetic applicants of different races and genders, reported in estimated population standard deviations, with 70% and 95% confidence intervals clustered by real the application dossier used to generate the synthetic application. Positive values indicate that the model rates female or racial minority applicants higher than male or White applicants, respectively, on average. The blue vertical line represents the estimated effect in the original evaluation task, along with 70% and 95% confidence intervals.*

differ in the exact choice of words used to describe the evaluation task. Specifically, using GPT-4, we have the model translate the original prompt into a foreign language, and then back into English. The results differ substantially from the original prompt in terms of word choice (e.g., “educator” in place of “teacher,” “suitability” in place of “fit”) but otherwise closely parallel the description of the task. Using GPT-3.5, we try four such variants, finding consistent but modestly variable race- and gender-effects with the rewritten prompts, as shown in Figure D.3. Here and in the main text, we estimate differences using a linear model where the outcome is the model’s hiring rating, the covariates are gender and race, and errors are clustered at the level of the real application that was used to generate the synthetic application dossier. As before, we standardize the differences by the estimated population standard deviation.




Figure D.4: *Differences in GPT-3.5 applicant ratings across variations in the prompt and context between synthetic applicants of different races and genders, reported in estimated population standard deviations, with 70% and 95% confidence intervals clustered by real the application dossier used to generate the synthetic application. Positive values indicate that the model rates female or racial minority applicants higher than male or White applicants, respectively, on average.*

Here, the blue vertical line represents the estimated race- or gender-effect in the original evaluation task, along with 70% and 95% confidence intervals. As can be seen, estimated effects vary across conditions, but only slightly beyond what would be expected on the basis of estimation error alone.

To test the sensitivity to our results to variations in the evaluation task itself, we also conduct the following variants of our main experiment, again using GPT-3.5:

- **“No Scratch”**: We modify the original task, removing the elicitation of a summary of the candidate’s qualifications.
- **“No Transcripts”**: We elicit hiring recommendations from the model as in the original task, but omit the interview transcripts from the input to the model.
- **“Other District”**: In applicant responses to interview questions and the description of the evaluation task itself, we substitute in place of the name of the actual school district to which

applicants applied (and minor variations of the name that commonly occur in the transcripts) the name of an alternate, mostly White school district in West Virginia. We also replace the city and state where the actual school district is located.

- **“EEOC Guidance”:** We include a brief instruction to adhere to the Equal Opportunity Employment Commission’s guidelines on disparate impact and disparate treatment in the original task description.

The results are shown in Figure D.4. As in the case of variations in the wording of the prompt, the estimated race- and gender-effects vary modestly, but in many cases within the range we would expect from estimation error alone.

D.4.3 Impact of Blinding

Blinding is a natural strategy for mitigating race- and gender-effects in model evaluations. To understand the impact of blinding, we repeat the adverse impact analysis in the main text, substituting résumés and transcripts from which name, pronouns, and other obvious indicators of race and gender have been removed for the unmodified résumé text and transcripts. We find that while model evaluations do differ, the adverse impact ratio would change relatively little compared to the adverse impact ratio if the model evaluated unblinded application materials.

D.4.4 Prediction of Race and Gender from Redacted Materials

Redacting an applicant’s listed name, pronouns, title, and college removes some information about their race and gender from the application materials; however, it does not remove all available information. For instance, an applicant may have held a job that is highly correlated with gender, or speak a dialect of English strongly associated with a certain race or ethnic group. To test how much information about race and gender the redacted and unredacted application materials contain, using both the redacted and unredacted application materials for each candidate, we follow the following steps. First, we divide the materials into two pieces: an applicant’s résumé, and the transcript of their interview. We then embed both pieces into 256 dimensions (512 dimensions total) using OpenAI’s `text-embedding-large`. Next, we split half the data into a training set which we use to




Figure D.5: *Adverse impact ratios for $GPT-3.5$ hiring recommendations at different hiring thresholds, with pivotal 95% bootstrapped confidence intervals, when the model is presented with redacted application materials. The results are substantially similar to those in the main text: we find near parity at the lowest threshold, but some evidence of disparities at higher thresholds.*

train a penalized logistic regression model to predict applicants’ races and genders using (1) just the embedding of the résumé, (2) just the embedding of the transcripts, and (3) both embeddings. Finally, we calculate the out-of-sample AUC using 10-fold cross-validation on the held-out testing data. Comparing between redacted and unredacted application materials, we find that there is little difference in the AUC for predicting race and gender whether or not redacted materials are used; see Figure D.6.

D.4.5 Additional Adverse Impact Ratio Analyses

We repeat the adverse impact analysis in the main text, shown in Figure 5.1, across all models. The results are shown in Figures D.7 through D.10. The adverse selection ratios vary substantially across models, in large part because different models exhibit different levels of “severity” (i.e., tend to give generally higher or lower scores); however, almost all of them would potentially raise issues




Figure D.6: *Out-of-sample AUC for predicting race and gender from redacted and unredacted application materials. The AUC is generally very high and is not altered significantly even when name, pronouns, and other obvious indicators of race and gender have been removed.*

under the EEOC “four-fifths” rule for at least some selection thresholds.

D.4.6 Alternative Analysis of Name Effects

To understand the potential impact of the particular names we use, over and above the effects of the demographic groups they signal, we re-analyze our results using a mixed-effects model that includes random effects for the particular name used in each synthetic application. We find that in all three of our primary analyses (across models, across framings, and across prompt variants), the effects of particular names are small and our estimates of the difference in ratings between synthetic White and Black, Hispanic, and Asian applicants and male and female applicants are largely unchanged.

D.4.7 Selection Ratios

To understand how the differences in mean model scores between synthetic applicants belonging to different races and genders shown in Figure 5.2 translate into differences in actual hiring outcomes, we calculate the difference in the proportion of applicants hired at each possible threshold (i.e., 2, 3, 4, or 5) across races and genders. The results are shown in Figures D.14 and D.15. We find that the differences in mean model scores correspond to differences in hiring rates on the order of a few percentage points, with the largest disparities occurring at the highest hiring thresholds.

D.5 Models

In our experiments, we used the following model versions through the OpenAI API and on the AWS Bedrock service, current as of July 27, 2024:

- OpenAI
 - **GPT-3.5:** gpt-3.5-turbo-0125
 - **GPT-4:** gpt-4-0125-preview
 - **GPT-4o:** gpt-4o-2024-05-13
 - **GPT-4o Mini:** gpt-4o-mini-2024-07-18
- Anthropic
 - **Claude Instant:** claude-instant-v1
 - **Claude 2:** claude-v2:1
 - **Claude 3 Sonnet:** claude-3-sonnet-20240229-v1:0
 - **Claude 3.5 Sonnet:** claude-3-5-sonnet-20240620-v1:0
 - **Claude 3 Haiku:** claude-3-haiku-20240307-v1:0
- Mistral
 - **Mistral 7B:** mistral-7b-instruct-v0:2
 - **Mixtral 8x7B:** mixtral-8x7b-instruct-v0:1

Model	Error Rate
GPT-3.5	0.0%
GPT-4	0.0%
GPT-4o Mini	0.0%
GPT-4o	0.0%
Claude 3 Sonnet	0.1%
Claude 3 Haiku	0.1%
Claude 3.5 Sonnet	0.2%
Claude 2	0.8%
Claude Instant	1.5%
Mistral 7B	1.7%
Mixtral 8x7B	25.7%

Table D.1: Proportion of unparseable responses by model.

In all cases, we used the default settings for the models, including the default temperature for token sampling. Preliminary experiments during the design phase indicated that within-applicant variance in ratings stemming from non-determinism was around 0.5 points on the 1–5 scale, meaning that our planned experiments would be well-powered to detect differences of 0.1 points.

We elicit model responses as structured JSON objects. Because different models correctly respond with structured JSON when prompted to do so at different rates, we use GPT-4o Mini as a “bridge” model to convert the responses of the other models into valid JSON to facilitate analysis. A manual check of 100 randomly selected responses did not reveal any errors in this conversion process.

However, we find that the models differ in their ability to produce valid responses to the prompt. As shown in Table D.1, the most advanced models essentially never produce unparseable responses—i.e., responses which do not contain hiring ratings or from which GPT-4o Mini is otherwise unable to extract valid structured data—while less advanced models do so in a small fraction of cases, with the exception of Mixtral 8x7B, which produces unparseable responses in nearly a quarter of cases.

The proportion of unparseable responses for each race and gender group is shown in Figures D.16 and D.17. The occurrence of unparseable responses appears to be uncorrelated with our dependent variables of interest, and, because of their low frequency, can safely be treated as missing at random. The exception is Mixtral 8x7B, which produces unparseable responses 1.5 (± 0.8) percentage points more often for female synthetic applicants than male applicants, 1.5 (± 1.1) percentage points more often for Black applicants than White applicants, and 6.3 (± 1.1) percentage points more often for Asian applicants than White applicants.




Figure D.7: Adverse gender impact ratios for LLM hiring recommendations at different hiring thresholds, with pivotal 95% bootstrapped confidence intervals, when the model is presented with unredacted application materials, for all the models we study. Results vary substantially across models, but almost all would potentially raise issues under the EEOC “four-fifths” rule for at least some selection thresholds. For visual clarity, we truncate the y-axis at 2.5, though some error bars extend beyond this range.




Figure D.8: *Adverse racial impact ratios for LLM hiring recommendations at different hiring thresholds, with pivotal 95% bootstrapped confidence intervals, when the model is presented with unredacted application materials, for all the models we study. Results vary substantially across models, but almost all would potentially raise issues under the EEOC “four-fifths” rule for at least some selection thresholds. For visual clarity, we truncate the y-axis at 1.7, though some error bars extend beyond this range. **Claude Instant** does not give any minority candidate a 5-point rating, and so the adverse impact ratio is zero at that threshold.*




Figure D.9: Adverse gender impact ratios for LLM hiring recommendations at different hiring thresholds, with pivotal 95% bootstrapped confidence intervals, when the model is presented with redacted application materials, for all the models we study. Results vary substantially across models, but almost all would potentially raise issues under the EEOC “four-fifths” rule for at least some selection thresholds. For visual clarity, we truncate the y-axis at 2.5, though some error bars extend beyond this range.




Figure D.10: Adverse racial impact ratios for LLM hiring recommendations at different hiring thresholds, with pivotal 95% bootstrapped confidence intervals, when the model is presented with redacted application materials, for all the models we study. Results vary substantially across models, but almost all would potentially raise issues under the EEOC “four-fifths” rule for at least some selection thresholds. For visual clarity, we truncate the y-axis at 1.7, though some error bars extend beyond this range. *Claude Instant* does not give any minority candidate a 5-point rating, and so the adverse impact ratio is zero at that threshold.




Figure D.11: Differences in mean applicant ratings across LLMs between synthetic applicants of different races and genders, adjusted for the particular name used to signal that race and gender, and reported in estimated population standard deviations, with 70% and 95% confidence intervals clustered by the real application dossier used to generate the synthetic application. Positive values indicate that the model rates female or racial minority applicants higher than male or White applicants, respectively, on average.




Figure D.12: Differences in mean GPT-3.5 applicant ratings across variations in the wording of the prompt between synthetic applicants of different races and genders, adjusted for the particular name used to signal that race and gender, and reported in estimated population standard deviations, with 70% and 95% confidence intervals clustered by real the application dossier used to generate the synthetic application. Positive values indicate that the model rates female or racial minority applicants higher than male or White applicants, respectively, on average. The blue vertical line represents the estimated effect in the original evaluation task, along with 70% and 95% confidence intervals.




Figure D.13: Differences in mean *GPT-3.5* applicant ratings across variations in the prompt and context between synthetic applicants of different races and genders, adjusted for the particular name used to signal that race and gender, and reported in estimated population standard deviations, with 70% and 95% confidence intervals clustered by real the application dossier used to generate the synthetic application. Positive values indicate that the model rates female or racial minority applicants higher than male or White applicants, respectively, on average.




Figure D.14: Differences in the percentage of synthetic applicants hired at each possible hiring threshold across LLMs between synthetic applicants of different genders, with 70% and 95% confidence intervals. The y-axis represents percentage point differences in hiring rates. Differences are relative to synthetic male applicants.




Figure D.15: Differences in the percentage of synthetic applicants hired at each possible hiring threshold across LLMs between synthetic applicants of different races, with 70% and 95% confidence intervals. The y-axis represents percentage point differences in hiring rates. Differences are relative to synthetic White applicants.




Figure D.16: *Proportion of unparseable responses by model and race of synthetic applicant.*




Figure D.17: *Proportion of unparseable responses by model and gender of synthetic applicant.*

Appendix to Chapter 6

E.1 Mathematical Appendix

We structure the mathematical theory of the robust outcome test in six parts. First, in Appendix E.1.1, to provide intuition for the general result, we prove Theorem 6.1. Then, in Appendix E.1.2, we extend the formal problem in the main text to a more general setting that accommodates arbitrary utilities and quasi-rational decision makers before stating Theorem E.2, a generalization of Theorem 6.1. In Appendix E.1.3, we give an overview of important properties of risk-decision curves, a key theoretical object introduced in our general setting; in particular, we show that risk-decision curves are well-defined and that the notion of generation, introduced in Appendix E.1.2, is also well-defined and non-vacuous. Next, we revisit important properties of stochastic orderings in Appendix E.1.4, including stochastic dominance and the MLRP. We prove Theorem E.2 in Appendix E.1.5. Finally, in Appendix E.1.6, we give an overview of how to account for sampling error in the robust outcome test.

E.1.1 Proof of Theorem 6.1

We begin by noting that Theorem 6.1 cannot hold without *some* hypothesis on the risk distributions. Figure E.1 illustrates one possible failure of the robust outcome test in our running lending example. Here, there is no discrimination, because lending decisions are made according to a uniform threshold. However, lending rates are lower and repayment rates higher for the blue group. One intuitive way of capturing the issue is that it is easier for loan officers to determine whether an applicant will default in the blue group than in the red group because the risk distribution of the blue group has higher variance than that of the red group. The MLRP formalizes and generalizes this intuition,

capturing the key properties needed for the robust outcome test to be correct.

First, we note that the MLRP is equivalent to the—in our setting, more intuitive—monotone conditional probability (MCP) condition stated in Theorem 6.1. Suppose that the conditional distributions of $R \mid G = g$ have positive densities on $(0, 1)$ given by $f_{R \mid G=g}(r)$. The MLRP simply states that the likelihood ratio

$$\frac{f_{R \mid G=1}(r)}{f_{R \mid G=0}(r)}$$

is a monotonic function of r . Recalling the monotonicity condition in Theorem 6.1, observe that

$$g(r) \stackrel{\text{def}}{=} \Pr(G = 1 \mid R = r) = \frac{p \cdot f_{R \mid G=1}(r)}{(1-p) \cdot f_{R \mid G=0}(r) + p \cdot f_{R \mid G=1}(r)},$$

where $p \stackrel{\text{def}}{=} \Pr(G = 1)$. Note that if

$$h(q) \stackrel{\text{def}}{=} \frac{1-p}{p} \cdot \frac{q}{1-q},$$

then

$$(h \circ g)(r) = \frac{f_{R \mid G=1}(r)}{f_{R \mid G=0}(r)},$$

i.e., the likelihood ratio. As a consequence, since $h(q)$ is monotonically increasing on $(0, 1)$, the MLRP holds in this case if and only if $\Pr(G = 1 \mid R = r)$ is a monotonic function of r .

To understand how the MLRP connects to the proof of Theorem 6.1, let $g^{\text{lwr}} \in \{0, 1\}$ denote the lower “base rate” group and $g^{\text{upr}} \in \{0, 1\}$ the higher base rate group, i.e., g^{lwr} and g^{upr} are such that

$$\Pr(Y = 1 \mid G = g^{\text{lwr}}) \leq \Pr(Y = 1 \mid G = g^{\text{upr}}).$$

(For avoidance of doubt, “base rate” refers throughout to the unconditional outcome rate of a group, not the prevalence of a group.) Satisfying the MLRP implies the following two useful properties. First, the risk-distribution of group $G = g^{\text{upr}}$ is “right-shifted” relative to that of group $G = g^{\text{lwr}}$, i.e.,

$$\Pr(R \geq r \mid G = g^{\text{upr}}) \geq \Pr(R \geq r \mid G = g^{\text{lwr}}) \quad \text{for all } r \in [0, 1]; \quad (\text{E.1})$$

that is, the MLRP implies that the distributions also satisfy stochastic dominance. Secondly, the lower base rate group retains a lower outcome rate even after conditioning on risk being above some

threshold $t < 1$. More specifically,¹

$$\Pr(Y = 1 \mid G = g^{\text{low}}, t \leq R) \leq \Pr(Y = 1 \mid G = g^{\text{up}}, t \leq R). \quad (\text{E.2})$$

See Theorems 1.C.1 and 1.C.5 in Shaked and Shanthikumar (2007) for proof of these properties.

This is enough to allow us to prove Theorem 6.1.

Proof of Theorem 6.1. We proceed by proving the contrapositive: if $t_1 \leq t_0$, then either the decision rate of group $G = 0$ will be no larger than that of group $G = 1$, or the outcome rate will be no smaller, i.e.,

$$\Pr(D = 1 \mid G = 0) \leq \Pr(D = 1 \mid G = 1) \quad (\text{E.3})$$

or

$$\Pr(Y = 1 \mid G = 0, D = 1) \geq \Pr(Y = 1 \mid G = 1, D = 1). \quad (\text{E.4})$$

We will show that, depending on whether $G = 0$ is the lower base rate group or the higher base rate group, either the benchmark test in (E.3) or the outcome test in (E.4), respectively, will point in the correct direction.

Suppose that group $G = 0$ has the *lower* base rate. (E.1) implies that

$$\Pr(R \geq t_0 \mid G = 1) \geq \Pr(R \geq t_0 \mid G = 0).$$

Furthermore, reducing the decision threshold from t_0 to t_1 can only increase the decision rate for group $G = 1$. Thus, in this case, the decision rate for group $G = 0$ cannot exceed that of group $G = 1$, i.e.,

$$\Pr(D = 1 \mid G = 1) = \Pr(R \geq t_1 \mid G = 1) \geq \Pr(R \geq t_0 \mid G = 0) = \Pr(D = 0 \mid G = 0),$$

showing that (E.3) holds.

On the other hand, suppose that group $G = 0$ has the *higher* base rate. The outcome test looks at

¹In general, the MLRP implies—and is equivalent to—uniform conditional stochastic dominance; here we derive Eq. (E.2) from uniform conditional stochastic dominance using the fact that, by the law of iterated expectations and the definition of R ,

$$\Pr(Y = 1 \mid G = g, t \leq R) = \mathbb{E}[\Pr(Y = 1 \mid X) \mid G = g, t \leq R] = \mathbb{E}[R \mid G = g, t \leq R].$$

the outcome rates of the groups *after* conditioning on receiving a positive decision. More specifically,

$$\Pr(Y = 1 \mid G = g, D = 1) = \Pr(Y = 1 \mid G = g, t_g \leq R)$$

by the definition of D in (6.2). Now, since $\Pr(D = 1 \mid G = g) > 0$, it follows that $t_g < 1$ for $g \in \{0, 1\}$. Therefore, by (E.2), we have that

$$\Pr(Y = 1 \mid G = 0, t_0 \leq R) \geq \Pr(Y = 1 \mid G = 1, t_0 \leq R).$$

Again, lowering the decision threshold from t_0 to t_1 only reduces the outcome rate for group $G = 1$, i.e.,

$$\Pr(Y = 1 \mid G = 0, D = 1) \geq \Pr(Y = 1 \mid G = 1, D = 1),$$

showing that (E.4) holds, and completing the proof. $\square$

In proving Theorem 6.1, the key insight is that, under the MLRP, whether group $G = 0$ is the lower or higher base rate group, either the benchmark or the standard outcome test will correctly detect the relative ordering of t_0 and t_1 —though we do not know which one. As a result, when both tests indicate discrimination against the same group, the conclusion is unambiguous. Theorem E.2 below extends this argument to a much more general setting. The chief technical obstacles there are: (1) showing that the standard outcome test still points in the right direction for the higher base rate group, and (2) accounting for quasi-rational decision makers and more complex risk distributions without densities.

E.1.2 General Utilities and Quasi-Rational Decision Makers

The correctness of the robust outcome test holds in a more general setting than the one presented in Section 6.1 that allows for both quasi-rational decision makers, as well as more general bases for their decisions. The general theorem and its assumptions are most naturally presented in the language of measure theory, which we adopt throughout.

We again imagine a population of individuals belonging to one of two groups $G \in \{0, 1\}$. To avoid trivialities, we assume that neither group is empty, i.e.,

$$\Pr(G = g) > 0 \quad \text{for } g \in \{0, 1\}. \tag{E.5}$$

As above, we assume that decision makers make binary decisions $D \in \{0, 1\}$ for each individual. We also assume that a non-zero proportion of individuals in each group receives decision $D = 1$, i.e.,

$$\Pr(D = 1 \mid G = g) > 0 \quad \text{for } g \in \{0, 1\}. \quad (\text{E.6})$$

In our running lending example, loan officers base decisions on an applicant's probability of repayment, i.e., on $U \stackrel{\text{def}}{=} \Pr(Y = 1 \mid X)$, where $Y = 1$ denotes repayment of the loan and X denotes loan applicants' observable features when lending decisions occur. In a more realistic model, instead of directly on $\Pr(Y = 1 \mid X)$, loan officers might instead base their decisions on a calibrated risk score U . We can further generalize to an arbitrary utility U , dissociating from a particular outcome Y and covariates X . For example, U could represent the lender's expected return on lending to an applicant, rather than just the applicant's probability of default. However, for consistency with Section 6.1, we still refer to distributions of U as *risk* distributions. We assume that the expectation of U is well-defined, i.e.,

$$\mathbb{E}[|U|] < \infty. \quad (\text{E.7})$$

To ensure that decisions can be compared across groups, we assume the following analogue of the common overlap assumption (Rosenbaum and Rubin, 1983b):

$$0 < \Pr(G = g \mid U) < 1 \quad \text{a.s.} \quad (\text{E.8})$$

In contrast to the usual setting in the outcome test literature, we do not require there to be a single decision maker—or, more generally, a collection of identical decision makers—making decisions based on group-specific thresholds t_g , $g \in \{0, 1\}$. Instead, we capture the decision process through *risk-decision curves*:

$$d_g(u) \stackrel{\text{def}}{=} \Pr(D = 1 \mid U = u, G = g). \quad (\text{E.9})$$

The risk-decision curves encode the proportion of individuals who receive a positive decision among those who belong to group $G = g$ with utility $U = u$. We can ask whether one group receives positive decisions more frequently at every level of utility, representing a double-standard. Depending on whether a positive decision is “desirable,” as in lending, or “undesirable,” as in policing, we understand $d_g(u) < d_{g'}(u)$ as either discrimination or preferential treatment.

In Section 6.1, we assumed particularly simple risk-decision curves of the following form:

$$d_g(u) = \mathbf{1}(u \geq t_g),$$

where $\mathbf{1}(\cdots)$ denotes an indicator function. Here we generalize beyond threshold rules to cases in which the decision makers collectively exhibit a form of bounded rationality.

Definition E.1 (Bounded Rationality). We say that the risk-decision curve $d_g(u)$ exhibits *bounded rationality* when it is right-continuous and non-decreasing U -a.s.²

The definition of bounded rationality itself is very general, although the proof of our main theorem requires an additional restriction defined below on the risk-decision curves. In practice, we expect risk-decision curves to be continuous; however, requiring only right continuity allows for the possibility that there are thresholds where decision makers have a discontinuous increase in their probability of choosing $D = 1$, as, e.g., would be the case for rational decision makers at the threshold t_g .

Since it is a.s. bounded between zero and one, a risk-decision curve $d_g(u)$ exhibiting bounded rationality can be seen as the cumulative distribution function (CDF) of a distribution H_g on the extended real numbers $\bar{\mathbb{R}} = \mathbb{R} \cup \{-\infty, \infty\}$ where

$$\Pr(H_g = -\infty) = \lim_{u \rightarrow -\infty} d_g(u), \quad \Pr(H_g \leq t) = d_g(t), \quad \text{and} \quad \Pr(H_g = \infty) = \lim_{u \rightarrow \infty} 1 - d_g(u). \quad (\text{E.10})$$

We say that $d_g(u)$ *generates* H_g . Because $d_g(u)$ is defined only U -a.s., $d_g(u)$ may not generate a *unique* distribution H_g . But, for our purposes, the different generated distributions are largely interchangeable, and so we will often refer to *the* generated distribution H_g . (For the well-definedness of H_g and related considerations, see Appendix E.1.3 below.)

The final condition we require is that the risk-decision curves generate some pair of distributions satisfying the MLRP. This property holds if, e.g., the risk-decision curves are threshold rules. (See Appendix E.1.4 below for discussion of the MLRP and non-continuous distributions.) We now state the general version of our main result.

Theorem E.2. *Suppose that the following two conditions hold:*

²By “ U -a.s.,” we mean that a property holds for all $u \in \mathbb{R} \setminus S$ where $\Pr(U \in S) = 0$.

- $\Pr(G = 1 \mid U = u)$ is U -a.s. monotone,
- The risk-decision curves $d_g(u)$ for $g \in \{0, 1\}$ generate distributions satisfying the MLRP.

Under these conditions, if

$$\Pr(D = 1 \mid G = 0) > \Pr(D = 1 \mid G = 1),$$

and

$$\mathbb{E}[U \mid D = 1, G = 0] < \mathbb{E}[U \mid D = 1, G = 1],$$

then $d_0(u) \geq d_1(u)$ U -a.s., where the equality is strict with positive probability.

We defer the proof to Appendix E.1.5. As with threshold decision rules, the robust outcome test holds in this more general setting even under the modest violations of the MLRP that we see in practice; see Appendix E.3 for the results of a simulation study analogous to the one presented in Section 6.2.2 above.

Theorem E.2 assumes that the risk-decision curves generate distributions satisfying the MLRP. This property holds in a variety of settings, including when decisions are made according to risk thresholds t_g as in Section 6.1. It is also satisfied by logistic risk-decision curves of the following form:

$$d_g(u) = \frac{1}{1 + \exp(\lambda \cdot [t_g - u])}. \quad (\text{E.11})$$

Here, as in the bounded rationality literature (e.g., McKelvey and Palfrey, 1995), $\lambda > 0$ represents the decision makers' degree of "rationality," with $\lambda \rightarrow \infty$ recovering threshold decision rules in the limit; and t_g represents a "soft" threshold at which decision makers become more likely to make decision $D = 1$ than $D = 0$. Many other families of risk-decision curves satisfy the MLRP, such as the CDFs of normal distributions with the same variance or beta distributions with the same total count. More generally, the CDFs of log, logit, or other monotonic transformations of normal, beta, or gamma distributions whose densities cross once satisfy the MLRP. (See Appendix E.1.4 below for more detailed discussion of families of distributions satisfying the MLRP, and Figure E.4 for an illustration of logit-normal CDF risk-decision curves satisfying the MLRP.) Intuitively, risk-decision curves satisfy our assumption when, as in (E.11), decision makers are similarly sensitive to risk across groups—a property that is strictly weaker than the "rationality" generally assumed in the

outcome test literature.

E.1.3 Risk-Decision Curves

We begin by noting that risk-decision curves are well-defined. Recall the definition in (E.9):

$$d_g(u) \stackrel{\text{def}}{=} \Pr(D = 1 \mid U = u, G = g).$$

The risk-decision curves $d_g(u)$ exist and are well-defined U -a.s. by the Doob-Dynkin lemma because of the non-triviality assumptions in Eqs. (E.5) and (E.6) and the overlap assumption in (E.8).

The following lemma shows that the definition of generation is not vacuous.

Lemma E.3. *Any risk-decision curve $d_g(u)$ exhibiting bounded rationality generates some distribution H_g .*

Proof. First, we recall that $d_g(u)$ is non-decreasing, bounded between zero and one, and right-continuous U -a.s. In other words, there exists a U -null set S such that these properties hold for all $u \in \mathcal{R} \stackrel{\text{def}}{=} \mathbb{R} \setminus S$. For $k \in \{1, \dots, 2^n\}$, let

$$\ell(k, n) \stackrel{\text{def}}{=} \inf \{ u \in \mathcal{R} : d_g(u) \geq k \cdot 2^{-n} \}, \quad (\text{E.12})$$

and define $F_n : \mathbb{R} \rightarrow [0, 1]$ as follows:

$$F_n(u) \stackrel{\text{def}}{=} \frac{1}{2^n} \sum_{k=1}^{2^n} \mathbf{1}(\ell(k, n) \leq u). \quad (\text{E.13})$$

The function $F_n(u)$ approximates $d_g(u)$. In particular, for all $u \in \mathcal{R}$,

$$0 \leq d_g(u) - F_n(u) < 2^{-n}. \quad (\text{E.14})$$

To see why, suppose that $u^* \in \mathcal{R}$ is arbitrary and satisfies

$$(k-1) \cdot 2^{-n} \leq d_g(u^*) < k \cdot 2^{-n}.$$

Then it immediately follows from (E.12) that $\ell(k-1, n) \leq u^*$. (If $k = 1$, then $\ell(k-1, n)$ is, strictly speaking, not defined by (E.12), but the proof differs only in one minor detail noted below.)

Separately, by right continuity, there exists $\delta > 0$ such that for all $u \in [u^*, u^* + \delta) \cap \mathcal{R}$,

$$d_g(u) < k \cdot 2^{-n}.$$

By monotonicity, for all $u \in \mathcal{R}$ less than or equal to u^* ,

$$d_g(u) \leq d_g(u^*) < k \cdot 2^{-n}.$$

Therefore, for all $u \in (-\infty, u^* + \delta) \cap \mathcal{R}$,

$$d_g(u) < k \cdot 2^{-n},$$

and so

$$u^* < u^* + \delta \leq \inf \{ u \in \mathcal{R} : d_g(u) \geq k \cdot 2^{-n} \} = \ell(k, n).$$

In particular, it follows that

$$\{ u \in \mathcal{R} : (k-1) \cdot 2^{-n} \leq d_g(u) < k \cdot 2^{-n} \} = [\ell(k-1, n), \ell(k, n)) \cap \mathcal{R}.$$

(If $k = 1$, simply replace the half-open interval with the open interval $(-\infty, \ell(1, n))$.) Since $F_n(u)$ is exactly equal to $(k-1) \cdot 2^{-n}$ on the latter set, (E.14) follows.

Now, since $\ell(k, n) = \ell(2^m \cdot k, n + m)$, it also follows from the definition of $F_n(u)$ in (E.13) that for all $n_0, n_1 \in \mathbb{N}$ and $u \in \mathbb{R}$,

$$|F_{n_0}(u) - F_{n_1}(u)| < 2^{-\min(n_0, n_1)}.$$

That is, the sequence $(F_n(u))_{n \in \mathbb{N}}$ converges uniformly on all of $\mathbb{R}$. Since $F_n(u)$ is non-decreasing, bounded between zero and one, and right-continuous for all n , it follows that the pointwise limit has these properties as well. That is,

$$F_{H_g}(u) \stackrel{\text{def}}{=} \lim_{n \rightarrow \infty} F_n(u)$$

is non-decreasing, bounded between zero and one, and right-continuous. Therefore $F_{H_g}(u)$ is the CDF of a random variable H_g .

However, using (E.14), we also have that

$$\sup |F_{H_g}(u) - d_g(u)| = \sup \lim_{n \rightarrow \infty} |F_n(u) - d_g(u)| \leq \sup \lim_{n \rightarrow \infty} 2^{-n} = 0,$$

where the suprema are taken over $u \in \mathcal{R}$. Therefore $F_{H_g}(u) = d_g(u)$ for all $u \in \mathcal{R}$. Since $\Pr(U \in \mathcal{R}) = 1$, it follows that $d_g(u) = F_{H_g}(u)$ U -a.s., i.e., $d_g(u)$ generates the distribution of H_g . $\square$

In light of Lemma E.3, we note two important points about risk decision curves. First, a risk-decision curve can potentially generate more than one distribution. In particular, CDFs are defined on all of $\mathbb{R}$, whereas risk-decision curves are only defined U -a.s. Therefore, if U is not supported on all of $\mathbb{R}$, it may be the case that there is not a unique distribution H_g satisfying

$$d_g(u) = \Pr(H_g \leq u) \quad U\text{-a.s.} \quad (\text{E.15})$$

In particular, $d_g(u)$ does not determine the distribution of H_g on regions outside the support of U . The functions $f(\cdot)$ we consider in the proof of Theorem E.2 are constant on regions outside the support of U , however, meaning that the distributions of $f(H_g)$ and $f(\tilde{H}_g)$ are nevertheless identical for distinct H_g and $\tilde{H}_g$ generated by $d_g(u)$. Consequently, the different distributions $d_g(u)$ generates are, for our purposes, interchangeable. For this reason, we will on occasion refer to “the” distribution generated by $d_g(u)$, despite the ambiguity. (In particular, it is enough in Theorem E.2 that the risk-decision curves generate *any* pair of distributions satisfying the MLRP.)

Second, as noted in (E.10), a generated distribution H_g is not necessarily a.s. finite. Throughout, we do not assume that *any* random variable is a.s. finite unless explicitly mentioned, as was the case for D , G , and U in Appendix E.1.2 above.

E.1.4 Stochastic Orderings

As noted in Section 6.1, some hypotheses on the risk distributions are needed in Theorem E.2 to rule out scenarios like the one shown in Figure E.1. These hypotheses take the form of stochastic ordering relations, which we review here. The simplest and most important stochastic ordering is stochastic dominance.

Definition E.4 (Stochastic Dominance). Let F_0 and F_1 be arbitrary CDFs. We say that F_0 (first-order) *stochastically dominates* F_1 , written $F_0 \succeq_1 F_1$, if for any $t \in \mathbb{R}$, $F_0(t) \leq F_1(t)$. For

random variables X and Y , we write $X \succeq_1 Y$ if $F_X \succeq_1 F_Y$, i.e., if

$$\Pr(X \leq t) \leq \Pr(Y \leq t) \quad \text{for all } t \in \mathbb{R}. \quad (\text{E.16})$$

We say that a pair of random variables or distributions is *stochastically ordered* if one stochastically dominates the other.

Stochastic dominance has the useful property that increasing functions of the dominating random variable have greater expectation.

Lemma E.5. *$X \succeq_1 Y$ if and only if $\mathbb{E}[f(X)] \geq \mathbb{E}[f(Y)]$ for any non-decreasing $f(x)$ for which the expectations are well-defined.*

For proof of Lemma E.5, see, e.g., Shaked and Shanthikumar (2007). (The statement in Shaked and Shanthikumar implicitly assumes that X and Y are a.s. finite, but the proof does not use this fact.)

To prove Theorem E.2, we will need the following strengthening of stochastic dominance.

Definition E.6 (Monotone Likelihood Ratio Property). Let F_0 and F_1 be arbitrary cumulative distribution functions. Let f_0 and f_1 be the Radon-Nikodym derivatives (i.e., densities) of F_0 and F_1 with respect to their sum measure μ .³ We say that F_0 and F_1 have the *monotone likelihood ratio property* (MLRP) with F_0 dominating, denoted $F_0 \succeq_{\text{lr}} F_1$, if

$$\frac{f_0(x)}{f_1(x)}$$

is non-decreasing μ -a.e., where, without loss of generality, we define the ratio to be ∞ when $f_1(x) = 0$.

For random variables X and Y , we write $X \succeq_{\text{lr}} Y$ if $F_X \succeq_{\text{lr}} F_Y$.

The MLRP has many useful consequences and appears widely in the literature. In addition to implying stochastic dominance (see Theorem 1.C.1 in Shaked and Shanthikumar, 2007, and below), the MLRP holds for many familiar parametric families used to model risk distributions. For example:

³I.e., for any interval $(a, b]$,

$$\mu[(a, b]] = F_0(b) + F_1(b) - (F_0(a) + F_1(a))$$

with μ suitably extended to all Borel sets by the Carathéodory extension theorem. Both F_0 and F_1 are absolutely continuous with respect to μ .

- For $X \sim \text{Beta}(\alpha_0, \beta_0)$ and $Y \sim \text{Beta}(\alpha_1, \beta_1)$, $X \succeq_{\text{lr}} Y$ if and only if $\alpha_0 \geq \alpha_1$ and $\beta_0 \leq \beta_1$;
- For $X \sim \mathcal{N}(\mu_0, \sigma_0^2)$ and $Y \sim \mathcal{N}(\mu_1, \sigma_1^2)$, $X \succeq_{\text{lr}} Y$ if and only if $\mu_0 \geq \mu_1$ and $\sigma_0 = \sigma_1$;
- For $X \sim \text{Gamma}(\alpha_0, \beta_0)$ and $Y \sim \text{Gamma}(\alpha_1, \beta_1)$, $X \succeq_{\text{lr}} Y$ if and only if $\alpha_0 \geq \alpha_1$ and $\beta_0 \leq \beta_1$;
- For $X \sim \text{Binom}(n, p_0)$ and $Y \sim \text{Binom}(n, p_1)$, $X \succeq_{\text{lr}} Y$ if and only if $p_0 \geq p_1$.

We omit the routine verification of these facts. In the case of beta, normal, and gamma distributions, the MLRP is also equivalent to the densities of the risk distributions intersecting once.

Lemma E.7. *Suppose $X \sim \text{Beta}(\alpha_0, \beta_0)$ and $Y \sim \text{Beta}(\alpha_1, \beta_1)$ have different distributions. Then X and Y satisfy the MLRP if and only if $f_X(p) = f_Y(p)$ for a unique $p \in (0, 1)$.*

The proof for normal and gamma distributions is similar.

Proof. We prove the lemma in three steps: First, we show that the densities must cross at least once. Next, we show that if the densities intersect more than once, then the MLRP does *not* hold. Finally, we show that if the densities intersect exactly once, then the MLRP *does* hold. We argue using the density ratio

$$r(p) \stackrel{\text{def}}{=} \frac{f_X(p)}{f_Y(p)} = \frac{B(\alpha_1, \beta_1)}{B(\alpha_0, \beta_0)} \cdot p^{\alpha_0 - \alpha_1} \cdot (1 - p)^{\beta_0 - \beta_1}. \quad (\text{E.17})$$

Note that, taking derivatives, $r(p)$ is strictly monotonic if $\alpha_0 - \alpha_1$ and $\beta_0 - \beta_1$ have opposite signs or one of them is zero. (They cannot both be zero since the densities differ by assumption.) If $\alpha_0 - \alpha_1$ and $\beta_0 - \beta_1$ are both negative or both positive, it follows from (E.17) that

$$\lim_{p \rightarrow 0} r(p) = \lim_{p \rightarrow 1} r(p) = \infty \quad \text{or} \quad \lim_{p \rightarrow 0} r(p) = \lim_{p \rightarrow 1} r(p) = 0, \quad (\text{E.18})$$

respectively. Therefore $r(p)$ is non-monotonic in this case. In particular, irrespective of the parameters, $r(p)$ is either non-monotonic or strictly monotonic.

We now show that the densities must cross at least once. That is, there must exist $p_0 < p_* < p_1$ such that $f_X(p_*) = f_Y(p_*)$ and

$$\text{sgn}(f_X(p_0) - f_Y(p_0)) = -\text{sgn}(f_X(p_1) - f_Y(p_1)) \neq 0.$$

If not, then, without loss of generality, $f_X(p) \leq f_Y(p)$ for all $p \in (0, 1)$, with strict inequality holding for some p . Thus, by continuity, it follows that $\int_0^1 f_Y(p) - f_X(p) dp > 0$. This contradicts the fact that

$$\int_0^1 f_Y(p) - f_X(p) dp = \int_0^1 f_Y(p) dp - \int_0^1 f_X(p) dp = 1 - 1 = 0.$$

Therefore, there must exist p_0 and p_1 such that $f_X(p_0) - f_Y(p_0)$ and $f_X(p_1) - f_Y(p_1)$ have opposite signs. By the intermediate value theorem, therefore, there exists some p_* between p_0 and p_1 such that $f_X(p_*) = f_Y(p_*)$, yielding a crossing.

Next, suppose that $f_X(q_i) = f_Y(q_i)$ for $0 < q_0 < q_1 < 1$, so that the densities intersect at least twice. Then

$$r(q_0) = \frac{f_X(q_0)}{f_Y(q_0)} = 1 = \frac{f_X(q_1)}{f_Y(q_1)} = r(q_1),$$

so $r(p)$ is not strictly monotonic, and hence by the argument above is non-monotonic. Therefore the MLRP does not hold by definition.

Finally, suppose that there is exactly one intersection at p_* . This intersection must therefore be a crossing. Assume without loss of generality that $r(p_0) < 1 < r(p_1)$. Then, by continuity, $r(p) < 1$ for all $p \in (0, p_*)$ and $r(p) > 1$ for all $p \in (p_*, 1)$. It follows immediately that $\alpha_0 - \alpha_1$ and $\beta_0 - \beta_1$ cannot both be positive, since then by (E.18), $r(0) = r(1) = 0$, and so by the intermediate value theorem there would be $p > p_*$ such that $r(x) < 1$. Similarly, they cannot both be negative. Therefore $\alpha_0 - \alpha_1$ and $\beta_0 - \beta_1$ have opposite signs, or one of them is zero, meaning that $r(p)$ is monotonic, i.e., that the MLRP holds. $\square$

The MLRP is also preserved by monotone transformations: If $X \succeq_{\text{lr}} Y$, then $g(X) \succeq_{\text{lr}} g(Y)$ for monotonically increasing functions $g(x)$; see Theorem 2.5 in Keilson and Sumita (1982). Thus, the MLRP also holds, e.g., with the analogous statements for log-normal, log-gamma, and logit-normal distributions. Like stochastic dominance, the MLRP is preserved when taking sums of random variables, so long as the densities are log-concave; see Lemma 1.1 in Shanthikumar and Yao (1986). For further examples of conditions under which the MLRP holds, see Section 1.C of Shaked and Shanthikumar (2007). In summary, the range of parametric families for which the MLRP holds is considerable.

In addition, the MLRP has many different equivalent formulations, some of which we note below.

Lemma E.8. *Let X and Y be arbitrary random variables. Let $Z \sim \text{Bernoulli}(p)$ for $p > 0$ be independent of X and Y , and let $W \stackrel{\text{def}}{=} Z \cdot X + (1 - Z) \cdot Y$. Then, the following are equivalent:*

1. *Monotone Likelihood Ratio Property (MLRP): The distributions of X and Y have the MLRP, i.e., $X \succeq_{\text{lr}} Y$;*
2. *Uniform Conditional Stochastic Dominance (UCSD): For all events E such that $\Pr(X \in E) > 0$ and $\Pr(Y \in E) > 0$,*

$$[X \mid E] \succeq_1 [Y \mid E],$$

i.e., the distribution of X conditional on E stochastically dominates the distribution of Y conditional on E .

3. *Klinostochastic Dominance (KSD):⁴*

$$\frac{\int_{\mathbb{R}} \mathbf{1}(x \leq t) \cdot h(x) dF_X(x)}{\int_{\mathbb{R}} h(x) dF_X(x)} \leq \frac{\int_{\mathbb{R}} \mathbf{1}(y \leq t) \cdot h(y) dF_Y(y)}{\int_{\mathbb{R}} h(y) dF_Y(y)}, \quad (\text{E.19})$$

for all $t \in \mathbb{R}$ and any $h(x)$ such that $h(X)$ and $h(Y)$ are a.s. non-negative and such that

$$0 < \mathbb{E}[h(X)] < \infty, \quad 0 < \mathbb{E}[h(Y)] < \infty.$$

4. *Monotone conditional probability (MCP): The conditional probability*

$$g(w) \stackrel{\text{def}}{=} \Pr(Z = 1 \mid W = w)$$

is monotone W -a.s.

Proof. The equivalence between the MLRP and UCSD was originally given as Theorems 1.1 and 1.3 in Whitt (1980)—though, see the note to Theorem 8 in Rüschendorf (1991). The equivalence between the MLRP and KSD was first shown in Theorems 7 and 8 in Rüschendorf (1991). (As above, the statements in these references assume that X and Y are a.s. finite, but the proofs do not use this fact.)

We now show the equivalence between MCP and the MLRP. The proof closely follows the sketch of the equivalence for distributions with densities in the proof of Theorem 6.1 above. Let $p = \Pr(Z = 1)$.

⁴This is the $\mathbb{B}^1$ -ordering in the notation of Rüschendorf (1991).

We begin by observing that if

$$\begin{aligned}
0 &< \Pr(W \in E) \\
&= p \cdot \Pr(W \in E \mid Z = 1) + (1 - p) \cdot \Pr(W \in E \mid Z = 0) \\
&= p \cdot \Pr(X \in E) + (1 - p) \cdot \Pr(Y \in E),
\end{aligned}$$

then

$$0 < \Pr(X \in E) + \Pr(Y \in E)$$

and *vice versa*. Therefore, the sum measure μ and the measure induced by F_W are mutually absolutely continuous. Consequently, by the chain rule for Radon-Nikodym derivatives,

$$\frac{f_X}{f_Y} = \frac{\frac{dF_X}{d\mu}}{\frac{dF_Y}{d\mu}} = \frac{\frac{dF_X}{d\mu} \cdot \frac{d\mu}{dF_W}}{\frac{dF_Y}{d\mu} \cdot \frac{d\mu}{dF_W}} = \frac{\frac{dF_X}{dF_W}}{\frac{dF_Y}{dF_W}}. \quad (\text{E.20})$$

Therefore, it suffices to compute the likelihood ratio with respect to F_W .

Now, note that

$$F_W = p \cdot F_{W|Z=1} + (1 - p) \cdot F_{W|Z=0} = p \cdot F_X + (1 - p) \cdot F_Y.$$

In particular, it follows from the linearity of the Radon-Nikodym derivative and the fact that the Radon-Nikodym derivative of a measure with respect to itself is unity that

$$1 - p \cdot \frac{dF_X}{dF_W} = (1 - p) \cdot \frac{dF_Y}{dF_W}. \quad (\text{E.21})$$

Next, consider the measure $\nu : E \mapsto \Pr(W \in E \mid Z = 1)$ —i.e., the distribution of X —and observe that

$$p \cdot \nu[E] = \Pr(W \in E, Z = 1) = \mathbb{E}[\Pr(Z = 1, W \in E \mid W)] = \mathbb{E}[\Pr(Z = 1 \mid W) \cdot \mathbf{1}(W \in E)].$$

Here the first equality follows from the definition of conditional probability, the second from the law of iterated expectation, and the third from the $\sigma(W)$ -measurability of $\mathbf{1}(W \in E)$. Now, by the Doob-Dynkin lemma, there exists a $\sigma(W)$ -measurable function $g(w)$ such that $\Pr(Z = 1 \mid W) = g(W)$ a.s. Therefore, the previous expression equals

$$\mathbb{E}[g(W) \cdot \mathbf{1}(W \in E)] = \int_E g(w) dF_W(w),$$

where the last equality is the fundamental property of pushforward measures; see, e.g., Theorem 3.6.1 in Bogachev (2007). In particular, since

$$p \cdot \Pr(W \in E \mid Z = 1) = \int_E p \cdot \frac{dF_{W|Z=1}}{dF_W}(w) dF_w = \int_E p \cdot \frac{dF_X}{dF_W}(w) dF_w,$$

by the uniqueness of the Radon-Nikodym derivative,

$$g(w) = \Pr(Z = 1 \mid W = w) = p \cdot \frac{dF_X}{dF_W}(w) \quad W\text{-a.s.}$$

Consider the function $h(q) : [0, 1] \rightarrow [0, \infty]$ given by

$$h(q) \stackrel{\text{def}}{=} \frac{1-p}{p} \cdot \frac{q}{1-q},$$

where, without loss of generality, we set $h(1) = \infty$. Now,

$$h \circ g = h \left(p \cdot \frac{dF_X}{dF_W} \right) = \frac{1-p}{p} \cdot \frac{p \cdot \frac{dF_X}{dF_W}}{1 - p \cdot \frac{dF_X}{dF_W}} = \frac{1-p}{p} \cdot \frac{p \cdot \frac{dF_X}{dF_W}}{(1-p) \cdot \frac{dF_Y}{dF_W}} = \frac{\frac{dF_X}{dF_W}}{\frac{dF_Y}{dF_W}},$$

which is the likelihood ratio. The second equality here follows from (E.21). However, since $h(q)$ is monotonically increasing, $h \circ g$ is non-decreasing if and only if $g(w)$ is. Therefore, by (E.20), MCP holds if and only if the MLRP holds. $\square$

Of these characterizations, the most important for proving Theorem E.2 is klinostochastic dominance (KSD). For instance, we see by taking $h(x) = 1$ in (E.19) that KSD (and hence the MLRP) implies stochastic dominance. KSD is closely related to the notion of tilted distributions. For a random variable X and an X -a.s. non-negative function $h(x)$ such that $\mathbb{E}[h(X)]$ is finite and positive, we denote by $h \circ X$ or $h \circ F_X$ the *tilt* of the distribution of X by the weight function h , i.e., the probability distribution with CDF

$$F_{h \circ X}(t) = \frac{\int_{\mathbb{R}} \mathbf{1}(x \leq t) \cdot h(x) dF_X}{\int_{\mathbb{R}} h(x) dF_X} = \frac{\mathbb{E}[\mathbf{1}(X \leq t) \cdot h(X)]}{\mathbb{E}[h(X)]}. \quad (\text{E.22})$$

KSD simply means that the stochastic ordering of two distributions is invariant under tilts: By Lemma E.8, $X \succeq_{\text{lr}} Y$ if and only if

$$h \circ X \succeq_1 h \circ Y$$

for any appropriate weight function $h(x)$ (i.e., $h(x)$ such that $h(X)$ and $h(Y)$ are a.s. non-negative and $\mathbb{E}[h(X)]$ and $\mathbb{E}[h(Y)]$ are finite and positive).

E.1.5 Proof of correctness of the robust outcome test

KSD is the key ingredient needed to prove our main result: Under appropriate ordering assumptions on the risk distributions and risk-decision curves, the robust outcome test is correct.

Proof. We prove the contrapositive. That is, assume instead that

$$\Pr(d_0(U) < d_1(U)) > 0 \quad \text{or} \quad d_0(u) = d_1(u) \quad U\text{-a.s.} \quad (\text{E.23})$$

Note that by our second ordering assumption, $d_0(u)$ and $d_1(u)$ generate distributions with the MLRP. Since the MLRP implies stochastic dominance, it follows that either $d_0(u) \leq d_1(u)$ U -a.s., or $d_1(u) \leq d_0(u)$ U -a.s. By (E.23), it follows that $d_0(u) \leq d_1(u)$.

Using the fact that $d_0(u) \leq d_1(u)$, as in the proof of Theorem 6.1, we will show that either:

1. The decision rate is at least as large for group $G = 1$ as for group $G = 0$, i.e.,

$$\Pr(D = 1 \mid G = 0) \leq \Pr(D = 1 \mid G = 1); \quad (\text{E.24})$$

2. The outcome rate is no bigger for group $G = 1$ than for group $G = 0$, i.e.,

$$\mathbb{E}[U \mid D = 1, G = 0] \geq \mathbb{E}[U \mid D = 1, G = 1]. \quad (\text{E.25})$$

Toward that end, let $U_0 \sim U \mid G = 0$ and $U_1 \sim U \mid G = 1$ be the conditional risk distributions, and let $d_0(u)$ and $d_1(u)$ generate distributions H_0 and H_1 with the MLRP. Then, there are two possibilities, which we will treat separately: Either $U_1 \succeq_{\text{lr}} U_0$, or $U_0 \succeq_{\text{lr}} U_1$.

Case 1 ($U_1 \succeq_{\text{lr}} U_0$). Suppose first that $U_1 \succeq_{\text{lr}} U_0$. Then, in particular, $U_1 \succeq_1 U_0$. Since $d_0(u) \leq d_1(u)$, we have that

$$\begin{aligned} \Pr(D = 1 \mid G = 0) &= \mathbb{E}[d_0(U) \mid G = 0] \\ &\leq \mathbb{E}[d_0(U) \mid G = 1] \\ &\leq \mathbb{E}[d_1(U) \mid G = 1] \\ &= \Pr(D = 1 \mid G = 1), \end{aligned}$$

i.e., the decision rate is at least as high for group $G = 1$ as for group $G = 0$, which is (E.24). Here, the equalities follow from the law of iterated expectations and the definition of $d_g(u)$ in (E.9). The

first inequality follows from stochastic dominance and the fact that $d_0(u)$ is non-decreasing—because it exhibits bounded rationality—and the second inequality from the assumption that $d_0(u) \leq d_1(u)$.

Case 2 ($U_0 \succeq_{lr} U_1$). Next, suppose that $U_0 \succeq_{lr} U_1$. The proof in this case is similar, although more delicate. We will use KSD to show that for a fixed risk-decision curve, the outcome rate is higher for group $G = 0$ than group $G = 1$, and then we will use KSD again in a different way to show that for a fixed group $G = g$, the risk-decision curve $d_0(u)$ results in a higher outcome rate than the risk-decision curve $d_1(u)$.

In particular, we will need to consider the distributions $d_{g'} \odot U_g$ for all $g, g' \in \{0, 1\}$, so we begin by verifying that these distributions are well-defined and have finite expectations.

Well-definedness and finite expectation of $d_{g'} \odot U_g$. Let $d(u)$ be any risk-decision curve, i.e., a U -a.s. non-decreasing function taking values in $[0, 1]$. Then $\mathbb{E}[d(U) \mid G = g]$ is the corresponding decision rate for group $G = g$. In order for $d \odot U_g$ to be well-defined, the definition of a tilted distribution in (E.22) requires that

$$0 < \int_{\mathbb{R}} d(u) dF_{U_g} = \mathbb{E}[d(u) \mid G = g] < \infty,$$

i.e., that the decision rate is positive. (For avoidance of doubt, by $\int \cdots dF_X$, we mean throughout the integral taken with respect to the pushforward measure $E \mapsto \Pr(X \in E)$ on $\bar{\mathbb{R}}$, not the Riemann-Stieltjes integral with respect to the integrator F_X .)

Finiteness follows immediately from the fact that $0 \leq d(u) \leq 1$ U -a.s. The positivity of $\mathbb{E}[d(U) \mid G = g]$ follows from our assumptions when $d(u) = d_{g'}(u)$. To see this, note that when $g = g'$, $\mathbb{E}[d_{g'}(U) \mid G = g] > 0$ by the non-triviality assumption in (E.6). On the other hand, when $g \neq g'$, observe that the overlap assumption in (E.8) implies that U_0 and U_1 are mutually absolutely continuous. Since $\mathbb{E}[d_{g'}(U) \mid G = g'] > 0$, there must exist some set E and $\epsilon > 0$ such that $d_{g'}(u) > \epsilon$ for all $u \in E$ and $\Pr(U \in E \mid G = g') > 0$. Therefore, in particular,

$$\mathbb{E}[d_{g'}(U) \mid G = g] \geq \epsilon \cdot \Pr(U \in E \mid G = g) > 0. \quad (\text{E.26})$$

Here, the first inequality follows from the fact that $d_g(u) \geq \epsilon \cdot \mathbf{1}(u \in E)$. The second follows from mutual absolute continuity and the fact that $\Pr(U \in E \mid G = g') > 0$. In short, the tilted distributions $d_{g'} \odot U_g$ are well-defined for all $g, g' \in \{0, 1\}$.

Next, we verify that the expectations of the tilted distributions $d_{g'} \circ U_g$ are all finite. These expectations are the outcome rates of the risk-decision curve $d_{g'}(u)$ over the conditional risk distribution of group $G = g$ for all $g, g' \in \{0, 1\}$. Let $d(u)$ be an arbitrary risk-decision curve, and assume that the corresponding decision rate $\mathbb{E}[d(u) \mid G = g]$ is positive. Then, the absolute value of the outcome rate for group $G = g$ can be bounded as follows:

$$\left| \frac{\int_{\mathbb{R}} u \cdot d(u) dF_{U_g}}{\int_{\mathbb{R}} d(u) dF_{U_g}} \right| \leq \frac{\int_{\mathbb{R}} |u| dF_{U_g}}{\int_{\mathbb{R}} d(u) dF_{U_g}} = \frac{\mathbb{E}[|U| \mid G = g]}{\mathbb{E}[d(U) \mid G = g]} \leq \frac{\mathbb{E}[|U|]}{\Pr(G = g) \cdot \mathbb{E}[d(U) \mid G = g]}. \quad (\text{E.27})$$

Here the first inequality follows from the fact that $0 \leq d(u) \leq 1$ U -a.s. combined with the triangle inequality for integrals. The second inequality follows from the fact that for any non-negative random variable X and positive probability event E , $\mathbb{E}[X \mid E] \leq \mathbb{E}[X] / \Pr(E)$. We note that $\mathbb{E}[|U|]$ is finite by the assumption in (E.7) and $\Pr(G = g) > 0$ by the assumption in (E.5). In particular, when $d(u) = d_{g'}(u)$, $\mathbb{E}[d(U) \mid G = g]$ is positive by (E.26). Therefore the outcome rates are finite.

Comparison of outcome rates for $d_0(u)$ and $d_1(u)$. Having verified that the decision rates are positive and the outcome rates are finite for groups $G = 0$ and $G = 1$ under both risk-decision curves $d_0(u)$ and $d_1(u)$, we move on to comparing these rates.

Since $U_0 \succeq_{\text{lr}} U_1$, we have by KSD in Lemma E.8 that $d_1 \circ U_0 \succeq_1 d_1 \circ U_1$. It follows from Lemma E.5 that the former outcome rate is greater than or equal to the latter outcome rate. To see this, note that $d_1(u)$ exhibits bounded rationality and is therefore non-decreasing. Additionally, the outcome rate

$$\frac{\int_{\mathbb{R}} u \cdot d_1(u) dF_{U_g}}{\int_{\mathbb{R}} d_1(u) dF_{U_g}}$$

is finite for $g \in \{0, 1\}$ by (E.27). Applying Lemma E.5, we therefore have that

$$\mathbb{E}[U \mid D = 1, G = 1] = \frac{\int_{\mathbb{R}} u \cdot d_1(u) dF_{U_1}}{\int_{\mathbb{R}} d_1(u) dF_{U_1}} \leq \frac{\int_{\mathbb{R}} u \cdot d_1(u) dF_{U_0}}{\int_{\mathbb{R}} d_1(u) dF_{U_0}}. \quad (\text{E.28})$$

Next, we show that switching the risk-decision curve from $d_1(u)$ to $d_0(u)$ can only increase the outcome rate. To see this, we proceed in four steps. First, note that by the definition of H_g , we have that

$$\frac{\int_{\mathbb{R}} u \cdot d_g(u) dF_{U_0}(u)}{\int_{\mathbb{R}} d_g(u) dF_{U_0}(u)} = \frac{\int_{\mathbb{R}} u \cdot \left[\int_{\mathbb{R}} \mathbf{1}(s \leq u) dF_{H_g}(s) \right] dF_{U_0}(u)}{\int_{\mathbb{R}} \left[\int_{\mathbb{R}} \mathbf{1}(s \leq u) dF_{H_g}(s) \right] dF_{U_0}(u)}.$$

Second, by (E.27), we can apply the Fubini-Tonelli theorem to the latter expression, yielding

$$\frac{\int_{\mathbb{R}} u \cdot d_g(u) dF_{U_0}(u)}{\int_{\mathbb{R}} d_g(u) dF_{U_0}(u)} = \frac{\int_{\mathbb{R}} \left[\int_{\mathbb{R}} u \cdot \mathbf{1}(u \geq s) dF_{U_0}(u) \right] dF_{H_g}(s)}{\int_{\mathbb{R}} \left[\int_{\mathbb{R}} \mathbf{1}(u \geq s) dF_{U_0}(u) \right] dF_{H_g}(s)}.$$

Third, multiplying and dividing the outer integrand in the numerator on the right-hand side by $\int_{\mathbb{R}} \mathbf{1}(u \geq s) dF_{U_0}(u)$ (and, when $\Pr(U \geq s \mid G = 0) = 0$, defining this ratio without loss of generality to be zero) yields that

$$\frac{\int_{\mathbb{R}} u \cdot d_g(u) dF_{U_0}(u)}{\int_{\mathbb{R}} d_g(u) dF_{U_0}(u)} = \frac{\int_{\mathbb{R}} \left[\frac{\int_{\mathbb{R}} u \cdot \mathbf{1}(u \geq s) dF_{U_0}(u)}{\int_{\mathbb{R}} \mathbf{1}(u \geq s) dF_{U_0}(u)} \right] \cdot \left[\int_{\mathbb{R}} \mathbf{1}(u \geq s) dF_{U_0}(u) \right] dF_{H_g}(s)}{\int_{\mathbb{R}} \left[\int_{\mathbb{R}} \mathbf{1}(u \geq s) dF_{U_0}(u) \right] dF_{H_g}(s)}. \quad (\text{E.29})$$

Fourth and finally, from the form of the integral, we see that the right-hand side of (E.29) is an expectation. In particular, it is the expectation of a non-decreasing function $f(s)$:

$$f(s) \stackrel{\text{def}}{=} \frac{\int_{\mathbb{R}} u \cdot \mathbf{1}(u \geq s) dF_{U_0}(u)}{\int_{\mathbb{R}} \mathbf{1}(u \geq s) dF_{U_0}(u)} = \mathbb{E}[U \mid U \geq s, G = 0].$$

The expectation in (E.29) is taken with respect to the distribution $h \circ H_g$, where

$$h(s) \stackrel{\text{def}}{=} \int_{\mathbb{R}} \mathbf{1}(u \geq s) dF_{U_0}(u) = \Pr(U \geq s \mid G = 0).$$

(Since $h(s) > 0$ if and only if $\Pr(U \geq s \mid G = 0) > 0$, $f(s)$ is well-defined $h \circ H_g$ -a.s.) By (E.29) the expectation of $f(s)$ with respect to the distribution $h \circ H_g$ equals

$$\frac{\int_{\mathbb{R}} u \cdot d_g(u) dF_{U_0}(u)}{\int_{\mathbb{R}} d_g(u) dF_{U_0}(u)},$$

and so the expectation is consequently finite by (E.27).

Putting these pieces together, since $H_0 \succeq_{\text{lr}} H_1$ by hypothesis, it follows that $h \circ H_0 \succeq_1 h \circ H_1$. Thus, since $f(s)$ is non-decreasing and has finite expectation, it follows by Lemma E.5 that its expectation with respect to $h \circ H_0$ is at least as large as its expectation with respect to $h \circ H_1$. In other words, the expression in (E.29) with $g = 0$ is greater than or equal to the expression with $g = 1$, i.e.,

$$\frac{\int_{\mathbb{R}} u \cdot d_1(u) dF_{U_0}}{\int_{\mathbb{R}} d_1(u) dF_{U_0}} \leq \frac{\int_{\mathbb{R}} u \cdot d_0(u) dF_{U_0}}{\int_{\mathbb{R}} d_0(u) dF_{U_0}} = \mathbb{E}[U \mid D = 1, G = 0].$$

Combining this with (E.28) gives that the outcome rate is at least as small for group $G = 1$ as group $G = 0$, i.e.,

$$\mathbb{E}[U \mid D = 1, G = 1] \leq \mathbb{E}[U \mid D = 1, G = 0],$$

which is (E.25).

□

E.1.6 A Note on Sampling Error

To clarify our main conceptual points, we have focused throughout on the infinite data regime, where decision and outcome rates are estimated with negligible error. In practice, however, sampling error can be an important concern.

For notational simplicity, let DR_g and OR_g denote the true decision and outcome rates for group g —i.e., $\Pr(D = 1 \mid G = g)$ and $\mathbb{E}[U \mid D = 1, G = g]$. Let

$$n_g \stackrel{\text{def}}{=} \sum_{i=1}^n \mathbf{1}(G_i = g) \quad \text{and} \quad n_{g,d} \stackrel{\text{def}}{=} \sum_{i=1}^n \mathbf{1}(G_i = g, D_i = d)$$

denote the number of individuals in group $G = g$ and, respectively, the number of observations both in group $G = g$ and receiving decision $D = d$. Consistent with our notation in Section 6.2.2, let $\widehat{\text{DR}}_g$ and $\widehat{\text{OR}}_g$ for $g \in \{0, 1\}$ denote the empirical decision and outcome rates, i.e.,

$$\widehat{\text{DR}}_g \stackrel{\text{def}}{=} \frac{n_{g,1}}{n_g}, \quad \widehat{\text{OR}}_g \stackrel{\text{def}}{=} \frac{\sum_{i=1}^n U_i \cdot \mathbf{1}(G_i = g, D_i = 1)}{n_{g,1}}.$$

Set

$$\widehat{\text{SVar}}_{U,g} \stackrel{\text{def}}{=} \frac{\sum_{i=1}^n \mathbf{1}(G_i = g, D_i = 1) \cdot (U_i - \widehat{\text{OR}}_g)^2}{n_{g,1} - 1}.$$

Finally, let $\Delta_{\text{DR}} \stackrel{\text{def}}{=} \text{DR}_1 - \text{DR}_0$ denote the true difference in decision rates, with Δ_{OR} , $\hat{\Delta}_{\text{DR}}$, and $\hat{\Delta}_{\text{OR}}$ defined analogously.

To quantify uncertainty in $\hat{\Delta}_{\text{DR}}$ and $\hat{\Delta}_{\text{OR}}$, first observe that $\widehat{\text{DR}}_0$, $\widehat{\text{DR}}_1$, $\widehat{\text{OR}}_0$, and $\widehat{\text{OR}}_1$ are approximately independent: Dependency arises only because n_0 , n_1 , $n_{0,1}$, and $n_{1,1}$ are not fixed.⁵ As a result, we can estimate the standard errors of $\hat{\Delta}_{\text{DR}}$ and $\hat{\Delta}_{\text{OR}}$ as follows:

$$\widehat{\text{SErr}}_{\Delta_{\text{DR}}} \stackrel{\text{def}}{=} \sqrt{\frac{\widehat{\text{DR}}_0 \cdot (1 - \widehat{\text{DR}}_0)}{n_0} + \frac{\widehat{\text{DR}}_1 \cdot (1 - \widehat{\text{DR}}_1)}{n_1}}, \quad \widehat{\text{SErr}}_{\Delta_{\text{OR}}} \stackrel{\text{def}}{=} \sqrt{\frac{\widehat{\text{SVar}}_{U,0}}{n_{0,1}} + \frac{\widehat{\text{SVar}}_{U,1}}{n_{1,1}}}.$$

From this approximate independence, it also follows immediately that the most compact central

⁵In particular, using the delta method and central limit theorem, it is straightforward to show that as $n \rightarrow \infty$, $(\hat{\Delta}_{\text{DR}} - \Delta_{\text{DR}}) / \widehat{\text{SErr}}_{\Delta_{\text{DR}}}$ and $(\hat{\Delta}_{\text{OR}} - \Delta_{\text{OR}}) / \widehat{\text{SErr}}_{\Delta_{\text{OR}}}$ converge in distribution to independent standard normals.

confidence region is given by

$$C\left(\hat{\Delta}_{\text{DR}}, \hat{\Delta}_{\text{OR}}; \alpha\right) \stackrel{\text{def}}{=} \left\{ (\Delta_{\text{DR}}, \Delta_{\text{OR}}) : \left(\frac{\Delta_{\text{DR}} - \hat{\Delta}_{\text{DR}}}{\widehat{\text{SErr}}_{\Delta_{\text{DR}}}} \right)^2 + \left(\frac{\Delta_{\text{OR}} - \hat{\Delta}_{\text{OR}}}{\widehat{\text{SErr}}_{\Delta_{\text{OR}}}} \right)^2 \leq x_{1-\alpha} \right\}, \quad (\text{E.30})$$

where $x_{1-\alpha}$ is the $(1 - \alpha)$ -quantile of the $\chi^2(2)$ distribution. Equivalently, $C\left(\hat{\Delta}_{\text{DR}}, \hat{\Delta}_{\text{OR}}; \alpha\right)$ corresponds to thresholding the joint density of independent

$$\mathcal{N}\left(\hat{\Delta}_{\text{DR}}, \widehat{\text{SErr}}_{\Delta_{\text{DR}}}\right) \quad \text{and} \quad \mathcal{N}\left(\hat{\Delta}_{\text{OR}}, \widehat{\text{SErr}}_{\Delta_{\text{OR}}}\right)$$

distributions above an appropriate threshold t_α to obtain the region of highest density. This results in an elliptical confidence region for Δ_{DR} and Δ_{OR} centered at $(\hat{\Delta}_{\text{DR}}, \hat{\Delta}_{\text{OR}})$, with the major and minor axes determined by the estimated standard errors of the sample difference in decision and outcome rates. Asymptotically valid hypothesis tests can be straightforwardly obtained by, for example, taking the maximum p -value of one-tailed tests of $\Delta_{\text{DR}} \leq 0$ and $\Delta_{\text{OR}} \geq 0$.

E.2 Data and Risk Estimation

Our empirical results are based on data drawn from five sources:

- **Police Stops:** Administrative data gathered under the California Racial Identity and Profiling Act (Guerrero et al., 2024);
- **Lending:** Lending data drawn from applicants to a large online financial technology platform;
- **Recidivism:** COMPAS recidivism prediction scores collected in Broward County, Florida (Angwin et al., 2022);
- **Contraband:** Administrative data from the New York Police Department’s Stop, Question, and Frisk program (Floyd v. City of New York, 2013);
- **Bar passage:** Law school admissions and graduation data from the Law School Admissions Council’s Longitudinal Bar Passage Study (Wightman, 1998).

Data sources, pre-processing steps, risk model specifications, exclusion criteria, and other details are given below.

E.2.1 Police Stops

The police stops dataset used in this analysis is drawn from data gathered in 2022 pursuant to California’s Racial Identity and Profiling Act (RIPA). Under RIPA, law enforcement officers must record a wide range of information about stops they make of vehicles and pedestrians, including the stop circumstances (e.g., date and time), demographics of the stopped individual (e.g., race and age), reasons for the stop (e.g., traffic violation or matching the description of a person of interest), outcomes of the stop (e.g., weapons or other contraband found), actions taken during the stop (e.g., arrest or use of force), and other information. First enacted in 2015, RIPA requirements were extended to all law enforcement agencies in the state in 2022.

The full dataset comprises more than 4.5 million records of police stops from 536 law enforcement agencies (Guerrero et al., 2024). Following Grossman et al. (2024a), we restrict our analysis to non-consensual, discretionary stops not conducted in an educational context. That is, we exclude:

1. Stops undertaken because the officer had knowledge of an arrest warrant or some form of mandatory supervision (i.e., probation, parole, or post-release community supervision);
2. Consensual encounters with motorists or pedestrians, which are recorded inconsistently across agencies;
3. Stops that occurred in an educational context, or for education-related reasons, such as truancy, as the threshold for conducting stops of this nature can be lower than for other types of stops. (See, e.g., *The People v. William G.*, ref. *The People v. William G.*, 1985.)

To ensure that sufficient data exist to accurately estimate the decision and outcome rates, we further restrict our analysis to agencies that conducted at least 1,000 stops of Black, Hispanic, and White individuals in 2022. This leaves us with a final dataset of approximately 2.8 million stops from 56 law enforcement agencies.

E.2.2 Lending

Data Source Our lending data come from an online financial technology platform that facilitates lending to individuals for a variety of purposes based on a proprietary underwriting model that

utilizes both traditional and non-traditional information to assess lending risk. Our training data consist of approximately 300,000 borrowers who applied for loans between January 2019 and July 2021 for whom repayment outcomes are known. While race is not recorded in the training data, it is imputed using names and geographic information using Bayesian Improved Surname Geocoding or “BISG” (Elliott et al., 2009). For each individual, we impute the most likely race using the BISG model, and then subset to individuals whose imputed race is Black, Hispanic, or White. We then fit our model on the approximately 260,000 individuals remaining. To understand the distribution of risk among loan applicants, we apply our fitted risk model to a separate set of 130,000 individuals who applied for loans in early- to mid-August 2021, but who did not necessarily receive a loan.

Risk Model We fit a gradient-boosted decision tree model to predict the likelihood that an individual will default before the end of the loan term using a bevy of traditional and non-traditional covariates, including credit score, available at the time of their application. The model achieves an AUC of 73% on a 10% set of held-out training data, and is well-calibrated, as shown in Figure E.2.

E.2.3 Recidivism

Data Source The recidivism dataset used in this analysis comes from ProPublica’s analysis of COMPAS recidivism prediction scores from Broward County, Florida (Angwin et al., 2022). The dataset contains COMPAS scores for approximately 12,000 individuals who were assessed for recidivism risk pre-trial in 2013 and 2014, as well as information on whether they were rearrested within two years of their initial assessment. After subsetting to the collection of individuals for whom outcomes are known, and whose race is recorded as Black, Hispanic, or White, we are left with a final dataset of approximately 11,000 observations.

Risk Model To predict probability of recidivism, we use raw COMPAS scores, which range continuously from approximately -3 to 3. To obtain calibrated probabilities, we fit a logistic regression model to predict recidivism from the interaction of the raw COMPAS score and race. The model achieves an AUC of 68% on the held-out training data, which is consistent with the

AUC reported for COMPAS scores in a variety of jurisdictions (Angwin et al., 2022; Brennan et al., 2009; Dieterich et al., 2011, 2017, 2018; Equivant, 2019; Farabee et al., 2010; Flores et al., 2016; Lansing, 2012; Reich et al., 2016). The model is also well-calibrated, as shown in Figure E.2.

E.2.4 Contraband

Data Source Under *Terry v. Ohio*, Police officers may stop and question pedestrians about whom they have “reasonable suspicion” of criminal activity, and may conduct a “frisk” of the individual’s outer clothing if they believe the stopped individual is in possession of a weapon (Terry v. Ohio, 1968). Between 2003 and 2013, the New York Police Department (NYPD) conducted over 100,000 such *Terry* stops per year. Officers recorded details of these stops, including stop circumstances (e.g., time and location), reasons for suspicion (e.g., suspicious bulge or furtive movements), stop outcomes (e.g., arrest or contraband found), and demographic information about the stopped individual. We use records of the approximately 2.7 million stops conducted between 2008 and 2013. Our training set consists of the approximately 800,000 stops conducted between 2008 and 2010 in which the stopped individual was frisked, and consequently for which it is known whether or not the individual was in possession of a weapon. We split the training data into an 80% training set, which we use to fit our model predicting weapon possession, and a 20% holdout set, which we use to evaluate the performance of our model. To understand the overall distribution of risk among stopped individuals, we apply our fitted risk model to the approximately 1.1 million stops conducted in 2011 and 2012.

Risk Model We fit a gradient-boosted decision tree model to predict the likelihood that an individual who is frisked will be found to be in possession of a weapon using 33 covariates known at the time of the stop and before the frisk decision is made. These covariates include the time and location of the stop, the officer’s reason for making the stop, and additional circumstances related to the stop, such as whether the individual was in proximity to a crime scene. The model achieves an AUC of 79% on the held-out training data. The model is reasonably well-calibrated for risks below 9% (approximately the 99th percentile of the risk distribution), although it is modestly

over-predictive of risk for White individuals; see Figure E.2.

E.2.5 Admissions

Data Source We draw our bar passage dataset from the Law School Admissions Council’s Longitudinal Bar Passage Study (Wightman, 1998), which studied bar passage rates among 163 of 172 ABA-accredited law schools in the United States. Study participants represent around 23,000 out of the approximately 40,000 law school students who entered law school in the fall of 1991. The dataset includes information about the law schools attended by the students, their undergraduate GPAs, their LSAT scores, their gender and race, and whether they passed the bar exam on their first or second attempt.

Risk Model We fit a logistic regression model to predict the likelihood that an individual eventually passed the bar exam using the individual’s LSAT score, undergraduate GPA, a measure of socioeconomic status, gender, and race. The model achieves an AUC of 79%, and is well-calibrated, as shown in Figure E.2: For Black and Hispanic applicants, the model is well-calibrated across a range of estimated risks, and for White applicants, it is well-calibrated for risks greater than roughly 90% (approximately the 3rd percentile of risk for White applicants).

E.3 Simulation Study

E.3.1 Violations of the MLRP

We present the results of the simulation study described in the main text for Hispanic individuals in Figure E.3. We further extend our simulation study to quasi-rational decision policies of the following form:

$$d_g(r) = F(r; t_g, \sigma) \stackrel{\text{def}}{=} \Phi \left(\frac{\text{logit}^{-1}(r) - \text{logit}^{-1}(t_g)}{\sigma} \right), \quad (\text{E.31})$$

where $\Phi(x)$ is the standard normal CDF. Here, $F(x; t, \sigma)$ is the cumulative distribution function of the logit-normal distribution with center t (on the probability scale) and scale parameter σ . To model a realistic level of bounded rationality, we set σ to be one half of the standard deviation of the

overall distribution of estimated risk on the logit scale. (In particular, we set σ to be 0.62 logits for the lending dataset, 0.36 logits for the recidivism dataset, 0.45 logits for the contraband dataset, and 0.51 logits for the admissions dataset.) As before, we then sweep over all percentiles t_g (excluding the 0th and 100th percentiles); see Figure E.4 for examples of the quasi-rational risk-decision curves used in the simulation, where t_g is set to the 1/3, 1/2, and 2/3 quantiles, respectively. At each percentile, we estimate the decision and outcome rates for group $G = g$ as follows:⁶

$$\widehat{\text{DR}}_g \stackrel{\text{def}}{=} \frac{1}{n_g} \sum_{i=1}^n \mathbf{1}(G = g) \cdot d_g(\hat{R}_i), \quad \widehat{\text{OR}}_g \stackrel{\text{def}}{=} \frac{\sum_{i=1}^n \mathbf{1}(G = g) \cdot d_g(\hat{R}_i) \cdot \hat{R}_i}{\sum_{i=1}^n \mathbf{1}(G = g) \cdot d_g(\hat{R}_i)}, \quad (\text{E.32})$$

where $d_g(r)$ is the risk-decision curve being tested for group $G = g$, $\hat{R}_i$ is the estimated risk for individual i , and n_g is the number of individuals in group g . (We note that this is equivalent to the method used in the main text, differing only in that the risk-decision curve takes values in $[0, 1]$.) We then estimate the difference in decision and outcome rates between groups g_1 and g_0 as

$$\widehat{\text{DR}}_{g_1} - \widehat{\text{DR}}_{g_0}, \quad \widehat{\text{OR}}_{g_1} - \widehat{\text{OR}}_{g_0},$$

and test for discrimination using the robust and standard outcome tests accordingly.

Figures E.5 and E.6 show that broadly similar, if even more pronounced, results hold when the risk-decision curves are quasi-rational as opposed to fully rational. In particular, the robust outcome test rarely detects discrimination when decisions are made in the same way for both groups, and usually detects discrimination when there is a double standard. In particular, in cases where the base rates differ substantially between groups and the decision maker has limited sensitivity to risk, the standard outcome test can essentially always “detect” that there is a lower decision standard for the lower base rate group, regardless of the actual decision standards.

E.3.2 Differences in Base Rates

The range of cases in which the robust outcome test detects discrimination rather than returning an inconclusive result varies across our examples. In the simulations shown in Figures 6.3, E.3, E.5, and E.6, the robust outcome test is inconclusive in more scenarios in examples where the base rates of the groups are more different.

⁶Our notation in Eq. (E.32) differs slightly from Section 6.2.2 and Appendix E.1.6 to take advantage of the fact that the risk decision curves are known exactly.




Figure E.1: A stylized illustration of a pair of risk distributions and lending decisions for which the robust outcome test would be incorrect. The lending rates correspond to the areas of the colored regions (viz., 50% for the red group and 38% for the blue group), and the repayment rates are shown by the dotted red and blue lines (viz., 62% for the red group and 67% for the blue group). Applying a uniform lending threshold (50%) results in a lower lending rate to applicants from the blue group, as well as a higher repayment rate, leading the robust outcome test to erroneously conclude that the thresholds differ for the two groups.

To gauge whether differences in base rates reduce the robust outcome test’s sensitivity more broadly, we additionally simulate a range of discriminatory and non-discriminatory scenarios where we allow base rates to differ to a greater or lesser extent. More precisely, we consider risk distributions of the form

$$R \mid G = g \sim \text{Beta}(\mu_g, \nu),$$

i.e., where the risk distribution for group $G = g$, $g \in \{0, 1\}$, is beta distributed with mean μ_g and total count parameter ν .⁷ We consider all possible threshold decision policies with $t_0, t_1 \in \{1\%, 2\%, \dots, 99\%\}$, and report the proportion of scenarios in which the robust outcome test is conclusive, i.e., indicates discrimination against one group or the other. We sweep over $\mu_0 \in \{10\%, 20\%, \dots, 90\%\}$, $\mu_1 \in \{1\%, 2\%, \dots, 99\%\}$, and $\nu \in \{1, 2, 4, \dots, 64\}$. The results are shown in Figure E.7. Across total count parameters and base rates μ_0 for group $G = 0$, a pattern consistent with our observations using empirical risk distributions holds: Increasing the gap between the high and low base rate groups results in an inconclusive robust outcome test more often.

⁷In the standard parameterization of the beta distribution, $\alpha = \mu \cdot \nu$ and $\beta = (1 - \mu) \cdot \nu$.




Figure E.2: Calibration curves for the risk models used in our analysis. The x-axis indicates the estimated proportion of cases in which the predicted event (viz., loan repayment, recidivism, weapon possession, or bar passage) occurs, and the y-axis indicates the actual proportion of cases in which the event occurs. The diagonal line indicates perfect calibration. The smooth curves represent logistic regression GAM models fit to the data using a thin plate regression spline basis. The shaded regions indicate 95% confidence intervals.




Figure E.3: Results of a simulation study comparing the standard and robust outcome tests. The x -axis indicates the decision threshold for White individuals, and the y -axis indicates the threshold for Hispanic individuals. The upper-left and lower-right triangular regions correspond to scenarios where decision makers discriminate against either Hispanic or White individuals, as indicated by the “×” and “•” symbols, respectively; non-discriminatory scenarios are shown by the dashed diagonal line. Red regions indicate where the tests suggest discrimination against Hispanic individuals, blue regions indicate where the tests suggest discrimination against White individuals, and yellow regions indicate where the robust outcome test is inconclusive. The gray areas represent simulation scenarios that are not feasible because a threshold lies outside the support of the risk distribution of the corresponding group. Across simulations, the standard outcome test indicates discrimination when, in reality, there is none—and often indicates discrimination against the group that in actuality was favored. In contrast, the robust outcome test is nearly always directionally accurate, though it sometimes returns an inconclusive result.




Figure E.4: Examples of the quasi-rational risk-decision curves used in the simulation study. The x -axis indicates the estimated risk. The overall distribution of risk is shown by the gray histogram. The red, green, and blue curves show the quasi-rational risk-decision curves centered, respectively, at the $1/3$, $1/2$, and $2/3$ quantiles—indicated by red, green, and blue dashed vertical lines.




Figure E.5: Results of a simulation study comparing the standard and robust outcome tests with quasi-rational (i.e., logit-normal CDF) risk-decision curves. The x -axis indicates the “center”—i.e., t , in Eq. (E.31)—for White individuals, and the y -axis indicates the center for Black individuals. The upper-left and lower-right triangular regions correspond to scenarios where decision makers discriminate against either Black or White individuals, as indicated by the “ $\times$ ” and “ $\cdot$ ” symbols, respectively; non-discriminatory scenarios are shown by the dashed diagonal line. Red regions indicate where the tests suggest discrimination against Black individuals, blue regions indicate where the tests suggest discrimination against White individuals, and yellow regions indicate where the robust outcome test is inconclusive. Across simulations, the standard outcome test indicates discrimination when, in reality, there is none—and often indicates discrimination against the group that in actuality was favored. In contrast, the robust outcome test is nearly always directionally accurate, though it sometimes returns an inconclusive result.




Figure E.6: Results of a simulation study comparing the standard and robust outcome tests with quasi-rational (i.e., logit-normal CDF) risk-decision curves. The x -axis indicates the “center”—i.e., t , in Eq. (E.31)—for White individuals, and the y -axis indicates the center for Hispanic individuals. The upper-left and lower-right triangular regions correspond to scenarios where decision makers discriminate against either Hispanic or White individuals, as indicated by the “ $\times$ ” and “ $\cdot$ ” symbols, respectively; non-discriminatory scenarios are shown by the dashed diagonal line. Red regions indicate where the tests suggest discrimination against Hispanic individuals, blue regions indicate where the tests suggest discrimination against White individuals, and yellow regions indicate where the robust outcome test is inconclusive. Across simulations, the standard outcome test indicates discrimination when, in reality, there is none—and often indicates discrimination against the group that in actuality was favored. In contrast, the robust outcome test is nearly always directionally accurate, though it sometimes returns an inconclusive result.




Figure E.7: Results of a simulation study comparing how often the robust outcome test is conclusive (i.e., indicates that there is discrimination against one of the groups) as a function of the difference between the base rates of the two groups. The facet indicates μ_0 , the base rate of group $G = 0$; the x-axis indicates μ_1 , the base rate of group $G = 1$; and the y-axis indicates the proportion of scenarios (i.e., choices of t_0 and t_1) in which the robust outcome test is conclusive. Across simulation settings, a larger gap is associated with a greater proportion of inconclusive results, reflecting the reduced sensitivity of the robust outcome test.

Appendix to Chapter 7

F.1 A Comparison to Alternative Ignorability Conditions

To better understand subset ignorability, we compare it to alternative conditions that recently have been proposed in the context of discrimination studies. In particular, we compare subset ignorability to a set of assumptions introduced by Knox et al. (2020), which they call treatment ignorability, mediator ignorability, and mediator monotonicity. We show that this set of assumptions, like subset ignorability, is sufficient—but not necessary—to ensure the SATE_M is nonparametrically identified by data on second-stage decisions. Importantly, however, the Knox et al. conditions are unlikely to be satisfied in important examples of potentially discriminatory decision making where subset ignorability holds (either exactly or approximately) and the SATE_M accordingly can be estimated, like those situations presented in Sections 7.3 and 7.4.

Aside from the Knox et al. conditions, it is instructive to compare subset ignorability to sequential ignorability (Imai et al., 2010a,b), a popular and often useful concept that was introduced to formalize causal mediation analysis, and one that is closely related to the Knox et al. conditions. Sequential ignorability is strictly stronger than subset ignorability, meaning that the former implies the latter but that the converse does not hold. In the setting of discrimination studies, there is little reason to believe sequential ignorability—or reasonable approximations of it—would be satisfied, and we primarily discuss the idea to clarify its distinction from subset ignorability.

The alternative ignorability conditions considered here were developed in the context of a single treatment. Therefore, to facilitate a direct comparison between subset ignorability and these alternatives, we adopt this single-treatment perspective throughout the Appendix. As discussed in the main text, there are substantive issues with positing a single manipulation of (perceived) race, gender, or other immutable characteristics in many multi-stage settings. Formally, however, it is

straightforward to collapse Z and D to a single treatment—which we call Z —that affects both the first-stage and the second-stage decisions. In particular, we now assume the potential outcomes $M(z)$ and $Y(z, m)$ satisfy the consistency relations $M = M(Z)$ and $Y = Y(Z, M)$. We emphasize that in this new framing, the definition of subset ignorability in Eq. (7.4) remains the same and that Theorem 7.4 likewise holds unaltered—since neither explicitly references the first-stage potential outcomes.¹

We start by formally considering sequential ignorability, following Imai et al. (2010a,b).

Definition F.1 (Sequential ignorability). We say that *sequential ignorability* is satisfied when the following two conditional independence criteria hold:

$$\{Y(z', m), M(z)\} \perp\!\!\!\perp Z \mid X, \quad (\text{F.1})$$

$$Y(z', m) \perp\!\!\!\perp M \mid Z, X, \quad (\text{F.2})$$

for $z, z' \in \{w, b\}$ and $m \in \{0, 1\}$.

The two key conditional independence assumptions we list are the same as in the definition of sequential ignorability given by Imai et al. (2010a,b), but to facilitate direct comparison with other ignorability criteria, we omit from our definition the accompanying overlap conditions. Also, for ease of exposition, we present the definition in the setting of binary treatment and mediator variables, though the original was more general. In the context of our running example, sequential ignorability means that: (a) conditional on the observed covariates X , the potential outcomes for charging $Y(z, m)$ and arrest $M(z)$ are jointly independent of an individual’s actual race Z ; and (b) conditional on the observed covariates X and an individual’s race Z , the arrest decision M is independent of the potential charging outcomes $Y(z, m)$.

¹As noted in Footnote 12 above, since we have restricted to the context of a single treatment Z , many of the quantities we consider are not expressible via the *do*-calculus, though they are still expressible in terms of potential outcomes. We emphasize that these potential outcomes should be understood in the conventional sense (Pearl, 2009): $Y(z, 1)$ represents what would have resulted for an individual if, counterfactually, one had intervened on M so that $M = 1$ and Z so that $Z = z$. Although directly manipulating the first-stage decision so that $M = 1$ may be implausible in some situations—for instance, it may be challenging in practice to intervene on an arresting officer’s decision—no issue arises in our setting as we are only concerned with the outcomes $Y(z, 1)$ for individuals who would be arrested in the absence of such an intervention (i.e., where it is already the case that $M = 1$). Moreover, while the FFRCISTG framework (Richardson and Robins, 2013; Robins, 1986) may consider these to be “cross-world” counterfactual quantities, we note that recent extensions of these frameworks discussed in Robins et al. (2020) could accommodate our estimand and identifying assumptions by allowing for the race variable to be split into race variables that are time- and context-specific, as we did in the main body of the article.

Theorem F.5, below, shows that sequential ignorability implies subset ignorability, but also, importantly, that sequential ignorability is a strictly stronger condition. To understand why, consider the stylized model of Section 7.3.1, in which one has all of the information that drives a prosecutor’s charging decision—satisfying subset ignorability—but not all of the information that drives an officer’s arrest decision. For example, suppose the prosecutor has access to the officer’s report, but not the arrested individual’s actual behavior. In this case, one would in general expect the first condition of sequential ignorability—in Eq. (F.1)—to be violated. In particular, without detailed data on what an officer observes, there is little reason to think the arrest potential outcomes, $M(z)$, would be independent of an individual’s race, even controlling for factors available to the prosecutor.

We next formally present the definitions of treatment ignorability, mediator ignorability, and mediator monotonicity proposed by Knox et al., starting with treatment ignorability.

Definition F.2 (Treatment ignorability). *Treatment ignorability* is the combination of the following two conditional independence criteria: for $z, z' \in \{w, b\}$ and $m \in \{0, 1\}$,

$$M(z) \perp\!\!\!\perp Z \mid X, \tag{F.3}$$

$$Y(z', m) \perp\!\!\!\perp Z \mid M(w), M(b), X. \tag{F.4}$$

In the context of arrest and charging decisions, treatment ignorability means that: (a) the potential outcomes for the arrest decision $M(z)$ are independent of race Z , after conditioning on the observed covariates X ; and (b) the potential outcomes for the charging decision $Y(z', m)$ are independent of race Z after conditioning on both the covariates X and the arrest potential outcomes $M(w)$ and $M(b)$.

The first condition of treatment ignorability is similar to the first condition of sequential ignorability, and it is unlikely to hold in our setting for the same reason. In general, given only information about what motivates the second-stage decision (e.g., charging, in our case) one cannot say much about what occurs in the first stage (e.g., arrest). But, critically, such information about the first stage is not necessary to estimate the SATE_M , which only quantifies discrimination in the second-stage decision. Theorem 7.4 makes that statement precise, showing that subset ignorability—which does not consider first-stage potential outcomes—is sufficient to ensure the

SATE_M is nonparametrically identified by the second-stage data.

The second criterion of treatment ignorability appears similar in spirit to subset ignorability, but it conditions on the potential outcomes $M(w)$ and $M(b)$ rather than on the actual outcome M . In practice, that distinction may not be too significant; in theory, however, the difference between the two is large. As we show in Theorem F.5 below, treatment ignorability alone—even with its strong assumption on the first stage—is not sufficient to ensure the SATE_M is identified by the second-stage data.

Finally, we consider mediator ignorability and the related mediator monotonicity condition.

Definition F.3 (Mediator ignorability). For $z \in \{w, b\}$ and $m \in \{0, 1\}$, *mediator ignorability* is satisfied when

$$Y(z, m) \perp\!\!\!\perp M(w) \mid Z = z, M(b) = 1, X. \quad (\text{F.5})$$

Definition F.4 (Mediator monotonicity). *Mediator monotonicity* is satisfied when


$$M(b) \geq M(w). \quad (\text{F.6})$$

In our running example, mediator ignorability means that the charging potential outcomes $Y(z, m)$ are independent of *one* of the arrest potential outcomes— $M(w)$, the arrest decision for (counterfactually) white individuals—conditional on the observed covariates X , and among individuals of race $Z = z$, who would be arrested if they were Black. The asymmetry in this condition stems from the additional mediator monotonicity constraint considered by Knox et al.: $M(b) \geq M(w)$, meaning that an individual who would be arrested if white would also be arrested if Black. The monotonicity condition is perhaps intuitively plausible given our understanding of racial discrimination, but the conditional independence assumption of mediator ignorability appears harder to interpret.

Having introduced the key definitions, we now present our main analytic result, Theorem F.5, which summarizes and formalizes our discussion of the various ignorability assumptions and their connections to estimating discrimination. In particular, we show that sequential ignorability is a strictly stronger assumption than subset ignorability, and recapitulate (from Theorem 7.4) that subset ignorability is a sufficient condition for the difference-in-means estimator Δ_n to yield consistent estimates of the SATE_M . Further, we show that treatment ignorability is not a necessary condition

for Δ_n to yield consistent estimates. We show this by explicitly constructing examples for which $\Delta_n \xrightarrow{\text{a.s.}} \text{SATE}_M$, but which violate the treatment ignorability condition. We additionally show that treatment ignorability is not a sufficient condition to guarantee consistency, despite its formal resemblance to the (sufficient) subset ignorability condition. To do so, we construct a family of observationally equivalent examples that satisfy treatment ignorability but which have different values of the SATE_M . Accordingly, no estimator, including Δ_n , can yield a consistent estimate of the SATE_M for every instance in the family. Importantly, the more conventional assumption of subset ignorability is sufficient to ensure the SATE_M can be identified from data on the second-stage decisions.

Theorem F.5. *Assume overlap holds, meaning that $\Pr(Z = z \mid X = x, M = 1) > 0$ for all x and z . Then we have the following collection of implications and nonimplications:*



Proof. Theorem 7.4 shows that subset ignorability implies that Δ_n is a consistent estimator of the SATE_M . We show the remaining seven implications and nonimplications in turn, starting with the claim that sequential ignorability implies subset ignorability. In particular, we prove that the conjunction of treatment ignorability, mediator ignorability, and mediator monotonicity implies that Δ_n is a consistent estimator of the SATE_M —a fact initially suggested by Knox et al.

Case 1 (Sequential ignorability implies subset ignorability). The first condition of sequential ignorability, in Eq. (F.1), states that $Y(z, m)$ and $M(z')$ are jointly independent of Z given X : $\{Y(z, m), M(z')\} \perp\!\!\!\perp Z \mid X$. From this, it immediately follows that $Y(z, m)$ alone is independent of Z given X : $Y(z, m) \perp\!\!\!\perp Z \mid X$. Now, because $Y(z, m) \perp\!\!\!\perp M \mid Z, X$ —which is the second condition

of sequential ignorability, in Eq. (F.2)—we have that $Y(z, m) \perp\!\!\!\perp \{Z, M\} \mid X$, by the contraction property of conditional independence. Therefore, by the weak-union property,

$$Y(z, m) \perp\!\!\!\perp Z \mid M, X. \quad (\text{F.7})$$

Subset ignorability now follows, as it is the special case in which $M = 1$ in Eq. (F.7).

Case 2 (Subset ignorability does not imply sequential ignorability). Sequential ignorability is an intuitively stronger condition than subset ignorability, as the former requires that Z is independent of the mediator potential outcomes $M(z)$ given X . Indeed, the synthetic example given in Section 7.3 satisfies subset ignorability but violates sequential ignorability.

To formally establish our claim, we construct an even simpler example that satisfies subset ignorability but not sequential ignorability. First, suppose that $Y(z, 1) = 1$ and $Y(z, 0) = 0$, deterministically for $z \in \{b, w\}$. In particular, using the language of our policing and prosecution application, everyone who is arrested is charged, regardless of race. We further set $X = 1$, which effectively means that there are no contextual variables. Finally, we set

$$\begin{aligned} \Pr(Z = z, M(b) = m_b, M(w) = m_w) \\ = \Pr(Z = z) \cdot \Pr(M(b) = m_b \mid Z = z) \cdot \Pr(M(w) = m_w \mid Z = z), \end{aligned} \quad (\text{F.8})$$

where $\Pr(Z = z) = \frac{1}{2}$, $\Pr(M(z) = 1 \mid Z = w) = \frac{1}{2}$, and $\Pr(M(z) = 1 \mid Z = b) = 1$. Note that $M = M(Z)$ and $Y = Y(Z, M)$, and so the above description fully defines the joint distribution on all the relevant variables.

Now, because $Y(z, 1) = 1$, we trivially have that $Y(z, 1) \perp\!\!\!\perp Z \mid M$, meaning that subset ignorability is satisfied. But, because $M(z) \not\perp\!\!\!\perp Z$, sequential ignorability is violated.

Case 3 (Consistency of Δ_n does not imply subset ignorability holds). At a high level, even if the potential outcomes $Y(z, 1)$ are not independent of Z —violating subset ignorability— Δ_n can still be a consistent estimator when there is appropriate cancellation. For a concrete illustration of this in the context of our two-stage arrest and charging application, consider a simple example in which: (a) there are no contextual variables (i.e., $X = 1$); (b) the population is evenly split across race groups (i.e., $\Pr(Z = z) = \frac{1}{2}$); (c) everyone in the population is arrested (i.e., $M = 1$); and (d) the prosecutor’s *potential* decisions depend on an arrestee’s *actual* race. Specifically, we set $Y(z, 0) = 0$

and $Y(z, 1)$ to be a Bernoulli random variable distributed as follows:

$$\Pr(Y(z, 1) = 1 \mid Z) = \begin{cases} 1 & z = b \wedge Z = b, \\ 0 & z = w \wedge Z = b, \\ \frac{1}{2} & Z = w. \end{cases} \quad (\text{F.9})$$

Because $Y = Y(Z, M)$, the above relationships completely specify the joint distribution of Y , Z , M , and X .

Subset ignorability is violated in this example since, by Eq. (F.9), $Y(z, 1) \not\perp\!\!\!\perp Z$. (Because X and M are constant, we need not condition on them when considering the subset ignorability criterion.) We further have,

$$\begin{aligned} \text{SATE}_M &= \mathbb{E}[Y(b, 1) \mid M = 1] - \mathbb{E}[Y(w, 1) \mid M = 1] \\ &= (\mathbb{E}[Y(b, 1) \mid Z = b] - \mathbb{E}[Y(w, 1) \mid Z = b]) \cdot \Pr(Z = b) \\ &\quad + (\mathbb{E}[Y(b, 1) \mid Z = w] - \mathbb{E}[Y(w, 1) \mid Z = w]) \cdot \Pr(Z = w) \\ &= (1 - 0) \cdot \frac{1}{2} + \left(\frac{1}{2} - \frac{1}{2}\right) \cdot \frac{1}{2} \\ &= \frac{1}{2}. \end{aligned}$$

Finally,

$$\begin{aligned} \lim_{n \rightarrow \infty} \Delta_n &\stackrel{\text{a.s.}}{=} \mathbb{E}[Y \mid Z = b, M = 1] - \mathbb{E}[Y \mid Z = w, M = 1] \\ &= 1 - \frac{1}{2} \\ &= \text{SATE}_M. \end{aligned}$$

Thus, even though subset ignorability is violated in this example, Δ_n yields a consistent estimate of the SATE_M .

Case 4 (Consistency of Δ_n does not imply treatment ignorability holds). Consider the example described in Case 2. As discussed there, subset ignorability is satisfied in that example and so, by Theorem 7.4, Δ_n is a consistent estimator of the SATE_M . However, that example does not satisfy treatment ignorability, as $M(z) \not\perp\!\!\!\perp Z$, contrary to Eq. (F.3). (Because X is constant, we need not condition on it when evaluating the treatment ignorability criterion.)

Case 5 (Consistency of Δ_n does not imply that treatment ignorability, mediator ignorability, and mediator monotonicity hold). This is directly implied by Case 4.

Case 6 (Treatment ignorability does not imply Δ_n is a consistent estimator of the SATE_M). We show, more generally, that the SATE_M is not identifiable under treatment ignorability alone. To do so, we construct a family of observationally equivalent examples that satisfy treatment ignorability but which have different values of SATE_M . As a result, no estimator—including Δ_n —can consistently estimate the SATE_M for every example in this family.

We construct the family of examples as follows. First, as in the other cases, we set $X = 1$, so that there are effectively no contextual variables, and we set $Y(z, 0) = 0$, meaning that if an individual were not arrested, that individual could not be charged. Second, we set $M(b) = 1$, meaning that everyone in the population would be arrested if they were Black. Finally, we set

$$\begin{aligned} \Pr(Y(z, 1) = y_{zm}, M(w) = m_w, Z = z) \\ = \Pr(Y(z, 1) = y_{zm} \mid M(w) = m_w) \cdot \Pr(M(w) = m_w) \cdot \Pr(Z = z), \end{aligned} \quad (\text{F.10})$$

where $\Pr(Z = z) = \frac{1}{2}$, $\Pr(M(w) = m_w) = \frac{1}{2}$, and, for $\alpha \in [0, 1]$,

$$\Pr(Y(z, 1) = 1 \mid M(w)) = \begin{cases} \alpha & M(w) = 0 \wedge z = w, \\ 1 & \text{otherwise.} \end{cases} \quad (\text{F.11})$$

The examples we construct thus, differ only in the choice of α .

Now, regardless of α , these examples all satisfy treatment ignorability. To see this, note that $M(w) \perp\!\!\!\perp Z$ by Eq. (F.10) and $M(b) \perp\!\!\!\perp Z$ since $M(b)$ is constant. Consequently, the first condition of treatment ignorability is satisfied. Eq. (F.10) further implies that $Y(z, 1) \perp\!\!\!\perp Z \mid M(w)$ and, since $Y(z, 0)$ is constant, $Y(z, 0) \perp\!\!\!\perp Z \mid M(w)$, establishing the second condition of treatment ignorability. (Because $M(b)$ and X are constant, we need not condition on them when considering the two treatment ignorability conditions.)

We next show that all these examples are observationally equivalent. Intuitively, observational equivalence stems from the fact that the only difference between the examples is in the distribution of $Y(w, 1)$ for those individuals with $M(w) = 0$. But for those with $M(w) = 0$, who would not be arrested if they were white, we never observe $Y(w, 1)$.

Now, to rigorously establish observational equivalence, we must show that $\Pr(X = x, Y = y, Z =$

$z \mid M = 1$) does not depend on the value of α . Because X is constant, we need only consider $\Pr(Y = y, Z = z \mid M = 1)$. First, observe that

$$\begin{aligned}\Pr(M = 1) &= \Pr(M(w) = 1, Z = w) + \Pr(M(b) = 1, Z = b) \\ &= \Pr(M(w) = 1) \cdot \Pr(Z = w) + \Pr(Z = b) \\ &= \frac{3}{4}.\end{aligned}$$

Further, note that

$$\Pr(Y = y, Z = z, M = 1) = \Pr(Y(z, 1) = y, Z = z, M(z) = 1),$$

and consider the case $z = b$. Then, because $Y(b, 1) = 1$ and $M(b) = 1$,

$$\Pr(Y = y, Z = b, M = 1) = \begin{cases} 0 & y = 0, \\ \frac{1}{2} & y = 1. \end{cases} \quad (\text{F.12})$$

Now consider the case $z = w$. By Eq. (F.10),

$$\begin{aligned}\Pr(Y(w, 1) = y, Z = w, M(w) = 1) \\ &= \Pr(Y(w, 1) = y \mid M(w) = 1) \cdot \Pr(M(w) = 1) \cdot \Pr(Z = w) \\ &= \Pr(Y(w, 1) = y \mid M(w) = 1) \cdot \frac{1}{4}.\end{aligned}$$

By Eq. (F.11), $\Pr(Y(w, 1) = 1 \mid M(w) = 1) = 1$, and so,

$$\Pr(Y = y, Z = w, M = 1) = \begin{cases} 0 & y = 0, \\ \frac{1}{4} & y = 1. \end{cases} \quad (\text{F.13})$$

Finally, combining Eqs. (F.12) and (F.13) with the fact that $\Pr(M = 1) = \frac{3}{4}$, we have

$$\Pr(Y = y, Z = z \mid M = 1) = \begin{cases} 0 & y = 0, \\ \frac{2}{3} & y = 1 \wedge z = b, \\ \frac{1}{3} & y = 1 \wedge z = w. \end{cases}$$

In particular, $\Pr(Y = y, Z = z \mid M = 1)$ does not depend on α , and so the examples are all observationally equivalent.

We conclude the proof by showing that the SATE_M differs across these examples. First, it remains to calculate $\Pr(M(w) = m_w \mid M = 1)$. To do so, note that

$$\begin{aligned}
\Pr(M(w) = 1, M = 1) &= \Pr(M(w) = 1, Z = w) + \Pr(M(w) = 1, M(b) = 1, Z = b) \\
&= \Pr(M(w) = 1) \cdot \Pr(Z = w) \\
&\quad + \Pr(M(w) = 1) \cdot \Pr(M(b) = 1) \cdot \Pr(Z = b) \\
&= \frac{1}{2} \cdot \frac{1}{2} + \frac{1}{2} \cdot 1 \cdot \frac{1}{2} \\
&= \frac{1}{2},
\end{aligned}$$

and so, since $\Pr(M = 1) = \frac{3}{4}$, it follows that $\Pr(M(w) = 1 \mid M = 1) = \frac{2}{3}$.

Consequently, we have

$$\begin{aligned}
\text{SATE}_M &= \mathbb{E}[Y(b, 1) \mid M = 1] - \mathbb{E}[Y(w, 1) \mid M = 1] \\
&= \Pr(M(w) = 1 \mid M = 1) \\
&\quad \cdot (\mathbb{E}[Y(b, 1) \mid M(w) = 1, M = 1] - \mathbb{E}[Y(w, 1) \mid M(w) = 1, M = 1]) \\
&\quad + \Pr(M(w) = 0 \mid M = 1) \\
&\quad \cdot (\mathbb{E}[Y(b, 1) \mid M(w) = 0, M = 1] - \mathbb{E}[Y(w, 1) \mid M(w) = 0, M = 1]) \\
&= \frac{2}{3} \cdot (1 - 1) + \frac{1}{3} \cdot (1 - \mathbb{E}[Y(w, 1) \mid M(w) = 0, Z = b]) \\
&= \frac{1 - \alpha}{3},
\end{aligned}$$

where second to last equality follows from Eq. (F.11) and the fact that the event $\{M(w) = 0 \wedge M = 1\}$ equals $\{M(w) = 0 \wedge Z = b\}$; the final equality also follows from Eq. (F.11), as well as the fact that $Y(z, 1) \perp\!\!\!\perp Z \mid M(w)$. We have thus, constructed a family of observationally equivalent examples that satisfy treatment ignorability but which have different SATE_M , implying that the SATE_M is not in general identifiable under treatment ignorability alone.

Case 7 (Treatment, mediator ignorability, and mediator monotonicity jointly imply Δ_n is a consistent estimator of the SATE_M). The proof is in two pieces. First, we derive an expression for the SATE_M holding X constant, and then prove the general claim.

Supposing $X = x$ is constant, recall that by definition $M = 1$ if and only if $M(z) = 1$ where $Z = z$. By mediator monotonicity, $M(b) \geq M(w)$. Therefore, the event $\{M = 1\}$ can be partitioned

into the following two events:

- $E_1 = \{M(b) = 1 \wedge M(w) = 1\}$,
- $E_2 = \{Z = b \wedge M(b) = 1 \wedge M(w) = 0\}$.

Recall the definition of the SATE_M in Definition 7.1. It follows from the law of total expectation that:

$$\begin{aligned} \text{SATE}_M &= \mathbb{E}[Y(b, 1) - Y(w, 1) \mid M = 1] \\ &= \mathbb{E}[Y(b, 1) - Y(w, 1) \mid E_1] \cdot \Pr(E_1 \mid M = 1) \\ &\quad + \mathbb{E}[Y(b, 1) - Y(w, 1) \mid E_2] \cdot \Pr(E_2 \mid M = 1) \end{aligned} \tag{F.14}$$

Now, we examine each of these summands in turn. First, consider the E_1 term:

$$\mathbb{E}[Y(b, 1) - Y(w, 1) \mid E_1] = \mathbb{E}[Y(b, 1) \mid E_1] - \mathbb{E}[Y(w, 1) \mid E_1]$$

By the definition of $E_1 = \{M(b) = 1 \wedge M(w) = 1\}$ and the second treatment ignorability condition, Eq. (F.4), we are free to condition both terms on the right hand side by levels of Z , yielding

$$\mathbb{E}[Y(b, 1) \mid Z = b, E_1] - \mathbb{E}[Y(w, 1) \mid Z = w, E_1] = \mathbb{E}[Y \mid Z = b, E_1] - \mathbb{E}[Y \mid Z = w, E_1], \tag{F.15}$$

where equality follows from replacing potential outcomes by their realized values according to the definition of $Y = Y(M, Z)$.

Next, consider the E_2 term. Again,

$$\mathbb{E}[Y(b, 1) - Y(w, 1) \mid E_2] = \mathbb{E}[Y(b, 1) \mid E_2] - \mathbb{E}[Y(w, 1) \mid E_2].$$

It follows from mediator ignorability, Eq. (F.5), and the definition of E_2 that

$$\begin{aligned} \mathbb{E}[Y(w, 1) \mid E_2] &= \mathbb{E}[Y(w, 1) \mid Z = b, M(b) = 1, M(w) = 0] \\ &= \mathbb{E}[Y(w, 1) \mid Z = b, M(b) = 1, M(w) = 1] \\ &= \mathbb{E}[Y(w, 1) \mid Z = w, M(b) = 1, M(w) = 1], \end{aligned}$$

where the last equality follows from treatment ignorability, Eq. (F.4). Replacing potential outcomes

with their realizations, it follows that

$$\mathbb{E}[Y(b, 1) - Y(w, 1) \mid E_2] = \mathbb{E}[Y \mid Z = b, E_2] - \mathbb{E}[Y \mid Z = w, E_1]. \quad (\text{F.16})$$

Now, we substitute Eqs. (F.15) and (F.16) into Eq. (F.14).

$$\begin{aligned} \text{SATE}_M &= (\mathbb{E}[Y \mid Z = b, E_1] - \mathbb{E}[Y \mid Z = w, E_1]) \cdot \Pr(E_1 \mid M = 1) \\ &\quad + (\mathbb{E}[Y \mid Z = b, E_2] - \mathbb{E}[Y \mid Z = w, E_1]) \cdot \Pr(E_2 \mid M = 1) \\ &= (\mathbb{E}[Y \mid Z = b, E_1] - \mathbb{E}[Y \mid Z = w, M = 1]) \cdot \Pr(E_1 \mid M = 1) \\ &\quad + (\mathbb{E}[Y \mid Z = b, E_2] - \mathbb{E}[Y \mid Z = w, M = 1]) \cdot \Pr(E_2 \mid M = 1) \\ &= (\mathbb{E}[Y \mid Z = b, E_1] \cdot \Pr(E_1 \mid M = 1) + \mathbb{E}[Y \mid Z = b, E_2] \cdot \Pr(E_2 \mid M = 1)) \\ &\quad - (\mathbb{E}[Y \mid Z = w, M = 1] \cdot (\Pr(E_1 \mid M = 1) + \Pr(E_2 \mid M = 1))) \\ &= \mathbb{E}[Y \mid Z = b, M = 1] - \mathbb{E}[Y \mid Z = w, M = 1], \end{aligned} \quad (\text{F.17})$$

where the second equality follows from the fact that $\{M = 1 \wedge Z = w\} = \{E_1 \wedge Z = w\}$ by mediator monotonicity, and the last equality follows from the facts that $\{M = 1 \wedge Z = b \wedge E_1\} = \{Z = b \wedge E_1\}$, $\{M = 1 \wedge Z = b \wedge E_2\} = \{Z = b \wedge E_2\}$, and $\Pr(E_1 \mid M = 1) + \Pr(E_2 \mid M = 1) = 1$.

Now, suppose that X is not constant. Conditioning Y , Z , and M on $X = x$, it follows from the law of total expectation that

$$\begin{aligned} &\mathbb{E}[Y(b, 1) - Y(w, 1) \mid M = 1] \\ &= \sum_x \mathbb{E}[Y(b, 1) - Y(w, 1) \mid M = 1, X = x] \cdot \Pr(X = x \mid M = 1) \\ &= \sum_x \mathbb{E}[Y \mid Z = b, M = 1, X = x] \cdot \Pr(X = x \mid M = 1) \\ &\quad - \mathbb{E}[Y \mid Z = w, M = 1, X = x] \cdot \Pr(X = x \mid M = 1), \end{aligned} \quad (\text{F.18})$$

where the second equality follows from Eq. (F.17), using the fact that X is constant on each of the events $\{X = x\}$. Eq. (F.18) is identical to the expression in the statement of Theorem 7.4, and so the estimator Δ_n converges almost surely to the quantity on the right-hand side of Eq. (F.18) by precisely the same argument as there.

□




Figure F.1: A causal DAG considered by Knox et al. In the context of our charging example, Z indicates race, M indicates arrest decisions, Y indicates charging decisions, and Q is an unobserved confounder. Even when one restricts to distributions compatible with this DAG, the Knox et al. conditions are not necessary to nonparametrically identify the SATE_M from data on second-stage decisions.

F.2 Analysis of a Restricted Family of Distributions

Theorem F.5 shows that treatment ignorability, mediator ignorability, and mediator monotonicity are jointly sufficient but not necessary to identify the SATE_M from data on second-stage decisions. We show that this non-necessity holds even if we restrict to distributions compatible with a particular causal DAG considered by Knox et al., shown in Figure F.1, where an unobserved confounder Q directly influences the first-stage decisions M (e.g., arrests) and the second-stage decisions Y (e.g., charging). To do so, we explicitly construct a counterexample in which: (a) the joint distribution of random variables is compatible with this causal DAG; (b) mediator ignorability is violated; and (c) subset ignorability is satisfied, which in turn implies that the stratified difference-in-means Δ_n is a consistent estimator of the SATE_M , by Theorem 7.4.

Proposition F.6. *There exists a structural causal model (SCM) compatible with the causal DAG in Figure F.1 which violates mediator ignorability but satisfies subset ignorability.*

Proof. We start by explicitly constructing an SCM that is (faithfully) compatible with the DAG in Figure F.1. Our SCM has the following independent exogenous variables:

$$U_Z \sim \text{UNIF}(\{w, b\}),$$

$$U_Q \sim \text{UNIF}(\{1, 2, 3, 4\}),$$

$$U_M \sim \text{UNIF}((0, 1)),$$

$$U_Y \sim \text{UNIF}((0, 1)),$$

where U_Z and U_Q are uniformly distributed over the specified discrete sets, and U_M and U_Y are

uniform over the unit interval. Now, the structural equations are given by:

$$\begin{aligned}
f_Z(u_z) &= u_z, \\
f_Q(u_q) &= u_q, \\
f_M(z, q, u_m) &= \mathbb{1} \left(u_m \leq \frac{1}{2} \cdot (1 + \mathbb{1}(z = b)) \right. \\
&\quad \left. \cdot (\mathbb{1}(q = 1) + \mathbb{1}(z = b \wedge q = 3) + \mathbb{1}(z = w \wedge q = 2)) \right), \\
f_Y(z, m, q, u_y) &= m \cdot \mathbb{1} \left(u_y \leq (1 + \mathbb{1}(z = b)) \cdot \frac{\mathbb{1}(q = 1)}{2} \right),
\end{aligned}$$

where $\mathbb{1}$ denotes the indicator function and $\wedge$ denotes conjunction (i.e., the **and** operator). For avoidance of doubt, $Z = f_Z(U_Z)$, $Q = f_Q(U_Q)$, $M = f_M(Z, Q, U_M)$, and $Y = f_Y(Z, M, Q, U_Y)$. Further, the potential arrest outcomes are given by $M(z) = f_M(z, Q, U_M)$, and the bivariate potential charge outcomes are given by $Y(z, m) = f_Y(z, m, Q, U_Y)$.

Mediator ignorability is violated. First, note that $Z \perp\!\!\!\perp \{Y(b, 1), M(w), M(b)\}$, because Z is a function of U_Z , and $\{Y(b, 1), M(w), M(b)\}$ are functions of U_Y , U_Q , and U_M , which are jointly independent of U_Z . Now, applying this fact and conditioning on Q , we have that,

$$\begin{aligned}
&\Pr(Y(b, 1) = 1 \mid M(w) = m_w, M(b) = 1, Z = z) \\
&= \Pr(Y(b, 1) = 1 \mid M(w) = m_w, M(b) = 1) \\
&= \sum_{q=1}^4 \Pr(Y(b, 1) = 1 \mid M(w) = m_w, M(b) = 1, Q = q) \\
&\quad \cdot \Pr(Q = q \mid M(w) = m_w, M(b) = 1) \\
&= \sum_{q=1}^4 \Pr(f_Y(b, 1, q, U_Y) = 1 \mid M(w) = m_w, M(b) = 1, Q = q) \\
&\quad \cdot \Pr(Q = q \mid M(w) = m_w, M(b) = 1).
\end{aligned}$$

Next, observe that $f_Y(b, 1, q, U_Y) = \mathbb{1}(q = 1)$, and so

$$\begin{aligned}
& \Pr(Y(b, 1) = 1 \mid M(w) = m_w, M(b) = 1, Z = z) \\
&= \Pr(Q = 1 \mid M(w) = m_w, M(b) = 1) \\
&= \frac{\Pr(M(w) = m_w, M(b) = 1 \mid Q = 1) \cdot \Pr(Q = 1)}{\sum_{q=1}^4 \Pr(M(w) = m_w, M(b) = 1 \mid Q = q) \cdot \Pr(Q = q)} \\
&= \frac{\Pr(M(w) = m_w, M(b) = 1 \mid Q = 1)}{\sum_{q=1}^4 \Pr(M(w) = m_w, M(b) = 1 \mid Q = q)}. \tag{F.19}
\end{aligned}$$

The second equality above follows from Bayes' rule, and the third follows from the fact that $\Pr(Q = q) = 1/4$.

Finally, we compute $\Pr(M(w) = m_w, M(b) = 1 \mid Q = q)$. Note that

$$\begin{aligned}
M(w) &= f_M(w, Q, U_M) \\
&= \mathbb{1} \left(U_M \leq \frac{\mathbb{1}(Q = 1) + \mathbb{1}(Q = 2)}{2} \right) \\
&= \begin{cases} \mathbb{1}(U_M \leq 1/2) & Q \in \{1, 2\}, \\ 0 & \text{otherwise.} \end{cases} \tag{F.20}
\end{aligned}$$

Likewise,

$$\begin{aligned}
M(b) &= f_M(b, Q, U_M) \\
&= \mathbb{1}(U_M \leq (\mathbb{1}(Q = 1) + \mathbb{1}(Q = 3))) \\
&= \begin{cases} 1 & Q \in \{1, 3\}, \\ 0 & \text{otherwise.} \end{cases} \tag{F.21}
\end{aligned}$$

As a result,

$$\Pr(M(w) = m_w, M(b) = 1 \mid Q = q) = \begin{cases} 1/2 & q = 1, \\ 1 & q = 3 \wedge m_w = 0, \\ 0 & \text{otherwise.} \end{cases}$$

Thus, by Eq. (F.19),

$$\Pr(Y(b, 1) = 1 \mid M(w) = 1, M(b) = 1, Z = z) = 1,$$

while

$$\Pr(Y(b, 1) = 1 \mid M(w) = 0, M(b) = 1, Z = z) = \frac{1}{3}.$$

Therefore, $Y(b, 1) \not\perp\!\!\!\perp M(w) \mid M(b) = 1, Z = z$, meaning that mediator ignorability does not hold.

Subset ignorability holds. Similar to the above, we have that

$$\begin{aligned} Y(b, 1) &= f_Y(b, 1, Q, U_Y) \\ &= \mathbb{1}(U_Y \leq \mathbb{1}(Q = 1)) \\ &= \begin{cases} 1 & Q = 1 \\ 0 & \text{otherwise,} \end{cases} \end{aligned} \tag{F.22}$$

and

$$\begin{aligned} Y(w, 1) &= f_Y(w, 1, Q, U_Y) \\ &= \mathbb{1}(U_Y \leq \mathbb{1}(Q = 1)/2) \\ &= \begin{cases} \mathbb{1}(U_Y \leq 1/2) & Q = 1 \\ 0 & \text{otherwise.} \end{cases} \end{aligned} \tag{F.23}$$

Now, as before, $Z \perp\!\!\!\perp \{Q, M(w), M(b)\}$, since Z is a function of U_Z , and $\{Q, M(w), M(b)\}$ are functions of U_Q and U_M , which are jointly independent of U_Z . Consequently,

$$\begin{aligned} \Pr(Q = 1 \mid M = 1, Z = z) &= \Pr(Q = 1 \mid M(z) = 1, Z = z) \\ &= \Pr(Q = 1 \mid M(z) = 1) \\ &= \frac{\Pr(M(z) = 1 \mid Q = 1) \cdot \Pr(Q = 1)}{\sum_{q=1}^4 \Pr(M(z) = 1 \mid Q = q) \cdot \Pr(Q = q)} \\ &= \frac{1}{2}, \end{aligned} \tag{F.24}$$

where the last equality follows from Eqs. (F.20) and (F.21), together with the fact that $\Pr(Q = q) = 1/4$, and that $U_M \perp\!\!\!\perp Q$.

Finally, conditioning on Q , we have

$$\begin{aligned}
& \Pr(Y(b, 1) = 1 \mid M = 1, Z = z) \\
&= \sum_{q=1}^4 \Pr(Y(b, 1) = 1 \mid M = 1, Z = z, Q = q) \cdot \Pr(Q = q \mid M = 1, Z = z) \\
&= \Pr(Q = 1 \mid M = 1, Z = z) \\
&= \frac{1}{2},
\end{aligned}$$

where the second equality follows from Eq. (F.22), and the third from Eq. (F.24). Similarly,

$$\begin{aligned}
& \Pr(Y(w, 1) = 1 \mid M = 1, Z = z) \\
&= \sum_{q=1}^4 \Pr(Y(w, 1) = 1 \mid M = 1, Z = z, Q = q) \cdot \Pr(Q = q \mid M = 1, Z = z) \\
&= \Pr(U_Y \leq 1/2 \mid M = 1, Z = z, Q = 1) \cdot \Pr(Q = 1 \mid M = 1, Z = z) \\
&= \Pr(U_Y \leq 1/2) \cdot \Pr(Q = 1 \mid M = 1, Z = z) \\
&= \frac{1}{4},
\end{aligned}$$

where the second equality follows from Eq. (F.23), the third from the fact that $U_Y \perp\!\!\!\perp \{M, Z, Q\}$, and the fourth from Eq. (F.24). Therefore, $\Pr(Y(b, 1) = y \mid M = 1, Z = b) = \Pr(Y(b, 1) \mid M = 1, Z = w)$ and similarly for $Y(w, 1)$. In particular, this means that $Y(z, 1) \perp\!\!\!\perp Z \mid M = 1$, and so subset ignorability holds. $\square$

F.3 Extending Theorem 7.4 to Allow for Continuous Covariates

Theorem 7.4 in the main text shows that subset ignorability—together with overlap—implies the SATE_M is nonparametrically identified, where, for simplicity, we proved the result for discrete covariates X . We now extend that result to allow for continuous covariates. At a conceptual level, the extension is straightforward: we first condition on X , then appeal to subset ignorability to condition on Z , and, finally, use consistency to replace potential outcomes by their observed values. In the general case, however, typically $\Pr(X = x) = 0$, and so one must take care to define expressions that nominally condition on these probability-zero events.

Recall that in the discrete case, the primary conditional expectations, treated as functions of z

and x , are of the form

$$\begin{aligned}\mathbb{E}[Y|Z = z, X = x, M = 1] &= \sum_y y \frac{\Pr(Y = y, Z = z, X = x \mid M = 1)}{\Pr(Z = z, X = x \mid M = 1)} \\ &= \sum_y y \frac{\Pr(Y = y, Z = z, X = x \mid M = 1)}{\Pr(Z = z \mid X = x, M = 1) \Pr(X = x \mid M = 1)}.\end{aligned}\quad (\text{F.25})$$

Overlap ensures that the denominator in (F.25) is nonzero, and, accordingly, that the conditional expectation is well-defined. In the continuous case, to address conditioning on probability-zero events, conditional probabilities are defined as random variables rather than simple numeric quantities (cf. Billingsley (2008)). Further, if the random variables $\Pr(Z = z \mid X, M = 1) > 0$ a.s. for $z \in \{w, b\}$ —a condition that we call generalized overlap—then the expression $\mathbb{E}[Y \mid Z = z, X = x, M = 1]$ is a well-defined function of z and x , as in the discrete case, up to a set of measure zero with respect to the pushforward measure $\mu_{X|M=1}$ for each fixed z .^{2,3}

We now state and prove the extension of Theorem 7.4, with the understanding that the conditional probabilities and expectations below are defined according to the usual measure-theoretic conventions.

Theorem F.7. *Suppose $Y(z, 1)$, Z , M , and X satisfy subset ignorability, and that generalized overlap holds—that is, for $z \in \{b, w\}$, $\Pr(Z = z \mid X, M = 1) > 0$ a.s. Then, the SATE_M equals*

$$\begin{aligned}&\int_{\mathcal{X}} \mathbb{E}[Y \mid Z = b, X = x, M = 1] dF_{X|M=1} \\ &\quad - \int_{\mathcal{X}} \mathbb{E}[Y \mid Z = w, X = x, M = 1] dF_{X|M=1},\end{aligned}\quad (\text{F.26})$$

where $\mathcal{X}$ denotes the range of X and $dF_{X|M=1}$ denotes integration over $\mathcal{X}$ with respect to the pushforward measure $\mu_{X|M=1}$.

²The pushforward measure $\mu_{X|M=1}$ is the measure on $\mathcal{X}$ —the range of X —given by $\mu_{X|M=1}[A] = \Pr(X \in A \mid M = 1)$ for measurable $A \subseteq \mathcal{X}$.

³To see this, first note that, in general, $\mathbb{E}[Y \mid Z = z, X = x, M = 1]$ is uniquely defined up to a set of measure zero with respect to the pushforward measure $\mu_{Z,X|M=1}$. Now, for fixed z , suppose, toward a contradiction, that $f_1(x)$ and $f_2(x)$ are two versions of $\mathbb{E}[Y \mid Z = z, X = x, M = 1]$ that differ on a set A such that $\Pr(X \in A \mid M = 1) > 0$. Then, by the generalized overlap condition, $\Pr(Z = z, X \in A \mid M = 1) = \int_A \Pr(Z = z \mid X = x, M = 1) dF_{X|M=1} > 0$, contradicting the fact that $f_1(x) \neq f_2(x)$ only on a null set with respect to the pushforward measure $\mu_{Z,X|M=1}$.

Proof. By conditioning on X , we have,

$$\begin{aligned}
\text{SATE}_M &= \mathbb{E}[Y(b, 1) - Y(w, 1) \mid M = 1] \\
&= \int_{\mathcal{X}} \mathbb{E}[Y(b, 1) - Y(w, 1) \mid X = x, M = 1] dF_{X|M=1} \\
&= \int_{\mathcal{X}} \mathbb{E}[Y(b, 1) \mid X = x, M = 1] - \mathbb{E}[Y(w, 1) \mid X = x, M = 1] dF_{X|M=1}. \tag{F.27}
\end{aligned}$$

Now, subset ignorability gives that

$$\mathbb{E}[Y(z, 1) \mid X = x, M = 1] = \mathbb{E}[Y(z, 1) \mid X = x, Z = z, M = 1] \text{ a.s.}, \tag{F.28}$$

where generalized overlap ensures that the right-hand side of Eq. (F.28) is well-defined up to a set of measure zero with respect to $dF_{X|M=1}$. Substituting this expression into Eq. (F.27), and then appealing to consistency to replace potential outcomes with their observed values, we have

$$\begin{aligned}
\text{SATE}_M &= \int_{\mathcal{X}} \mathbb{E}[Y(b, 1) \mid X = x, Z = b, M = 1] \\
&\quad - \mathbb{E}[Y(w, 1) \mid X = x, Z = w, M = 1] dF_{X|M=1} \\
&= \int_{\mathcal{X}} \mathbb{E}[Y(Z, M) \mid X = x, Z = b, M = 1] \\
&\quad - \mathbb{E}[Y(Z, M) \mid X = x, Z = w, M = 1] dF_{X|M=1} \\
&= \int_{\mathcal{X}} \mathbb{E}[Y \mid X = x, Z = b, M = 1] \\
&\quad - \mathbb{E}[Y \mid X = x, Z = w, M = 1] dF_{X|M=1}.
\end{aligned}$$

□

All of the quantities in Eq. (F.26) (i.e., the distribution of X and the conditional expectations) are functions of observables, establishing that the SATE_M is identified by data on second-stage decisions. One may adopt a variety of approaches to estimate the terms in Eq. (F.26), including model-based strategies, as we do in Section 7.4. One may also adopt nonparametric estimation strategies, wherein continuous covariates are appropriately binned into discrete sets. For further treatment of these issues, see, for example, Gelman et al. (2013), Friedman et al. (2001), and Tsybakov (2008).


F.4 Summary Statistics of Prosecution Dataset

We present summary statistics, disaggregated by demographic group, of the dataset used to conduct the empirical analysis of prosecutorial charging decisions in Section 7.4.


Table F.1: *Breakdown of charges, prior arrests, and prior convictions, and weapons involvement among individuals arrested for a felony offense between 2013 and 2019 in a major U.S. county, as analyzed in Section 7.4. Note that only Black and non-Hispanic white individuals are analyzed in our race-based analysis, so within-row sums differ between race and gender.*

		Female	Male	Black	White
Charge:	arson	50	218	79	124
	assault	2,062	9,160	4,131	3,662
	burglary	664	3,334	1,762	1,535
	burglary (auto)	125	1,216	737	338
	driving under the influence	26	90	17	50
	drug-related offense	736	5,643	2,815	1,881
	forcible rape	3	112	49	26
	forgery	60	187	103	84
	hit-and-run	10	42	18	17
	kidnapping	4	24	12	1
	manslaughter (vehicular)	2	8	4	3
	motor vehicle theft	83	323	161	141
	other felony offense	287	2,115	877	910
	other sex offense	8	180	85	54
	robbery	458	2,186	1,536	589
	theft	877	3,385	1,723	1,539
	weapons offense	80	882	513	218
	willful homicide	53	178	65	60
Prior arrests:	0	3,192	13,491	5,881	5,682
	1	1,005	5,486	2,972	1,977
	2	553	3,315	1,929	1,134

		Female	Male	Black	White
	3	321	2,199	1,215	786
	4	191	1,558	895	540
	5	116	1,066	591	364
	6	78	697	381	242
	7	40	469	248	157
	8	29	316	180	103
	9	18	204	111	84
	10+	45	482	283	163
Prior convictions:	0	4,451	20,775	9,937	8,290
	1	748	5,221	2,961	1,773
	2	278	2,055	1,111	730
	3	78	766	429	261
	4	23	312	171	116
	5+	10	154	77	62
Weapons:	no weapon involved	4,442	24,932	12,810	9,434
	weapon involved	1,146	4,351	1,876	1,798



(a) Distribution of age by gender



(b) Distribution of age by race

Figure F.2: *Breakdown of age by gender and by race of individuals arrested for a felony offense between 2013 and 2019 in a major U.S. county, as analyzed in Section 7.4.*

Appendix to Chapter 8

Table G.1: To corroborate our main analysis, we compare the ratio of stopped drivers per capita by race with our regression coefficients. We note that within a single city, the per capita ratio of stopped drivers is not expected to be unity due to white drivers on average driving more (Kim, W., V. Anorve, and B. C. Tefft. “American Driving Survey, 2014–2017.” (2019)). Furthermore, neither Oklahoma City nor Houston coded Hispanic drivers in their stop data, and so we excluded them here. We find that proportionally more stopped speeders tend to be white in cities with negative values for β_{RACE} and proportionally more tend to be non-white in cities with positive β_{RACE} coefficients. This is consistent with the interpretation of greater values of β_{RACE} as concentration of speeding enforcement on speeding by non-white drivers.


City	Black/white ratio	Hispanic/white ratio	Coefficient (SE)
Mesa	1.54	0.47	2.16 (0.60)
Chicago	2.44	0.79	1.92 (0.093)
Plano	1.55	0.50	1.06 (1.14)
San Antonio	1.12	0.92	1.03 (0.27)
Aurora	0.99	0.06	0.49 (0.55)
Madison	2.28	1.00	0.40 (0.40)
Tulsa	0.72	0.04	-0.62 (0.39)
Oklahoma City	0.83		-0.83 (0.37)
Wichita	0.96	0.58	-1.08 (0.48)
Houston	0.59		-1.18 (0.23)

Table G.2: Regression table for the main analysis. The variable “Speeding over 15 KPH” is the normalized percent of time drivers spent driving in a beat at least 15 KPH over the speed limit, and “% Non-White” is the percentage of residents of a beat who are not non-Hispanic white. City names represent city fixed effects, years represent year fixed effects, and interactions are indicated by “×.”


Term	Estimate	Standard Error	P-value
Speeding Over 15 KPH	0.542	0.018	0.000
Aurora	0.420	0.371	0.258
Chicago	-4.586	0.098	0.000
Houston	1.931	0.205	0.000

Term	Estimate	Standard Error	P-value
Madison	0.826	0.146	0.000
Mesa	-0.306	0.320	0.338
Oklahoma	0.821	0.250	0.001
Plano	2.451	0.542	0.000
San Antonio	-1.060	0.234	0.000
Tulsa	0.578	0.244	0.018
Wichita	1.361	0.268	0.000
Aurora $\times$ % Non-White	0.492	0.549	0.370
Chicago $\times$ % Non-White	1.918	0.093	0.000
Houston $\times$ % Non-White	-1.183	0.228	0.000
Madison $\times$ % Non-White	0.395	0.399	0.321
Mesa $\times$ % Non-White	2.159	0.595	0.000
Oklahoma $\times$ % Non-White	-0.833	0.374	0.026
Plano $\times$ % Non-White	1.058	1.138	0.353
San Antonio $\times$ % Non-White	1.032	0.274	0.000
Tulsa $\times$ % Non-White	-0.617	0.392	0.115
Wichita $\times$ % Non-White	-1.076	0.483	0.026
Aurora $\times$ 2014	-0.105	0.296	0.722
Chicago $\times$ 2014	0.418	0.098	0.000
Houston $\times$ 2014	0.227	0.155	0.144
Madison $\times$ 2014	-0.200	0.156	0.200
Mesa $\times$ 2014	-0.127	0.272	0.641
Oklahoma $\times$ 2014	0.007	0.234	0.976
Plano $\times$ 2014	0.025	0.320	0.937
San Antonio $\times$ 2014	-0.146	0.146	0.320
Tulsa $\times$ 2014	-0.171	0.232	0.461
Wichita $\times$ 2014	0.075	0.256	0.771
Aurora $\times$ 2015	-0.231	0.296	0.435

Term	Estimate	Standard Error	P-value
Chicago $\times$ 2015	-0.022	0.098	0.819
Houston $\times$ 2015	0.072	0.155	0.644
Madison $\times$ 2015	0.072	0.156	0.646
Mesa $\times$ 2015	0.253	0.272	0.352
Oklahoma $\times$ 2015	-0.162	0.234	0.488
San Antonio $\times$ 2015	-0.336	0.147	0.022
Tulsa $\times$ 2015	-0.381	0.232	0.100
Wichita $\times$ 2015	0.100	0.256	0.695
Aurora $\times$ 2016	-0.362	0.296	0.221
Chicago $\times$ 2016	0.293	0.098	0.003
Houston $\times$ 2016	-0.045	0.155	0.770
Madison $\times$ 2016	-0.202	0.156	0.196
Oklahoma $\times$ 2016	-0.129	0.234	0.583
San Antonio $\times$ 2016	-0.214	0.146	0.143
Tulsa $\times$ 2016	-0.042	0.232	0.857
Wichita $\times$ 2016	0.023	0.256	0.929
Chicago $\times$ 2017	0.649	0.097	0.000
Oklahoma $\times$ 2017	-0.424	0.234	0.070
San Antonio $\times$ 2017	-0.387	0.147	0.008
Chicago $\times$ 2018	0.544	0.097	0.000
Chicago $\times$ 2019	0.537	0.097	0.000



(a) Absolute speeding thresholds



(b) Relative speeding thresholds

Figure G.1: Robustness of our main analysis to the definition of speeding. We show, for each city, the value of β_{RACE} as we vary the threshold above the speed limit we consider speeding when we compute d_i , on an absolute (G.1a) and relative (G.1b) scale. We find that our estimates of the coefficients are largely stable as we shift the threshold.





Figure G.2: Estimates of β_{RACE} in our speeding violation analysis using quasipoisson instead of negative binomial regression. We find that our results are largely unchanged, except that Houston moves from having a negative coefficient in the main model to 0.036 (SE: 0.181, $p = 0.843$).



(a) Demographics.

(b) Speeding violations and stops.

Figure G.3: A comparison of Aurora, CO's demographic composition, the location of speeding events, as measured by telematics data, and the location of speeding stops. (G.3a): The demographics of residents, where red points represent non-Hispanic white residents and blue points represent non-white residents. (G.3b): The density of speeding stops recorded by police officers, indicated by gray dots, where the area of the points is proportional to the number of stops at a single location; stops for speeding are overlaid with a heatmap of speeding events, as estimated from telematics data.






Figure G.4: The demographics of residents of Chicago, IL, where red points represent non-Hispanic white residents, green points Black residents, blue points Hispanic residents, purple points Asian and Pacific Islander residents, and orange points all other residents. (Traffic stops in Chicago indicate beat but were not geocoded.)




(a) Demographics.




(b) Speeding violations and stops.

Figure G.5: A comparison of Houston, TX's demographic composition, the location of speeding events, as measured by telematics data, and the location of speeding stops. (G.5a): The demographics of residents, where red points represent non-Hispanic white residents and blue points represent non-white residents. (G.5b): The density of speeding stops recorded by police officers, indicated by gray dots, where the area of the points is proportional to the number of stops at a single location; stops for speeding are overlaid with a heatmap of speeding events, as estimated from telematics data.




(a) Demographics.



(b) Speeding violations and stops.


Figure G.6: A comparison of Madison, WI's demographic composition, the location of speeding events, as measured by telematics data, and the location of speeding stops. (G.6a): The demographics of residents, where red points represent non-Hispanic white residents and blue points represent non-white residents. (G.6b): The density of speeding stops recorded by police officers, indicated by gray dots, where the area of the points is proportional to the number of stops at a single location; stops for speeding are overlayed with a heatmap of speeding events, as estimated from telematics data.




(a) Demographics.

(b) Speeding violations and stops.

Figure G.7: A comparison of Oklahoma City, OK's demographic composition, the location of speeding events, as measured by telematics data, and the location of speeding stops. (G.7a): The demographics of residents, where red points represent non-Hispanic white residents and blue points represent non-white residents. (G.7b): The density of speeding stops recorded by police officers, indicated by gray dots, where the area of the points is proportional to the number of stops at a single location; stops for speeding are overlayed with a heatmap of speeding events, as estimated from telematics data.




(a) Demographics.




(b) Speeding violations and stops.

Figure G.8: A comparison of Plano, TX's demographic composition, the location of speeding events, as measured by telematics data, and the location of speeding stops. (G.8a): The demographics of residents, where red points represent non-Hispanic white residents and blue points represent non-white residents. (G.8b): The density of speeding stops recorded by police officers, indicated by gray dots, where the area of the points is proportional to the number of stops at a single location; stops for speeding are overlaid with a heatmap of speeding events, as estimated from telematics data.




(a) Demographics.



(b) Speeding violations and stops.


Figure G.9: A comparison of San Antonio, TX's demographic composition, the location of speeding events, as measured by telematics data, and the location of speeding stops. (G.9a): The demographics of residents, where red points represent non-Hispanic white residents and blue points represent non-white residents. (G.9b): The density of speeding stops recorded by police officers, indicated by gray dots, where the area of the points is proportional to the number of stops at a single location; stops for speeding are overlayed with a heatmap of speeding events, as estimated from telematics data.



(a) Demographics.

(b) Speeding violations and stops.

Figure G.10: A comparison of Tulsa, OK's demographic composition, the location of speeding events, as measured by telematics data, and the location of speeding stops. (G.10a): The demographics of residents, where red points represent non-Hispanic white residents and blue points represent non-white residents. (G.10b): The density of speeding stops recorded by police officers, indicated by gray dots, where the area of the points is proportional to the number of stops at a single location; stops for speeding are overlaid with a heatmap of speeding events, as estimated from telematics data.



(a) Demographics.

(b) Speeding violations and stops.

Figure G.11: A comparison of Wichita, KS's demographic composition, the location of speeding events, as measured by telematics data, and the location of speeding stops. (G.11a): The demographics of residents, where red points represent non-Hispanic white residents and blue points represent non-white residents. (G.11b): The density of speeding stops recorded by police officers, indicated by gray dots, where the area of the points is proportional to the number of stops at a single location; stops for speeding are overlayed with a heatmap of speeding events, as estimated from telematics data.




Figure G.12: Average demographic composition of police beats in each city, weighted by the beat's total residential population (blue) and the number of CMT drivers resident in the beat (red). That is, the blue dots represent $[\sum_i p_i \cdot m_i] / [\sum_i p_i]$ and the red dots $[\sum_i c_i \cdot m_i] / [\sum_i c_i]$, where p_i represents the number of residents of beat i , m_i the proportion of residents of beat i that are non-white, and c_i the number of CMT drivers resident in beat i . The average difference is 5%, and the maximum difference is 10%, in Houston.




Figure G.13: The proportion of driving time spent at least 15 KPH above the speed limit according to TomTom data, versus the proportion of non-white residents, where each point represents a police beat in the 10 cities we analyze. The flat line of best fit (in red) indicates that neighborhood composition and speeding behavior are largely unrelated ($r = 0.03$). (Cf. Fig. 8.2.)




Figure G.14: Comparison of estimates for β_{RACE} in the main analysis when using CMT (blue) and TomTom (Red), as the data source. In general the estimates using TomTom suggest higher amounts of unjustified disparate impact, with the caveat that the TomTom data cover a shorter time period and cannot be post-stratified.




Figure G.15: The value of $\hat{\beta}_{\text{Race}}$ in our regression for each city using TomTom data, adjusting for the true prevalence of speeding (left) versus raw, unadjusted results (right). In both the adjusted and unadjusted regressions, we find significant heterogeneity across cities, suggesting that the nature and degree of policing practices varies considerably between cities. The mean of the unadjusted city-level coefficients is 0.58 (SE: 0.18, $p = 0.002$), meaning that speeding enforcement tends to be higher in areas with greater proportions of non-white residents. After adjusting for speeding using TomTom data, our average city-level coefficient is 1.16 (SE: 0.19, $p < 0.001$), meaning that disparities in speeding enforcement grow after adjusting for driving behavior. (Cf. Fig. 8.3, with the caveat that the TomTom data cover a shorter time period and cannot be post-stratified.)






Figure G.16: The city-level values of $\hat{\beta}_{Race}$ in our regression for all traffic violations using TomTom data. We find uniformly positive coefficients with average 2.08 (SE: 0.17, $p < 0.001$), meaning that in every city we analyze, there were more recorded traffic violations in beats with higher proportions of non-white residents when not adjusting for true prevalence of violations. (Cf. Fig. 8.6, with the caveat that the TomTom data cover a shorter time period and cannot be post-stratified.)



(a) Absolute speeding thresholds



(b) Relative speeding thresholds

Figure G.17: Robustness of our main analysis to the definition of speeding using TomTom data. We show, for each city, the value of β_{RACE} as we vary the threshold above the speed limit we consider speeding when we compute d_i , on an absolute (G.17a) and relative (G.17b) scale. We find that our estimates of the coefficients are largely stable as we shift the threshold. (Cf. Fig. G.1, with the caveat that the TomTom data cover a shorter time period and cannot be post-stratified.)




Figure G.18: Estimates of β_{RACE} in our speeding violation analysis using quasipoisson instead of negative binomial regression and TomTom data instead of CMT data. We find that the coefficients of Madison, Mesa, and Plano are all no longer statistically significant in this specification. (Cf. Fig. G.2, with the caveat that the TomTom data cover a shorter time period and cannot be post-stratified.)

References

Chapter 1: Introduction

- Alberto Alesina and Eliana La Ferrara. A test of racial bias in capital sentencing. *American Economic Review*, 104(11):3397–3433, 2014.
- Julia Angwin, Jeff Larson, Surya Mattu, and Lauren Kirchner. Machine bias. In *Ethics of data and analytics*, pages 254–264. Auerbach Publications, 2022.
- Shamena Anwar and Hanming Fang. An alternative test of racial prejudice in motor vehicle searches: Theory and evidence. *American Economic Review*, 96(1):127–151, 2006.
- Sarah F Anzia and Christopher R Berry. The Jackie (and Jill) Robinson effect: Why do congresswomen outperform congressmen? *American Journal of Political Science*, 55(3):478–493, 2011.
- Peter S. Arcidiacono. Report of Peter S. Arcidiacono. Expert report submitted in *Students for Fair Admissions, Inc. v. President and Fellows of Harvard College*, No. 14-cv-14176 (D. Mass.), 2018.
- David Arnold, Will Dobbie, and Crystal S Yang. Racial bias in bail decisions. *The Quarterly Journal of Economics*, 133(4):1885–1932, 2018.
- David Arnold, Will Dobbie, and Peter Hull. Measuring racial discrimination in bail decisions. *American Economic Review*, 112(9):2992–3038, 2022.
- Kenneth Arrow. The theory of discrimination. *Discrimination in labor markets*, 3(10):3–33, 1973.
- Ian Ayres. Three tests for measuring unjustified disparate impacts in organ transplantation: The problem of “included variable” bias. *Perspectives in Biology and Medicine*, 48(1):68–S87, 2005.
- Ian Ayres. Testing for discrimination and the problem of “included variable bias”. *Working paper*, 2010. Available at <http://islandia.law.yale.edu/ayers/ayresincludedvariablebias.pdf>.
- Solon Barocas, Moritz Hardt, and Arvind Narayanan. *Fairness and Machine Learning: Limitations and Opportunities*. MIT Press, 2023.

- Gary S Becker. *The economics of discrimination*. University of Chicago Press, 1957.
- Gary S Becker. Nobel lecture: The economic way of looking at behavior. *Journal of Political Economy*, 101(3):385–409, 1993.
- James A Berkovec, Glenn B Canner, Stuart A Gabriel, and Timothy H Hannan. Race, redlining, and residential mortgage loan performance. *The Journal of Real Estate Finance and Economics*, 9:263–294, 1994.
- James A Berkovec, Glenn B Canner, Stuart A Gabriel, and Timothy H Hannan. Discrimination, competition, and loan performance in FHA mortgage lending. *Review of Economics and Statistics*, 80(2):241–250, 1998.
- Marianne Bertrand and Sendhil Mullainathan. Are Emily and Greg more employable than Lakisha and Jamal? A field experiment on labor market discrimination. *American economic review*, 94(4):991–1013, 2004.
- Marianne Bertrand, Dolly Chugh, and Sendhil Mullainathan. Implicit discrimination. *American Economic Review*, 95(2):94–98, 2005.
- Bostock v. Clayton County. 590 U.S. 644, 2020.
- Joy Buolamwini and Timnit Gebru. Gender shades: Intersectional accuracy disparities in commercial gender classification. In *Conference on fairness, accountability and transparency*, pages 77–91. PMLR, 2018.
- David Card. Report of David Card. Expert report submitted in *Students for Fair Admissions, Inc. v. President and Fellows of Harvard College*, No. 14-cv-14176 (D. Mass.), 2018.
- Tianqi Chen and Carlos Guestrin. XGBoost: A scalable tree boosting system. In *Proceedings of the 22nd ACM SIGKDD International Conference on Knowledge Discovery and Data Mining*, KDD ’16, pages 785–794. Association for Computing Machinery, 2016.
- Adam Chilton, Jonathan Masur, and Kyle Rozema. Political ideology and the law review selection process. *American Law and Economics Review*, 22(1):211–240, 2020.
- Alexandra Chouldechova. Fair prediction with disparate impact: A study of bias in recidivism prediction instruments. *Big Data*, 5(2):153–163, 2017.
- Civil Rights Act of 1991. 42 U.S.C. § 2000e-2(m), 1991.
- Madison Coots, Kristin A Linn, Sharad Goel, Amol S Navathe, and Ravi B Parikh. Racial bias in clinical and population health algorithms: a critical review of current debates. *Annual Review of Public Health*, 46(1):507–523, 2025.
- Sam Corbett-Davies, Emma Pierson, Avi Feller, Sharad Goel, and Aziz Huq. Algorithmic decision making and the cost of fairness. In *Proceedings of the 23rd ACM SIGKDD International Conference on Knowledge Discovery and Data Mining*, pages 797–806, 2017.

- William Wentworth Daniel. *Racial Discrimination in England: Based on the P.E.P. Report*. Penguin Special S257. Penguin Books, Harmondsworth, 1968.
- Cynthia Dwork, Moritz Hardt, Toniann Pitassi, Omer Reingold, and Richard Zemel. Fairness through awareness. In *Proceedings of the 3rd Innovations in Theoretical Computer Science Conference*, pages 214–226, 2012.
- Jennifer L Eberhardt. *Biased: Uncovering the hidden prejudice that shapes what we see, think, and do*. Viking, 2019.
- Equal Credit Opportunity Act. 15 U.S.C. § 1691, 1974.
- Floyd v. City of New York. 959 F. Supp. 2d 540, S.D.N.Y, 2013.
- Roland G Fryer Jr. An empirical analysis of racial differences in police use of force. *Journal of Political Economy*, 127(3):1210–1261, 2019.
- S Michael Gaddis. How black are Lakisha and Jamal? racial perceptions from names used in correspondence audit studies. *Sociological Science*, 4:469, 2017.
- Andrew Gelman, Jeffrey Fagan, and Alex Kiss. An analysis of the New York City Police Department’s “stop-and-frisk” policy in the context of claims of racial bias. *Journal of the American Statistical Association*, 102(479):813–823, 2007.
- Sharad Goel, Justin M Rao, Ravi Shroff, et al. Precinct or prejudice? Understanding racial disparities in New York City’s stop-and-frisk policy. *The Annals of Applied Statistics*, 10(1): 365–394, 2016.
- Sharad Goel, Maya Perelman, Ravi Shroff, and David Alan Sklansky. Combatting police discrimination in the age of big data. *New Criminal Law Review: An International and Interdisciplinary Journal*, 20(2):181–232, 2017.
- Clifton Green, Narasimhan Jegadeesh, and Yue Tang. Gender and job performance: Evidence from Wall Street. *Financial Analysts Journal*, 65(6):65–78, 2009.
- D James Greiner and Donald B Rubin. Causal effects of perceived immutable characteristics. *Review of Economics and Statistics*, 93(3):775–785, 2011.
- Griggs v. Duke Power Co. 401 U.S. 424, 1971.
- Jeffrey Grogger and Greg Ridgeway. Testing for racial profiling in traffic stops from behind a veil of darkness. *Journal of the American Statistical Association*, 101(475):878–887, 2006.
- Joshua Grossman, Julian Nyarko, and Sharad Goel. Reconciling legal and empirical conceptions of disparate impact: An analysis of police stops across California. *Journal of Law & Empirical Analysis*, 1(1):118–133, 2024.
- Moritz Hardt, Eric Price, and Nati Srebro. Equality of opportunity in supervised learning. *Advances*

- in *Neural Information Processing Systems*, 29:3315–3323, 2016.
- James J Heckman. Detecting discrimination. *Journal of Economic Perspectives*, 12(2):101–116, 1998.
- Deborah Hellman. Measuring algorithmic fairness. *Virginia Law Review*, 106(4):811–866, 2020.
- Paul W Holland. Statistics and causal inference. *Journal of the American statistical Association*, 81(396):945–960, 1986.
- Lily Hu and Issa Kohler-Hausmann. What’s sex got to do with machine learning? In *Proceedings of the 2020 ACM Conference on Fairness, Accountability, and Transparency*, 2020.
- Aziz Z Huq. What is discriminatory intent? *Cornell Law Review*, 103:1211–1292, 2019.
- Kosuke Imai and Zhichao Jiang. Principal fairness for human and algorithmic decision-making. *arXiv preprint arXiv:2005.10400*, 2020.
- Jon Kleinberg, Sendhil Mullainathan, and Manish Raghavan. Inherent trade-offs in the fair determination of risk scores. In *Proceedings of Innovations in Theoretical Computer Science (ITCS)*, 2017.
- Patrick Kline, Evan K Rose, and Christopher R Walters. Systemic discrimination among large US employers. *The Quarterly Journal of Economics*, 137(4):1963–2036, 2022.
- John Knowles, Nicola Persico, and Petra Todd. Racial bias in motor vehicle searches: Theory and evidence. *Journal of Political Economy*, 109(1):203–229, 2001.
- Dean Knox, Will Lowe, and Jonathan Mummolo. Administrative records mask racially biased policing. *American Political Science Review*, 2020.
- Allison Koenecke, Andrew Nam, Emily Lake, Joe Nudell, Minnie Quartey, Zion Mengesha, Connor Toups, John R Rickford, Dan Jurafsky, and Sharad Goel. Racial disparities in automated speech recognition. *Proceedings of the National Academy of Sciences*, 117(14):7684–7689, 2020.
- Emmanuel Martinez and Lauren Kirchner. The secret bias hidden in mortgage-approval algorithms. The Markup, August 2021. URL <https://themarkup.org/denied/2021/08/25/the-secret-bias-hidden-in-mortgage-approval-algorithms>.
- Rajiv Movva, Divya Shanmugam, Kaihua Hou, Priya Pathak, John Guttag, Nikhil Garg, and Emma Pierson. Coarse race data conceals disparities in clinical risk score performance. In *Proceedings of the 8th Machine Learning for Healthcare Conference*, volume 219 of *Proceedings of Machine Learning Research*, pages 443–472. PMLR, 2023.
- Roland Neil and Christopher Winship. Methodological challenges and opportunities in testing for racial discrimination in policing. *Annual Review of Criminology*, 2:73–98, 2019.
- Ziad Obermeyer, Brian Powers, Christine Vogeli, and Sendhil Mullainathan. Dissecting racial bias

- in an algorithm used to manage the health of populations. *Science*, 366(6464):447–453, 2019.
- Brendan O’Flaherty and Rajiv Sethi. *Shadows of Doubt: Stereotypes, Crime, and the Pursuit of Justice*. Harvard University Press, 2019.
- Patient Protection and Affordable Care Act. 42 U.S.C. § 300gg, 2010.
- Personnel Adm’r of Mass. v. Feeney. 442 U.S. 256, 1979.
- Edmund S Phelps. The statistical theory of racism and sexism. *The American Economic Review*, 62(4):659–661, 1972.
- Emma Pierson, Camelia Simoiu, Jan Overgoor, Sam Corbett-Davies, Daniel Jenson, Amy Shoemaker, Vignesh Ramachandran, Phoebe Barghouty, Ravi Shroff, Cheryl Phillips, and Sharad Goel. A large-scale analysis of racial disparities in police stops across the United States. *Nature Human Behaviour*, 4(5), 2020.
- Devin G Pope and Justin R Sydnor. What’s in a picture? Evidence of discrimination from Prosper.com. *Journal of Human Resources*, 46(1):53–92, 2011.
- Price Waterhouse v. Hopkins. 490 U.S. 228, 1989.
- Inioluwa Deborah Raji and Joy Buolamwini. Actionable auditing: Investigating the impact of publicly naming biased performance results of commercial AI products. In *Proceedings of the 2019 AAAI/ACM Conference on AI, Ethics, and Society*, AIES ’19, pages 429–435. Association for Computing Machinery, 2019. doi: 10.1145/3306618.3314244.
- M Marit Rehavi and Sonja B Starr. Racial disparity in federal criminal sentences. *Journal of Political Economy*, 122(6):1320–1354, 2014.
- Richard Rothstein. *The Color of Law: A Forgotten History of How Our Government Segregated America*. Liveright Publishing, 2017.
- Maya Sen and Omar Wasow. Race as a bundle of sticks: Designs that estimate effects of seemingly immutable characteristics. *Annual Review of Political Science*, 19(1):499–522, 2016.
- SFFA v. Harvard. Students for Fair Admissions, Inc., Petitioner, v. President and Fellows of Harvard College. Students for Fair Admissions, Inc., Petitioner, v. University of North Carolina, et al., 2023. https://www.supremecourt.gov/opinions/22pdf/20-1199_16gn.pdf.
- Scott B Smart and Joel Waldfogel. A citation-based test for discrimination at economics and finance journals, 1996.
- Terry v. Ohio. 392 U.S. 1, 1968.
- Title VII of the Civil Rights Act of 1964. 42 U.S.C. § 2000e-2(e)(1), 1964.
- United States v. Brignoni-Ponce. 422 U.S. 873, 1975.

U.S. Census Bureau. American community survey, 2007 1-year estimates, 2007.

Vill. of Arlington Heights v. Metro. Hous. Dev. Corp. 429 U.S. 252, 1977.

Samuel Walker. Searching for the denominator: Problems with police traffic stop data and an early warning system solution. *Justice Research and Policy*, 3(1):63–95, 2001.

Washington v. Davis. 426 U.S. 229, 1976.

Ronald E Wienk. *Measuring racial discrimination in American housing markets: The housing market practices survey*, volume 444. Department of Housing and Urban Development, Office of Policy Development, 1979.

Chapter 2: The Measure and Mismeasure of Fairness

- Alekh Agarwal, Alina Beygelzimer, Miroslav Dudík, John Langford, and Hanna Wallach. A reductions approach to fair classification. In *International Conference on Machine Learning*, 2018.
- Rahul Aggarwal, Kirsten Bibbins-Domingo, Robert W. Yeh, Yang Song, Nicholas Chiu, Rishi K. Wadhwa, Changyu Shen, and Dhruv S. Kazi. Diabetes screening by race and ethnicity in the United States: Equivalent body mass index and age thresholds. *Annals of Internal Medicine*, 175(6):765–773, 2022.
- Robert M Anderson and William R Zame. Genericity with infinitely many parameters. *Advances in Theoretical Economics*, 1(1):1–62, 2001.
- Julia Angwin, Jeff Larson, Surya Mattu, and Lauren Kirchner. Machine bias: There’s software used across the country to predict future criminals. and it’s biased against blacks. *ProPublica*, 5 2016.
- Shamena Anwar and Hanming Fang. An alternative test of racial prejudice in motor vehicle searches: Theory and evidence. *The American Economic Review*, 2006.
- Idir Arab, Paulo Eduardo Oliveira, and Tilo Wiklund. Convex transform order of beta distributions with some consequences. *Statistica Neerlandica*, 75(3):238–256, 2021.
- Matthew Arnold, Rachel KE Bellamy, Michael Hind, Stephanie Houde, Sameep Mehta, Aleksandra Mojsilović, Ravi Nair, K Natesan Ramamurthy, Alexandra Olteanu, David Piorkowski, et al. Factsheets: Increasing trust in ai services through supplier’s declarations of conformity. *IBM Journal of Research and Development*, 63(4/5):6–1, 2019.
- Imanol Arrieta-Ibarra, Paman Gujral, Jonathan Tannen, Mark Tygert, and Cherie Xu. Metrics of calibration for probabilistic predictions. *Journal of Machine Learning Research*, 2022.
- Kenneth Arrow. The theory of discrimination. In *Discrimination in labor markets*. Princeton University Press, 1973.
- Ian Ayres. Outcome tests of racial disparities in police practices. *Justice Research and Policy*, 4(1-2):131–142, 2002.
- Jack M Balkin and Reva B Siegel. The American civil rights tradition: Anticlassification or antisubordination. *Issues in Legal Scholarship*, 2(1), 2003.
- Michelle Bao, Angela Zhou, Samantha Zottola, Brian Brubach, Sarah Desmarais, Aaron Horowitz, Kristian Lum, and Suresh Venkatasubramanian. It’s COMPASlicated: The messy relationship between RAI datasets and algorithmic fairness benchmarks. *arXiv preprint arXiv:2106.05498*, 2021.
- Chelsea Barabas, Madars Virza, Karthik Dinakar, Joichi Ito, and Jonathan Zittrain. Interventions over predictions: Reframing the ethical debate for actuarial risk assessment. In *Conference on Fairness, Accountability and Transparency*, pages 62–76, 2018.

- Solon Barocas and Andrew D Selbst. Big data’s disparate impact. *Cal. L. Rev.*, 104:671, 2016.
- Solon Barocas, Moritz Hardt, and Arvind Narayanan. *Fairness and Machine Learning: Limitations and Opportunities*. fairmlbook.org, 2019. <http://www.fairmlbook.org>.
- Gary S Becker. *The Economics of Discrimination*. University of Chicago Press, 1957.
- Benji. The sum of an uncountable number of positive numbers. Mathematics Stack Exchange, 2020. URL <https://math.stackexchange.com/q/20661>. (version: 2020-05-29).
- Richard Berk. *Criminal Justice Forecasts of Risk: A Machine Learning Approach*. Springer Science & Business Media, 2012.
- Richard Berk, Hoda Heidari, Shahin Jabbari, Michael Kearns, and Aaron Roth. Fairness in criminal justice risk assessments: The state of the art. *Sociological Methods & Research*, 50(1):3–44, 2021.
- Marianne Bertrand and Sendhil Mullainathan. Are Emily and Greg more employable than Lakisha and Jamal? A field experiment on labor market discrimination. *American Economic Review*, 94(4):991–1013, 2004.
- Patrick Billingsley. *Probability and Measure*. Wiley Series in Probability and Mathematical Statistics. John Wiley & Sons, Inc., New York, third edition, 1995. ISBN 0-471-00710-2. A Wiley-Interscience Publication.
- Avrim Blum and Kevin Stangl. Recovering from biased data: Can fairness constraints improve accuracy? *arXiv preprint arXiv:1912.01094*, 2019.
- Henk Brozius. Conditional expectation - $E[f(X, Y)|Y]$. Mathematics Stack Exchange, 2019. URL <https://math.stackexchange.com/q/3247577>. (Version: 2019-06-01).
- Joy Buolamwini and Timnit Gebru. Gender shades: Intersectional accuracy disparities in commercial gender classification. In *Conference on fairness, accountability and transparency*, pages 77–91. PMLR, 2018.
- William Cai, Johann Gaebler, Nikhil Garg, and Sharad Goel. Fair allocation through selective information acquisition. In *Proceedings of the AAAI/ACM Conference on AI, Ethics, and Society*, pages 22–28, 2020.
- William Cai, Ro Encarnacion, Bobbie Chern, Sam Corbett-Davies, Miranda Bogen, Stevie Bergman, and Sharad Goel. Adaptive sampling strategies to construct equitable training datasets. In *Proceedings of the Conference on Fairness, Accountability, and Transparency*, 2022a.
- William Cai, Johann Gaebler, Justin Kaashoek, Lisa Pinals, Samuel Madden, and Sharad Goel. Measuring racial and ethnic disparities in traffic enforcement with large-scale telematics data. *PNAS Nexus*, 1(4):pgac144, 2022b.
- Toon Calders and Sicco Verwer. Three naive Bayes approaches for discrimination-free classification. *Data Mining and Knowledge Discovery*, 21(2):277–292, 2010.

- Alycia N Carey and Xintao Wu. The causal fairness field guide: Perspectives from social and formal sciences. *Frontiers in Big Data*, 5, 2022.
- James H Carr and Isaac F Megbolugbe. *The Federal Reserve Bank of Boston study on mortgage lending revisited*. Fannie Mae Office of Housing Policy Research, 1993.
- Centers for Disease Control and Prevention. National Health and Nutrition Examination Survey data. Technical report, National Center for Health Statistics, 2011-2018. URL <https://wwwn.cdc.gov/nchs/nhanes/continuousnhanes/overview.aspx?BeginYear=2011>.
- Jessica P Cerdeña, Marie V Plaisime, and Jennifer Tsai. From race-based to race-conscious medicine: How anti-racist uprisings call us to act. *The Lancet*, 396(10257):1125–1128, 2020.
- Silvia Chiappa. Path-specific counterfactual fairness. In *Proceedings of the AAAI Conference on Artificial Intelligence*, volume 33, pages 7801–7808, 2019.
- Kasia S Chmielinski, Sarah Newman, Matt Taylor, Josh Joseph, Kemi Thomas, Jessica Yurkofsky, and Yue Chelsea Qiu. The dataset nutrition label (2nd gen): Leveraging context to mitigate harms in artificial intelligence. *arXiv preprint arXiv:2201.03954*, 2022.
- Alex Chohlas-Wood, Joe Nudell, Keniel Yao, Zhiyuan Lin, Julian Nyarko, and Sharad Goel. Blind justice: Algorithmically masking race in charging decisions. In *Proceedings of the 2021 AAAI/ACM Conference on AI, Ethics, and Society*, pages 35–45, 2021.
- Alex Chohlas-Wood, Madison Coots, Emma Brunskill, and Sharad Goel. Learning to be fair: A consequentialist approach to equitable decision-making. *arXiv preprint arXiv:2109.08792*, 2023a.
- Alex Chohlas-Wood, Madison Coots, Sharad Goel, and Julian Nyarko. Designing equitable algorithms. *Nature Computational Science*, 3, 2023b.
- Alexandra Chouldechova. Fair prediction with disparate impact: A study of bias in recidivism prediction instruments. *Big Data*, 5(2):153–163, 2017.
- Alexandra Chouldechova and Aaron Roth. A snapshot of the frontiers of fairness in machine learning. *Communications of the ACM*, 63(5):82–89, 2020.
- Alexandra Chouldechova, Diana Benavides-Prado, Oleksandr Fialko, and Rhema Vaithianathan. A case study of algorithm-assisted decision making in child maltreatment hotline screening decisions. In *Conference on Fairness, Accountability and Transparency*, pages 134–148, 2018.
- Jens Peter Reus Christensen. On sets of Haar measure zero in Abelian Polish groups. *Israel Journal of Mathematics*, 13(3-4):255–260, 1972.
- T Anne Cleary. Test bias: Prediction of grades of Negro and white students in integrated colleges. *Journal of Educational Measurement*, 5(2):115–124, 1968.
- Ruth Colker. Anti-subordination above all: Sex, race, and equal protection. *NYUL Rev.*, 61:1003, 1986.

- Madison Coots, Soroush Saghaian, David Kent, and Sharad Goel. Reevaluating the role of race and ethnicity in diabetes screening. *arXiv preprint arXiv:2306.10220*, 2023.
- Sam Corbett-Davies and Sharad Goel. The measure and mismeasure of fairness: A critical review of fair machine learning. *arXiv preprint arXiv:1808.00023v2*, 2018.
- Sam Corbett-Davies, Emma Pierson, Avi Feller, Sharad Goel, and Aziz Huq. Algorithmic decision making and the cost of fairness. In *Proceedings of the 23rd ACM SIGKDD International Conference on Knowledge Discovery and Data Mining*, pages 797–806, 2017.
- Amanda Coston, Alan Mishler, Edward H Kennedy, and Alexandra Chouldechova. Counterfactual risk assessments, evaluation, and fairness. In *Proceedings of the 2020 Conference on Fairness, Accountability, and Transparency*, pages 582–593, 2020.
- Andrew Cotter, Heinrich Jiang, and Karthik Sridharan. Two-player games for efficient non-convex constrained optimization. In *Algorithmic Learning Theory*, pages 300–332. PMLR, 2019.
- Stewart J D’Alessio and Lisa Stolzenberg. Race and the probability of arrest. *Social forces*, 81(4): 1381–1397, 2003.
- Richard B Darlington. Another look at “cultural fairness”. *Journal of Educational Measurement*, 8 (2):71–82, 1971.
- Matthew DeMichele, Peter Baumgartner, Michael Wenger, Kelle Barrick, Megan Comfort, and Shilpi Misra. The Public Safety Assessment: A re-validation and assessment of predictive utility and differential prediction by race and gender in Kentucky, 2018. URL <https://papers.ssrn.com/abstract=3168452>.
- Cynthia Dwork, Moritz Hardt, Toniann Pitassi, Omer Reingold, and Richard Zemel. Fairness through awareness. In *Proceedings of the 3rd Innovations in Theoretical Computer Science Conference*, pages 214–226, 2012.
- Harrison Edwards and Amos Storkey. Censoring representations with an adversary. In *Proceedings of the International Conference in Learning Representations*, 2016.
- Robin S Engel and Rob Tillyer. Searching for equilibrium: The tenuous nature of the outcome test. *Justice Quarterly*, 25(1):54–71, 2008.
- Michael Feldman, Sorelle A Friedler, John Moeller, Carlos Scheidegger, and Suresh Venkatasubramanian. Certifying and removing disparate impact. In *Proceedings of the 21th ACM SIGKDD International Conference on Knowledge Discovery and Data Mining*, pages 259–268. ACM, 2015.
- Fisher v. University of Texas. 579 U.S., 2016.
- Owen M Fiss. Groups and the Equal Protection Clause. *Philosophy & Public Affairs*, pages 107–177, 1976.
- Johann Gaebler, William Cai, Guillaume Basse, Ravi Shroff, Sharad Goel, and Jennifer Hill. A

- causal framework for observational studies of discrimination. *Statistics and Public Policy*, 2022.
- Sainyam Galhotra, Karthikeyan Shanmugam, Prasanna Sattigeri, and Kush R Varshney. Causal feature selection for algorithmic fairness. *Proceedings of the 2022 International Conference on Management of Data (SIGMOD)*, 2022.
- George C Galster. The facts of lending discrimination cannot be argued away by examining default rates. *Housing Policy Debate*, 4(1):141–146, 1993.
- Timnit Gebru, Jamie Morgenstern, Briana Vecchione, Jennifer Wortman Vaughan, Hanna Wallach, Hal Daumé Iii, and Kate Crawford. Datasheets for datasets. *Communications of the ACM*, 64(12):86–92, 2021.
- Sharad Goel, Maya Perelman, Ravi Shroff, and David Alan Sklansky. Combatting police discrimination in the age of big data. *New Criminal Law Review: An International and Interdisciplinary Journal*, 20(2):181–232, 2017.
- Claudia Goldin and Cecilia Rouse. Orchestrating impartiality: The impact of “blind” auditions on female musicians. *American Economic Review*, 90(4):715–741, 2000.
- D James Greiner and Donald B Rubin. Causal effects of perceived immutable characteristics. *Review of Economics and Statistics*, 93(3):775–785, 2011.
- Jeffrey Grogger and Greg Ridgeway. Testing for racial profiling in traffic stops from behind a veil of darkness. *Journal of the American Statistical Association*, 101(475):878–887, 2006.
- Moritz Hardt, Eric Price, and Nati Srebro. Equality of opportunity in supervised learning. *Advances in Neural Information Processing Systems*, 29:3315–3323, 2016.
- Ursula Hébert-Johnson, Michael Kim, Omer Reingold, and Guy Rothblum. Multicalibration: Calibration for the (computationally-identifiable) masses. In *International Conference on Machine Learning*, pages 1939–1948. PMLR, 2018.
- Jennifer L Hill. Bayesian nonparametric modeling for causal inference. *Journal of Computational and Graphical Statistics*, 20(1):217–240, 2011.
- Paul W Holland. Statistics and causal inference. *Journal of the American Statistical Association*, 81(396):945–960, 1986.
- Sarah Holland, Ahmed Hosny, Sarah Newman, Joshua Joseph, and Kasia Chmielinski. The dataset nutrition label. *Data Protection and Privacy*, 12(12):1, 2020.
- Lily Hu and Issa Kohler-Hausmann. What’s sex got to do with machine learning? In *Proceedings of the 2020 ACM Conference on Fairness, Accountability, and Transparency*, 2020.
- Brian R Hunt, Tim Sauer, and James A Yorke. Prevalence: a translation-invariant “almost every” on infinite-dimensional spaces. *Bulletin of the American Mathematical Society*, 27(2):217–238, 1992.

- Idaho H.B. 118. H.B. 118, 65th Leg., 1st Reg. Sess., 2019. <https://legislature.idaho.gov/wp-content/uploads/sessioninfo/2019/legislation/H0118.pdf>.
- Kosuke Imai and Zhichao Jiang. Principal fairness for human and algorithmic decision-making. *arXiv preprint arXiv:2005.10400*, 2020.
- Kosuke Imai, Zhichao Jiang, James Greiner, Ryan Halen, and Sooahn Shin. Experimental evaluation of algorithm-assisted human decision-making: Application to pretrial public safety assessment. *arXiv preprint arXiv:2012.02845*, 2020.
- Guido W Imbens and Donald B Rubin. *Causal Inference in Statistics, Social, and Biomedical Sciences*. Cambridge University Press, 2015.
- Christopher Jung, Sampath Kannan, Changhwa Lee, Mallesh Pai, Aaron Roth, and Rakesh Vohra. Fair prediction with endogenous behavior. In *Proceedings of the 21st ACM Conference on Economics and Computation*, pages 677–678, 2020a.
- Jongbin Jung, Connor Concannon, Ravi Shroff, Sharad Goel, and Daniel G Goldstein. Simple rules to guide expert classifications. *Journal of the Royal Statistical Society: Series A (Statistics in Society)*, 183(3):771–800, 2020b.
- Jongbin Jung, Ravi Shroff, Avi Feller, and Sharad Goel. Bayesian sensitivity analysis for offline policy evaluation. In *Proceedings of the AAAI/ACM Conference on AI, Ethics, and Society*, pages 64–70, 2020c.
- Faisal Kamiran, Indrė Žliobaitė, and Toon Calders. Quantifying explainable discrimination and removing illegal discrimination in automated decision making. *Knowledge and Information Systems*, 35(3):613–644, 2013.
- John G. Kemeny and J. Laurie Snell. *Finite Markov Chains*. Undergraduate Texts in Mathematics. Springer-Verlag, New York-Heidelberg, 1976. Reprinting of the 1960 original.
- Niki Kilbertus, Mateo Rojas-Carulla, Giambattista Parascandolo, Moritz Hardt, Dominik Janzing, and Bernhard Schölkopf. Avoiding discrimination through causal reasoning. In *Proceedings of the 31st International Conference on Neural Information Processing Systems*, pages 656–666, 2017.
- Jon Kleinberg, Himabindu Lakkaraju, Jure Leskovec, Jens Ludwig, and Sendhil Mullainathan. Human decisions and machine predictions. *The Quarterly Journal of Economics*, 133(1):237–293, 2017a.
- Jon Kleinberg, Sendhil Mullainathan, and Manish Raghavan. Inherent trade-offs in the fair determination of risk scores. In *Proceedings of Innovations in Theoretical Computer Science (ITCS)*, 2017b.
- Achim Klenke. *Probability Theory: A Comprehensive Course*. Universitext. Springer, Cham, 2020. Third edition.
- John Knowles, Nicola Persico, and Petra Todd. Racial bias in motor vehicle searches: Theory and

- evidence. *Journal of Political Economy*, 109(1), 2001.
- Allison Koenecke, Andrew Nam, Emily Lake, Joe Nudell, Minnie Quartey, Zion Mengesha, Connor Toups, John R Rickford, Dan Jurafsky, and Sharad Goel. Racial disparities in automated speech recognition. *Proceedings of the National Academy of Sciences*, 117(14):7684–7689, 2020.
- Allison Koenecke, Eric Giannella, Robb Willer, and Sharad Goel. Popular support for balancing equity and efficiency in resource allocation: A case study in online advertising to increase welfare program awareness. In *Proceedings of the International AAAI Conference on Web and Social Media*, volume 17, pages 494–506, 2023.
- Matt J Kusner, Joshua Loftus, Chris Russell, and Ricardo Silva. Counterfactual fairness. In *Advances in Neural Information Processing Systems*, pages 4066–4076, 2017.
- John M. Lee. *Introduction to Smooth Manifolds*, volume 218 of *Graduate Texts in Mathematics*. Springer, New York, second edition, 2013.
- Annie Liang, Jay Lu, and Xiaosheng Mu. Algorithmic design: Fairness versus accuracy. In *Proceedings of the 23rd ACM Conference on Economics and Computation*, pages 58–59, 2022.
- Joshua R Loftus, Chris Russell, Matt J Kusner, and Ricardo Silva. Causal reasoning for algorithmic fairness. *arXiv preprint arXiv:1805.05859*, 2018.
- Kristian Lum and William Isaac. To predict and serve? *Significance*, 13(5):14–19, 2016.
- Charles F Manski, John Mullahy, and Atheendar Venkataramani. Using measures of race to make clinical predictions: Decision making, patient health, and fairness. Technical report, National Bureau of Economic Research, 2022.
- Vishwali Mhasawade and Rumi Chunara. Causal multi-level fairness. In *Proceedings of the 2021 AAAI/ACM Conference on AI, Ethics, and Society*, pages 784–794, 2021.
- Alan Mishler, Edward H Kennedy, and Alexandra Chouldechova. Fairness in risk assessment instruments: Post-processing to achieve counterfactual equalized odds. In *Proceedings of the 2021 ACM Conference on Fairness, Accountability, and Transparency*, pages 386–400, 2021.
- Pat Muldowney, Krzysztof Ostaszewski, and Wojciech Wojdowski. The Darth Vader rule. *Tatra Mountains Mathematical Publications*, 52(1):53–63, 2012.
- Sendhil Mullainathan and Jann Spiess. Machine learning: An applied econometric approach. *Journal of Economic Perspectives*, 31(2):87–106, 2017.
- Razieh Nabi and Ilya Shpitser. Fair inference on outcomes. In *Proceedings of the AAAI Conference on Artificial Intelligence*, 2018.
- Hamed Nilforoshan, Johann D Gaebler, Ravi Shroff, and Sharad Goel. Causal conceptions of fairness and their consequences. In *International Conference on Machine Learning*, pages 16848–16887. PMLR, 2022.

- Abigail Nurse. Anti-subordination in the Equal Protection Clause: A case study. *NYUL Rev.*, 89: 293, 2014.
- Julian Nyarko, Sharad Goel, and Roseanna Sommers. Breaking taboos in fair machine learning: An experimental study. In *Equity and Access in Algorithms, Mechanisms, and Optimization*. Association for Computing Machinery, 2021.
- Ziad Obermeyer, Brian Powers, Christine Vogeli, and Sendhil Mullainathan. Dissecting racial bias in an algorithm used to manage the health of populations. *Science*, 366(6464):447–453, 2019.
- William Ott and James Yorke. Prevalence. *Bulletin of the American Mathematical Society*, 42(3): 263–290, 2005.
- Scott E Page. Making the difference: Applying a logic of diversity. *Academy of Management Perspectives*, 21(4):6–20, 2007.
- Jessica K Paulus and David M Kent. Predictably unequal: Understanding and addressing concerns that algorithmic clinical prediction may increase health disparities. *NPJ Digital Medicine*, 3(1): 1–8, 2020.
- Judea Pearl. Direct and indirect effects. In *Proceedings of the Seventeenth Conference on Uncertainty and Artificial Intelligence, 2001*, pages 411–420. Morgan Kaufman, 2001.
- Judea Pearl. Causal inference in statistics: An overview. *Statistics surveys*, 3:96–146, 2009a.
- Judea Pearl. *Causality*. Cambridge University Press, second edition, 2009b.
- Dino Pedreshi, Salvatore Ruggieri, and Franco Turini. Discrimination-aware data mining. In *Proceedings of the 14th ACM SIGKDD International Conference on Knowledge Discovery and Data Mining*, pages 560–568, 2008.
- Edmund S Phelps. The statistical theory of racism and sexism. *The American Economic Review*, 62(4):659–661, 1972.
- Emma Pierson, Sam Corbett-Davies, and Sharad Goel. Fast threshold tests for detecting discrimination. In *The 21st International Conference on Artificial Intelligence and Statistics (AISTATS)*, 2018.
- Emma Pierson, Camelia Simoiu, Jan Overgoor, Sam Corbett-Davies, Daniel Jenson, Amy Shoemaker, Vignesh Ramachandran, Phoebe Barghouty, Cheryl Phillips, Ravi Shroff, and Sharad Goel. A large-scale analysis of racial disparities in police stops across the United States. *Nature Human Behaviour*, 4(7):736–745, 2020.
- Richard A Primus. Equal protection and disparate impact: Round three. *Harv. L. Rev.*, 117:494, 2003.
- Rajeev Ramchand, Rosalie Liccardo Pacula, and Martin Y Iguchi. Racial differences in marijuana-users’ risk of arrest in the United States. *Drug and alcohol dependence*, 84(3):264–272, 2006.

- M. M. Rao. *Conditional measures and applications*, volume 271 of *Pure and Applied Mathematics (Boca Raton)*. Chapman & Hall/CRC, Boca Raton, FL, second edition, 2005.
- Guy N Rothblum and Gal Yona. Decision-making under miscalibration. *arXiv preprint arXiv:2203.09852*, 2022.
- Walter Rudin. *Principles of Mathematical Analysis*, volume 3. McGraw-Hill New York, 1976.
- Walter Rudin. *Real and Complex Analysis*. McGraw-Hill Book Co., New York, third edition, 1987. ISBN 0-07-054234-1.
- Walter Rudin. *Functional Analysis*. International Series in Pure and Applied Mathematics. McGraw-Hill, Inc., New York, second edition, 1991. ISBN 0-07-054236-8.
- Laleh Seyyed-Kalantari, Haoran Zhang, Matthew McDermott, Irene Y Chen, and Marzyeh Ghassemi. Underdiagnosis bias of artificial intelligence algorithms applied to chest radiographs in under-served patient populations. *Nature Medicine*, 27(12):2176–2182, 2021.
- SFFA v. Harvard. Students for Fair Admissions, Inc., Petitioner, v. President and Fellows of Harvard College. Students for Fair Admissions, Inc., Petitioner, v. University of North Carolina, et al., 2023. https://www.supremecourt.gov/opinions/22pdf/20-1199_16gn.pdf.
- Moshe Shaked and J George Shanthikumar. *Stochastic Orders*. Springer, 2007.
- Ravi Shroff. Predictive analytics for city agencies: Lessons from children’s services. *Big Data*, 5(3):189–196, 2017.
- Reva B Siegel. Equality talk: Antisubordination and anticlassification values in constitutional struggles over Brown. *Harv. L. Rev.*, 117:1470, 2003.
- C. E. Silva. *Invitation to Ergodic Theory*, volume 42 of *Student Mathematical Library*. American Mathematical Society, Providence, RI, 2008.
- Camelia Simoiu, Sam Corbett-Davies, and Sharad Goel. The problem of infra-marginality in outcome tests for discrimination. *The Annals of Applied Statistics*, 11(3):1193–1216, 2017.
- Jennifer Skeem, John Monahan, and Christopher Lowenkamp. Gender, risk assessment, and sanctioning: The cost of treating women like men. *Law and human behavior*, 40(5):580, 2016.
- Jennifer L. Skeem and Christopher T. Lowenkamp. Risk, race, and recidivism: Predictive bias and disparate impact. *Criminology*, 54(4):680–712, 2016.
- Rhys Steele. Space of vector measures equipped with the total variation norm is complete. Mathematics Stack Exchange, 2019. URL <https://math.stackexchange.com/q/3197508>. (Version: 2019-04-22).
- Julia Stoyanovich and Bill Howe. Nutritional labels for data and models. *A Quarterly bulletin of the Computer Society of the IEEE Technical Committee on Data Engineering*, 42(3), 2019.

- Yixin Wang, Dhanya Sridhar, and David M Blei. Equal opportunity and affirmative action via counterfactual predictions. *arXiv preprint arXiv:1905.10870*, 2019.
- Watson v. Fort Worth. *Watson v. Forth Worth Bank & Trust*, 1988. 487 U.S. 977.
- Hilde Weerts, Miroslav Dudík, Richard Edgar, Adrin Jalali, Roman Lutz, and Michael Madaio. Fairlearn: Assessing and improving fairness of ai systems, 2023.
- Teresa Scotton Williams. Some issues in the standardized testing of minority students. *Journal of Education*, pages 192–208, 1983.
- Blake Woodworth, Suriya Gunasekar, Mesrob I Ohannessian, and Nathan Srebro. Learning non-discriminatory predictors. In *Conference on Learning Theory*, pages 1920–1953. PMLR, 2017.
- Yongkai Wu, Lu Zhang, Xintao Wu, and Hanghang Tong. PC-fairness: A unified framework for measuring causality-based fairness. *Advances in Neural Information Processing Systems*, 32, 2019.
- Muhammad Bilal Zafar, Isabel Valera, Manuel Gomez Rodriguez, and Krishna P Gummadi. Fairness beyond disparate treatment & disparate impact: Learning classification without disparate mistreatment. In *Proceedings of the 26th International Conference on World Wide Web*, pages 1171–1180, 2017a.
- Muhammad Bilal Zafar, Isabel Valera, Manuel Gomez Rodriguez, Krishna P Gummadi, and Adrian Weller. From parity to preference-based notions of fairness in classification. In *Proceedings of the 31st International Conference on Neural Information Processing Systems*, pages 228–238, 2017b.
- Muhammad Bilal Zafar, Isabel Valera, Manuel Gomez Rogriguez, and Krishna P Gummadi. Fairness constraints: Mechanisms for fair classification. In *Artificial Intelligence and Statistics*, pages 962–970, 2017c.
- Michael Zanger-Tishler, Julian Nyarko, and Sharad Goel. Risk scores, label bias, and everything but the kitchen sink. *arXiv preprint arXiv:2305.12638*, 2023.
- Rich Zemel, Yu Wu, Kevin Swersky, Toni Pitassi, and Cynthia Dwork. Learning fair representations. In *International Conference on Machine Learning*, pages 325–333, 2013.
- Junzhe Zhang and Elias Bareinboim. Fairness in decision-making—the causal explanation formula. In *Thirty-Second AAAI Conference on Artificial Intelligence*, 2018.
- Lu Zhang, Yongkai Wu, and Xintao Wu. A causal framework for discovering and removing direct and indirect discrimination. In *Proceedings of the 26th International Joint Conference on Artificial Intelligence*, pages 3929–3935, 2017.

Chapter 3: Fair Allocation through Selective Information Acquisition

- Apaam et al. 2017 FDIC National survey of unbanked and underbanked households. *US Federal Deposit Insurance Corporation*, 2018.
- Michiel A Bakker, Alejandro Noriega-Campero, Duy Patrick Tu, Prasanna Sattigeri, and Kush R Varshney. On fairness in budget-constrained decision making. In *KDD '19: Explainable AI (XAI)*, 2019.
- Dimitris Bertsimas and John N Tsitsiklis. *Introduction to linear optimization*, volume 6. Athena Scientific Belmont, MA, 1997.
- Tolga Bolukbasi, Kai-Wei Chang, James Y Zou, Venkatesh Saligrama, and Adam T Kalai. Man is to computer programmer as woman is to homemaker? Debiasing word embeddings. In *Advances in Neural Information Processing Systems*, pages 4349–4357, 2016.
- Steven J Brams and Alan D Taylor. *Fair Division: From cake-cutting to dispute resolution*. Cambridge University Press, 1996.
- Joy Buolamwini and Timnit Gebru. Gender shades: Intersectional accuracy disparities in commercial gender classification. In *Conference on fairness, accountability and transparency*, pages 77–91, 2018.
- Aylin Caliskan, Joanna J Bryson, and Arvind Narayanan. Semantics derived automatically from language corpora contain human-like biases. *Science*, 356(6334):183–186, 2017.
- Alexandra Chouldechova. Fair prediction with disparate impact: A study of bias in recidivism prediction instruments. *Big Data*, 5(2):153–163, 2017.
- Sam Corbett-Davies and Sharad Goel. The measure and mismeasure of fairness: A critical review of fair machine learning. *arXiv preprint arXiv:1808.00023*, 2018.
- Sam Corbett-Davies, Emma Pierson, Avi Feller, Sharad Goel, and Aziz Huq. Algorithmic decision making and the cost of fairness. In *Proceedings of the 23rd ACM SIGKDD International Conference on Knowledge Discovery and Data Mining*, pages 797–806. ACM, 2017.
- George B Dantzig. Discrete-variable extremum problems. *Operations research*, 5(2):266–288, 1957.
- Jeffrey Dastin. Amazon scraps secret ai recruiting tool that showed bias against women, Oct 2018. URL <https://www.reuters.com/article/us-amazon-com-jobs-automation-insight/amazon-scraps-secret-ai-recruiting-tool-that-showed-bias-against-women-idUSKCN1MK08G>.
- Hadi Elzayn, Shahin Jabbari, Christopher Jung, Michael Kearns, Seth Neel, Aaron Roth, and Zachary Schutzman. Fair algorithms for learning in allocation problems. In *Proceedings of the*

- Conference on Fairness, Accountability, and Transparency*, pages 170–179. ACM, 2019.
- Nikhil Garg, Londa Schiebinger, Dan Jurafsky, and James Zou. Word embeddings quantify 100 years of gender and ethnic stereotypes. *Proceedings of the National Academy of Sciences*, 115(16): E3635–E3644, 2018.
- Timnit Gebru, Jamie Morgenstern, Briana Vecchione, Jennifer Wortman Vaughan, Hanna Wallach, Hal Daumeé III, and Kate Crawford. Datasheets for datasets. *arXiv preprint arXiv:1803.09010*, 2018.
- Moritz Hardt, Eric Price, and Nati Srebro. Equality of opportunity in supervised learning. In *Advances In Neural Information Processing Systems*, pages 3315–3323, 2016.
- Hans Hofmann. Statlog (German Credit Data) data set. *UCI Repository of Machine Learning Databases*, 1994.
- Jon Kleinberg, Sendhil Mullainathan, and Manish Raghavan. Inherent trade-offs in the fair determination of risk scores. In *Proceedings of Innovations in Theoretical Computer Science (ITCS)*, 2017.
- Lydia T Liu, Sarah Dean, Esther Rolf, Max Simchowitz, and Moritz Hardt. Delayed impact of fair machine learning. *arXiv preprint arXiv:1803.04383*, 2018.
- Prem Melville, Maytal Saar-Tsechansky, Foster Provost, and Raymond Mooney. Active feature-value acquisition for classifier induction. In *Fourth IEEE International Conference on Data Mining (ICDM’04)*, pages 483–486. IEEE, 2004.
- Prem Melville, Maytal Saar-Tsechansky, Foster Provost, and Raymond Mooney. An expected utility approach to active feature-value acquisition. In *Fifth IEEE International Conference on Data Mining (ICDM’05)*, pages 4–pp. IEEE, 2005.
- Hervé Moulin. *Fair division and collective welfare*. MIT press, 2004.
- Alejandro Noriega-Campero, Michiel A Bakker, Bernardo Garcia-Bulle, and Alex’Sandy’ Pentland. Active fairness in algorithmic decision making. In *Proceedings of the 2019 AAAI/ACM Conference on AI, Ethics, and Society*, pages 77–83. ACM, 2019.
- Brendan O’Donoghue, Eric Chu, Neal Parikh, and Stephen Boyd. Conic optimization via operator splitting and homogeneous self-dual embedding. *Journal of Optimization Theory and Applications*, 169(3):1042–1068, 2016.
- Alexandra Olteanu, Carlos Castillo, Fernando Diaz, and Emre Kiciman. Social data: Biases, methodological pitfalls, and ethical boundaries. *Frontiers in Big Data*, 2:13, 2019.
- Maytal Saar-Tsechansky, Prem Melville, and Foster Provost. Active feature-value acquisition. *Management Science*, 55(4):664–684, 2009.
- William Thomson. Fair allocation rules. In *Handbook of social choice and welfare*, volume 2, pages

393–506. Elsevier, 2011.

Chapter 4: Measuring Disparate Impact in Human and Machine Decisions

- A framework for considering the value of race and ethnicity in estimating disease risk. *Annals of Internal Medicine*, 178(1):98–107, 2025. doi: 10.7326/M23-3166. PMID: 39622056.
- David S. Abrams. The law and economics of stop-and-frisk. *Loyola University Chicago Law Journal*, 46:369–381, 2014.
- David Arnold, Will Dobbie, and Peter Hull. Measuring racial discrimination in algorithms. In *AEA Papers and Proceedings*, volume 111, pages 49–54, 2021.
- David Arnold, Will Dobbie, and Peter Hull. Measuring racial discrimination in bail decisions. *American Economic Review*, 112(9):2992–3038, September 2022. doi: 10.1257/aer.20201653. URL <https://www.aeaweb.org/articles?id=10.1257/aer.20201653>.
- Kenneth Arrow. The theory of discrimination. *Discrimination in labor markets*, 3(10):3–33, 1973.
- Ian Ayres. Outcome tests of racial disparities in police practices. *Justice Research and Policy*, 4(1-2):131–142, 2002.
- Ian Ayres. Three tests for measuring unjustified disparate impacts in organ transplantation: The problem of “included variable” bias. *Perspectives in Biology and Medicine*, 48(1):68–S87, 2005.
- Ian Ayres. Testing for discrimination and the problem of “included variable bias”. *Working paper*, 2010. Available at <http://islandia.law.yale.edu/ayers/ayresincludedvariablebias.pdf>.
- Yuntao Bai, Andy Jones, Kamal Ndousse, Amanda Askell, Anna Chen, Nova DasSarma, Dawn Drain, Stanislav Fort, Deep Ganguli, Tom Henighan, Nicholas Joseph, Saurav Kadavath, Jackson Kernion, Tom Conerly, Sheer El-Showk, Nelson Elhage, Zac Hatfield-Dodds, Danny Hernandez, Tristan Hume, Scott Johnston, Shauna Kravec, Liane Lovitt, Neel Nanda, Catherine Olsson, Dario Amodei, Tom Brown, Jack Clark, Sam McCandlish, Chris Olah, Ben Mann, and Jared Kaplan. Training a helpful and harmless assistant with reinforcement learning from human feedback, 2022. URL <https://arxiv.org/abs/2204.05862>.
- Ana I Balsa, Thomas G McGuire, and Lisa S Meredith. Testing for statistical discrimination in health care. *Health Services Research*, 40(1):227–252, 2005.
- Gary S Becker. *The economics of discrimination*. University of Chicago Press, 1957.
- Gary S Becker. Nobel lecture: The economic way of looking at behavior. *Journal of Political Economy*, 101(3):385–409, 1993.
- James A Berkovec, Glenn B Canner, Stuart A Gabriel, and Timothy H Hannan. Mortgage discrimination and FHA loan performance. In *Mortgage Lending, Racial Discrimination, and Federal Policy*, pages 289–305. Routledge, 2018.

- J Aislinn Bohren, Peter Hull, and Alex Imas. Systemic discrimination: Theory and measurement. *The Quarterly Journal of Economics*, pages 1743–1799, 05 2025. ISSN 0033-5533. doi: 10.1093/qje/qjaf022. URL <https://doi.org/10.1093/qje/qjaf022>.
- Stephen P Boyd and Lieven Vandenbergh. *Convex optimization*. Cambridge University Press, 2004.
- Shubham Chaudhuri and Rajiv Sethi. Statistical Discrimination with Peer Effects: Can Integration Eliminate Negative Stereotypes? *The Review of Economic Studies*, 75(2):579–596, 04 2008. ISSN 0034-6527. doi: 10.1111/j.1467-937X.2008.00468.x. URL <https://doi.org/10.1111/j.1467-937X.2008.00468.x>.
- Alexandra Chouldechova, Diana Benavides-Prado, Oleksandr Fialko, and Rhema Vaithianathan. A case study of algorithm-assisted decision making in child maltreatment hotline screening decisions. In *Conference on Fairness, Accountability and Transparency*, pages 134–148, 2018.
- Carlos Cinelli and Chad Hazlett. Making sense of sensitivity: Extending omitted variable bias. *Journal of the Royal Statistical Society Series B-Statistical Methodology*, 82(1):39–67, 2020.
- Carlos Cinelli, Andrew Forney, and Judea Pearl. A crash course in good and bad controls. *Sociological Methods & Research*, 53(3):1071–1104, 2024.
- Connecticut v. Teal. 457 U.S. 440, 1982.
- Madison Coots, Kristin A Linn, Sharad Goel, Amol S Navathe, and Ravi B Parikh. Racial bias in clinical and population health algorithms: a critical review of current debates. *Annual Review of Public Health*, 46, 2025.
- Sam Corbett-Davies, Johann D. Gaebler, Hamed Nilforoshan, Ravi Shroff, and Sharad Goel. The measure and mismeasure of fairness. *Journal of Machine Learning Research*, 24(312):1–117, 2023. URL <http://jmlr.org/papers/v24/22-1511.html>.
- Amanda Coston, Alan Mishler, Edward H Kennedy, and Alexandra Chouldechova. Counterfactual risk assessments, evaluation, and fairness. In *Proceedings of the 2020 Conference on Fairness, Accountability, and Transparency*, pages 582–593, 2020.
- Will Dobbie, Jacob Goldin, and Crystal S Yang. The effects of pre-trial detention on conviction, future crime, and employment: Evidence from randomly assigned judges. *American Economic Review*, 108(2):201–240, 2018.
- Jacob Dorn and Kevin Guo. Sharp sensitivity analysis for inverse propensity weighting via quantile balancing. *Journal of the American Statistical Association*, pages 1–13, 2022.
- Hadi Elzayn, Evelyn Smith, Thomas Hertz, Arun Ramesh, Jacob Goldin, Daniel E Ho, and Robin Fisher. *Measuring and mitigating racial disparities in tax audits*. Stanford Institute for Economic Policy Research (SIEPR), 2023.
- Thomas J Espenshade, Chang Y Chung, and Joan L Walling. Admission preferences for minority

- students, athletes, and legacies at elite universities. *Social Science Quarterly*, 85(5):1422–1446, 2004.
- Executive Order. Restoring equality of opportunity and meritocracy, April 2025. URL <https://www.whitehouse.gov/presidential-actions/2025/04/restoring-equality-of-opportunity-and-meritocracy/>. The White House.
- Jeffrey Fagan. Report of Jeffrey Fagan in the case of Floyd v. the City of New York, 2010. Available at <https://www.law.columbia.edu/sites/default/files/microsites/policing-litigation-conference/files/Fagan%20Report%20with%20Technical%20Appendices.pdf>.
- FDIC Consumer Compliance Examination Manual: V. Lending—Equal Credit Opportunity Act (ECOA)*. Federal Deposit Insurance Corporation, Washington, DC, December 2024. Retrieved May 15, 2025, from <https://www.fdic.gov/resources/supervision-and-examinations/consumer-compliance-examination-manual/documents/5/v-7-1.pdf>.
- Floyd v. City of New York. 959 F. Supp. 2d 540, S.D.N.Y, 2013.
- Johann Gaebler, William Cai, Guillaume Basse, Ravi Shroff, Sharad Goel, and Jennifer Hill. A causal framework for observational studies of discrimination. *Statistics and Public Policy*, 9:26–48, 2022.
- Johann D Gaebler and Sharad Goel. A simple, statistically robust test of discrimination. *Proceedings of the National Academy of Sciences*, 122(10):e2416348122, 2025.
- Johann D. Gaebler, Sharad Goel, Aziz Huq, and Prasanna Tambe. Auditing large language models for race & gender disparities: Implications for artificial intelligence-based hiring. *Behavioral Science & Policy*, 10(2):46–55, 2024. doi: 10.1177/23794607251320229. URL <https://doi.org/10.1177/23794607251320229>.
- Nikhil Garg, Hannah Li, and Faidra Monachou. Standardized tests and affirmative action: The role of bias and variance. In *Proceedings of the 2021 ACM Conference on Fairness, Accountability, and Transparency*, pages 261–261, 2021.
- Andrew Gelman, Jeffrey Fagan, and Alex Kiss. An analysis of the New York City police department’s “stop-and-frisk” policy in the context of claims of racial bias. *Journal of the American Statistical Association*, 102(479):813–823, 2007.
- Sharad Goel, Justin M Rao, and Ravi Shroff. Personalized risk assessments in the criminal justice system. *The American Economic Review*, 106(5):119–123, 2016a.
- Sharad Goel, Justin M Rao, and Ravi Shroff. Precinct or prejudice? Understanding racial disparities in New York City’s stop-and-frisk policy. *Annals of Applied Statistics*, 10, 2016b.
- Sharad Goel, Maya Perelman, Ravi Shroff, and David Sklansky. Combatting police discrimination in the age of big data. *New Criminal Law Review*, 20, 2017.
- Sharad Goel, Ravi Shroff, Jennifer Skeem, and Christopher Slobogin. The accuracy, equity, and

- jurisprudence of criminal risk assessment. In *Research Handbook on Big Data Law*. Edward Elgar Publishing, 2021.
- D James Greiner and Donald B Rubin. Causal effects of perceived immutable characteristics. *Review of Economics and Statistics*, 93(3):775–785, 2011.
- J. Grossman, S. Tomkins, L. Page, et al. The disparate impacts of college admissions policies on Asian American applicants. *Scientific Reports*, 14:4449, 2024a. doi: 10.1038/s41598-024-55119-0. URL <https://doi.org/10.1038/s41598-024-55119-0>.
- Joshua Grossman, Julian Nyarko, and Sharad Goel. Reconciling legal and empirical conceptions of disparate impact: An analysis of police stops across California. *Journal of Law & Empirical Analysis*, 1(1):118–133, 2024b. doi: 10.1177/2755323X241243168. URL <https://doi.org/10.1177/2755323X241243168>.
- Susan S Grover. The business necessity defense in disparate impact discrimination cases. *Ga. L. Rev.*, 30:387, 1995.
- Moritz Hardt, Eric Price, and Nati Srebro. Equality of opportunity in supervised learning. *Advances in Neural Information Processing Systems*, 29, 2016.
- Melody Huang and Samuel D Pimentel. Variance-based sensitivity analysis for weighting estimators results in more informative bounds. *Biometrika*, 112(1):asae040, 08 2024. ISSN 1464-3510. doi: 10.1093/biomet/asae040. URL <https://doi.org/10.1093/biomet/asae040>.
- Jongbin Jung, Connor Concannon, Ravi Shroff, Sharad Goel, and Daniel G Goldstein. Simple rules to guide expert classifications. *Journal of the Royal Statistical Society: Series A (Statistics in Society)*, 183(3):771–800, 2020.
- Nathan Kallus, Xiaojie Mao, and Angela Zhou. Assessing algorithmic fairness with unobserved protected class using data combination. *Management Science*, 68(3):1959–1981, 2022.
- William Karush. Minima of functions of several variables with inequalities as side constraints. *M. Sc. Dissertation. Dept. of Mathematics, Univ. of Chicago*, 1939.
- Jon Kleinberg, Himabindu Lakkaraju, Jure Leskovec, Jens Ludwig, and Sendhil Mullainathan. Human decisions and machine predictions. *The Quarterly Journal of Economics*, 133(1):237–293, 2018.
- John Knowles, Nicola Persico, and Petra Todd. Racial bias in motor vehicle searches: Theory and evidence. *Journal of Political Economy*, 109(1), 2001.
- Kathleen Kruse and Debra A. Pinals. Rights to reasonable medical treatment and consideration of race-related treatment outcomes. *Journal of the American Academy of Psychiatry and the Law Online*, 45(4):497–499, 2017. ISSN 1093-6793. URL <https://jaapl.org/content/45/4/497>.
- H. W. Kuhn and A. W. Tucker. Nonlinear programming. In *Proceedings of the Second Berkeley Symposium on Mathematical Statistics and Probability, 1950*, pages 481–492. University of

- California Press, Berkeley-Los Angeles, Calif., 1951.
- John MacDonald and Steven Raphael. Effect of scaling back punishment on racial and ethnic disparities in criminal case outcomes. *Criminology & Public Policy*, 19(4):1139–1164, 2020. doi: <https://doi.org/10.1111/1745-9133.12495>. URL <https://onlinelibrary.wiley.com/doi/abs/10.1111/1745-9133.12495>.
- Sandra G Mayson. Bias in, bias out. *Yale Law Journal*, 128:2218, 2018.
- Cory McCartan, Robin Fisher, Jacob Goldin, Daniel E. Ho, and Kosuke Imai. Estimating racial disparities when race is not observed. *Journal of the American Statistical Association*, 0(0):1–14, 2025. doi: 10.1080/01621459.2025.2526695. URL <https://doi.org/10.1080/01621459.2025.2526695>.
- John Monahan and Jennifer L Skeem. Risk assessment in criminal sentencing. *Annual review of clinical psychology*, 12:489–513, 2016.
- Alicia H. Munnell, Geoffrey M. B. Tootell, Lynn E. Browne, and James McEneaney. Mortgage lending in Boston: Interpreting HMDA data. *The American Economic Review*, 86(1):25–53, 1996.
- Hamed Nilforoshan, Johann D Gaebler, Ravi Shroff, and Sharad Goel. Causal conceptions of fairness and their consequences. In *International Conference on Machine Learning*, pages 16848–16887. PMLR, 2022.
- Long Ouyang, Jeff Wu, Xu Jiang, Diogo Almeida, Carroll L. Wainwright, Pamela Mishkin, Chong Zhang, Sandhini Agarwal, Katarina Slama, Alex Ray, John Schulman, Jacob Hilton, Fraser Kelton, Luke Miller, Maddie Simens, Amanda Askell, Peter Welinder, Paul Christiano, Jan Leike, and Ryan Lowe. Training language models to follow instructions with human feedback. Technical report, 2022.
- Edmund S Phelps. The statistical theory of racism and sexism. *The American Economic Review*, 62(4):659–661, 1972.
- Emma Pierson, Sam Corbett-Davies, and Sharad Goel. Fast threshold tests for detecting discrimination. In *The 21st International Conference on Artificial Intelligence and Statistics (AISTATS)*, 2018.
- Emma Pierson, Camelia Simoiu, Jan Overgoor, Sam Corbett-Davies, Daniel Jenson, Amy Shoemaker, Vignesh Ramachandran, Phoebe Barghouty, Cheryl Phillips, Ravi Shroff, et al. A large-scale analysis of racial disparities in police stops across the United States. *Nature Human Behaviour*, 4(7):736–745, 2020.
- Solomon W Polachek. Earnings over the life cycle: The Mincer earnings function and its applications. *Foundations and Trends in Microeconomics*, 4(3):165–272, 2008.
- Manish Raghavan. What should we do when our ideas of fairness conflict? *Communications of the ACM*, 67(1):88–97, 2023.

- M Marit Rehavi and Sonja B Starr. Racial disparity in federal criminal charging and its sentencing consequences. *U of Michigan Law & Econ, Empirical Legal Studies Center Paper*, (12-002), 2012.
- Paul R Rosenbaum. *Observational studies*. Springer, 2002.
- Paul R Rosenbaum and Donald B Rubin. Assessing sensitivity to an unobserved binary covariate in an observational study with binary outcome. *Journal of the Royal Statistical Society. Series B (Methodological)*, 45:212–218, 1983a.
- Paul R Rosenbaum and Donald B Rubin. The central role of the propensity score in observational studies for causal effects. *Biometrika*, 70(1):41–55, 1983b.
- Sartaj Sahni. Computationally related problems. *SIAM Journal on computing*, 3(4):262–279, 1974.
- SFFA v. Harvard. Students for Fair Admissions, Inc., Petitioner, v. President and Fellows of Harvard College. Students for Fair Admissions, Inc., Petitioner, v. University of North Carolina, et al., 2023. https://www.supremecourt.gov/opinions/22pdf/20-1199_16gn.pdf.
- Ravi Shroff. Predictive analytics for city agencies: Lessons from children’s services. *Big Data*, 5(3): 189–196, 2017.
- Camelia Simoiu, Sam Corbett-Davies, and Sharad Goel. The problem of infra-marginality in outcome tests for discrimination. *The Annals of Applied Statistics*, 11(3):1193–1216, 2017.
- Scott Smart and Joel Waldfogel. A citation-based test for discrimination at economics and finance journals. Technical report, National Bureau of Economic Research, 1996.
- João M. Souto-Maior and Ravi Shroff. Differences in academic preparedness do not fully explain black–white enrollment disparities in advanced high school coursework. *Sociological Science*, 11(6):138–163, 2024. ISSN 2330-6696. doi: 10.15195/v11.a6. URL <http://dx.doi.org/10.15195/v11.a6>.
- U.S. Equal Employment Opportunity Commission. Bona fide occupational qualifications. <https://www.eeoc.gov/laws/guidance/cm-625-bona-fide-occupational-qualifications>, 2009.
- Tyler J VanderWeele and Whitney R Robinson. Confounding and mediating variables. *Epidemiology*, 25(4):473–484, 2014.
- Washington v. Davis. 426 U.S. 229, 1976.
- Michael Zanger-Tishler, Julian Nyarko, and Sharad Goel. Risk scores, label bias, and everything but the kitchen sink. *Science Advances*, 10(13):eadi8411, 2024. doi: 10.1126/sciadv.adi8411. URL <https://www.science.org/doi/abs/10.1126/sciadv.adi8411>.
- Yao Zhang and Qingyuan Zhao. Bounds and semiparametric inference in l^∞ - and l^2 -sensitivity analysis for observational studies. *arXiv preprint arXiv:2211.04697*, 2022.
- Qingyuan Zhao, Dylan S Small, and Bhaswar B Bhattacharya. Sensitivity analysis for inverse

probability weighting estimators via the percentile bootstrap. *Journal of the Royal Statistical Society Series B: Statistical Methodology*, 81(4):735–761, 2019.

Chapter 5: Auditing Large Language Models for Race & Gender Disparities: Implications for Artificial Intelligence–Based Hiring

- Josh Achiam, Steven Adler, Sandhini Agarwal, Lama Ahmad, Ilge Akkaya, Florencia Leoni Aleman, Diogo Almeida, Janko Altschmidt, Sam Altman, Shyamal Anadkat, et al. Gpt-4 technical report. *arXiv preprint arXiv:2303.08774*, 2023.
- Jiafu An, Difang Huang, Chen Lin, and Mingzhu Tai. Measuring gender and racial biases in large language models. *arXiv preprint arXiv:2403.15281*, 2024.
- Anthropic. Model Card and Evaluations for Claude Models. <https://www-cdn.anthropic.com/bd2a28d2535bfb0494cc8e2a3bf135d2e7523226/Model-Card-Claude-2.pdf>, July 2023.
- Anthropic. The Claude 3 Model Family: Opus, Sonnet, Haiku. https://www-cdn.anthropic.com/de8ba9b01c9ab7cbabf5c33b80b7bbc618857627/Model_Card_Claude_3.pdf, March 2024a.
- Anthropic. Claude 3.5 Sonnet: Model Card Addendum. https://www-cdn.anthropic.com/fed9cc193a14b84131812372d8d5857f8f304c52/Model_Card_Claude_3_Addendum.pdf, June 2024b.
- Lena Armstrong, Abbey Liu, Stephen MacNeil, and Danaë Metaxa. The silicone ceiling: Auditing gpt’s race and gender biases in hiring. *arXiv preprint arXiv:2405.04412*, 2024.
- Ian Ayres, Mahzarin Banaji, and Christine Jolls. Race effects on eBay. *The RAND Journal of Economics*, 46(4):891–917, 2015.
- Jack Bandy. Problematic machine behavior: A systematic literature review of algorithm audits. *Proceedings of the acm on human-computer interaction*, 5(CSCW1):1–34, 2021.
- Richard Berk, Hoda Heidari, Shahin Jabbari, Michael Kearns, and Aaron Roth. Fairness in criminal justice risk assessments: The state of the art. *Sociological Methods & Research*, 50(1):3–44, 2021.
- Marianne Bertrand and Sendhil Mullainathan. Are Emily and Greg more employable than Lakisha and Jamal? A field experiment on labor market discrimination. *American economic review*, 94(4):991–1013, 2004.
- Tom Brown, Benjamin Mann, Nick Ryder, Melanie Subbiah, Jared D Kaplan, Prafulla Dhariwal, Arvind Neelakantan, Pranav Shyam, Girish Sastry, Amanda Askell, et al. Language models are few-shot learners. *Advances in neural information processing systems*, 33:1877–1901, 2020.
- Silvia Chiappa. Path-specific counterfactual fairness. In *Proceedings of the AAAI Conference on Artificial Intelligence*, volume 33, pages 7801–7808, 2019.
- Alex Chohlas-Wood, Joe Nudell, Keniel Yao, Zhiyuan Lin, Julian Nyarko, and Sharad Goel. Blind justice: Algorithmically masking race in charging decisions. In *Proceedings of the 2021 AAAI/ACM Conference on AI, Ethics, and Society*, pages 35–45, 2021.

- Alexandra Chouldechova. Fair prediction with disparate impact: A study of bias in recidivism prediction instruments. *Big Data*, 5(2):153–163, 2017.
- Alexandra Chouldechova and Aaron Roth. A snapshot of the frontiers of fairness in machine learning. *Communications of the ACM*, 63(5):82–89, 2020.
- Sam Corbett-Davies, Emma Pierson, Avi Feller, Sharad Goel, and Aziz Huq. Algorithmic decision making and the cost of fairness. In *Proceedings of the 23rd ACM SIGKDD International Conference on Knowledge Discovery and Data Mining*, pages 797–806, 2017.
- Sam Corbett-Davies, Johann D Gaebler, Hamed Nilforoshan, Ravi Shroff, and Sharad Goel. The measure and mismeasure of fairness. *The Journal of Machine Learning Research*, 24(1):14730–14846, 2023.
- Amanda Coston, Alan Mishler, Edward H Kennedy, and Alexandra Chouldechova. Counterfactual risk assessments, evaluation, and fairness. In *Proceedings of the 2020 Conference on Fairness, Accountability, and Transparency*, pages 582–593, 2020.
- Cynthia Dwork, Moritz Hardt, Toniann Pitassi, Omer Reingold, and Richard Zemel. Fairness through awareness. In *Proceedings of the 3rd innovations in theoretical computer science conference*, pages 214–226, 2012.
- Tyna Eloundou, Alex Beutel, David G. Robinson, Keren Gu-Lemberg, Anna-Luisa Brakman, Pamela Mishkin, Meghan Shah, Johannes Heidecke, Lilian Weng, and Adam Tauman Kalai. First-person fairness in chatbots. *arXiv preprint arXiv:2410.19803*, 2024.
- European Commission. The Digital Services Act package — Shaping Europe’s digital future. <https://digital-strategy.ec.europa.eu/en/policies/digital-services-act-package>, 2024.
- European Parliament. Regulation of the European Parliament and of the Council Laying Down Harmonised Rules on Artificial Intelligence. <https://eur-lex.europa.eu/legal-content/EN/TXT/?uri=celex%3A52021PC0206>, 2024.
- S Michael Gaddis. Understanding the “how” and “why” aspects of racial-ethnic discrimination: A multimethod approach to audit studies. *Sociology of Race and Ethnicity*, 5(4):443–455, 2019.
- S Michael Gaddis and Raj Ghoshal. Searching for a roommate: A correspondence audit examining racial/ethnic and immigrant discrimination among millennials. *Socius*, 6:2378023120972287, 2020.
- Lara Groves, Jacob Metcalf, Alayna Kennedy, Briana Vecchione, and Andrew Strait. Auditing work: Exploring the New York City algorithmic bias audit regime. *arXiv preprint arXiv:2402.08101*, 2024.
- Amit Haim, Alejandro Salinas, and Julian Nyarko. What’s in a name? Auditing large language models for race and gender bias. *arXiv preprint arXiv:2402.14875*, 2024.
- Moritz Hardt, Eric Price, and Nati Srebro. Equality of opportunity in supervised learning. *Advances*

- in *Neural Information Processing Systems*, 29:3315–3323, 2016.
- Paul W Holland. Statistics and causal inference. *Journal of the American statistical Association*, 81 (396):945–960, 1986.
- Lily Hu and Issa Kohler-Hausmann. What’s sex got to do with fair machine learning? *arXiv preprint arXiv:2006.01770*, 2020.
- Kosuke Imai and Zhichao Jiang. Principal fairness for human and algorithmic decision-making. *arXiv preprint arXiv:2005.10400*, 2020.
- Albert Q Jiang, Alexandre Sablayrolles, Arthur Mensch, Chris Bamford, Devendra Singh Chaplot, Diego de las Casas, Florian Bressand, Gianna Lengyel, Guillaume Lample, Lucile Saulnier, et al. Mistral 7B. *arXiv preprint arXiv:2310.06825*, 2023.
- Albert Q Jiang, Alexandre Sablayrolles, Antoine Roux, Arthur Mensch, Blanche Savary, Chris Bamford, Devendra Singh Chaplot, Diego de las Casas, Emma Bou Hanna, Florian Bressand, et al. Mixtral of experts. *arXiv preprint arXiv:2401.04088*, 2024.
- Jongbin Jung, Sam Corbett-Davies, Johann D Gaebler, Ravi Shroff, and Sharad Goel. Mitigating included-and omitted-variable bias in estimates of disparate impact. *arXiv preprint arXiv:1809.05651*, 2018.
- Sonia K Kang, Katherine A DeCelles, András Tilcsik, and Sora Jun. Whitened résumés: Race and self-presentation in the labor market. *Administrative science quarterly*, 61(3):469–502, 2016.
- Niki Kilbertus, Mateo Rojas-Carulla, Giambattista Parascandolo, Moritz Hardt, Dominik Janzing, and Bernhard Schölkopf. Avoiding discrimination through causal reasoning. In *Proceedings of the 31st International Conference on Neural Information Processing Systems*, pages 656–666, 2017.
- Jon Kleinberg, Sendhil Mullainathan, and Manish Raghavan. Inherent trade-offs in the fair determination of risk scores. In *Proceedings of Innovations in Theoretical Computer Science (ITCS)*, 2017.
- Matt J Kusner, Joshua Loftus, Chris Russell, and Ricardo Silva. Counterfactual fairness. In *Advances in Neural Information Processing Systems*, pages 4066–4076, 2017.
- Louis Lippens. Computer says ‘no’: Exploring systemic bias in ChatGPT using an audit approach. *Computers in Human Behavior: Artificial Humans*, 2(1):100054, 2024.
- Louis Lippens, Siel Vermeiren, and Stijn Baert. the state of hiring discrimination: a meta-analysis of (almost) all recent correspondence experiments. *European Economic Review*, 151:104315, 2023.
- Joshua R Loftus, Chris Russell, Matt J Kusner, and Ricardo Silva. Causal reasoning for algorithmic fairness. *arXiv preprint arXiv:1805.05859*, 2018.
- Sarah Lyons-Padilla, Hazel Rose Markus, Ashby Monk, Sid Radhakrishna, Radhika Shah, Norris A “Daryn” Dodson IV, and Jennifer L Eberhardt. Race influences professional investors’ financial

- judgments. *Proceedings of the National Academy of Sciences*, 116(35):17225–17230, 2019.
- Meta. Introducing Meta Llama 3: The most capable openly available LLM to date. AI at Meta Blog, 2024. <https://ai.meta.com/blog/meta-llama-3/>.
- Danaë Metaxa, Joon Sung Park, Ronald E Robertson, Karrie Karahalios, Christo Wilson, Jeff Hancock, Christian Sandvig, et al. Auditing algorithms: Understanding algorithmic systems from the outside in. *Foundations and Trends® in Human–Computer Interaction*, 14(4):272–344, 2021.
- Razieh Nabi and Ilya Shpitser. Fair inference on outcomes. In *Proceedings of the AAAI Conference on Artificial Intelligence*, 2018.
- Huy Nghiem, John Prindle, Jieyu Zhao, and Hal Daumé, III. “You gotta be a doctor, Lin”: An investigation of name-based bias of large language models in employment recommendations. *arXiv preprint arXiv:2406.12232*, 2024.
- N.Y.C. Local Law No. 144. N.Y.C. Admin. Code, §20-870, 2021.
- OpenAI. Hello GPT-4o. <https://openai.com/index/hello-gpt-4o/>, May 2024a.
- OpenAI. GPT-4o mini: advancing cost-efficient intelligence. <https://openai.com/index/hello-gpt-4o/>, July 2024b.
- Manish Raghavan. What should we do when our ideas of fairness conflict? *Communications of the ACM*, 67(1):88–97, 2023.
- Manish Raghavan, Solon Barocas, Jon Kleinberg, and Karen Levy. Mitigating bias in algorithmic hiring: Evaluating claims and practices. In *Proceedings of the 2020 conference on fairness, accountability, and transparency*, pages 469–481, 2020.
- Christian Sandvig, Kevin Hamilton, Karrie Karahalios, and Cedric Langbort. Auditing algorithms: Research methods for detecting discrimination on internet platforms. *Data and discrimination: converting critical concerns into productive inquiry*, 22(2014):4349–4357, 2014.
- Maya Sen and Omar Wasow. Race as a bundle of sticks: Designs that estimate effects of seemingly immutable characteristics. *Annual Review of Political Science*, 19(1):499–522, 2016.
- Latanya Sweeney. Discrimination in online ad delivery: Google ads, black names and white names, racial discrimination, and click advertising. *Queue*, 11(3):10–29, 2013.
- Prasanna Tambe, Peter Cappelli, and Valery Yakubovich. Artificial intelligence in human resources management: Challenges and a path forward. *California Management Review*, 61(4):15–42, 2019.
- Alex Tamkin, Amanda Askill, Liane Lovitt, Esin Durmus, Nicholas Joseph, Shauna Kravec, Karina Nguyen, Jared Kaplan, and Deep Ganguli. Evaluating and mitigating discrimination in language model decisions. *arXiv preprint arXiv:2312.03689*, 2023.
- Kevin Tobia. Disparate statistics. *The Yale Law Journal*, pages 2382–2420, 2017.

- Akshaj Kumar Veldanda, Fabian Grob, Shailja Thakur, Hammond Pearce, Benjamin Tan, Ramesh Karri, and Siddharth Garg. Are Emily and Greg still more employable than Lakisha and Jamal? investigating algorithmic hiring bias in the era of ChatGPT. *arXiv preprint arXiv:2310.05135*, 2023.
- Yixin Wang, Dhanya Sridhar, and David M Blei. Equal opportunity and affirmative action via counterfactual predictions. *arXiv preprint arXiv:1905.10870*, 2019.
- White House. Executive Order on the Safe, Secure, and Trustworthy Development and Use of Artificial Intelligence. <https://www.whitehouse.gov/briefing-room/presidential-actions/2023/10/30/executive-order-on-the-safe-secure-and-trustworthy-development-and-use-of-artificial-intelligence>, October 2023.
- Ronald E Wienk. *Measuring racial discrimination in American housing markets: The housing market practices survey*, volume 444. Department of Housing and Urban Development, Office of Policy Development, 1979.
- Muhammad Bilal Zafar, Isabel Valera, Manuel Gomez Rodriguez, and Krishna P Gummadi. Fairness beyond disparate treatment & disparate impact: Learning classification without disparate mistreatment. In *Proceedings of the 26th International Conference on World Wide Web*, pages 1171–1180, 2017.

Chapter 6: A Simple, Statistically Robust Test of Discrimination

- Alberto Alesina and Eliana La Ferrara. A test of racial bias in capital sentencing. *American Economic Review*, 104(11):3397–3433, 2014.
- Julia Angwin, Jeff Larson, Surya Mattu, and Lauren Kirchner. Machine bias. In *Ethics of data and analytics*, pages 254–264. Auerbach Publications, 2022.
- Kate Antonovics and Brian G Knight. A new look at racial profiling: Evidence from the Boston police department. *The Review of Economics and Statistics*, 91(1):163–177, 2009.
- Shamena Anwar and Hanming Fang. An alternative test of racial prejudice in motor vehicle searches: Theory and evidence. *American Economic Review*, 96(1):127–151, 2006.
- Sarah F Anzia and Christopher R Berry. The Jackie (and Jill) Robinson effect: Why do congresswomen outperform congressmen? *American Journal of Political Science*, 55(3):478–493, 2011.
- David Arnold, Will Dobbie, and Crystal S Yang. Racial bias in bail decisions. *The Quarterly Journal of Economics*, 133(4):1885–1932, 2018.
- Susan Athey and Jonathan Levin. The value of information in monotone decision problems. *Research in Economics*, 72(1):101–116, 2018.
- Ian Ayres. Outcome tests of racial disparities in police practices. *Justice Research and Policy*, 4(1-2):131–142, 2002.
- Ian Ayres. Testing for discrimination and the problem of “included variable bias”. *Working paper*, 2010. Available at <http://islandia.law.yale.edu/ayers/ayresincludedvariablebias.pdf>.
- Robert Bartlett, Adair Morse, Richard Stanton, and Nancy Wallace. Consumer-lending discrimination in the fintech era. *Journal of Financial Economics*, 143(1):30–56, 2022.
- Gary S Becker. The economics of discrimination. *University of Chicago Press Economics Books*, 1957.
- Gary S Becker. Nobel lecture: The economic way of looking at behavior. *Journal of Political Economy*, 101(3):385–409, 1993.
- James A Berkovec, Glenn B Canner, Stuart A Gabriel, and Timothy H Hannan. Race, redlining, and residential mortgage loan performance. *The Journal of Real Estate Finance and Economics*, 9:263–294, 1994.
- James A Berkovec, Glenn B Canner, Stuart A Gabriel, and Timothy H Hannan. Discrimination, competition, and loan performance in FHA mortgage lending. *Review of Economics and Statistics*, 80(2):241–250, 1998.

- V. I. Bogachev. *Measure Theory. Vol. I, II*. Springer-Verlag, Berlin, 2007. ISBN 978-3-540-34513-8; 3-540-34513-2. doi: 10.1007/978-3-540-34514-5. URL <https://doi.org/10.1007/978-3-540-34514-5>.
- Tim Brennan, William Dieterich, and Beate Ehret. Evaluating the predictive validity of the COMPAS risk and needs assessment system. *Criminal Justice and behavior*, 36(1):21–40, 2009.
- James H Carr and Isaac F Megbolugbe. The Federal Reserve Bank of Boston study on mortgage lending revisited. *Journal of Housing Research*, pages 277–313, 1993.
- Adam Chilton, Jonathan Masur, and Kyle Rozema. Political ideology and the law review selection process. *American Law and Economics Review*, 22(1):211–240, 2020.
- Alex Chohlas-Wood, Marissa Gerchick, Sharad Goel, Aziz Z Huq, Amy Shoemaker, Ravi Shroff, and Keniel Yao. Identifying and measuring excessive and discriminatory policing. *U. Chi. L. Rev.*, 89:441, 2022.
- Billy R Close and Patrick Leon Mason. Searching for efficient enforcement: Officer characteristics and racially biased policing. *Review of Law & Economics*, 3(2):263–321, 2007.
- Sam Corbett-Davies, Emma Pierson, Avi Feller, Sharad Goel, and Aziz Huq. Algorithmic decision making and the cost of fairness. In *Proceedings of the 23rd ACM SIGKDD International Conference on Knowledge Discovery and Data Mining*, pages 797–806, 2017.
- Sam Corbett-Davies, Johann D. Gaebler, Hamed Nilforoshan, Ravi Shroff, and Sharad Goel. The measure and mismeasure of fairness. *Journal of Machine Learning Research*, 24(312):1–117, 2023.
- Decio Coviello and Nicola Persico. An economic analysis of Black-White disparities in the New York Police Department’s stop-and-frisk program. *The Journal of Legal Studies*, 44(2):315–360, 2015.
- W. Dieterich, T. Brennan, and W. L. Oliver. Predictive validity of the COMPAS Core Risk Scales: A probation outcomes study conducted for the Michigan Department of Corrections. Tech. rep., Northpointe Inc., Traverse City, MI, 2011.
- W. Dieterich, C. Mendoza, D. Hubbard, J. Ferro, and T. Brennan. COMPAS risk scales validation study: An outcomes study conducted for the Santa Barbara County Probation Department. Tech. rep., Northpointe, Traverse City, MI, 2017.
- W. Dieterich, C. Mendoza, D. Hubbard, J. Ferro, and T. Brennan. COMPAS risk scales validation study: An outcomes study conducted for the Riverside County Probation Department. Tech. rep., Northpointe, Traverse City, MI, 2018.
- Marc N Elliott, Peter A Morrison, Allen Fremont, Daniel F McCaffrey, Philip Pantoja, and Nicole Lurie. Using the Census Bureau’s surname list to improve estimates of race/ethnicity and associated disparities. *Health Services and Outcomes Research Methodology*, 9:69–83, 2009.
- Robin S Engel. A critique of the “outcome test” in racial profiling research. *Justice Quarterly*, 25(1):1–36, 2008.

- Robin S. Engel and Rob Tillyer. Searching for equilibrium: The tenuous nature of the outcome test. *Justice Quarterly*, 25(1):54–71, 2008.
- Charles R Epp, Steven Maynard-Moody, and Donald Haider-Markel. *Pulled over: How police stops define race and citizenship*. University of Chicago Press, 2014.
- Equivant. Practitioner’s guide to COMPAS Core, April 2019. URL <https://www.equivant.com/wp-content/uploads/Practitioners-Guide-to-COMPAS-Core-040419.pdf>. Accessed: 6 March, 2024.
- David Farabee, Sheldon Zhang, Robert EL Roberts, and Joy Yang. COMPAS validation study, 2010.
- Benjamin Feigenberg and Conrad Miller. Would eliminating racial disparities in motor vehicle searches have efficiency costs? *The Quarterly Journal of Economics*, 137(1):49–113, 2022.
- Anthony W Flores, Kristin Bechtel, and Christopher T Lowenkamp. False positives, false negatives, and false analyses: A rejoinder to “Machine bias: There’s software used across the country to predict future criminals. And it’s biased against blacks”. *Fed. Probation*, 80:38, 2016.
- Floyd v. City of New York. 959 F. Supp. 2d 540, S.D.N.Y, 2013.
- Roland G Fryer Jr. An empirical analysis of racial differences in police use of force. *Journal of Political Economy*, 127(3):1210–1261, 2019.
- Johann Gaebler, William Cai, Guillaume Basse, Ravi Shroff, Sharad Goel, and Jennifer Hill. A causal framework for observational studies of discrimination. *Statistics and Public Policy*, 9(1): 26–48, 2022.
- George C Galster. The facts of lending discrimination cannot be argued away by examining default rates. *Housing Policy Debate*, 4(1):141–146, 1993.
- Andrew Gelman, Jeffrey Fagan, and Alex Kiss. An analysis of the New York City Police Department’s “stop-and-frisk” policy in the context of claims of racial bias. *Journal of the American Statistical Association*, 102(479):813–823, 2007.
- Sharad Goel, Justin M Rao, and Ravi Shroff. Precinct or prejudice? Understanding racial disparities in New York City’s stop-and-frisk policy. *Annals of Applied Statistics*, 10(1):365–394, 2016.
- Sharad Goel, Maya Perelman, Ravi Shroff, and David Alan Sklansky. Combatting police discrimination in the age of big data. *New Criminal Law Review*, 20(2):181–232, 2017.
- Clifton Green, Narasimhan Jegadeesh, and Yue Tang. Gender and job performance: Evidence from Wall Street. *Financial Analysts Journal*, 65(6):65–78, 2009.
- Jeffrey Grogger and Greg Ridgeway. Testing for racial profiling in traffic stops from behind a veil of darkness. *Journal of the American Statistical Association*, 101(475):878–887, 2006.

- Joshua Grossman, Julian Nyarko, and Sharad Goel. Racial bias as a multi-stage, multi-actor problem: An analysis of pretrial detention. *Journal of Empirical Legal Studies*, 20(1):86–133, 2023.
- Joshua Grossman, Julian Nyarko, and Sharad Goel. Reconciling legal and empirical conceptions of disparate impact: An analysis of police stops across California. *Journal of Law and Empirical Analysis*, 1, 2024a.
- Joshua Grossman, Sabina Tomkins, Lindsay C Page, and Sharad Goel. The disparate impacts of college admissions policies on Asian American applicants. *Scientific Reports*, 14, 2024b.
- Sanford J Grossman and Oliver D Hart. An analysis of the principal-agent problem. In *Foundations of Insurance Economics: Readings in Economics and Finance*, pages 302–340. Springer, 1992.
- Andrea Guerrero, William Armaline, Chad Bianco, DJ Criner, John Dobard, Sean Duryee, Lawanda Hawkins, Brian Eric Kennedy, Lily Khadjavi, Manjusha P. Kulkarni, Melanie Ochoa, Rich Randolph, Angela Sierra, Tamani Taylor, Sean Thuilliez, Cha Vang, and Ronaldo Villeda. Ripa board report 2024. Technical report, California Department of Justice, Racial Identity and Profiling Act Advisory Board, 2024. URL <https://oag.ca.gov/system/files/media/ripa-board-report-2024.pdf>. Accessed: 2024-02-27.
- Aziz Z Huq. Judging discriminatory intent. *Cornell Law Review*, 103, 2017.
- Jongbin Jung, Sam Corbett-Davies, Johann Gaebler, Ravi Shroff, and Sharad Goel. Mitigating included- and omitted-variable bias in estimates of disparate impact. *arXiv preprint arXiv:1809.05651*, 2024.
- Samuel Karlin and Herman Rubin. The theory of decision procedures for distributions with monotone likelihood ratio. *The Annals of Mathematical Statistics*, pages 272–299, 1956.
- Julian Keilson and Ushio Sumita. Uniform stochastic ordering and related inequalities. *Canadian Journal of Statistics*, 10(3):181–198, 1982.
- John Knowles, Nicola Persico, and Petra Todd. Racial bias in motor vehicle searches: Theory and evidence. *Journal of Political Economy*, 109(1):203–229, 2001.
- Sharon Lansing. New York State COMPAS-probation risk and need assessment study: Examining the recidivism scale’s effectiveness and predictive accuracy. *Retrieved March*, 1:2013, 2012.
- John MacDonald and Steven Raphael. Effect of scaling back punishment on racial and ethnic disparities in criminal case outcomes. *Criminology & Public Policy*, 19(4):1139–1164, 2020.
- Richard D McKelvey and Thomas R Palfrey. Quantal response equilibria for normal form games. *Games and Economic Behavior*, 10(1):6–38, 1995.
- Paul R Milgrom. Good news and bad news: Representation theorems and applications. *The Bell Journal of Economics*, pages 380–391, 1981.

- Roland Neil and Christopher Winship. Methodological challenges and opportunities in testing for racial discrimination in policing. *Annual Review of Criminology*, 2:73–98, 2019.
- Nicola Persico. Information acquisition in auctions. *Econometrica*, 68(1):135–148, 2000.
- Nicola Persico and Petra Todd. Generalising the hit rates test for racial bias in law enforcement, with an application to vehicle searches in Wichita. *The Economic Journal*, 116(515):F351–F367, 2006.
- Emma Pierson, Sam Corbett-Davies, and Sharad Goel. Fast threshold tests for detecting discrimination. In *Proceedings of the Twenty-First International Conference on Artificial Intelligence and Statistics*, 2018.
- Emma Pierson, Camelia Simoiu, Jan Overgoor, Sam Corbett-Davies, Daniel Jenson, Amy Shoemaker, Vignesh Ramachandran, Phoebe Barghouty, Cheryl Phillips, Ravi Shroff, and Sharad Goel. A large-scale analysis of racial disparities in police stops across the United States. *Nature Human Behaviour*, 4(7):736–745, 2020.
- Devin G Pope and Justin R Sydnor. What’s in a picture? Evidence of discrimination from Prosper.com. *Journal of Human Resources*, 46(1):53–92, 2011.
- Warren A Reich, Sarah Picard-Fritsche, Virginia Barber Rioja, and Merrill Rotter. Evidence-based risk assessment in a mental health court. *Center for Court Innovation*, 2016.
- Greg Ridgeway. Analysis of racial disparities in the New York Police Department’s stop, question, and frisk practices. *RAND Safety and Justice*, 2007.
- Jutta Roosen and David A Hennessy. Testing for the monotone likelihood ratio assumption. *Journal of Business & Economic Statistics*, 22(3):358–366, 2004.
- Paul R Rosenbaum and Donald B Rubin. The central role of the propensity score in observational studies for causal effects. *Biometrika*, 70(1):41–55, 1983.
- Ludger Rüschendorf. On conditional stochastic ordering of distributions. *Advances in Applied Probability*, 23(1):46–63, 1991.
- Moshe Shaked and J George Shanthikumar. *Stochastic Orders*. Springer, 2007.
- J George Shanthikumar and David D Yao. The preservation of likelihood ratio ordering under convolution. *Stochastic Processes and their Applications*, 23(2):259–267, 1986.
- Camelia Simoiu, Sam Corbett-Davies, and Sharad Goel. The problem of infra-marginality in outcome tests for discrimination. *Annals of Applied Statistics*, pages 1193–1216, 2017.
- Scott B Smart and Joel Waldfogel. A citation-based test for discrimination at economics and finance journals, 1996.
- João M Souto-Maior and Ravi Shroff. Differences in academic preparedness do not fully explain

- Black–White enrollment disparities in advanced high school coursework. *Sociological Science*, 11: 138–163, 2024.
- Sonja B Starr and M Marit Rehavi. Mandatory sentencing and racial disparity: Assessing the role of prosecutors and the effects of *Booker*. *Yale LJ*, 123:2, 2013.
- Terry v. Ohio. 392 U.S. 1, 1968.
- The People v. William G. 40 Cal.3d 550, 1985.
- Ward Whitt. Uniform conditional stochastic order. *Journal of Applied Probability*, 17(1):112–123, 1980.
- Linda F Wightman. LSAC National Longitudinal Bar Passage Study. *LSAC Research Report Series*, 1998.

Chapter 7: A Causal Framework for Observational Studies of Discrimination

Ian Ayres. Outcome tests of racial disparities in police practices. *Justice Research and Policy*, 4 (1-2):131–142, 2002.

Ana I Balsa, Thomas G McGuire, and Lisa S Meredith. Testing for statistical discrimination in health care. *Health Services Research*, 40(1):227–252, 2005.

Sandy Baum and Eban Goodstein. Gender imbalance in college applications: Does it lead to a preference for men in the admissions process? *Economics of Education Review*, 24(6):665–675, 2005.

Nathan Berg and Donald Lien. Measuring the effect of sexual orientation on income: Evidence of discrimination? *Contemporary Economic Policy*, 20(4):394–414, 2002.

Marianne Bertrand and Sendhil Mullainathan. Are Emily and Greg more employable than Lakisha and Jamal? A field experiment on labor market discrimination. *American Economic Review*, 94 (4):991–1013, 2004.

Patrick Billingsley. *Probability and Measure*. John Wiley & Sons, 2008.

Nicole Bohme Carnegie, Masataka Harada, and Jennifer L Hill. Assessing sensitivity to unmeasured confounding using a simulated potential confounder. *Journal of Research on Educational Effectiveness*, 9(3):395–420, 2016.

Alex Chohlas-Wood, Joe Nudell, Keniel Yao, Zhiyuan Lin, Julian Nyarko, and Sharad Goel. Blind justice: Algorithmically masking race in charging decisions. In *Proceedings of the 2021 AAAI/ACM Conference on AI, Ethics, and Society*, pages 35–45, 2021.

Carlos Cinelli and Chad Hazlett. Making sense of sensitivity: Extending omitted variable bias. *J. R. Statist. Soc. B*, 2020.

Jerome Cornfield, William Haenszel, E Cuyler Hammond, Abraham M Lilienfeld, Michael B Shimkin, and Ernst L Wynder. Smoking and lung cancer: Recent evidence and a discussion of some questions. *Journal of the National Cancer Institute*, 22(1):173–203, 1959.

Vincent Dorie, Masataka Harada, Nicole Bohme Carnegie, and Jennifer Hill. A flexible, interpretable framework for assessing sensitivity to unmeasured confounding. *Statistics in Medicine*, 35(20): 3453–3470, 2016.

Benjamin G Edelman and Michael Luca. Digital discrimination: The case of airbnb.com. *Harvard Business School NOM Unit Working Paper*, 14(054), 2014.

AlexanderM Franks, Alexander D’Amour, and Avi Feller. Flexible sensitivity analysis for observational studies without observable implications. *Journal of the American Statistical Association*,

- pages 1–33, 2019.
- Jerome Friedman, Trevor Hastie, and Robert Tibshirani. *The Elements of Statistical Learning*. Springer Series in Statistics, New York, 2001.
- Roland G Fryer Jr. An empirical analysis of racial differences in police use of force. *Journal of Political Economy*, 127(3):1210–1261, 2019.
- Andrew Gelman, Jeffrey Fagan, and Alex Kiss. An analysis of the New York City Police Department’s “stop-and-frisk” policy in the context of claims of racial bias. *Journal of the American Statistical Association*, 102(479):813–823, 2007.
- Andrew Gelman, John B Carlin, Hal S Stern, David B Dunson, Aki Vehtari, and Donald B Rubin. *Bayesian Data Analysis*. CRC press, 2013.
- Sharad Goel, Maya Perelman, Ravi Shroff, and David Alan Sklansky. Combatting police discrimination in the age of big data. *New Criminal Law Review: An International and Interdisciplinary Journal*, 20(2):181–232, 2017.
- Claudia Goldin and Cecilia Rouse. Orchestrating impartiality: The impact of “blind” auditions on female musicians. *American Economic Review*, 90(4):715–741, 2000.
- Deena Greenberg, Carl Gershenson, and Matthew Desmond. Discrimination in evictions: Empirical evidence and legal challenges. *Harv. CR-CLL Rev.*, 51:115, 2016.
- D James Greiner and Donald B Rubin. Causal effects of perceived immutable characteristics. *Review of Economics and Statistics*, 93(3):775–785, 2011.
- Jeffrey Grogger and Greg Ridgeway. Testing for racial profiling in traffic stops from behind a veil of darkness. *Journal of the American Statistical Association*, 101(475):878–887, 2006.
- James J Heckman. Sample selection bias as a specification error. *Econometrica: Journal of the econometric society*, pages 153–161, 1979.
- James J Heckman and Steven N Durlauf. Comment on “An empirical analysis of racial differences in police use of force” by Roland G. Fryer Jr. *Journal of Political Economy*, 2020.
- Jennifer L Hill. Bayesian nonparametric modeling for causal inference. *Journal of Computational and Graphical Statistics*, 20(1):217–240, 2011.
- Paul W Holland. Statistics and causal inference. *Journal of the American Statistical Association*, 81(396):945–960, 1986.
- Kosuke Imai, Luke Keele, and Dustin Tingley. A general approach to causal mediation analysis. *Psychol. Methods*, 15(4):309–334, 2010a.
- Kosuke Imai, Luke Keele, and Teppei Yamamoto. Identification, inference and sensitivity analysis for causal mediation effects. *Statist. Sci.*, 25(1):51–71, 2010b.

- Guido W Imbens. Sensitivity to exogeneity assumptions in program evaluation. *The American Economic Review*, 93(2):126–132, 2003.
- Guido W Imbens and Donald B Rubin. *Causal Inference in Statistics, Social, and Biomedical Sciences*. Cambridge University Press, 2015.
- Jongbin Jung, Sam Corbett-Davies, Ravi Shroff, and Sharad Goel. Omitted and included variable bias in tests for disparate impact. *arXiv preprint arXiv:1809.05651*, 2018.
- Jongbin Jung, Ravi Shroff, Avi Feller, and Sharad Goel. Bayesian sensitivity analysis for offline policy evaluation. In *Proceedings of the AAAI/ACM Conference on AI, Ethics, and Society*, pages 64–70, 2020.
- Dean Knox, Will Lowe, and Jonathan Mummolo. Administrative records mask racially biased policing. *American Political Science Review*, 2020.
- John MacDonald and Steven Raphael. Effect of scaling back punishment on racial and ethnic disparities in criminal case outcomes. *Criminology & Public Policy*, 2021.
- Lawrence C McCandless and Paul Gustafson. A comparison of Bayesian and Monte Carlo sensitivity analysis for unmeasured confounding. *Statistics in Medicine*, 2017.
- Lawrence C McCandless, Paul Gustafson, and Adrian Levy. Bayesian sensitivity analysis for unmeasured confounding in observational studies. *Statistics in Medicine*, 26(11):2331–2347, 2007.
- Alicia H. Munnell, Geoffrey M. B. Tootell, Lynn E. Browne, and James McEneaney. Mortgage lending in Boston: Interpreting HMDA data. *The American Economic Review*, 86(1):25–53, 1996.
- Judea Pearl. *Causality*. Cambridge university press, 2009.
- Judea Pearl. Conditioning on post-treatment variables. *Journal of Causal Inference*, 3(1):131–137, 2015.
- Judea Pearl, Madelyn Glymour, and Nicholas P Jewell. *Causal Inference in Statistics: A Primer*. John Wiley & Sons, 2016.
- Emma Pierson, Camelia Simoiu, Jan Overgoor, Sam Corbett-Davies, Daniel Jenson, Amy Shoemaker, Vignesh Ramachandran, Phoebe Barghouty, Ravi Shroff, Cheryl Phillips, and Sharad Goel. A large-scale analysis of racial disparities in police stops across the United States. *Nature Human Behaviour*, 4(5), 2020.
- M Marit Rehavi and Sonja B Starr. Racial disparity in federal criminal sentences. *Journal of Political Economy*, 122(6):1320–1354, 2014.
- Thomas Richardson and James Robins. Single world intervention graphs (SWIGs): A unification of the counterfactual and graphical approaches to causality. Technical report, Technical Report 128, Center for Statistics and the Social Sciences, Univ. Washington, Seattle, WA., 2013.

- Christopher Robertson, Shima Baradaran Baughman, and Megan S. Wright. Race and Class: A Randomized Experiment with Prosecutors. *Journal of Empirical Legal Studies*, 16(4):807–847, 2019.
- James Robins. A new approach to causal inference in mortality studies with a sustained exposure period— application to control of the healthy worker survivor effect. *Mathematical Modelling*, 7(9):1393 – 1512, 1986.
- James M Robins, Thomas S Richardson, and Ilya Shpitser. An interventionist approach to mediation analysis. *arXiv preprint arXiv:2008.06019*, 2020.
- P. R Rosenbaum and D. B. Rubin. Assessing sensitivity to an unobserved binary covariate in an observational study with binary outcome. *Journal of the Royal Statistical Society*, 45(2):212–218, 1983a.
- Paul R Rosenbaum and Donald B Rubin. The central role of the propensity score in observational studies for causal effects. *Biometrika*, 70(1):41–55, 1983b.
- Jasjeet Sekhon. The Neyman–Rubin model of causal inference and estimation via matching methods. In *The Oxford Handbook of Political Methodology*. Oxford University Press, 2008.
- Katherine B Spencer, Amanda K Charbonneau, and Jack Glaser. Implicit bias and policing. *Social and Personality Psychology Compass*, 10(1):50–63, 2016.
- Alexandre B Tsybakov. *Introduction to nonparametric estimation*. Springer Science & Business Media, 2008.
- David Word and Colby Perkins. *Building a Spanish Surname List for the 1990’s: A New Approach to an Old Problem*. Population Division, US Bureau of the Census Washington, DC, 1996.
- David Word, Charles Coleman, Robert Nunziata, and Robert Kominski. Demographic aspects of surnames from Census 2000. Technical report, U.S. Census Bureau Population Division, 2008. URL <http://www2.census.gov/topics/genealogy/2000surnames/surnames.pdf>.
- Qingyuan Zhao, Luke J Keele, Dylan S Small, and Marshall M Joffe. A note on posttreatment selection in studying racial discrimination in policing. *American Political Science Review*, page 1–14, 2021.

Chapter 8: Measuring Racial and Ethnic Disparities in Traffic Enforcement with Large-Scale Telematics Data

Revisions to the Standards for the Classification of Federal Data on Race and Ethnicity, October 1997.

Shamena Anwar and Hanming Fang. An alternative test of racial prejudice in motor vehicle searches: Theory and evidence. *American Economic Review*, 96(1):127–151, 2006.

Deborah N Archer. Transportation policy and the underdevelopment of Black communities. *Iowa Law Review*, 106, 2021.

Ian Ayres. Three tests for measuring unjustified disparate impacts in organ transplantation: The problem of “included variable” bias. *Perspectives in Biology and Medicine*, 48(1):68–S87, 2005.

Hunter M Boehme, Deanna Cann, and Deena A Isom. Citizens’ perceptions of over-and under-policing: A look at race, ethnicity, and community characteristics. *Crime & Delinquency*, 68(1): 123–154, 2022.

Sarah Brayne. Surveillance and system avoidance: Criminal justice contact and institutional attachment. *American Sociological Review*, 79(3):367–391, 2014.

Camille Zubrinsky Charles. The dynamics of racial residential segregation. *Annual review of sociology*, 29(1):167–207, 2003.

Department of Transportation, Federal Highway Administration. Functional highway classifications. Technical report, 2017. URL https://www.fhwa.dot.gov/planning%20/processes/statewide/related/highway_functional_classifications/section03.cfm. Last accessed Feb. 05, 2021.

Charles R Epp, Steven Maynard-Moody, and Donald P Haider-Markel. *Pulled over: How police stops define race and citizenship*. University of Chicago Press, 2014.

Johann Gaebler, William Cai, Guillaume Basse, Ravi Shroff, Sharad Goel, and Jennifer Hill. A causal framework for observational studies of discrimination. *Statistics and Public Policy*, 9:26–48, 2022.

Andrew Gelman, Jeffrey Fagan, and Alex Kiss. An analysis of the New York City police department’s “stop-and-frisk” policy in the context of claims of racial bias. *Journal of the American Statistical Association*, 102(479):813–823, 2007.

Andrew Gelman, Yu-Sung Su, Masanao Yajima, Jennifer Hill, Maria Grazia Pittau, Jouni Kerman, Tian Zheng, and Vicent Dorie. *Package ‘arm’*. 2015.

Sharad Goel, Maya Perelman, Ravi Shroff, and David Alan Sklansky. Combatting police discrimination in the age of big data. *New Criminal Law Review*, 20(2):181–232, 2017.

- Jeffrey Grogger and Greg Ridgeway. Testing for racial profiling in traffic stops from behind a veil of darkness. *Journal of the American Statistical Association*, 101(475):878–887, 2006.
- Bernard E Harcourt and Tracey L Meares. Randomization and the Fourth Amendment. *U. Chi. L. Rev.*, 78:809, 2011.
- Alexes Harris. *A pound of flesh: Monetary sanctions as punishment for the poor*. Russell Sage Foundation, 2016.
- Aziz Z Huq, Sharad Goel, Alex Cholas-Wood, Ravi Shroff, Marissa Gerchick, Amy Shoemaker, and Keniel Yao. Identifying and measuring excessive and discriminatory policing. *University of Chicago Law Review*, 89(441), 2022.
- Jongbin Jung, Sam Corbett-Davies, Ravi Shroff, and Sharad Goel. Omitted and included variable bias in tests for disparate impact. Technical report, 08 2018. ArXiv preprint:1809.05651.
- David Karas. Highway to inequity: The disparate impact of the interstate highway system on poor and minority communities in American cities. *New Visions for Public Affairs*, 7(April):9–21, 2015.
- Issa Kohler-Hausmann. *Misdemeanorland*. Princeton University Press, 2018.
- James E Lange, Mark B Johnson, and Robert B Voas. Testing the racial profiling hypothesis for seemingly disparate traffic stops on the New Jersey Turnpike. *Justice Quarterly*, 22(2):193–223, 2005.
- Lynn Langton and Matthew R Durose. *Police behavior during traffic and street stops, 2011*. US Department of Justice, Office of Justice Programs, Bureau of Justice, 2013.
- Kirk Miller. Police stops, pretext, and racial profiling: Explaining warning and ticket stops using citizen self-reports. *Journal of Ethnicity in Criminal Justice*, 6(2):123–149, 2008.
- Raymond A Mohl. Citizen activism and freeway revolts in Memphis and Nashville: The road to litigation. *Journal of Urban History*, 40(5):870–893, 2014.
- Brendan O’Flaherty and Rajiv Sethi. *Shadows of doubt: Stereotypes, crime, and the pursuit of justice*. Harvard University Press, 2019.
- Aur lie Ouss and John Rappaport. Is police behavior getting worse? Data selection and the measurement of policing harms. *Journal of Legal Studies*, 49(1):153–198, 2020.
- Devah Pager, Rebecca Goldstein, Helen Ho, and Bruce Western. Criminalizing poverty: The consequences of court fees in a randomized experiment. *American Sociological Review*, 87, 2022.
- Emma Pierson, Camelia Simoiu, Jan Overgoor, Sam Corbett-Davies, Daniel Jenson, Amy Shoemaker, Vignesh Ramachandran, Phoebe Barghouty, Cheryl Phillips, Ravi Shroff, et al. A large-scale analysis of racial disparities in police stops across the United States. *Nature Human Behaviour*, 4(7):736–745, 2020.

- Steven Raphael, Michael A Stoll, Kenneth A Small, and Clifford Winston. *Can boosting minority car-ownership rates narrow inter-racial employment gaps?* Brookings Institution Press, 2001.
- Greg Ridgeway. Assessing the effect of race bias in post-traffic stop outcomes using propensity scores. *Journal of Quantitative Criminology*, 22(1):1–29, 2006.
- Greg Ridgeway and John M MacDonald. Doubly robust internal benchmarking and false discovery rates for detecting racial bias in police stops. *Journal of the American Statistical Association*, 104(486):661–668, 2009.
- Jeff Rojek, Richard Rosenfeld, and Scott Decker. The influence of driver’s race on traffic stops in Missouri. *Police Quarterly*, 7(1):126–147, 2004.
- Matt E Ryan. Frisky business: Race, gender and police activity during traffic stops. *European Journal of Law and Economics*, 41(1):65–83, 2016.
- Camelia Simoiu, Sam Corbett-Davies, and Sharad Goel. The problem of infra-marginality in outcome tests for discrimination. *Annals of Applied Statistics*, 11(3):1193–1216, 2017.
- Michael R Smith and Matthew Petrocelli. Racial profiling? A multivariate analysis of police traffic stop data. *Police Quarterly*, 4(1):4–27, 2001.
- U.S. Census Bureau. American Community Survey, 2013–2018.
- Rob Voigt, Nicholas P Camp, Vinodkumar Prabhakaran, William L Hamilton, Rebecca C Hetey, Camilla M Griffiths, David Jurgens, Dan Jurafsky, and Jennifer L Eberhardt. Language from police body camera footage shows racial disparities in officer respect. *Proceedings of the National Academy of Sciences*, 114(25):6521–6526, 2017.